

Counterexamples to List and Line Decodability Bounds over Prime Fields

Justin Thaler

Working draft — September 18, 2026

Abstract

We give counterexamples to proposed list-decoding and proximity-gap bounds for Reed–Solomon codes over prime fields. The central line-level phenomenon is that, even strictly below the characteristic-based Elias radius, two far words can have every other affine combination close to the code. The nearby polynomial can change with the mixing coefficient; a nearby mixture need not reflect common agreement of the inputs.

Write $\theta = 1 - \rho - \eta$ for the tested radius at rate ρ . At every fixed limiting rate, we construct polynomial-size prime alphabets $p = \Theta((n/\log n)^\beta)$, for any fixed $\beta > 12/5$, with $\eta = \Theta(1/\log n)$. Both endpoints have distance $\theta + (1/2 + o(1))\eta$, while all $p - 2$ interior affine mixtures have distance exactly θ . The domains are selected unions of multiplicative cosets. A second family works over every sufficiently large odd prime, with $n = \Theta(\log p)$ and endpoint separation exactly $2\eta/3$. An explicit length-48, dimension-24 example over $\mathbb{F}_{65537}$ has $\eta = 1/16$ and endpoint separation $1/24$; complete coverage is independently verified with integer arithmetic. These statements sharpen the line profile at the already known $1/\log p$ capacity-margin scale, rather than establishing a new gap order.

A separate paired-domain construction has an almost maximally far point, all other affine parameters nearby, and pairwise disjoint nearby lists. It extends to batches consisting entirely of far inputs. Other constructions violate numerical line bounds with singleton lists and efficient witness recovery, so decoding ambiguity is not required for these failures.

We also disprove uniform list-size prescriptions. At rate $1/4$ and fixed gap $1/8$, full-length prime-field codes have lists of size $n/2$. Integer-interval domains give $\log L = \Omega_\rho(\eta^{-2}/\log(1/\eta))$ strictly below Elias, ruling out every uniform $\exp(O_\rho(1/\eta))$ bound. At vanishing rate, separate examples have superlinear numbers of exceptional challenges with singleton lists in the regime of the Dao–Kominers–Thaler bounds; these do not establish general tightness.

We distinguish fixed-gap from shrinking-gap claims, prime from extension-field challenges, and chosen from prescribed domains when applying the results to the conjectures in the S-two whitepaper. The results do not by themselves give a concrete-security bound for an entire deployed protocol.

1 Introduction

Many hash-based SNARKs ask the prover to commit to evaluations of low-degree polynomials. The verifier then checks a small number of positions and random linear combinations of the committed words. Analyzing these checks requires quantitative information about the underlying code. How many low-degree polynomials can be close to one word? How many points on an affine line can be close to the code? When do the nearby polynomials imply a common agreement set for the original words?

These are coding-theoretic questions, but the numerical answers enter soundness calculations. A bound that is polynomial in the block length may be adequate for an asymptotic theorem and




Figure 1: Two far inputs, but all $p - 2$ other affine combinations are nearby (Corollary 4.16). The endpoint separation r/n approaches half the capacity gap $\eta = (2r + 1)/n$. The diagram suppresses most field parameters.

inadequate for a claimed concrete security level. Likewise, a bound proved for random evaluation domains does not automatically apply to a prescribed subgroup. The field size, domain, and order of quantifiers must be explicit.

This paper gives counterexamples to quantitative list and line bounds over prime fields. The main results are independent of any particular proof system. We first prove the general statements, then give finite circle-code examples and apply them to the two conjectures underlying the S-two security analysis.

A useful starting point is the root bound. A nonzero polynomial of degree less than k has fewer than k roots, so two distinct Reed–Solomon codewords cannot agree in k positions. This proves a minimum-distance guarantee. It does not answer the first question above: many polynomials may each agree with an arbitrary received word in more than k positions, using different agreement sets.

Write n for the number of evaluation points and $\rho = k/n$ for the rate. We consider agreement on at least $(\rho + \eta)n$ positions, equivalently relative distance at most $1 - \rho - \eta$. The gap $\eta > 0$ is the excess agreement over the rate. Throughout the paper, we distinguish the largest list at one received word, the number of nearby parameters on an affine line, and correlated agreement across that line. None of these quantities can be substituted for another without a theorem relating them.

The main quantitative message is that the proposed dependence on η fails even when $n = o(p)$. At any fixed rational rate, the proposed list bound has $\log L = O_\rho(1/\eta)$, whereas our short-domain examples require

$$\log L = \Omega_\rho\left(\frac{1}{\eta^2 \log(1/\eta)}\right).$$

Thus the exponent must grow almost quadratically in the reciprocal gap, rather than linearly. The examples remain strictly below the characteristic-based Elias radius.

The clearest line-level message is that *even below Elias, two far words can have every affine combination other than the two endpoints close to the code*. Corollary 4.16 places both endpoints almost half a gap outside the tested radius, while a random mixture is nearby with probability $1 - 2/p$. More generally, for any fixed batch size t , all t inputs can lie almost $(1 - 1/t)\eta$ outside the radius while a random affine combination is nearby with probability $1 - O(t/p)$. The same obstruction holds for unrestricted linear combinations and for multilinear mixtures (Corollary 4.17).

The underlying construction gives an even sharper exceptional point: *a line can be close to the code everywhere except at one point, and that point can be only one coordinate short of being maximally far*. At every fixed rational rate, Corollary 4.15 gives prime-field codes for which all $p - 1$ nonzero parameters are nearby, but parameter zero has distance $1 - \rho - 1/n$. Every word is within $1 - \rho$ of a Reed–Solomon code, so this is almost the largest distance possible. Its separation from the tested radius is $(1 - o(1))\eta$, almost the entire capacity gap. The radius remains strictly below Elias. Every nonzero point has exactly the tested distance, and no codeword is nearby at two distinct nonzero parameters. The line changes only $2r$ coordinates to improve the distance by exactly $2r/n$, attaining the Hamming-distance bound. No pair of codewords explains the line on the required common coordinate set.

The obstruction is that the nearby polynomial changes with the mixing coefficient. In these complete-coverage examples, no polynomial is close at two distinct nearby parameters. A root bound controls agreements between fixed polynomials; it does not force the polynomials selected at different line parameters to fit together. The construction realizes this behavior through products of locator values on the evaluation domain, without any proper subfield.

The $c_1 = c_2 = 1$ numerical prescription predicts a nearby fraction tending to zero; the actual fraction is $1 - 1/p$. Thus testing a random linear combination can find a nearby word with probability arbitrarily close to one despite an almost maximally far input. Theorem 4.41 separately defeats any fixed constants c_1, c_2 , with a fixed positive separation fraction depending on c_2 .

Decoding ambiguity is unnecessary for numerical failure. Separately, Theorem 4.24 gives explicit powers-of-two domains where every nearby point has a unique codeword, recoverable in polynomial time in $\log p$, and the number of nearby parameters exceeds the numerical prescription exponentially. The exceptional point is again one coordinate short of maximum distance. At half rate these examples can have only two distance levels: every parameter outside the nearby bank is at the far level, including both displayed endpoints. Thus the obstruction persists even when proximity along the line is easy to decide and every accepted parameter has an unambiguous witness. Corollary 4.28 strengthens this conclusion: for every fixed $\varepsilon > 0$, suitable infinitely many prime fields admit unique-witness examples whose far point lies at least $(1 - \varepsilon)\eta$ outside the nearby radius. These examples use scaled prime-order orbits and retain efficient witness recovery. They have a sparse nearby set; they are separate from the complete-coverage construction. Theorem 4.30 obtains the same relative separation and uniqueness on domains of length $\Theta(\log p)$ by random sampling, with a polynomial-in- p violation of the numerical bound, but without a witness-recovery guarantee. Theorem 4.32 goes further: two padding orbits preserve uniqueness and almost full-gap separation even when the proposed exponent constant is doubled to $c_2 = 2$. It provides efficient witness recovery over sufficiently large splitting primes. Corollary 4.34 obtains the same numerical strengthening at length $\Theta(\log p)$ without that recovery guarantee; its finite half-rate certificates include separation $38\eta/39$ and a violation factor greater than 2^9 over $\mathbb{F}_{2^{44497-1}}$. A dimension tradeoff preserves uniqueness for every fixed $c_2 < 3$, with separation greater than two thirds of the gap (Corollary 4.35).

In the full-coverage examples the gap shrinks as $\Theta(1/\log p)$, and the strongest relative separation approaches the Elias boundary. These results do not establish superlinear growth in n at one fixed positive gap on short prime-field domains; that stronger question remains open. Separate interval-padding constructions give explicit density/separation tradeoffs and can make the number of missing coordinates grow.

In the fixed-gap line examples, every nearby word can have a unique nearby codeword. Thus their many nearby parameters do not arise from decoding ambiguity at individual points of the line. This uniqueness is a statement about the displayed line, not about all received words for the code.

A separate construction keeps the code’s global maximum list size linear in n , yet has exponentially many uniquely nearby points on a line; every other point on that line is as far from the code as any

word can be (Theorem 5.1).

1.1 The list-size obstruction

Our first result concerns a particularly simple evaluation domain: $0, 1, \dots, n-1$ in a prime field.

Theorem 1.1 (Exponential lists on integer intervals, informal). *Fix any rational rate $0 < \rho < 1$. For arbitrarily large n , there is a prime p , a rate- ρ Reed–Solomon code over $\mathbb{F}_p$, and a received word with $2^{\Omega_\rho(n)}$ nearby codewords, where*

$$\eta = \Theta_\rho\left(\frac{1}{\sqrt{n \log n}}\right), \quad \log_2 p = \Theta_\rho(\sqrt{n \log n}).$$

Every listed codeword agrees with the received word on exactly $(\rho + \eta)n$ positions. The associated distance is strictly below the characteristic-based Elias radius.

The field is large, but its logarithm is sublinear in the block length. The list is much larger than any bound of the form $c_1 2^{c_2 H_2(\rho)/\eta}$ with constants depending only on the rate: the logarithm of that proposed bound is $O_\rho(\sqrt{n \log n})$, whereas the logarithm of our list size is $\Omega_\rho(n)$. Here H_2 is binary entropy. For a fixed rate it is simply a positive constant.

The qualification about the field and domain matters. This is a statement about what can be guaranteed for *every* prime-field Reed–Solomon code. It does not give exponential lists at a fixed positive gap, or when the prime is required to be polynomial in n .

Finite and uniform refinements. The moment count also gives explicit certificates. At length 64 and dimension 32, the exact enumeration in Section 7.1 proves a list of 5,552,914,238,035 codewords with 34 agreements. Over the M31 prime, a further interval example at length 157 and dimension 63 has at least 28,169,451,256,663,519,418 codewords with 68 agreements, exceeding the finite prescription with $c_1 = c_2 = 1$ by more than $2^{34.09967}$ (Section 7.2). Proposition 2.13 gives an analytic refinement: for each fixed number m of canceled coefficients and fixed limiting rate, concentration improves the full-range counting guarantee by a factor of order $n^{m/2}$. Classical Gram orthogonalization gives explicit constants; Proposition 2.14 strengthens the finite bound, and Section 2.7 obtains the Gaussian central-density constant when m is fixed. These refinements retain the exponent above. The exact checks certify the finite arithmetic; the proofs establish the asymptotic statements.

The same construction has a stronger formulation in terms of the gap. Corollary 2.3 rules out every uniform bound whose logarithm is $o_\rho(\eta^{-2}/\log(1/\eta))$. Moreover, over every sufficiently large prime field, with $b = \log_2 p$, it gives a below-Elias list with

$$\log_2 L \geq (1/2 - o(1))b^2/\log_2 b.$$

Here $n = (H_2(\rho)^{-1} + o(1))b^2/\log_2 b$ and $\eta = (H_2(\rho) + o(1))/b$. These sharper parameter statements follow by optimizing the length in the existing moment construction.

1.2 Why the construction works

There is a useful way to reverse the usual root-counting argument. Choose a set A of $r = k + s$ evaluation points, and form its root polynomial

$$F_A(X) = \prod_{a \in A} (X - a).$$

Suppose many choices of A give the same coefficients in degrees $r, r-1, \dots, k$. Let W be this common high-degree part. Then $G_A = W - F_A$ has degree less than k , and it agrees with W at every point of A . Thus each such set produces a member of one decoding list. The problem has become a counting question about leading coefficients.

Those coefficients are determined by the first s power sums of A . When $A \subseteq \{0, \dots, n-1\}$, each power sum is a relatively small *integer*. There are $\binom{n}{r}$ possible sets, but far fewer possible vectors of s power sums. Some vector must occur many times. The crucial point is that we count these vectors over the integers *before choosing the prime*. The count therefore does not grow with the field size.

For a small example, take

$$A = \{0, 1, 6, 7\}, \quad B = \{0, 2, 5, 7\}.$$

Both sets have size four and sum fourteen. Their root polynomials therefore begin with $X^4 - 14X^3$. Subtracting either polynomial from $W = X^4 - 14X^3$ leaves a polynomial of degree at most two. Over any prime larger than seven, on the domain $\{0, \dots, 7\}$ each agrees with W at four positions. The general construction repeats this idea with many more sets and more moments. This example illustrates the mechanism; its two codewords are not the quantitative lower bound of Theorem 1.1.

1.3 Lines of received words

Proof systems often test a random linear combination of several words. This leads to a related question. Suppose many words on an affine line $f_z = f_0 + zf_1$ are close to the code. Must the nearby codewords fit together into a single affine line of codewords? Must f_0 and f_1 agree with two codewords on a large common set of coordinates?

We give three constructions, because these questions have different quantifiers.

Many nearby points without correlated agreement. Corollary 4.15 gives complete nonzero coverage with an almost maximally far exception. On a paired core, even locator polynomials leave codeword candidates after their common leading term cancels. A padding block supplies several pairs of new agreements at once. Multiplicative-character estimates show that the candidates' joint evaluation images fill almost all possible vectors, so random padding directions cover every nonzero line parameter.

Theorem 4.41 gives complete coverage against arbitrary fixed numerical constants. Matching three times as many moments permits a cubic change of the seed domain without losing the required common locator prefix. A geometric irreducibility argument and a finite-field point-count estimate then control outside collisions.

The interval constructions also give elementary alternatives with quantified nearby density. Division by $X - \tau$ yields $\Omega(p/\log p)$ nearby labels at logarithmic length (Theorem 3.4); padding gives the stronger density bounds in Theorems 4.43 and 4.45. Corollary 4.48 makes the number of coordinates outside the radius grow while retaining density tending to one. These variants use padded interval domains. The complete-coverage constructions use random paired domains or cubic images of intervals with added coordinates.

An obstruction that persists at a fixed gap. The gap in Theorem 1.1 shrinks with n . To address asymptotic claims at a fixed gap, we start with one fixed large list and replace each evaluation point by B points on which X^B has the same value. Polynomial composition preserves the rate, relative agreement, and list size. Suitable prime fields exist for every B .

An anchored version removes one common agreement point and divides by its linear factor. Padding then produces many nearby challenges while keeping every common coefficient agreement

below the threshold. Theorem 4.6 gives actual failure of correlated agreement at a fixed rate and gap, with a linear count whose coefficient has logarithm of order $\eta^{-2}/\log(1/\eta)$. Every nearby word on these lines has exactly one nearby codeword, while the parameter-zero word is one coordinate beyond the radius. The obstruction persists without local list ambiguity. The same exceptional-count lower bound also has realizations with $p = n^{O_{\rho,\eta}(1)}$ (Corollary 4.7); uniqueness is not asserted for those smaller-field realizations. No remainder that is $o(n)$ along this family can absorb a smaller linear coefficient. The simpler selection argument in Theorem 4.4 is retained as a preliminary construction.

Small lists with many isolated nearby line points. Large lists are not the only obstruction. At every fixed rate $\rho = k/n$, we also construct codes whose maximum list size is exactly $n - k - 1$, together with a line having $\binom{n}{k+1}$ nearby parameters. Each nearby word has a unique nearby codeword, yet no parametrized codeword line contains more than $n - k$ of those witnesses. The gap is $1/n$. At rate one-half these bounds are $n/2 - 1$, $\binom{n}{n/2+1}$, and $n/2$, respectively. The construction combines an elementary line spanned by consecutive powers with a domain chosen using a known list-decoding theorem for random Reed–Solomon codes. Proposition 5.4 shows that its nearby-parameter count is optimal among lines without correlated agreement at this radius.

What remains open at a fixed positive gap. Our fixed-gap lower bound forces a large coefficient of n , but does not force a superlinear exponent. Proposition 4.9 identifies a sufficient route: a source list of size L at length N gives $\Omega(\min\{NL, p\})$ full-support MCA exceptions over the same prime field, after exactly halving both rate and gap. Thus lists of size N^c over superpolynomial prime fields would force exceptions of order N^{c+1} . The missing ingredient is a growing fixed-gap list on domains short enough relative to the prime; full-length lists with $N = p - 1$ cannot give superlinear scalar-label counts. The factor-of-two parameter change is not essential: Corollary 4.10 uses a small positive padding fraction to preserve any strict margin above a continuous sufficient agreement curve. Thus the same exponent transfer can stay inside an open regime where a sharper upper bound is available.

Sharp algebraic counts need not give sharp proximity gaps. For a differential equation of fixed order d , fixed jet degree, and fixed challenge degree, Appendix I proves that the actual regular polynomial–challenge locus has at most $O(D^{d+1})$ isolated points for message degree D . An explicit Wronskian family attains $\binom{D+d+1}{d+1}$ such points, all regular, with distinct labels over suitable prime fields. Nevertheless, that family has only constantly many nearby labels on any received line at a fixed positive agreement gap. At every fixed order d , Appendix J also bounds the degree of each positive-dimensional actual stratum by $O(D^d)$, and gives examples attaining this exponent. The first-order case in Appendix H is linear in D . These results distinguish sharpness of algebraic candidate counts from the still-open fixed-gap superlinear proximity-gap lower bound over sufficiently large prime fields. Proposition I.4 extends this obstruction to arbitrary candidate families with bounded-degree rational residuals after a common normalization: their full-support exceptional count is linear, regardless of how many algebraic candidates they contain. Theorem H.5 gives a further first-order restriction: bad labels with a positive fraction of nonsingular agreement coordinates number only $O(n)$ at a fixed positive gap. For $R(X, z)P' = \mathcal{A}(X, z, P)$ this bounds all bad labels linearly when $\deg_X R$ also stays a fixed positive fraction below the agreement threshold. A nonlinear two-solution family attains $n/2$ bad labels at fixed rate and gap. The general first-order conjecture still requires controlling singular agreement coordinates. Appendix M removes that restriction for a fixed nonlinear Riccati equation, proving a linear full-support bound and a sharp ordinary-list bound with prime-field equality families. Appendix P gives a different criterion for

challenge-dependent equations: if the highest coefficient equations determine all nonconstant message coefficients by a triangular linear system, bounded challenge and nonlinear degrees imply only $O(n)$ bad labels at fixed agreement excess. This includes candidates singular at every agreement coordinate. The criterion applies to a specified differential equation covering the candidates; it does not construct such an equation for an arbitrary received line or settle the general first-order conjecture. Appendix Q removes both the triangular pivot restriction and any bound on the derivative coefficient's degree for bounded-challenge monic Riccati families. Above agreement $5D/3$ with fixed linear slack, these families have only $O(n)$ bad labels, even when the singular coordinates move with the challenge. This includes the quarter-rate first-order agreement regime. A separate construction attains linearly many bad labels with every agreement coordinate singular, showing that singularity alone does not force a superlinear exceptional count.

Appendix R controls a broader singular mechanism: dependence of degree at most three on the derivative, with bounded dependence on the value and the challenge and unrestricted coefficient degrees in X . If the specialized equation does not vanish identically in both the derivative and the challenge at any domain coordinate, then agreement $A > \sqrt{Dn}/2$ with fixed normalized slack gives only $O(n)$ bad full-support labels. The proof uses both the value and derivative at singular matches, and includes polynomial solutions that appear only at special challenge values. Thus a large persistent singular set is not by itself an obstruction in this regime; an uncontrolled set of coordinates carrying no equation remains a limitation. As with the preceding results, one common identity must contain the candidates.

For one explicit equation we can describe the obstruction exactly. Its complete polynomial solution family contains a linear-size list at agreement $3/8$, but every received word has only $O(\varepsilon^{-2})$ members of that family at agreement $3/8 + \varepsilon$. A majority construction attains $\Omega(\varepsilon^{-2})$ for sufficiently large primes (Theorem R.5). Thus the growing list in this family has an exact agreement threshold. This conclusion concerns the classified solution family, rather than the full Reed–Solomon code. We also control every degree-two rational first-integral family in characteristic $p > 10D$: at agreement $D + \eta n$, its list size is at most $\max\{40, \lfloor 1/\eta \rfloor\}$ (Theorem R.6). This includes denominators depending quadratically on the candidate value. The proof derives a height bound from five polynomial sections and applies the classical square-value theorem of Pasten and Wang; it gives a family-specific list bound, not a general proximity-gap bound.

For weighted monic-cubic first integrals with denominator depending only on X , we prove a corresponding bound of $O(1/\eta)$ at agreement $D + \eta n$, in characteristic zero or $p > \max\{3, 2D\}$ (Theorem R.9). On prime-field domains this includes every rate below $1/2$. The proof treats repeated basepoints and positive-dimensional critical loci: nonconstant Frobenius critical values force too few distinct denominator roots to support a large list. These conclusions concern the specified first-integral families; a cubic differential equation alone does not suffice. More generally, every fixed monic value degree has an $O_b(1/\eta)$ list bound when the original critical locus is finite (Theorem R.7). In characteristic zero, or $p > b(b-1)D$, a larger list forces a constant fiber with a repeated factor (Corollary R.8). This is a necessary condition on the first integral, not a classification of the remaining families.

A separate exact construction gives a quarter-rate Reed–Solomon code of length 28 and a received word whose complete nearest-codeword list has exactly seven members, each agreeing on exactly half the coordinates (Theorem U.1). It holds over arbitrarily large prime fields. More generally, a polynomial pullback preserves this exact list size at length $14m$ and degree bound $3m$ for every $m \geq 1$. The construction follows from two explicit cyclotomic identities; completeness follows from saturated pairwise root counts and a seven-point interpolation obstruction. At its underlying cubic parameters, seven is optimal for every fourteen-point domain in characteristic zero and in every sufficiently large characteristic (Theorem V.1). This sharp upper bound concerns the

cubic parameters, not arbitrary quarter-rate length-twenty-eight codes.

Relaxing the length gives a complete eight-member nearest list: at $n = 32$ and dimension $k = 7$, all eight codewords have exactly 14 agreements, and no other codeword attains that agreement (Theorem W.6). This holds over arbitrarily large prime fields and lies strictly above the first-order curve at rate $7/32$. The construction gives an explicit rational sixteen-node cubic incidence configuration and proves completeness by exact interpolation and a quadratic norm calculation. It is a finite seed; growing the list while preserving this agreement regime remains open. Shared intersections of two additional cubics strengthen this to a complete ten-member nearest list on eighteen nodes, with seven agreements each. A degree-five pullback gives at least ten candidates at $(n, k, A) = (90, 16, 35)$, strictly above the first-order curve (Theorem W.7). Completeness is proved for the cubic source, but is not asserted for this pullback. A triangle of shared intersections similarly gives a complete eleven-member cubic list on nineteen nodes, with seven agreements each; its degree-thirteen pullback has at least eleven candidates at $(n, k, A) = (247, 40, 91)$ above the same curve (Theorem W.8). These finite extensions do not yet give an unbounded-list family. A separate incidence exchange gives a complete nine-member quartic list with eight agreements on eighteen nodes. Its degree-twelve pullback has at least nine nearby codewords at $(n, k, A) = (216, 49, 96)$, again above the first-order curve (Theorem X.1). This step removes two selected incidences before lifting and restores both at a fresh intersection; it does not assert an indefinitely repeatable extension.

Another construction verifies two successive one-pole extensions at quarter rate. Over arbitrarily large prime fields, it gives six distinct polynomials of degree at most fourteen, each agreeing with one word on thirty of sixty coordinates (Theorem T.1). The proof uses an exact finite-field witness and two independent nonsingular lifting certificates. This is a fixed-size construction; unbounded iteration remains open.

1.4 Finite examples on circle codes

We also study a different regime: the fixed prime $p = 2^{31} - 1$ and domains on the norm-one circle. These are structured evaluation domains used by Circle STARKs. At the common parameters

$$n = 2^{26}, \quad \rho = \frac{1}{2}, \quad \eta = \frac{33}{1024} - \frac{2}{n},$$

we obtain the following finite lower bounds:

Object constructed	Guaranteed size
One base-field decoding list	$> 2^{40}$
Distinct nearby parameters on an extension-field line	$> 2^{74}$

The line has no correlated agreement at the specified threshold. The constructions give existence with explicit numerical certificates; they do not enumerate these very large lists. Section 6 explains the idea, and Appendix A gives the coefficient calculations.

1.5 Context and organization

Proximity gaps are established tools in the analysis of low-degree tests [8, 25]. In particular, existing theorems give proximity guarantees in specified radius regimes. Our fixed-rate asymptotic examples have $\eta \rightarrow 0$ and eventually lie beyond the Johnson radius $1 - \sqrt{\rho}$, although they remain below the characteristic-based Elias radius. The distinction between these radii is essential when comparing the results with proved soundness bounds for a low-degree test.

The connection between equal subset moments, elementary symmetric functions, and Reed–Solomon decoding is established prior work; Gandikota, Ghazi, and Grigorescu use this connection in decoding hardness results [16]. The integer moment pigeonhole argument is also classical: Borwein, Erdélyi, and Kós use derivative vectors of binary polynomials in the proof of their Theorem 2.7 [15]. Restricting that argument to fixed-weight subsets and retaining a large fiber, followed by the standard locator correspondence, already yields the asymptotic exponent $\log L = \Omega_\rho(\eta^{-2}/\log(1/\eta))$ used here. Our interval counts give sharper finite bounds and an application to the specific uniform prime-field conjecture; we do not claim a new moment counting mechanism. The affine-line distance profiles require additional constructions. The number-theoretic splitting theorem used in the lift is also classical.

Prime-field list constructions using multiplicative cosets also appear in work of Guruswami and Rudra, presented in Rudra’s thesis [17, Section 6.4.3]. Those full-field examples already use cancellation of leading coefficients through unions of multiplicative orbits. Their displayed coset bank has logarithmic size $O(1/\eta)$; our interval construction gives the larger, almost-quadratic reciprocal-gap exponent above. Neither comparison supplies growing lists at one fixed positive gap.

The generic list-profile upper bound is due to Brakensiek, Gopi, and Makam [13]. At every fixed positive rate, Roth’s ordinary list-decoding bound [14] gives the matching lower bound for all sufficiently large lengths. We include an explicit construction that also covers the remaining finite parameters $k \geq 2$, without claiming priority for that construction. The separation here combines these list bounds with a line containing many isolated nearby points.

Brakensiek, Chen, Putterman, Zhang, and Zheng [12] give polynomial-in- p time list decoding up to capacity for arbitrary prime-field evaluation sets at fixed rate and fixed slack, subject to their field-size hypotheses. In particular, their all-rates statement assumes $p \geq C(R, \delta)n$; their formal Corollary 5.1 gives runtime and list size $p^{O_{R,\delta}(1)}$. Such polynomial list bounds allow growing lists; they do not imply the uniform reciprocal-gap list bounds or the affine-line conclusions considered here.

Goyal, Guruswami, Sun, and Wootters [26] give further recent proximity-gap and correlated-agreement bounds for random Reed–Solomon domains. Their results allow general received curves and retain explicit slack and field-size hypotheses. They provide useful context for the shrinking-gap regime studied here, but do not subsume the specific actual-list-size separation in Theorem 5.1.

Krachun, Kazanin, and Haböck [21] construct large lists and affine-line obstructions on multiplicative subgroups of prime fields, with a gap of order $1/\log n$. Our interval construction addresses a different family of domains and gives lists larger than every fixed exponential in the reciprocal gap. The paired-domain construction adds a different line phenomenon: two quantitatively far inputs have all $p - 2$ non-endpoint mixtures nearby, and the underlying distance profile can have one almost maximally far point and $p - 1$ nearby points. These statements use specially constructed domains and a shrinking gap. The paired-domain theorem uses a superpolynomial alphabet in the block length. Theorem N.37 gives the two-far profile instead over a polynomial-size prime alphabet, with every non-endpoint mixture nearby at a radius strictly below Elias. Its domain is a selected union of multiplicative cosets, not necessarily a subgroup. Krachun–Kazanin–Haböck already obtain polynomial prime alphabets and below-Elias inverse-logarithmic gaps; those features alone are not new. An origin shift in their quotient construction also gives two far endpoints with the same asymptotic half-gap separation and exact source/common agreement. The strengthening here is coverage of *all* $p - 2$ interior mixtures in the polynomial-alphabet, below-Elias regime, rather than only the large subset certified there. Thus the comparisons concern quantitative coverage, separation, and reciprocal-gap dependence; prime-field counterexamples themselves are already established prior work. Our finite circle examples concern the evaluation spaces of Circle STARKs [28].

The Proximity Prize’s orbit-pencil upper certificate also uses multiplicative fibers and a fixed

core [29]. Our circle construction expresses the analogous core through a common anchor and division at two circle points. Appendix A.6 makes this change of presentation explicit. The contest’s field, code, and score are different from the circle specialization considered here.

The distinction between fixed and shrinking gap separates our results from the fixed-gap polynomial-in- n list bound stated in the recent preprint of Jeronimo [19]. Our exponential-list family has a shrinking gap. Our fixed-gap family preserves a fixed finite list of witnesses as n grows; it does not bound the full decoding list by that constant. Neither construction contradicts such a polynomial bound.

The current Dao–Kominers–Thaler paper [43] gives, above its first-order curve by margin η_1 , list bound $O_\rho(n/\eta_1^2)$ and full agreement-set MCA exceptional count $O_\rho(n^2/\eta_1^4)$ in sufficiently large characteristic. Its capacity bounds have exponents depending on the fixed capacity gap. There are two qualitative intrinsic lower bounds: Proposition 2.5 gives $n/2$ list elements over prime fields at fixed rate $1/4$ and gap $1/8$; Corollary 2.7 gives $\Omega(n^2)$ full agreement-set MCA exceptions over quadratic extension fields of growing prime characteristic, at fixed rate $1/8$ and gap $1/16$. The latter rules out a universal linear capacity MCA bound in that larger field class. Theorem N.1 also gives $\Omega(n^2)$ nearby labels with no ordinary correlated agreement, over quadratic extensions at exact rate $3/13$ and capacity gaps bounded below by $3/26$ (not fixed exactly). Corollary N.3 obtains quadratic ordinary-CA counts at exact rate $1/8$ and gap $1/16$ over $\mathbb{F}_{p^2}$, with at least $\lceil n^2/192 \rceil$ exceptions (Proposition N.4). Thus quadratic length dependence is necessary for ordinary correlated-agreement exceptions in the large-characteristic capacity regime. In particular, at these parameter pairs no uniform $o(n^2)$ exceptional-count bound holds in this field class. Corollary N.5 extends the exact parameters to rate b/d and gap $1/d$ for $b \in \{2, 3, 4, 5\}$ and $d \geq b + 7$, over the fixed-degree extensions $\mathbb{F}_{p^{2b-1}}$, including rate $2/9$ and gap $1/9$ over $\mathbb{F}_{p^2}$. These results do not establish first-order-regime tightness or force an exponent growing as the gap shrinks. Over prime ambient fields, even a superlinear exceptional-count lower bound at one fixed positive gap remains open here. The large gap-dependent coefficient in our linear lower bound must not be confused with a growing exponent of n . There is now a prime-field lower bound at error rates covered by the Dao–Kominers–Thaler upper bounds, albeit at vanishing code rate. Theorem N.8 constructs a code of dimension three and length $n = 2L(L - 1) + 1$, with more than $n^{3/2}/4$ exceptional challenges. Both source agreements and common agreement are $2L - 2$, but every exceptional word has a singleton list at threshold $T = 2L - 1$. This threshold is above the first-order curve and below exact Johnson: $T^2 = 2n - 1$. The construction works over every sufficiently large prime field, with a stated polynomial field-size bound, and an explicit finite certificate gives the uniform upper comparison $31000n^2$ for full-agreement-set MCA exceptions. Thus *even over prime fields, many exceptional challenges can coexist with singleton lists at error rates covered by the first-derivative upper bounds*. Corollary N.9 strengthens the field-size accounting: for an infinite sequence of primes, $n = \Theta(p^{2/3})$ and more than $p/36$ challenges have singleton threshold lists. Thus the exceptional probability can stay bounded away from zero despite these lists being singletons. The rate is $3/n$, and the source common agreement misses the tested threshold by only one coordinate. This is a superlinear lower bound in the applicable regime, not fixed-rate or fixed-gap tightness.

The one-coordinate restriction can be removed without increasing the message dimension. Theorem N.11 gives, for every sufficiently large prime, two words whose individual and common agreements are exactly A , while a constant fraction of their linear combinations have exactly one codeword with at least $A + d$ agreements. Here d grows on the order of $\log n / \log \log n$. For example, one choice gives $d \sim \log n / (2 \log \log n)$ and $\Omega(n \log \log n)$ singleton exceptional challenges. The construction replaces polynomial padding by independent random values with the two source lists excluded, and uses a multiplicative Sidon set to realize the core over small prime fields. Its scope remains vanishing rate and vanishing fractional gap; it does not prove fixed-rate tightness. Moreover,

on these positive-density examples $p = n \text{ polylog}(n)$, so the uniform $O(n^2)$ exceptional-count upper bound exceeds the entire field and gives no nontrivial probability bound. The construction therefore does not establish sharpness of that quadratic upper bound.

Allowing the message dimension to grow gives a polynomial source separation. Theorem N.13 constructs, for every fixed $0 < \alpha \leq 1/3$, examples with dimension and source separation $\Theta(n^\alpha)$ and $\Theta(n^{3(1-\alpha)/2})$ singleton exceptional challenges. For example, separation $\Theta(n^{1/5})$ coexists with $\Theta(n^{6/5})$ exceptions. The construction groups coordinates into complete fibers of $x \mapsto x^s$ and makes their fresh values agree. A separate argument excludes all new polynomial witnesses introduced by the larger code dimension. The fractional source gap and code rate still tend to zero; this does not prove fixed-rate tightness. Over sufficiently large primes, a deterministic variant makes every list above the source agreement empty or a singleton, while attaining the separate conic-witness upper bound up to a constant factor (Proposition N.17). This is a statement about all points on the constructed line, not about all received words: other received words for the same code can have at least L witnesses at the tested threshold.

These statements should not be read as the first superlinear prime-field proximity-gap counterexamples. Diamond and Gruen [45] refute the capacity conjecture using random-word distance distributions. Krachun, Kazanin, and Haböck [21] give prime-field counterexamples at distance $1 - \rho - \Theta_\rho(1/\log n)$ from the code, including fixed positive rates. Their distance from capacity is a different parameter from our source separation $d/n = \Theta(\log n/(n \log \log n))$. Our examples lie in the range of decoding radii covered by the first-derivative upper bounds of Dao, Kominers, and Thaler [43]. Their purpose is to lower-bound the number of exceptional challenges within that range, even when the relevant lists are singletons. This does not supersede those fixed-rate near-capacity results. Throughout this comparison, “below Johnson” for an agreement threshold means a decoding radius beyond the Johnson radius.

An independent projective construction in Theorem N.19 gives the same $3/2$ exponent over extension fields with an exact two-level agreement spectrum: all exceptional words have agreement $2(p^2 + p + 1)$ and every other word has agreement $2(p + 1)$. Its affine-witness companion in Corollary N.24 asymptotically attains the universal triple-count bound with constant one. That separate sharpness statement is not above the rate-only first-order curve.

Proposition N.20 improves the characteristic scale of the extension-field construction from $\Theta(n^{1/4})$ to $\Theta(\sqrt{n})$, retaining the $n^{3/2}$ exceptional-count scale and $\sqrt{n}$ agreement-gap scale. Explicitly, over $\mathbb{F}_{p^4}$ it has $n = 9p^2$, dimension three, and threshold $4p$, while both source agreements and common agreement are exactly $2p$. There are exactly $(p + 1)(p - 1)^2$ singleton exceptions and one two-word exception. The construction satisfies an explicit finite DKT certificate whose characteristic condition is only $p > 3$. Its ambient field has size $n^2/81$. This improves the characteristic and field-size parameters; it does not increase the known exceptional-count exponent or yield prime-alphabet, fixed-rate, or general-bound tightness.

For ordinary common agreement, Theorem N.21 gives a stronger probability statement over $\mathbb{F}_{p^3}$. With $n = 2p^2 - p - 1$ and dimension three, every affine mixture has agreement at least $2p - \lceil 4\sqrt{p} \rceil - 2$, while the two inputs have ordinary common agreement at most $p + 2\sqrt{p}$. Exactly $p^3 - p$ of the p^3 affine parameters have singleton lists; the remaining p have lists of size p at the stated threshold. The threshold satisfies an explicit finite first-order certificate and is below the exact Johnson agreement threshold for primes $p \geq 4099$. Thus the common-agreement deficit has order $\sqrt{n}$ even when every mixture is nearby and all but a $1/p^2$ fraction have unique witnesses. The inputs themselves are nearby, so this result addresses common agreement rather than separation of individually far sources. The ambient field is an extension field of size $\Theta(n^{3/2})$; the construction does not instantiate the fixed better.codes benchmark.

Theorem N.22 upgrades this construction to p individually far inputs over the same field. Their

unique nearby quadratic is the same polynomial. Deleting just its core agreement fiber, of size at most p , lowers their individual agreements to at most $p + 2\sqrt{p}$. Every other nearby quadratic loses at most two matches, or has a disjoint core fiber. At threshold $2p - \lceil 4\sqrt{p} \rceil - 4$, the exact affine-line spectrum becomes p empty lists, $p^3 - 2p$ singleton lists, and p lists of size p . The length remains $2p^2 - O(p)$, and the finite first-order certificate still applies for $p \geq 4099$. Thus *two far inputs have a nearby mixture with probability exactly* $1 - p^{-2}$, with source separation and characteristic both of order $\sqrt{n}$. The far parameters form an affine $\mathbb{F}_p$ -line, so combinations with $\mathbb{F}_p$ coefficients stay far; the high failure probability uses coefficients from $\mathbb{F}_{p^3}$. This is a vanishing-rate extension-field result, not a prime-alphabet assertion.

For list size itself, Theorem N.25 gives a matched constant-factor comparison. At length $N = 2(p^2 + p + 1)$ over $\mathbb{F}_{p^3}$, dimension three, and threshold $T = \lceil 19p/10 \rceil$ for $p \geq 53$, an explicit word has exactly $N/2$ nearby quadratics, while a finite DKT certificate bounds every word's list by $22N$. Thus the uniform list bound has the correct linear dependence on N along this parameter family. The threshold stays a constant times $\sqrt{N}$ away from both the first-order curve and the Johnson boundary. Corollary N.26 extends the matched linear growth to every fixed threshold $T = \lceil cp \rceil$ with $\sqrt{3} < c < 2$: for $p \geq 401$, the same word has $N/2$ witnesses and the uniform upper bound is $137N$. Thus linear list growth is necessary throughout the asymptotic interior between the two boundaries, rather than only at one chosen threshold. More generally, Corollary N.27 gives $\Theta_c(N)$ lists for every fixed $\sqrt{2} < c < 2$, using the same finite support machinery with larger fixed derivative caps. The additional range below $\sqrt{3}$ is a finite-degree consequence, rather than a change to the rate-only first-order curve. The example uses classical projective-plane incidence geometry; its rate and normalized agreement vanish. This establishes neither prime-alphabet tightness nor tightness of the received-line exceptional count.

For a structured algebraic realization at fixed rate, Theorem N.28 uses length $N = p^5$ over $\mathbb{F}_{p^5}$ and gives at least $N - 3$ nearby native challenges, with common-agreement gap $\Theta_\rho(N^{3/5})$; on an infinite sequence of primes every challenge is nearby. It combines the locator compiler of [42] with internal padding, a published scattered-set exclusion, and an exact cubic criterion for the remaining labels. Here characteristic is $N^{1/5}$, but is smaller than the message degree. The first source is nearby and the direction is far. As an existence statement, its count and asymptotic gap are subsumed by Crites–Stewart [23, Corollary 1], even with the same full domain and exact common agreement. It is not a new existential regime or a large-characteristic, above-first-order comparison.

A norm-quotient specialization gives a stronger gap with two far sources (Theorem N.31). For every fixed extension degree $d \geq 2$, length $q = p^d$, and exact dimension $J = \lfloor \rho q \rfloor$, put $m = (q - 1)/(p - 1)$. Each source has agreement $J + m - 1$, and their common agreement is also $J + m - 1$, while all $q - 2$ nonendpoint affine combinations have at least $J + 2m - 1$ agreements. Thus the common-agreement loss is about $1/p$, half the approximately $2/p$ distance below capacity. This combines the standard quotient construction with Katz's character-sum estimate; it still has a smaller capacity margin than the inverse-logarithmic counterexamples in [21]. Its full-field domain and extension alphabet do not establish the practical benchmark or prime-alphabet tightness. Moreover, Crites–Stewart already give every affine challenge nearby at inverse-logarithmic margin over arbitrary fields and domains. The additional property of the norm construction is its two far sources with exact source and common agreement; novelty of that refinement is not established here. By choosing only logarithmically many norm fibers, Theorem N.34 obtains the same two-far-source profile with a common-agreement loss of order $1/\log n$, asymptotically half its distance below capacity. Here $n = \Theta(p^{d-1} \log p)$; for quadratic extensions the characteristic is $\Theta(n/\log n)$. This matches the known order of capacity margin, while retaining the extension-alphabet and specially chosen-domain limitations just described. Theorem N.37 also removes the extension-alphabet requirement: for every

fixed $\beta > 12/5$, it gives prime alphabets $p = \Theta((n/\log n)^\beta)$ on selected unions of multiplicative cosets, with the same two-far profile and inverse-logarithmic gap. More precisely, for every $C > 1/h(\rho)$, where $h(\rho) = -\rho \log \rho - (1 - \rho) \log(1 - \rho)$, its distance below capacity is $(2/C + o(1))/\log p$ and its source/common-agreement loss is $(1/C + o(1))/\log p$. Choosing $C < 2/h(\rho)$ places the tested radius strictly below Elias. The proof first covers almost all multiplicative labels with fixed-size products, then fills the remaining holes by adding a constant number of pairs of alternatives. Each representation uses exactly one factor from each pair, preserving its required degree. This retains the earlier paired theorem’s below-Elias two-far behavior while allowing a polynomial prime alphabet. It does not establish fixed-rate first-order tightness or settle a prescribed-domain benchmark. An exact below-Elias certificate over $\mathbb{F}_{65537}$ has length 48, dimension 24, source/common agreement 25, and agreement exactly 27 at every nonzero pencil label. Its distance below capacity is $1/16$, and its source-to-near distance loss is $1/24$. Corollary N.38 gives the same ratio $2/3$ between loss and capacity margin over every sufficiently large odd prime, with $n = \Theta(\log p)$ and radius strictly below Elias. This last family has exponential alphabet size in n ; it is distinct from the polynomial-alphabet family above. Novelty of the two-far refinement relative to the wider literature remains unestablished.

Likewise, optimality of a specified interpolation dimension count is a statement about that proof method, not an intrinsic list or MCA lower bound. The capacity gap and the margin above the first-order curve are different parameters. A separate support-count audit establishes a restricted form of method tightness for the current first-order formulas. At rate $1/4$, every monomial support downward in the value variable, with full derivative-weighted coefficient prefixes, necessarily incurs inverse-margin powers two in the squarefree list budget and four in the regular-family MCA budget, under either current challenge-counting test. The row-test proof uses a positive-feasibility reduction that deletes, truncates, and rearranges columns; ordinary sorting alone need not work. For that test the result also covers total-degree-graded, translation-stable nonmonomial jet spaces. A separately justified filtered row test gives the same costs when the maximum jet degree is at most the code’s polynomial degree bound. The proof uses a corrected finite-length margin and assumes characteristic zero or characteristic exceeding every retained total jet degree. It does not cover arbitrary filtered nonmonomial spaces, multiple equations, or counting actual component degrees. The proof and precise hypotheses appear in the companion note, *A Sharp Quarter-Rate Threshold for a First-Derivative Interpolation Dimension Count*, maintained with this manuscript.

Sections 2–5 prove the general list and line results. Section 6 gives finite circle examples. Section 8 explains what these results imply for hash-based SNARK analyses, and Section 9 applies them to S-two. Section 10 states the remaining coding-theoretic questions. The first appendices contain the circle proofs, quotient calculations, and finite certificates. The later appendices study restrictions on possible constructions: rational fibers, binary-affine locator families, and polynomial differential equations. These additional results address the open fixed-gap problem; the main list and line counterexamples do not depend on them.

2 Large lists from subsets with equal moments

The construction starts with a simple observation: if two polynomials differ by a polynomial whose roots are a set A , then they agree at every point of A . We will choose many sets A whose vanishing polynomials have the same leading coefficients. Subtracting those common coefficients will give many low-degree polynomials close to one received word.

Notation. For a field F and a set $D \subset F$ of n distinct points, write

$$\text{RS}(F, D, k) = \{(P(x))_{x \in D} : \deg P < k\}.$$

This code has dimension k and rate $\rho = k/n$, provided $k \leq n$. The relative Hamming distance between two words is the fraction of coordinates on which they differ. A list at radius θ consists of all codewords within that distance of a received word. Write $\ell_C(\theta)$ for the largest such list over all received words. We parameterize the radius by

$$\theta = 1 - \rho - \eta.$$

Thus a nearby codeword agrees with the received word on at least a $\rho + \eta$ fraction of the coordinates. We call $\eta > 0$ the gap from $1 - \rho$.

All logarithms are base two, and $H_2(t) = -t \log_2 t - (1 - t) \log_2 (1 - t)$ is binary entropy. For characteristic p , define

$$H_p(t) = \frac{H_2(t) + t \log_2 (p - 1)}{\log_2 p}.$$

The characteristic-based Elias radius is the solution of $H_p(r_E) = 1 - \rho$ on the increasing branch $[0, 1 - 1/p]$. We will keep our examples strictly below this radius. The simple estimate

$$H_p(\theta) \leq \theta + \frac{1}{\log_2 p} \quad (1)$$

shows that it suffices to have $\eta \log_2 p > 1$. In our prime-field examples $p > n$ and $\rho \geq 1/n$, so $0 \leq \theta < 1 - 1/p$ as required.

2.1 Why matching moments produces a list

Consider the 34-element subsets A of $\{0, \dots, 63\}$. Their vanishing polynomials begin

$$F_A(X) = \prod_{a \in A} (X - a) = X^{34} - \left(\sum_{a \in A} a \right) X^{33} + \left(\sum_{a < b \in A} ab \right) X^{32} + \dots$$

Subsets with the same sum and the same sum of squares have the same three leading coefficients, because

$$\sum_{a < b \in A} ab = \frac{(\sum_{a \in A} a)^2 - \sum_{a \in A} a^2}{2}.$$

Let W be those common leading terms. Each polynomial $W - F_A$ has degree less than 32 and agrees with W at exactly the 34 points of A . Consequently, a large class of subsets sharing these two moments gives a large list in a rate-one-half Reed–Solomon code.

For the general construction we fix the first s moments. It is convenient to use *binomial moments*, $\sum_{a \in A} \binom{a}{j}$, because their integer ranges can be counted exactly. They determine the usual power sums and hence the leading coefficients in exactly the same way.

Lemma 2.1 (A finite list bound). *Let $n \geq 2$, $k \geq 1$, $s \geq 1$, and $t = k + s < n$. Let $p > n$ be prime and $D = \{0, \dots, n - 1\} \subset \mathbb{F}_p$. Put*

$$R_j = \binom{n}{j+1} - \binom{n-t}{j+1} - \binom{t}{j+1} + 1 \quad (1 \leq j \leq s), \quad L = \left\lceil \frac{\binom{n}{t}}{\prod_{j=1}^s R_j} \right\rceil.$$

There is a received word with at least L distinct codewords of $\text{RS}(\mathbb{F}_p, D, k)$ at distance exactly $1 - t/n$. If $p > \max_j (R_j - 1)$, its entire list at radius $1 - t/n$ corresponds exactly to one class of subsets with equal integer binomial moments.

Proof. For each t -subset $A \subset D$, record $T_j(A) = \sum_{a \in A} \binom{a}{j}$ for $1 \leq j \leq s$. The minimum and maximum values of T_j occur on the first and last t integers in D . The identity $\sum_{a=0}^{b-1} \binom{a}{j} = \binom{b}{j+1}$ gives their values as

$$\binom{t}{j+1} \quad \text{and} \quad \binom{n}{j+1} - \binom{n-t}{j+1}.$$

There are therefore at most $\prod_j R_j$ moment vectors. By pigeonhole, at least L subsets share one vector.

Equal binomial moments imply equal power sums through degree s : each X^j is an integer linear combination of $\binom{X}{0}, \dots, \binom{X}{j}$, and all subsets have size t . Newton identities then imply equal elementary symmetric functions through degree s . These identities are valid modulo p since $p > n > s$. Thus the monic polynomials $F_A(X) = \prod_{a \in A} (X - a)$ have the same coefficients in degrees $t, t-1, \dots, t-s = k$. Let W be their common degree- $\geq k$ part. Then $P_A = W - F_A$ has degree less than k and agrees with W exactly on A : outside A , each factor of $F_A(x) = \prod_{a \in A} (x - a)$ is nonzero in the field because $p > n$ makes the evaluation points distinct. Different subsets give different codewords, since these polynomials have degree less than n .

For the last assertion, any codeword P in the list makes $W - P$ a monic degree- t polynomial with at least t roots in D . It must therefore be F_A for a t -subset A . Its first s binomial moments are congruent to those of the selected class modulo p . Each lies in an integer interval of width $R_j - 1 < p$, so the congruences are equalities. $\square$

2.2 An exponential list and a subexponential proposed bound

There are exponentially many subsets of a fixed density. Fixing a modest number of moments divides them into far fewer classes. The following choice of parameters makes this discrepancy precise.

Theorem 2.2 (No uniform $\exp(O_\rho(1/\eta))$ list bound). *Fix a rational rate $\rho \in (0, 1)$. For arbitrarily large n , there are prime-field Reed–Solomon codes of exact rate ρ , on $D = \{0, \dots, n-1\}$, with*

$$\eta = \Theta_\rho\left(\frac{1}{\sqrt{n \log n}}\right), \quad \log_2 p = \Theta_\rho(\sqrt{n \log n}), \quad \log_2 \ell_C(1 - \rho - \eta) = \Omega_\rho(n).$$

The radius lies strictly below the characteristic-based Elias radius. Consequently no finite constants $c_1, c_2 > 0$ make

$$\ell_C(1 - \rho - \eta) \leq c_1 2^{c_2 H_2(\rho)/\eta} \tag{2}$$

hold uniformly over prime fields, evaluation domains, and these radii. The same sequence works for every fixed pair c_1, c_2 .

Proof. Let n grow through multiples of the denominator of ρ and set $k = \rho n$. Choose a constant $c > 0$ with $c^2 < 4H_2(\rho)$, and put

$$s = \left\lfloor c \sqrt{\frac{n}{\log_2 n}} \right\rfloor, \quad t = k + s, \quad \eta = s/n.$$

For large n , the hypotheses of Lemma 2.1 hold. Retaining the factorials in the binomial coefficients gives $R_j \leq 2 \binom{n}{j+1} \leq 2(3n/(j+1))^{j+1}$. Hence

$$\log_2 \prod_{j=1}^s R_j \leq s + \sum_{j=1}^s (j+1) \log_2 \frac{3n}{j+1} = \frac{s^2}{2} \log_2 \frac{n}{s} + O(s^2 + s \log n) = \left(\frac{c^2}{4} + o(1) \right) n.$$

The middle equality follows by summing $j \log j$, or by integral comparison. The final equality uses $\log_2(n/s) = \frac{1}{2} \log_2 n + O(\log \log n)$. On the other hand, $\log_2 \binom{n}{t} = H_2(\rho)n + o(n)$. The lemma gives

$$\log_2 L \geq \left(H_2(\rho) - \frac{c^2}{4} - o(1) \right) n = \Omega_\rho(n). \quad (3)$$

Choose any constant $A > H_2(\rho)/c$. Bertrand's postulate supplies a prime with $\log_2 p = A\sqrt{n \log_2 n} + O(1)$. Eventually $p > n$ and $\eta \log_2 p \rightarrow cA > H_2(\rho)$. Since $H_2(1 - \rho - \eta) \rightarrow H_2(\rho)$, the sharper entropy estimate $H_p(\theta) \leq \theta + H_2(\theta)/\log_2 p$ proves the Elias condition. In contrast, the logarithm of the right-hand side of (2) is

$$\log_2 c_1 + c_2 H_2(\rho) \frac{n}{s} = O_{\rho, c_1, c_2}(\sqrt{n \log n}) = o(n).$$

It is eventually smaller than the logarithm of the constructed list. $\square$

The rate is fixed, while the gap decreases and the prime grows. The theorem leaves open stronger bounds under restrictions on the domains or on the relation between field size and block length.

Corollary 2.3 (The obstruction as a function of gap and field size). *Fix a rational rate $\rho \in (0, 1)$ and write $H = H_2(\rho)$. For every sufficiently small real $\varepsilon > 0$, there is a length*

$$n = (H + o(1))/(\varepsilon^2 \log_2(1/\varepsilon))$$

such that, for every prime $p > n$, the interval-domain code $C = \text{RS}(\mathbb{F}_p, \{0, \dots, n-1\}, \rho n)$ satisfies

$$\log_2 \ell_C(1 - \rho - \varepsilon) \geq \frac{H^2/2 - o(1)}{\varepsilon^2 \log_2(1/\varepsilon)}. \quad (4)$$

This radius is strictly below the characteristic-based Elias radius whenever $\varepsilon \log_2 p > H_2(1 - \rho - \varepsilon)$.

Moreover, for every sufficiently large prime p , writing $b = \log_2 p$, there is such a code of length $n = (H^{-1} + o(1))b^2/\log_2 b$ and a radius strictly below Elias with

$$\log_2 \ell_C(1 - \rho - \varepsilon) \geq (1/2 - o(1)) \frac{b^2}{\log_2 b}, \quad \varepsilon = (H + o(1))/b. \quad (5)$$

All asymptotics fix ρ . These are consequences of the preceding interval construction.

Proof. Choose n to be a nearest multiple of the denominator of ρ to $H/(\varepsilon^2 \log_2(1/\varepsilon))$, and put $k = \rho n$, $m = \lceil \varepsilon n \rceil$, and $t = k + m$. For sufficiently small ε , we have $1 \leq m < t < n$. The range estimate in the proof of Theorem 2.2, applied to Lemma 2.1, gives

$$\log_2 \ell_C(1 - t/n) \geq Hn - \frac{m^2}{2} \log_2(n/m) - O_\rho(m^2 + m \log n + \log n).$$

Here $\log_2 \binom{n}{t} = Hn + O_\rho(m + \log n)$, and all error terms are $o(n)$ for the chosen parameters. Substitution gives (4): its two leading terms are respectively $H^2/(\varepsilon^2 \log_2(1/\varepsilon))$ and half that quantity. Since $m/n \geq \varepsilon$, the constructed radius is at most the claimed radius; list-size monotonicity handles every real prescribed gap, including the rounding of m .

The entropy estimate $H_p(\theta) \leq \theta + H_2(\theta)/\log_2 p$ proves the stated sufficient condition for strict distance below Elias. For the final claim, start with any sufficiently large prime, put $A_b = H + b^{-1/2}$, and take $\varepsilon = A_b/b$ in the first construction. Then $p > n$ eventually, $H_2(1 - \rho - \varepsilon) = H + O_\rho(1/b) < A_b$, and $\log_2(1/\varepsilon) = \log_2 b + O_\rho(1)$. Equation (4) now gives (5) and the asserted length. The fields grow and the gap shrinks; these statements impose no conclusion at a fixed field or fixed gap. $\square$

Corollary 2.4 (A second-order refinement of the gap bound). *Fix a rational rate $\rho \in (0, 1)$, and put*

$$H = H_2(\rho), \quad \gamma_G = \frac{3}{2 \ln 2} - 2 = 0.16404256 \dots$$

For every sufficiently small $\varepsilon > 0$, choose n to be a nearest multiple of the denominator of ρ to

$$\frac{H}{\varepsilon^2(\log_2(1/\varepsilon) + \gamma_G)}.$$

For every prime $p > n$, the interval-domain code $C = \text{RS}(\mathbb{F}_p, \{0, \dots, n-1\}, \rho n)$ satisfies

$$\log_2 \ell_C(1 - \rho - \varepsilon) \geq \frac{H^2}{2\varepsilon^2(\log_2(1/\varepsilon) + \gamma_G)} - O_\rho(1/\varepsilon). \quad (6)$$

The radius is strictly below the characteristic-based Elias radius if $\varepsilon \log_2 p > H_2(\rho + \varepsilon)$. In particular, for every sufficiently large prime p , writing $b = \log_2 p$, there is such a code at a radius strictly below Elias with

$$\log_2 \ell_C(1 - \rho - \varepsilon) \geq \frac{b^2}{2(\log_2 b - \log_2 H + \gamma_G)} - O_\rho(b), \quad \varepsilon = \frac{H}{b - H'_2(\rho)}. \quad (7)$$

These bounds retain the leading statements of Corollary 2.3; they refine the next term and do not remove its logarithmic denominator.

Proof. Set $m = \lceil \varepsilon n \rceil$ and $t = \rho n + m$. The uniform growing-moment Gram estimate from Section 2.5, converted to base-two logarithms, reads

$$\log_2 D_G = \frac{m^2}{2}(\log_2(n/m) + \gamma_G) + O_\rho(m \log n + m^4/n^2).$$

Here either of the two proved Gram denominators may be used. At the chosen length, $\log n \ll m \ll n$, and rounding m or n changes the displayed leading expression by $O_\rho(1/\varepsilon)$. Also $\log_2 \binom{n}{t} = Hn + O_\rho(m + \log n)$. Consequently the list bound is at least

$$Hn - \frac{\varepsilon^2 n^2}{2}(\log_2(1/\varepsilon) + \gamma_G) - O_\rho(1/\varepsilon).$$

The chosen n maximizes this quadratic up to rounding and proves (6). Since $m/n \geq \varepsilon$, list-size monotonicity gives the claimed radius. The entropy estimate used in Corollary 2.3 gives the stated Elias condition.

For the field-size assertion let $h = H'_2(\rho)$ and take $\varepsilon = H/(b - h)$. For sufficiently large b this is positive and $\rho + \varepsilon < 1$. Strict concavity gives

$$H_2(\rho + \varepsilon) < H + h\varepsilon = b\varepsilon,$$

so the radius is strictly below Elias. The chosen n is $\Theta_\rho(b^2/\log b) < p$. Substitution in (6) gives

$$\frac{(b - h)^2}{2(\log_2(b - h) - \log_2 H + \gamma_G)} - O_\rho(b).$$

Replacing $b - h$ by b changes this by $O_\rho(b/\log b)$, which is absorbed in the error in (7). $\square$

In particular, a bound with $\log_2 \ell_C(1 - \rho - \varepsilon) = o_\rho(\varepsilon^{-2}/\log(1/\varepsilon))$ cannot hold uniformly over prime fields and evaluation domains, even below Elias. In terms of field size, the exhibited lists are superpolynomial in p .

Full-length comparison at a fixed gap. The short-domain condition is a substantive distinction. The following elementary comparison already rules out a list bound depending only on rate and gap when the domain occupies almost the whole prime field. We make no novelty claim for this full-length observation.

Proposition 2.5. *Let $p \equiv 1 \pmod{8}$ be prime, $n = p - 1$, and $k = n/4$. On $\mathbb{F}_p^\times$, there is a received word with $n/2$ distinct degree- $(k - 1)$ polynomials, each agreeing on exactly $3n/8$ coordinates. Thus the rate is $1/4$ and the gap is $1/8$. For $p > 256$, the radius $5/8$ is strictly below the characteristic-based Elias radius.*

Proof. Let χ be the quadratic character, with $\chi(0) = 0$, and put $e = (p + 1)/2 = 2k + 1$. For $a \neq 0$ define

$$G_a(X) = \sum_{j=0}^k \binom{e}{2j+1} a^{e-2j-1} X^j, \quad P_a(X) = G_a(X) - X^k, \quad w(x) = \frac{1 + \chi(x)}{2} - x^k.$$

The degree- k terms cancel, and $[X^{k-1}]P_a = -a^2/8$. All powers of a are even, so the parameters a modulo sign give exactly $n/2$ distinct polynomials of degree $k - 1$. The binomial identity

$$G_a(s^2) = \frac{(a + s)^e - (a - s)^e}{2s}$$

counts their agreements. If $x = s^2$ is nonsquare, take s in $\mathbb{F}_{p^2}$. Then $s^p = -s$, and $u = (a + s)^e$ satisfies $u^2 = a^2 - x$ and $(a - s)^e = u^p$. Hence $G_a(x) = 0$ exactly when $\chi(a^2 - x) = 1$. There are $(p - 1)/4$ such nonsquare x .

For square $x = s^2$ with $s \in \mathbb{F}_p^\times$, Euler's criterion gives

$$G_a(s^2) = \frac{(a + s)\chi(a + s) - (a - s)\chi(a - s)}{2s}.$$

Away from $s = \pm a$, this equals one exactly when both characters are one. The weighted count $\frac{1}{4} \sum_s (1 + \chi(a + s))(1 + \chi(a - s)) = (p - 1)/4$ uses $\sum_s \chi(a^2 - s^2) = -1$. Removing $s = 0$ and correcting the half-weights at $s = \pm a$ cancel: at the endpoints the actual value is $\chi(2a) = \chi(a)$. Dividing by two for s and $-s$ gives $(p - 1)/8$ square agreements. The nonsquare count above follows from the same quadratic-character identity, applied to $(1 - \chi(x))(1 + \chi(a^2 - x))/4$ with its endpoints removed. Thus the total is $3n/8$. Finally, $H_p(5/8) \leq 5/8 + 1/\log_2 p < 3/4$ for $p > 256$. $\square$

Here $n/p \rightarrow 1$. In contrast, Theorem 2.2 and Corollary 2.3 give exponentially large lists on short integer intervals, with the stated quantitative dependence on the gap. The full-length example does not settle the fixed-gap, short-domain proximity-gap question. Exact checks of the displayed agreement counts are in `research/dickson_fixed_gap/`.

Lemma 2.6 (Splitting of pairwise differences). *For the polynomials above, if $a^2 \neq b^2$, every root of $G_a - G_b$ belongs to $\mathbb{F}_p$.*

Proof. Let t^2 be a root in an algebraic closure, with $y = G_a(t^2) = G_b(t^2)$. The case $t = 0$ is immediate. For $c \in \{a, b\}$, put $U_c = (c + t)^e$ and $V_c = (c - t)^e$. Then

$$U_c - V_c = 2ty, \quad U_c^2 - V_c^2 = 2c(t + t^p), \quad U_c^2 + V_c^2 = 2c^2 + 2t^{p+1}.$$

If $y = 0$, then $t^p = -t$, so $t^2 \in \mathbb{F}_p$. Otherwise put $H = (t + t^p)/(ty)$. We have $U_c + V_c = cH$, whence

$$c^2 H^2 + 4t^2 y^2 = 4c^2 + 4t^{p+1}.$$

Subtracting the equations for a and b gives $H^2 = 4$, and then $t^2 y^2 = t^{p+1}$. Thus $(t + t^p)^2 = 4t^{p+1}$, or $(t^p - t)^2 = 0$. Again $t^2 \in \mathbb{F}_p$. $\square$

Corollary 2.7 (Quadratic exceptions in growing prime characteristic). *For every prime $p \equiv 1 \pmod{8}$, $p \geq 17$, put $n = 2(p-1)$. Over $\mathbb{F}_{p^2}$ there is a Reed–Solomon code of length n and dimension $n/8$, and a received line with at least $\lceil n^2/20 \rceil$ full agreement-set MCA exceptional labels at agreement $3n/16$. Thus the capacity gap is exactly $1/16$, and $p = n/2 + 1 > n/8 - 1$.*

Proof. Set $N = p - 1$. Proposition 2.5 supplies $L = N/2$ candidates at agreement $A = 3N/8$, with dimension $k = N/4$. Choose any nonsquare anchor. Exactly $\ell = N/4$ candidates agree there: indeed $\sum_{a \in \mathbb{F}_p} \chi(a^2 - x) = -1$ for nonsquare x , none of the summands is zero, and the $a = 0$ summand is -1 . There are therefore $N/2$ admissible nonzero parameters, or $N/4$ modulo sign. As in Proposition 4.9, divide each anchored candidate minus the received anchor value by the anchor’s linear factor. The resulting quotients have degree at most $k - 2$. By Lemma 2.6, their evaluations are pairwise distinct at every point outside $\mathbb{F}_p$.

Append any $N + 1 = p$ distinct points of $\mathbb{F}_{p^2} \setminus \mathbb{F}_p$, with direction one, and retain direction zero on the $N - 1$ old coordinates. Independent uniform translations of the padding value sets yield an expected number of counted labels equal to

$$p^2[1 - (1 - \ell/p^2)^p] \geq p^2[1 - \exp(-N/(4p))] \geq Np[1 - \exp(-1/4)] > N^2/5 = n^2/20.$$

The second inequality uses concavity of $1 - \exp(-u/4)$ on $[0, 1]$; the last uses $1 - e^{-1/4} \geq 1/4 - 1/32 = 7/32 > 1/5$. Every counted label has a selected quotient with at least $A - 1 \geq k$ old agreements and one new agreement. A degree- $< k$ direction polynomial matching this full support would vanish on the old agreements, hence identically, but would equal one on the new agreement. This is impossible. The dimension and agreement threshold remain k and A . $\square$

This elementary consequence rules out a universal linear capacity MCA bound in the large-characteristic field class of [43]. Its ambient field is a quadratic extension, not a prime field, and its agreement lies below the first-order regime. It therefore neither proves tightness of the first-order quadratic upper bound nor resolves the prime-field superlinear-exception target. It concerns full agreement-set MCA; absence of ordinary correlated agreement is not asserted. Exact extension-field replays at $p = 17, 41$ are recorded in `research/prime_field_tightness/quadratic_extension_verification.json`.

Finite examples. At $p = 2^{127} - 1$, $n = 64$, $k = 32$, $s = 2$, Lemma 2.1 gives

$$L > 2^{35}.$$

The exact moment-class enumeration in Section 7.1 strengthens this to $L \geq 5552914238035 > 2^{42}$. Here $\eta = 1/32$, so the choice $c_1 = c_2 = 1$ in (2) gives only 2^{32} ; the improved ratio exceeds 1292. The exact inequality $p^2 > 2^{64}$ verifies the sufficient Elias condition. At $p = 2^{521} - 1$, $n = 4096$, $k = 2048$, $s = 8$, the same formula gives $L > 2^{3600}$, compared with 2^{512} . The exact inequality $p^8 > 2^{4096}$ verifies Elias. At $c_1 = 1$, this example forces $c_2 > 7$. A larger example over the same prime, using Gram concentration, forces $c_2 > 27.31865$. Already over $p = 2^{31} - 1$, an interval-domain example forces $c_2 > 2.00475$ at rate one-half. Appendix D.1 records the exact parameters and counts. As a small check of the algebra, among the 12,870 eight-subsets of $\{0, \dots, 15\}$, the binomial-moment class $(56, 252)$ has exactly 22 members. Over $p = 2^{31} - 1$, they give 22 distinct degree- < 6 polynomials, each agreeing with the common word at exactly eight positions.

2.3 Sharper finite bounds from conditional moments

The product of moment ranges in Lemma 2.1 counts many combinations without using how the moments depend on each other. For two moments, one can first fix the sum of the selected integers, then bound the spread of the second binomial moment within that class. Both steps can be computed exactly without enumerating the subsets.

Proposition 2.8 (Conditional range and dispersion bounds). *Let $3 \leq t < n$, $k = t - 2$, and $p > n$ be prime. For an integer q , let $\mathcal{A}_q$ be the family of t -subsets of $\{0, \dots, n - 1\}$ with sum q , and put*

$$C_q = |\mathcal{A}_q|, \quad Y_A = \sum_{a \in A} \binom{a}{2}.$$

For $C_q > 0$, let l_q and u_q be the minimum and maximum of Y_A in this family. There is a received word for $\text{RS}(\mathbb{F}_p, \{0, \dots, n - 1\}, k)$ whose list at radius $1 - t/n$ has size at least

$$\left\lceil \frac{C_q}{u_q - l_q + 1} \right\rceil. \quad (8)$$

More generally, choose an integer $c \in [l_q, u_q]$ and an integer $r \geq 0$. Define

$$D_q(c) = \sum_{A \in \mathcal{A}_q} (Y_A - c)^2, \quad W_q(c, r) = \min(u_q, c + r) - \max(l_q, c - r) + 1.$$

Then a possibly different received word has a list of size at least

$$\left\lceil \frac{\max\{0, C_q - \lfloor D_q(c)/(r + 1)^2 \rfloor\}}{W_q(c, r)} \right\rceil. \quad (9)$$

Every codeword produced by these bounds agrees in exactly t positions.

Proof. Pigeonhole over the integer values of Y_A proves (8). For the second bound, every subset with $|Y_A - c| > r$ contributes at least $(r + 1)^2$ to $D_q(c)$. Thus at most $\lfloor D_q(c)/(r + 1)^2 \rfloor$ subsets lie outside the specified interval. The remaining subsets occupy at most $W_q(c, r)$ integer values, proving (9) by pigeonhole. In either case, a class shares its first two binomial moments. The polynomial construction in Lemma 2.1 applies. $\square$

Exact computation. Process the integers $0, \dots, n - 1$ in order. For each subset size and first-moment sum, store its count, the minimum and maximum of Y , and the two totals $S = \sum Y$ and $Q = \sum Y^2$. On adding the integer j , put $b = \binom{j}{2}$. A predecessor record with count C contributes

$$C, \quad S + bC, \quad Q + 2bS + b^2C$$

to the new count and totals; its minimum and maximum each increase by b . Merge the contributions from taking and omitting j . The desired squared deviation is then $D_q(c) = Q - 2cS + c^2C$. There are $O(nt^2)$ possible size–sum states and n stages, giving $O(n^2t^2)$ integer operations and $O(nt^2)$ records. All stored integers have $O(n + \log n)$ bits. No assumption that the second moment fills its entire range is needed.

A stronger length-64 certificate. For $n = 64$ and $t = 34$, conditioning on $q = 1061$ gives $C_q = 8558724088132137$, $l_q = 17784$, and $u_q = 25714$. The range bound alone gives 1079148163931 codewords. The dispersion bound is stronger. For $q = 1071$, take $c = 22138$ and $r = 1076$. The exact counts are

$$\begin{aligned} C_q &= 8634334451040106, \\ D_q(c) &= 3337241714848691877182, \\ C_q - \left\lfloor \frac{D_q(c)}{1077^2} \right\rfloor &= 5757225839350346, \quad W_q(c, r) = 2153. \end{aligned}$$

Consequently,

$$\ell_C(15/32) \geq \left\lceil \frac{5757225839350346}{2153} \right\rceil = 2674048230075 > 2^{41}, \quad C = \text{RS}(\mathbb{F}_p, \{0, \dots, 63\}, 32). \quad (10)$$

This holds for every prime $p > 64$. At $p = 2^{127} - 1$, the original certificate $p^2 > 2^{64}$ keeps the radius strictly below Elias. The dispersion bound exceeds the product-range bound by a factor of 53, and exceeds $2^{1/\eta} = 2^{32}$ by a factor of 622. These are certified list lower bounds; neither is an exact maximum list size or an enumeration of the large moment class.

The script `checks/verify_conditional_moments.py` checks all counts, extrema, and moment totals against exhaustive enumeration for 66 smaller parameter pairs, covering 32,180 subsets. It independently verifies both displayed conditional counts using the coefficient identity

$$[y^t] \prod_{j=0}^{n-1} (1 + yx^j) = x^{\binom{t}{2}} \prod_{j=1}^t \frac{1 - x^{n-t+j}}{1 - x^j}.$$

Summing the conditional moment totals also agrees with the direct one-element and two-element inclusion counts over all t -subsets.

Signed polynomial weights. The interval argument can be sharpened further: instead of charging every retained moment value equally, assign a polynomial weight to each value. Negative weights outside favored ranges are allowed and improve the bound.

Proposition 2.9 (Polynomial-weight refinement). *Use the notation of Proposition 2.8, fix an integer c , and put*

$$M_j = \sum_{A \in \mathcal{A}_q} (Y_A - c)^j.$$

For any polynomial $P(Z) = \sum_{j=0}^d b_j Z^j$ with rational coefficients, define

$$S_P = \sum_{y=l_q}^{u_q} \max\{P(y - c), 0\}.$$

If $S_P > 0$, there is a received word whose list at radius $1 - t/n$ has size at least

$$\max \left\{ 0, \left\lceil \frac{\sum_{j=0}^d b_j M_j}{S_P} \right\rceil \right\}. \quad (11)$$

Each resulting codeword agrees in exactly t positions.

Proof. Let N_y count the members of $\mathcal{A}_q$ with second moment y , and let $L = \max_y N_y$. Since $0 \leq N_y \leq L$,

$$\sum_j b_j M_j = \sum_y N_y P(y - c) \leq L \sum_y \max\{P(y - c), 0\} = L S_P.$$

The equal-moment construction then gives the asserted codeword list. No positivity assumption on P over the entire moment range is needed. $\square$

Taking $P(Z) = h^2 - Z^2$ gives the particularly simple certificate

$$L \geq \left\lceil \frac{C_q h^2 - D_q(c)}{h(4h^2 - 1)/3} \right\rceil \quad \text{when } C_q h^2 > D_q(c).$$

Here the denominator sums positive weights over all integers, so it remains a valid upper bound if the actual moment range is shorter. The preceding $q = 1071$, $c = 22138$ data with $h = 1077$ already give $L \geq 4009211058641$ without computing any further moments.

An exact length-64 certificate. Continue with $n = 64$, $t = 34$, $q = 1071$, and $c = 22138$. The exact second-moment range is $[17969, 26129]$. Define the monic degree-20 polynomial $P(Z) = \prod_{r \in R} (Z - r)$, where

$$R = \{-4169, -2688, -2687, -2062, -2059, -1507, -1472, \\ -1038, -868, -643, 647, 870, 1039, 1464, 1498, \\ 2035, 2038, 2637, 2638, 3991\}.$$

Exact evaluation of (11) yields

$$\ell_C(15/32) \geq 5074503250115 > 2^{42}, \quad C = \text{RS}(\mathbb{F}_p, \{0, \dots, 63\}, 32), \quad p > 64. \quad (12)$$

This exceeds the original product-range certificate by a factor of 101, the preceding dispersion certificate by a factor of 1.89, and 2^{32} by a factor of 1181. At $p = 2^{127} - 1$, the same below-Elias certificate applies. These remain list lower bounds, not exact maximum list sizes.

For reproducibility, extend the count recurrence to raw moment totals $T_r = \sum Y_A^r$ for $0 \leq r \leq 20$. Adding an integer a with $b = \binom{a}{2}$ contributes

$$T'_r = \sum_{j=0}^r \binom{r}{j} b^{r-j} T_j, \quad M_r = \sum_{j=0}^r \binom{r}{j} (-c)^{r-j} T_j.$$

The positive-weight denominator is an exact sum over 8161 integers. The files in `moment_certificates/` record all raw and centered moments, polynomial coefficients, numerator, and denominator. The polynomial was discovered numerically, then replaced by the displayed integer-root polynomial. Its proof and certificate verification use only integer arithmetic; numerical optimizer tolerances play no role. The checker compares every moment through degree 20 with exhaustive small-subset enumeration and independently checks the length-64 moments using complementary 30-subsets.

A degree-40 refinement. At the same n, t, q, c , take the monic polynomial with roots

$$R_{40} = \{-4169, -3591, -3590, -3214, -3213, -2844, -2843, -2475, -2474, -2107, \\ -2106, -1742, -1735, -1391, -1358, -1065, -968, -765, -565, -476, \\ 482, 571, 769, 969, 1065, 1352, 1385, 1721, 1729, 2081, \\ 2082, 2436, 2437, 2787, 2788, 3135, 3136, 3484, 3485, 4184\}.$$

The exact numerator and positive-weight denominator in (11) give the stronger certificate

$$\ell_C(15/32) \geq 5285900426578, \quad C = \text{RS}(\mathbb{F}_p, \{0, \dots, 63\}, 32), \quad p > 64. \quad (13)$$

The radius and below-Elias qualification are unchanged. This improves (12) by more than 4.16% and exceeds the original product-range certificate by a factor of 105. The files in `research/finite_weights/` record the integer roots, expanded coefficients, exact moments, numerator, and denominator. Independent checks compare all moments through degree 40 with the complementary-subset identity and exhaustive smaller examples; the support extrema are recomputed separately. Earlier degree-24 and 28 certificates are retained for comparison. Numerical optimization discovers candidate weights; all claimed lower bounds and the final integer-root refinement use exact integer arithmetic.

The weighted inequality applies to any finite signature map, but its application to a prescribed evaluation domain requires statistics for that domain's coefficient signatures. The interval moments above do not by themselves supply such statistics for a multiplicative NTT subgroup.

Why an exponent substitution does not transfer the certificate. On a multiplicative subgroup generated by g , equal integer moments of exponent indices need not imply equal coefficient signatures. For $n \geq 2^{m+1}$, the positive and negative coefficient supports of $Q(X) = \prod_{j=0}^m (1 - X^{2^j})$ have equal integer moments through degree m , because Q has a root of multiplicity $m+1$ at 1. But their first field moments differ by $Q(g) \neq 0$ when g has order n . Thus replacing each index a by g^a loses the required coefficient equality.

Nor does the bare field-signature pigeonhole estimate certify a large list below Elias. For $p > n$, $t = k + m < n$, and $k, m \geq 1$, it gives only $\lceil \binom{n}{t} / p^m \rceil$. Writing $\theta = 1 - t/n$, the strict Elias inequality $H_p(\theta) < 1 - k/n$ implies

$$\binom{n}{t} / p^m < (p/(p-1))^{n\theta} < e,$$

so this guarantee is at most three. These observations limit the two transfer arguments, rather than the true list size on a subgroup. The exact checks in `research/structured_domains/` accompany these elementary obstructions and a separate complete-fiber transfer lemma.

What full polynomial fibers can accomplish. The following elementary composition argument explains why replacing a power map by another full balanced polynomial map gives no additional fiber partitions on a multiplicative subgroup.

Lemma 2.10 (Full polynomial fibers on a subgroup). *Let F be a field and $n \geq 1$. Suppose $D = \mu_n \subset F$ has n distinct points, and let $R \in F[X]$ have degree $B \geq 1$. If every fiber of $R|_D$ has exactly B points, then $B \mid n$ and $R(X) = aX^B + b$, with $a \neq 0$. Conversely, every such polynomial has B points per fiber when $B \mid n$.*

Proof. Put $M = |R(D)| = n/B$, $a = \text{lc}(R)$, and $P(Y) = \prod_{y \in R(D)} (Y - y)$. Comparing roots and leading coefficients gives

$$P(R(X)) = a^M (X^n - 1).$$

Write $b = R(0)$, $T = R - b$, and $Q(Y) = P(Y + b)$. If T has a nonzero coefficient c_j below degree B , choose its largest positive index j . The coefficient of X^{n-B+j} in T^M is $Ma^{M-1}c_j \neq 0$, since the distinct-root hypothesis implies $\text{char } F \nmid n$. Every lower power of T in $Q(T)$ has degree at most $n - B$. This contradicts $Q(T) = a^M (X^n - 1)$, so $T = aX^B$. The converse follows from the kernel size of $x \mapsto x^B$ on μ_n . $\square$

Rational maps permit different fibers. The lemma exhausts full polynomial fibers, but rational maps admit other partitions. If $2d \mid n$, choose a nonsquare $c \in \mu_{n/d}$. The rational map

$$R(x) = x^d + c/x^d$$

has degree $B = 2d$ and no pole on μ_n . After setting $u = x^d$, its fibers pair u with c/u : equality of two outputs is equivalent to $(u - v)(uv - c) = 0$. These pairs have no fixed points, so every fiber has exactly $2d$ elements.

Rational substitution also preserves a useful polynomial degree bound. In general, write $R = A/V$ in reduced form with $\max\{\deg A, \deg V\} = B$. Suppose it has no domain poles and exactly B points over each of its N image labels. For $1 \leq k \leq N$, replace a polynomial f of degree less than k by

$$V^{k-1}f(A/V),$$

whose degree is at most $B(k-1)$. Apply the same nonzero coordinate scaling to the composed received word. Each agreement repeats exactly B times, preserving relative Hamming distance and distinctness. The images lie in a Reed–Solomon code of dimension $B(k-1) + 1$. A large list still requires a separate certificate on the image domain.

Corollary 2.11 (Balanced fibers of Galois rational maps). *Let F be a finite field of characteristic $p > n \geq 3$, with $\mu_n \subset F$. Suppose $R \in F(X)$ has degree $B \geq 1$, has no poles on μ_n , and has exactly B points in every fiber of $R|_{\mu_n}$. If $\overline{F}(X)/\overline{F}(R)$ is Galois, then $F(X)/F(R)$ is Galois and, for some $h \in \text{PGL}_2(F)$, either*

$$R(X) = h(X^B), \quad B \mid n,$$

or

$$R(X) = h(X^d + c/X^d), \quad B = 2d \mid n,$$

where c is a nonsquare in $\mu_{n/d}$. Here h has no pole on the corresponding image of μ_n .

Proof. Each domain fiber already contains the full B distinct geometric preimages of its image. The geometric deck group G , of order B , therefore preserves μ_n and acts freely there.

A Möbius map $(aX + b)/(cX + d)$ preserving μ_n satisfies $(aX + b)^n - (cX + d)^n = \lambda(X^n - 1)$ with $\lambda \neq 0$; its intermediate coefficients give $a^i b^{n-i} = c^i d^{n-i}$ for $1 \leq i < n$. Comparing $i = 1, 2$ rules out four nonzero entries, and the remaining cases give precisely $X \mapsto \zeta X$ and $X \mapsto \zeta/X$, with $\zeta \in \mu_n$. These maps are defined over F , proving arithmetic descent.

If G consists of rotations, its invariant field is $F(X^B)$. Otherwise its rotations have order d , and a reflection $X \mapsto \alpha/X$ gives the invariant $X^d + \alpha^d/X^d$, of degree $2d = |G|$. Writing $c = \alpha^d$, freeness is equivalent to the absence of $u \in \mu_{n/d}$ with $u^2 = c$: such a u lifts to a fixed point of some reflection. The displayed invariant generates $F(X)^G$; any other generator R differs by an output Möbius map over F . $\square$

Thus this Galois subclass reduces to power and twisted-inversion quotients. Full balanced fibers alone do not imply the Galois hypothesis: the explicit cubic over $\mathbb{F}_{13}$ on μ_6 in `research/gram_norm/BALANCED_FIBER_REVIEW.md` has four simple critical points and is not a Galois cover.

For the large-index congruence classes below, full balance forces the Galois hypothesis. In fact, a small number of uncovered domain points is allowed. This includes the field and subgroup in the pinned `better.codes` instance.

Theorem 2.12 (Nearly complete rational fibers at large Frobenius index). *Let F be a finite field of odd characteristic p , containing μ_n , and suppose*

$$p = \ell n + \epsilon, \quad \epsilon \in \{1, -1\}, \quad \ell \geq 6.$$

Let $R \in F(X)$ have degree $B \geq 2$ and no poles on μ_n . Suppose $T \subseteq \mu_n$ is a union of complete geometric fibers, each consisting of B distinct points, and $|\mu_n \setminus T| \leq c$. If $n > 6(B - 1 + c)$, then $\overline{F}(X)/\overline{F}(R)$ is Galois, and all its deck transformations have the form $X \mapsto \zeta X$ or $X \mapsto \zeta/X$, with $\zeta \in \mu_n$. Consequently $B \mid 2n$; when n is a power of two, $B \mid n$. If also $c < B/2$, then the deck group acts freely on μ_n , every fiber of $R|_{\mu_n}$ has B points, and $B \mid n$. In this case R has one of the two forms in Corollary 2.11.

The proof, in Appendix E, uses the four-term unit equation on each collision component and a Frobenius-index argument establishing the required linear independence. At $p = 2130706433$ and $n = 262144$, one has $\ell = 8128$ and $\epsilon = 1$. For full balance ($c = 0$), the theorem covers every integer $B \leq 43691$. Since full balance forces $B \mid n$, it covers every possible degree through 32768. The remaining degrees have at most four fibers and therefore at most sixteen unions of complete fibers. This classifies the small-degree fixed-map full-fiber partitions at these parameters; it gives neither a new numerical better.codes bound nor a restriction on constructions using partial fibers or varying maps. The extension also excludes degrees 513, 514, 516 with complete-fiber coverage defects at most 1, 4, 16, respectively: each satisfies both inequalities, but none divides 262144. These are exclusions of proposed packet maps, not attained improvements in agreement or score.

On this dyadic domain the divisibility conclusion does not need $c < B/2$. For any non-dyadic degree B , the complete-fiber count is therefore at most

$$\min \left\{ \left\lfloor \frac{262144}{B} \right\rfloor, 1 + \left\lfloor \frac{218452}{B} \right\rfloor \right\}.$$

Indeed, the uncovered count must satisfy $c \geq 43692 - B$, and the covered count is a multiple of B . At degrees 513, 514, 516, this forces at least 43606, 43180, 43360 uncovered points, respectively.

2.4 A uniform concentration bound for several moments

The improvement is not limited to a computed length-64 example. Concentrating the moment signatures near their means saves a polynomial factor for every fixed number of canceled coefficients.

Proposition 2.13 (Concentration of triangular moment signatures). *Let $1 \leq m < t < n$, $k = t - m$, and let $p > n$ be prime. Choose rational coefficients c_{ji} and functions*

$$b_j(a) = \binom{a}{j} + \sum_{i=1}^{j-1} c_{ji} \binom{a}{i}, \quad 1 \leq j \leq m, \quad 0 \leq a < n.$$

For a uniformly chosen t -subset A of $\{0, \dots, n-1\}$, write $B_j(A) = \sum_{a \in A} b_j(a)$ and define

$$\bar{b}_j = \frac{1}{n} \sum_{a=0}^{n-1} b_j(a), \quad V_j = \frac{t(n-t)}{n-1} \left(\frac{1}{n} \sum_{a=0}^{n-1} b_j(a)^2 - \bar{b}_j^2 \right), \quad h_j = \left\lceil \sqrt{(m+2)V_j} \right\rceil.$$

Then $C = \text{RS}(\mathbb{F}_p, \{0, \dots, n-1\}, k)$ satisfies

$$\ell_C(1 - t/n) \geq \left\lceil \frac{2 \binom{n}{t}}{(m+2) \prod_{j=1}^m (2h_j + 1)} \right\rceil. \quad (14)$$

The constructed codewords agree in exactly t positions.

Proof. The inclusion probabilities t/n and $t(t-1)/(n(n-1))$ give $\mathbb{E}B_j = t\bar{b}_j$ and $\text{Var}(B_j) = V_j$. Here $V_j > 0$: b_j is a nonconstant polynomial of degree $j < n$ on n distinct integers, and $0 < t < n$. Chebyshev's inequality and a union bound show that at least a $2/(m+2)$ fraction of all subsets satisfy $|B_j - t\bar{b}_j| \leq h_j$ for every j .

For the first coordinate there are at most $2h_1 + 1$ possible integer values. Given the earlier coordinates, triangular inversion fixes the earlier binomial moments, and B_j differs from the integer $\sum_{a \in A} \binom{a}{j}$ by a fixed rational number. Hence its interval contains at most $2h_j + 1$ possible values. There are at most $\prod_j (2h_j + 1)$ retained signatures. Pigeonhole and triangular inversion give a class sharing its first m integer binomial moments. Lemma 2.1 applies. Rational changes of coordinates are used only for counting; no reduction of their denominators modulo p is required. $\square$

For fixed m and $t/n \rightarrow \rho \in (0, 1)$, choosing all $c_{ji} = 0$ gives $V_j = \Theta_{\rho, m}(n^{2j+1})$ and therefore

$$\ell_C(1 - t/n) = \Omega_{\rho, m} \left(\frac{\binom{n}{t}}{n^{m(m+2)/2}} \right).$$

The product of complete moment ranges instead has order $n^{m(m+3)/2}$ in its denominator. Thus concentration improves that guarantee by a factor of order $n^{m/2}$. This comparison fixes m ; it does not replace the separate estimates when m grows with n .

Two moments with explicit constants. For $m = 2$, take $b_1(a) = a$ and $b_2(a) = \binom{a}{2} - (n-2)a/2$. These centered population functions are uncorrelated, and their subset-sum variances are

$$V_1 = \frac{t(n-t)(n+1)}{12}, \quad V_2 = \frac{t(n-t)(n+1)(n^2-4)}{720}.$$

Consequently (14) implies

$$\ell_C(1 - t/n) \geq \left(\frac{3\sqrt{60}}{8\rho(1-\rho)} + o(1) \right) \frac{\binom{n}{t}}{n^4} \quad \text{if } t/n \rightarrow \rho \in (0, 1).$$

At $n = 64, t = 34$, the exact variances are 5525 and 376805, and $h_1 = 149, h_2 = 1228$. This gives 1102772374154 codewords without the conditional dynamic program. The degree-20 certificate is stronger at that length. The script `moment_certificates/verify_concentration.py` checks the variance identities, triangular signature counts, and finite bounds against exhaustive small-subset enumeration.

2.5 Gram moments and an ellipsoid concentration bound

Orthogonalizing the binomial moments gives explicit variances, and a signed quadratic weight improves the elementary concentration constant. Let $1 \leq m < t < n$. For $1 \leq j \leq m$, let b_j be the degree- j polynomial with leading coefficient $1/j!$ orthogonal to every lower-degree polynomial on $\{0, \dots, n-1\}$ with uniform weight. These are the classical Gram polynomials, with the normalization $b_j(x) = \binom{x}{j} + \sum_{i < j} c_{ji} \binom{x}{i}$. Their classical squared norms [34, eqs. (1.1)–(1.2)] and the sampling-without-replacement variance identity give, for a uniform t -subset A and $B_j(A) = \sum_{a \in A} b_j(a)$,

$$\mathbb{E}B_j = 0, \quad V_j := \text{Var}(B_j) = \frac{t(n-t)}{n-1} \frac{\prod_{i=1}^j (n^2 - i^2)}{(2j+1) \binom{2j}{j}^2 (j!)^2}. \quad (15)$$

The coordinates B_j are pairwise uncorrelated. Put $\kappa_m = \pi^{m/2}/\Gamma(m/2 + 1)$, the volume of the unit ball. The continuous inequality below is the classical maximum-density/covariance bound [37, Proposition III.1 and proof of Corollary III.2]; unit-cube smoothing gives the exact variance correction $1/12$. No log-concavity assumption is needed.

Proposition 2.14 (Finite smoothed ellipsoid bound). *For every prime $p > n$, the code $C = \text{RS}(\mathbb{F}_p, \{0, \dots, n-1\}, t-m)$ satisfies*

$$\ell_C(1 - t/n) \geq \left\lceil \frac{\binom{n}{t}}{\kappa_m(m+2)^{m/2} \prod_{j=1}^m \sqrt{V_j + 1/12}} \right\rceil. \quad (16)$$

The constructed codewords agree in exactly t coordinates.

Proof. Write $N = \binom{n}{t}$, and let L be the largest class sharing its vector $B = (B_1, \dots, B_m)$. Because the zeroth moment is fixed, these vectors lie in an affine lattice $b_0 + T\mathbb{Z}^m$, with T lower triangular and diagonal entries one. The translates of $[-1/2, 1/2]^m$ centered at its lattice points tile $\mathbb{R}^m$: conditional rounding determines the first integer coordinate, then the second, and so on, uniquely.

Add to B an independent uniform vector U from this unit cube. The resulting vector $Y = B + U$ has density at most L/N , mean zero, and coordinate variances $W_j = V_j + 1/12$. Thus $Z_j = Y_j/\sqrt{W_j}$ satisfies $\mathbb{E}\|Z\|^2 = m$ and has density at most $H = (L/N) \prod_j \sqrt{W_j}$. The signed quadratic weight gives

$$2 = \mathbb{E}(m + 2 - \|Z\|^2) \leq H \int_{\mathbb{R}^m} (m + 2 - \|z\|^2)_+ dz = 2H\kappa_m(m+2)^{m/2}.$$

Rearranging bounds L as in (16). Triangular inversion makes these classes exactly the classes sharing their first m integer binomial moments. The construction in Lemma 2.1 converts a class into the stated list. The rational changes of coordinates are used only for real-valued counting, so their denominators impose no further condition on p . $\square$

Asymptotics for a fixed number of moments. For fixed m and $t/n \rightarrow \rho \in (0, 1)$, (16) gives

$$\ell_C(1 - t/n) \geq \left(\frac{\prod_{j=1}^m j! \binom{2j}{j} \sqrt{2j+1}}{\kappa_m(m+2)^{m/2} [\rho(1-\rho)]^{m/2}} + o(1) \right) \frac{\binom{n}{t}}{n^{m(m+2)/2}}. \quad (17)$$

For the same variances, the leading constant is larger than that from the optimized rectangular Chebyshev bound by $(m+2)2^{m-1}/\kappa_m$. This comparison concerns elementary concentration; it does not assert improvement over conditional counting.

Two moments. Here

$$V_1 = \frac{t(n-t)(n+1)}{12}, \quad V_2 = \frac{t(n-t)(n+1)(n^2-4)}{720},$$

$$\ell_C(1 - t/n) \geq \left\lceil \frac{\binom{n}{t}}{4\pi\sqrt{(V_1 + 1/12)(V_2 + 1/12)}} \right\rceil.$$

The asymptotic leading constant is $3\sqrt{60}/(\pi\rho(1-\rho))$, a factor $8/\pi$ larger than the rectangular constant. At $n = 64, t = 34$, the variances are 5525 and 376805. Rational upper bounds for the denominator certify 2,825,885,263,137 codewords from the smoothed bound. Direct exact summation of the positive part of the quadratic weight $4 - B_1^2/V_1 - B_2^2/V_2$ on the full moment lattice improves this to 2,825,905,919,527. Both exceed the elementary rectangular certificate 1,102,772,374,154; the existing degree-20 conditional certificate 5,074,503,250,115 remains stronger.

A growing number of moments. The gain also holds when $\log n \ll m \ll n$, provided t/n stays in a compact subinterval of $(0, 1)$. Let

$$D_E = \kappa_m(m+2)^{m/2} \prod_{j=1}^m \sqrt{V_j + 1/12}, \quad D_B = \frac{m+2}{2} \prod_{j=1}^m (2\lceil \sqrt{(m+2)V_j} \rceil + 1)$$

be the smoothed ellipsoid and rectangular Gram denominators, respectively. For either $D_\star \in \{D_E, D_B\}$, natural logarithms give the uniform estimate

$$\log D_\star = \frac{m^2}{2} \log(n/m) + (3/4 - \log 2)m^2 + O(m \log n + m^4/n^2).$$

Indeed, Stirling's bounds evaluate $\prod_{j=1}^m \sqrt{2j+1}(2j)!/j!$; the remaining grid correction is $O(m^4/n^2)$. Since $m = o(n)$, eventually $17m^2 \leq n^2 + 4$. The recurrence $V_j/V_{j-1} = (n^2 - j^2)/(4(2j-1)(2j+1))$ then gives $V_j \geq V_1 = \Theta(n^3)$, controlling both the smoothing and rounding errors. The complete-range denominator has the same leading expression with $3/4$ in place of $3/4 - \log 2$ and error $O(m \log n + m^3/n)$. Both errors are $o(m^2)$ in the stated regime. Thus either Gram guarantee improves the unrounded range guarantee by $2^{m^2+o(m^2)}$. At $m \sim c\sqrt{n/\log_2 n}$, this saves $c^2 n / \log_2 n + o(n/\log n)$ bits in the guaranteed logarithmic list size, while preserving the leading condition $c^2 < 4H_2(\rho)$. No central limit theorem in growing dimension is required. For example, exact arithmetic gives the following lower bounds on $\lfloor \log_2 L \rfloor$ from the complete ranges and the *rectangular* Gram denominator D_B , at rate $1/2$ and $t = n/2 + m$:

n	m	Complete ranges	Gram concentration
128	4	42	60
512	8	183	253
4096	20	1915	2333
16384	40	7241	8869

The finite certificates and uniform estimates are recorded in `research/growing_m/`. These comparisons concern interval domains; field-size conditions for a below-Elias radius must still be imposed.

2.6 A finite certificate without conditioning on the first moment

Proposition 2.15 (An exact two-moment interval certificate). *Among the 34-element subsets $S \subseteq \{0, \dots, 63\}$, some class with fixed values of $\sum_{a \in S} a$ and $\sum_{a \in S} \binom{a}{2}$ contains at least*

$$5,133,798,314,667$$

subsets. Consequently, for every prime $p > 64$, $C = \text{RS}(\mathbb{F}_p, \{0, \dots, 63\}, 32)$ satisfies $\ell_C(30/64) \geq 5,133,798,314,667$, with each constructed codeword agreeing in exactly 34 coordinates.

Proof. Taking complements preserves class sizes, so work with 30-subsets. Put

$$X_a = 2a - 63, \quad Y_a = \frac{(2a - 63)^2 - 1365}{4}, \quad x = \frac{1}{2} \sum_{a \in S} X_a, \quad y = \frac{1}{2} \sum_{a \in S} Y_a.$$

The pair $(x, y) \in \mathbb{Z}^2$ is an invertible affine transformation of the two moment sums. Set

$$T = 341x^2 + 5y^2, \quad D = 1,884,025, \quad Q = T/D, \quad A = \frac{851547}{580691}.$$

The ancillary file `radial_certificates/certificate_degree12.json` specifies a rational polynomial h of degree 12. Use the degree-25 radial weight

$$w(Q) = (A - Q)h(Q/A)^2.$$

It is nonpositive outside $Q < A$. If L denotes the largest class, then

$$\sum_{|S|=30} w(Q(S)) \leq L \sum_{(x,y) \in \mathbb{Z}^2} \max\{w(Q(x,y)), 0\}.$$

The numerator is evaluated exactly from the moments through degree 25 of T , obtained by an integer dynamic program on the 32 population pairs $(u, v), (-u, v)$. The denominator is an exact sum over the 210,135 lattice points with $T \leq 2,762,804$. After clearing all coefficient denominators, the ancillary file records the integer numerator and denominator, whose ratio has ceiling 5,133,798,314,667. The standard-library script `radial_certificates/verify_publication.py` recomputes all 676 mixed moments, checks the saved polynomial and ratio, and independently recomputes the denominator by summing expanded monomials over the full signed lattice. Numerical optimization was used only to discover the rational polynomial; it is not an input to the verification. Lemma 2.1 converts the common-moment class into the stated Reed–Solomon list. This improves the degree-20 conditional certificate by approximately 1.17%. The higher-degree conditional certificates above give larger bounds at these parameters. The radial argument offers a separate way to use both moment coordinates at once; no optimality is claimed. $\square$

2.7 A stronger asymptotic constant for fixed dimension

Fix $m \geq 1$, and suppose only that $\min(t, n - t) \rightarrow \infty$. Write V_j for the Gram subset-sum variances above, and let L_n be the largest class of t -subsets sharing their first m binomial moments. Then

$$\liminf_{n \rightarrow \infty} \frac{L_n \prod_{j=1}^m \sqrt{V_j}}{\binom{n}{t}} \geq (2\pi)^{-m/2}. \quad (18)$$

This follows from a finite-population central limit theorem and ordinary pigeonhole counting; a local central limit theorem is unnecessary.

Proof. Let G_j be the mean-zero Gram polynomial of degree j and leading coefficient $1/j!$. For a uniformly chosen t -subset A , put $Z_j = \sum_{a \in A} G_j(a)/\sqrt{V_j}$. Orthogonality gives $\mathbb{E}Z = 0$ and $\text{Cov}(Z) = I_m$. For any fixed nonzero $u \in \mathbb{R}^m$, the population values $a_x = \sum_j u_j G_j(x)/\sqrt{V_j}$ have mean zero and finite-population variance

$$v_n = \frac{1}{n-1} \sum_x a_x^2 = \frac{n \|u\|^2}{t(n-t)}.$$

For fixed m , $\max_x |G_j(x)| = O_m(n^j)$ and the population variance of G_j is $\Theta_m(n^{2j})$. Cauchy–Schwarz therefore gives

$$\frac{\max_x a_x^2}{\min(t, n-t)v_n} = O_{u,m}(1/\min(t, n-t)) \rightarrow 0.$$

Hájek’s finite-population central limit theorem [36], in the sufficient form of [35, Theorem 1], followed by the Cramér–Wold criterion, therefore gives $Z \Rightarrow N(0, I_m)$.

The Gram signatures belong to an affine unit lower triangular image of $\mathbb{Z}^m$. Its centered axis-aligned unit cubes tile $\mathbb{R}^m$, as in the finite bound above. After dividing coordinate j by $\sqrt{V_j}$,

each cell has volume $1/\prod_j \sqrt{V_j}$ and lies within $\epsilon_n = \frac{1}{2}\sqrt{\sum_j V_j^{-1}} = o(1)$ of its center. Consequently a ball of radius R contains at most

$$\kappa_m(R + \epsilon_n)^m \prod_j \sqrt{V_j}$$

possible signatures, where κ_m is unit-ball volume. For fixed $R > 0$, pigeonhole counting and convergence in distribution give

$$\liminf_n \frac{L_n \prod_j \sqrt{V_j}}{\binom{n}{t}} \geq \frac{\Pr(\|N(0, I_m)\| \leq R)}{\kappa_m R^m}.$$

Letting $R \downarrow 0$ after this limit proves (18) by continuity of the Gaussian density. $\square$

If additionally $t/n \rightarrow \rho \in (0, 1)$, then with $k = t - m$ and $p > n$ the moment-class construction gives

$$\ell_{\text{RS}(\mathbb{F}_p, \{0, \dots, n-1\}, k)}(1 - t/n) \geq \left(\frac{\prod_{j=1}^m [\sqrt{2j+1} \binom{2j}{j} j!]}{(2\pi\rho(1-\rho))^{m/2}} - o(1) \right) \frac{\binom{n}{t}}{n^{m(m+2)/2}}.$$

The number of moments is fixed here. This statement supplies neither a finite-length error bound nor the mass of a prescribed individual signature. It is an application of classical limit theory, with no priority claim.

3 From large lists to nearby codewords on a line

List-decoding asks about one received word. Line-decodability asks how nearby codewords can vary along an affine family of received words $f(z) = f_0 + zf_1$. Its selected-witness formulation asks whether many chosen nearby codewords must themselves follow one affine rule in z .

Definition 3.1 (Selected-witness line-decodability). Let $C \subseteq F^n$ be linear. Say that C is line-decodable at radius θ , threshold a , and concurrency b if the following holds. For every affine word line $f(z) = f_0 + zf_1$, every set $S \subset F$ with $|S| > a$, and every choice of codewords $p_z \in C$ within distance θ of $f(z)$, some $P_0, P_1 \in C$ satisfy $p_z = P_0 + zP_1$ for more than b parameters $z \in S$.

The word “selected” matters: the property must hold for every choice of the nearby witnesses. Also, equality uses the *same* parameter z on both sides. We call $\{(z, P_0 + zP_1) : z \in F\}$ a parametrized codeword line. Separately, (f_0, f_1) has *correlated agreement* on t coordinates if two codewords agree with f_0, f_1 simultaneously on one set of t coordinates.

We first turn a large list into a family that has many nearby codewords but no such correlated agreement. Division by $X - \tau$ lowers each polynomial’s degree by one. The value at the chosen pole τ determines which parameter on the word line receives that polynomial.

Lemma 3.2 (A rational word line). Let $D \subset \mathbb{F}_p$ have size n , and let $1 \leq d < t < n$. Suppose N distinct polynomials R_i of degree at most d agree with a common word w on at least t coordinates. Write $Q = p - n > 0$. There is a pole $\tau \in \mathbb{F}_p \setminus D$ for which

$$f_0(x) = \frac{w(x)}{x - \tau}, \quad f_1(x) = \frac{1}{x - \tau}$$

has a set S of distinct parameters, with a selected nearby codeword $P_z \in \text{RS}(\mathbb{F}_p, D, d)$ at each parameter, such that

$$|S| \geq \frac{NQ}{Q + d(N - 1)}. \quad (19)$$

Each selected codeword has at least t agreements. No pair of codewords agrees jointly with (f_0, f_1) on more than d coordinates, and every parametrized codeword line contains at most $\lfloor (n-d)/(t-d) \rfloor < n$ selected pairs. If $N \geq Q$, then $|S| \geq Q/(d+1)$. If $p > n + d\binom{N}{2}$, all N parameters can be made distinct.

Proof. For an admissible pole, assign

$$z_i = -R_i(\tau), \quad P_i(X) = \frac{R_i(X) - R_i(\tau)}{X - \tau}.$$

Then $\deg P_i < d$, and $f_0 + z_i f_1 - P_i = (w - R_i)/(X - \tau)$ on D , preserving the agreement set.

Two parameters collide only at a root of the nonzero polynomial $R_i - R_j$, which has degree at most d . Averaging over the Q admissible poles gives one with at most $dN(N-1)/(2Q)$ colliding pairs. If the multiplicities of its distinct parameters are $m_1, \dots, m_v$, then

$$\sum_j m_j = N, \quad \sum_j m_j^2 \leq N + \frac{dN(N-1)}{Q}.$$

Cauchy–Schwarz gives (19). Select one witness per parameter. If $N \geq Q$, the denominator is at most $(d+1)N$. Under the stronger field condition, the union of D and all pairwise root sets has size less than p , so a collision-free pole exists.

Any agreement of f_1 with a degree- $< d$ polynomial P_1 is a root of $(X - \tau)P_1(X) - 1$. This polynomial has degree at most d and is nonzero, as its value at τ is -1 . Thus even f_1 alone has at most d agreements with any codeword. Finally fix a parametrized codeword line $P_0 + zP_1$. Let $b \leq d$ be the number of coordinates where both coefficient words agree. At every other coordinate the two affine expressions agree for at most one parameter. If v selected witnesses lie on this line, counting agreements gives $v(t-b) \leq n-b$. The ratio $(n-b)/(t-b)$ increases with b , so $v \leq (n-d)/(t-d)$. $\square$

Theorem 3.3 (No uniform finite $n \exp(O_\rho(1/\eta))$ line threshold). *Fix rational $\rho \in (0, 1)$ and finite $c_1, c_2 > 0$. There are arbitrarily large prime-field Reed–Solomon codes of exact rate ρ on consecutive-integer domains that fail selected-witness line-decodability at concurrency n and threshold*

$$a = c_1 2^{c_2 H_2(\rho)/\eta} n, \quad \theta = 1 - \rho - \eta < r_E(\rho).$$

One can take $\eta = \Theta_\rho(1/\sqrt{n \log n})$ and $\log_2 p = \Theta_{\rho, c_2}(\sqrt{n \log n})$.

Proof. Let $d = \rho n$ be the target dimension. Choose $c^2 < 4H_2(\rho)$ and set $s = \lfloor c\sqrt{n/\log_2 n} \rfloor$, $t = d + s + 1$, and $\eta = (s+1)/n$. Apply Lemma 2.1 to the source code of dimension $d+1$. It supplies $N = 2^{\Omega_\rho(n)}$ degree- $\leq d$ polynomials with common agreement size t ; the one-dimensional change does not affect the estimate (3). Choose

$$A > \max\{H_2(\rho)/c, c_2 H_2(\rho)/c\}, \quad \log_2 p = A\sqrt{n \log_2 n} + O(1)$$

using Bertrand’s postulate. Eventually $N > p > n$. Lemma 3.2 gives at least $(p-n)/(d+1)$ selected parameters, whose logarithmic count is $A\sqrt{n \log_2 n} - O(\log n)$. The proposed threshold has logarithm

$$\log_2(c_1 n) + c_2 H_2(\rho) \frac{n}{s+1} = (c_2 H_2(\rho)/c + o(1)) \sqrt{n \log_2 n},$$

which is smaller. Every parametrized codeword line contains fewer than n selected witnesses. Finally $\eta \log_2 p \rightarrow cA > H_2(\rho)$ verifies Elias by the same entropy-continuity argument as above. $\square$

A finite line example. Take $p = 2^{127} - 1$, $n = 64$, target dimension $d = 32$, and agreement size $t = 34$. The first-moment sum 1071 has exactly

$$N = 8,634,334,451,040,106 > 2^{52}$$

34-subsets, as certified in Section 2.3. Their vanishing polynomials have two common leading terms, giving source polynomials of degree at most 32. The exact inequality $p > 64 + 32\binom{N}{2}$ makes all parameters distinct. Here $\eta = 1/32$, the threshold with $c_1 = c_2 = 1$ is 2^{38} , and each parametrized codeword line contains at most 16 selected witnesses. The radius is below Elias. In a smaller exhaustive check, $n = 12$, $d = 6$, $t = 8$, $p = 2^{31} - 1$ yield a moment class of size 33. The pole $\tau = 20$ gives 33 distinct parameters, and the maximum number of selected pairs on one parametrized codeword line is exactly 3, matching the lemma's bound.

3.1 Every large prime admits a short interval-domain counterexample

Holding the number of canceled moments fixed gives a useful alternative to Theorem 3.3. It strengthens the field quantifier and reduces the required block length to logarithmic in the prime.

Theorem 3.4 (Logarithmic-length prime-field lines). *Fix a rational $\rho \in (0, 1)$, write $H = H_2(\rho)$, and fix an integer $r \geq 2$. For every sufficiently large prime p , writing $b = \log_2 p$, there is an interval-domain Reed–Solomon code over $\mathbb{F}_p$ with dimension $k = \rho n$ and*

$$n = H^{-1}b + O_{\rho,r}(\log b), \quad t = k + r, \quad \eta = r/n,$$

and a received line with at least

$$\left(\frac{H}{\rho(1-\rho)} - o(1) \right) \frac{p}{\log_2 p} \quad (20)$$

distinct nearby labels at agreement threshold t . There is no correlated agreement on more than k coordinates, and every parametrized codeword line contains at most $\lfloor (n-k)/r \rfloor$ selected witnesses. The radius is strictly below the characteristic-based Elias radius.

In particular, given any finite $c_1, c_2 > 0$, taking an integer $r > \max\{1, c_2\}$ violates the finite line threshold $c_1 n 2^{c_2 H/\eta}$ over every sufficiently large prime field.

Proof. Put $a = (r-1)(r+2)/2$ and choose

$$n = \frac{b + (a+2)\log_2 b}{H} + O_\rho(1)$$

by rounding upward to a multiple of the denominator of ρ . Apply Lemma 2.1 with source dimension $k+1$ and $s = r-1$. Its moment ranges satisfy $\prod_{j=1}^{r-1} R_j = O_r(n^a)$. Moreover, $\binom{n}{t} \geq 2^{nH_2(t/n)}/(n+1)$ and $nH_2(t/n) = Hn + O_{\rho,r}(1)$. Thus a bank of N distinct polynomials of degree at most k , each with exactly t agreements, has

$$N \geq c_{\rho,r} \frac{2^{Hn}}{n^{a+1}} \geq c'_{\rho,r} p b.$$

In particular $N/p \rightarrow \infty$. Select exactly N members if the class is larger.

We sharpen the collision estimate in Lemma 3.2 by subtracting shared agreement roots. Let I_x be the support incidences, so $\sum_x I_x = Nt$. Across the $Q = p - n$ admissible poles, the total number of pair collisions is at most

$$k \binom{N}{2} - \sum_x \binom{I_x}{2} \leq \frac{N^2(k - t^2/n) + N(t - k)}{2}.$$

Indeed each difference already vanishes on its shared support, and $\sum_x I_x^2 \geq N^2 t^2/n$. Averaging and applying Cauchy–Schwarz to the value multiplicities at one pole gives at least

$$\frac{NQ}{Q + N(k - t^2/n) + (t - k)}$$

distinct labels. Since $k - t^2/n = \rho(1 - \rho)n - 2\rho r - r^2/n$ and $N/p \rightarrow \infty$, this is (20). Lemma 3.2 also gives the correlated-agreement and concurrency assertions.

Finally, $\eta \log_2 p \rightarrow rH > H$, whereas $H_2(1 - \rho - \eta) \rightarrow H$, proving strict Elias admissibility. The logarithm of the proposed threshold is $(c_2/r)b + O_{\rho,r,c_1,c_2}(\log b)$, while the constructed count has logarithm $b - \log_2 b + O_\rho(1)$. Their difference tends to infinity when $r > c_2$. $\square$

The gap here is $\Theta_{\rho,r}(1/\log p)$; the theorem does not assert a superlinear exception count at a fixed positive gap. It uses the same moment classes and rational division as the earlier results, with a different parameter choice.

Finite examples entirely over prime fields. The finite smoothed ellipsoid bound of Proposition 2.14, followed by the same shared-support collision count, gives the following existence certificates. The final column is a strict lower bound on $\log_2(J/(n2^{H_2(k/n)/\eta}))$, with $\eta = (t - k)/n$.

Prime	n	k	t	Nearby labels J	Excess bits
$2^{61} - 1$	136	25	29	$> 8.9248 \cdot 10^{16}$	25.81662
$2^{127} - 1$	257	72	78	$> 2.7727 \cdot 10^{36}$	76.40420

Both examples lie strictly below the characteristic-based Elias radius and have no correlated agreement at their displayed thresholds. All words, witnesses, and challenges are in the stated prime field. The count certifies existence of a suitable moment class and pole; it does not enumerate those large objects. These are ratios of coding counts to the finite prescription with $c_1 = c_2 = 1$.

For exact reproducibility, let $\mathcal{P}(u, v)$ be the minimum number of colliding pairs for u objects in v bins, as in Section C.2. If the concentration bound certifies a source bank of N supports, put

$$B = k \binom{N}{2} - \mathcal{P}(Nt, n), \quad J = \min\{j \geq 1 : \mathcal{P}(N, j) \leq \lfloor B/(p - n) \rfloor\}.$$

For the length-257 example, the certified source and label counts are

$$\begin{aligned} N &= 12326787713559741838360124456674964971, \\ J &= 2772727901311387298989842831775757339. \end{aligned}$$

The standard-library checker in `research/logarithmic_length_lines/` uses rational upper bounds for the concentration denominator, exact integer collision arithmetic, and rational logarithm intervals for the reported decimal comparisons. It also verifies primality by Lucas–Lehmer, checks the strict sufficient Elias inequality $p^{t-k}t^t(n-t)^{n-t} > n^n$, and exhaustively tests 60 small-field support classes over 1,252 admissible poles. The discovery scan is not an optimality proof.

4 A fixed-gap obstruction that survives sublinear remainders

The preceding line theorem uses gaps that shrink with length. It therefore does not settle a proposed bound of the form $c_1 2^{c_2 H_2(\rho)/\eta} n + o(n)$ when the remainder is specified separately for each fixed gap. We now obtain a fixed-gap family.

There are two steps. First, keep one finite list example and replace each evaluation point by the same number of preimages under a power map. Degrees and agreement counts grow by the same factor, preserving both rate and gap. Second, use each nearby codeword on n distinct parameters of a word line. This forces a linear discrepancy from any proposed coefficient smaller than the list size.

Lemma 4.1 (A list example over the integers). *Fix rational $\rho \in (0, 1)$ and finite $c_1, c_2 \geq 1$. There are fixed integers $1 \leq k < t < m$, with $k/m = \rho$, distinct t -subsets $A_1, \dots, A_L \subset \{1, \dots, m\}$, and integer polynomials $W, G_1, \dots, G_L$ satisfying*

$$\deg G_i < k, \quad W(X) - G_i(X) = \prod_{a \in A_i} (X - a), \quad L > c_1 2^{c_2 H_2(\rho)/\eta}, \quad \eta = (t - k)/m.$$

The data is chosen before choosing a reducing prime.

Proof. Apply the moment construction over the integers. Equality of binomial moments implies equality of power sums over $\mathbb{Q}$, and Newton identities imply equality of the leading coefficients as integers. To see that a sufficiently large list exists without fixing a field, write $\rho = a/b$ and take $m = bu^3$, $k = au^3$, $s = u$, $t = k + s$. The logarithm of the lower bound in Lemma 2.1 is $bH_2(\rho)u^3 - O_\rho(u^2 \log u)$, whereas the logarithm of the proposed bound is $\log_2 c_1 + c_2 H_2(\rho)bu^2$. Choose a sufficiently large, now fixed, integer u and a class exceeding the proposed bound. Translating the original domain $\{0, \dots, m - 1\}$ by one, and substituting $X - 1$ in each polynomial, preserves the identities and avoids zero. $\square$

Lemma 4.2 (Replacing each point by a full power-map fiber). *Fix integer data as in Lemma 4.1. For every positive integer B and every bound T_0 , there is a prime $p > T_0$ and a domain $D_B \subset \mathbb{F}_p$ of size mB such that $x \mapsto x^B$ maps D_B onto $\{1, \dots, m\}$ with exactly B preimages per point. The polynomials $G_i(X^B)$ give distinct codewords in $\text{RS}(\mathbb{F}_p, D_B, kB)$, each agreeing with $W(X^B)$ on exactly tB coordinates.*

Proof. Let $K_B/\mathbb{Q}$ be the splitting field of $F_B(X) = \prod_{a=1}^m (X^B - a)$. The completely split case of Chebotarev's density theorem gives infinitely many primes that split completely in this finite Galois extension.¹ Exclude primes at most $\max\{T_0, m, B\}$ and primes dividing the nonzero discriminant of F_B . At any remaining such prime, all roots of F_B reduce to distinct elements of $\mathbb{F}_p$. Let D_B be these roots. The roots of each $X^B - a$ form one fiber of size B , and the fibers are disjoint.

Now $\deg G_i(X^B) \leq (k - 1)B < kB$, and

$$W(X^B) - G_i(X^B) = \prod_{a \in A_i} (X^B - a).$$

The difference vanishes on exactly the t selected fibers. The codewords remain distinct: for $A_i \neq A_j$, a point in a fiber indexed by $A_i \setminus A_j$ makes exactly one of the two displayed products zero. No further coefficient-preservation condition is needed. $\square$

Only the existence of arbitrarily large completely split primes is used; we require no bound on the least such prime. The construction may therefore use fields much larger than those in Theorem 2.2.

Lemma 4.3 (A list forces a selected-witness threshold). *Let $C \subset F^n$ be a nonzero linear code. Suppose a word w has L distinct codewords $q_1, \dots, q_L$ within distance θ , where $L \leq n$ and $|F| \geq Ln$. There is a nonconstant affine word line with Ln distinct parameters and a selected nearby codeword at each, such that no parametrized codeword line contains more than n selected pairs. Thus line-decodability with concurrency n fails at every positive threshold smaller than Ln .*

¹See J. S. Milne, *Algebraic Number Theory*, Theorem 8.31 and Corollary 8.32, [course notes](#). Only infinitude of completely split primes is used.

Proof. Choose nonzero $v \in C$ and disjoint sets $S_1, \dots, S_L \subset F$, each of size n . Put $f(z) = w + zv$ and select $p_z = q_i + zv$ for $z \in S_i$. Translation preserves distance, so every selected codeword is nearby. Fix $P_0, P_1 \in C$. If $P_1 = v$, the equality $P_0 + zP_1 = q_i + zv$ holds only for the branch with $P_0 = q_i$, if any, and that branch has n parameters. If $P_1 \neq v$, the equation $P_0 - q_i = z(v - P_1)$ has at most one solution per branch, giving at most $L \leq n$ selected pairs. The argument compares pairs at the same z ; it does not require the unparametrized images of the lines to be disjoint. $\square$

Theorem 4.4 (Failure at a fixed gap with any sublinear remainder). *For every rational $\rho \in (0, 1)$ and finite $c_1, c_2 \geq 1$, there are a fixed gap $\eta \in (0, 1 - \rho)$, an integer $L > c_1 2^{c_2 H_2(\rho)/\eta}$, and a sequence of prime-field Reed–Solomon codes of lengths $n_B \rightarrow \infty$ and exact rate ρ , all with $\theta = 1 - \rho - \eta$ strictly below their characteristic-based Elias radii, such that the following holds. Each code admits Ln_B selected nearby witnesses on a nonconstant affine word line, but no parametrized codeword line contains more than n_B selected pairs. Consequently, for every sequence $R_B = o(n_B)$, all sufficiently large codes fail line-decodability with concurrency n_B and threshold*

$$a_B = c_1 2^{c_2 H_2(\rho)/\eta} n_B + R_B. \quad (21)$$

The code family is chosen independently of the remainder sequence.

Proof. Fix the integer example from Lemma 4.1; hence m, k, t, L, η are now constants. Let B grow with $n_B = mB \geq L$. For each B , choose the least completely split prime from Lemma 4.2 satisfying

$$p_B > \max\{Ln_B, m, 2^{1/\eta}\}$$

and the separability condition in that lemma. This fixes a prime sequence without reference to R_B . The lifted code has rate $k/m = \rho$ and contains a list of size L at distance exactly $1 - t/m = \theta$. It is nonzero, and the hypotheses of Lemma 4.3 hold. Moreover

$$H_{p_B}(\theta) \leq \theta + 1/\log_2 p_B < \theta + \eta = 1 - \rho,$$

which verifies the Elias condition. Put $A = c_1 2^{c_2 H_2(\rho)/\eta}$. Since $L - A > 0$ is fixed, eventually $R_B < (L - A)n_B$, so $a_B < Ln_B$. Lemma 4.3 contradicts line-decodability at this threshold. $\square$

The gap can depend on the proposed constants, but is fixed along the resulting family. Even a remainder chosen separately for that family cannot absorb its positive linear discrepancy. This last construction concerns the arbitrary selection of nearby witnesses: it has correlated agreement, since (q_i, v) agrees jointly with (w, v) wherever q_i agrees with w . Also, its selected count Ln never exceeds $\ell_C(\theta)n$. Thus the obstruction rules out the proposed quantitative coefficient, not a line threshold expressed using a valid bound on the actual maximum list size.

4.1 Anchored padding gives genuine correlated-agreement failure

The preceding selection argument is not needed if we allow a small padding of the lifted domain. The following construction keeps the rate and gap exactly fixed and makes correlated agreement impossible at the threshold, for every challenge.

Lemma 4.5 (Anchored list-to-line padding). *Let D_0 be a set of m distinct nonzero integers containing 1. Let $1 \leq k < t < m$, and let $A_1, \dots, A_L$ be distinct t -subsets of D_0 , all containing 1, whose monic locator polynomials have the same coefficients in degrees $t, t-1, \dots, k$. Choose an integer $1 \leq r < t - k$. For every positive integer B there is an arbitrarily large prime p and a Reed–Solomon code with*

$$n_B = B(m + r), \quad k_B = Bk, \quad A_B = Bt,$$

admitting a nonconstant received line $f + zg$ with at least $L(Br + 1)$ distinct nearby challenges at agreement threshold A_B . Every nearby word on the line has exactly one nearby codeword, and the parameter-zero word has maximum agreement exactly $A_B - 1$. No degree- $< k_B$ polynomials F, G jointly agree with f, g on A_B coordinates. The radius can be chosen strictly below Elias. If $k \geq 2$, one may instead set $k_B = B(k - 1)$ and allow $1 \leq r \leq t - k$, with the same conclusion.

Proof. Write $F_i(X) = \prod_{a \in A_i} (X - a)$, let W be their common degree- $\geq k$ part, and put $G_i = W - F_i$. Then $\deg G_i < k$ and $G_i(1) = W(1)$ for every i . Use the proof of Lemma 4.2 for the splitting field of $\prod_{a \in D_0} (X^B - a)$, with a prime-size bound specified below, and remove the point 1 from its union of full fibers. On the remaining set D_{old} of size $mB - 1$, put

$$f(x) = F_{\text{old}}(x) := \frac{W(x^B) - W(1)}{x - 1}, \quad g(x) = 0.$$

The candidate polynomials are

$$P_i(X) = \frac{G_i(X^B) - W(1)}{X - 1}.$$

Their numerators vanish at 1, so they are polynomials of degree at most $B(k - 1) - 1 < k_B$, with a zero numerator interpreted in the usual way. Each agrees with F_{old} at exactly $Bt - 1 = A_B - 1$ old coordinates, since their difference is $F_i(X^B)/(X - 1)$. The candidates remain distinct after reduction, since their old agreement supports are distinct. Also $\deg F_{\text{old}} = A_B - 1$.

Let $\mathcal{T}$ contain all degree- $< k_B$ polynomials interpolating F_{old} on some k_B old coordinates. Then $|\mathcal{T}| \leq M := \binom{mB-1}{k_B}$, and every selected P_i belongs to $\mathcal{T}$. Add $q_B = Br + 1$ distinct points x_j outside $D_{\text{old}} \cup \{1\}$ at which the values $\{P(x_j) : P \in \mathcal{T}\}$ are pairwise distinct and avoid $F_{\text{old}}(x_j)$. There are enough such points if

$$p > mB + (k_B - 1) \binom{M}{2} + (A_B - 1)M + q_B,$$

by the root bounds for $P - P'$ and $F_{\text{old}} - P$. Choose nonzero values g_j successively so that all labels

$$\frac{P(x_j) - F_{\text{old}}(x_j)}{g_j}, \quad P \in \mathcal{T},$$

are distinct across all P, j . At step j , at most $(j - 1)M^2$ nonzero values are forbidden, so $p - 1 > (q_B - 1)M^2$ suffices. Set $f = F_{\text{old}}$ on the entire domain and $g(x_j) = g_j$ on the new coordinates. Every selected P_i is nearby at q_B distinct nonzero labels, giving Lq_B such challenges at length n_B .

Any nearby polynomial has at least $A_B - q_B \geq k_B$ old agreements, so belongs to $\mathcal{T}$. It has at most $A_B - 1$ old agreements and, by label injectivity, at most one new agreement. Thus it has exactly $A_B - 1$ old agreements and one new agreement. The same injectivity prevents two polynomials from being nearby at the same challenge. This proves uniqueness for the entire line, not merely for the selected witnesses.

Since $f = F_{\text{old}}$ globally has degree $A_B - 1$, it agrees with any degree- $< k_B$ polynomial on at most $A_B - 1$ coordinates. The selected candidates attain that number at parameter zero, so this word is exactly one coordinate beyond the nearby radius. The same degree bound makes joint agreement on A_B coordinates impossible.

All the prime-size conditions, and $((t - k)/(m + r)) \log_2 p > 1$, can be imposed before applying the infinitude of completely split primes. The last condition verifies the strict Elias inequality. For the sharper dimension $k_B = B(k - 1)$, the candidates still fit, and $A_B - q_B \geq k_B$ still holds when $r \leq t - k$, preserving the exhaustive candidate-pool argument. The larger capacity gap preserves the Elias inequality. $\square$

Theorem 4.6 (The linear coefficient at a fixed positive gap). *Fix rational $\rho \in (0, 1)$ and put $H = H_2(\rho)$. For every sufficiently small rational gap $\eta > 0$, there are prime-field Reed–Solomon codes of lengths $n_B \rightarrow \infty$, exact rate ρ , and exact agreement fraction $\rho + \eta$, and received lines with at least $C_\rho(\eta)n_B$ distinct nearby challenges, where*

$$\log_2 C_\rho(\eta) \geq \frac{H^2/2 - o_\rho(1)}{\eta^2 \log_2(1/\eta)}. \quad (22)$$

No pair of codewords has correlated agreement at that threshold, every nearby word on the line has a unique nearby codeword, and the parameter-zero word has distance $1 - \rho - \eta + 1/n_B$. The nearby radius is strictly below the characteristic-based Elias radius. The $o_\rho(1)$ is as rational η tends to zero; $C_\rho(\eta)$ is fixed as B grows.

Proof. The proof of Lemma 4.5 allows arbitrary integer length n , dimension K , and padding count $q = n - mB + 1$, provided

$$q \geq 1, \quad K \geq B(k - 1), \quad K - 1 + q < Bt.$$

Indeed, the candidates have degree at most $B(k - 1) - 1$, and the global received polynomial has degree $Bt - 1$. The inequality $Bt - q \geq K$ preserves the full candidate-pool argument proving uniqueness at every nearby parameter.

Put $a = \rho + \eta$ and choose

$$m = \left\lfloor \frac{H}{\eta^2 \log_2(1/\eta)} \right\rfloor, \quad t = \left\lceil \frac{am}{1 - \eta/2} \right\rceil, \quad k = \left\lfloor \frac{\rho t}{a} \right\rfloor, \quad s = t - k.$$

For sufficiently small η , one has $1 \leq k < t < m$ and

$$am < t < \frac{am}{1 - \eta}, \quad \frac{t}{m} \rightarrow \rho, \quad s = (1 + o(1))\eta m.$$

The strict upper inequality absorbs the rounding error because $\eta m \rightarrow \infty$. Among the t -subsets containing 1, the same binomial-moment count as in Lemma 2.1 gives a class of size

$$L \geq \left\lceil \frac{\binom{m-1}{t-1}}{\prod_{j=1}^s R_j} \right\rceil,$$

where R_j are the full moment-range sizes for length m . They still bound the ranges after the anchor is fixed. The moment-range estimate used in Theorem 2.2 gives

$$\log_2 \prod_{j=1}^s R_j \leq \frac{s^2}{2} \log_2(m/s) + O(s^2 + s \log m) = (H/2 + o(1))m.$$

Since $\binom{m-1}{t-1} = (t/m)\binom{m}{t}$, it follows that

$$\log_2 L \geq (H/2 - o(1))m.$$

Let B grow through multiples making $n_B = Bt/a$ and $K_B = \rho n_B$ integers, and set $q_B = n_B - mB + 1$. Then $K_B \geq Bk$ and

$$K_B - 1 + q_B = B((1 + \rho)t/a - m) < Bt.$$

The flexible padding construction therefore has exact rate ρ , exact gap η , and more than $C_\rho(\eta)n_B$ nearby labels for $C_\rho(\eta) = (1 - am/t)L$. Since $1 - am/t \sim \eta/2$, this prefactor costs only $O(\log(1/\eta))$ in its logarithm. Substituting the chosen m proves (22). Arbitrarily large splitting primes ensure $n_B = o(p)$ and strict Elias. The rationality of η is used only to keep the displayed rate and agreement fraction exactly fixed at integer lengths. $\square$

The same parameters also give a list lower bound at every sufficiently small rational gap. Keep all mB fiber points and add $n_B - mB$ distinct points outside the fibers. The received polynomial $W(X^B)$ has the L nearby candidates $G_i(X^B)$, each of degree at most $B(k-1) < K_B$. Each difference is the degree- Bt locator $F_i(X^B)$, so each candidate has exactly Bt agreements. Thus, at exact rate ρ and exact gap η , a displayed list has logarithm at least $(H^2/2 - o_\rho(1))/(\eta^2 \log_2(1/\eta))$ on arbitrarily long domains with $n_B = o(p)$ and strict Elias. This is a lower bound on one list; it does not assert growth of that list with B .

In particular, no uniform exceptional-count bound $c_1 2^{c_2/\eta} n + o(n)$ can hold, even when its remainder is specified separately at each fixed gap. Choose a gap for which $C_\rho(\eta) > c_1 2^{c_2/\eta}$; the positive linear discrepancy then absorbs any such remainder. Every exhibited nearby challenge also violates the full common-witness conclusion, because it has no common agreement set of the required size at all. The construction remains linear in n at each fixed gap and does not establish a quadratic lower bound there.

Uniqueness is restricted to the displayed line. Once $q_B \geq L$, the same code has another received word with exactly the L selected codewords nearby: retain the old word, assign each selected codeword its own padding coordinate, and set all remaining padding values outside the corresponding evaluation image of $\mathcal{T}$. Injectivity of those evaluation maps excludes every other candidate. Thus line-wise uniqueness does not imply a small global maximum list.

It also supplies a selected-witness counterexample with concurrency at most n_B . Indeed, on any proposed affine codeword graph let a coordinates agree identically with the received line. Then $a < A_B$, and each other coordinate agrees at at most one selected challenge. If the graph contains v selected nearby pairs, $vA_B \leq va + n_B - a$, giving $v \leq n_B - A_B + 1 \leq n_B$. Exact anchored moment counts and prime-field fixtures of lengths 10, 20, 40 and 11, 22, 44 appear in [research/fixed_gap_padding/](#).

Corollary 4.7 (Polynomial-size prime fields). *For each fixed rational $\rho \in (0, 1)$ and every sufficiently small rational $\eta > 0$, the list lower bound and the exceptional-count lower bound $C_\rho(\eta)n$ above also hold along growing lengths with*

$$n = o(p), \quad p = n^{O_{\rho, \eta}(1)}.$$

The rate and gap are exact, the radius is strictly below Elias, and no pair of codewords has correlated agreement at the threshold. The parameter-zero word is exactly one coordinate beyond the nearby radius. This assertion does not include uniqueness of nearby codewords.

Proof. Fix the integer seed with m distinct nodes $a_1, \dots, a_m$ and anchor $a_* = 1$. If $B \mid p-1$, a standard multiplicative-character expansion counts the shifts c for which every $c + a_i$ is a nonzero B th power. The product-character Weil bound [2, Lemma 17], with constant additive phase, gives

$$\#\{c : c + a_i \in (\mathbb{F}_p^*)^B \text{ for all } i\} \geq \frac{p}{B^m} - m\sqrt{p} - m.$$

Indeed, extend all characters by zero at zero. The all-principal term is $p - m$; every other term has absolute value at most $m\sqrt{p} + m$, allowing for the excluded roots. The resulting bound is positive if $p > 16m^2 B^{2m}$.

By Linnik's least-prime theorem [3], there are absolute constants C, L_0 such that the least prime congruent to 1 modulo B^{2m+2} satisfies

$$B^{2m+2} < p \leq CB^{(2m+2)L_0}.$$

For all sufficiently large B , the nodes remain distinct modulo p and the preceding shift exists. Write $W_c(Y) = W(Y - c)$ and $G_{i,c}(Y) = G_i(Y - c)$, and choose $\alpha^B = c + a_*$. Remove α from the union

of the full fibers and replace the anchored quotients by

$$\frac{W_c(X^B) - W(a_*)}{X - \alpha}, \quad \frac{G_{i,c}(X^B) - W(a_*)}{X - \alpha}.$$

Their degrees and exact old-agreement counts are unchanged. For the padding, separate only the L selected candidates, rather than the full interpolation pool $\mathcal{T}$. The sufficient bounds are then

$$p > mB + (K_B - 1) \binom{L}{2} + (Bt - 1)L + q_B, \quad p - 1 > (q_B - 1)L^2.$$

Both are eventually implied by the displayed lower bound on p , since the seed and L are fixed and $K_B, q_B = O(B)$. This gives $Lq_B > C_\rho(\eta)n_B$ distinct nearby labels. The global degree- $Bt - 1$ received word gives the far point and excludes correlated agreement. Restoring the anchor gives the list-only construction as before. Finally $n_B = \Theta(B)$, so $n_B = o(p)$, the field size is polynomial in n_B , and strict Elias holds eventually. The exponent may depend on ρ, η ; the displayed choice is $O_\rho(\eta^{-2}/\log(1/\eta))$ as $\eta \rightarrow 0$. Separating only the selected candidates does not establish whole-line uniqueness. $\square$

Concrete fixed-gap coefficients. The earlier finite integer classes also supply explicit seeds. Among L_0 supports of size t on m points, some point belongs to at least $L = \lceil tL_0/m \rceil$ supports. If that point is a , the integer affine change $x \mapsto 2(x - a) + 1$ makes it the anchor 1 and makes every domain point nonzero. Multiplying transformed locator polynomials by 2^t preserves monicity, integer coefficients, and the common prefix. Using the sharper dimension $B(k - 1)$, the two previously certified classes give

(m, k, t)	r	L	coefficient $C = rL/(m + r)$
$(64, 32, 34)$	2	2949985688957	2949985688957/33
$(157, 63, 68)$	5	12200781436007129430	10167317863339274525/27

The resulting rates and gaps are $(31/66, 1/22)$ and $(31/81, 1/27)$, respectively. In the second family, more than $3.76 \cdot 10^{17}n_B$ distinct challenges are nearby, while correlated agreement at the threshold is impossible. These are growing-prime-field families obtained from the integer classes; they are not statements about retaining the M31 alphabet or a prescribed subgroup domain.

4.2 Padding with collisions improves the finite-field bounds

The preceding constructions required distinct labels. Allowing collisions and averaging their union gives a stronger bound when the field is fixed. The mechanism is simple: many codewords already agree with a fixed core on $A - 1$ coordinates, just one short of the threshold. Each added coordinate lets some of them acquire the missing agreement, at different line parameters. With enough added coordinates, these opportunities cover almost all parameters. Keeping the starting word equal to a polynomial of degree $A - 1$ makes it one coordinate too far from the code and rules out any common agreement explanation. Write $\mathcal{P}(h, N)$ for the minimum number of colliding pairs for h objects in N bins, as in Section C.2.

Lemma 4.8 (Averaged padding). *Let $D_{\text{old}} \subset \mathbb{F}_p$ have N points, and let w be a polynomial of degree $A - 1$. Suppose $P_1, \dots, P_L$ are distinct polynomials of degree less than $K < A$, each agreeing*

with w at exactly $A - 1$ old coordinates, and each pairwise difference has degree at most d . Choose $1 \leq q \leq p - N$ with $K - 1 + q < A$, and put

$$R = p - N, \quad T = d \binom{L}{2} - \mathcal{P}(L(A - 1), N), \quad M = \frac{L^2 R}{LR + 2T}.$$

Then $T \geq 0$. On a domain of size $n = N + q$, there is a received line with at least

$$J = \left\lceil p \left(1 - \left(1 - \frac{M}{p} \right)^q \right) \right\rceil \quad (23)$$

nearby parameters at agreement threshold A , and no pair of codewords has correlated agreement at that threshold.

Proof. For $x \notin D_{\text{old}}$, let s_x be the number of distinct values among the $P_i(x)$, with multiplicities $r_{x,v}$. Each polynomial pair has at most d coincidences. Their old common agreements consume at least $\mathcal{P}(L(A - 1), N)$ such coincidences, so

$$\sum_{x \notin D_{\text{old}}} \sum_v r_{x,v}^2 \leq LR + 2T.$$

Cauchy–Schwarz gives $R^{-1} \sum_x s_x \geq M$. Choose the q points x_j with largest s_x . Set $f = w, g = 0$ on the old coordinates, and $f(x_j) = b_j, g(x_j) = 1$ on the new ones. For independent uniform offsets b_j , the expected size of the union of labels $P_i(x_j) - b_j$ is

$$p \left(1 - \prod_{j=1}^q (1 - s_{x_j}/p) \right) \geq p (1 - (1 - M/p)^q),$$

by arithmetic–geometric mean. Some offsets attain its ceiling. Every label in the union gives a candidate with at least A agreements. For common witnesses F, G , if $G = 0$ the joint agreement is at most $A - 1$ by the degree of $w - F$; otherwise it is at most $K - 1 + q < A$. $\square$

What a growing fixed-gap list would imply. The same averaging mechanism gives a useful transfer without requiring the received word to have a low-degree polynomial representation. The conclusion here concerns the entire agreement support of a nearby candidate; it does not assert the absence of ordinary correlated agreement.

Proposition 4.9 (List amplification with exact parameter scaling). *Let $k + 1 \leq A \leq N$, $k \geq 2$, and let $p \geq 2N + 1$ be prime. Suppose a word on N distinct points of $\mathbb{F}_p$ has L distinct degree- $< k$ candidates, each with at least A agreements. There is a Reed–Solomon code of length $2N$ and dimension k over the same prime field, and a received line $f + zg$, with at least*

$$J \geq \left\lceil \frac{pAL}{p + 3AL} \right\rceil \geq \frac{1}{4} \min\{AL, p\}$$

distinct labels z for which some nearby candidate has at least A agreements and g has no degree- $< k$ interpolant on its entire agreement support. Thus source rate ρ and gap η become exactly $\rho/2$ and $\eta/2$.

Proof. Some coordinate a belongs to at least $\ell \geq AL/N$ candidate agreement supports. Retain those candidates and form

$$Q_P(X) = \frac{P(X) - w(a)}{X - a}, \quad f(x) = \frac{w(x) - w(a)}{x - a}, \quad g(x) = 0 \quad (x \neq a).$$

The Q_P are distinct of degree at most $k - 2$, and have at least $A - 1 \geq k$ agreements on the $N - 1$ retained coordinates. At each of the $R = p - N$ unused coordinates let s_x be the number of distinct values of the Q_P . Counting pairwise coincidences and applying Cauchy–Schwarz gives

$$\frac{1}{R} \sum_x s_x \geq \mu := \frac{\ell R}{R + (\ell - 1)(k - 2)}.$$

Choose the $q = N + 1$ unused coordinates with largest s_x , put $g = 1$ there, and choose their f values independently and uniformly. The expected union of the label sets $\{Q_P(x) - f(x)\}$ is at least $p(1 - (1 - \mu/p)^q)$, exactly as in Lemma 4.8. Every label in this union adds a new agreement to one selected candidate. A degree- $< k$ interpolant for g on that candidate's support would vanish at its at least k old agreements, hence vanish identically, contradicting its value 1 at the new agreement.

For the count, $R \geq p/2$ gives

$$\frac{q\mu}{p} \geq \frac{N\ell}{p + 2(k - 2)\ell} \geq \frac{N\ell}{p + 2N\ell} \geq \frac{AL}{p + 2AL}.$$

Using $1 - (1 - u)^q \geq 1 - e^{-qu} \geq qu/(1 + qu)$ proves the claimed bound. The new length is $N - 1 + q = 2N$, while its dimension and agreement threshold remain k and A , proving exact parameter scaling. $\square$

At fixed positive rate and gap, this gives $J = \Omega(\min\{NL, p\})$. In particular, a list lower bound N^c over superpolynomial prime fields would force $\Omega(N^{c+1})$ full-support exceptions at half the rate and gap. A universal linear full-support bound at the latter parameters would imply constant source list sizes whenever $p/N \rightarrow \infty$. This identifies a sufficient route to exponent lower bounds; it does not supply the growing lists. The full-length lists with $N = p - 1$ fall outside its field-size hypothesis.

Corollary 4.10 (Amplification within an open agreement regime). *Under the source hypotheses of Proposition 4.9, fix rational $0 < \tau \leq 1$ with τN integral. There is a line of length $(1 + \tau)N$, dimension k , and agreement threshold A with at least*

$$\frac{\tau}{5} \min\{AL, p\}$$

full-support MCA-bad labels. Its rate and agreement fraction are $\rho/(1 + \tau)$ and $a/(1 + \tau)$, where $\rho = k/N$ and $a = A/N$. If $a > h(\rho)$ for a continuous sufficient agreement curve h , then every sufficiently small positive τ preserves that strict inequality at the output parameters. If $a^2 < \rho$, the output also remains below Johnson.

Proof. Use the same proof with $q = \tau N + 1$ padding coordinates. Before specializing q , its count gives

$$J \geq \frac{pq\ell}{p + (2(k - 2) + q)\ell} \geq \frac{pqAL}{Np + (2(k - 2) + q)AL}.$$

Here the first inequality follows from $\mu \geq \ell p/(p + 2(k - 2)\ell)$ and $p(1 - (1 - \mu/p)^q) \geq pq\mu/(p + q\mu)$; the second uses monotonicity in ℓ . Since $q \geq \tau N$ and $2(k - 2) + q < 4N$, this is at least $\tau pAL/(p + 4AL) \geq (\tau/5) \min\{AL, p\}$. The old $A - 1 \geq k$ agreements still exclude a full-support direction witness. The parameter identities are exact. The strict inequality $a/(1 + \tau) > h(\rho/(1 + \tau))$ follows by continuity at $\tau = 0$. Finally $a^2/(1 + \tau)^2 < \rho/(1 + \tau)$ whenever $a^2 < \rho$. $\square$

Thus, within any such open regime, source lists of size $\Omega(N^c)$ over sufficiently large primes would force $\Omega(n^{c+1})$ full-support exceptions at nearby parameters in the same regime. This remains a conditional statement: it requires the growing source lists. For arbitrary lengths, taking $q = \lfloor \tau N \rfloor + 1$ preserves the strict margin eventually, with parameters converging to the displayed pair.

The exact finite checks in `research/prime_field_tightness/check_prime_boundary_amplification.py` include parameter halving over $\mathbb{F}_{101}$ and examples where full-support failure coexists with ordinary correlated agreement. The variable-padding checks in `research/prime_field_tightness/check_sparse_padding.py` independently exhaust all direction witnesses for the selected bad supports.

Lemma 4.11 (Padding around a far point). *Keep the core hypotheses and R, T, M of Lemma 4.8, but allow any $1 \leq q \leq p - N$, without requiring $K - 1 + q < A$. There is a nonconstant received line on $n = N + q$ coordinates with at least*

$$J = \left\lceil (p-1) \left(1 - \left(1 - \frac{M}{p-1} \right)^q \right) \right\rceil \quad (24)$$

nearby parameters at agreement threshold A . Its parameter-zero word has distance exactly $(n - A + 1)/n$, one coordinate beyond the nearby radius. No pair of codewords has correlated agreement at threshold A .

Proof. Each $P_i - w$ has degree $A - 1$ and already has $A - 1$ roots on the core. It therefore has no roots outside the core. At an available point x , the difference image $S_x = \{P_i(x) - w(x) : 1 \leq i \leq L\}$ is a subset of $\mathbb{F}_p^*$ of size s_x . The preceding collision bound still gives average image size at least M , so in particular $M \leq p - 1$. Choose the q points with largest images.

Set $f = w$ on the entire domain, let $g = 0$ on the core, and choose $g(x_j)$ independently uniformly in $\mathbb{F}_p^*$ on the padding. The candidate P_i is nearby at the nonzero parameter

$$z = \frac{P_i(x_j) - w(x_j)}{g(x_j)}.$$

For each fixed $z \neq 0$, these independent multiplicative translates hit z with probabilities $s_{x_j}/(p-1)$. The expected union size is

$$(p-1) \left(1 - \prod_j (1 - s_{x_j}/(p-1)) \right) \geq (p-1) (1 - (1 - M/(p-1))^q).$$

Some direction attains the ceiling. Finally $f = w$ agrees with every codeword on at most $A - 1$ coordinates by its degree, and each selected P_i attains that number. This proves the exact distance at zero and also rules out correlated agreement, regardless of the chosen direction. $\square$

4.3 Every point but one, with an almost maximally far exception

Pairing coordinates as $x, -x$ gives a stronger separation for the $c_1 = c_2 = 1$ prescription. The exceptional point can be just one coordinate short of the maximum possible Reed–Solomon distance, while its separation from the nearby radius approaches the entire capacity gap. The construction uses random paired domains and a multiplicative-character estimate, without a matching-moment class.

Lemma 4.12 (Sampling with blocks of paired coordinates). *Let r, m, D, q be positive integers with $D < m$ and $m + rq \leq (p - 1)/2$, where $p > \max\{(4r + 2)^2, 2m^2\}$ is prime. Put*

$$L = \binom{m}{D}, \quad \delta = \frac{p^r}{L} + \frac{(4r + 2)^{2r+2} + 2m^2}{p - 2m^2}.$$

There is an expected $O(n^2)$ -field-operation randomized procedure outputting a length- $n = 2(m + rq)$, dimension- $K = 2D - 1$ Reed-Solomon code and a received line. Every output has maximum agreement exactly $2D = K + 1$ at parameter zero. For every $\gamma > 0$, the probability that some nonzero parameter has fewer than $2D + 2r$ agreements is at most

$$\frac{q\delta}{\gamma} + (p - 1) \left(\frac{\gamma}{1 + \gamma} \right)^q. \quad (25)$$

Whenever all nonzero parameters are nearby, there is no correlated agreement at that threshold.

Proof. Choose $m + rq$ distinct orbits $\{a, -a\}$ uniformly from $\mathbb{F}_p^*$. Use the first m as the core and divide the remaining orbits into q blocks of r orbits each. For every D -subset I of the core, write

$$H_I(X) = \prod_{i \in I} (X^2 - a_i^2).$$

Fix one reference support I_0 , put $w = H_{I_0}$, and consider the candidates $P_I = w - H_I$. Their leading terms cancel and all odd coefficients vanish, so $\deg P_I \leq 2D - 2 < K$. Each candidate has exactly $2D$ core agreements. On each padding pair choose a shared uniform nonzero direction, independently across pairs. Set $f = w$ everywhere and $g = 0$ on the core.

We bound the image of $I \mapsto (-H_I(x_1), \dots, -H_I(x_r))$ for a random outside block. Let $U = p - 1$ and $V = U^r$. Fix representatives $x_1, \dots, x_r$ with distinct nonzero squares. Independently sampling seed variables from

$$T = \mathbb{F}_p \setminus \{0, \pm x_1, \dots, \pm x_r\}$$

induces a probability measure μ on $(\mathbb{F}_p^*)^r$ by $a \mapsto (x_j^2 - a^2)_{j=1}^r$. Every nonprincipal Fourier coefficient of μ has absolute value at most

$$\frac{2r\sqrt{p} + 2r + 1}{p - 2r - 1} \leq \frac{4r + 2}{\sqrt{p}} =: \lambda < 1.$$

Indeed, factor $x_j^2 - a^2$ into the two distinct linear factors at $\pm x_j$ and apply the product-character bound [2, Lemma 17], with constant additive phase. At least one of the $2r$ characters is nonprincipal. The extra $2r + 1$ allows for the deleted arguments. The displayed hypotheses imply the second inequality.

For two supports with u indices on each side of their symmetric difference, independent seeds give vector-collision probability

$$\frac{1}{V} \sum_{\chi} |\hat{\mu}(\chi)|^{2u} \leq \frac{1}{V} + \lambda^{2u}.$$

Conditioning all m seeds to have distinct squares costs at most $1/a_0$, where $a_0 = 1 - 2m^2/p > 0$, by the union bound. After this conditioning, pairs with $1 \leq u \leq r$ cannot collide: cancel common factors and regard the residual products as distinct monic degree- u polynomials in $Y = X^2$. Their difference has degree at most $u - 1 < r$, so cannot vanish at the r distinct values x_j^2 . For $u \geq r + 1$, the conditional collision probability is at most $(V^{-1} + \lambda^{2r+2})/a_0$.

Let E be the collision energy of the block image, counting ordered support pairs, and set $\Delta = VE/L^2 - 1 \geq 0$. Counting the diagonal pairs separately gives

$$\mathbb{E}\Delta \leq \frac{V}{L} + \frac{1 + V\lambda^{2r+2}}{a_0} - 1 \leq \delta.$$

This expectation applies to a uniform core and an outside block. Equivalently one may choose the block first and the admissible core second: each has the same number of compatible choices. Every block in the sampled domain has this marginal law, so Markov's inequality and a union bound give $\Pr(\text{some block has } \Delta > \gamma) \leq q\delta/\gamma$. No independence among these events is assumed.

When $\Delta \leq \gamma$, Cauchy–Schwarz gives an image of size at least $V/(1 + \gamma)$. Conditional on the core and blocks, the random direction vectors are independent and uniform in $(\mathbb{F}_p^*)^r$. For any fixed nonzero label z , multiplication of a direction vector by z preserves this distribution. Thus, if all blocks are good, the expected number of uncovered labels is at most the second term of (25). A block hit supplies $2r$ padding agreements. Combining the two failure events proves the bound.

The polynomial w is monic of degree $2D$, and zero already agrees at its $2D$ core roots. This proves the exact far distance for every output and rules out correlated agreement. Sampling distinct orbits and forming and evaluating w take expected $O(n^2)$ field operations. $\square$

Exact rates and parity. The same construction realizes every fixed rational rate, with at most one additional coordinate. Take n through odd multiples of the reduced denominator of ρ , and set $K = \rho n$. If K is odd, use $D = (K + 1)/2$ and add zero if n is odd, with $f(0) = w(0), g(0) = 0$. No selected candidate agrees there, since $H_I(0) \neq 0$. If K is even, then n is odd. Take $D = K/2$, multiply w and all candidates by X , and add zero with $f(0) = g(0) = 0$. On a padding pair use $xh, -xh$ as its directions. Dimension, threshold and exact far agreement all increase by one. In both cases the nearby threshold is $K + 2r + 1$, and the exact far agreement is $K + 1$. The probability bound is unchanged.

Theorem 4.13 (Fixed separation arbitrarily close to the full gap). *Fix rational $\rho \in (0, 1)$, an integer $r \geq 1$, $c_1 > 0$, and $0 < c_2 < 1 + 1/(2r)$. For every sufficiently large prime p , a code of exact rate ρ and length $n = \Theta(\log p)$ has a received line with every nonzero parameter within radius $\theta = 1 - \rho - \eta$, where $\eta = (2r + 1)/n$, and parameter zero at distance*

$$1 - \rho - 1/n = \theta + \frac{2r}{2r + 1}\eta.$$

The radius is strictly below Elias, there is no correlated agreement, and $c_1 n 2^{c_2 H_2(\rho)/\eta} = o(p)$. The code and line can be sampled in expected polynomial time in $\log p$, with success probability $1 - p^{-\Omega(1)}$.

Proof. Choose $\rho < \beta < 1$ close to ρ , put $\alpha = \rho/\beta$, and choose

$$\frac{2r}{\alpha H_2(\beta)} < C < \frac{2r + 1}{\max\{1, c_2\} H_2(\rho)}.$$

Take $n = C \log_2 p + O(1)$ with the rate parity above, $m = \alpha n/2 + O(r)$, and $D = \rho n/2 + O(1)$, adjusting m so the padding forms complete blocks. Then $L = \binom{m}{D} \geq p^{r+\epsilon}$ for some fixed $\epsilon > 0$, and $q = \Theta(\log p)$. Hence $\delta = p^{-\Omega(1)}$. A sufficiently small fixed γ makes both terms of (25) $p^{-\Omega(1)}$. The degree and threshold identities give the distances. Finally $\eta \log_2 p \rightarrow (2r + 1)/C > \max\{1, c_2\} H_2(\rho)$ proves strict Elias and the numerical comparison. $\square$

4.3.1 Completing the joint image and determining exact distances

A few additional core choices can fill the entire product image. This removes the need for multiple padding blocks and gives an exact distance profile on an affine space. The completion step follows the classical random-translate strategy of Erdős and Rényi [1, Theorem 2]; here a character bound controls nonuniform translates arising from new evaluation points.

Lemma 4.14 (Completion by pairs of alternatives). *Fix r nonzero padding representatives with distinct squares. Let m_0, D_0, t be positive integers with $D_0 < m_0$, put $n = 2(m_0 + 2t + r)$ and $C = 4r + 2$, and assume $p > \max\{(n+1)^2, C^2\}$. Define*

$$L_0 = \binom{m_0}{D_0}, \quad d_0 = \frac{p^r}{L_0} + \frac{C^{2r+2} + 2m_0^2}{p - 2m_0^2}, \quad \Phi(h) = \frac{h^2 + C^2 h/p}{1 - 4/p}.$$

If $(p-1)^r \Phi^{ot}(d_0) < 1$, there is a core of $m_0 + 2t$ sign-orbits and a family of $(D_0 + t)$ -subsets whose locator values on the padding fill $(\mathbb{F}_p^)^r$. More generally, uniform successive sampling of the core succeeds with failure probability at most*

$$\frac{d_0}{h_0} + \sum_{i=1}^t \frac{\Phi(h_{i-1})}{h_i} \quad (26)$$

for any positive thresholds $h_0, \dots, h_t$ with $(p-1)^r h_t < 1$.

Proof. Choose m_0 base orbits uniformly outside the padding. The energy argument in Lemma 4.12 applies with the padding fixed: the mean missing fraction of the product image of all D_0 -subsets is at most d_0 .

At each subsequent step add two new core orbits, represented by a, b . Each candidate chooses exactly one of these two. If S is the current product image and $v(a) = (x_j^2 - a^2)_{j=1}^r$, the new image is $Sv(a) \cup Sv(b)$. For its complement B of density h , the new missing set is the intersection of two translates of B .

First sample a, b independently outside zero, the padding, and all previous core orbits. At most $n+1 < \sqrt{p}$ field elements are excluded. The same character estimate gives every nonprincipal Fourier coefficient of the distribution μ of $v(a)$ absolute value at most $C/\sqrt{p}$. Using normalized Fourier coefficients on $(\mathbb{F}_p^*)^r$, Parseval gives

$$\mathbb{E} \frac{|Bv(a) \cap Bv(b)|}{(p-1)^r} = \sum_{\chi} |\widehat{1_B}(\chi)|^2 |\widehat{\mu}(\chi)|^2 \leq h^2 + \frac{C^2}{p} h.$$

The available set has size at least $p/2$. Conditioning $a \neq \pm b$ therefore costs at most $1/(1-4/p)$, so the conditional mean missing fraction is at most $\Phi(h)$. This bound holds uniformly over every previous choice of the core and its image.

Successive averaging proves the existence assertion: at each step choose a pair attaining at most the conditional mean. Once the missing fraction is less than $1/(p-1)^r$, no vector is missing. Alternatively, Markov's inequality at the base step and conditionally at each later step, followed by a union bound, gives (26). The sampler does not need to compute any of the images. $\square$

Corollary 4.15 (Exact distances and almost the entire capacity gap). *Fix rational $\rho \in (0, 1)$, put $b = \log_2 p$ and $H = H_2(\rho)$, and let $r = r(p) \geq 1$ satisfy $r \log b = o(b)$. Over every sufficiently large prime field, a code of exact rate ρ and length*

$$n = \frac{(2r + 4/5)b}{H} + O_\rho(1)$$

has an affine space with independent directions $e_1, \dots, e_r$, each supported on a disjoint pair of coordinates, such that

$$\text{dist} \left(f + \sum_{j=1}^r z_j e_j, C \right) = 1 - \rho - \frac{1 + 2 \text{wt}(z)}{n} \quad (z \in \mathbb{F}_p^r). \quad (27)$$

Here $\text{wt}(z)$ counts its nonzero entries. At radius $\theta = 1 - \rho - \eta$, $\eta = (2r+1)/n$, the nearby parameters are exactly $(\mathbb{F}_p^*)^r$, and their decoding lists are pairwise disjoint. The radius is strictly below Elias, $n2^{H/\eta} = o(p)$, and there is no correlated agreement. The code and affine space can be sampled in expected polynomial time in $\log p$, with success probability $1 - p^{-\Omega(1)}$.

In particular, every line through f with all direction coefficients nonzero is nearby at all $p-1$ nonzero parameters and has direction weight exactly $2r$. Its exceptional point has distance $1 - \rho - 1/n$, only one coordinate short of maximum distance. If $r \rightarrow \infty$, the separation is $(1 - o(1))\eta$ and $\eta \sim H/b$. Taking $r = \lfloor \log_2 b \rfloor$ gives length $\Theta_\rho(b \log b)$.

Proof. Take $t = 16(r+1)$ and choose n as displayed with the exact-rate parity specified above. Let $a \in \{0, 1\}$ indicate the extra zero coordinate, and let $\epsilon \in \{0, 1\}$ indicate multiplication by X . Set

$$m_0 = (n - a)/2 - 2t - r, \quad D_0 = (\rho n + 1 - \epsilon)/2 - t.$$

These are positive with $D_0 < m_0$ eventually. Since $m_0 = n/2 - O(r)$ and $D_0 = \rho n/2 - O(r)$, the binomial entropy estimate and smoothness of H_2 near the fixed rate give

$$\log_2 L_0 = \frac{nH}{2} + O_\rho(r + \log n) = rb + 2b/5 + o(b).$$

Thus $L_0 \geq p^{r+1/3}$ eventually, while $C^{2r+2} + m_0^2 = p^{o(1)}$. The initial mean missing fraction is at most $2p^{-1/3}$. With probability at least $1 - 2p^{-5/24}$, it is at most $p^{-1/8}$. Whenever $h \leq p^{-1/8}$, for sufficiently large p ,

$$\Phi(h) \leq 4p^{-1/8}h.$$

Consequently the probability that a step fails to reduce the missing fraction by a factor $p^{-1/16}$ is at most $4p^{-1/16}$. A zero missing fraction remains zero. A union bound over t steps gives failure probability at most $2p^{-5/24} + 4tp^{-1/16} = p^{-\Omega(1)}$; on success the final missing fraction is at most $p^{-1/8-t/16} < p^{-r}$, hence is zero.

Fix a reference locator w in the completed family. Its degree is $K+1$ after the parity adjustment, and zero attains its $K+1$ roots, so $\text{dist}(f, C) = (n - K - 1)/n$. The full product image supplies a candidate matching any prescribed nonzero shifts on any subset of padding pairs: extend the unspecified components to arbitrary nonzero values. If z has weight j , this gives $K+1+2j$ agreements. Conversely the received word differs from f on exactly $2j$ positions, so the Hamming-distance inequality gives the matching lower bound. This proves (27).

A codeword within θ at a full-support parameter must agree on all $2r$ padding positions: otherwise changing those positions back to f would give a distance smaller than its established minimum. Distinct full-support parameters have different padding values, so their lists are disjoint. The far point rules out correlated agreement.

Finally $\eta b - H = \Theta_\rho(1/r)$ dominates the $O_\rho(1/b)$ entropy-radius change, proving strict Elias, and

$$\log_2 \left(\frac{n2^{H/\eta}}{p} \right) = \log_2 n - \frac{b}{5(2r+1)} + O_\rho(1/r) \rightarrow -\infty.$$

The construction samples $O(n)$ field elements and forms and evaluates one locator in $O(n^2)$ field operations. It does not recover a nearby witness for an arbitrary supplied parameter. $\square$

Corollary 4.16 (Two far inputs and almost-certain nearby mixtures). *In Corollary 4.15, let $r \geq 2$. When r is even, there are words u, v with*

$$\text{dist}(u, C) = \text{dist}(v, C) = \theta + \frac{r}{2r+1}\eta$$

such that

$$\text{dist}((1-z)u + zv, C) = \theta \quad \text{for every } z \in \mathbb{F}_p \setminus \{0, 1\}.$$

Thus exactly the two endpoints are outside the tested radius, and a uniform affine mixture is nearby with probability $1 - 2/p$. For $r \rightarrow \infty$, both endpoints are almost half a gap outside that radius, although the numerical prescription predicts a nearby fraction tending to zero. All the same rate, domain-length, strict-Elias, and sampling guarantees apply.

More generally, for every divisor $t \geq 2$ of r , the same code has t words $v_1, \dots, v_t$, each at distance

$$\theta + \frac{2r}{2r+1} \left(1 - \frac{1}{t}\right) \eta,$$

whose uniform affine combination is nearby with exact probability

$$\left(1 - \frac{1}{p}\right)^t - \frac{(-1)^t}{p^t} = 1 - O(t/p).$$

Here the coefficients are uniform subject to summing to one. For each fixed batch size t , taking $r \rightarrow \infty$ through multiples of t makes every input almost $(1 - 1/t)\eta$ outside the radius.

Proof. For the first statement, partition the r directions into two equal sets A, B and put $u = f + \sum_{j \in A} e_j$, $v = f + \sum_{j \in B} e_j$. Their affine mixture has parameter coordinates $1 - z$ on A and z on B . At the endpoints exactly $r/2$ coordinates are nonzero; at every other parameter all r are nonzero. The exact distance profile (27) proves every claim about the line. Its $p - 2$ nearby parameters still exceed the prescribed bound by a factor tending to infinity.

For the second statement partition the directions into t groups A_j of size r/t and take $v_j = f + \sum_{i \in A_j} e_i$. Each word has parameter weight r/t , giving the stated distance. An affine combination has coefficient λ_j repeated on A_j , so it is nearby exactly when all t coefficients are nonzero. The number of nonzero-coordinate vectors with sum one is $((p-1)^t - (-1)^t)/p$: this follows by inclusion-exclusion, or by separating zero and nonzero target sums in a recurrence on t . Dividing by the p^{t-1} affine coefficient vectors proves the formula. $\square$

Corollary 4.17 (Linear and multilinear mixtures of far inputs). *Use the t words of Corollary 4.16 and assume $n - K \geq 4r + 1$, as holds throughout its sufficiently large-field asymptotic regime. A uniform unrestricted linear combination is nearby with exact probability*

$$\frac{1}{p} + \left(1 - \frac{1}{p}\right) \left[\left(1 - \frac{1}{p}\right)^t - \frac{(-1)^t}{p^t} \right].$$

If $t = 2^\ell$, index the words by $b \in \{0, 1\}^\ell$ and use the multilinear coefficients

$$\lambda_b(x) = \prod_{j:b_j=0} (1 - x_j) \prod_{j:b_j=1} x_j.$$

For uniform $x \in \mathbb{F}_p^\ell$, the mixture is nearby with exact probability $(1 - 2/p)^\ell$. For every fixed ℓ , all inputs can lie almost $(1 - 2^{-\ell})\eta$ outside the tested radius while this probability tends to one.

Proof. Let s be the sum of the linear coefficients. When $s \neq 0$, scaling by s^{-1} preserves distance to the linear code and produces a uniform affine combination, with the probability already computed. When $s = 0$, the f terms cancel and the mixture is supported on at most $2r$ coordinates. Its distance to the zero codeword is at most $2r/n \leq \theta$. Since s is uniform, the first formula follows. For the multilinear statement the coefficients sum to one. All are nonzero exactly when every x_j avoids zero and one. Apply the exact affine profile and independence of the ℓ parameters. $\square$

Corollary 4.18 (Prescribed exceptional distances on curves). *In the affine space of Corollary 4.15, fix $e \geq 1$, distinct parameters $z_1, \dots, z_t \in \mathbb{F}_p$, and integers $1 \leq a_i \leq r$ with $\sum_i a_i \leq er$. There is a polynomial curve of degree at most e whose exact distance is*

$$\theta + 2a_i/n \quad \text{at } z_i, \quad \theta \quad \text{at every other parameter.}$$

Conversely, every degree-at-most- e curve in this affine space that has any point at distance θ satisfies

$$\sum_{z \in \mathbb{F}_p} (\text{dist}(f_z, C) - \theta) \leq \frac{2er}{n} = e \left(\eta - \frac{1}{n} \right).$$

For lines, two equal endpoint separations cannot exceed half this budget. The two equal separations in Corollary 4.16 attain that bound. This converse is a property of this construction, not a bound for all Reed–Solomon curves.

Proof. Distribute a_i copies of each z_i consecutively around r bins, continuing cyclically between successive parameters. Since $a_i \leq r$, no bin receives the same parameter twice. Since $\sum_i a_i \leq er$, each bin receives at most e parameters. Let its direction coordinate be the product of $z - z_i$ over its assigned parameters, or one if empty. Exactly a_i coordinates vanish at z_i , and none vanish elsewhere. Equation (27) proves the claimed distances. For the converse, each direction coordinate is a polynomial of degree at most e . None is identically zero, since there is a point at distance θ . Each therefore vanishes at at most e field parameters. Sum the zero-coordinate counts and multiply by $2/n$. $\square$

Corollary 4.19 (Interpolation of far inputs). *Let $2 \leq t < p$ divide r , and choose distinct $x_1, \dots, x_t \in \mathbb{F}_p$. The t far words of Corollary 4.16 have a degree-at-most- $(t-1)$ interpolating curve $F(z)$ such that*

$$\text{dist}(F(x_i), C) = \theta + \frac{2r(t-1)}{nt} \quad (1 \leq i \leq t), \quad \text{dist}(F(z), C) = \theta \quad (z \notin \{x_1, \dots, x_t\}).$$

A uniform curve parameter is nearby with probability $1 - t/p$. For fixed t and growing r , all t specified inputs are almost $(1 - 1/t)\eta$ outside the radius. Their equal excess distances attain the degree- $(t-1)$ total budget above.

Proof. Use the Lagrange coefficients $\lambda_j(z) = \prod_{i \neq j} (z - x_i)/(x_j - x_i)$ and set $F(z) = \sum_j \lambda_j(z)v_j$. The coefficients sum to one, are all nonzero away from the interpolation nodes, and form a unit vector at each node. Apply the exact affine distance profile. The total excess distance is $2r(t-1)/n$, as asserted. $\square$

Corollary 4.20 (Witness variation and extension-field challenges). *For a completed line above, a polynomial codeword graph of parameter degree $e \geq 1$ supplies nearby witnesses at at most*

$$\left\lceil \frac{e(n - K - 1)}{2r} \right\rceil$$

parameters. Consequently a polynomial graph selecting a nearby codeword at every nonzero parameter has degree at least $\lceil 2r(p-1)/(n-K-1) \rceil = \Omega_p(p/\log p)$.

Let $E/\mathbb{F}_p$ be any field extension and extend the Reed–Solomon code’s coefficient field to E . In the same affine space, at the same radius, the nearby parameter vectors are exactly $(\mathbb{F}_p^*)^r \subseteq E^r$, and their decoding lists are unchanged. In particular a completed line has exactly the $p-1$ nonzero base-field nearby parameters over E . Every parameter in $E \setminus \mathbb{F}_p$ has the same distance as the exceptional point, namely $1 - \rho - 1/n$.

Proof. Apply Lemma A.8 with $a_0 = K+1$ and $T = K+2r+1$. If a graph supplies a nearby codeword at several parameters, choose those codewords as the selected witnesses in that lemma. The asymptotic degree bound follows from $n = \Theta_\rho(r \log p)$.

The distance of f over E is still $(n-K-1)/n$, by the same root bound. To reach the tested radius after changing at most $2r$ padding positions, every one must change and must match the nearby codeword. That codeword also agrees with f on at least $K+1$ core positions. Interpolation on any K of them forces its coefficients into $\mathbb{F}_p$, since both the evaluation points and the values of f lie in $\mathbb{F}_p$. The padding equalities then force all parameter coordinates into $\mathbb{F}_p^*$. Conversely every such vector already has a base-field nearby codeword, and the same interpolation argument excludes additional extension-field codewords from its list.

Finally fix a line parameter $z \notin \mathbb{F}_p$. There is an $\mathbb{F}_p$ -linear functional $\lambda : E \rightarrow \mathbb{F}_p$ with $\lambda(1) = 1$ and $\lambda(z) = 0$. Applying it coefficientwise to an extension-field codeword shows that every agreement with $f + zg$ is an agreement of a base-field codeword with f . There are at most $K+1$ such agreements. The zero codeword attains $K+1$ on the reference core roots (and the optional common zero coordinate), since g vanishes there. This proves the exact distance. $\square$

The polynomial-degree conclusion concerns algebraic witness selection; it is not a lower bound on decoding running time. The extension-field statement also limits the scope of the example: over a finite extension, the nearby fraction on the line is $(p-1)/|E|$, not $1 - 1/|E|$.

Exact half-rate completion certificates give the following instances. The direction weight is the number of coordinates changed along the line, and the last column bounds $n2^{1/\eta}/p$.

Prime	n	r	Direction weight	Failure <	Prescription <
$2^{61} - 1$	158	1	2	existence	2^{-1}
$2^{127} - 1$	294	1	2	existence	2^{-20}
$2^{521} - 1$	1090	1	2	existence	2^{-147}
$2^{521} - 1$	2142	2	4	existence	2^{-81}
$2^{521} - 1$	2414	2	4	2^{-64}	2^{-26}
$2^{1279} - 1$	7742	3	6	existence	2^{-160}
$2^{1279} - 1$	8014	3	6	2^{-64}	2^{-121}
$2^{1279} - 1$	25970	10	20	2^{-64}	2^{-27}

Rows marked existence use successive averaging; the others have the stated randomized success guarantees. The 61-bit example has distance $78/158$ at zero and exactly $76/158$ at every nonzero parameter. Its line changes only two coordinates. All rows lie strictly below Elias.

The checkers in `research/paired_domain_warp/` bound the recurrence with dyadic or finer fixed-point upper bounds and replay the resulting integer inequalities independently using weaker character and conditioning constants. The search finds sufficient parameters; no optimality is claimed. They separately verify primality, strict Elias, and the numerical comparison. Small-field fixtures verify the entire affine distance profile at all 1369 points of a two-dimensional space over

$\mathbb{F}_{37}$; these fixtures illustrate the geometry and are not claimed to lie below Elias. The stored length-2518 sample has a four-coordinate direction, generator failure below 2^{-88} , and a known nearest codeword, zero, at parameter one. Its coordinates and distances at zero and one are replayed deterministically. The full-profile probability is over the generator's randomness, rather than a deterministic certificate for the individual sample. The following cubic construction separately defeats arbitrary fixed numerical constants c_1, c_2 .

4.3.2 Unambiguous nearby points and an almost maximally far input

Complete nonzero coverage is not needed to refute the numerical bound. A different choice of paired domain makes every nearby point uniquely decodable, while retaining the almost maximally far exceptional point. The following argument uses only polynomial zero bounds.

Lemma 4.21 (A criterion for exact, unambiguous line lists). *Use one padding pair ± 1 and nonzero core representatives $a_1, \dots, a_m$ with distinct squares different from one, and the exact-rate parity construction above, with D selected core orbits. Suppose the only $\varepsilon \in \{-1, 0, 1\}^m$ satisfying*

$$\sum_i \varepsilon_i a_i = 0, \quad \prod_i \left(\frac{1 - a_i}{1 + a_i} \right)^{\varepsilon_i} = 1$$

is the zero vector. Set the direction on the padding to one in the even-locator case, and to x in the case multiplied by X . At agreement threshold $K + 3$, the nearby parameters are exactly

$$-\prod_{i \in I} (1 - a_i^2), \quad |I| = D,$$

and the list at each parameter consists exactly of the candidate codewords from that product's fiber. If these products are distinct, there are exactly $\binom{m}{D}$ nearby parameters and each has a unique nearby codeword. Parameter zero has exact distance $1 - K/n - 1/n$; every nearby parameter has exact distance $1 - K/n - 3/n$.

Proof. The usual locators supply the displayed parameter for every support. Conversely, a nearby codeword must match both padding coordinates and $K + 1$ nonpadding coordinates, because the far-point agreement bound is $K + 1$. Write F for the reference polynomial, including the optional factor X , and let Q be any nearby codeword polynomial. Then $H = F - Q$ is monic of degree $K + 1$, has zero coefficient of X^K , and is the locator of exactly $K + 1$ nonpadding roots. Thus those roots sum to zero.

Remove every sign-orbit whose two elements are roots. Each remaining nonzero root is $\varepsilon_i a_i$ for a nonzero entry of a ternary vector ε . The root-sum condition gives its first equation. The padding equalities give $H(1)/H(-1) = (-1)^{K+1}$. Factoring this ratio into root contributions gives its second equation; a possible zero root contributes one. The hypothesis therefore leaves no single roots. When $K + 1$ is even, zero cannot be a root, and when $K + 1$ is odd, zero must be a root. Consequently H is exactly one of the paired locators in the candidate family, including the optional factor X . This proves the exact list description; distinct products then give uniqueness and the stated count. The root bound at zero and the two-coordinate distance inequality prove the asserted distances. $\square$

Theorem 4.22 (Numerical failure with no local decoding ambiguity). *Fix a rational rate $\rho \in (0, 1)$, $c_1 > 0$, and $0 < c_2 < 3/2$. Over every sufficiently large prime field, there is an exact-rate code of length*

$$n = \frac{1}{2} \log_2 p + O_\rho(1)$$

and a received line with exactly

$$J = 2^{(H_2(\rho)/2 + o(1))n}$$

nearby parameters at radius $\theta = 1 - \rho - 3/n$. Every nearby point has a unique nearby codeword. Parameter zero has distance $1 - \rho - 1/n$, so its separation is $2/3$ of the capacity gap. The radius is strictly below Elias, there is no correlated agreement, and

$$\frac{J}{c_1 n 2^{c_2 H_2(\rho)/(3/n)}} \rightarrow \infty.$$

Uniform sampling of the core constructs the code and line in expected polynomial time in $\log p$, with failure probability at most $p^{-1/2+o(1)}$.

Proof. Choose n with the exact-rate parity above and take $m = (n - a)/2 - 1$, where $a \in \{0, 1\}$ records the optional zero coordinate. Thus $m = \frac{1}{4} \log_2 p + O_\rho(1)$ and $D = \rho m + O_\rho(1)$. First sample the a_i independently and uniformly from $\mathbb{F}_p$. The probability of a zero, a padding collision, or repeated squares is at most $(m^2 + 2m)/p$.

For two distinct supports, equality of their products is a nonzero polynomial equation of degree at most $2m$. A union bound over at most $\binom{m}{D}^2 \leq 4^m$ support pairs bounds the probability of any such collision by $2m4^m/p$.

Fix a nonzero ternary vector with support size u . For $u = 1, 2$, its sum equation is impossible on a valid domain. For $u \geq 3$, change signs of the selected variables and eliminate the last using their sum. The product-ratio equation becomes a nonzero polynomial of degree at most u in the remaining $u - 1$ variables. To see it is nonzero, set all but three original variables to zero and set the remaining three to $X, Y, -X - Y$: the difference of the two products is, up to sign, $2XY(X + Y)$. The probability of the two equations is therefore at most u/p^2 , by the polynomial zero bound. Summing over the ternary vectors gives at most $m3^m/p^2$.

The total bad probability is at most

$$\frac{m^2 + 2m + 2m4^m}{p} + \frac{m3^m}{p^2} = p^{-1/2+o(1)}.$$

Rejection sampling for a valid domain preserves this asymptotic bound. Lemma 4.21 applies on every remaining output. The binomial entropy estimate gives $\log_2 J = nH_2(\rho)/2 + O_\rho(\log n)$, and hence the displayed ratio diverges for $c_2 < 3/2$. Since $(3/n) \log_2 p \rightarrow 6$, the sufficient Elias condition holds with room to spare. The far point rules out correlated agreement. Forming and evaluating the reference locator costs $O(n^2)$ field operations; the sampler need not verify the two combinatorial conditions or recover a witness at a supplied parameter. $\square$

4.3.3 A deterministic family with unique and recoverable witnesses

The unambiguous construction can be made deterministic, with an efficient procedure deciding whether a parameter is nearby and recovering its unique witness. Its core representatives are simply powers of two.

Lemma 4.23 (Decoding products of $4^i - 1$). *For distinct subsets $I \subseteq \{1, \dots, m\}$, the positive integers $T_I = \prod_{i \in I} (4^i - 1)$ are distinct. Given T_I , its support can be recovered using a polynomial number of bit operations in m and $\log T_I$.*

Proof. Write $s = \sum_{i \in I} i$ and $U_I = \prod_{i \in I} (1 - 4^{-i})$, so $T_I = 4^s U_I$. The elementary product inequality gives $2/3 < U_I \leq 1$. For a nonempty support, T_I therefore has exactly $2s$ binary digits, determining s and hence the rational number U_I .

Process indices in increasing order, removing previously selected factors from U_I . At index i , the remaining product is at most $1 - 4^{-i}$ if i is selected. If i is not selected, it is at least

$$\prod_{j=i+1}^m (1 - 4^{-j}) \geq 1 - \sum_{j=i+1}^m 4^{-j} > 1 - \frac{1}{3}4^{-i} > 1 - 4^{-i}.$$

These comparisons determine every membership bit, proving injectivity. They can be implemented with integers: if the remaining integer product is t and the remaining index sum is s , compare $t4^i$ with $(4^i - 1)4^s$. On selecting i , divide t by $4^i - 1$ and subtract i from s . For an arbitrary input, require exact divisibility at every selection and verify the final product and support size. All intermediate integers have polynomial bit length. $\square$

Theorem 4.24 (Deterministic, efficiently decoded counterexamples). *Let $1 \leq D < m$, let p be prime with*

$$p > 2^{D(2m-D+1)},$$

and use core representatives $a_i = 2^i$ for $1 \leq i \leq m$, with padding ± 1 and the exact-rate parity convention above. At threshold $K + 3$, the resulting line has exactly $\binom{m}{D}$ nearby parameters, each with a unique nearby codeword. Parameter zero has distance $1 - K/n - 1/n$, and each nearby parameter has distance $1 - K/n - 3/n$. The code, line, and a procedure deciding membership in the nearby set and recovering the unique witness are deterministic and polynomial in $\log p$.

Consequently, for every fixed rational rate $\rho \in (0, 1)$ and every $c_1 > 0$, $0 < c_2 < 3/2$, all sufficiently large prime fields have such examples at exact rate ρ and length $n = \Theta_\rho(\sqrt{\log p})$, strictly below Elias, with no correlated agreement and

$$\frac{\binom{m}{D}}{c_1 n 2^{c_2 H_2(\rho)/(3/n)}} = 2^{\Omega_{\rho, c_2}(n)}.$$

Proof. Every positive D -subset product is less than

$$\prod_{i=m-D+1}^m 4^i = 2^{D(2m-D+1)} < p.$$

The displayed field bound also gives $p > 2^{m+1}$. No nonzero ternary sum of $2, 4, \dots, 2^m$ can vanish modulo p : its integer absolute value is less than p , and its largest power exceeds the sum of all smaller powers. Thus the signed-relation hypothesis of Lemma 4.21 holds. The nearby labels are

$$z_I = (-1)^{D+1} T_I \pmod{p}.$$

Lemma 4.23 and the no-wrap bound show that these labels are distinct. They also give a decoder: lift $(-1)^{D+1} z$ to $\{0, \dots, p-1\}$, recover and check its support, then output the reference locator minus that support's locator, with the optional factor X . The exact list classification guarantees that rejection means the parameter is outside the tested radius. The degree and distance assertions follow from the same criterion.

For exact rate ρ , choose the admissible n near a sufficiently small constant multiple of $\sqrt{\log_2 p}$. Then $m = n/2 + O(1)$, $D = \rho n/2 + O_\rho(1)$, and the displayed field bound holds. As before, $\log_2 \binom{m}{D} = nH_2(\rho)/2 + O_\rho(\log n)$, which proves the numerical separation for $c_2 < 3/2$. Moreover $(3/n) \log_2 p \rightarrow \infty$, proving strict Elias. The far point rules out correlated agreement. Forming the locators, evaluating them, and performing the integer comparisons use polynomial time in $\log p$. $\square$

For example, exact arithmetic gives the following deterministic rows. Every nearby parameter has a unique witness recoverable by the preceding integer decoder; the last column lower-bounds the ratio to the $c_1 = c_2 = 1$ numerical prescription.

Prime	n	K	D	Nearby count	Ratio $>$
$2^{521} - 1$	74	15	8	$\binom{36}{8}$	1
$2^{1279} - 1$	82	41	21	$\binom{40}{21}$	2^3
$2^{2203} - 1$	106	53	27	$\binom{52}{27}$	2^6
$2^{9689} - 1$	226	113	57	$\binom{112}{57}$	2^{25}

These are proofs for the specified domains, not guarantees over random generation. The checkers in `research/paired_domain_warp/powers_two/` verify primality, the no-wrap bound, strict Elias, and the numerical comparisons with integers, and test support recovery exhaustively on small instances. Uniqueness here concerns every nearby point on this line; it does not assert a global list-size bound for the code.

Proposition 4.25 (An exact two-level distance profile). *In the unadjusted case $n = 2m + 2$, $K = 2D - 1$, assume additionally*

$$p > 2^{m(m+1)/2+2m+2}.$$

Every parameter outside the nearby bank has distance exactly $(n - K - 1)/n$. Thus the line has only two distance levels, and the integer decoder determines which level holds. Both f and $f + g$ are at the far level. The three half-rate rows in the preceding table satisfy this extra bound.

Proof. The zero codeword supplies $K + 1$ reference-core agreements at every parameter, so distance is always at most $(n - K - 1)/n$. Suppose a codeword has at least $K + 2$ agreements. At least one padding position must agree, by the root bound. If exactly one padding position agrees, its $K + 1$ core roots have sum zero; integer dissociation forces them into pairs. Its locator is even, so the other padding position agrees too, placing the parameter in the nearby bank.

It remains to consider two padding agreements and only K known core roots. The difference polynomial is monic of degree $K + 1$, with zero coefficient of X^K . Its last root is therefore minus the sum of those K roots. Cancel complete sign-pairs among the known roots and call the remaining signed powers of two $y_1, \dots, y_u$. Their number is odd. If $u = 1$, the last root completes its pair, again giving a bank locator. For $u \geq 3$, put $S = \sum_i y_i$. Equality of the two padding values would require

$$F = (1 + S) \prod_i (1 - y_i) - (1 - S) \prod_i (1 + y_i) = 0 \pmod{p}.$$

This integer is nonzero. Indeed, writing e_j for the elementary symmetric functions of the y_i gives

$$F/2 = (Se_2 - e_3) + (Se_4 - e_5) + \dots.$$

If the two smallest-magnitude signed powers are $u = \pm 2^a$ and $v = \pm 2^b$, with $a < b$, the monomial with uniquely smallest 2-adic valuation is u^2v in the first parenthesis, with valuation $2a + b$. All other degree-three terms and all higher-degree terms have larger valuation. Thus cancellation is impossible. Finally,

$$|F| \leq 2(1 + |S|) \prod_i (1 + |y_i|) < 2^{m(m+1)/2+2m+2} < p,$$

so F cannot vanish modulo p either. Consequently every improvement over the far distance is already a nearby-bank point. Finally, neither zero nor one is a bank label: each positive product lies strictly between one and $p - 1$. Thus both displayed endpoints are at the far level. $\square$

4.3.4 Unique witnesses with almost the entire gap as separation

Larger prime-order orbits strengthen the separation while preserving uniqueness and efficient witness recovery. The evaluation domain is a union of scaled small orbits, with a bounded number of extra coordinates. It is not a prescribed multiplicative subgroup.

Lemma 4.26 (Rigidity of root sums on scaled prime-order orbits). *Let d be prime, $p \equiv 1 \pmod{d}$, and let $\omega \in \mathbb{F}_p$ have order d . Fix $m \geq 2$ and $0 \leq c < d$, and put*

$$A = 2^{m+1}(d + 2^c).$$

If $p > A^{d-1}$, the following points are distinct:

$$\{2^i \omega^j : 1 \leq i \leq m, 0 \leq j < d\} \cup \{2^{m+t} : 1 \leq t \leq c\} \cup \{\omega^j : 0 \leq j < d\}.$$

Every zero-sum subset of the first two sets is a union of complete d -element orbits from the first set; it contains no extra coordinate.

Proof. Let ζ be a complex primitive d th root of unity. The polynomial $\Phi_d(X) = 1 + X + \dots + X^{d-1}$ is irreducible over $\mathbb{Q}$ by Eisenstein after translating X by one. For an integer polynomial P of degree at most $d-1$, if $P(\omega) = 0$ but $P(\zeta) \neq 0$, its nonzero integer resultant with Φ_d is divisible by p . On the other hand, if the sum of the absolute values of its coefficients is at most A ,

$$|\text{Res}(\Phi_d, P)| = \prod_{j=1}^{d-1} |P(\zeta^j)| \leq A^{d-1} < p,$$

a contradiction.

A difference between two distinct displayed formal points is nonzero over $\mathbb{C}$ and has coefficient sum less than A , so their reductions cannot collide. Now lift a zero-sum subset of nonpadding points to $P(\zeta)$. Its coefficient sum is less than A , hence $P(\zeta) = 0$. Its degree is at most $d-1$, so all its coefficients must be equal. Each coefficient records a sum of distinct powers of two. Uniqueness of binary expansion forces the same selected indices in every orbit position. Thus each core orbit is either fully selected or absent, and an extra coordinate, which occurs only in position zero, is absent. $\square$

Theorem 4.27 (Explicit orbit lines with unique witnesses). *With the preceding notation, take $1 \leq D < m$ and assume also*

$$p > 2^{dD(2m-D+1)/2}.$$

There is a length- $n = d(m+1) + c$, dimension- $K = dD - 1$ Reed-Solomon code and a received line with exactly $\binom{m}{D}$ nearby parameters at radius

$$\theta = 1 - \frac{K}{n} - \frac{d+1}{n}.$$

Every nearby point has a unique codeword, recoverable in polynomial time in $\log p$. The exceptional point has exact distance $1 - K/n - 1/n$; every nearby point has exact distance θ . The direction has weight d , attaining the distance-improvement bound. There is no correlated agreement at the tested threshold. Given ω , construction and decoding are deterministic. An order- d root can be found by a zero-error randomized procedure in expected polynomial time in $\log p$.

Proof. Use the three displayed sets as the core, extra coordinates, and padding, respectively. For D -subsets $I \subseteq \{1, \dots, m\}$ put

$$H_I(X) = \prod_{i \in I} (X^d - 2^{di}).$$

Fix a reference $w = H_{I_0}$, let f be its evaluation word, and let g equal one on the padding and zero elsewhere. The polynomial $w - H_I$ has degree at most $dD - d < K$ and supplies the nearby parameter

$$z_I = -H_I(1) = (-1)^{D+1} \prod_{i \in I} (2^{di} - 1).$$

It matches $dD = K + 1$ core coordinates and all d padding coordinates. The root bound gives exact far agreement $K + 1$ for f ; changing only d coordinates proves the exact nearby distance.

Conversely, any nearby codeword must match all padding coordinates and $K + 1 = dD$ nonpadding coordinates. Its difference from w is monic of degree dD , has zero coefficient of X^{dD-1} , and is the locator of those dD distinct roots. Their sum is zero, so Lemma 4.26 forces them to be exactly D complete core orbits. Thus every nearby codeword belongs to the displayed family.

The positive integer products are less than $2^{dD(2m-D+1)/2} < p$. The proof and greedy decoder of Lemma 4.23 apply with base 2^d in place of four. The normalized product is still greater than $2/3$, and the product of all later factors exceeds the membership threshold. This proves injectivity, the exact count, and unique witness recovery from the signed integer lift of z . Finally, sampling $h \in \mathbb{F}_p^*$ and testing $h^{(p-1)/d} \neq 1$ finds ω in expected $d/(d-1)$ trials. Every accepted root has order d , so this step has no correctness failure event. $\square$

Corollary 4.28 (Unambiguous numerical failure with nearly full separation). *Fix a rational rate $\rho = a/b \in (0, 1)$ in lowest terms and a prime $d \nmid a$. For every sufficiently large prime $p \equiv 1 \pmod{d}$, the construction has exact rate ρ , length $n = \Theta_{\rho,d}(\sqrt{\log p})$, and gap $\eta = (d+1)/n$, with*

$$\text{dist}(f, C) = \theta + \frac{d}{d+1} \eta.$$

Every nearby point has a unique, efficiently recoverable witness. The radius is strictly below Elias and, for every fixed $c_1 > 0$ and $0 < c_2 < 1 + 1/d$,

$$\frac{\binom{m}{D}}{c_1 n 2^{c_2 H_2(\rho)/\eta}} = 2^{\Omega_{\rho,d,c_2}(n)}.$$

In particular, for every fixed $\varepsilon > 0$, infinitely many prime fields admit such counterexamples to the $c_2 = 1$ prescription, with separation at least $(1 - \varepsilon)\eta$ and no decoding ambiguity anywhere on the nearby set.

Proof. Take $n = bt$ with $at \equiv -1 \pmod{d}$, let c be the residue of n in $\{0, \dots, d-1\}$, and put

$$m = (n - c)/d - 1, \quad D = (\rho n + 1)/d.$$

Choosing n near a sufficiently small constant multiple of $\sqrt{\log_2 p}$ ensures both field bounds, with $1 \leq D < m$ eventually. Then

$$\log_2 \binom{m}{D} = \frac{n}{d} H_2(\rho) + O_{\rho,d}(\log n).$$

Comparing this with $c_2 n H_2(\rho)/(d+1)$ proves the ratio. Since $\eta \log_2 p \rightarrow \infty$, strict Elias holds. Choosing a sufficiently large prime d not dividing a gives the final separation claim. There are infinitely many primes congruent to one modulo this d . The gap still tends to zero; no fixed-positive-gap conclusion is asserted. $\square$

Exact half-rate certificates illustrate the tradeoff:

Prime	d	n	K	Separation	Ratio $>$
$2^{9689} - 1$	3	274	137	$3\eta/4$	2^9
$2^{19937} - 1$	5	518	259	$5\eta/6$	2^2
$2^{44497} - 1$	7	908	454	$7\eta/8$	1

Here the ratio compares the exact nearby count $\binom{m}{D}$ with $n2^{1/\eta}$. The respective (m, D, c) are $(90, 46, 1)$, $(102, 52, 3)$, and $(128, 65, 5)$. The accompanying integer certificates verify the field bounds, primality, strict Elias, and these comparisons. Unlike the complete-coverage examples, these lines have a sparse nearby bank. Their witness classification and decoder have no sampling failure.

Smaller finite fields using roots of two. When the field also contains a d th root of two, the same decoder can use the factors $2^i - 1$ instead of $2^{di} - 1$. This reduces the integer product bound substantially. The following statement requires both specified roots in the field; it does not replace the prime-progression hypothesis of Corollary 4.28.

Proposition 4.29 (A root-of-two refinement). *Let d be prime and let $\omega, \alpha \in \mathbb{F}_p$ satisfy $\text{ord}(\omega) = d$ and $\alpha^d = 2$. Take $m \geq 2$, $1 \leq D < m$, $0 \leq c < d$, and put*

$$n = d(m+1) + c, \quad K = dD - 1, \quad A = n2^{\lceil (m+c+1)/d \rceil}.$$

Assume

$$p > A^{d(d-1)}, \quad p > 2^{D(2m-D+3)/2}.$$

Use the core $\{\alpha^i \omega^j : 2 \leq i \leq m+1, 0 \leq j < d\}$, extra coordinates $\alpha^{m+2}, \dots, \alpha^{m+c+1}$, and padding $\{\omega^j : 0 \leq j < d\}$. These points are distinct. The corresponding locator line has exactly $\binom{m}{D}$ nearby parameters at radius $1 - K/n - (d+1)/n$; each has a unique codeword, recoverable in polynomial time in $\log p$. Parameter zero has exact distance $1 - K/n - 1/n$. The construction and decoder are deterministic given the two roots.

Proof. Let $u = 2^{1/d} > 0$ and let ζ be a complex primitive d th root. The field $L = \mathbb{Q}(u, \zeta)$ has degree $d(d-1)$: the degrees d and $d-1$ of its two subfields are relatively prime, so their product both divides and bounds $[L : \mathbb{Q}]$. The ring $\mathbb{Z}[u, \zeta]$ is free on the basis $u^s \zeta^j$ for $0 \leq s < d$, $0 \leq j < d-1$. The given finite-field roots define a ring homomorphism to $\mathbb{F}_p$. Consequently a nonzero element mapping to zero has nonzero integer norm divisible by p , as follows by reducing its multiplication matrix modulo p .

For a lifted subset sum or a difference of two displayed formal coordinates, every conjugate has absolute value at most A . Thus its norm has absolute value at most $A^{d(d-1)} < p$. A modular zero must therefore already be zero in L . Distinct formal coordinates have different positive magnitudes unless their indices agree, in which case their root-of-unity factors are distinct. This proves distinctness after reduction.

For a subset sum of nonpadding coordinates, group exponents of u according to their residue modulo d , using $u^{qd+s} = 2^q u^s$. The powers $1, u, \dots, u^{d-1}$ are linearly independent over $\mathbb{Q}(\zeta)$, so each resulting polynomial in ζ vanishes. Its degree is at most $d-1$, forcing all its coefficients to agree. Within each residue class, each coefficient is a sum of distinct powers of two. Binary uniqueness again forces whole core orbits and excludes all extra coordinates.

The all-codeword classification from Theorem 4.27 now applies, with locators

$$H_I(X) = \prod_{i \in I} (X^d - 2^i), \quad I \subseteq \{2, \dots, m+1\}, \quad |I| = D.$$

Their signed labels lift to positive integers $T_I = \prod_{i \in I} (2^i - 1) < 2^{D(2m-D+3)/2} < p$. Writing $s = \sum_{i \in I} i$, the normalized product satisfies

$$\frac{1}{2} < \prod_{i \in I} (1 - 2^{-i}) < 1,$$

so T_I has exactly s binary digits. At index i , membership is equivalent to the remaining normalized product being at most $1 - 2^{-i}$; the product of all later factors is strictly greater than that threshold. The greedy decoder therefore works in base two. Starting the indices at two is essential to this normalization argument. The remaining degree, count, and distance conclusions are unchanged. $\square$

For a Mersenne prime $p = 2^b - 1$ with $d \neq b$, one can take $\alpha = 2^v$, where $dv \equiv 1 \pmod{b}$. The check $\alpha^d = 2$ is immediate. An order- d root is still required, so $d \mid p - 1$ remains a separate condition.

Exact certificates give the following specific domains. Rates are K/n and need not be one half. Each comparison uses that row's exact binary entropy, evaluated through integer powers rather than floating-point logs.

Prime	d	n	K	Separation	Ratio $>$
$2^{521} - 1$	2	80	29	$2\eta/3$	2^3
$2^{1279} - 1$	3	198	68	$3\eta/4$	2^4
$2^{9689} - 1$	5	905	324	$5\eta/6$	2^{14}
$2^{23209} - 1$	13	3588	1338	$13\eta/14$	2^2
$2^{44497} - 1$	19	7068	2830	$19\eta/20$	1

The ratio is $\binom{m}{D} / (n 2^{H_2(K/n)/\eta})$, with $\eta = (d+1)/n$. The respective (m, D, c) are $(39, 15, 0)$, $(65, 23, 0)$, $(180, 65, 0)$, $(275, 103, 0)$, and $(371, 149, 0)$. All five primes have independent Lucas–Lehmer replays; saved roots, domain distinctness, the two field bounds, strict Elias, and representative witness evaluations are checked independently. These are deterministic specific-domain certificates once the certified roots are supplied. The final row has a uniquely recoverable nearby witness at every accepted parameter, even though the far point is 19/20 of the capacity gap outside the radius and the numerical prescription is too small.

Shorter domains by randomizing the orbit representatives. The preceding decoder uses the integer size of the binary-scaled coordinates. If efficient recovery at an arbitrary supplied parameter is not required, random representatives give shorter domains and a larger nearby bank relative to the field.

Theorem 4.30 (Unique orbit lines of logarithmic length). *Fix a rational rate $\rho = a/b \in (0, 1)$ in lowest terms, a prime $d \nmid a$, and $0 < \alpha < 1$. For every sufficiently large prime $p \equiv 1 \pmod{d}$, there are exact-rate Reed–Solomon codes of length*

$$n = \alpha \log_2 p + O_{\rho,d}(1)$$

and received lines with exactly

$$J = 2^{(H_2(\rho)/d + o(1))n}$$

nearby parameters at radius $\theta = 1 - \rho - (d+1)/n$. Every nearby point has a unique codeword and exact distance θ ; parameter zero has exact distance $1 - \rho - 1/n$. Consequently its separation is

$d/(d+1)$ of the capacity gap. The radius is strictly below Elias, correlated agreement fails, and for every fixed $c_1 > 0$ and $0 < c_2 < 1 + 1/d$,

$$\frac{J}{c_1 n 2^{c_2 H_2(\rho)/((d+1)/n)}} = p^{\Omega_{\rho,d,\alpha,c_2}(1)}.$$

A randomized procedure constructs the code and line in expected polynomial time in $\log p$, with failure probability at most $p^{\alpha-1+o(1)}$. This theorem does not assert efficient witness recovery from an arbitrary nearby parameter.

Proof. Choose n, c, m, D by the exact-rate congruences in Corollary 4.28, now with n near $\alpha \log_2 p$. Fix an order- d root ω . Sample independent uniform $a_1, \dots, a_m, b_1, \dots, b_c \in \mathbb{F}_p$ and use the domain

$$\{a_i \omega^j : 1 \leq i \leq m, 0 \leq j < d\} \cup \{b_1, \dots, b_c\} \cup \{\omega^j : 0 \leq j < d\}.$$

The probability of a collision among its n formal points is at most $\binom{n}{2}/p$.

For $p > d^{d-1}$, no nonempty proper subset of $\{1, \omega, \dots, \omega^{d-1}\}$ sums to zero. Indeed, its lift to $\mathbb{Z}[\zeta_d]$ is nonzero by irreducibility of Φ_d , and its nonzero norm has absolute value at most $d^{d-1} < p$. Thus every subset of the nonpadding domain other than a union of complete core orbits has a sum that is a nonzero linear form in the sampled variables. The probability that any such subset has sum zero is at most

$$\frac{2^{dm+c}}{p} = \frac{2^{n-d}}{p}.$$

For two distinct D -subsets I, J , the polynomial

$$\prod_{i \in I} (1 - a_i^d) - \prod_{j \in J} (1 - a_j^d)$$

is nonzero and has degree at most dD . The probability of any product collision is therefore at most

$$\frac{dD}{p} \binom{\binom{m}{D}}{2}.$$

On every output avoiding these three bad events, the root-locator classification in Theorem 4.27 applies verbatim: every nearby codeword comes from exactly D whole orbits, and the distinct products make all these codewords occur at distinct parameters.

Before conditioning on a distinct domain, the total bad probability is at most

$$B = \frac{\binom{n}{2} + 2^{n-d} + dD \binom{\binom{m}{D}}{2}}{p}. \quad (28)$$

This is $p^{\alpha-1+o(1)}$, since $\binom{m}{D}^2 \leq 2^{2m}$ and $2m \leq n - d$ for $d \geq 2$. Rejecting invalid domains changes the bound by at most the factor $(1 - \binom{n}{2}/p)^{-1} = 1 + o(1)$. The same entropy calculation proves the numerical ratio, and $((d+1)/n) \log_2 p \rightarrow (d+1)/\alpha$ gives strict Elias with room to spare. The sampler checks domain validity and forms the reference polynomial; it does not enumerate subsets or test the two combinatorial good events. $\square$

For example, the exact union bound gives the following half-rate rows. The ratio compares J with $n2^{1/\eta}$; the probability column bounds failure after rejection sampling for distinct coordinates.

Prime	d	n	K	Separation	Ratio $>$	Failure $<$
$2^{521} - 1$	3	466	233	$3\eta/4$	2^{24}	2^{-57}
$2^{521} - 1$	5	468	234	$5\eta/6$	2	2^{-57}
$2^{1279} - 1$	7	1146	573	$7\eta/8$	2^4	2^{-139}

These are guarantees for the sampling procedure. The arithmetic certificates do not deterministically certify the combinatorial good events for a particular sampled domain.

Two padding orbits and a stronger exponent constant. Allowing much larger splitting primes gives a different tradeoff. One can retain unique, efficiently recoverable nearby witnesses and almost full-gap separation even when the numerical prescription's exponent constant is doubled. This construction does not retain logarithmic length in the field size.

Lemma 4.31 (Rigidity for nearby radical radii). *Let d be prime, let q be a power of two with $q - 1 > 3d^{d-1}$, and let $M \geq 1$. Choose an integer $L \geq 1$ with $q^L > 4(dM)^{d-2}$. Write*

$$t = q^{1/d} > 0, \quad R_i = \frac{q^{i+1} - 1}{q^i - 1}, \quad a_i = R_i^{1/d} > 0 \quad (L \leq i < L + M),$$

and let ζ be a complex primitive d th root of unity. Every zero-sum subset of the orbits $a_i\{1, \zeta, \dots, \zeta^{d-1}\}$ is a union of whole orbits. The same holds if some orbits are replaced by single points: no such single point can occur in a zero-sum subset.

Proof. Write $a_i = t + \delta_i$. Since $R_i = q + (q - 1)/(q^i - 1)$, the difference-of-powers identity gives

$$\frac{t}{3d}q^{-i} < \delta_i \leq \frac{t}{d}q^{-i}.$$

Indeed its denominator lies between dt^{d-1} and $d(q + 1)^{(d-1)/d}$; use $q^i - 1 \geq (1 - 1/q)q^i$ for the upper bound and $3(q - 1) \geq q + 1$ for the lower bound.

A selected root sum is $\sum_i a_i C_i$, where C_i is a subset sum of the d th roots of unity. Always $|C_i| \leq d$; if $C_i \neq 0$, its nonzero cyclotomic integer norm gives $|C_i| \geq d^{-(d-2)}$. Similarly, a nonzero $\sum_i C_i$ has absolute value at least $(dM)^{-(d-2)}$, since every other conjugate has absolute value at most dM . In that case

$$\left| \sum_i \delta_i C_i \right| \leq t \sum_{i \geq L} q^{-i} < 2tq^{-L} < \frac{1}{2}t(dM)^{-(d-2)},$$

so it cannot cancel $t \sum_i C_i$. If $\sum_i C_i = 0$, take the least i with $C_i \neq 0$. Its perturbation has absolute value greater than $tq^{-i}/(3d^{d-1})$, while the entire later tail has absolute value at most $tq^{-i}/(q - 1)$, strictly smaller by hypothesis. Thus all C_i must vanish. For prime d , irreducibility of Φ_d then forces each selected subset to be empty or a whole orbit. A singleton cannot qualify. $\square$

Theorem 4.32 (Unique witnesses against a doubled exponent). *Fix a rational rate $\rho = a/b \in (0, 1)$ in lowest terms and a prime $d \nmid a$. For arbitrarily large exact-rate lengths n , and over infinitely many sufficiently large prime fields for each such length, there are Reed–Solomon codes and received lines with gap*

$$\eta = \frac{2d + 1}{n}$$

and exactly $J = 2^{(H_2(\rho)/d + o(1))n}$ nearby parameters. Every nearby point has a unique codeword at exact distance $\theta = 1 - \rho - \eta$, while parameter zero has exact distance

$$1 - \rho - 1/n = \theta + \frac{2d}{2d+1}\eta.$$

The radius is strictly below Elias and correlated agreement fails. For every fixed $c_1 > 0$ and $0 < c_2 < 2 + 1/d$,

$$\frac{J}{c_1 n 2^{c_2 H_2(\rho)/\eta}} = 2^{\Omega_{\rho,d,c_2}(n)}.$$

Given the specified finite-field roots, the code, line, and a decoder for every nearby parameter can be constructed in polynomial time in $n + \log p$. No comparable bound on finding a suitable splitting prime is asserted. In particular, $c_2 = 2$ is compatible with separation at least $(1 - \varepsilon)\eta$ for any fixed $\varepsilon > 0$.

Proof. Choose $n = bu$ with $au \equiv -1 \pmod{d}$, let c be its residue in $\{0, \dots, d-1\}$, and set

$$m = (n - c)/d - 2, \quad D = (\rho n + 1)/d, \quad M = m + c.$$

For large n we have $1 \leq D < m$. Choose q, L as in Lemma 4.31. Use the first m radical orbits as core, the next c radical points as extra coordinates, and the two padding orbits μ_d and $t\mu_d$. These are distinct in characteristic zero, as their magnitudes are distinct except within each orbit.

Reduce without new root relations. Let E be the Galois field generated by ζ, t and all a_i . Its degree is at most $N = (d-1)d^{M+1}$. Put

$$W = \prod_{i=L}^{L+M-1} (q^i - 1), \quad B = (Wn(q+1))^N.$$

All displayed coordinates become algebraic integers after multiplication by W . A nonzero difference of two coordinates or a nonzero subset sum of nonpadding coordinates has, after this multiplication, nonzero integer norm of absolute value at most B : each conjugate has absolute value at most $Wn(q+1)$. Choose a completely split prime $p > B$. The finite-field reductions are therefore distinct and retain exactly the root-sum rigidity just proved. Arbitrarily large such primes exist by the splitting-prime fact used in Lemma 4.2; only infinitude is needed. We may also require

$$p > q^{S_{\max}}, \quad S_{\max} = D(2L + 2m - D - 1)/2.$$

Keep one index-sum class. Among the D -subsets of $\{L, \dots, L + m - 1\}$, choose a sum S of maximum multiplicity. Its class has size

$$J \geq \frac{\binom{m}{D}}{D(m-D)+1}.$$

For every support I put $H_I(X) = \prod_{i \in I} (X^d - R_i)$. Fix any reference support of size D , set $w = H_{I_0}$, and let f be its evaluation word. Set the line direction to zero on the nonpadding coordinates, to q^S on μ_d , and to one on $t\mu_d$. The identities

$$\frac{H_I(1)}{H_I(t)} = q^{\sum_{i \in I} i}, \quad -H_I(t) = \frac{(-1)^{D+1}(q-1)^D}{\prod_{i \in I} (q^i - 1)}$$

show that $w - H_I$ is a nearby codeword whenever $\sum I = S$. Its degree is at most $dD - d < K = dD - 1$, and it matches $K + 1$ core coordinates and all $2d$ padding coordinates. The root bound and the support of the direction give the exact distances.

Conversely, any nearby codeword must match all padding and $K + 1$ nonpadding coordinates. Its monic difference from w has degree dD and zero penultimate coefficient. Root-sum rigidity forces its roots to be D whole core orbits. The two padding values then imply $q^{\sum I} = q^S$ in $\mathbb{F}_p$. Since both integer powers are less than p , the index sums are equal. This classifies every nearby codeword.

Recover the unique witness. For a nonzero parameter z , lift $(-1)^{D+1}(q-1)^D/z$ to $\{0, \dots, p-1\}$. Every genuine nearby label yields $T_I = \prod_{i \in I} (q^i - 1) < q^{S_{\max}} < p$. The base- q greedy decoder from Lemma 4.23 recovers the support uniquely. Check its cardinality and index sum; rejection proves the parameter is outside the tested radius. Dynamic programming finds a largest index-sum class in polynomial time in m ; the integer products, finite-field locators, and decoder use polynomial time in $n + \log p$.

Finally $m = n/d + O_d(1)$ and $D = \rho n/d + O(1)$, so both the lower bound on J and the trivial upper bound $J \leq \binom{m}{D}$ give $\log_2 J = nH_2(\rho)/d + O_{\rho,d}(\log n)$. Comparison with $c_2 nH_2(\rho)/(2d+1)$ proves the numerical ratio. The splitting primes can be chosen arbitrarily large, ensuring strict Elias. Choosing a large prime d not dividing a proves the final separation statement. The gap tends to zero, and the nearby set is sparse. $\square$

A logarithmic-length version over prescribed prime progressions. The large splitting-prime requirement can be removed if efficient recovery from an arbitrary nearby parameter is not required. Regard q as a finite-field parameter, rather than a fixed small integer. Character sums supply enough parameters for which all required roots exist; polynomial zero bounds exclude unintended root sums and label collisions.

Lemma 4.33 (A finite two-orbit existence certificate). *Let d be prime, $1 \leq D < m$, $0 \leq c < d$, and put*

$$M = m + c, \quad n = d(m + 2) + c, \quad K = dD - 1, \quad U = (M + 1)^2, \quad T = M(M + 1)/2,$$

$$A_0 = \binom{m}{D}, \quad R = D(m - D), \quad S_{\max} = D(2m - D + 1)/2.$$

Let $p \equiv 1 \pmod{d}$ be prime with $p > \max\{n, d^{d-1}\}$. Define

$$Q_0 = \frac{p - U}{d^{M+1}} - U(\sqrt{p} + 1),$$

$$B_0 = (T + 1)((2^d - 1)^M - 1) + S_{\max} \binom{A_0}{2} + \frac{R(R + 1)}{2}.$$

If $Q_0 > B_0$, the two-orbit construction has a valid parameter $q \in \mathbb{F}_p$ for which every nearby point has a unique codeword and

$$J \geq \frac{\binom{m}{D}}{D(m - D) + 1}$$

nearby parameters at radius $1 - K/n - (2d + 1)/n$. The far and nearby distances are exact as in Theorem 4.32. This is an existence certificate; it does not specify a successful parameter q or an arbitrary-parameter witness decoder.

Proof. Let $R_i(Q) = (Q^{i+1} - 1)/(Q^i - 1)$ for $1 \leq i \leq M$. Exclude $Q = 0$ and all zeros of their numerators and denominators. There are at most U excluded elements. On the remaining set, the order of q exceeds $M + 1$, the $R_i(q)$ are distinct and nonzero, and none is 1 or q . Thus the two padding orbits and all core/extra orbits will be disjoint whenever the roots exist.

Count parameters with all roots. Fix a character χ of order d . The indicator that $q, R_1(q), \dots, R_M(q)$ are all d th powers is the product of their d -character averages. Every nonprincipal term is a nontrivial character sum. To check this, take the largest i with a nonzero exponent of R_i . At a primitive $(i+1)$ st root of unity, its valuation is that exponent, nonzero modulo d , while every earlier factor is regular and nonzero. If no such i exists, the nonzero exponent of Q has the same property at zero. Here $p > M+1$ ensures separability.

Replace denominator character powers by positive powers modulo d . The resulting polynomial is not a constant times a d th power and has at most U distinct roots in the algebraic closure. The classical multiplicative-character Weil bound [40, Lemma 2], followed by deletion of at most U arguments, bounds the absolute value of each nonprincipal term by $U(\sqrt{p}+1)$. Character averaging therefore gives at least Q_0 parameters with all required roots.

Exclude unintended zero-sum subsets. Fix an order- d root $\omega \in \mathbb{F}_p$. In the formal series field at infinity choose branches $A_i(Q)^d = R_i(Q)$ with

$$\frac{A_i(Q)}{Q^{1/d}} = 1 + \frac{1}{d}Q^{-i} + O(Q^{-i-1}).$$

For constants C_i , not all zero, the function $\sum_i C_i A_i(Q)$ is nonzero. If $\sum_i C_i \neq 0$ its leading term proves this; otherwise the least index with $C_i \neq 0$ contributes the first nonzero term $C_i Q^{1/d-i}/d$. A proper nonempty subset sum of μ_d is nonzero in $\mathbb{F}_p$: the cyclotomic norm argument bounds its possible nonzero integer norm by $d^{d-1} < p$.

For a support of u nonzero coefficients C_i , take the product of $\sum C_i \epsilon_i A_i(Q)$ over all $\epsilon_i \in \mu_d$. To see explicitly that the product is rational, first form it in independent variables Y_i . Invariance under each substitution $Y_i \mapsto \omega Y_i$ forces every variable exponent in every monomial to be divisible by d . Substitute $Y_i^d = R_i(Q)$. Multiplication by $W_U(Q)^{d^u}$, where $W_U(Q) = \prod_{C_i \neq 0} (Q^i - 1)$, clears all denominators. The resulting polynomial is nonzero because every factor is nonzero in the formal series field. Its degree is at most $d^u(T+1)$: $\deg W_U \leq T$, and each radical has pole order $1/d$ at infinity. Every specialization at which any of these root sums vanishes is a zero of this polynomial.

There are $(2^d - 2)/d$ rotation classes of proper nonempty subsets of μ_d ; primality of d makes every such orbit have size d . Empty and full subsets both contribute coefficient zero. Counting coefficient patterns up to these independent rotations gives the bound

$$(T+1) \sum_{u=1}^M \binom{M}{u} \left(\frac{2^d - 2}{d} \right)^u d^u = (T+1)((2^d - 1)^M - 1)$$

on bad parameters. This also controls any chosen singleton extra coordinates, since their nonzero subset coefficients are included. Outside this set, every nonpadding zero-sum subset is a union of whole core orbits and contains no extra coordinate.

Exclude collisions of the labels and index sums. For distinct D -subsets I, J of $\{1, \dots, m\}$,

$$\prod_{i \in I} (Q^i - 1) - \prod_{j \in J} (Q^j - 1)$$

is a nonzero polynomial of degree at most $S_{\max}$ over any field. If the index sums differ, the leading degrees prove this. If they agree, cancel common factors; in the expansions at infinity, the smallest index in the symmetric difference has coefficient 1 or -1 and cannot cancel. A union bound gives at most $S_{\max} \binom{A_0}{2}$ bad parameters. Finally, exclude roots of $Q^h - 1$ for $1 \leq h \leq R$, costing at most $R(R+1)/2$ parameters. Equal D -support index sums are then the only way their powers of q can agree.

Since $Q_0 > B_0$, at least one root-admitting parameter avoids all bad sets. Choose a largest D -support index-sum class and use the two padding directions q^S and one. The all-codeword locator argument of Theorem 4.32 now applies. Distinct product labels give unique nearby codewords. The integer lift decoder is not available here, since the products may wrap modulo p . $\square$

Corollary 4.34 (Doubling survives at logarithmic length). *Fix rational $\rho = a/b \in (0, 1)$ and a prime $d \nmid a$. Put*

$$C = \max\{\log_2(2^d - 1), 2H_2(\rho)\}, \quad \alpha_* = \frac{d}{\log_2 d + C}.$$

For any fixed $0 < \alpha < \alpha_$ and every sufficiently large prime $p \equiv 1 \pmod{d}$, the conclusions of Theorem 4.32 concerning distances, uniqueness, and numerical violation hold at length*

$$n = \alpha \log_2 p + O_{\rho,d}(1),$$

with no witness-recovery guarantee. In particular they beat every fixed $c_1 > 0$ and $0 < c_2 < 2 + 1/d$ by a factor $p^{\Omega(1)}$. Uniform sampling conditioned on the existence of the required roots has failure probability at most

$$p^{-\gamma+o(1)}, \quad \gamma = 1 - \alpha(\log_2 d + C)/d > 0.$$

Obtaining such a sample by rejection takes expected $d^{M+1}(1 + o(1))$ trials, rather than polynomial time in $\log p$. The simple uniform choice $\alpha = 1/8$ gives $\gamma \geq 3/4$.

Proof. Choose the exact-rate congruences as before with n near $\alpha \log_2 p$. Then $m = \alpha(\log_2 p)/d + O_{\rho,d}(1)$, $M = m + O_d(1)$, and $D = \rho m + O_{\rho,d}(1)$. The principal root-admitting count is $p^{1-\alpha \log_2 d/d+o(1)}$, while the character-sum error is $p^{1/2+o(1)}$. Since $C > \log_2 d$, the hypothesis on α makes the former dominate. The combined bad-parameter count is at most $p^{\alpha C/d+o(1)}$: the root-sum term contributes $p^{\alpha \log_2(2^d-1)/d+o(1)}$ and the product-collision term contributes $p^{2\alpha H_2(\rho)/d+o(1)}$. Thus Lemma 4.33 applies, with the displayed conditional failure ratio. The root-admitting count is $p/d^{M+1}(1 + o(1))$, proving the rejection cost. The same entropy comparison proves the numerical ratio. Finally $\eta \log_2 p \rightarrow (2d + 1)/\alpha > 1$, giving strict Elias. For $\alpha = 1/8$, use $C \leq d$ and $\log_2 d \leq d$. $\square$

Corollary 4.35 (A tradeoff up to exponent constant three). *In either two-orbit construction, one may instead take $K = dD - s$ for any $1 \leq s < d$. The nearby count and uniqueness conclusions persist, with*

$$\eta = \frac{2d + s}{n}, \quad \text{dist}(f, C) = 1 - \rho - s/n = \theta + \frac{2d}{2d + s}\eta.$$

The numerical prescription is violated exponentially for every fixed $c_2 < 2 + s/d$. Thus any fixed $0 < c_2 < 3$ is compatible with unique nearby witnesses and separation greater than two thirds of the gap. Efficient witness recovery remains available in the integer construction; the function-field construction retains its logarithmic length and existence guarantees.

Proof. All constructed differences $w - H_I$ have degree at most $dD - d < K$. For any codeword, the monic difference from w still has degree $dD = K + s$ and zero penultimate coefficient. The same root-sum argument classifies every codeword with $dD + 2d$ agreements. The far word has exactly dD agreements, so both displayed distances follow. For an exact rational rate $\rho = a/b$, choose $n = bu$ with $au \equiv -s \pmod{d}$; this is possible when $d \nmid a$. The support count still has logarithm $nH_2(\rho)/d + O(\log n)$, which gives the threshold $c_2 < (2d + s)/d$. Take $s = d - 1$ and then a sufficiently large prime d for the final claim. In the logarithmic-length version the strict Elias condition still holds; in the integer version choose the splitting prime sufficiently large. $\square$

Remark 4.36 (The dimension cutoff in the integer construction). For $D \geq 2$, lowering K further to $dD-d$ leaves at most one nearby parameter at the same agreement threshold $dD+2d$. Indeed the sums $\sum_{i \in I} R_i$ of distinct D -supports are distinct. To see this over the rationals, write $R_i = q+(q-1)/(q^i-1)$ and take the least index in the symmetric difference. Its nonconstant term exceeds the sum of all later terms, since $(q-1)^2 > q$. Multiplication by W clears denominators, and the nonzero difference has absolute value at most $Wn(q+1) < p$; thus distinctness survives reduction. The coefficient of X^{dD-d} in H_I is $-\sum_{i \in I} R_i$. For $w - H_I$ to have degree below $dD-d$, this sum must equal the reference support's sum, forcing that support alone. The all-codeword classification excludes any other nearby witnesses. This is a limitation of further dimension reduction in this construction, not an upper bound for other constructions.

Lemma 4.37 (A concentrated index-sum class). *Let $A = \binom{m}{D}$ and $V = D(m-D)(m+1)$. The largest index-sum class of D -subsets of $\{1, \dots, m\}$ has size J satisfying*

$$J \geq \left\lceil \frac{A}{\sqrt{V+1}} \right\rceil.$$

The same bound holds for any translated interval of indices. In particular $J = \Omega(A/m^{3/2})$ at fixed density D/m .

Proof. For a uniform D -subset, its index sum S has variance $V/12$. This follows by writing S as the sum of index-weighted selection indicators: their means are D/m and distinct pairs have joint mean $D(D-1)/(m(m-1))$. Let U be independent and uniform on $[-1/2, 1/2]$. The density of $Y = S + U$ is at most $\mu = J/A$, and its variance is $(V+1)/12$. Any real random variable with density at most μ has variance at least $1/(12\mu^2)$: at distance t from its mean, its tail probability is at least $1 - 2\mu t$, and hence

$$\text{Var}(Y) \geq \int_0^{1/(2\mu)} 2t(1 - 2\mu t) dt = \frac{1}{12\mu^2}.$$

Consequently $J^2(V+1) \geq A^2$, proving the claim. Translation changes all index sums by the same amount. $\square$

Finite existence certificates against $c_2 = 2$. Table 1 gives exact half-rate certificates for Lemma 4.33. Every row has positive $Q_0 - B_0$, is strictly below Elias, and beats the prescription even with $c_1 = 1, c_2 = 2$. The last row places the exceptional point $86/87$ of the gap outside the tested radius. These are existence certificates for a suitable parameter and domain; no particular domain is specified by the table.

The certificates use $(m, D, c) = (171, 87, 0), (285, 144, 1), (1325, 664, 3), (2361, 1182, 5), (1895, 949, 17),$ and $(4419, 2211, 41)$, respectively. They replace $\sqrt{p}$ by its integer ceiling and check all bounds using integer arithmetic. Lucas–Lehmer establishes each prime; $(2d+1)(b-1) > n$ gives a strict sufficient Elias certificate. Using $J_0 = \lceil \binom{m}{D} / \sqrt{D(m-D)(m+1)+1} \rceil$ from Lemma 4.37, a displayed exponent w is certified by

$$J_0^{2d+1} > n^{2d+1} 2^{2n+w(2d+1)}.$$

An independent replay verifies these inequalities. Small-field checks separately enumerate every parameter in two fixtures over the field of order 1,000,003, verify the excluded-parameter counts, and classify every potentially nearby codeword on selected lines by interpolation. Integer fixtures also check actual radical domains, the conservative norm bound, and greedy witness recovery. These smaller fixtures test the geometry and proofs; they are not additional violations of the numerical prescription.

$p = 2^b - 1$: b	d	n	K	separation/ η	ratio $>$
521	2	346	173	4/5	2^9
1279	3	862	431	6/7	2^{13}
9689	5	6638	3319	10/11	2^{85}
23209	7	16546	8273	14/15	2^{119}
44497	19	36060	18030	38/39	2^9
216091	43	190144	95072	86/87	2^6

Table 1: Unique-witness two-orbit existence certificates. The ratio is $J/(n2^{2H_2(1/2)/\eta})$, with $\eta = (2d+1)/n$. The table does not claim efficient sampling or decoding.

A length-226 half-rate certificate. An exact index-sum count gives a shorter example over the 338-bit prime

$$p = 204255 \cdot 2^{320} + 1, \quad n = 226, \quad K = 113.$$

Primality has the compact certificate $11^{(p-1)/2} \equiv -1 \pmod{p}$: every prime divisor of p must then be 1 modulo 2^{320} , and hence exceed $\sqrt{p}$. For $m = 111$, $D = 57$, the index-sum class $S = 3192$ contains exactly

$$J = 441734689284929317660654122848$$

supports. This is the coefficient of $y^{57}x^{3192}$ in $\prod_{i=1}^{111}(1+yx^i)$, replayed by two independent recurrences. The finite existence inequality is positive, so some domain over this explicit field has these uniquely nearby parameters, gap $5/226$, and far separation $4\eta/5$. Exact integer comparisons give

$$\frac{J}{n2^{2/\eta}} > \frac{119}{100}, \quad \frac{J}{n2^{1/\eta}} > 2^{45}.$$

The second inequality concerns the original $c_1 = c_2 = 1$ prescription. This certificate specifies the field and the support class, but not a successful domain parameter q .

At rate $19/44$, a separate certificate has $p = 121 \cdot 2^{320} + 1$, $n = 220$, $K = 95$, and $\eta = 1/44$. Here $(m, D, S) = (108, 48, 2616)$ and the exact class contains 31495815243247545087807036124 supports. The ratio exceeds $21/20$ for $c_2 = 2$ and 2^{43} for $c_2 = 1$, with $c_1 = 1$ in both cases. Primality is certified by $3^{(p-1)/2} \equiv -1 \pmod{p}$. The non-half-rate comparison is also entirely integral, using

$$2^{nH_2(K/n)} = \frac{n^n}{K^K(n-K)^{n-K}}.$$

The same existence and domain-specification qualification applies.

4.3.5 Deterministic banks over the 31-bit Mersenne prime

The preceding small-bit full-coverage certificates use averaging. At $p = 2^{31} - 1$, direct enumeration instead certifies the following large banks for concrete domains. These counts do not depend on a sampling success event.

Proposition 4.38 (Explicit M31 banks, verified by enumeration). *Over $\mathbb{F}_{2^{31}-1}$ there are the following Reed–Solomon codes and received lines. Put $\rho = K/n$, $\eta = 3/n$, and $B = n2^{H_2(\rho)/\eta}$.*

n	K	Candidate supports	Certified nearby parameters	Comparison
58	27	40,116,600	39,763,294	$> B$
62	31	145,422,675	140,916,078	$> B$
64	31	300,540,195	281,212,602	$> B$
68	33	1,166,803,110	897,817,238	$> 2B$

Parameter zero has exact distance $(n - K - 1)/n$, and every certified parameter has exact distance $(n - K - 3)/n$. The direction changes only two coordinates. All four radii are strictly below Elias and have no correlated agreement. Thus all four violate the $c_1 = c_2 = 1$ numerical prescription; the last also violates $c_1 = 2, c_2 = 1$.

Construction and computational verification. Let $m = n/2 - 1$ and $D = (K + 1)/2$. Use the first m entries of the following list as core representatives, with ± 1 as the padding pair:

848979899, 738121439, 699266884, 128913239, 806814741, 629503248,
179005809, 134797481, 207154190, 1023774369, 339671428, 287462841,
393054946, 39961481, 1026875673, 850367968, 913861226, 264277071,
48753053, 60664611, 1054291373, 882987304, 134769212, 35446086,
175650241, 263723351, 718238076, 6429846, 428509650, 268381772,
130012564, 236350235, 900863422.

The domain consists of both signs of these representatives and of 1. Set $w(X) = \prod_{i=1}^D (X^2 - a_i^2)$, let f be its evaluation word, and let g equal one on ± 1 and zero elsewhere. A D -subset I supplies the codeword $w - H_I$ at parameter

$$z_I = - \prod_{i \in I} (1 - a_i^2) \neq 0.$$

It has $K + 1$ core and two padding agreements. The root bound at zero and the two-coordinate distance inequality prove both exact distances.

Two independent programs in `research/paired_domain_warp/m31_exact/` enumerate the image of these subset products. One joins product lists from two halves of the core and uses Mersenne reduction; the other recursively enumerates complementary supports, using inverse factors and ordinary modular reduction. Their bitmaps give the same four counts in the table. The verifier also recovers explicit supports and codeword coefficients for selected parameters. The counts are exact for this candidate family and lower-bound the total number of nearby parameters on the line.

The comparison with $c_1 B$ is checked by the integer inequality

$$J^3 K^K (n - K)^{n-K} > c_1^3 n^{n+3},$$

where J is the certified count. In the half-rate row, $B = 62 \cdot 2^{62/3} < 103199661$. Finally $\eta \log_2 p > 90/n > 1$ proves strict Elias in every row. The code's prescribed domain is the displayed paired set; a power-of-two length in the third row does not make it an FFT-domain example. $\square$

Exact local lists in the half-rate row. For $n = 62, K = 31$, the signed-relation condition of Lemma 4.21 is also checked exhaustively. Splitting its 30 variables into two groups gives $3^{15} = 14348907$ assignments per group; the only matching pair is the empty assignment. Two independent implementations agree. The lemma therefore excludes every nearby codeword outside the enumerated locator family. Two further independent counts give its complete list-size distribution:

List size	0	1	2	3	4	5
Parameters	2,006,567,569	136,494,714	4,337,028	83,442	891	3

In particular, even the 136494714 uniquely nearby parameters exceed $B < 103199661$. Every point on this line has at most five nearby codewords at the tested radius. This is an exact statement about the displayed line, not a global list-size upper bound for the code.

4.4 A cubic change of domain covers every nonzero parameter

The outside-collision estimate can be improved by changing the domain before adding the padding. Matching three times as many integer moments allows a two-parameter cubic change of coordinates. The residual collision equations are absolutely irreducible surfaces; averaging their point counts makes the outside evaluation images occupy almost the entire field. The padding can then cover every nonzero parameter.

Lemma 4.39 (Irreducibility of a cubic collision equation). *Let S, T be disjoint sets of distinct field elements with $|S| = |T| = u \geq 1$. Then*

$$F(X, U, V) = \prod_{a \in S} (X - a^3 - Va^2 - Ua) - \prod_{a \in T} (X - a^3 - Va^2 - Ua)$$

is absolutely irreducible, of degree at most u .

Proof. The case $u = 1$ is a nonconstant linear polynomial. Work over the algebraic closure and homogenize using W . The forms $\ell_a = X - aU - a^2V - a^3W$ define planes in $\mathbb{P}^3$ in general position: every four are independent by the Vandermonde determinant. Write the two families as A_i, B_j . Each base line $A_i = B_j = 0$ is generically smooth on $\prod_i A_i - \prod_j B_j = 0$, and hence belongs to a unique irreducible component. The lines indexed by (i, j) and (i', j) meet at $A_i = A_{i'} = B_j = 0$. No other plane passes through this point. The derivative of the second product is nonzero there, so the point is smooth and the two lines belong to the same component. The same argument works for (i, j) and (i, j') . Consequently all base lines lie on one component G .

Restrict G to $A_i = 0$. The restriction is nonzero, since A_i does not divide the difference of products. It vanishes on the u distinct lines cut out by the B_j , so their product divides it. Thus $\deg G \geq u$, the degree of the original homogeneous polynomial. The latter is irreducible; the smooth base points also exclude a repeated factor. Dehomogenizing at $W = 1$ preserves irreducibility, because W is not a factor. The resulting polynomial is nonconstant. $\square$

Lemma 4.40 (Cubic deformation and complete padding). *Let $p > 2m^2$ be prime. On an m -point seed domain in $\mathbb{F}_p$, suppose $L \geq 2$ distinct t -subsets contain a common anchor and have equal power sums of orders $1, \dots, 3s$. Assume $s \geq 0$ and $K = t - s - 1 \geq 1$. Set*

$$N = m - 1, \quad R = p - N, \quad U_0 = p - 1, \quad B = \frac{p + (t - 1)(t - 2)\sqrt{p} + 3t^4}{p - \binom{m}{2}}, \quad M = \frac{LR}{R + (L - 1)B}.$$

If $1 \leq q \leq R$ and

$$U_0(1 - M/U_0)^q < 1, \tag{29}$$

there is a length- $n = N + q$, dimension- K prime-field Reed–Solomon code and a nonconstant received line such that every nonzero parameter has at least t agreements with the code, whereas parameter zero has maximum agreement exactly $t - 1$. There is no correlated agreement at threshold t .

Proof. Translate the seed so the anchor is zero; equal power sums are preserved. For $(a, b) \in \mathbb{F}_p^2$ use the map $\phi_{a,b}(x) = x^3 + bx^2 + ax$. At least $p^2 - \binom{m}{2}p$ maps are injective on the seed: each collision of two distinct nodes specifies one affine line in the parameter plane. The transformed supports have equal power sums through order s , because $\phi_{a,b}(x)^j$ has degree at most $3s$ for $j \leq s$. Newton's identities therefore give equal first s locator coefficients. For transformed locators H_i , choose one H_0 and put $w = H_0/X$, $P_i = (H_0 - H_i)/X$. These are distinct polynomials with $\deg P_i < K$, and each agrees with the degree- $(t-1)$ polynomial w on exactly $t-1$ coordinates of the transformed core.

For a support pair, cancel its common nodes before forming the difference of products. At an outside point, collision of P_i and P_j is exactly the equation in Lemma 4.39. Its degree is at most t . The affine hypersurface estimate of Cafure and Matera [39, Corollary 5.6] bounds its number of triples (X, a, b) by

$$p^2 + (t-1)(t-2)p^{3/2} + 3t^4p.$$

The hypothesis $p > 2m^2$ implies the characteristic condition in that estimate; a linear residual equation has exactly p^2 points. Summing over support pairs and averaging over injective maps gives one map with at most $\binom{L}{2}B$ outside pair collisions.

For this map, the Cauchy-Schwarz image calculation in Lemma 4.8 gives average outside image size at least M . The differences $P_i(x) - w(x)$ are all nonzero outside the core. Choose the q points with largest images and independent uniform nonzero directions there, with direction zero on the core. The expected number of uncovered nonzero parameters is at most the left-hand side of (29). Some choice leaves none. The global degree of w bounds its agreement with every codeword by $t-1$, attained on the core. This proves the exact far-point distance and excludes correlated agreement at threshold t . $\square$

Theorem 4.41 (A full punctured line at fixed rate). *Fix rational $\rho \in (0, 1)$, constants $c_1, c_2 \geq 1$, and an integer $u > c_2$. For every sufficiently large prime p , there are a code of exact rate ρ , length $n = \Theta(\log p)$, and a received line such that all $p-1$ nonzero parameters are within radius $\theta = 1 - \rho - \eta$, while parameter zero has distance exactly*

$$\theta + \eta/u, \quad \eta = u/n.$$

The radius is strictly below Elias, there is no correlated agreement, and $c_1 n 2^{c_2 H_2(\rho)/\eta} = o(p)$.

Proof. Choose rational $\rho < \beta < 1$, close enough to ρ that, with $\alpha = \rho/\beta$, there is a constant C satisfying

$$\frac{1}{\alpha H_2(\beta)} < C < \frac{u}{c_2 H_2(\rho)}.$$

This is possible because $u > c_2$. Put $s = u - 1$, $b = \log_2 p$, and choose $n = Cb + O(1)$ through common denominator multiples for ρ, α . Set $K = \rho n$, $m = \alpha n$, and $t = K + s + 1$. The anchored t -subsets of an integer interval of length m number $2^{mH_2(t/m) - O(\log m)}$. Their first $3s$ integer moments have only $m^{O_s(1)}$ possible vectors. Thus some anchored class has $L \geq p^{1+\epsilon}$ for a fixed $\epsilon > 0$.

In Lemma 4.40, $B = 1 + o(1)$ and $M/(p-1) \rightarrow 1$. Take $q = n - m + 1 = \Theta(\log p)$. Then $(p-1)(1 - M/(p-1))^q \rightarrow 0$, proving complete nonzero coverage. The zero parameter is $1/n = \eta/u$ beyond the radius. Finally $\eta \log_2 p \rightarrow u/C > c_2 H_2(\rho) \geq H_2(\rho)$ gives strict Elias and the stated numerical comparison. $\square$

For every selection of nearby witnesses, an affine codeword graph contains at most $n - K - u + 1 \leq n$ selected pairs, by Lemma A.8. Thus complete coverage also violates the numerical line-decodability threshold at concurrency n .

For $c_1 = c_2 = 1$, one may take $u = 2$: every nonzero point is nearby while the remaining point is $\eta/2$ outside. This is a shrinking-gap construction on cubic images of intervals with added coordinates. It does not prove superlinear nearby counts at a fixed positive gap or transfer the result to a prescribed evaluation subgroup. Its global lists remain large: on the original core, match $3(s+1)$ moments for t -subsets, and then apply the same cubic map. The count still exceeds p , and the $s+1$ preserved moments give degree-below- K candidates at threshold t . Thus it does not replace the separate actual-list-size separation. The displayed line necessarily has local ambiguity as well: at any one padding coordinate, $L > p$ candidates produce only $p-1$ nonzero labels.

Exact finite certificates give the following half-rate instances. Every nonzero parameter is nearby in every row; the last column compares $n2^{1/\eta}/p$ with the $c_1 = c_2 = 1$ prescription.

Prime	n	K	Agreement threshold	Far separation	Prescription $<$
$2^{127} - 1$	214	107	109	$\eta/2$	2^{-12}
$2^{1279} - 1$	3834	1917	1921	$\eta/4$	2^{-308}
$2^{9689} - 1$	18162	9081	9083	$\eta/2$	2^{-593}

The checker `verify_finite_certificate.py` in `research/cubic_domain_warp/` verifies primality, strict Elias, the list count and (29). The independent `audit_finite_independent.py` uses elementary anchored moment boxes and the weaker point-count bound of [39, Theorem 5.2], with no construction-helper imports. Both replays verify the missing-label expectation is less than one. For the two larger fields, $M/(p-1) > 3/4$ and $2q \geq \log_2(p+1)$ suffice; for the 127-bit field, each replay checks the rational power directly. The finite geometry checker verifies general position, smooth base intersections, small-field point counts, and the degree and moment identities on all 5692 injective cubic maps of a 17-node fixture. These are existence certificates for the large examples, not enumerations of their domains, directions or nearby codewords.

4.4.1 Sampling the code and line without enumerating a moment class

The large moment class and the largest evaluation images need not be found explicitly. Random choices produce the code and received line with a quantitative success guarantee.

Proposition 4.42 (A direct randomized construction). *Use the parameters of Lemma 4.40, with seed $\{0, \dots, m-1\}$ and anchor zero. Put*

$$Q = \binom{m-1}{t-1}, \quad \mu = Q/\mathcal{V},$$

where $\mathcal{V}$ is an upper bound on the number of first-3s integer binomial-moment vectors of the anchored supports. Choose an anchored support uniformly, an injective cubic map uniformly, q distinct outside coordinates uniformly, and independent uniform nonzero padding directions. Let w be the locator of the chosen transformed support after removing its anchor, and output $w + zg$. For every $L_* \geq 2$ and $\gamma > 0$, the probability that some nonzero parameter is not nearby is at most

$$\frac{L_*}{\mu} + \frac{\delta}{\gamma} + U_0 \left(\frac{\gamma}{1+\gamma} \right)^q, \quad \delta = \frac{U_0}{L_*} + \frac{U_0 B}{R} - 1. \quad (30)$$

Every output has the exact far-point distance and absence of correlated agreement in Lemma 4.40.

Proof. The chance that the sampled support belongs to a moment class of size less than L_* is at most $\mathcal{V}L_*/Q = L_*/\mu$, by summing the sizes of such classes. Equal binomial moments imply equal

power sums through the same order. Condition on any class of size $L \geq L_*$. For an injective map, let $r_{x,v}$ be the multiplicity of the nonzero difference value $v = P_i(x) - w(x)$, and put

$$E = \sum_{x \text{ outside}} \sum_{v \neq 0} r_{x,v}^2, \quad \Delta = \frac{U_0 E}{RL^2} - 1.$$

Cauchy–Schwarz gives $\Delta \geq 0$. The collision average proved in Lemma 4.40 gives

$$\mathbb{E}E \leq LR + L(L-1)B, \quad \mathbb{E}\Delta \leq \frac{U_0}{L} + \frac{U_0(L-1)B}{RL} - 1 \leq \delta.$$

Hence $\Pr(\Delta > \gamma) \leq \delta/\gamma$. If $\Delta \leq \gamma$, the average outside image size is at least $L^2 R/E \geq U_0/(1+\gamma)$.

For sampling distinct coordinates, the mean product of their nonnegative missing fractions is at most the corresponding power of the population mean. Indeed, replacing two entries by their average increases each elementary symmetric mean: with all other entries fixed, the only changing term is a nonnegative multiple of their product. Iterating this averaging proves the inequality. The expected number of uncovered labels is therefore at most $U_0(\gamma/(1+\gamma))^q$, and the probability of any uncovered label is at most this expectation. Combining the three bad events proves (30). The degree of w and its selected roots give the far-point claim for every output, independently of the coverage event. $\square$

For the parameters in Theorem 4.41, the elementary moment ranges give $\mu \geq p^{1+\epsilon}$ for a fixed $\epsilon > 0$. Taking $L_* = p^{1+\epsilon/2}$ in the analysis gives $\delta = p^{-\Omega(1)}$. Since $q = \Theta(\log p)$, a sufficiently small fixed γ makes the last term in (30) $p^{-\Omega(1)}$ as well. Thus the algorithm succeeds with probability $1 - p^{-\Omega(1)}$. It does not need to compute L_* , the moment class, or any image size. Rejection sampling makes the cubic map uniform among injective maps in expected $O(1)$ trials. Forming and evaluating w takes $O(n^2)$ field operations, so the code and line are output in expected polynomial time in $\log p$. This does not give a procedure for recovering a nearby codeword at a supplied nonzero parameter.

For an exact finite example, take $p = 2^{521} - 1$, $(n, K, m, t) = (990, 495, 900, 497)$, $s = 1$, $L_* = p^{2^{150}}$, and $\gamma = 2^{-8}$. The exact bound in (30) is less than 2^{-137} . An independent integer replay bounds each of its three terms by 2^{-130} , and hence their sum by 2^{-128} . The radius is strictly below Elias, the far point is $\eta/2$ outside, and the numerical prescription fraction is below 2^{-16} . The programs `generate_sample.py` and `verify_sample.py` in `research/cubic_domain_warp/` output and replay a stored sample. Its domain, line values, and exact far-point properties are checked deterministically. Complete coverage has the stated probability guarantee over the generator’s randomness; it is not deterministically certified for the individual stored sample.

4.5 Elementary padded-interval refinements

The following bounds use the original interval core directly. They give additional finite examples and tradeoffs without the cubic deformation.

Theorem 4.43 (An arbitrarily large fraction of nearby parameters). *Fix rational $0 < \rho < \beta < 1$, put $H = H_2(\beta)$, and write $b = \log_2 p$. For every sufficiently large prime p , there is a Reed–Solomon code of exact rate ρ , length*

$$n = \left(\frac{2\beta}{\rho H} + o(1) \right) b, \quad \eta = \frac{\rho H / \beta + o(1)}{\sqrt{b \log_2 b}},$$

and a received line with at least

$$\left(1 - \exp\left(-\frac{\beta - \rho}{\rho \beta (1 - \beta)} \right) - o(1) \right) p$$

nearby parameters at gap η . The parameter-zero word has distance $1 - \rho - \eta + 1/n$. No pair of codewords has correlated agreement at the threshold, and the nearby radius is strictly below Elias. Consequently, for any fixed $\delta > 0$, choosing β sufficiently close to 1 gives at least $(1 - \delta)p$ nearby parameters at length $\Theta_{\rho,\delta}(\log p)$ and gap $\Theta_{\rho,\delta}(1/\sqrt{\log p \log \log p})$.

Proof. Put $\alpha = \rho/\beta$. Choose n through common denominator multiples for ρ, α , with $n = 2b/(\alpha H) + O(1)$, and set

$$m = \alpha n, \quad K = \rho n, \quad k = K + 1, \quad s = \left\lfloor 2(1 - \varepsilon_b) \sqrt{b/\log_2 b} \right\rfloor, \quad t = k + s,$$

where $\varepsilon_b \rightarrow 0$ sufficiently slowly, for instance $\varepsilon_b = \sqrt{\log_2 \log_2 b / \log_2 b}$ for large b . The anchored integer moment count gives a class of size L with

$$\log_2 L \geq m H_2(t/m) - \frac{s^2}{2} \log_2(m/s) - O_{\rho,\beta}(s^2 + s \log m + \log m) \geq b + \Omega(\varepsilon_b b).$$

Thus $L/p \rightarrow \infty$. Remove the anchor and divide as above, giving $N = m - 1$ old points and a received polynomial of degree $A - 1$, where $A = t$. Add $q = n - m + 1$ coordinates, and use $d = K - 1$ in the collision bound. Balanced occupancies give

$$\frac{2T}{L^2} = d - \frac{(A - 1)^2}{N} + o(1) = (\rho(1 - \beta) + o(1))n, \quad \frac{M}{p} = \frac{1 + o(1)}{\rho(1 - \beta)n}.$$

Apply Lemma 4.11. Since $q/n \rightarrow 1 - \rho/\beta$ and $(p - 1)/p \rightarrow 1$, substitution into (24) gives the asserted constant fraction and the exact distance at zero. The exact gap is $(s + 1)/n$, and $\eta \log_2 p \rightarrow \infty$ verifies strict Elias. Finally the exponent in the displayed fraction tends to $-\infty$ as $\beta \rightarrow 1$ with ρ fixed. $\square$

Thus almost every parameter can be nearby while no common agreement explanation exists. Here “almost every” means any fixed fraction below one; the length and gap constants depend on that fraction. The gap still shrinks. This does not prove superlinear growth at one fixed positive gap, or uniqueness of nearby codewords. At these parameters, the proposed numerical threshold $c_1 n 2^{c_2 H_2(\rho)/\eta}$ is only $p^{o(1)}$ for fixed c_1, c_2 , whereas the actual count is a constant fraction of p .

Theorem 4.44 (Separation by a fixed fraction of the capacity gap). *Fix rational $\rho \in (0, 1)$, a desired error fraction $\delta > 0$, and finite $c_1, c_2 \geq 1$. There is a constant $\kappa > 0$ and, for every sufficiently large prime p , a code of exact rate ρ and a received line with*

$$n = \Theta(\log p), \quad \eta = \Theta(1/\log p), \quad J \geq (1 - \delta)p,$$

whose parameter-zero word has distance

$$1 - \rho - \eta + \kappa \eta.$$

The nearby radius is strictly below Elias, there is no correlated agreement at the threshold, and $c_1 n 2^{c_2 H_2(\rho)/\eta} = o(p)$. The constants, including κ , may depend on ρ, δ, c_1, c_2 . This does not assert density tending to one with the same positive κ .

Proof. Choose rational $\rho < \beta < 1$ so that

$$\lambda := \frac{\beta - \rho}{\rho\beta(1 - \beta)} > \ln(2/\min\{\delta, 1\}).$$

Put $\alpha = \rho/\beta$, choose a fixed $C > 1/(\alpha H_2(\beta))$, and choose a fixed positive integer s satisfying

$$\frac{s+1}{C} > \max\{H_2(\rho), c_2 H_2(\rho)\}.$$

Write $b = \log_2 p$. Choose $n = Cb + O(1)$ through common denominator multiples for ρ, α , and set $m = \alpha n$, $K = \rho n$, $k = K + 1$, and $t = k + s$. The anchored support count, divided by the number of first- s moment vectors, gives

$$\log_2 L \geq m H_2(t/m) - O_s(\log m) = \alpha C H_2(\beta) b - O_{\rho, \beta, s, C}(\log b) > b + \Omega(b).$$

Thus $L/p \rightarrow \infty$. Apply exactly the anchored core and far-point padding argument of Theorem 4.43, now with this fixed s . Its collision calculation is unchanged to leading order, so $J/p \geq 1 - e^{-\lambda} - o(1) > 1 - \delta$ eventually. The exact gap is $\eta = (s+1)/n$, and the parameter-zero word is $1/n = \eta/(s+1)$ beyond the nearby radius. Set $\kappa = 1/(s+1)$. The choice of s gives strict Elias and

$$c_1 n 2^{c_2 H_2(\rho)/\eta} = p^{c_2 H_2(\rho)C/(s+1)+o(1)} = o(p).$$

The global degree bound excludes correlated agreement as before. $\square$

The dependence of the far-point separation on the desired density can be quantified. For fixed ρ, c_2 , as $\delta \downarrow 0$ the construction permits

$$\kappa = \Omega_{\rho, c_2} \left(\frac{\log \log(1/\delta)}{\log(1/\delta)} \right). \quad (31)$$

Indeed, first take $\delta = 2^{-v}$ with integer $v \geq 2$. Choose a power of two q at least $\max\{2/(1-\rho), 2\rho(v+1)/(1-\rho)\}$, and set $\beta = 1 - 1/q$, $C = 2q/(\rho \log_2 q)$, and $s+1 = \lfloor Cc_2 \rfloor + 2$. Since $H_2(\beta) > (\log_2 q)/q$, the support-count condition holds. Also $(s+1)/C > c_2 \geq c_2 H_2(\rho)$ and $\lambda > (1-\rho)q/(2\rho) \geq v+1$, so $e^{-\lambda} < \delta/2$. Taking the smallest such power of two gives $q = \Theta_\rho(v)$ and $\kappa = 1/(s+1) = \Omega_{\rho, c_2}(\log v/v)$. Rounding an arbitrary δ down to the next dyadic value changes only constant factors. This quantifies the tradeoff across separate families; it still does not give density tending to one at fixed positive κ .

For a concrete interpretation, take $\rho = 1/2$, $\beta = 15/16$, $C = 6$, and $s = 6$. Then $\alpha = 8/15$ and $C\alpha H_2(\beta) > 1$, since $15^{15} < 2^{59}$. This gives

$$n = (6 + o(1)) \log_2 p, \quad \eta = \frac{7/6 + o(1)}{\log_2 p}, \quad \frac{J}{p} \geq 1 - e^{-224/15} - o(1) > 0.999999 \quad (32)$$

for sufficiently large p . Yet the line contains a word at distance $1/2 - \eta + \eta/7$, and the $c_1 = c_2 = 1$ numerical prescription has fraction $p^{-1/7+o(1)} \rightarrow 0$. Thus a point can lie a fixed fraction of the capacity gap outside the radius even when more than 99.9999% of the line is nearby.

There are exact finite certificates of this behavior. Over $p = 2^{4423} - 1$, the parameters

$$n = 26520, \quad K = 13260, \quad A = 13267, \quad \eta = 7/n$$

give $J/p > 0.99999974$, a parameter-zero word exactly $\eta/7$ outside the radius, and numerical prescription fraction below 2^{-619} . A larger separation is also possible: choosing $\beta = 3/4$, $C = 15/8$, and $s = 1$ gives asymptotic nearby fraction at least $1 - e^{-8/3} > 0.93$ with separation $\eta/2$. For $p = 2^{2203} - 1$, the finite parameters $(n, K, A) = (4128, 2064, 2066)$ give $J/p > 0.93155073$ and prescription fraction below 2^{-126} , with exact separation $\eta/2$. The checker `verify_gap_scale_far.py` in `research/fixed_gap_padding/` verifies these statements with exact arithmetic, including primality and strict Elias. The domains are padded intervals; these statements do not transfer automatically to a prescribed evaluation subgroup.

Theorem 4.45 (Nearby density tending to one). *Fix rational $\rho \in (0, 1)$ and a real $c > H_2(\rho)$. Write $b = \log_2 p$ and $\ell = \log_2 b$. For every sufficiently large prime p , there is a Reed–Solomon code of exact rate ρ and a received line with parameters*

$$n = \left(\frac{\sqrt{2}}{c} + o(1) \right) \frac{b^{3/2}}{\sqrt{\ell}}, \quad \eta = \frac{c + o(1)}{b},$$

for which the number J of nearby parameters satisfies

$$\frac{J}{p} \geq 1 - \exp \left[- \left(\frac{1 - \rho}{2\sqrt{2}c} - o(1) \right) \sqrt{b\ell} \right]. \quad (33)$$

The parameter-zero word has distance $1 - \rho - \eta + 1/n$. No pair of codewords has correlated agreement at the threshold, and the nearby radius is strictly below Elias. In particular, $J/p \rightarrow 1$ and $n = o(p)$.

Proof. Choose n a denominator multiple for ρ within $O(1)$ of $\sqrt{2}b^{3/2}/(c\sqrt{\ell})$, put $K = \rho n$, and set

$$s = \lfloor cn/b \rfloor, \quad h = \lceil (4 + \varepsilon_b)b/\ell \rceil, \quad k = K + 1, \quad t = k + s, \quad m = t + h,$$

where $\varepsilon_b = \sqrt{\log_2 \ell / \ell}$. For large p these satisfy $1 \leq k < t < m < n < p$. The anchored t -subsets of an m -point integer interval number $\binom{m-1}{h}$. Since

$$\log_2((m-1)/h) = \ell/2 + (\log_2 \ell)/2 + O_{\rho,c}(1),$$

the bound $\binom{m-1}{h} \geq ((m-1)/h)^h$ gives

$$\log_2 \binom{m-1}{h} \geq (2 + \varepsilon_b/2)b - O_{\rho,c}(b/\ell).$$

The first s binomial moments have a number of possible vectors with logarithm at most

$$\frac{s^2}{2} \log_2(m/s) + O(s^2 + s \log m) = b + O_{\rho,c}(b/\ell + \sqrt{b\ell}).$$

Thus an anchored class has $\log_2 L \geq b + \Omega(\varepsilon_b b)$, so $L/p \rightarrow \infty$.

Remove the anchor and divide, using $N = m - 1 = K + s + h$ old coordinates, threshold $A = t$, and $q = n - N$ padding coordinates. With $d = K - 1$,

$$d - \frac{(A-1)^2}{N} = h - s - 1 - \frac{h^2}{K + s + h} = (1 + o(1))h.$$

The balanced collision count therefore gives $M/p = (1 + o(1))/h$. Apply Lemma 4.11. Since $1 - x \leq e^{-x}$,

$$\frac{J}{p-1} \geq 1 - e^{-qM/(p-1)}, \quad \frac{qM}{p-1} = \left(\frac{1 - \rho}{2\sqrt{2}c} + o(1) \right) \sqrt{b\ell}.$$

The loss from replacing $p-1$ by p is $1/p$ and is absorbed by the $o(1)$ in (33). The lemma also gives the exact far-point distance and excludes correlated agreement. The exact gap is $(s+1)/n$. Finally $\eta b \rightarrow c > H_2(\rho)$ gives strict Elias. $\square$

For any fixed c_1, c_2 , choosing $c > \max\{H_2(\rho), c_2 H_2(\rho)\}$ makes the proposed threshold

$$c_1 n 2^{c_2 H_2(\rho)/\eta} = p^{c_2 H_2(\rho)/c + o(1)} = o(p),$$

whereas the nearby fraction tends to one. This does not assert that all parameters are nearby, and the gap is not fixed. This family also has very large off-line lists. Matching one additional moment on the core loses only $O(\sqrt{b \log b})$ bits from the preceding count and gives more than p codewords near another received word at the same threshold. Thus the comparison here is with the proposed numerical list prescription, not with the code's actual maximum list size; Section 5 addresses the latter distinction.

At rate $1/2$, exact finite certificates give the following contrast. The last column is the finite prescription $n2^{1/\eta}/p$, with $c_1 = c_2 = 1$ and without an unspecified remainder.

Prime	n	K	A	Proven $J/p >$	Prescription $<$
$2^{31} - 1$	92	46	50	0.74303810	2^{-1}
$2^{61} - 1$	216	108	113	0.91025173	2^{-10}
$2^{127} - 1$	468	234	240	0.96160200	2^{-40}
$2^{521} - 1$	2800	1400	1411	0.99485642	2^{-255}

Every row lies strictly below Elias, has a parameter-zero word exactly one coordinate beyond the radius, and has no common agreement pair at the threshold. These are coding-count comparisons on padded interval domains, not protocol-security estimates or prescribed-subgroup results. The exact checker `research/fixed_gap_padding/verify_far_point_padding.py` verifies the density and prescription inequalities. The related checker `verify_near_unit_density.py` in that directory also certifies nearby fractions above 0.99999545, 0.99999976, and 0.99999998 for primes $2^{2203} - 1$, $2^{3217} - 1$, and $2^{4423} - 1$, respectively. Finite list counts use the stronger of the binomial-moment box and Gram bounds; the asymptotic proof needs only the box bound. All certificates use Lucas–Lehmer primality, integer Elias inequalities, and rational lower bounds on the union size.

A separate arithmetic audit, `audit_density_independent.py` in the same directory, rederives all four displayed bounds and the 2203-bit bound without importing the construction checker. The separate small-field replay also verifies the far point and the complete agreement profile.

A far point with a growing coordinate separation. The one-coordinate separation in Theorem 4.45 can be increased, while retaining nearby density tending to one. The convergence estimate becomes weaker. We give the counting argument explicitly because it distinguishes relative distance from the number of coordinates by which a word misses the radius.

Lemma 4.46 (Several padding matches). *Suppose a polynomial w of degree $D \geq K \geq 2$ and L distinct polynomials P_i of degree below K each have exactly D agreements on a core of size N . Put $R = p - N$, $U = p - 1$, $d = K - 1$, and*

$$T = d \binom{L}{2} - \mathcal{P}(LD, N).$$

Here $\mathcal{P}$ is the balanced collision count in Lemma 4.8. Let $1 \leq r \leq q \leq R$, and define

$$v_a = \frac{\binom{r}{a} \binom{q-r}{r-a}}{\binom{q}{r}}, \quad B = v_0 + \sum_{a=1}^r v_a U^a \left(\frac{1}{L} + \frac{2T(d-1)_{a-1}}{L^2(R)_a} \right),$$

where $(x)_a$ denotes a falling factorial and out-of-range binomial coefficients are zero. There is an extension to $n = N + q$ distinct domain points and a line $w + zg$, with $g = 0$ on the core, having at least $\lceil U/B \rceil$ nearby nonzero parameters at agreement threshold $D + r$. At parameter zero the maximum agreement is exactly D , and there is no correlated agreement on $D + r$ coordinates.

Proof. Choose the q new points uniformly without replacement from outside the core, and choose their directions independently and uniformly from $\mathbb{F}_p^*$. For fixed $z \neq 0$, let X count pairs (i, S) , where S is an r -subset of the new positions on which $P_i = w + zg$. Every $P_i - w$ has all its roots in the core. Consequently $\mathbb{E}X = L \binom{q}{r} / U^r$.

Let c_{ij} count outside roots of $P_i - P_j$. The root bound and the balanced core count give $c_{ij} \leq d$ and $\sum_{i < j} c_{ij} \leq T$. Two r -subsets whose intersection has size a contribute probability U^{-2r+a} if $i = j$, and $U^{-2r+a} (c_{ij})_a / (R)_a$ otherwise. Hence

$$\frac{\mathbb{E}X^2}{(\mathbb{E}X)^2} = \sum_{a=0}^r v_a U^a \left(\frac{1}{L} + \frac{1}{L^2} \sum_{i \neq j} \frac{(c_{ij})_a}{(R)_a} \right) \leq B,$$

using $(c_{ij})_a \leq c_{ij}(d-1)_{a-1}$ for $a \geq 1$. Cauchy–Schwarz gives $\Pr(X > 0) \geq 1/B$. Average over the U nonzero parameters to obtain one domain and direction with the claimed count. The global polynomial w has degree D , so its maximum agreement with the code is at most D , attained by the selected polynomials on the core. This also excludes correlated agreement at the larger threshold. $\square$

The common core also gives a uniform bound

$$c_{ij} \leq e := d - \max\{0, 2D - N\} \quad (34)$$

whenever there are at least two selected polynomials. Indeed, their core agreement sets intersect in at least $2D - N$ positions. For $e \geq 1$, one may replace $(d-1)_{a-1}$ by $(e-1)_{a-1}$ in Lemma 4.46. The same proof also permits the simpler bound

$$\frac{\mathbb{E}X^2}{(\mathbb{E}X)^2} \leq \frac{U^r}{L} + \left(\frac{U}{R-r+1} \right)^r \exp\left(\frac{r^2 e}{q-r+1} \right). \quad (35)$$

To see this, let Z be the intersection size of two independent uniform r -subsets of a q -set, so $\Pr(Z = a) = v_a$. Bound $(c_{ij})_a \leq e^a$, and use

$$\mathbb{E}e^Z = \sum_{j=0}^r (e-1)^j \binom{r}{j} \frac{(r)_j}{(q)_j} \leq \left(1 + \frac{r(e-1)}{q-r+1} \right)^r \leq \exp\left(\frac{r^2 e}{q-r+1} \right).$$

Corollary 4.47 (An unbounded number of coordinates outside). *Fix rational $\rho \in (0, 1)$ and $c > H_2(\rho)$. With $b = \log_2 p$ and $\ell = \log_2 b$, there are codes and lines of exact rate ρ satisfying*

$$n = \left(\frac{\sqrt{2}}{c} + o(1) \right) \frac{b^{3/2}}{\sqrt{\ell}}, \quad \eta = \frac{c + o(1)}{b}, \quad \frac{J}{p} \geq 1 - O_{\rho,c}(\ell^{-1/2}).$$

The parameter-zero word is exactly r/n beyond the nearby radius, where $r = \lfloor b^{1/6} \rfloor \rightarrow \infty$. The radius is strictly below Elias and there is no correlated agreement at the threshold. The relative separation r/n still tends to zero. For every choice of nearby witnesses, each affine codeword graph contains at most $\lfloor (n - A + r)/r \rfloor = O_{\rho,c}(n/b^{1/6}) = o(n)$ of the selected pairs, where $A = (\rho + \eta)n$. More generally, a codeword graph of degree at most $e \geq 1$ contains at most $\lfloor e(n - A + r)/r \rfloor$ pairs.

Proof. Take n, K, s as in Theorem 4.45, but now put

$$\varepsilon = \ell^{-1/4}, \quad h = \lceil (3 + \varepsilon)(r + 1)b/\ell \rceil, \quad t = K + 1 + s, \quad m = t + h.$$

Since $r \sim b^{1/6}$, we have $h = o(n)$ and $\log_2((m-1)/h) = \ell/3 + (\log_2 \ell)/2 + O_{\rho,c}(1)$. The anchored support count therefore has logarithm at least $(1+\varepsilon/3)(r+1)b - O_{\rho,c}(rb/\ell)$. The moment-vector count still has logarithm $b + O_{\rho,c}(b/\ell + \sqrt{b\ell})$. Thus $L \geq p^r 2^{\Omega(\varepsilon b)}$. Remove the anchor as before, setting $D = K + s$, $N = D + h$, and $q = n - N$. The threshold is now $A = D + r$, so $\eta = (s+r)/n \sim c/b$. The outside-pair cap in (34) is $e = h - s - 1 \sim h$, and

$$\frac{r^2 e}{q - r + 1} = O_{\rho,c}(\ell^{-1/2}), \quad \left(\frac{U}{R - r + 1} \right)^r = 1 + o(\ell^{-1/2}).$$

Together with $U^r/L \leq 2^{-\Omega(\varepsilon b)}$, the bound (35) gives a second-moment ratio at most $1 + O_{\rho,c}(\ell^{-1/2})$. The averaging argument proves the nearby count, and the degree argument gives the exact far-point distance and excludes correlated agreement. Finally $\eta b \rightarrow c > H_2(\rho)$ gives strict Elias. Apply Lemma A.8 with $a_0 = D = A - r$ and $T = A$ to obtain the graph-intersection bounds. $\square$

Thus the far point need not merely miss the radius by one coordinate. This remains a shrinking-gap statement with vanishing relative separation. As before, choosing $c > \max\{H_2(\rho), c_2 H_2(\rho)\}$ makes the numerical prescription $o(p)$ despite nearby density tending to one. The checker `verify_multi_match_far_padding.py` in `research/fixed_gap_padding/` exhausts 55,000 choices of domain and direction over $\mathbb{F}_{11}$ and verifies the exact first and second moments for $r = 1, 2, 3$, as well as the uniform pair cap and 1968 hypergeometric identities. These small fixtures check the mechanism, not Elias.

The relative separation may vanish arbitrarily slowly. The parameter choice in Corollary 4.47 gives a faster approach to density one. A more slowly growing seed surplus permits a larger separation relative to η . The following is another consequence of the same padding lemma, with explicit dependence on a freely chosen schedule.

Corollary 4.48 (Slowly vanishing relative separation). *Fix rational $\rho \in (0, 1)$ and $c > H_2(\rho)$. Write $b = \log_2 p$, and let $s = s(b)$ be any positive integer function satisfying*

$$s \rightarrow \infty, \quad s^{5/3} \log_2 b = o(b).$$

For every sufficiently large prime, there are a code of exact rate ρ and a received line with

$$n = (1/c + o(1))sb, \quad \eta = (c + o(1))/b, \quad \frac{J}{p} \geq 1 - \frac{3c/(1-\rho) + o(1)}{\log_2 s}.$$

The parameter-zero word is exactly $r/n = \kappa\eta$ outside the radius, where

$$r = \lfloor s^{1/3} \rfloor \rightarrow \infty, \quad \kappa = \frac{r}{s+r} \sim s^{-2/3}.$$

The radius is strictly below Elias, and there is no correlated agreement. Every affine codeword graph contains only $o(n)$ selected witnesses, for every selection of nearby witnesses. If also $c > c_2 H_2(\rho)$, the numerical prescription $c_1 n 2^{c_2 H_2(\rho)/\eta}$ is $o(p)$.

Proof. Choose $n = sb/c + O(1)$ through denominator multiples for ρ , and put

$$K = \rho n, \quad h = \left\lceil \frac{3(r+1)b}{\log_2 s} \right\rceil, \quad D = K + s, \quad t = D + 1, \quad m = t + h, \quad N = m - 1.$$

Here $h = o(n)$ and

$$\log_2(N/h) = \frac{2}{3} \log_2 s + \log_2 \log_2 s + O_{\rho,c}(1).$$

Thus the anchored support count has logarithm $\log_2 \binom{N}{h} \geq h \log_2(N/h) \geq (2 - o(1))rb$. The first s binomial moment ranges have total logarithm $O(s^2 \log m) = O(s^2 \log b) = o(rb)$, by the schedule condition. The anchored class therefore has $\log_2 L \geq (2 - o(1))rb$.

Remove the anchor and use Lemma 4.46 with $q = n - N$ new coordinates. The uniform pair cap is $e = h - s - 1 \sim h$, and

$$\frac{r^2 e}{q - r + 1} = \frac{3c/(1 - \rho) + o(1)}{\log_2 s}.$$

In (35), $U^r/L \leq 2^{-(1-o(1))rb}$ and $(U/(R - r + 1))^r = 1 + o(1/\log s)$, since n, r are polynomial in b . The second-moment ratio is consequently at most $\exp((3c/(1 - \rho) + o(1))/\log_2 s)$. Inverting and averaging gives the stated nearby count. The global degree- D word gives exact far separation r/n ; $\eta = (s + r)/n$ gives the displayed value of κ and strict Elias. The incidence bound in Lemma A.8 gives at most $\lfloor (n - D)/r \rfloor = o(n)$ witnesses on each graph. Finally the numerical prescription is $p^{c_2 H_2(\rho)/c + o(1)} = o(p)$ under the stated condition. $\square$

If one only requires a one-coordinate far point, the nearby density can approach one much faster. Under the stronger schedule condition $s^2 \log_2 b = o(b)$, keep $n = sb/c + O(1)$ but take $r = 1$ and $h = \lceil 2b/\log_2 s \rceil$. Then

$$\kappa = \frac{1}{s + 1}, \quad \frac{J}{p} \geq 1 - \exp\left[-\left(\frac{1 - \rho}{2c} - o(1)\right) s \log_2 s\right].$$

Indeed, the support logarithm is at least $(2 - o(1))b$, the moment cost is $o(b)$, and the averaged image bound satisfies $M/p \sim 1/h$. Lemma 4.11 gives the displayed exponent from $q/h \sim (1 - \rho)s \log_2 s/(2c)$. The $1/p$ loss is negligible because $s \log s = o(b)$. This choice preserves a larger relative separation than the earlier near-unit-density parameter choice, at the cost of a slower approach to density one.

In particular, apply Corollary 4.48 to *any* prescribed function $\varepsilon(p) > 0$ tending to zero. Taking

$$s = \lfloor \min\{\varepsilon(p)^{-3/4}, b^{1/4}\} \rfloor$$

satisfies the schedule condition and gives $\kappa > \varepsilon(p)$ eventually. Hence requiring a point at least $\varepsilon(p)\eta$ outside the radius does not prevent nearby density tending to one. This corollary uses a vanishing separation fraction. The cubic-domain construction in Section 4.4 gives complete nonzero coverage with a fixed positive separation fraction. The qualitative statement can also be obtained by diagonalization from fixed-fraction families; the corollary supplies explicit lengths and growing coordinate deficits. For example, $s = \lfloor \log_2 b \rfloor$ gives $n \sim b \log_2 b/c$, $\kappa \sim (\log_2 b)^{-2/3}$, and $J/p \geq 1 - O(1/\log_2 \log_2 b)$.

These codes still have global lists larger than p . On the core, take $(D + r)$ -subsets and match $s + r$ moments. Their support count has logarithm $(2 - o(1))rb$, while the moment cost is still $O((s + r)^2 \log b) = o(rb)$; their common locator prefix yields degree-below- K candidates at the desired threshold. Extend the received word over the padding. Thus this remains a comparison with the numerical prescription, not an actual-list-size separation.

The checker `verify_slow_separation.py` audits sufficient integer parameter inequalities and an exact prime-field instance; the general schedule is justified by the preceding proof.

For Lemma 4.46, exact finite certificates also give the following half-rate examples. Here r is the number of coordinates by which the parameter-zero word misses the radius; the last column is the same numerical prescription $n2^{1/\eta}/p$, with $c_1 = c_2 = 1$ and no unspecified remainder.

Prime	n	K	A	r	Proven $J/p >$	Prescription $<$
$2^{127} - 1$	694	347	356	2	0.25107817	2^{-40}
$2^{127} - 1$	830	415	426	3	0.02273637	2^{-40}
$2^{521} - 1$	5840	2920	2936	2	0.59499256	2^{-100}
$2^{521} - 1$	8140	4070	4091	3	0.28599616	2^{-100}

All four radii are strictly below Elias. The certificate `verify_multi_match_finite.py` in the same directory verifies primality, the anchored Gram list count, the rational second-moment bound, and the displayed inequalities using exact arithmetic. It also certifies two- and three-coordinate far points over $2^{61} - 1$ and $2^{1279} - 1$. These are existence statements on padded interval domains, not prescribed-domain or protocol-security estimates.

Exact finite comparisons. For a seed with parameters (m, k, t) and $s = t - k$, the same proof uses $N = m - 1$, $K = k - 1$, $A = t$, $q = s + 1$, and $d = k - 2$. The Gram bound supplies L_0 supports, and anchor averaging gives $L = \lceil tL_0/m \rceil$. Exact evaluation of (23) yields the following certificates. The last column is a strict lower bound on $\log_2(J/(n2^{H_2(K/n)/\eta}))$, where $\eta = (A - K)/n$.

Prime	n	K	A	Nearby labels J	Excess bits
$2^{31} - 1$	72	12	15	381222100	6.73559
$2^{61} - 1$	155	35	40	$> 3.2687 \cdot 10^{17}$	27.01590
$2^{127} - 1$	273	70	76	$> 1.6423 \cdot 10^{37}$	78.16602
$2^{521} - 1$	1065	381	393	$> 3.1200 \cdot 10^{155}$	422.98809

All radii are strictly below Elias. These compare coding counts with the prescription $c_1 = c_2 = 1$; they are not protocol-security estimates. The padded domains are not prescribed subgroups. The exact checker `research/fixed_gap_padding/verify_average_padding.py` verifies rational Gram bounds, collision arithmetic, outward logarithm intervals, Lucas–Lehmer primality, and integer Elias inequalities. Two small-field fixtures also exhaust all 5202 offset tuples and 167042 codeword pairs. The finite search is not an optimality proof.

The far-point variant is checked in `research/fixed_gap_padding/verify_far_point_padding.py`: over $\mathbb{F}_{17}$, it exhausts all 65536 nonzero padding directions and then the full agreement profile and all codeword pairs of a selected line. The parameter-zero word on that line has maximum agreement $A - 1$; 15 of its 16 nonzero parameters are nearby.

4.6 An elementary low-rate boundary case

For comparison, if the rate may tend to zero, even the constant code admits a line with exactly one far point. The construction below uses only elementary counting and works on every domain of the specified size, with a field polynomial in the length.

Theorem 4.49 (A punctured nearby line for the constant code). *For every finite $c_1, c_2 \geq 1$, there are a fixed integer $D \geq 2$ and a constant $a > 1$ such that, for every sufficiently large prime p , every dimension-one Reed–Solomon code of length $n = \lfloor p^{1/a} \rfloor$ has the following properties, regardless of its evaluation domain. Put $\rho = 1/n$, $\eta = D/n$, and $\theta = 1 - \rho - \eta$. There is a nonconstant received line $f + zg$ for which*

$$\text{dist}(f + zg, C) \leq \theta \quad \text{for every } z \neq 0, \quad \text{dist}(f, C) = \theta + \frac{\eta}{D}.$$

There is no correlated agreement at the threshold, θ is strictly below Elias, and $c_1 n 2^{c_2 H_2(\rho)/\eta} = o(p)$. Here $n = o(p)$ and $p = \Theta(n^a)$, but $\rho \rightarrow 0$.

Proof. Choose an integer $D > c_2$ and a real $1 + c_2/D < a < 2$. Some monic degree- D polynomial f over $\mathbb{F}_p$ has $\Omega_D(p)$ distinct values with exactly D preimages. Indeed, associate each D -subset of $\mathbb{F}_p$ with its monic locator F , and group the locators by $f = F - F(0)$. There are at most p^{D-1} groups, so one contains at least $\binom{p}{D}/p^{D-1} = \Omega_D(p)$ locators. Within a group, distinct locators give distinct full fibers of f .

Select $L = \lfloor n/(2D) \rfloor$ such fibers as the core, of size $N = DL$, and take their L values as constant codewords. On the core put $g = 0$; add $q = n - N$ distinct outside coordinates. At every outside point x , the L differences between the selected constants and $f(x)$ are distinct and nonzero. Choose $g(x)$ independently and uniformly in $\mathbb{F}_p^*$ on the new coordinates. Each fixed $z \neq 0$ misses all the selected values at every new coordinate with probability

$$\left(1 - \frac{L}{p-1}\right)^q.$$

If it does not miss, one constant has its D core agreements and one new agreement. The expected number of uncovered nonzero parameters is therefore at most

$$(p-1) \left(1 - \frac{L}{p-1}\right)^q \leq p \exp\left(-\frac{qL}{p-1}\right) < 1,$$

since $qL/(p-1) = \Omega_D(n^{2-a})$. Some direction thus makes every nonzero parameter nearby at threshold $A = D + 1$. At zero, the degree of f bounds the agreement with any constant by D , attained on the selected fibers. This gives the exact far-point distance and excludes correlated agreement at $D + 1$.

The sufficient Elias inequality has logarithm

$$\log\left(\frac{p^D A^A (n-A)^{n-A}}{n^n}\right) = (aD - (D+1)) \log n + O(1) > 0.$$

The small-rate entropy expansion gives $c_1 n 2^{c_2 H_2(1/n)/\eta} = n^{1+c_2/D+o(1)} = o(p)$. Finally, the constant code is unchanged by a coordinate permutation. Relabeling the constructed words therefore transfers the statement to any prescribed evaluation domain of the same size. $\square$

For $c_1 = c_2 = 1$, take $D = 2$, $a = 5/3$, and $n = \lfloor p^{3/5} \rfloor$. An explicit core polynomial is

$$f(X) = X^2 - (N+1)X, \quad N = 2\lfloor n/4 \rfloor.$$

The core $\{1, \dots, N\}$ consists of $N/2$ full fibers, namely the pairs $\{x, N+1-x\}$. The unique far point is $\eta/2$ beyond the radius, where $\eta = 2/n$. The exact checker `verify_constant_code_density.py` in `research/fixed_gap_padding/` certifies all $p-1$ nonzero parameters for $p = 2^{61} - 1$, $n = 104159249330$, and for two larger primes. It also verifies the fiber-count identity by exhaustive small-field enumeration. The large-field results are existence certificates, not enumerations of those enormous lines.

The constant code's actual maximum list size at this radius is exactly $\lfloor n/(D+1) \rfloor$: distinct constant values use disjoint agreement sets, and assigning blocks of $D+1$ positions attains the bound. Thus its actual list size times n is $\Theta(n^2) > p$ in this construction. This is not an actual-list counterexample. Its additional feature is that it works on every domain of the given size, whereas the fixed-rate construction in Theorem 4.41 uses chosen domains.

5 Small lists do not bound the number of nearby line points

The previous constructions produced large lists. We now show a different obstruction: even a code with small lists can have an affine line with exponentially many nearby points, each with a unique nearby codeword. The degree argument is elementary. To keep the code's *maximum* list size small, we choose its evaluation domain using a generic-rank theorem for Reed–Solomon agreement matrices and a finite-field union bound. The underlying line is the consecutive-power construction of Ben-Sasson, Carmon, Haböck, Kopparty, and Saraf [25, Theorem 7.1], also used in [21]. Here we choose a domain on which that line coexists with a small global maximum list size at the target radius.

Recall that $\ell_C(\theta)$ is the maximum number of codewords within relative Hamming distance θ of any received word. A selected pair (z, p_z) lies on a parametrized codeword line if $p_z = P_0 + zP_1$ for two fixed codewords $P_0, P_1 \in C$; the scalar z must be the original parameter. Correlated agreement on t coordinates means that $P_0 = f_0$ and $P_1 = f_1$ simultaneously on those coordinates.

Theorem 5.1 (An actual-list-size separation). *Fix any rational rate $0 < \rho < 1$. For every sufficiently large n with $k = \rho n$ integral, there are a prime p , an n -point domain $D \subset \mathbb{F}_p$, and the rate- ρ code $C = \text{RS}(\mathbb{F}_p, D, k)$ with the following properties at*

$$\theta = 1 - \rho - \frac{1}{n}.$$

- (i) *The actual maximum list size satisfies $\ell_C(\theta) = n - k - 1$.*
- (ii) *An affine line $f_0 + zf_1$ has exactly $M = \binom{n}{k+1}$ parameters within distance θ of C . Each has a unique nearby codeword p_z . All remaining $p - M$ points have distance $1 - \rho = \theta + 1/n$, the maximum possible distance from this code.*
- (iii) *Every parametrized codeword line contains at most $n - k$ of these selected pairs (z, p_z) , and this bound is attained. The pair (f_0, f_1) has no correlated agreement on $k + 1$ coordinates.*
- (iv) *The radius θ lies strictly below the characteristic- p Elias radius.*

Consequently, along this family, line-decodability with concurrency n fails at threshold $\ell_C(\theta)n + R_n$ for every $R_n = o(n)$.

Proof. Choose the domain. For sufficiently large n , we have $k \geq 2$ and $a = n - k \geq 2$; take a prime $p \geq n + k2^{10an}$. Sample n distinct domain points uniformly from $\mathbb{F}_p$. We impose three conditions on this domain.

First, put $L = a - 1 \geq 1$. We use the generic-matrix argument underlying [10, Sections 2.2–2.4 and 3.1], with an explicit finite-field union bound. If $L + 1$ distinct codewords lie in a radius- θ ball, their total (unnormalized) distance from its center is at most $(L + 1)(a - 1) = La = L(n - k)$. The non-strict inequality in [10, Lemma 2.3] permits equality and implies that some subcollection of t codewords, $2 \leq t \leq L + 1$, has a k -weakly-partition-connected agreement hypergraph. Its vertices are these codewords, and its i th edge records which agree with the center at coordinate i . Weak partition connectivity means that every partition $\mathcal{P}$ of the vertices satisfies $\sum_i \max\{|\mathcal{P}(e_i)| - 1, 0\} \geq k(|\mathcal{P}| - 1)$, where $\mathcal{P}(e_i)$ denotes the parts meeting e_i .

The associated reduced intersection matrix encodes equalities of polynomial values within each edge, using a spanning star, with the coefficient block for the last polynomial removed. It has $(t - 1)k$ columns and entries of degree at most $k - 1$ in the formal domain points. It has full symbolic column rank by [10, Theorem 2.11]. However, at a domain realizing the proposed subcollection, the

coefficient differences from its last codeword give a nonzero kernel vector. Thus a fixed nonzero full-column minor must vanish there. This minor has total degree at most $(t-1)k(k-1) \leq Lk^2$.

There are at most $\sum_{t=2}^{L+1} 2^{tn} \leq 2^{(L+2)n}$ such labeled hypergraphs. For independent uniform domain points, Schwartz–Zippel and a union bound give failure probability at most $2^{(L+2)n} Lk^2/p$. Since $p > n^2$, the probability that those points are distinct is greater than $1/2$. Conditioning on distinctness therefore gives failure probability at most

$$\frac{2^{(L+2)n+1} Lk^2}{p} < 2^{-2an}.$$

For the last inequality, use $L = a-1$, $p > k2^{10an}$, $2 \leq a < n$, and $1 + \log_2((a-1)k) < (7a-1)n$ for $n \geq 3$. Consequently, except with probability 2^{-2an} , every radius- θ list has size at most $a-1$.

Second, require all $(k+1)$ -subsets of the domain to have distinct sums. For two different index sets T, U , condition on every domain point except one indexed by $T \setminus U$. Equality of the two sums fixes at most one value of this remaining point, among $p-n+1$ possibilities. The probability of any collision is therefore at most $\binom{M}{2}/(p-n+1)$. Our choice of p gives

$$2^{-2an} + \frac{\binom{M}{2}}{p-n+1} < 1.$$

Indeed, $M \leq 2^n$ and $p-n+1 > k2^{10an}$ make the second failure probability less than $2^{2n-1-10an} \leq 2^{-2an}$. Finally, impose the matching-list construction from Proposition 5.2 below, simultaneously at all integer radii strictly below $n-k$. Its proof bounds failure by $(a-1)^2/[2(p-n+1)] < 2^{-2an}$. Together with the preceding two bounds, total failure is less than $3 \cdot 2^{-2an} < 1$. Fix a domain satisfying all three conditions. The matching construction in particular gives $\ell_C(\theta) \geq a-1$, so the bound at our target radius is exact.

Construct the nearby points. For each $(k+1)$ -subset $T \subset D$, put

$$s_T = \sum_{a \in T} a, \quad V_T(X) = \prod_{a \in T} (X - a).$$

Evaluate the affine family

$$f_z(X) = X^{k+1} - zX^k$$

on D . At parameter $z = s_T$, choose

$$p_{s_T}(X) = X^{k+1} - s_T X^k - V_T(X).$$

Its two leading coefficients cancel, so $\deg p_{s_T} < k$. Its difference from f_{s_T} vanishes exactly on T , giving distance $1 - (k+1)/n = \theta$.

Conversely, if a polynomial P of degree less than k agrees with f_z on at least $k+1$ coordinates, then $f_z - P$ is monic of degree $k+1$ and must equal V_T for its root set $T \subset D$. Comparing the coefficient of X^k gives $z = s_T$. Distinct subset sums therefore identify both T and P uniquely. There are exactly M nearby parameters, each with one nearby codeword. Any received word agrees with some degree-below- k polynomial on at least k coordinates, by interpolation. Every remaining parameter therefore has maximum agreement exactly k , hence distance $1 - \rho$. Since $p > M$, such maximally distant points exist on the same line.

Bound concurrency. Suppose $P_0 + zP_1$ contains m selected pairs, and let b count the coordinates where both $X^{k+1} = P_0(X)$ and $-X^k = P_1(X)$. Since $\deg P_1 < k$, the second equality holds on at most k coordinates; hence $b \leq k$, which already rules out the claimed common agreement. At those

b coordinates all m parameters may agree. At every other coordinate two distinct affine functions of z agree at most once. Counting the $k + 1$ agreements of each selected witness yields

$$m(k + 1) \leq bm + n - b, \quad m \leq \frac{n - b}{k + 1 - b} \leq n - k.$$

The bound is attained by fixing a k -subset $A \subset D$ and writing $V_A(X) = \prod_{a \in A} (X - a)$ and $s_A = \sum_{a \in A} a$. The codeword line

$$[X^{k+1} - (X + s_A)V_A(X)] + z[V_A(X) - X^k]$$

has both coefficient polynomials of degree less than k and contains the selected witness for each $T = A \cup \{a\}$ with $a \in D \setminus A$, at its parameter $z = s_A + a$. There are $n - k$ such parameters.

Check the radius and threshold. The entropy estimate $H_p(\theta) \leq \theta + 1/\log_2 p < 1 - \rho$ places θ strictly below the Elias radius. Finally, $M = 2^{H_2(\rho)n - O_\rho(\log n)}$, whereas $\ell_C(\theta)n \leq (n - k - 1)n = O_\rho(n^2)$. Thus $M > \ell_C(\theta)n + R_n$ eventually for every $R_n = o(n)$, without the required concurrency. $\square$

Attribution of the exact profile. The upper bound below follows from [13, Theorem 1.5]. For the lower bound, Roth's Theorem 4 [14] suffices whenever $a - 2 \leq \binom{k+1}{2}$. Indeed, set $L_s = \lfloor (a - 1)/s \rfloor$. If $L_s \geq 2$, put $L = L_s - 1$ and write $a = (L + 1)u + v$ with $1 \leq v \leq L + 1$. Then $u \geq 1$ and $L \leq a - 2 \leq \binom{k+1}{2} \leq \binom{k-1+u+v}{k-1}$. If the maximum list size at radius $(a - s)/n$ were at most L , Roth's bound would imply $a - s \leq La/(L + 1)$, contrary to $L_s s \leq a - 1$. The case $L_s = 1$ is immediate. The sufficient condition holds eventually at every fixed rate in $(0, 1)$. Thus the fixed-rate exact profile is a consequence of published results. The proof below supplies a direct lower construction for every finite $k \geq 2$ and uses [10] for the stated finite-field probability.

Proposition 5.2 (Exact generic list sizes at every integer radius below $n - k$). *Let $k \geq 2$, $a = n - k \geq 2$, and let $p \geq n + k2^{10an}$ be prime. For a uniformly random ordered set of n distinct evaluation points, with probability greater than $1 - 2^{1-2an}$, the Reed–Solomon code C simultaneously satisfies*

$$\ell_C\left(\frac{a - s}{n}\right) = \left\lfloor \frac{a - 1}{s} \right\rfloor \quad (1 \leq s < a).$$

At radius zero the maximum list size is one. The domain in Theorem 5.1 can be chosen to have this entire profile, as well as distinct sums of its $(k + 1)$ -subsets.

Proof. Upper bounds from one rank condition. Use the good-domain event in the proof of Theorem 5.1: every k -weakly-partition-connected agreement hypergraph on at most a vertices has a reduced intersection matrix of full column rank. Its failure probability is less than 2^{-2an} . For each s , put $L_s = \lfloor (a - 1)/s \rfloor$. If $L_s + 1$ distinct codewords are each at distance at most $a - s$ from one word, their total distance is at most

$$(L_s + 1)(a - s) \leq L_s a,$$

since $(L_s + 1)s \geq a$. Lemma 2.3 of [10] then produces one of the excluded hypergraphs on at most $L_s + 1 \leq a$ vertices. Thus the same rank condition supplies all upper bounds simultaneously.

Matching lists from a chain of affine polynomials. Set aside $k - 2$ coordinates H and put $F(X) = \prod_{h \in H} (X - h)$. There remain $a + 2$ coordinates. Fix s and write $L = L_s$. Choose two anchors A, B , another s coordinates, and a seed z_0 . Give the anchors and the other s coordinates received value zero, and give z_0 value one. The zero polynomial already has $s + 2$ agreements. For

each $1 \leq j \leq L-1$, choose a fresh set Z_j of s coordinates and designate one of them z_j . Let $A_j = A$ for odd j and $A_j = B$ for even j , and define

$$c_j = \frac{w(z_{j-1})}{z_{j-1} - A_j}, \quad q_j(X) = c_j(X - A_j), \quad w(x) = q_j(x) \quad (x \in Z_j).$$

This uses $s + 3 + (L-1)s = Ls + 3 \leq a + 2$ coordinates. Give any remaining coordinates value zero. All denominators and slopes are nonzero, and q_j agrees with w at the anchor A_j , the preceding point z_{j-1} , and all s points of Z_j . Each candidate therefore has at least $s + 2$ agreements.

It remains to ensure that the candidates are distinct. Candidates with different anchors have different roots. For $i < j$ with the same anchor, the equality $q_i = q_j$ is equivalent to

$$\prod_{r=i}^{j-1} \frac{z_r - A_r}{z_r - A_{r+1}} = 1.$$

Condition on every domain coordinate except z_i . All factors but one are fixed and nonzero, so the equality takes the form $C(z_i - A_i)/(z_i - A_{i+1}) = 1$ with $C \neq 0$. Because the anchors differ, this equation has at most one solution. Under distinct sampling its probability is therefore at most $1/(p - n + 1)$. There are at most $L^2/4$ pairs with the same anchor, so candidate repetition has probability at most $L^2/[4(p - n + 1)]$. On its complement, multiplication by F gives L distinct polynomials of degree less than k . Set the full received word to zero on H and to $F(x)w(x)$ elsewhere. Each candidate has at least $(k - 2) + (s + 2) = k + s$ agreements, proving the matching lower bound.

Choose one domain for all radii. The sum of the lower-construction failure probabilities is at most

$$\frac{1}{4(p - n + 1)} \sum_{s=1}^{a-1} L_s^2 < \frac{(a-1)^2}{2(p - n + 1)} < 2^{-2an},$$

using $L_s \leq (a-1)/s$ and $\sum_{s \geq 1} s^{-2} < 2$. Together with the single rank-condition failure bound, this proves the probability assertion. Requiring distinct $(k+1)$ -subset sums adds less than 2^{-2an} to the failure probability, as in the preceding proof; the total remains less than one. $\square$

The upper bounds here use the known generic-rank theorem; the lower bounds are explicit interpolation constructions. The domain is generic and the field bound is large. The assertion does not give this list profile on arbitrary prescribed evaluation domains. Dimension $k = 1$ is a real exception: at agreement $1 + s$, the constant code has maximum list size $\lfloor n/(s+1) \rfloor$.

Corollary 5.3 (Extension fields and interleaving). *The domain in Proposition 5.2 has the same exact list profile over every extension of its base field and for every finite interleaving order. Interleaved distance is measured by columns: a coordinate agrees only when all components agree. No further exceptional-domain event is needed.*

Proof. A nonzero full-column minor over the base field remains nonzero over an extension. For interleaved words, agreement of entire columns gives the same agreement hypergraph as in the scalar upper-bound proof. Apply its full-rank matrix to each component of the polynomial coefficient differences. Each component difference vanishes, so two allegedly distinct tuples would coincide. The scalar lower list embeds in one component, with zero in the others, and remains a list over every extension field. This proves equality. The argument does not apply to entrywise distance averaged over all components. $\square$

The companion `checks/astra_rs_all_surplus_check.py` verifies 3,654 large-field parameter fixtures on 378 domains, each reused for all surpluses, including nondivisible surpluses and cases with $k < 2s$. It also checks 450 small-field fixtures, retaining 55 with repeated candidate lines, and enumerates the full lists at four constructed centers. The 79,079 allocation and probability checks, and 3,003 combined lower-failure checks, pass. Separately, `checks/actual_list_global_small_check.py` exhausts all center classes modulo the code in four small dimension-two examples. Three structured domains have a larger list than the generic profile, confirming the need for the domain qualification.

Sharpness of the global list bound at rates at least one-half. Write $a = n - k \geq 2$. If $k \geq a - 2$, every n -point Reed–Solomon domain has a list of size at least $a - 1$ at radius $(a - 1)/n$. Indeed, set $L = a - 1$ and $h = k - a + 2 \geq 0$, and partition the domain into H , $\{\alpha_1, \dots, \alpha_L\}$, and $\{\beta_1, \dots, \beta_L\}$, where $|H| = h$. Let $F(X) = \prod_{x \in H} (X - x)$ and

$$Q_i(X) = F(X) \left[X^{L-1} - \prod_{j \neq i} (X - \alpha_j) \right] \quad (1 \leq i \leq L).$$

The leading terms cancel, so each has degree at most $h + L - 2 = k - 1$ (for $L = 1$, the sole polynomial is zero). For $i \neq j$, their values at α_i differ by the nonzero quantity $F(\alpha_i) \prod_{r \neq i} (\alpha_i - \alpha_r)$, so they are distinct. Choose a received word equal to zero on H , to $F(\alpha_j) \alpha_j^{L-1}$ at α_j , and to $Q_i(\beta_i)$ at β_i . Each Q_i agrees on H , on every α_j with $j \neq i$, and on β_i , giving $h + L = k + 1$ agreements. This supplies the matching lower bound on every domain in this range. The chain construction in Proposition 5.2 supplies a matching lower bound on a generic domain for every $k \geq 2$.

A sharper collision estimate. For the $s = 1$ chain, the probability of a repeated line is at most

$$\frac{\lfloor \max(a - 5, 0)^2 / 4 \rfloor}{p - n + 1}.$$

For a same-anchor pair $i < j$, condition on all domain points except z_i . Equality of the two slopes has the form $C(z_i - A_i)/(z_i - A_{i+1}) = 1$ with fixed $C \neq 0$. If $C = 1$, equality is impossible; otherwise it fixes at most one of the $p - n + 1$ remaining values of z_i . Pairs with $j - i = 2$ cannot collide: their cleared equation is, up to sign, $(A - B)(z_{i+1} - z_i) = 0$. Among $a - 2$ alternating lines, the number of same-anchor pairs separated by at least four steps is $\lfloor \max(a - 5, 0)^2 / 4 \rfloor$, proving the estimate. In particular, the chain is distinct on every domain for $a \leq 6$. The independent checker `checks/actual_list_chain_conditional_audit.py` exhausts 36,200 allowed values across 1,400 conditional pair cases over $\mathbb{F}_{37}$, verifying the one-solution bound and all 720 gap-two exclusions in those fixtures.

The gap in Theorem 5.1 is $\eta = 1/n$. It therefore rules out a *uniform*, familywise threshold based on the actual list size; it does not rule out an asymptotic statement made separately at each fixed positive gap, with arbitrary dependence of its remainder on that gap. This differs from the separate fixed-gap construction, which defeats a prescribed numerical list coefficient. The present construction instead has provably small actual lists. Its ingredients are the vanishing-polynomial construction and a domain chosen to satisfy the generic-rank guarantee simultaneously with distinct subset sums.

5.1 A general upper bound and a sharp endpoint

The affine-line count in Theorem 5.1 attains a general upper bound at the same radius. The upper bound follows by puncturing to the unique-decoding regime and applying the correlated-agreement

theorem of Ben-Sasson et al. [8, Theorems 1.2 and 1.4]. We include the argument to explain the extremality of the example; no novelty is claimed for this corollary of the known theorem. Stronger bounds are known in parts of the unique-decoding and Johnson regimes [25, Corollary 1.4 and Theorem 1.5]; the elementary bound below is useful for its general range and its sharp endpoint.

Proposition 5.4 (An unconditional puncturing bound). *Let $C = \text{RS}(F, D, k)$, where F is finite, $|D| = n$, and $a = n - k$. Let r be an integer with $a/2 \leq r < a$, and put $s = a - r$ and $d = 2r - a$. Suppose that $f_0, f_1 \in F^D$ have no correlated agreement with degree-below- k polynomials on $n - r$ coordinates. Then the total number of nearby witness pairs satisfies*

$$\sum_{z \in F} |\{P : \deg P < k, \text{dist}(f_0 + zf_1, P|_D) \leq r/n\}| \leq \left\lfloor (n - d) \frac{\binom{n}{d}}{\binom{r}{d}} \right\rfloor = \left\lfloor s \frac{\binom{n}{d+1}}{\binom{r}{d+1}} \right\rfloor.$$

Proof. Delete any d coordinates. The resulting code has length $n - d$, dimension k , and redundancy $a - d = 2s$. Correlated agreement on $(n - d) - s = n - r$ remaining coordinates would also give correlated agreement in the original domain, contrary to the hypothesis. The unique-radius theorem therefore permits at most $n - d$ parameters whose punctured words have at most s errors. At each parameter the nearby polynomial is unique, because the punctured code has minimum distance $2s + 1$. Its relative radius is $s/(n - d) = (1 - k/(n - d))/2$, the included endpoint of the cited unique-radius bound.

For each nearby witness pair (z, P) in the original code, let its error support be E , of size at most r , and extend E to an r -element coordinate set R . Deleting any d -subset of R leaves at most $r - d = s$ errors. Thus this pair is counted by at least $\binom{r}{d}$ punctures. Restriction preserves distinct polynomials because $n - d \geq k$. There are $\binom{n}{d}$ punctures, each counting at most $n - d$ witness pairs. Double counting proves the first bound. The second follows from $r - d = s$ and the adjacent binomial identities. The case $d = 0$ uses the empty puncture. $\square$

At $r = a - 1$, the bound is exactly $\binom{n}{k+1}$. Theorem 5.1 attains it, while its global maximum list size is only $n - k - 1$. Consequently, among affine lines without correlated agreement at this radius, the number of nearby parameters in that theorem is as large as possible. This extremality concerns an algebraic count, not an algorithm for producing its witnesses.

Proposition 5.5 (Pairs with unexplained agreement sets). *Retain the parameters of Proposition 5.4, but allow f_0, f_1 to have correlated agreement. For a nearby pair (z, P) , let $S = \{x \in D : P(x) = f_0(x) + zf_1(x)\}$ be its entire agreement set. The number of such pairs for which no two degree-below- k polynomials jointly explain f_0, f_1 on S is at most*

$$\left\lfloor (n - d) \frac{\binom{n}{d}}{\binom{r}{d}} \right\rfloor.$$

Proof. First consider a length- N puncture of redundancy $2s$. If its coefficient words have no correlated agreement on $N - s$ coordinates, the unique-radius theorem and uniqueness of the nearby polynomial bound all nearby pairs by N . Otherwise, choose a jointly agreeing codeword path $Q_0 + zQ_1$, and let J be its full coefficient agreement set. This path is nearby at every parameter and hence is the unique nearby polynomial. Outside J , an accidental agreement can occur at at most one parameter per coordinate. There are at most $N - |J| \leq s$ such parameters. At every other parameter the entire agreement set is J and is jointly explained. Thus at most N pairs in any puncture have unexplained entire agreement sets.

For an original pair with an unexplained set S , interpolate the two coefficient words on any k positions of S . At least one further position fails one of these interpolants. The resulting $(k + 1)$ -set

$K \subset S$ is itself not jointly explainable. Its complement has size $a - 1 \geq r$, so the error support $E = D \setminus S$ can be extended to an r -set $R \subseteq D \setminus K$. Every d -subset of R leaves at most s errors after puncturing and preserves K . The punctured agreement set therefore remains unexplained. Each original pair is counted in at least $\binom{r}{d}$ punctures, each of which counts at most $n - d$ such pairs. Restriction is injective on the polynomial witnesses, so double counting gives the bound. $\square$

Corollary 5.6 (Polynomial curves). *Retain a, r, s, d from Proposition 5.4, and let $f(z) = \sum_{j=0}^e z^j f_j$ for $f_j \in F^D$ and $e \geq 0$. The number of nearby pairs (z, P) whose entire agreement set cannot be jointly explained by degree-below- k polynomials for all $e + 1$ coefficient words is at most*

$$\left\lceil e(n - d) \frac{\binom{n}{d}}{\binom{r}{d}} \right\rceil.$$

If no such joint agreement exists on $n - r$ coordinates, the same bound holds for all nearby pairs, summed over parameters.

Proof. This is an elementary consequence of the polynomial-curve unique-radius theorem [8, Theorem 6.1]. On a length- N puncture of redundancy $2s$, that theorem bounds nearby parameters by eN unless the coefficient words jointly agree with codewords on $N - s$ positions. Uniqueness makes this a witness bound. In the joint-agreement case, choose the corresponding codeword curve and its full coefficient agreement set J . It is nearby and unique at every parameter. Each coordinate outside J contributes at most e roots of its nonzero discrepancy polynomial. Thus at most $e(N - |J|) \leq es$ parameters have an agreement set larger than J ; all remaining sets are jointly explained.

For each original unexplained support, interpolate every coefficient word on the same k positions and retain one additional position where an interpolant fails. This gives an unexplained $(k + 1)$ -set, independent of e . The protected-support counting argument in Proposition 5.5 now applies with eN in place of N , yielding the displayed bound. For $e = 0$, every nearby support is explained by its nearby polynomial, so the count is zero. The root bound uses formal polynomials and requires no assumption that $e < |F|$. $\square$

The affine bound is also n at $2r = a$. Between the two endpoints it can be much larger than the actual maximum list size times n ; at a fixed gap beyond unique decoding it is generally exponential in n . It therefore leaves the fixed-gap question open. The exact checker `checks/verify_puncturing_bound.py` verifies the parameter and support-count identities and enumerates all 257 parameters on each of five small lines attaining the one-coordinate-gap bound. The independent `checks/astra_syndrome_determinants_check.py` checks the witness and entire-agreement-set bounds on 378 affine lines, including lines with correlated agreement, by exact decoding and 572,226 codeword-distance evaluations. For the curve corollary, `checks/verify_curve_puncturing.py` enumerates all field parameters and codewords on 300 deterministic small curves, including 60 with formal degree at least the field size. All 466,140 distance checks pass; the fixtures include 421 unexplained witness pairs on curves with some joint agreement.

5.2 An incidence bound for all nearby witnesses

Proposition 5.7 (A witness form of the circuit-incidence bound). *Let F be finite, let $1 \leq k < n$, and let $C \subseteq F^n$ be a linear $[n, k]$ MDS code. Let t, e be integers with $k + 1 \leq t \leq n$ and $e \geq 0$, and put $f(z) = \sum_{j=0}^e z^j f_j$ for $f_j \in F^n$. For each pair $(z, c) \in F \times C$, write $S(z, c) = \{i : c_i = f(z)_i\}$. Let $\mathcal{W}$ consist of the pairs with $|S(z, c)| \geq t$ whose entire agreement set does not jointly explain all*

the coefficient words $f_0, \dots, f_e$. Then, for every $k+1 \leq b \leq t$,

$$\sum_{(z,c) \in \mathcal{W}} \binom{|S(z,c)|-1}{b-1} \leq \min\{e, |F|\} \binom{n}{b}.$$

In particular,

$$|\mathcal{W}| \leq \left\lfloor \min\{e, |F|\} \min_{k+1 \leq b \leq t} \frac{\binom{n}{b}}{\binom{t-1}{b-1}} \right\rfloor.$$

If the coefficient words have no joint explanation on any t coordinates, the bound counts every nearby witness pair.

Proof. This extends the incidence proof of Jo [27, Lemma 4.1 and Theorem 4.2]. His local rejection lemma states that a noncodeword on m coordinates of an MDS code rejects at least $\binom{m-1}{b-1}$ tests of size $b \geq k+1$. For each pair in $\mathcal{W}$, apply that lemma to a coefficient word that is outside the code on $S(z, c)$.

Fix a b -set A rejecting at least one coefficient word. The polynomial $f(Z)|_A$ in the quotient space $F^A/(C|_A)$ is formally nonzero and has degree at most e . A linear functional nonzero on one of its coefficient vectors gives a nonzero scalar polynomial of degree at most e . Hence at most $\min\{e, |F|\}$ parameters make $f(z)|_A$ a code restriction. Each such restriction determines at most one global codeword, because $b \geq k$ and the code is MDS. Counting all incidences proves the weighted bound, and then $|S(z, c)| \geq t$ gives the stated count. The argument needs no restriction $e < |F|$. $\square$

For affine lines in the range of Proposition 5.4, put $t = k + s = n - r$. The minimum in Proposition 5.7 is attained at $b = k + 1$, yielding the additional bound

$$|\mathcal{W}| \leq \left\lfloor \frac{\binom{n}{k+1}}{\binom{k+s-1}{k}} \right\rfloor.$$

Indeed, the ratio of the expressions for successive test sizes is $b(n-b)/((b+1)(t-b))$, which exceeds one here because $(k+1)(r+1) > t$. At $(n, k, r) = (10, 2, 6)$ this improves the puncturing bound from 84 to 40. At $s = 1$ both bounds equal $\binom{n}{k+1}$, so the sharp endpoint is unchanged. The puncturing bound can be stronger near the unique radius; we may take the minimum of the two bounds. At fixed positive rate $k/n \rightarrow \rho$ and agreement fraction $t/n \rightarrow \tau$ with $0 < \rho < \tau < 1$, the circuit bound still has exponential rate $H(\rho) - \tau H(\rho/\tau) > 0$, with natural-log entropy H . It therefore does not resolve the fixed-gap question.

Corollary 5.8 (Sharp endpoint for polynomial curves). *Let $1 \leq k \leq n - 2$ and $e \geq 1$. Over a suitable finite field, there are an n -point Reed–Solomon domain and a degree- e curve with exactly $e \binom{n}{k+1}$ nearby parameter–witness pairs at radius $(n - k - 1)/n$. Each nearby word has a unique witness, and the coefficient words have no joint agreement on $k + 1$ coordinates. Every parametrized codeword curve of degree at most e contains at most $e(n - k)$ of these pairs, and this bound is attained.*

Proof. Choose a field of characteristic not dividing e and a domain whose $(k+1)$ -subset sums s_T are distinct. Choose b outside those sums, and pass to a finite extension splitting all polynomials $Z^e + b - s_T$. Each has e distinct roots, and different polynomials have disjoint root sets. Set $h(z) = z^e + b$ and $f(z)(x) = x^{k+1} - h(z)x^k$. As in the affine construction, agreement with a degree-below- k polynomial on $k+1$ coordinates forces $h(z) = s_T$ and uniquely determines the witness. This gives exactly $e \binom{n}{k+1}$ pairs. The coefficient of z^e is $-X^k$, which precludes joint agreement on $k+1$ coordinates.

For a parametrized codeword curve $Q(z)$ of degree at most e , let b_0 count coordinates on which $f(z) = Q(z)$ identically in z . Their leading coefficients show $b_0 \leq k$. At each other coordinate the discrepancy has at most e roots. If Q contains m selected pairs, then $m(k+1) \leq b_0 m + e(n-b_0)$, so $m \leq e(n-k)$. The affine codeword line attaining concurrency $n-k$, composed with h , attains equality. If the domain also satisfies the exact generic list profile, that profile persists over the splitting field by the field-extension argument above. $\square$

6 Finite examples on the norm-one circle

The interval construction shows that a uniform list bound fails when the field and block length vary. We now give finite examples over one fixed prime, $p = 2^{31} - 1$. The domains have multiplicative structure, and the construction uses that structure instead of small integer moments. The two methods share a simpler underlying idea: many root polynomials can have the same outer coefficients while differing in their middle coefficients.

6.1 Why the circle helps

Let $i^2 = -1$ in $F_2 = \mathbb{F}_{p^2}$. The norm-one circle is the cyclic group

$$U = \{u \in F_2^\times : u^{p+1} = 1\}.$$

Every element can be written $u = x + iy$ with $x, y \in \mathbb{F}_p$ and $x^2 + y^2 = 1$. On this group $u^p = u^{-1}$. Thus conjugating a polynomial's coefficients and reversing its exponents are closely related operations. For $n \mid p+1$, write μ_n for the subgroup of size n . Our evaluation domains are μ_n and its rotations $g\mu_n$.

For even $K = 2m$, consider the dimension- K circle space

$$C_K = \{A(x) + yB(x) : \deg A < m, \deg B < m\}.$$

Its coefficients may lie in $\mathbb{F}_p$ or in an extension, as specified. In the coordinate u , the space contains all Laurent polynomials in the interior band $-(m-1), \dots, m-1$, with the appropriate coefficient field condition. A Laurent polynomial simply allows negative powers of u ; these are well defined because every evaluation point is nonzero. The interior band leaves room to divide out two common zeros without leaving C_K .

Write $n = MB$. The power map $u \mapsto u^B$ sends μ_n onto μ_M , with B points in each fiber. To choose agreement positions, choose an h -element subset $A \subseteq \mu_M$ containing 1, and form

$$F_A(Y) = \prod_{a \in A} (Y - a) = \sum_{j=0}^h f_{A,j} Y^j.$$

The roots have norm one, so, with $c = f_{A,0}$,

$$f_{A,j} = c f_{A,h-j}^p.$$

Fixing c and r leading coefficients therefore also fixes r trailing coefficients. There are $\binom{M-1}{h-1}$ possible subsets, but at most Mp^{2r} choices of this coefficient data. Some class has at least

$$\left\lceil \frac{\binom{M-1}{h-1}}{Mp^{2r}} \right\rceil$$

members. Subtracting their common outer terms leaves distinct low-degree polynomial parts. Agreement occurs on the selected fibers, so the count of roots in μ_M becomes a count of coordinates in μ_n .

Two additional steps yield the two results below. For a base-field list, one rescales the polynomial parts to make their values lie in $\mathbb{F}_p$. The common anchor $1 \in A$ then permits division at the two circle points $u = 1, -1$. For an affine family, one instead divides by $Y - \tau$, where $\tau \in \mathbb{F}_{p^4} \setminus \mathbb{F}_{p^2}$. As in Lemma 3.2, this produces a scalar label for each polynomial part and a direction word with too little agreement to permit correlated agreement. A collision count selects a pole with many distinct labels. Appendix A gives both arguments in full.

6.2 Two constructions at the same parameters

Theorem 6.1 (Finite circle list and affine-family bounds). *Let $p = 2^{31} - 1$, $E = \mathbb{F}_{p^4}$, and*

$$n = 2^{26}, \quad K = 2^{25}, \quad \rho = \frac{1}{2}, \quad \eta = \frac{33}{1024} - \frac{2}{n}, \quad T = (\rho + \eta)n, \quad \theta = 1 - \rho - \eta.$$

On $D = \mu_n$, and on any rotated domain $g\mu_n \subseteq U$, the following statements hold.

(i) *There is an $\mathbb{F}_p$ -valued received word with more than*

$$2^{40}$$

distinct codewords from the base-field circle space C_K agreeing with it in at least T positions.

(ii) *There are words $f_0, f_1 \in E^D$ with no correlated agreement at threshold T with the extension-field circle code C_K , but more than*

$$2^{74}$$

distinct scalars $z \in E$ for which $f_0 + zf_1$ has a codeword agreeing in at least T positions. Witnesses can be selected so that every parametrized affine codeword graph contains at most fifteen selected pairs (z, P_z) .

The radius θ is strictly below the characteristic- p Elias radius at rate $1/2$.

Proof. For (i), apply Theorem A.2 with $M = 2048$, $d = 1024$, $r = 32$, $h = 1090$, and $B = 2^{15}$. For (ii), apply Theorem A.4 with $M = 1024$, $r = 15$, $h = 545$, and $B = 2^{16}$. Both choices give $K = n/2$ and the same certified agreement threshold $T = hB - 2$. The exact coefficient counts and collision bound are evaluated in (44) and Section A.1.2. Corollary A.9 gives the selected-graph bound. The strict Elias comparison has the rational certificate in Section A.1.3. $\square$

These are two different received-word constructions with the same length, rate, and radius. Neither the theorem nor its proof requires enumerating the very large coefficient classes. The formulas and integer lower bounds are explicit; the choice of a sufficiently large class, and of a suitable pole in (ii), is guaranteed by counting.

Transfer to ordinary Reed–Solomon codes. The second construction also gives an ordinary Reed–Solomon example. Remove the two anchor coordinates and use the stereographic coordinate $t = y/(1 + x) \in \mathbb{F}_p$. Multiplication by the nonzero weight $(1 + t^2)^{m-1}$ identifies the interior Laurent band with ordinary polynomials of degree at most $K - 2$. The resulting code has

$$n' = n - 2, \quad k' = K - 1, \quad \rho' = \frac{1}{2}, \quad \eta' = \frac{T}{n - 2} - \frac{1}{2} = \frac{33 \cdot 2^{16} - 1}{n - 2}.$$

It has the same scalar labels and at least T agreements per selected witness, with no correlated agreement at threshold T . Theorem A.10 proves that the coordinate change and scaling preserve these equalities. The new gap is η' , not η : puncturing preserves the number of certified agreements but changes the block length. Comparisons for this Reed–Solomon code must use n' and η' .

6.3 Stronger finite examples at other rates

The preceding examples use rate $1/2$ to make the two constructions directly comparable. Neither construction is restricted to that rate. In the list construction, choose any even packet dimension d and put $h = d + 2r + 2$. The resulting rate is d/M ; the gap remains $(2r + 2)/M - 2/n$. In particular, no coefficient reserve is needed for the following larger list.

Corollary 6.2 (A quarter-rate base-field list). *Let $p = 2^{31} - 1$, $n = 2^{26}$, $K = 2^{24}$, and*

$$\rho = \frac{1}{4}, \quad \eta = \frac{1}{32} - \frac{2}{n}.$$

On μ_n and its rotations, some $\mathbb{F}_p$ -valued word has more than 2^{43} codewords in C_K agreeing on at least $(\rho + \eta)n$ coordinates. The radius $1 - \rho - \eta$ is strictly below the characteristic- p Elias radius. If L denotes the certified list size, then

$$\frac{L}{2^{H_2(\rho)/\eta}} > 200,000.$$

Proof. Use Theorem A.2 with $M = 64$, $B = 2^{20}$, $d = 16$, $r = 0$, and $h = 18$. The count is $\lceil \binom{63}{17}/64 \rceil$. Appendix D.2 gives exact certificates for the count, radius, and displayed comparison. $\square$

Corollary 6.3 (An eighth-rate line without correlated agreement). *Let $p = 2^{31} - 1$, $E = \mathbb{F}_{p^4}$, $n = 2^{26}$, $K = 2^{23}$, and*

$$\rho = \frac{1}{8}, \quad \eta = \frac{5}{256} - \frac{2}{n}, \quad T = (\rho + \eta)n.$$

On μ_n and its rotations, there are words $f_0, f_1 \in E^D$ with no correlated agreement at threshold T with C_K , but more than 2^{75} distinct scalars z for which $f_0 + zf_1$ has a codeword agreeing on at least T coordinates. If J is the certified scalar count, then

$$\frac{J}{n2^{H_2(\rho)/\eta}} > 2^{22}.$$

The radius is strictly below the characteristic- p Elias radius. Witnesses can be selected so that every parametrized affine codeword graph contains at most 55 selected pairs.

Proof. Apply Theorem A.4 with $M = 256$, $B = 2^{18}$, $r = 1$, and $h = 37$. Its polynomial degree is $d = 33$, so its collision factor is $d - 1 = 32$. The second word agrees with every outer-band codeword at at most $a_0 = K + B = 33B < T = 37B - 2$ coordinates. Lemma A.8 bounds the number of selected pairs on one graph by

$$\left\lfloor \frac{223B}{4B - 2} \right\rfloor = 55.$$

The exact scalar count, Elias comparison, and ratio certificate are in Appendix D.2. $\square$

The small packet group explains the improvement in the list construction: only 64 possible constant coefficients divide the $\binom{63}{17}$ anchored root sets. Reciprocity fixes the opposite end of each root polynomial for free. The degree allowance at rate $1/4$ is already enough to contain the remaining

middle terms after removing the anchors. The line construction uses a different rate and one pair of reserved coefficients. Its scalar count exceeds the numerical line prescription by more than four million. Both examples transfer to ordinary Reed–Solomon codes by Theorem A.10, with length $n - 2$ and dimension $K - 1$. At these rates the exact transferred rate is $(K - 1)/(n - 2)$, rather than precisely $1/4$ or $1/8$.

6.4 The quadratic field already suffices

The larger scalar count above uses a pole in $\mathbb{F}_{p^4} \setminus \mathbb{F}_{p^2}$. An outside pole is not essential to the counterexample. It is enough to avoid the packet labels μ_M . Choosing the pole inside $\mathbb{F}_{p^2}$ gives fewer available poles, but also keeps the entire construction over the quadratic field.

For example, take $p = 2^{31} - 1$, $n = 2^{26}$, $K = 2^{23}$, and

$$\rho = \frac{1}{8}, \quad \eta = \frac{3}{128} - \frac{2}{n}.$$

Corollary A.5, with $M = 128$, $B = 2^{19}$, $r = 0$, and $h = 19$, gives an $\mathbb{F}_{p^2}$ -valued pair with no correlated agreement at threshold $(\rho + \eta)n$, and more than 2^{57} nearby scalar combinations. Its certified scalar count exceeds $n2^{H_2(\rho)/\eta}$ by more than 400. The radius remains strictly below Elias, and the same example transfers to an ordinary Reed–Solomon code with a prime-field evaluation domain and quadratic alphabet. Appendix D.2 records the exact count and inequalities.

6.5 Shorter codes give larger separations

The number of selected scalar labels depends on the packet group and the reserved coefficients, but not on the number B of coordinates in each packet. Reducing B therefore preserves that number while reducing the length factor in the proposed line threshold. The gap also changes, so the comparison must be recomputed.

Corollary 6.4 (Compact eighth-rate examples). *Let $p = 2^{31} - 1$ and $\rho = 1/8$. Each row below gives a circle code on μ_n , or any rotation, and a line with no correlated agreement at threshold $T = (\rho + \eta)n$. Write J for its certified number of distinct nearby parameters.*

Alphabet	n	K	η	$J/(n2^{H_2(\rho)/\eta})$
$\mathbb{F}_{p^2}$	2048	256	23/1024	$> 2^{22}$
$\mathbb{F}_{p^4}$	8192	1024	79/4096	$> 2^{34}$

Both radii are strictly below the characteristic- p Elias radius. In the first row, the words, witnesses, and labels lie in $\mathbb{F}_{p^2}$; the absence of correlated agreement holds even against the larger circle code over $\mathbb{F}_{p^4}$.

Proof. Use $(M, B, r, h) = (128, 16, 0, 19)$ in Corollary A.5 and $(256, 32, 1, 37)$ in Theorem A.4. The respective agreement counts are 302 and 1182, exceeding the direction-word root bounds $K + B = 272$ and 1056. The label counts are unchanged from the preceding eighth-rate examples. Appendix D.3 gives integer certificates for the two new radii and comparisons. $\square$

7 Stronger finite certificates

The following refinements strengthen the finite examples without changing our asymptotic claims. All comparisons in this section use the literal constants $c_1 = c_2 = 1$: the list prescription is $2^{H_2(\rho)/\eta}$ and the line prescription is $n2^{H_2(\rho)/\eta}$. The evaluation domain and alphabet are specified separately in each example.

7.1 Exact enumeration at length 64

An exact enumeration strengthens the finite interval example further. Among the 34-subsets $A \subseteq \{0, \dots, 63\}$, exactly

$$5,552,914,238,035$$

satisfy $\sum_{a \in A} a = 1071$ and $\sum_{a \in A} \binom{a}{2} = 22139$. Their monic root polynomials share the coefficients of degrees 34, 33, and 32, the last being $\binom{1071}{2} - 22139$. Subtracting their common prefix therefore gives distinct codewords of degree at most 31, each agreeing with the prefix word at exactly 34 coordinates. This proves the displayed list lower bound over every prime field with $p > 64$. The count uses an exact take-or-omit recurrence in subset size, first moment, and binomial second moment, with reachability pruning. An independent ascending-label recurrence on complementary 30-subsets verifies every histogram entry; all 129 previously computed moments also agree.

This is a 5.05% improvement on the degree-40 polynomial-weight certificate 5285900426578. At $p = 2^{127} - 1$, the original strict below-Elias comparison still applies. The enumeration counts one integer moment class exactly; it does not claim that this is the largest possible Reed–Solomon decoding list.

7.2 Conditional radial weights and a displaced boundary

Proposition 7.1 (Conditional signed radial weights). *Let $\mathcal{A}$ be a family of t -subsets of $\{0, \dots, n-1\}$ sharing their first binomial moment, and put $C = |\mathcal{A}|$. For $j = 2, 3, 4, 5$, choose*

$$b_j(x) = \binom{x}{j} + \sum_{i < j} c_{ji} \binom{x}{i}, \quad c_{ji} \in \mathbb{Q}.$$

For a uniform $A \in \mathcal{A}$, write $B_j = \sum_{a \in A} b_j(a)$. Let $U_2, \dots, U_5$ be independent uniform variables on $[-1/2, 1/2]$, independent of A , and choose positive rational scales s_j . Set

$$X_j = \frac{B_j - \mathbb{E}B_j + U_j}{s_j}, \quad S = \sum_{j=2}^5 X_j^2, \quad \mu_r = \mathbb{E}S^r.$$

For $w \in \mathbb{Q}[T]$ with integrable positive part on $[0, \infty)$, suppose $E = \mathbb{E}[w(S)] > 0$. Then some class sharing all five binomial moments has size at least

$$\left\lceil \frac{CE}{\pi^2 (\prod_{j=2}^5 s_j) \int_0^\infty w(u)_+ u \, du} \right\rceil.$$

Consequently, for every prime $p > n$, the ordinary Reed–Solomon code on $\{0, \dots, n-1\}$ with dimension $t-5$ has a list of that size at agreement threshold t .

Proof. Because the size and first moment are fixed, the remaining moment vectors lie in an affine lower-triangular lattice with diagonal entries one. Its translates of $[-1/2, 1/2]^4$ tile $\mathbb{R}^4$: successively rounding coordinates determines the lattice point uniquely. If L is the largest fiber, the density of $B+U$ is at most L/C , and that of X is at most $(L/C) \prod_j s_j$.

Radial integration in four dimensions gives

$$\int_{\mathbb{R}^4} w(\|x\|^2)_+ \, dx = \pi^2 \int_0^\infty w(u)_+ u \, du.$$

Replacing the signed weight by its positive part and applying the density bound proves the inequality. Triangular inversion fixes the first five integer binomial moments. Newton identities then fix the

leading six terms of the monic support locators. Subtracting those locators from their common prefix gives distinct polynomials of degree less than $t - 5$, agreeing exactly on their supports. Distinctness after reduction follows from uniqueness of the monic locator of a subset of distinct field elements. $\square$

Choosing the weight. The subclass $w(u) = (R - u)P(u)^2$ is convenient. For $P(u) = \sum_{i=0}^d c_i u^i$, its expectation and radial integral are quadratic forms with matrices

$$A_{ij}(R) = R\mu_{i+j} - \mu_{i+j+1}, \quad B_{ij}(R) = \frac{R^{i+j+3}}{(i+j+2)(i+j+3)}.$$

The matrix $B(R)$ is positive definite. At fixed R , its generalized Rayleigh quotient gives the largest ratio in this polynomial class. Numerical eigenvectors propose a weight; rational rounding and exact evaluation certify the selected weight independently of the optimization.

A verified M31 instance. Take $p = 2^{31} - 1$, $n = 157$, $t = 68$, and condition on $\sum_{a \in A} a = 5304$. The conditional population is

$$C = 4037124019196117766480536278705909334612974.$$

Using normalized population Gram polynomials of degrees 2 through 5, the scales

$$(s_2, s_3, s_4, s_5) = (5751, 76272, 754066, 5946077),$$

and a degree-seven weight

$$w(u) = - \prod_{i=1}^7 (u - r_i),$$

with positive ordered rational roots approximately

$$(r_1, \dots, r_7) \approx (2.78693, 5.20044, 5.74234, 11.14408, 11.19106, 20.21394, 20.21451),$$

gives a list of at least

$$28165503856884755099.$$

Exact rational logarithm intervals verify that the radius is strictly below Elias and that the list exceeds $2^{H_2(63/157)/(5/157)}$ by more than $2^{34.09946}$. The positive intervals are $[0, r_1]$, $[r_2, r_3]$, $[r_4, r_5]$, and $[r_6, r_7]$, so their integrals are exact polynomial endpoint differences. The certificate uses 280 scalar directions through degree fourteen and exact cube-noise moments. Of 560 subset-mode outputs, 351 were counted directly and 209 obtained through independently checked reflections. All complement, prior-moment, and projection-polynomial identities pass. The exact rational roots, rather than their displayed decimal approximations, are in `research/overnight_2026-09-16/certificates/radial_moment_optimum7_verification.json`.

Limit of the truncated radial calculation. An exact rational contraction certificate proves existence of ordered roots and a constant $\lambda > 0$ satisfying

$$\mu_j = \frac{\lambda}{j+2} \sum_{i=1}^7 (-1)^{i+1} r_i^{j+2}, \quad 0 \leq j \leq 7.$$

The density equal to λ/π^2 on these four radial bands has the prescribed moments. It therefore supplies an upper limit on every degree-at-most-seven radial weight certificate; the corresponding signed root polynomial attains it. The validated interval identifies the same rounded integer bound displayed above. This is an optimum of the truncated radial moment method, not an upper bound on the actual RS list.

Proposition 7.2 (Displacing a radial boundary). *Let $X \in \mathbb{R}^d$ have density at most M , where $d \in \{3, 4\}$. Let $0 < r_1 < \dots < r_7$, put $l(s) = \prod_{j=2}^7 (s - r_j)$, and define*

$$w_0(s) = l(s)(r_1 - s), \quad w_v(x) = l(|x|^2)(r_1 - |x|^2 + v \cdot x).$$

Suppose $|v|^2 r_2 < (r_2 - r_1)^2$. Write $a = v/2$ and $R = r_1 + |a|^2$. Then

$$\begin{aligned} \int_{\mathbb{R}^d} (w_v)_+ &= \int_{\mathbb{R}^d} (w_0(|x|^2))_+ - \int_{|x|^2 < r_1} w_0(|x|^2) dx \\ &\quad + \int_{|y|^2 < R} (R - |y|^2) l(|y + a|^2) dy. \end{aligned}$$

In particular, if $\mathbb{E}w_v(X) > 0$, then $M \geq \mathbb{E}w_v(X) / \int (w_v)_+$.

Proof. Completing the square gives $r_1 - |x|^2 + v \cdot x = R - |x - a|^2$. The assumed inequality places this displaced ball strictly inside $\{|x|^2 < r_2\}$. The other positive annuli remain unchanged. On each origin-centered annulus the added term $(v \cdot x)l(|x|^2)$ integrates to zero. Substituting $x = y + a$ on the displaced ball proves the integral identity. Finally, $\mathbb{E}w_v \leq \mathbb{E}(w_v)_+ \leq M \int (w_v)_+$. $\square$

For the normalized smoothed moment lattice, $M \leq (L/C) \prod_j s_j$. The required expectations use no additional subset counts: differentiating an identity $|x|^{2j+2} = \sum_u c_u (u \cdot x)^{2j+2}$ gives $x_i |x|^{2j} = \sum_u c_u u_i (u \cdot x)^{2j+1}$. The displaced-ball integral is exact by angular moments or Cartesian monomial integration. At $n = 157$, $k = 63$, $t = 68$, the verified four-dimensional certificate gives $L \geq 28169451256663519418$ over $\mathbb{F}_{2^{31}-1}$, with finite-prescription excess greater than 34.09967 bits. The three-dimensional certificate gives source class 508108952862977950462384886 and quartic-extension line bank 508108952659354448051620618, with excess greater than 50.90955 bits. The selected rational vectors, exact integrals, and separate raw-moment checks are recorded in the ancillary displaced-boundary certificates.

For the three-dimensional version, retain Gram coordinates of degrees 2, 3, 4 after conditioning on the first moment. The radial integration factor is 2π , and its radial density is $s^{1/2}$. The same triangular lattice argument certifies a class with four fixed integer moments. Its locator prefix is fixed through degree 64, so its residuals have degree at most 63. The rational-line lemma lowers that degree to 62. The quartic-extension instance therefore has length 157, dimension 63, and 68 agreements, with common agreement exactly 63 and selected-graph concurrency at most 18. Its evaluation points remain in $\mathbb{F}_p$; its received words and labels use $\mathbb{F}_{p^4}$. The Elias comparison uses the characteristic $p = 2^{31} - 1$.

7.3 Exact incidences and prime-alphabet lines

Proposition 7.3 (Coordinate and next-coefficient collision savings). *Let N distinct support locators share a prefix leaving residual degree at most k over $\mathbb{F}_p$, on an evaluation domain of size $n < p$. Let I_x count the supports containing x . If their next coefficients take at most K values, then some pole outside the domain has at most*

$$\left\lceil \frac{k \binom{N}{2} - \sum_x \binom{I_x}{2} - P(N, K)}{p - n} \right\rceil$$

colliding unordered pairs, where $P(N, K) = K \binom{a}{2} + ab$ for $N = aK + b$, $0 \leq b < K$.

Proof. At least $P(N, K)$ pairs share the next coefficient and have difference degree at most $k - 1$; the others have degree at most k . Shared support roots account for exactly $\sum_x \binom{I_x}{2}$ roots of these differences on the evaluation domain. Subtracting them and averaging over the remaining $p - n$ poles proves the assertion. $\square$

The least J for which $P(N, J)$ does not exceed this pair budget is a guaranteed number of distinct line labels. For the interval $\{0, \dots, 81\}$, the selected twelve-subset class has first sum 486, second binomial moment 12792, and $N = 221190005$. Its exact coordinate incidences and a rationally certified next-moment range give $J = 138752510$ over $\mathbb{F}_{2^{31}-1}$, with $k = 9$, $t = 12$, common agreement 9, and concurrency at most 24. This exceeds the stated finite line prescription by more than $2^{7.04613}$. The largest two-moment class has more supports, but yields a weaker line certificate.

Here the source locators share the prefix $X^{12} - 486X^{11} + 105063X^{10}$. The exact shared-root count is $\sum_x \binom{l_x}{2} = 43128963788215318$. Write $Z = \sum_{x \in A} \binom{x}{3}$. Newton identities give

$$e_3 = \binom{486}{3} + (2 - 486)12792 + 2Z.$$

Exact rational dual inequalities place Z between 222145 and 276503. Explicitly, for $\sigma \in \{1, -1\}$ and rational multipliers $\lambda = (\lambda_0, \lambda_1, \lambda_2)$, every selected support satisfies

$$\sigma Z \geq 12\lambda_0 + 486\lambda_1 + 12792\lambda_2 + \sum_{x=0}^{81} \min \left\{ 0, \sigma \binom{x}{3} - \lambda_0 - \lambda_1 x - \lambda_2 \binom{x}{2} \right\}.$$

This follows termwise from $z_x r_x \geq \min\{0, r_x\}$ for $z_x \in \{0, 1\}$. The choices $(\sigma, \lambda) = (1, (1421, -509, 107/3))$ and $(-1, (-9568, 2447/3, -130/3))$ give the exact lower bounds $666433/3$ and $-829511/3$, respectively. Integer rounding gives the stated interval for Z . Thus at least 449907051360 source pairs share the next coefficient, and the averaged allowed-pole collision budget is 82437495. Inverting the balanced-pair function gives the stated 138752510 labels. The coordinate counts have a separate formal-division check, using the same underlying histogram engine in a different order. The next-coefficient bounds use rational termwise inequalities, independently of the exploratory minimum/maximum recurrence.

7.4 Stronger native circle examples

Proposition 7.4 (A refined circle-pencil bank). *Let N distinct residual polynomials R_A be obtained from the anchored h -subsets of $G = \mu_M$ in one reciprocal coefficient class, and put $d = h - 2r - 2$. For integers $v \geq 0$ and $b \geq 1$, define*

$$\Pi(v, b) = (b - e) \binom{a}{2} + e \binom{a+1}{2}, \quad v = ab + e, \quad 0 \leq e < b.$$

Set

$$C = d \binom{N}{2} - \binom{N}{2} - \Pi(N(h-1), M-1) - (p+1-M)\Pi(N, p).$$

For poles in $\mathbb{F}_{p^2} \setminus \mu_{p+1}$ take $Q = p^2 - p - 1$; for poles in $\mathbb{F}_{p^4} \setminus \mathbb{F}_{p^2}$ take $Q = p^4 - p^2$. There is an admissible pole with at least

$$\min\{j \in \{1, \dots, N\} : \Pi(N, j) \leq \lfloor C/Q \rfloor\}$$

distinct labels. The agreement and absence-of-correlated-agreement conclusions of the original construction are unchanged.

Proof. The function $\Pi(v, b)$ is the minimum number of colliding unordered pairs in a partition into at most b classes. For each $y \in G$, let s_y count supports containing y . Every shared support point is a root of $R_A - R_{A'}$, so double counting gives

$$\sum_{A < A'} |A \cap A'| = \sum_{y \in G} \binom{s_y}{2} \geq \binom{N}{2} + \Pi(N(h-1), M-1).$$

At every norm-one point, reciprocity gives $R_A(y)^p = c^{-1}y^{-d}R_A(y)$ with the same c for the entire class. Thus at most p evaluation values occur there. Each of the other $p+1-M$ norm-one points contributes at least $\Pi(N, p)$ pair roots. These points and G are disjoint. Subtracting their contributions from the total degree budget $d \binom{N}{2}$ bounds the sum of pair collisions over admissible poles by C . Some pole has at most $\lfloor C/Q \rfloor$ collisions. Balancing its label multiplicities proves the assertion. $\square$

For $p = 2^{31} - 1$, $M = 128$, $h = 19$, and $r = 0$, this gives

$$J \geq 365153906390146587$$

with quadratic-field poles, improving the earlier 285003988493147037. At length 2^{26} and rate $1/8$, the bank exceeds the compared finite prescription by more than 2^9 . At length 2048, dimension 256, and 302 agreements, it exceeds that prescription by more than 2^{23} . Both radii are strictly below the characteristic-based Elias bound. These numerical statements have exact integer and rational certificates in `research/overnight_2026-09-16/certificates/circle_threshold_verification.json`.

Counting omitted-support collisions as well. There is a further exact saving at each $y \in G$. Exactly s_y residuals take the value $v(y)$, whereas the others occupy at most $p-1$ values in the same one-dimensional prime-field subspace. Their collision count is therefore at least

$$f_N(s_y) = \binom{s_y}{2} + \Pi(N - s_y, p-1).$$

This counts disjoint pairs and can replace the support-root term above. Its forward difference is $s - \lfloor (N - s - 1)/(p-1) \rfloor$, an increasing function, so balancing unknown incidences gives a valid lower bound. At reserve $r = 0$, an exact cyclic subset-product count supplies the class population and every s_y . For M dyadic, if $C_k(c)$ counts k -subsets of $\{1, \dots, M-1\}$ with exponent sum $c \pmod{M}$, then

$$C_k(c) = \frac{1}{M} \sum_{a|M} \mathcal{R}_a(c) (-1)^{k+\lfloor k/a \rfloor} \binom{M/a-1}{\lfloor k/a \rfloor},$$

where $\mathcal{R}_a$ is the Ramanujan sum. This follows by extracting the coefficient of z^k in $(1 - (-z)^a)^{M/a} / (1 + z)$ for each character. For a nonanchor exponent b , removing its generating factor gives

$$I_b = \sum_{j=0}^{h-2} (-1)^j C_{h-2-j}(c - (j+1)b).$$

Applying these exact populations and collisions gives the following banks. All rows have $p = 2^{31} - 1$, native circle domains, and radii strictly below the characteristic-based Elias bound.

Alphabet	n	K	T	Guaranteed labels	Excess bits
$\mathbb{F}_{p^2}$	2048	256	302	365153907657996934	> 23
$\mathbb{F}_{p^2}$	4096	64	86	334185312123996677	> 24.5949
$\mathbb{F}_{p^4}$	8192	1024	1182	73497635817388704528460	> 34.77733

The first row improves the same-parameter quadratic bank in Appendix D.3 by about 28.1%. The last two rows maximize the checked guarantee within the stated one-anchor dyadic family, using complete product classes at $r = 0$ and generic prefix pigeonholing at $r > 0$. This is an optimization of that formula, not of arbitrary circle lists or actual coefficient classes. The checked range is dyadic $M \geq 4$, $B \geq 2$, $MB \leq 2^{30}$, dyadic packet dimensions $K/B < M$, and every reserve satisfying the strict Elias condition and the construction’s support-size requirements. A separate coefficient-sum implementation checks all 703 computed product classes and 1406 collision certificates.

Scope and reproducibility. The interval-domain results and the native circle results above are different constructions. None lowers the pinned better.codes display of 116.13, and none is a revised security bound for the complete deployed Stwo protocol. Their contribution is a stronger set of counterexamples to the specified finite coding prescriptions. The proof statements, portable inputs, and verification commands are in `research/overnight_2026-09-16/README.md`.

8 Implications for hash-based SNARK security

The preceding results concern properties of codes. Their relevance to SNARKs comes from the role those properties play in a soundness proof. This section explains that connection before considering a particular protocol.

8.1 The coding statement used by a reduction

An interactive oracle proof can ask a prover to supply evaluations of polynomials and then test those evaluations through a small number of queries. In a hash-based argument, Merkle commitments bind the prover to the supplied words. The algebraic analysis must still explain what passing the tests says about those words. List decoding and proximity gaps are two ingredients in such analyses [7, 8].

For example, consider a code $C \subseteq E^n$, two fixed words $f_0, f_1 \in E^n$, and a uniformly random challenge $z \in E$. Define

$$B_\theta(f_0, f_1) = \{z \in E : \text{dist}(f_0 + zf_1, C) \leq \theta\}.$$

Suppose a coding theorem guarantees that, unless f_0 and f_1 have correlated agreement with two codewords, the set B_θ has size at most T . Then, for words without that correlated agreement,

$$\Pr_z[\text{dist}(f_0 + zf_1, C) \leq \theta] \leq \min\{1, T/|E|\}.$$

This is a bound on a specified algebraic event. Using it in a soundness proof also requires a reduction from verifier acceptance to that event, with the words fixed before the challenge is sampled. The evaluation field, coefficient field, and challenge field must be the fields covered by the coding theorem.

The size of a decoding list enters a different counting step. If an analysis associates a received word with at most L candidate codewords and proves a failure bound for each candidate, it can apply a union bound over those candidates. A numerical estimate for L is therefore part of the quantitative argument. Replacing the actual list size by a conjectured smaller value can invalidate that argument. It does not determine the probability of the verifier’s acceptance event.

8.2 What the counterexamples rule out

Three distinctions are needed to apply the results correctly.

Minimum distance and list size. The root bound proves the usual minimum-distance guarantee for Reed–Solomon codes. It does not by itself establish the proposed quantitative list bound beyond the unique-decoding radius. Theorem 2.2 rules out a uniform $\exp(O_\rho(1/\eta))$ bound over prime fields and evaluation domains, even below the characteristic-based Elias radius. An analysis that needs such a bound must impose further restrictions on the code or use a different quantitative estimate.

A list and the nearby points on a line. A maximum-list bound concerns one received word at a time. It does not say that the nearby codewords for different points on an affine line belong to a common small family. Theorem 5.1 makes this distinction explicit: the global maximum list size is $n - k - 1$, while one affine line has $\binom{n}{k+1}$ nearby points, each with a unique witness. Its gap is $1/n$. This theorem does not settle the analogous question with a separately fixed positive gap.

Uniform estimates and fixed-parameter estimates. A security calculation uses a particular field, domain, rate, and radius. A theorem quantified over all prime-field codes would apply to that code, but a counterexample to the universal theorem need not use that particular code. Conversely, a bound with constants depending on the field or the gap cannot be used uniformly without controlling that dependence. The fixed-gap lifts in Theorems 4.4 and 4.6 rule out prescribed numerical coefficients with remainders sublinear along their increasing-length families. The anchored version also prevents correlated agreement altogether at its threshold. Their fields grow, and their chosen seed lists remain fixed. These quantifiers are part of the result.

8.3 From coding counterexamples to protocol conclusions

A coding counterexample is enough to refute a coding premise under the same quantifiers. A conclusion about a complete SNARK requires more. Its oracle words may be constrained by an arithmetization, previous commitments, claimed evaluations, and later consistency tests. A soundness analysis must account for those constraints and the distribution and order of all challenges. Passing to a noninteractive protocol through Fiat–Shamir introduces a further cryptographic analysis.

The circle constructions give finite examples in polynomial spaces used by Circle STARKs. The quotient calculations in Appendices B and C examine additional algebraic conditions on derived words. Their statements retain their own sampling, radius, and conditioning hypotheses. In particular, the numerical raw-quotient example lies above the Elias radius, whereas the principal list and affine-line examples lie below it. No full-protocol security level follows from the quotient counts in this paper.

The application below uses this distinction to identify which quantitative premises of a published analysis are contradicted and which questions about the instantiated protocol remain open.

9 Application to S-two’s quantitative coding assumptions

The S-two whitepaper uses list- and line-decodability assumptions in its analysis of polynomial proximity tests [30]. This provides one application of the preceding coding results. We separate three questions: the asserted uniform list bound, the coefficient in a line-decodability threshold, and finite numerical instantiations. They require different counterexamples and different quantifiers.

9.1 The uniform list prescription

Conjecture 1 of the whitepaper takes the form

$$\ell_C(\theta) \leq c_1 2^{c_2 H_2(\rho)/\eta}, \quad \theta = 1 - \rho - \eta, \quad (36)$$

with constants depending only on the rate and a stated range below the characteristic-based Elias radius. Under its universal prime-field and evaluation-domain quantifiers, this prescription is false for every fixed pair of finite positive constants. Theorem 2.2 gives a single sequence with

$$\log_2 \ell_C(\theta) = \Omega_\rho(n), \quad 1/\eta = \Theta_\rho(\sqrt{n \log n}).$$

The logarithm of the right-hand side of (36) is sublinear in n , whereas the constructed list has a logarithm linear in n . Changing fixed constants cannot close that gap.

Even a polynomial dependence of the prefactor on the prime cannot repair a constant exponent coefficient. More precisely, suppose the prescription is required for every eligible code over each sufficiently large $\mathbb{F}_p$, allowing $c_1 = c_1(p) \leq p^K$ for a fixed K and $c_2 = c_2(p)$. Corollary 2.3, with $b = \log_2 p$, forces

$$c_2(p) \geq (1/2 - o(1)) \frac{\log_2 p}{\log_2 \log_2 p}.$$

Indeed its list logarithm is at least $(1/2 - o(1))b^2/\log_2 b$, whereas the proposed logarithm is at most $Kb + (1 + o(1))c_2(p)b$.

This conclusion uses growing prime fields and a decreasing gap. It does not assert exponential lists at a fixed positive gap, and it does not rule out stronger guarantees for a restricted class of evaluation domains or a restricted relation between field size and length.

Proposition 2.5 gives a separate obstruction at rate $1/4$ and gap $1/8$: when $n = p - 1$, even a finite list bound depending only on rate and gap is impossible. The short-domain theorem above shows that restricting to $n = o(p)$ does not rescue the proposed reciprocal-gap exponent.

9.2 The quantitative line coefficient and its remainder

Conjecture 2 in Appendix A.5 of the whitepaper has a threshold of the form $a = \ell(\theta)n + o(n)$, with concurrency number n . Substituting the proposed quantitative list bound gives

$$a = c_1 2^{c_2 H_2(\rho)/\eta} n + o(n). \quad (37)$$

Here the asymptotic remainder matters. A finite example, or a sequence whose gap decreases, need not refute a remainder statement made only for each separately fixed gap.

Theorem 4.4 addresses that issue directly. For any fixed constants, it preserves the prescribed rate and chooses a gap fixed along an increasing-length family. A preserved list has size $L > A$, where $A = c_1 2^{c_2 H_2(\rho)/\eta}$. The family has Ln selected nearby witnesses, but no parametrized affine codeword graph contains more than n selected pairs. Its discrepancy from the proposed coefficient is

$$(L - A)n.$$

Since $L - A > 0$ is fixed, no remainder that is $o(n)$ along this family can absorb the discrepancy. The family is selected independently of the remainder. Here $o(n)$ means sublinear along admissible increasing-length families, allowing the prime field to grow; the whitepaper does not explicitly state its field-uniformity convention. The theorem is stated for constants at least one; any smaller positive constants are covered by enlarging them to one.

Theorem 4.6 strengthens this conclusion: its lines have no correlated agreement at their exactly fixed rate and gap, yet have at least $C_\rho(\eta)n$ nearby challenges with $\log_2 C_\rho(\eta) = \Omega_\rho(\eta^{-2}/\log(1/\eta))$ at every sufficiently small rational gap. Their selected witnesses also have concurrency at most n . Thus the fixed-gap coefficient obstruction does not depend on selecting among persistent good witness lines.

The shrinking-gap constructions show how large the numerical failure can be. Corollary 4.15 makes every nonzero parameter nearby while the exception is $(1 - o(1))\eta$ outside the radius, only one coordinate short of maximum distance. The proposed numerical fraction tends to zero. Separately, Theorem 4.41 has a length-214, half-rate certificate over $2^{127} - 1$ that compares nearby fraction $1 - 1/p$ with a prescription below 2^{-12} ; the larger certificate over $2^{9689} - 1$ gives a prescription below 2^{-593} . These comparisons concern the prescribed numerical coefficient; they do not replace the fixed-gap argument needed to address a remainder controlled separately at each gap.

These arguments do not establish failure of a fixed-gap threshold expressed using the *actual* maximum list size $\ell_C(\theta)$. In the preliminary construction, for example, the selected count is $Ln \leq \ell_C(\theta)n$. Replacing a false numerical estimate for list size and retaining a theorem that uses actual list size are different operations.

Theorem 5.1 addresses the actual-list question separately. At every fixed rate it has $\ell_C(\theta) = n - k - 1$ but exponentially many uniquely nearby line points and maximum concurrency $n - k$. At rate one-half these bounds are $n/2 - 1$ and $n/2$, respectively. Thus even the threshold $\ell_C(\theta)n + o(n)$ fails along this shrinking-gap family. This rules out a familywise interpretation of that threshold; it does not settle a version whose remainder is controlled only after a positive gap has been fixed. The original conjecture does not specify this uniformity convention.

9.3 Finite comparisons on the circle code

The whitepaper’s concrete examples set $c_1 = c_2 = 1$ and omit the asymptotic remainder. At the common parameters of Theorem 6.1, the resulting comparison values are $2^{1/\eta}$ for a list and $n2^{1/\eta}$ for the line threshold. Let L_0 and B_0 denote the list-size and scalar-count lower bounds certified in Appendix A. They give

$$\frac{L_0}{2^{1/\eta}} > 500, \quad \frac{B_0}{n2^{1/\eta}} > 200,000.$$

These are ratios of coding-theoretic counts. They are not multipliers for a complete attack’s success probability and are not security-bit losses. In particular, the logarithm of a list-size ratio is not, by itself, a reduction in a proof system’s security level.

The two bounds also concern the displayed radius. A parameter-table row evaluated at a different gap must be checked at that gap. For example, the value 2^{16} at rate $1/2$ corresponds to $\eta = 1/16$, which differs from $33/1024 - 2/n$. A counterexample to a universal prescription does not automatically supply a counterexample to every individual numerical row.

The circle tensor space and canonical full-coset domains described above are the algebraic spaces appearing in the examined public proving source [31]; Proposition A.12 and Proposition A.13 explain their mathematical identification. This connection makes the fixed-prime examples pertinent to the coding assumptions used for that family of tests.

Concrete parameters and the challenge schedule. The inspected production configuration uses a nominal 96-bit target: 26 bits of query grinding, 70 queries, and blowup factor two. The source’s parameter counter computes $26 + 70 \cdot 1 = 96$, and the project documentation calls its default soundness conjectured [32]. This arithmetic records the parameter convention; a security

theorem must justify the contributions attributed to each stage. The whitepaper’s illustrated 100-bit configuration is a different choice [30, Section 5.5].

There is a further distinction between that illustrated calculation and the cited implementation. The former allocates grinding to intermediate challenge stages, while the latter performs query grinding after those challenges and omits the corresponding intermediate checks [33]. Grinding after a challenge cannot be credited to an earlier stage without a supporting argument. A quantitative analysis of the implementation must therefore use its actual challenge schedule and accepted parameters. This issue is separate from the counterexamples to the coding assumptions, and the comparison concerns only the pinned source revisions.

The remaining protocol question. A proximity test inside a proof system does not necessarily receive an arbitrary pair of words. Its inputs can be derived quotient polynomials constrained by earlier commitments, algebraic identities, and sampled evaluation points. Appendix C resolves one of these algebraic conditions: fixed tables and fixed, generally false claimed openings produce raw quotient pencils with a certified joint proximity event. The transfer also handles specified nonlinear batching schedules (Theorem C.4), with explicit quotient-layer realizations. Its numerical raw-quotient example lies above the Elias radius (Section C.4); it is distinct from the below-Elias counterexamples. It does not prove that these inputs satisfy the complete set of commitment, AIR, and FRI conditions. Nor does a count of exceptional scalars determine the cost of finding favorable claims and witnesses or arranging a preceding transcript. Thus the coding results invalidate the stated quantitative justifications under their stated quantifiers; turning them into an attack on a complete protocol is a separate problem.

10 Further coding-theoretic questions

The main open questions concern the restrictions needed for useful coding bounds and their role in concrete soundness analyses.

1. **Restrictions on field size and domain.** What quantitative bounds hold under a prescribed restriction $p \leq n^D$, with D fixed independently of the gap, or on a specified multiplicative subgroup or circle coset? Corollary 4.7 already realizes the fixed-gap coefficient obstruction with polynomial-size fields, but its field exponent depends on the gap. The domain-specific questions remain distinct.
2. **The correct dependence on the gap.** For a prescribed relation between p and n , what dependence on η is necessary? The exponential-list family has shrinking gap; the fixed-gap lifts retain fixed seed lists rather than exponentially growing ones. The anchored lift forces the logarithm of any universal linear proximity-gap coefficient to grow at least as $\eta^{-2}/\log(1/\eta)$ at every sufficiently small rational gap. Whether a fixed positive gap permits superlinear growth in length remains open here. Theorem G.1 excludes all candidate families with boundedly many distinct roots as a route to such growth in sufficiently large characteristic, even allowing varying root positions and multiplicities. A third-order equation within this family still has an actual solution variety of quadratic degree.
3. **Bounds for a prescribed code.** Which list and proximity-gap bounds hold for a fixed field and a fixed FFT domain at the radii used in a proof system? A generic-domain existence theorem and a counterexample on an integer interval leave this question open. Such a theorem should expose its finite constants and its dependence on the challenge field.

4. **Complete coverage with unique nearby codewords.** Can the full punctured-line behavior of Corollary 4.15 coexist with a unique nearby codeword at every nonzero parameter? The cubic construction necessarily has local ambiguity: at a single padding coordinate, more than p candidates produce only $p-1$ labels. It also has very large global lists. The paired constructions in Theorems 4.22, 4.24, and 4.30 have whole-line uniqueness and exponentially many nearby parameters, but do not give complete nonzero coverage. The fixed-gap construction gives whole-line uniqueness with a linear lower bound on the nearby count.
5. **Recovering nearby witnesses.** Can one efficiently recover nearby codewords for the full-coverage constructions, or enumerate enough of the equal-moment classes to match the counting lower bounds? The deterministic powers-of-two family in Theorem 4.24 and the orbit family in Theorem 4.27 have efficient witness-recovery algorithms. Lemma 4.12 and Proposition 4.42 output the code and line in expected polynomial time in $\log p$, but do not recover a witness for a supplied nonzero parameter. Theorem F.1 limits one possible description: an affine binary family with a fixed locator prefix has dimension at most $\lfloor \min\{t, n-t\}/(s+1) \rfloor$. At a fixed positive prefix fraction this permits only a constant-size family.
6. **Lines and actual list size.** What line-decodability statements follow from the actual list size at a fixed positive gap? Theorem 5.1 rules out the proposed linear threshold along a shrinking-gap family, even with unique nearby witnesses on the selected line.

A Circle-code constructions in detail

A.1 Base-field lists for circle evaluation codes

We first construct many codewords close to a single received word. The construction uses the roots of polynomials to choose agreement positions, then fixes a few coefficients so that the remaining polynomial parts have low degree. On the norm-one circle, the leading and trailing coefficients are linked. This lets one impose both sets of conditions at the cost of fixing only one set.

A.1.1 The norm-one circle and Laurent polynomials

Let $p \equiv 3 \pmod{4}$ be prime. Since -1 is not a square in $\mathbb{F}_p$, adjoining a root i of $i^2 = -1$ gives the quadratic extension $F_2 = \mathbb{F}_{p^2}$. Its nontrivial field automorphism is $a \mapsto a^p$, and $i^p = -i$. Thus, for $x, y \in \mathbb{F}_p$,

$$(x + iy)^p = x - iy, \quad (x + iy)^{p+1} = x^2 + y^2.$$

The solutions of $x^2 + y^2 = 1$ can therefore be written as the group

$$U = \{u \in F_2^\times : u^{p+1} = 1\}.$$

This group is cyclic of order $p+1$: it is the kernel of the norm map from $F_2^\times$ to $\mathbb{F}_p^\times$. For $n \mid p+1$, let $\mu_n \subseteq U$ be its subgroup of order n . Multiplying all points by $g \in U$ gives a rotated domain $g\mu_n$.

The coordinate $u = x + iy$ simplifies computations on the circle because $u^{-1} = x - iy$. In particular,

$$x = \frac{u + u^{-1}}{2}, \quad y = \frac{u - u^{-1}}{2i}.$$

A *Laurent polynomial* is a finite sum $\sum_j a_j u^j$ in which negative exponents are allowed. They cause no poles on U , since $u \neq 0$. Its support is the set of exponents with nonzero coefficients. We call a

Laurent polynomial with coefficients in F_2 *base-field-real* when

$$a_{-j} = a_j^p \quad \text{for every } j.$$

This coefficient condition makes its values lie in $\mathbb{F}_p$ on U : taking the p th power both conjugates the coefficients and replaces u by u^{-1} . The word “real” here denotes this finite-field condition, not a condition on real numbers.

For even $K = 2s$, the circle polynomial space used below is

$$\mathcal{C}_K = \{A(x) + yB(x) : A, B \in \mathbb{F}_p[x], \quad \deg A < s, \deg B < s\}.$$

It has dimension K as a polynomial-function space on the algebraic circle. Whenever the stated evaluation domain is large enough to make evaluation injective, evaluating these functions gives a dimension- K linear code. We use the term *circle FFT space* for this space; its tensor basis is explained in Proposition A.12. The constructions below lie in its slightly smaller interior Laurent subspace. Working in that subspace makes rotation straightforward.

For a length- n , dimension- K evaluation code, write $\rho = K/n$ for its rate. A word agreeing with a received word in at least $(\rho + \eta)n$ positions lies at relative Hamming distance at most $\theta = 1 - \rho - \eta$. Thus η measures the agreement beyond the rate. Throughout, H_2 denotes binary entropy, and

$$H_p(t) = t \log_p(p-1) - t \log_p t - (1-t) \log_p(1-t).$$

The characteristic- p Elias radius at rate ρ is the solution $r_E(\rho)$ of $H_p(r_E) = 1 - \rho$ on the increasing branch of H_p .

Lemma A.1 (Interior band of the circle FFT space). *Let $p \equiv 3 \pmod{4}$ be prime, let $i^2 = -1$ in $\mathbb{F}_{p^2}$, and write $u = x + iy$ on the norm-one circle. For even $K = 2s$, the dimension- K circle FFT space*

$$\mathcal{C}_K = \mathbb{F}_p[x]_{<s} + y\mathbb{F}_p[x]_{<s}$$

contains every base-field-real Laurent polynomial supported on $[-s+1, s-1]$. This interior subspace is preserved by rotations of the circle.

Proof. For $0 \leq j < s$, the functions $u^j + u^{-j}$ are polynomials in x of degree j , and $(u^j - u^{-j})/i$ are y times polynomials in x of degree $j-1$. The real combinations of these functions span the indicated interior band. Replacing u by a norm-one scalar times u preserves the band and reality. $\square$

Theorem A.2 (Reciprocal-reserve real circle lists). *Let $p \equiv 3 \pmod{4}$ be prime, let $n = MB$ divide $p+1$, and suppose $M \geq 4$ and $B \geq 2$ are powers of two. Choose an even integer d with $0 < d < M$, and an integer $r \geq 0$ such that*

$$h = d + 2r + 2 < M.$$

Set $K = dB$ and $s = K/2$. On $D = \mu_n$, and on any rotation of this domain, there exists an $\mathbb{F}_p$ -valued received word with at least

$$L = \left\lceil \frac{\binom{M-1}{h-1}}{Mp^{2r}} \right\rceil$$

distinct codewords in the exact space $\mathcal{C}_K$ of Lemma A.1, each agreeing with the received word at at least $hB - 2$ positions. The rate and gap are exactly

$$\rho = \frac{d}{M}, \quad \eta = \frac{2r+2}{M} - \frac{2}{n}, \quad \theta = 1 - \rho - \eta.$$

Every witness lies in the strict interior Laurent band, so no boundary-mode identification is needed to transport the construction to a coset.

Proof. For each h -subset $A \subset \mu_M$ containing 1, put

$$F_A(Y) = \prod_{a \in A} (Y - a) = \sum_{j=0}^h f_j Y^j, \quad c = f_0.$$

The integer h is even. Consequently $c \in \mu_M$, and the norm-one roots give the reciprocal coefficient identity

$$f_{h-j} = c f_j^p. \quad (\text{R1})$$

Partition all $\binom{M-1}{h-1}$ anchored subsets by c and the r leading coefficients $f_{h-1}, \dots, f_{h-r}$. There are at most $M p^{2r}$ classes. In a largest class, containing at least L supports, the bottom r coefficients are fixed as well, by (R1).

Write $T_A(Y) = F_A(Y)/Y^{r+1}$. Let $v(Y)$ consist of its common top $r+1$ and bottom $r+1$ Laurent terms. Thus these are precisely the terms coming from $f_0, \dots, f_r$ and $f_{h-r}, \dots, f_h$. The remainder

$$R_A(Y) = v(Y) - T_A(Y)$$

is an ordinary polynomial of degree at most $d = h - 2r - 2$. Distinct supports give distinct R_A , and the anchor gives $R_A(1) = v(1)$.

For norm-one Y , (R1) implies

$$T_A(Y)^p = c^{-1} Y^{-d} T_A(Y). \quad (\text{R2})$$

The involution in (R2) pairs the retained top and bottom blocks, so v and R_A satisfy the same identity. Choose $\alpha \neq 0$ with $\alpha^p/\alpha = c$, possible because c has norm one. Explicitly, $\alpha = 1 + c^{-1}$ works if $c \neq -1$, and $\alpha = i$ works otherwise. Define

$$V(u) = \alpha u^{-s} v(u^B), \quad S_A(u) = \alpha u^{-s} R_A(u^B).$$

Since $dB = 2s$, (R2) proves that V and S_A are fixed by base-field Frobenius and hence are $\mathbb{F}_p$ -valued on the circle. Both B and s are even, so every S_A has the same value $\beta = \alpha v(1) \in \mathbb{F}_p$ at $u = 1$ and $u = -1$.

Let $y(u) = (u - u^{-1})/(2i)$. Away from these anchors define

$$W(u) = \frac{V(u) - \beta}{y(u)}, \quad P_A(u) = \frac{S_A(u) - \beta}{y(u)}.$$

Give W arbitrary base-field values at the two anchors. For the witness, multiplication of its numerator by u^s gives the polynomial

$$\alpha(R_A(u^B) - v(1)u^s),$$

of degree at most K , divisible by $u^2 - 1$. Therefore P_A is a Laurent polynomial supported on $[-s+1, s-1]$, with removable values at the anchors. Its reality, including at the removable points, follows from the Frobenius identity of the rational expression. Lemma A.1 places it in the exact dimension- K circle space. When K is a power of two, this is the corresponding FFT space.

Off the anchors,

$$W(u) - P_A(u) = \frac{\alpha u^{-s} F_A(u^B)}{u^{B(r+1)} y(u)}.$$

Among the nonanchor coordinates, agreement holds exactly on the selected h fibers. The map $u \mapsto u^B$ has fibers of size B on D , so there are exactly $hB - 2$ nonanchor agreements, and at least this many agreements on the full domain. Distinct F_A give distinct P_A ; a nonzero difference in the interior band, after clearing its Laurent monomial, has degree at most $K - 2 < n$, so the codewords are distinct on D . Finally, rotation preserves the real interior band and all agreement counts. $\square$

A.1.2 A numerical example over a Mersenne prime

Take $p = 2^{31} - 1$ and $n = 2^{26}$. For $M = 2048$, $r = 32$, and $h = 1090$, Theorem A.2 gives the exact integer lower bound

$$L = \left\lceil \frac{\binom{2047}{1089}}{2048(2^{31} - 1)^{64}} \right\rceil = 1,199,728,024,606.$$

For comparison, the same theorem with no reserved leading coefficients has $M = 32$, $r = 0$, $h = 18$, and $L = 8,286,954$. Consider the candidate finite list-size bound $2^{H_2(\rho)/\eta}$. At rate $1/2$, its base-two logarithm is $1/\eta$. The following table compares that bound with the two lower bounds.

M	r	h	$\log_2 L$	$1/\eta$	$\log_2 \text{ratio}$
32	0	18	22.98241048	16.00000763	6.98240285
2048	32	1090	40.12584453	31.03033173	9.09551280

The last column is the base-two logarithm of the ratio between the proved list-size lower bound and the candidate upper bound. The pigeonhole argument establishes the existence of the large coefficient class; it does not require enumerating that class.

A.1.3 Strict below-Elias certificate

For the second row, the gap is

$$\eta = \frac{33}{1024} - \frac{2}{2^{26}} = 0.0322265326976776123046875.$$

It is slightly smaller than $1/\log_2 p$, so that convenient sufficient condition must not be used. Direct high-precision evaluation instead gives

$$H_p(1/2 - \eta) = 0.49993479960319703654 \dots < 1/2.$$

The conclusion also has an exact integer/rational certificate. Direct integer arithmetic verifies $p^{100} > 2^{3099}$, hence $\log_2 p > 3099/100$. The binary entropy satisfies $H_2(1/2 - \eta) \leq 1 - 2\eta^2$: its second derivative is at most -4 , and its derivative vanishes at $1/2$. Direct rational arithmetic gives

$$\eta \frac{3099}{100} > 1 - 2\eta^2.$$

Consequently

$$H_p(1/2 - \eta) < (1/2 - \eta) + \frac{H_2(1/2 - \eta)}{\log_2 p} < 1/2.$$

This radius is strictly below characteristic- p Elias. Independently of rounded logarithms, the exact checks $L^8 > 2^{321}$ and $1/\eta < 31031/1000$ certify a list excess greater than 9.094 bits.

A.1.4 A small exact-arithmetic example

An independent implementation over $\mathbb{F}_{31}[i]/(i^2 + 1)$ uses $n = 32$, $M = 16$, $B = 2$, $r = 1$, and $h = 12$. It enumerates all 1,365 anchored supports, finding 1,336 coefficient classes and a largest class of size three. For every member of that class it verifies all reciprocal coefficient identities, real-valued centering, both removable anchors, exact division by $u^2 - 1$, the strict Laurent band $[-7, 7]$, distinctness, and agreement at exactly 22 non-anchor domain positions. This small example checks the algebraic construction; its parameters are not claimed to violate the numerical candidate bound.

The finite-field checks use exact field arithmetic and integers; the displayed logarithms use 100-digit decimal arithmetic. Negative controls with $(M, r) = (512, 7)$ and $(1024, 15)$ have larger counts but are above Elias for this list construction and therefore do not give examples below Elias for this list construction.

A.2 Affine families from simultaneous coefficient constraints

We now seek a different phenomenon: many scalar combinations of two received words have nearby codewords, even though the two original words have no nearby codewords on a sufficiently large common set of positions. Such a scalar-parameterized affine family is also called a *pencil*.

Throughout this section, $p \equiv 3 \pmod{4}$ is prime, $F_2 = \mathbb{F}_{p^2} = \mathbb{F}_p(i)$ with $i^2 = -1$, and $E = \mathbb{F}_{p^4}$. A bar denotes the nontrivial automorphism of $F_2/\mathbb{F}_p$. The norm-one group

$$U = \{u \in F_2^\times : u\bar{u} = 1\}$$

has order $p + 1$. For each $n \mid p + 1$, write μ_n for its unique subgroup of order n . We use the split coordinate $u = x + iy$ on the circle; then $x = (u + u^{-1})/2$ and $y = (u - u^{-1})/(2i)$.

For $m \geq 1$, define the Laurent spaces

$$\mathcal{L}_m^- = \text{span}_E\{u^j : -(m-1) \leq j \leq m-1\}, \quad \mathcal{L}_m^+ = \text{span}_E\{u^j : -m \leq j \leq m\}.$$

If $D \subseteq U$, evaluation of either space on D will be denoted by the same symbol when the domain is unambiguous. We consider any code $C \subseteq E^D$ satisfying

$$\mathcal{L}_m^- \subseteq C \subseteq \mathcal{L}_m^+. \tag{38}$$

The dimension- $2m$ circle FFT space satisfies this condition, as shown in Section A.4. The constructions below place every witness in the smaller space, and rule out correlated agreement even in the larger one.

For words $f, g \in E^D$, let $\text{agr}_D(f, g)$ be the number of coordinates on which they agree. A pair (f_0, f_1) has correlated agreement at threshold T with C if there exist $P_0, P_1 \in C$ agreeing with the two words simultaneously on at least T coordinates. When no such pair exists, a scalar $z \in E$ is *globally bad at threshold T* if

$$\text{agr}_D(f_0 + zf_1, P_z) \geq T$$

for some $P_z \in C$. A *bank* is a set of distinct such scalars, together with one witness for each scalar. Here the bank records only which scalars have nearby codewords; it does not require those codewords to depend affinely on the scalar.

Lemma A.3 (Global failure and line-decodability). *Suppose (f_0, f_1) has no correlated agreement at threshold T on a domain of length n . Select a witness P_z with at least T agreements for each scalar in a bank. No affine line of codewords contains more than $n - T + 1$ of the selected pairs (z, P_z) . In particular, it contains at most n such pairs.*

Proof. Suppose $P_z = A + zB$ for q distinct selected scalars. Let a be the number of coordinates at which both $f_0 = A$ and $f_1 = B$. At every other coordinate, the affine equality $f_0 + zf_1 = A + zB$ can hold for at most one scalar. Hence

$$qT \leq qa + n - a.$$

Since $a \leq T - 1$, this gives $q(T - a) \leq n - a$, and therefore $q \leq n - T + 1$. $\square$

Consequently, a bank larger than a proposed line-decodability threshold contradicts the definition with concurrency number n . This observation will connect the explicit pencils to the finite threshold discussed below.

Theorem A.4 (Simultaneous coefficient reserve pencil). *Let $n = MB \mid p + 1$, where $B \geq 2$ is even. Let $r \geq 0$ and $2r + 4 \leq h \leq M - 1$, and put*

$$a = 2r + 3, \quad K = (h - a)B = 2m, \quad s = m - 1, \quad T = hB - 2.$$

Assume $K + B < T \leq n$, and let C be any code on $D = \mu_n$ satisfying (38). Define

$$N = \left\lceil \frac{\binom{M-1}{h-1}}{Mp^{2r}} \right\rceil, \quad Q = p^4 - p^2, \quad d = h - 2r - 2. \quad (39)$$

There exist words $\gamma_0, \gamma_1 \in E^D$ which have no correlated agreement at threshold T with C , and which have a bank of size at least

$$\left\lceil \frac{NQ}{Q + (d - 1)(N - 1)} \right\rceil. \quad (40)$$

Each selected witness lies in $\mathcal{L}_m^-$, and agrees on at least T coordinates outside $\{1, -1\}$. The same conclusion holds on any rotated domain $g\mu_n \subseteq U$.

Proof. Put $G = \mu_M$. For every h -element subset $A \subseteq G$ containing 1, write

$$F_A(Y) = \prod_{a \in A} (Y - a) = \sum_{j=0}^h f_{A,j} Y^j, \quad f_{A,h} = 1.$$

Partition these subsets by $c = f_{A,0}$ and the r coefficients $f_{A,h-1}, \dots, f_{A,h-r}$. The constant coefficient has at most M possible values, and each other coefficient lies in F_2 , so there are at most Mp^{2r} classes. Fix a class of size at least N , and retain exactly N distinct subsets in it.

The norm-one condition on the roots gives the identity

$$F_A(Y) = cY^h \overline{F_A}(1/Y), \quad f_{A,j} = c \overline{f_{A,h-j}}. \quad (41)$$

Thus the bottom r coefficients are common to the class as well. In $F_A(Y)/Y^{r+1}$, let $v(Y)$ be the sum of the common terms with indices $0 \leq j \leq r$ and $h - r \leq j \leq h$. Then

$$R_A(Y) := v(Y) - F_A(Y)/Y^{r+1}$$

is an ordinary polynomial of degree at most $d = h - 2r - 2$. The polynomials R_A are distinct, and all satisfy $R_A(1) = v(1)$.

Choose a pole $\tau \in E \setminus F_2$, to be specified below. Define

$$f_0(u) = \frac{v(u^B)}{u^B - \tau}, \quad f_1(u) = \frac{1}{u^B - \tau}, \quad \gamma_j(u) = u^{-s} \frac{f_j(u) - f_j(1)}{u^2 - 1}.$$

At $u = \pm 1$ use the removable values. There are no other denominator zeros on D : every u^B lies in F_2 , while $\tau \notin F_2$. For each A , put

$$z_A = -R_A(\tau), \quad Q_A(Y) = \frac{R_A(Y) - R_A(\tau)}{Y - \tau}, \quad P_A(u) = u^{-s} \frac{Q_A(u^B) - Q_A(1)}{u^2 - 1}.$$

The polynomial Q_A has degree at most $d - 1 = h - a = K/B$. Since B is even, its composed numerator vanishes at $u = \pm 1$. Dividing by $u^2 - 1$ leaves degree at most $K - 2$, and multiplying by u^{-s} places P_A in the Laurent band $[-s, s]$. Thus $P_A \in \mathcal{L}_m^- \subseteq C$. At every nonanchor point with $u^B \in A$, the identity $v(u^B) = R_A(u^B)$ holds; it also holds at the subtracted anchor value. Therefore $\gamma_0(u) + z_A \gamma_1(u) = P_A(u)$ at at least $hB - 2 = T$ coordinates outside $\{1, -1\}$.

To rule out correlated agreement, it suffices to rule out sufficient agreement with γ_1 alone. Put

$$S_B(u) = 1 + u^2 + \dots + u^{B-2}.$$

Cancellation yields

$$\gamma_1(u) = -\frac{u^{-s} S_B(u)}{(1 - \tau)(u^B - \tau)}. \quad (42)$$

For $P \in \mathcal{L}_m^+$, the function $H(u) = u^m P(u)$ is an ordinary polynomial of degree at most K . Every agreement with γ_1 is a root of

$$(1 - \tau)(u^B - \tau)H(u) + uS_B(u). \quad (43)$$

This polynomial is nonzero: the second summand has degree $B - 1$, whereas the first has degree at least B when $H \neq 0$. Its degree is at most $K + B$. Hence no word in the entire outer-band code agrees with γ_1 on more than $K + B < T$ coordinates. Formula (42) also holds at the removable anchor values.

It remains to select a pole with sufficiently many distinct labels. For $A \neq A'$, the nonzero polynomial $R_A - R_{A'}$ has degree at most d and has the root $1 \in F_2$. It therefore has at most $d - 1$ roots outside F_2 . Averaging over the Q possible poles, some pole has at most $(d - 1)N(N - 1)/(2Q)$ colliding unordered pairs. If the multiplicities of its distinct labels are $b_1, \dots, b_t$, then

$$\sum b_i = N, \quad \sum b_i^2 \leq N + (d - 1)N(N - 1)/Q.$$

Cauchy–Schwarz gives $t \geq NQ/[Q + (d - 1)(N - 1)]$, proving (40). Finally, substituting u/g transports all words to $g\mu_n$. Rotation preserves both Laurent bands and all agreement counts. $\square$

Corollary A.5 (Poles in the quadratic coefficient field). *Under the hypotheses of Theorem A.4, the defining words, selected witnesses, and scalar labels can all lie in $F_2 = \mathbb{F}_{p^2}$. With*

$$Q_2 = p^2 - M,$$

the guaranteed number of distinct labels is

$$\left\lceil \frac{NQ_2}{Q_2 + (d - 1)(N - 1)} \right\rceil.$$

The pair has no correlated agreement at threshold T , even against outer-band codewords over an extension of F_2 .

Proof. Choose the pole from $F_2 \setminus \mu_M$. The denominators only require $u^B \neq \tau$ and $1 \neq \tau$, so every pole in this set is allowed, including zero. All coefficients and values in the pencil formulas then belong to F_2 . A difference $R_A - R_{A'}$ still has the root 1, which is excluded from the pole set, and therefore has at most $d - 1$ roots among its Q_2 elements. The same collision average proves the count with Q_2 in place of Q . The nonzero polynomial in (43) has degree at most $K + B < T$ over any ambient extension. Its root bound, and hence the absence of correlated agreement, are unchanged. $\square$

Remark A.6 (Existence and explicitness). The formulas specify a concrete family of pencils. The large coefficient class and a sufficiently good pole are guaranteed by counting; a large coefficient class or a suitable pole need not have been enumerated to prove the lower bound. We distinguish this existence-certified bank from a bank whose labels and witnesses have all been listed explicitly.

A.2.1 A numerical affine-family example over a Mersenne prime

Let $p = 2^{31} - 1$ and take

$$M = 1024, \quad r = 15, \quad h = 545, \quad n = 2^{26}, \quad B = 2^{16}.$$

Then $a = 33$, $K = n/2$, $d = 513$, and

$$T = 545B - 2, \quad \theta = 1 - T/n = \frac{479}{1024} + \frac{2}{n}, \quad \eta = \frac{33}{1024} - \frac{2}{n}.$$

The theorem gives the integer bank lower bound

$$30,608,283,697,878,843,892,998, \tag{44}$$

whose base-two logarithm exceeds 74.6963402382.

At exact rate $1/2$, the finite prescription $n2^{H_2(1/2)/\eta}$ has logarithm less than 57.0303317264. The excess in the logarithm is greater than 17.6660085118 bits.

For clarity, define the characteristic- p Elias radius $r_E(\rho)$ by $H_p(r_E) = 1 - \rho$, with

$$H_p(t) = t \log_p(p-1) - t \log_p t - (1-t) \log_p(1-t).$$

The above distance is strictly below $r_E(1/2)$. Equivalently, the quantity

$$(1/2 - \theta) \log_2 p - H_2(\theta) - \theta \log_2(1 - 1/p)$$

is positive; its value exceeds 0.0020212123. The exact count and high-precision entropy comparison are reproduced in the accompanying `counterexample_parameters.py` and its machine-readable output. The probability that a uniform scalar in E belongs to the certified bank is at least the count in (44) divided by p^4 . This probability describes the scalar-combination event for the two words in the theorem.

Remark A.7 (The comparison radius matters). This example contradicts a universal finite threshold prescription at the displayed radius. It does not, by itself, contradict each numerical row of a parameter table that evaluates the prescription at a different radius. For example, substituting $\eta = 1/16$ into $2^{1/\eta}$ gives 2^{16} ; that evaluates the candidate at a different radius from the one used here.

A.2.2 A sharper bound on selected-witness concurrency

The single-word agreement bound in the reserve construction gives a much stronger concurrency conclusion than the generic bound in Lemma A.3.

Lemma A.8 (Incidence bound from a quantitative agreement deficit). *Let $C \subseteq E^D$ be a linear code on n coordinates, and let $0 \leq a_0 < T \leq n$ be integers. Suppose either the direction f_1 , or some point $f_0 + z_0 f_1$, agrees with every codeword in at most a_0 coordinates. Select codewords p_z at distinct scalars with $\text{agr}_D(f_0 + z f_1, p_z) \geq T$. Every parametrized codeword graph $Q(z) = \sum_{j=0}^e z^j P_j$, where $P_j \in C$ and $e \geq 1$, contains at most*

$$\left\lfloor e \frac{n - a_0}{T - a_0} \right\rfloor$$

of the selected pairs (z, p_z) . In particular, take $e = 1$ for affine codeword graphs.

Proof. Let a count coordinates where $f_0 + zf_1 = Q(z)$ identically as a polynomial in z . If the hypothesis concerns f_1 , comparison of linear coefficients gives $a \leq a_0$. If it concerns a point at z_0 , evaluation there gives the same bound because $Q(z_0) \in C$. At every other coordinate the equality has at most e roots. If the graph contains q selected pairs, counting agreements gives $qT \leq qa + e(n - a)$, so

$$q \leq e \frac{n - a}{T - a} = e \left(1 + \frac{n - T}{T - a} \right) \leq e \frac{n - a_0}{T - a_0}.$$

Taking the integer part proves the claim. $\square$

Corollary A.9 (At most fifteen selected witnesses on one graph). *Every parametrized affine codeword graph contains at most 15 selected pairs from the M31 bank in (44). The same upper bound holds after the two-anchor puncturing and stereographic scaling of Theorem A.10.*

Proof. The root certificate (43) bounds agreement of γ_1 with every outer-band codeword by $a_0 = K + B$. In the displayed M31 instance, $n = 1024B$, $K = 512B$, $T = 545B - 2$, and $B = 65536$. Lemma A.8 therefore gives

$$q \leq \left\lfloor \frac{511B}{32B - 2} \right\rfloor = 15.$$

Indeed, subtracting $15(32B - 2)$ from the numerator gives $31B + 30 > 0$, whereas $16(32B - 2) - 511B = B - 32 > 0$.

Puncturing removes only coordinates outside the certified T agreements, and cannot increase the maximum agreement with γ_1 beyond a_0 . The nonzero stereographic scaling preserves equalities; every transformed codeword lifts to the interior band, which is contained in the outer band. Thus the transformed code of length $n - 2$ satisfies

$$q \leq \left\lfloor \frac{511B - 2}{32B - 2} \right\rfloor = 15$$

by the same lemma. This is an upper bound on selected-witness concurrency, not a claim that a graph attaining fifteen witnesses has been exhibited. $\square$

A.3 Transfer to prime-domain Reed–Solomon codes

Theorem A.10 (Two-anchor puncturing and stereographic scaling). *Let $p \equiv 3 \pmod{4}$, let E contain $\mathbb{F}_{p^2}$, and let $D = \mu_n$, where $n \mid p + 1$ is even and $K = 2m < n$. Suppose words γ_0, γ_1 and a bank of distinct scalars have witnesses in $\mathcal{L}_m^-$, each with at least T agreements outside $\{1, -1\}$. Suppose the pair has no correlated agreement at threshold T with $\mathcal{L}_m^-$. Then there is a set $D' \subseteq \mathbb{F}_p$ of size $n - 2$ and a pair of words over E with the same bank and the same agreement threshold for the ordinary Reed–Solomon code*

$$\text{RS}_{<K-1}(E, D').$$

This code has exact rate $(K - 1)/(n - 2)$, and the transformed pair has no correlated agreement at threshold T .

Proof. Delete the two anchors. Writing $u = x + iy$, map each retained point to

$$t = \frac{y}{1 + x} \in \mathbb{F}_p.$$

Stereographic parametrization is injective away from $(-1, 0)$, which was deleted, and has inverse

$$x = \frac{1 - t^2}{1 + t^2}, \quad y = \frac{2t}{1 + t^2}, \quad u = \frac{1 + it}{1 - it}.$$

The resulting set D' has $n - 2$ distinct finite points. Since $p \equiv 3 \pmod{4}$, the function $w(t) = (1 + t^2)^{m-1}$ is nonzero on $\mathbb{F}_p$. For every $j \in [-(m-1), m-1]$,

$$w(t)u^j = (1 + it)^{m-1+j}(1 - it)^{m-1-j}. \quad (45)$$

These are polynomials of degree at most $2m - 2 = K - 2$, and they form a basis of that polynomial space. Indeed they are the images of the homogeneous monomial basis of binary forms of degree $2m - 2$ under the invertible change of variables $(X, Y) \mapsto (X + iY, X - iY)$. This proof is valid in positive characteristic and does not require any binomial coefficient to be nonzero.

Multiply each retained coordinate of each word and witness by $w(t)$. Equation (45) shows that the witness code becomes exactly $\text{RS}_{<K-1}(E, D')$. All equality tests and scalar labels are preserved by nonzero coordinate scaling, so every witness keeps its T guaranteed agreements. Any correlated agreement of the transformed pair would lift to two elements of $\mathcal{L}_m^-$ agreeing on the same original coordinates, a contradiction. Finally,

$$\frac{\dim \text{RS}_{<K-1}(E, D')}{|D'|} = \frac{K - 1}{n - 2}.$$

This equals $1/2$ when $K = n/2$. □

Corollary A.11 (Transfer of the reserve bank). *In Theorem A.4, suppose $h = M/2 + a$, where $a = 2r + 3$, so that $K = n/2$. The same bank exists for an ordinary Reed–Solomon code with a prime-field domain, extension alphabet E , exact rate $1/2$, length $n' = n - 2$, and gap*

$$\eta' = \frac{T}{n - 2} - \frac{1}{2} = \frac{aB - 1}{n - 2}.$$

Together with Lemma A.3, a bank larger than $(n - 2)2^{1/\eta'}$ contradicts the corresponding finite line-decodability threshold with concurrency number $n - 2$.

The M31 instance above satisfies this inequality as well. It is therefore an example under the literal prime-field-domain and extension-alphabet quantifiers of the finite prescription, rather than only a circle-code analogue. The transfer uses a change of coordinates and nonzero coordinate scaling. It therefore preserves equality sets and scalar labels exactly.

The accompanying exact-arithmetic fixture checks all 189 archived circle witnesses over $\mathbb{F}_{81914}$ after this transformation. They become degree-at-most-14 polynomials on 30 distinct prime-field coordinates, retain 189 distinct labels, and each has at least 20 agreements. The no-correlated-agreement argument uses the algebraic bound $18 < 20$; it is not represented as an exhaustive enumeration of all codewords.

A.4 Circle polynomial spaces and coset domains

For completeness, we verify two algebraic facts used by the constructions: the relation between the circle tensor basis and Laurent bands, and the relation between a full coset and its two half-cosets. Neither fact depends on an ordering of evaluation coordinates.

Proposition A.12 (The circle tensor basis). *For $K = 2m$ a power of two, the circle FFT tensor basis spans*

$$C_K = \{A(x) + yB(x) : \deg A < m, \deg B < m\}.$$

After scalar extension to E it satisfies $\mathcal{L}_m^- \subseteq C_K \subseteq \mathcal{L}_m^+$.

Proof. The tensor factors are $y, x, \pi(x), \pi^2(x), \dots$, with $\pi(x) = 2x^2 - 1$. Products not using y have distinct leading degrees $0, \dots, m-1$ and nonzero leading coefficients, so they form a basis of the polynomials in x of degree less than m . Multiplication by y gives the second summand. Substitution of $x = (u + u^{-1})/2$ and $y = (u - u^{-1})/(2i)$ proves the outer-band containment. The functions $1, x, \dots, x^{m-1}, y, xy, \dots, x^{m-2}y$ span the inner band. The same basis remains a basis after extending coefficients from $\mathbb{F}_p$ to E . $\square$

Proposition A.13 (Canonical domains). *Suppose n is even and $2n \mid p+1$. If $g \in U$ has order $2n$, the canonical coset $g\mu_n$ is the disjoint union*

$$\{g^{1+4j} : 0 \leq j < n/2\} \cup \{g^{-(1+4j)} : 0 \leq j < n/2\}.$$

Thus describing a canonical domain as two conjugate half-cosets changes only the ordering of its points.

Proof. In additive point-index notation modulo $2n$, the full odd coset is $\{1 + 2j : 0 \leq j < n\}$. Its half-coset representation is $\{1 + 4j : 0 \leq j < n/2\}$ and its negatives. These two half-cosets are disjoint and their union is exactly the full odd coset. Passing to the split multiplicative coordinate gives $g\mu_n$. $\square$

Consequently, a construction on a full multiplicative coset also applies when its coordinates are listed using the two half-cosets. Rotation preserves both Laurent bands.

Remark A.14 (Sizes and target spaces). The M31 circle has order $p+1 = 2^{31}$; a canonical coset uses a subgroup of twice its own size and has maximum length 2^{30} . Here n denotes evaluation length. For example, the rate-1/2 code at length 2^{26} has dimension 2^{25} . The full outer-band circle code has dimension $K+1$, rather than K ; the proof excluding correlated agreement still covers it, but its rate is $1/2 + 1/n$ and any comparison naming that code must use that rate.

A.5 What the finite counterexamples establish

The list construction and the affine-family construction address two different finite bounds. At rate $1/2$, the first gives a list larger than $2^{1/\eta}$. The second gives more than $n2^{1/\eta}$ exceptional scalars, with no correlated agreement at the prescribed threshold. The Reed–Solomon transfer shows that the latter phenomenon also occurs with a prime-field evaluation domain and an extension-field alphabet.

One must distinguish such finite statements from a proposed asymptotic threshold of the form

$$a = \ell(\theta)n + o(n), \quad \ell(\theta) \leq c_1 2^{c_2 H_2(\rho)/\eta}, \quad \theta = 1 - \rho - \eta.$$

The displayed numerical examples contradict the remainder-free choice $c_1 = c_2 = 1$ at their stated fields, lengths, and radii. A single finite example does not determine the unspecified quantifiers in $o(n)$ and does not refute every choice of constants c_1, c_2 .

For an asymptotic claim, the field, domain length, and gap must be allowed to vary in a stated way. In particular, an ordinary Reed–Solomon code with a prime-field evaluation domain has length at most p . Holding p fixed while allowing an arbitrary prime-dependent remainder leaves no unbounded sequence of such lengths. The separate fixed-gap lifting construction in this manuscript gives an increasing-length family for the quantitative coefficient: its discrepancy is linear and cannot be absorbed by a sublinear remainder for that family. That selected-graph construction has correlated agreement and is distinct from the globally non-correlated affine families constructed here. Neither argument refutes a line bound whose coefficient validly bounds the actual list size.

A.6 The common-core form of the circle pencil

The anchored construction in Theorem A.4 can be written in the fixed-core form familiar from the Proximity Prize's orbit pencil [29]. This is a change of presentation, with the same witnesses and agreement counts.

Use the theorem's notation, and fix one support A_0 . For each support $A = \{1\} \cup U$, write

$$F_A(Y) = (Y - 1)G_U(Y), \quad D_U(Y) = \frac{G_{U_0}(Y) - G_U(Y)}{Y^{r+1}}.$$

The common leading and trailing coefficients imply $\deg D_U \leq h - 2r - 3 = K/B$. Since $R_A - R_{A_0} = (Y - 1)D_U$, the old scalar labels and the new labels $t_U = D_U(\tau)$ satisfy

$$z_A - z_{A_0} = (1 - \tau)t_U.$$

Set $Q_U = (D_U - t_U)/(Y - \tau)$ and $S_B(u) = (u^B - 1)/(u^2 - 1)$. Then the new pencil and witnesses are

$$g_0 = \frac{u^{-s}S_B(u)G_{U_0}(u^B)}{u^{B(r+1)}(u^B - \tau)}, \quad g_1 = -\frac{u^{-s}S_B(u)}{u^B - \tau}, \quad W_U = u^{-s}S_B(u)Q_U(u^B).$$

Their difference is

$$g_0 + t_U g_1 - W_U = \frac{u^{-s}S_B(u)G_U(u^B)}{u^{B(r+1)}(u^B - \tau)}.$$

The zeros consist of the $h - 1$ varying fibers and the fixed core of $B - 2$ points with $u^B = 1$ and $u \neq \pm 1$. The agreement count is $(h - 1)B + B - 2 = hB - 2$. Also $\deg(S_B Q_U(u^B)) \leq K - 2$, so the centered witnesses lie in the required interior Laurent band.

Direct subtraction in the original formulas gives

$$g_0 = \gamma_0 + z_{A_0}\gamma_1 - P_{A_0}, \quad g_1 = (1 - \tau)\gamma_1, \quad W_U = P_A - P_{A_0}.$$

Thus an invertible affine change of scalar and a codeword translation identify the two presentations. The number of labels and the failure of correlated agreement are unchanged. In particular, introducing this common-core notation does not itself improve the existing numerical bound.

B A two-fiber pencil in raw quotient form

This section addresses one algebraic compatibility condition with the conjugate-line quotient, using the unweighted formula in² the source. It does not construct an accepted Stwo transcript. In particular, the claimed openings below are zero and generally false.

Write $p = 2^{31} - 1$, $H = \mathbb{F}_p[i]$, $E = H[w]$, where $i^2 = -1$ and $w^2 = \kappa = 2 + i$. Conjugation over H sends w to $-w$. Let $n = 2M \mid p + 1$, $K = 2m \mid p + 1$, with m even and $m + 3 \leq M - 1$. The circle code C_K contains every Laurent polynomial with exponents between $1 - m$ and $m - 1$, and is contained in the band $[-m, m]$.

Proposition B.1 (Fixed tables realizing a two-fiber pencil). *There are two $\mathbb{F}_p$ -valued tables on μ_n , fixed independently of the sample, and an out-of-domain circle sample over E , such that their normalized raw quotients r_0, r_1 with claimed values zero have at least*

$$J = \left\lceil \frac{NQ}{Q + m(N - 1)} \right\rceil, \quad N = \left\lceil \frac{\binom{M-1}{m+2}}{M} \right\rceil, \quad Q = \frac{p^2 - 1}{2},$$

²Source checked at Stwo revision 826591c6c371376810ca8213b5812d5daf6d5092, `core/constraints.rs` and `core/pcs/quotients.rs`.

distinct parameters α for which $r_0 + \alpha r_1$ agrees with C_K on at least $K + 4$ coordinates. The direction r_1 agrees with any codeword on at most $K + 2$ coordinates, so the pencil has no common agreement on $K + 4$ coordinates.

Proof. Choose a class of $(m + 2)$ -subsets $U \subset \mu_M \setminus \{1\}$ with common constant coefficient c in $G_U(Y) = \prod_{a \in U} (Y - a)$. There are at least N members. Fix U_0 , and set

$$D_U = (G_{U_0} - G_U)/Y, \quad \deg D_U \leq m.$$

These polynomials are distinct. Average collisions of their evaluations over the Q nonsquares $\tau \in H^*$. Every difference has at most m roots, so Cauchy–Schwarz gives a nonsquare τ with at least J distinct labels $t_U = D_U(\tau)$. Every element of μ_M is a square in H , so these poles avoid the domain.

Put $Y = u^2$ and

$$g_0 = \frac{u^{1-m} G_{U_0}(u^2)}{u^2(u^2 - \tau)}, \quad g_1 = -\frac{u^{1-m}}{u^2 - \tau}, \quad W_U = u^{1-m} \frac{D_U(u^2) - t_U}{u^2 - \tau}.$$

The last expression lies in the strict Laurent band. The difference $g_0 + t_U g_1 - W_U$ is $u^{1-m} G_U(u^2)/[u^2(u^2 - \tau)]$, giving $2(m + 2)$ agreements. For any codeword C , multiplying $g_1 - C$ by $u^m(u^2 - \tau)$ gives a nonzero polynomial of degree at most $K + 2$. Nonzeroness follows by evaluating at a root of $u^2 - \tau$, where the remaining term is $-u$. This proves the direction bound.

Choose $a_0 \neq 0$ with $a_0^p/a_0 = c$: take $a_0 = 1 + c^{-1}$ if $c \neq -1$, and $a_0 = i$ otherwise. The two fixed tables are

$$f_0(u) = a_0 u^{-(m+2)} G_{U_0}(u^2), \quad f_1(u) = u^m + u^{-m}.$$

Conjugating coefficients and using $u^p = u^{-1}$ proves that both are $\mathbb{F}_p$ -valued. Neither depends on τ .

Because κ and τ are nonsquares in H , choose $b \in H^*$ with $\kappa b^2 = \tau$. Set $a = wb$ and

$$z_x = (a + a^{-1})/2, \quad z_y = (a - a^{-1})/(2i).$$

These coordinates satisfy the circle equation and are anti-invariant under conjugation over H . Writing $z_x = wX_1, z_y = wY_1$, we have $Y_1 \neq 0$ and

$$U_z = (x - (X_1/Y_1)y)/2 = L_0 u^{-1}(u^2 - \tau), \quad L_0 = \frac{1}{2(1 - \tau)}.$$

For claimed opening zero, the raw conjugate-line formula is $q = (\bar{z}_y - z_y)f/[(-x)Y_1 - (-y)X_1]$, hence $q/w = f/U_z$. Thus, with $r_j = q_j/w$, direct multiplication gives

$$r_0 = \lambda_0 g_0, \quad r_1 = \lambda_1 g_1 + P_\tau,$$

where

$$\lambda_0 = a_0/L_0, \quad \lambda_1 = -(\tau^m + 1)/L_0, \quad P_\tau = \frac{u^{1-m}}{L_0} \sum_{j=0}^{m-1} \tau^{m-1-j} u^{2j} \in C_K.$$

Here $\lambda_1 \neq 0$: otherwise $\tau^{2m} = 1$, which would make τ a square in H , since $2m \mid p + 1$. For $\alpha_U = \lambda_0 t_U/\lambda_1$, the witness is $\lambda_0 W_U + \alpha_U P_\tau$. The scalar changes and codeword shift preserve distinct labels, agreement counts, and the direction bound. Multiplication by w also preserves membership in the E -linear code. $\square$

Finite instance and remaining protocol gap. Taking $(n, K, M, m) = (128, 32, 64, 16)$ gives agreement 36/128 and $N = 40,448,653,414,754$, yielding $J = 40,437,303,920,543$. The expression above is an exact lower bound; no enumeration of this bank is claimed. A separate deterministic fixture over $p = 127$, with the same extension polynomial, checks the raw quotient identities and agreement counts for multiple nonsquare poles, holding the two base-field tables fixed throughout. An additional fixture over the actual M31 field checks three supports at one nonsquare pole: all 1,536 identities pass, with 36 agreements for each support. This result concerns a specially chosen sample family. It does not establish its probability under Stwo’s sampler, map the actual challenge power schedule or its precise evaluation cosets, or address interpolation and extension of the raw quotient subdomain. Nor does it show that the tables and false openings survive the remaining proof checks. The bounds for larger fibers do not automatically apply to this two-fiber case.

There is a concrete reason to keep the sampler issue separate. Under an idealized uniform sample from the E -rational circle and an independent uniform batching parameter in E , this particular family is rare. The map $a \mapsto ((a + a^{-1})/2, (a - a^{-1})/(2i))$ bijects E^* with that circle. Samples used here have $a \in wH^*$, a fraction $1/(p^2 + 1)$ of the circle. All constructed batching labels lie in H , a fraction $1/p^2$ of E . Thus the event consisting of this sample family and these labels has probability at most $1/[p^2(p^2 + 1)]$, approximately 2^{-124} , even before limiting the label bank. This is an upper bound for the constructed event in the stated idealized model, not a soundness bound for Stwo. In particular, raw quotient compatibility by itself does not establish a practical attack or a reduction below a 96-bit security target.

C From large real lists to raw quotient pencils at typical samples

Sampling outside the evaluation domain to constrain candidate codewords is the established DEEP technique [9]. Here we study large explicit families under the conjugate-line quotient, fixed claims, and prescribed scalar schedules.

The special sample restriction in Appendix B can be removed. We first prove a bound with sample-dependent claimed openings; Theorem C.2 then fixes the claims before the sample and improves the joint probability bound. Sample-dependent claims are consistent with the inspected source order: column commitments precede the sample, while claimed values are mixed before the batching challenge. Both results concern raw quotient agreement; neither proves full proof acceptance.

Theorem C.1 (A list-to-quotient transfer). *Let $p = 2^{31} - 1$, $H = \mathbb{F}_{p^2}$ and $E = \mathbb{F}_{p^4}$, with the tower and normalization of Appendix B. Let $\mathcal{D}$ be a set of n base-field circle points and $K = 2m$. Suppose a fixed base-field word f has N distinct real Laurent polynomials F_j , supported in $[-m, m]$, each agreeing with f on at least $T > K + 2$ points of $\mathcal{D}$. Set*

$$S = (p^2 - 1)^2, \quad D_* = \left\lceil \frac{NS}{S + 2K(N - 1)} \right\rceil, \quad J = \left\lceil \frac{D_*}{p^2} \right\rceil.$$

For at least half of the S admissible circle samples described below, there is a claimed value v such that the two fixed tables $(f, 0)$, with claimed values $(v, 1)$, have normalized raw quotients r_0, r_1 for which at least J distinct $\alpha \in H$ satisfy

$$r_0 + \alpha r_1 \text{ agrees with } C_K(H) \text{ on at least } T \text{ coordinates.}$$

The direction r_1 has at most $K + 2$ agreements with any codeword, so these pencils have no common agreement on T coordinates. The claims are chosen after the sample and before α .

Proof. Parameterize the E -circle by $z_x = (a + a^{-1})/2$, $z_y = (a - a^{-1})/(2i)$, $a \in E^*$. Call a admissible when $a \notin H$ and $z_y \notin H$. For $a \notin H$, the second condition fails exactly when $a\bar{a} = -1$. The sets H^* and this norm fiber have sizes $p^2 - 1$ and $p^2 + 1$, and intersect at $\pm i$. Their complement therefore has size S .

For $j \neq k$, multiplying $F_j - F_k$ by u^m gives a nonzero polynomial of degree at most K . Thus at most K admissible values of a cause $F_j(z) = F_k(z)$. The expected number of unordered collisions is at most $\binom{N}{2}K/S$. By Markov's inequality, at least half the admissible samples have at most $N(N-1)K/S$ collisions. Cauchy-Schwarz then gives at least D_* distinct values among $F_j(z)$. One of the p^2 additive H -cosets in E contains at least J of them. Write this coset as $v + H$, and retain one index for each distinct value. The labels $\alpha_j = F_j(z) - v$ are distinct and belong to H .

Use the source functions U_z and ℓ_b defined in Appendix B and the preceding note. Explicitly, for $z_x = X_0 + wX_1$, $z_y = Y_0 + wY_1$, and $b = B_0 + wB_1$,

$$U_z = (x - X_0)/2 - (X_1/Y_1)(y - Y_0)/2, \quad \ell_b = B_0 + (B_1/Y_1)(y - Y_0).$$

The polynomial uU_z has degree two and distinct nonzero roots $a, \bar{a}$. The numerator $F_j - \ell_{F_j(z)}$ vanishes at both roots. Indeed $\ell_b(z) = b$, and conjugating over H supplies the other equality. Division therefore yields

$$P_j = \frac{F_j - \ell_{F_j(z)}}{U_z} \in \text{span}_H\{u^{1-m}, \dots, u^{m-1}\} \subset C_K(H).$$

Since $b \mapsto \ell_b$ is H -linear and $\ell_1 = 1$, we obtain

$$r_0 + \alpha_j r_1 - P_j = \frac{f - \ell_v - \alpha_j}{U_z} - P_j = \frac{f - F_j}{U_z}.$$

No denominator vanishes on the base-field circle, so the claimed agreements follow. Finally $r_1 = -1/U_z$. For any codeword P in the outer Laurent band, clearing $r_1 - P$ by $u^m(uU_z)$ gives a polynomial of degree at most $K + 2$. It is nonzero at $u = a$, where its value is $-a^{m+1}$. This proves the direction bound. $\square$

An explicit family and finite count. The construction in Appendix B supplies such a family: put $f = R_{U_0}$ and

$$R_U = a_0 u^{-(m+2)} G_U(u^2), \quad F_U = f - R_U.$$

The common leading and constant coefficients give support $F_U \subseteq [-m, m]$, while R_U has $K + 4$ roots on the domain. For $(n, K) = (512, 128)$, the exact counts give

$$N = \left\lceil \binom{255}{66} / 256 \right\rceil, \quad J \geq 18,014,398,492,704,768.$$

On every good sample the certified raw proximity event therefore has probability at least $J/p^4 \approx 2^{-70}$ under a uniform linear batching challenge. The good-sample fraction costs approximately one further bit. This is not a protocol security estimate.

Domains and the source sampler. The transfer theorem applies to arbitrary base-field circle domains. The displayed family transports to a coset $u_0\mu_n$ by replacing u with u/u_0 . Reality and Laurent support are preserved, and its quotient witnesses lie in the rotation-invariant interior band. Stwo's canonical domain is precisely such a coset, with u_0 of order $2n$. The source samples using stereographic coordinates $((1 - t^2)/(1 + t^2), 2t/(1 + t^2))$. For uniform $t \in E$, each admissible circle point has exactly one preimage; the two singular parameters and the omitted point are outside this admissible set. Thus, under independent uniform secure-field draws, the joint sample-and-batching proximity event has probability at least $SJ/(2p^8)$.

Scope and checks. Exact generic-sample fixtures pass over $\mathbb{F}_{127}$ and M31; a separate $\mathbb{F}_{23}$ fixture checks the fixed second claim 1 and distinct labels on additive H -cosets. Section C.6 gives a compact description of the large support family and constructs its fixed table efficiently. An efficient algorithm for finding a favorable claim and its witnesses is still not established. Additional columns, periodicity terms, claim constraints from the AIR, commitment consistency, FRI queries and grinding remain to be analyzed. In particular, neither the proximity probability nor the existence of a nearby codeword establishes an accepted false proof.

C.1 A fixed claim and a stronger joint probability bound

The preceding argument selects a favorable claimed value separately for each sample and discards up to half the admissible samples. Both losses can be avoided when the objective is a joint sample-and-batching bound. The following averaging argument fixes a single claim before either challenge. Its guarantee concerns the total number of successful pairs; it does not promise the same number of labels at every sample.

Theorem C.2 (Fixed-claim list-to-quotient transfer). *Under the hypotheses of the list-to-quotient transfer above, let*

$$A_* = \left\lceil \frac{NS^2}{p^2(S + K(N-1))} \right\rceil, \quad S = (p^2 - 1)^2.$$

There is one value $v_ \in wH$, independent of the sample and batching challenge, such that the fixed tables $(f, 0)$ with fixed claims $(v_*, 1)$ have at least A_* distinct pairs (a, α) with a admissible and $\alpha \in H$, for which*

$$r_{0,a} + \alpha r_{1,a} \text{ agrees with } C_K(H) \text{ in at least } T \text{ coordinates.}$$

For every admissible sample, the direction has at most $K + 2 < T$ agreements with any codeword. Under independent uniform draws $t, \alpha \in E$, with t mapped to the circle by the stereographic sampler, the probability of the certified raw proximity event is at least

$$\frac{A_*}{p^8} \geq \frac{NS^2}{p^{10}(S + K(N-1))}. \quad (46)$$

The two singular sampler parameters are counted as failures in this unconditioned bound.

Proof. Let $\mathcal{A}$ be the set of admissible parameters. For $a \in \mathcal{A}$ and $b \in E$, let

$$m_{a,b} = \#\{j : F_j(z(a)) = b\}, \quad \mathcal{I} = \{(a, b) : m_{a,b} > 0\}.$$

Then $\sum_{a,b} m_{a,b} = NS$. Each unordered pair of distinct Laurent polynomials collides at at most K parameters, by the same degree- K root bound used above. Consequently

$$\sum_{a,b} m_{a,b}^2 = NS + 2 \sum_a \sum_b \binom{m_{a,b}}{2} \leq NS + KN(N-1).$$

Cauchy–Schwarz, applied once to all incidences, gives

$$|\mathcal{I}| \geq \frac{(NS)^2}{NS + KN(N-1)} = \frac{NS^2}{S + K(N-1)}. \quad (47)$$

Fix the p^2 representatives wH for the additive cosets of H in E . Partition $\mathcal{I}$ according to the coset containing b . One fixed coset $v_* + H$ contains at least A_* incidences. For each of these

incidences choose one witness F_j and put $\alpha = b - v_* \in H$. This is an injective map from the retained incidences to pairs (a, α) , even when several polynomials have the same value at one sample. Both v_* and the tables are fixed throughout this count.

For the sample a , write $U_a = U_{z(a)}$ and use the interpolation operator $\ell_{a,b}$ from the preceding proof. The quotient witness is

$$P_{a,j} = \frac{F_j - \ell_{a,b}}{U_a} \in C_K(H), \quad r_{0,a} + \alpha r_{1,a} - P_{a,j} = \frac{f - F_j}{U_a}.$$

Here $r_{0,a} = (f - \ell_{a,v_*})/U_a$ and $r_{1,a} = -1/U_a$. Thus all retained incidences give the required agreements. The degree- $K + 2$ direction argument applies to every admissible sample, regardless of the chosen coset. Finally, every admissible a has exactly one nonsingular stereographic preimage $t \in E$. The A_* distinct pairs therefore occupy A_* of the p^8 equally likely pairs (t, α) , proving the probability claim. $\square$

The same finite instance with fixed claims. For the preceding $(n, K, T) = (512, 128, 132)$ instance,

$$N = \left\lceil \binom{255}{66} / 256 \right\rceil, \quad S = (p^2 - 1)^2, \quad p = 2^{31} - 1.$$

Exact integer arithmetic gives

$$A_* = 766247768292072891947921916124690454684786490150411277, \quad 2^{69} A_* > p^8.$$

The certified joint raw proximity probability is therefore strictly greater than 2^{-69} . The earlier half-sample bound is $SJ/(2p^8)$, approximately 2^{-71} at these same parameters; the improvement is approximately a factor of four. The earlier 2^{-70} figure was conditional on a good sample, so it should not be compared directly with this joint probability.

Quantifiers and limits. This strengthens the timing statement to *there exists one fixed claim, followed by random sample and batching draws*. It does not identify an efficiently computable v_* . Section C.6 constructs the table and describes the support bank implicitly; efficient witness recovery remains open. Individual samples may have fewer labels than the previous per-good-sample bound, or no certified label. The false claim 1 for the zero table remains part of the construction. The result concerns raw quotient agreement under the stated independent uniform challenge model. It does not establish a Fiat–Shamir attack, an accepted false proof, or a 69-bit security level for Stwo.

Reproducible checks. The companion script `checks/verify_fixed_claim.py` checks the displayed large-field counts using exact integers and rational numbers. An independent toy instance over $\mathbb{F}_{74}$ exhausts all 2,304 admissible samples. It checks the sampler bijection, counts distinct evaluation incidences and additive cosets, fixes one claim globally, and verifies every retained quotient identity and agreement count. A separate small family with actual evaluation collisions tests the multiplicity term in (47). These finite checks support the formulas; the proof above establishes the general statement.

C.2 Subtracting roots forced by shared agreements

The degree bound in Theorem C.2 counts all roots of a difference as possible sample collisions. Roots on the base-field evaluation domain cannot be admissible samples. Agreements with the same received word force many such roots and give a sharper bound.

Proposition C.3 (Agreement-aware collision bound). *Suppose that, for each of the N polynomials in the transfer theorem, we select T agreement coordinates from a common pool of n_0 base-field circle points. Write*

$$NT = n_0q + r, \quad 0 \leq r < n_0, \quad R = (n_0 - r) \binom{q}{2} + r \binom{q+1}{2}, \quad C_* = K \binom{N}{2} - R.$$

Then the total number of unordered polynomial-pair collisions over admissible samples is at most C_ . In particular, the assumed family implies $C_* \geq 0$. The total number I of distinct sample/value incidences satisfies*

$$I \geq \frac{N^2 S^2}{NS + 2C_*}. \quad (48)$$

Consequently Theorem C.2 holds with the stronger pair count

$$A_{\text{agr}} = \left\lceil \frac{N^2 S^2}{p^2(NS + 2C_*)} \right\rceil.$$

Proof. Let d_x count selected agreement sets containing coordinate x . Then $\sum_x d_x = NT$, and double counting gives

$$\sum_{j < k} |A_j \cap A_k| = \sum_x \binom{d_x}{2} \geq R.$$

The last inequality follows by balancing the integer counts: moving one unit from a count at least two larger than another strictly decreases the sum of binomial coefficients. At a coordinate of $A_j \cap A_k$, both polynomials equal the received word. These are distinct roots of $u^m(F_j - F_k)$, in addition to any admissible collision roots, and are disjoint from the admissible sample set. Hence pair (j, k) has at most $K - |A_j \cap A_k|$ admissible collisions. Summing proves the bound. The same incidence Cauchy–Schwarz argument as before yields (48); partition into the p^2 additive H -cosets gives the final statement. $\square$

In the displayed two-fiber family, all 132 agreement coordinates avoid the two points above $1 \in \mu_{256}$. Thus $n_0 = 510$ is valid, with $N = \lceil \binom{255}{66} / 256 \rceil$. The exact bound becomes

$$A_{\text{agr}} = 1045232662865816582797566371469227027334159065108360767.$$

The negative base-two logarithm of A_{agr}/p^8 is approximately 68.5521. This is a further factor of approximately 1.3641 over Theorem C.2 at the same parameters.

C.3 Fixed claims under nonuniform batching schedules

The coefficient multiplying the second quotient need not itself be a uniform field element. A second averaging argument handles a prescribed function of the original uniform challenge. Here the claimed value is allowed to range over all of E , rather than just the representatives wH .

Theorem C.4 (Transfer for a prescribed scalar schedule). *Use the same list, admissible samples, and raw quotients as above. Let $g_0, g_1 : E \rightarrow E$ be fixed functions, and set*

$$\mathcal{B} = \{\alpha \in E : g_0(\alpha) \neq 0, g_1(\alpha)/g_0(\alpha) \in H\}, \quad G = |\mathcal{B}|.$$

Let I_0 be any real lower bound for the number of distinct sample/value incidences, for example the right-hand side of (48). There is a single claim $v \in E$, fixed before both challenges, for which at least

$$A_{\text{sched}} = \left\lceil \frac{GI_0}{p^4} \right\rceil$$

distinct admissible pairs (a, α) satisfy

$$g_0(\alpha)r_{0,a} + g_1(\alpha)r_{1,a} \text{ agrees with } C_K(E) \text{ in at least } T \text{ coordinates,}$$

with fixed tables $(f, 0)$ and fixed claims $(v, 1)$. Under the same independent uniform challenge model, the joint probability is at least A_{sched}/p^8 .

Proof. For every distinct incidence (a, b) and every $\alpha \in \mathcal{B}$, assign the triple to

$$v = b - \phi(\alpha), \quad \phi(\alpha) = g_1(\alpha)/g_0(\alpha) \in H.$$

There are GI triples and p^4 possible values of v , so one fixed value receives at least A_{sched} triples. For this fixed v , the map to (a, α) is injective: these two coordinates uniquely determine $b = v + \phi(\alpha)$. Choose a polynomial witnessing each retained incidence. The raw quotient identity gives

$$g_0(\alpha)r_{0,a} + g_1(\alpha)r_{1,a} - g_0(\alpha)P_{a,j} = g_0(\alpha)\frac{f - F_j}{U_a}.$$

The witness $g_0(\alpha)P_{a,j}$ lies in the scalar-extended code $C_K(E)$. Nonzero multiplication preserves agreement coordinates. Counting the distinct sampler preimages proves the probability bound. No injectivity or uniform-output assumption about ϕ is used. $\square$

Additional terms that are codewords. The conclusion is preserved after adding any $B_{a,\alpha} \in C_K(E)$ to the batched word: add the same codeword to its witness. Thus other columns need not have zero quotients if their complete contribution is proved to lie in this code. A common preceding factor α^e , $e > 0$, removes at most the challenge $\alpha = 0$ from the certified set. These statements require code membership of the additional contribution; they do not discard arbitrary additional quotient or AIR terms.

Polynomial schedules over the subfield. If $g_0, g_1 \in H[X]$ and g_0 is a nonzero polynomial of degree $d_0 < p^2$, then $G \geq p^2 - d_0$: every $\alpha \in H$ outside the roots of g_0 qualifies. Thus even a nonlinear relative coefficient incurs no additional loss from the degree of g_1 . More precisely, for $g_0 = 1$, $g_1 = X^d$, $d \geq 1$, cyclicity of E^* gives

$$G = 1 + (p^2 - 1) \gcd(d, p^2 + 1). \quad (49)$$

Indeed a nonzero α qualifies exactly when $\alpha^{d(p^2-1)} = 1$. The assertion is about the number of original challenges, not a claim that α^d is uniformly distributed.

A finite dependence on the gap between powers. For $p = 2^{31} - 1$, each of $d = 1, 2, 5, 10$ divides $p^2 + 1$. The same list with the agreement-aware incidence bound gives:

Power gap d	One-sample columns	Qualifying challenges G	$-\log_2(A_{\text{sched}}/p^8)$
1	2	p^2	68.5521
2	3	$2p^2 - 1$	67.5521
5	6	$5p^2 - 4$	66.2301
10	11	$10p^2 - 9$	65.2301

The final column is approximate; exact integer certificates prove, respectively, probabilities greater than 2^{-69} , 2^{-68} , 2^{-67} , and 2^{-66} . All rows use the same $(n, K, T) = (512, 128, 132)$; the favorable fixed claim may differ by row. At the quotient layer the schedule $1, \alpha^d$ is realized by putting $d - 1$ zero columns with truthful zero claims between the two designated columns, with one sample per column at the same point. Every inserted numerator is identically zero, while its sample consumes one challenge power. This establishes the displayed quotient expression and column count; it does not show that any particular AIR allows this arrangement.

Corollary C.5 (The periodicity coefficient $\alpha + \alpha^3$). *For $g_0 = 1$ and $g_1 = X + X^3$, one has $G = 2p^2 - 1$. Using the agreement-aware bound at $(n, K, T) = (512, 128, 132)$ gives*

$$A_{\text{cubic}} = \left\lceil \frac{(2p^2 - 1)N^2S^2}{p^4(NS + 2C_*)} \right\rceil, \quad 2^{68}A_{\text{cubic}} > p^8.$$

Its joint raw-proximity lower bound is therefore strictly greater than 2^{-68} ; its negative base-two logarithm is approximately 67.5521.

Proof. Put $q = p^2$ and write $\alpha = x + wy$, with $x, y \in H$ and $w^2 = \kappa$. For $y = 0$ there are q qualifying challenges. For $y \neq 0$, equality of $\alpha + \alpha^3$ with its conjugate is equivalent to

$$3x^2 + \kappa y^2 = -1.$$

The element $3 \in \mathbb{F}_p^*$ is a square in H (here $p \neq 3$). Choose $s \in H$ with $s^2 = 3$. Since $i^2 = -1$,

$$\text{Norm}_{E/H}(sx + wiy) = 3x^2 + \kappa y^2.$$

The nonzero norm fiber at -1 has $q + 1$ elements. Exactly two have $y = 0$, since $-1/3$ is a square in H . Thus $q - 1$ further challenges qualify, giving $G = 2q - 1$. The numerical inequality is certified with exact integers by the companion script. $\square$

The source schedule that gives this coefficient. At the pinned proving revision [31], `build_samples_with_randomness_and_periodicity` assigns powers across trees, columns, and samples in that order. Two consecutive one-sample columns at the same point receive $1, \alpha$ if they are the first nonempty columns. No periodicity sample is added to either.

There is also an exact three-power realization of the cubic coefficient. Place the (f, v) column first with its sole sample at z . Place the zero column next, with claims 0 at a previous point z_- and 1 at z . Give the zero column the maximal lifted log size. Its periodicity point then coincides with z , and the source assigns the three powers $\alpha, \alpha^2, \alpha^3$ to its periodicity, previous, and current samples. Its previous-sample numerator is zero. The normalized sum is therefore exactly

$$r_{0,z} + (\alpha + \alpha^3)r_{1,z}.$$

If the previous point makes the source's conjugate-line formula singular, that sample must additionally be excluded. For a fixed base-field rotation $z_- = hz$, at most $p^2 - 1$ initially admissible samples are lost: the excluded parameters satisfy $a\bar{a} = -h^{-2}$, a norm fiber of size $p^2 + 1$, of which two parameters lie in H . Replacing S by the conservative $S' = S - (p^2 - 1)$ in the incidence bound gives a version covering both source denominators. The strict 2^{-68} comparison still holds.

This realizes the scalar schedule at the quotient layer for the stated column arrangement. It does not establish that an AIR admits these claims, that a complete transcript accepts, or that the favorable claim and corresponding witnesses can be found efficiently.

C.4 The radius of the raw-quotient example

The (512, 128, 132) example has rate $\rho = 1/4$, gap $\eta = 1/128$, and radius $\theta = 95/128$. It is *above* the characteristic-based Elias radius: $H_p(95/128) \approx 0.7687496752 > 3/4$. This comparison follows without floating-point assumptions from

$$p^{96}95^{95}33^{33} < (p-1)^{95}128^{128}.$$

It follows that this raw-quotient example is not itself a counterexample within the below-Elias range of the whitepaper’s list prescription. The below-Elias counterexamples elsewhere in the paper remain separate constructions. Neither the 2^{-69} bound nor its stronger variants should be interpreted as a protocol security level.

C.5 What a nearby codeword supplies at the FRI query stage

The source permits one further conditional statement. Its first FRI layer combines the quotient values computed at sampled positions with complementary values supplied by the FRI proof. Agreement is required at the sampled positions; the surrounding folding block need not consist entirely of agreement coordinates.

Lemma C.6 (Conditional query survival). *Let $R \in E^D$ be the quotient word computed from fixed committed tables and claimed openings, and suppose $P \in C_K(E)$ agrees with R on $A \subseteq D$. Assume that an honest FRI proof is formed for P on this same domain and degree bound, and that the original table decommitments are valid. If every sampled quotient-query position lies in A , the first-layer value substitution leaves the honest FRI verification unchanged. This holds for the source’s supported folding steps and leaf packing.*

Under q_{qry} independent uniform draws from D , with replacement before sorting and deduplication, the probability of this sufficient event is exactly

$$\left(\frac{|A|}{|D|}\right)^{q_{\text{qry}}}.$$

If the preceding raw-proximity event has probability at least π and supplies $|A| \geq T$ on a length- n domain, then this particular two-stage sufficient event has probability at least $\pi(T/n)^{q_{\text{qry}}}$, under independent future query draws.

Proof. The source function `compute_decommitment_positions_and_rebuild_evals` partitions the sorted query positions into folding blocks. At an actual query position it uses the provided quotient value; at any other position in that block it takes a value from the FRI witness. Form the latter values and all authentication paths using the honest proof for P . When every queried position is in A , the supplied quotient value is also P at that position. Thus every reconstructed block, including any packed leaf, equals its honest counterpart. Its Merkle check, subsequent fold, and all later FRI checks are unchanged.

The raw query draws all land in A with the displayed probability. Sorting and removing duplicates does not change that event. The final statement follows by conditioning on the earlier sample and batching challenges, including the selected witness, and then averaging. $\square$

Numerical consequence and scope. For the (512, 128, 132) example, $T/n = 33/128$. At 70 raw query draws the additional sufficient-event factor is

$$(33/128)^{70} \approx 2^{-136.8924}.$$

For illustration, multiplying this by the certified raw bound for the eleven-column, power-gap-ten schedule gives a two-stage lower bound of approximately $2^{-202.1225}$. This uses 70 as a stated query count; it does not identify this rate-1/4 example with a production parameter set. These are lower bounds for a particular sufficient event, not upper bounds on all possible strategies or proof-system soundness guarantees.

The lemma is conditional on constructing the codeword and honest FRI proof, providing valid table decommitments, and meeting the prescribed domain and degree metadata. It does not establish the AIR conditions, efficient recovery of witnesses or the favorable fixed claim, or a successful native verifier execution. Proof-of-work search and Fiat–Shamir retry costs are not included in this probability. In particular, an inverse one-trial probability is not automatically an attack work factor: reusing a successful prefix and searching later challenges requires its own analysis.

The companion `checks/verify_query_survival.py` enumerates all $8^3 = 512$ ordered query triples for an eight-coordinate word with five agreements. For each of folding steps 1, 2, 3, precisely 125 triples give unchanged reconstructed blocks, including cases where unqueried positions in the block are disagreement coordinates. This checks the reconstruction argument against an explicit vector commitment model; it does not replace a native Merkle or FRI test.

C.6 An explicit table with an exponentially large implicit list

The support family need not be enumerated to construct its received word. On μ_M , equality of the constant coefficients of the root polynomials is a subset-sum condition modulo M . This yields an efficient exact construction of the fixed table used above.

Proposition C.7 (Constructing the table without enumerating its list). *Let λ generate μ_M and consider the h -element subsets of $\{\lambda, \lambda^2, \dots, \lambda^{M-1}\}$. The sizes of all classes with equal constant coefficient, a largest class, and one support in that class can be computed with $O(M^2h)$ additions of integers of $O(M)$ bits, storing $O(Mh)$ such integers and support masks. For the displayed $M = 256$, $h = 66$ instance, the largest count is*

$$N_{\text{exact}} = 471959451942648482921857753174390587612228702156784439968875.$$

The residue 1 attains this count. It gives an explicitly computed 512-entry $M31$ -valued table on a canonical coset and an implicit family of N_{exact} distinct real Laurent polynomials with support in $[-64, 64]$, each agreeing with that table in exactly 132 coordinates.

Proof. Let $D_{j,k,r}$ be the number of k -subsets of $\{1, \dots, j\}$ whose sum is r modulo M . The initial condition is $D_{0,0,0} = 1$, with all other counts zero, and

$$D_{j,k,r} = D_{j-1,k,r} + D_{j-1,k-1,r-j \bmod M}.$$

Updating k in decreasing order permits an in-place computation using $O(Mh)$ counts. Whenever a count first becomes positive, retain one witness support mask from the responsible predecessor, adding j if the second term supplied it. These masks have at most M bits. Every count is at most 2^{M-1} . The stated operation and storage bounds follow. Also $\sum_r D_{M-1,h,r} = \binom{M-1}{h}$, which checks the total count.

The constant coefficient of $G_U(Y) = \prod_{j \in U} (Y - \lambda^j)$ is $(-1)^h \lambda^{\sum_{j \in U} j}$. Thus the residue classes counted by the recurrence are exactly the constant-coefficient classes. Choose a largest class and its retained support U_0 . The earlier construction sets

$$f(u) = a_0 u^{-(m+2)} G_{U_0}(u^2), \quad F_U(u) = a_0 u^{-(m+2)} (G_{U_0}(u^2) - G_U(u^2)).$$

Here $m = 64$, and a_0 is the explicit phase factor from Appendix B. The common leading and constant coefficients put every F_U in $[-64, 64]$. Its difference from f vanishes at precisely the two-point fibers of U , giving 132 agreements. Distinct supports give distinct polynomials. Rotation to the canonical coset preserves these statements. The table requires only G_{U_0} , not the other polynomials in the family. $\square$

A closed form and a fixed defining support. In this instance, every odd residue has exactly

$$N_{\text{exact}} = \frac{\binom{255}{66} + \binom{127}{33}}{256}$$

supports. This also follows from the roots-of-unity filter. For a character of order $d = 2^s > 1$, the support generating function is $(1 - t^d)^{256/d} / (1 + t)$; its degree-66 coefficient is $A_d = (-1)^{\lfloor 66/d \rfloor} \binom{256/d-1}{\lfloor 66/d \rfloor}$. Thus the count at residue r is

$$\frac{1}{256} \left(\binom{255}{66} + \sum_{d=2,4,\dots,256} c_d(r) A_d \right),$$

where $c_d(r)$ is $d/2$ if $d \mid r$, $-d/2$ if $d/2 \mid r$ but $d \nmid r$, and zero otherwise. For odd r only $d = 2$ contributes, giving the displayed closed form. For even r the $d = 2$ contribution has the opposite sign; the exact inequality $\sum_{d \geq 4} (d/2) |A_d| < 2 \binom{127}{33}$ shows its remaining terms cannot reach the odd-residue count. The script independently checks this character formula against all 256 dynamic-programming counts.

One may fix the defining support outright as $U_0 = \{2, 3, \dots, 68\} \setminus \{40\}$: it has 66 elements and sum $2305 \equiv 1 \pmod{256}$. Hence even the search for a defining support is unnecessary for this particular table.

Concrete artifact. The companion `checks/construct_m31_bank.py` performs 3,759,360 integer updates; the recorded run took approximately 0.6 seconds for this counting stage. It writes all 256 exact residue counts, a defining support, the canonical domain parameters, and the 512 base-field table entries to `checks/explicit_m31_bank.json`. A checked generic-sample fixture also supplies a normalized quotient witness in the interior Laurent band $[-63, 63]$ and verifies all 512 coordinate identities, with exactly 132 agreements.

The fixture's claimed value is chosen to realize one specified sample/challenge pair. It is not asserted to be the favorable fixed claim promised by Theorem C.4. The family is now explicit as a compact membership rule and exact count, and its fixed table is fully constructed. Finding the favorable claim and recovering a witness efficiently for a fresh challenge remain open. The numerical probability bounds above may continue to use the slightly smaller $N = \lceil \binom{255}{66} / 256 \rceil$, by retaining that many members of the class.

C.7 Witness recovery as a constrained subset product

The implicit description of the family permits an exact search for a witness after sampling. The following reduction isolates the remaining algorithmic problem and gives a complete, although expensive, baseline. The claim v is held fixed throughout the search.

Proposition C.8 (Exact recovery criterion and a complete search). *Use the two-fiber family of Section C.6, with $h = m + 2$, defining support U_0 , residue r , and relative sample parameter $a \notin H$.*

Assume $\mu_{2M} \subset H$. For any desired value $\beta \in E$, a support U in this family satisfies $F_U(a) = \beta$ if and only if

$$|U| = h, \quad \sum_{j \in U} j \equiv r \pmod{M}, \quad \prod_{j \in U} (a^2 - \lambda^j) = G_{U_0}(a^2) - \frac{a^h}{a_0} \beta. \quad (50)$$

Each factor on the left is nonzero. A zero target therefore has no solution. For a nonzero target, partition $\{1, \dots, M-1\}$ into parts of sizes s and $M-1-s$. There is a complete search using a table of at most

$$R = \sum_{k=0}^{\min(h, M-1-s)} \binom{M-1-s}{k}$$

right-hand subset records and at most

$$L = \sum_{k=0}^{\min(h, s)} \binom{s}{k}$$

dictionary lookups per target, in addition to constructing the subset products. Keeping all records with equal keys recovers every matching support; keeping one per key suffices to find a witness when one exists.

Proof. Substitute $F_U(a) = a_0 a^{-h} (G_{U_0}(a^2) - G_U(a^2))$ and rearrange. If $a^2 = \lambda^j$, then a is one of its two square roots in $\mu_{2M} \subset H$, contradicting $a \notin H$.

For a right-hand subset V , store its support under the key

$$\left(|V|, \sum_{j \in V} j \pmod{M}, \prod_{j \in V} (a^2 - \lambda^j) \right).$$

For each left-hand subset W , with cardinality k , residue t , and nonzero product P_W , look up $(h-k, r-t \pmod{M}, \tau/P_W)$, where τ is the target in (50). The parts are disjoint, so every returned union has the required cardinality, residue, and product. Conversely, each solution has exactly one such decomposition, which the search visits. The subset counts give the stated bounds. $\square$

For the scheduled transfer, set $\beta = v + g_1(\alpha)/g_0(\alpha)$ when $g_0(\alpha) \neq 0$ and this ratio belongs to H . A recovered support supplies F_U , and exact polynomial division then supplies the interior-band quotient witness from Theorem C.4. Preparing the subset-product table depends on a and must generally be repeated when the sample changes. It can be reused for multiple targets at the same sample. This reduction does not identify the favorable fixed claim from the averaging theorem.

Fresh-query check on a smaller valid circle code. The companion `checks/verify_witness_search.py` fixes $v = 0$ before generating samples. It uses $p = 23$, $n = 24$, $M = 12$, $h = 4$, $m = 2$, and the full base-field circle, with defining support $U_0 = \{1, 2, 3, 7\}$ and residue 1. The resulting family has 27 members of outer Laurent degree at most 2, each with 8 agreements. The tower $H = \mathbb{F}_{23}[i]$, $E = H[w]$, $w^2 = 2 + i$, is a field: the script checks $(2 + i)^{(23^2-1)/2} = -1$. This is a circle-code experiment; its length is not a Stwo FFT size.

The deterministic experiment chooses 128 distinct admissible samples using a recorded pseudorandom seed. At each sample it tests every one of the 529 challenges in H , together with eight specified outside- H controls. On all 68,736 queries, the returned support set agrees exactly with an independent enumeration of all 27 family members. There are 135 successful queries: 128 have

$\alpha = 0$, which always recovers $F_{U_0} = 0$, and seven have $\alpha \neq 0$. No outside- H control succeeds in this fixture. For all 135 recovered witnesses, exact division yields coefficients in H with Laurent support $[-1, 1]$; all 3,240 coordinate identities pass, with exactly eight agreements per witness. The split uses 31 left-hand and 57 right-hand records per sample. These challenges are deliberately exhaustive in H ; the observed success fraction is not an estimate for uniform challenges in E .

Cost at the displayed M31 parameters. For the direct implementation with $s = 127$, $M = 256$, and $h = 66$, the complete right-hand table would have

$$R = \sum_{k=0}^{66} \binom{128}{k} = 228210150822262702568674774219878072803 \approx 2^{127.4236}$$

records, before accounting for bytes per record. Thus this baseline does not give a practical M31 decoder. This is the cost of the specified complete enumeration, not a lower bound on witness-search complexity or a claim that more selective algorithms cannot improve it. The useful conclusion is an exact algebraic target and a checked recovery procedure for fresh queries with a claim fixed in advance.

D Checks and provenance

D.1 Exact interval certificates

The comparisons in the main text follow from the moment-class bounds and the exact certificates below. All examples have $k = n/2$, agreement $t = k + s$, and gap $\eta = s/n$. The last column gives a strict lower bound on $\log_2 L$.

p	n	s	Certified strict lower bound on $\log_2 L$
$2^{31} - 1$	154	5	61.74656
$2^{61} - 1$	548	9	242.16447
$2^{127} - 1$	64	2	42
$2^{127} - 1$	2158	17	995.98234
$2^{521} - 1$	4096	8	3634
$2^{521} - 1$	29176	56	14233.01899

For the length-64 row the original denominator product is 32,285,041, giving 50,186,958,429 code-words. The conditional dispersion certificate (10) strengthens this to 2,674,048,230,075; the exact enumeration in Section 7.1 further gives 5,552,914,238,035 $> 2^{42}$. The length-4096 row uses Lemma 2.1. The four decimal rows use Proposition 2.14, with rational upper bounds for its denominator and outward rational intervals for logarithms. Their checker is in `research/prime_exponent_coefficients/`. It verifies primality by Lucas–Lehmer and checks the strict sufficient Elias inequality

$$p^s t^t (n - t)^{n-t} > n^n$$

using integers. The two remaining rows satisfy the stronger sufficient condition $p^s > 2^n$. For the last row, $\eta = 1/521$, so a bound $2^{c_2/\eta}$ requires

$$c_2 > \frac{14233.01899}{521} > 27.31865.$$

The corresponding forced coefficients for the three smaller primes are 2.00475, 3.97715, and 7.84601. These comparisons set $c_1 = 1$; for a different prefactor, subtract $\eta \log_2 c_1$ from each necessary lower bound on c_2 . They concern interval domains, not the prescribed circle domain. The finite rational-line example uses the conditional first-moment count

$$N = 8,634,334,451,040,106$$

from Section 2.3. The exact inequality $2^{127} - 1 > 64 + 32 \binom{N}{2}$ certifies a pole with all N scalar labels distinct. These integers are retained here for reproducibility; the comparisons in the main text only need their rounded powers of two.

D.2 Certificates at additional rates

For Corollary 6.2, the exact guaranteed list size is

$$L = \left\lceil \binom{63}{17} / 64 \right\rceil = 15,827,733,944,904, \quad \eta = \frac{1,048,575}{33,554,432}.$$

Integer arithmetic verifies $p^{100} > 2^{3099}$, $3^{100000} > 2^{158496}$, and $L^{2000} > 2^{87695}$. Since $H_2(1/4) = 2 - \frac{3}{4} \log_2 3 < 0.81128$, these give

$$\log_2 L - \frac{H_2(1/4)}{\eta} > \frac{87695}{2000} - \frac{81128}{100000\eta} = \frac{937767635977}{52428750000} > 17.8865.$$

This implies the rounded ratio in the corollary. For the Elias check, write $\theta = 3/4 - \eta$. The exact rational inequality

$$\eta \frac{3099}{100} > 1 - 2(\theta - 1/2)^2 \geq H_2(\theta)$$

and $\log_2 p > 3099/100$ imply $H_p(\theta) < 3/4$. The same construction at rates $1/8$, $1/16$, and $3/4$ gives list ratios exceeding 2^{13} , 2^{15} , and 2^{13} , respectively. The rate-3/4 space has even dimension but is not a native power-of-two FFT size. No such restriction is needed for the mathematical code.

For Corollary 6.3, put $Q = p^4 - p^2$. The exact counts are

$$N = \left\lceil \frac{\binom{255}{36}}{256p^2} \right\rceil = 73,497,635,817,391,996,020,834,$$

$$J = \left\lceil \frac{NQ}{Q + 32(N-1)} \right\rceil = 73,497,635,817,383,868,141,347.$$

Here $\eta = 655359/33554432$. Integer arithmetic verifies $7^{100000} > 2^{280735}$, $7^{1000} < 2^{2808}$, and $J^{1000} > 2^{75960}$. Set

$$H^+ = 3 - \frac{7 \cdot 280735}{8 \cdot 100000} > H_2(1/8).$$

Concavity gives $H_2(7/8 - \eta) < H^+ + \eta(2808/1000)$, whereas exact rational arithmetic gives $\eta(3099/100 - 2808/1000) > H^+$. Together with $\log_2 p > 3099/100$, this certifies the strict Elias bound. Finally,

$$\log_2 J - 26 - \frac{H_2(1/8)}{\eta} > \frac{75960}{1000} - 26 - \frac{H^+}{\eta} = \frac{9064121863}{409599375} > 22.129.$$

After puncturing, the exact parameters are $\rho' = (2^{23} - 1)/(2^{26} - 2)$ and $\eta' = (5 \cdot 2^{18} - 1)/(2^{26} - 2)$. Exact rational logarithm intervals separately certify the Elias bound and a ratio greater than 2^{22} at these changed parameters.

For the quadratic-field example, put $Q_2 = p^2 - 128$. Then

$$N_2 = \left\lceil \left(\binom{127}{18} / 128 \right) \right\rceil = 25,460,924,804,916,891,029,$$

$$J_2 = \left\lceil \frac{N_2 Q_2}{Q_2 + 16(N_2 - 1)} \right\rceil = 285,003,988,493,147,037.$$

Here $\eta = 786431/33554432$. The same H^+ above satisfies $\eta(3099/100 - 2808/1000) > H^+$, proving the strict Elias bound. The integer inequality $J_2^{10000} > 2^{579837}$ gives

$$\log_2 J_2 - 26 - \frac{H_2(1/8)}{\eta} > \frac{579837}{10000} - 26 - \frac{H^+}{\eta} = \frac{13827665031}{1572862000} > 8.791.$$

Thus the ratio exceeds 400. An independent fixture over $\mathbb{F}_{312}$ checks all 945 allowable poles for an anchored coefficient class of 87 supports. It verifies the collision bound and all witness identities for a best pole and for the allowed pole zero; these yield 87 and 84 distinct labels, respectively.

D.3 Certificates for the compact examples

Corollary 6.4 uses the same integers J_2 and J as the two eighth-rate constructions above. The two new Elias comparisons follow from

$$p^{896} 873^{873} 151^{151} > (p-1)^{873} 1024^{1024},$$

$$p^{3584} 3505^{3505} 591^{591} > (p-1)^{3505} 4096^{4096}.$$

Taking logarithms gives $H_p(873/1024) < 7/8$ and $H_p(3505/4096) < 7/8$. Also,

$$J_2^2 > 2^{115}, \quad J^3 > 2^{227}, \quad 7^5 > 2^{14}.$$

The last inequality gives $H_2(1/8) < 11/20$, so the logarithms of the two ratios exceed

$$\frac{115}{2} - 11 - \frac{11}{20} \frac{1024}{23} > 22, \quad \frac{227}{3} - 13 - \frac{11}{20} \frac{4096}{79} > 34.$$

All these comparisons use exact integer or rational arithmetic.

D.4 Verification and source notes

The accompanying files give exact integer certificates for the finite bounds and exhaustive small-field checks of the constructions. They also include a small explicit test of the power-fiber lift. The proofs establish the asymptotic claims; the experiments check formulas and boundary cases. Large circle lists and suitable poles are obtained by counting, rather than enumerated.

The numerical certificates and construction checks are included with this draft. They use exact integer or finite-field arithmetic. The protocol-related models do not constitute native verifier executions. The source comparisons concern the cited repository revisions; they are not claims about later implementations. This is a working draft with machine-assisted internal checks, and it has not undergone external peer review.

E Nearly complete rational fibers and the Frobenius index

We prove Theorem 2.12. The unit-equation and Wronskian method follows the approach of Corvaja–Zannier [5, Proposition 2]; the linear-independence argument below uses the special relation $p = \ell n + \epsilon$. We give the needed height calculation explicitly.

Proof of Theorem 2.12. Work over $k = \overline{F}$. The hypotheses imply $B < p$, $2\ell B < p$, and $p > n$. In particular, R is separable.

The collision components and their boundary. The reduced curve $R(x) = R(y)$ in $\mathbb{P}^1 \times \mathbb{P}^1$ has bidegree (B, B) and contains the diagonal. Suppose it has a component other than a torus graph $x/y = \text{constant}$ or $xy = \text{constant}$. On its smooth projective normalization C , let $a = \deg x$ and $b = \deg y$. Neither projection is constant, and $1 \leq a, b \leq B - 1$ because the diagonal is a different component. Put $K = k(C) = k(x, y)$. Let S consist of the zeros and poles of x, y , together with the points above $x \in \mu_n \setminus T$ or $y \in \mu_n \setminus T$. The genus and boundary bounds give

$$\begin{aligned} g_C &\leq (a-1)(b-1), & |S| &\leq 2a + 2b + c(a+b), \\ \chi &:= 2g_C - 2 + |S| &\leq 2ab + c(a+b). \end{aligned} \tag{51}$$

The four units. Put $u = x^n$, $v = y^n$, and $w = (u-1)/(v-1)$. Because $\deg x, \deg y < p$, both coordinates are separating; so are u and v , since $p \nmid n$. In particular, $[K^p(v) : K^p] = p$. Outside S , $u-1$ and $v-1$ have exactly the same zeros, all simple. Indeed, outside S , $x \in \mu_n$ means $x \in T$; its complete geometric fiber consists of B distinct points in T , so $y \in T$, and conversely. Such fibers are unramified, and the collision curve is smooth with both coordinate projections unramified at their pairs. Consequently w is an S -unit, meaning that all its zeros and poles lie in S . The four S -units

$$-1, \quad u, \quad w, \quad -wv$$

sum to zero.

Independence over K^p . We claim that $w, 1, u$ are linearly independent over K^p . Otherwise there are $A_0, A_1, A_2 \in K^p$, not all zero, such that

$$P(U, V) = A_0(U-1) + (A_1 + A_2U)(V-1)$$

vanishes at (u, v) . This polynomial is nonzero and has bidegree at most $(1, 1)$. Remove the common factor of its two coefficients as a polynomial in U ; that factor has degree at most one in V and cannot vanish at v . The resulting primitive polynomial is linear in U , since an equation in v alone of degree at most one would contradict $[K^p(v) : K^p] = p$. It defines a graph $U = T(V)$. The function T is nonconstant, since $u \notin K^p$, and is therefore a Möbius transformation.

Set $c = x^p, d = y^p \in K^p$. The congruence defining ℓ gives

$$x = (c/u^\ell)^\epsilon, \quad y = (d/v^\ell)^\epsilon.$$

Hence (u, v) also annihilates the numerator $G(U, V)$ of

$$R((c/U^\ell)^\epsilon) - R((d/V^\ell)^\epsilon),$$

whose bidegree is at most $(\ell B, \ell B)$ over K^p . If the primitive P and G were coprime, their resultant in U would be a nonzero polynomial in V of degree at most $2\ell B < p$ vanishing at v . This is impossible. Thus P divides G , and

$$R((c/T(V)^\ell)^\epsilon) = R((d/V^\ell)^\epsilon) \quad (52)$$

as rational functions of V .

Over an algebraic closure of K^p , choose $\alpha^\ell = c$, $\beta^\ell = d$, and set

$$F_0(Z) = R(Z^{\ell\epsilon}), \quad M(Z) = \alpha/T(\beta/Z).$$

Equation (52) says $F_0 \circ M = F_0$. The deck group of F_0 has order at most $\ell B < p$ and contains the scalings by ℓ th roots of unity. The tame finite-subgroup classification for PGL_2 gives cyclic, dihedral, A_4 , S_4 , or A_5 [6]. The last three have no element of order at least six. In a cyclic or dihedral group, an element of order at least three has the rotation subgroup's pair of fixed points. Thus every deck map preserves the pair $\{0, \infty\}$ fixed by the scalings, and

$$M(Z) = \zeta Z \quad \text{or} \quad M(Z) = \zeta/Z.$$

Here $\zeta \in k$: since F_0 is defined over the algebraically closed field k , any deck map takes every k -point into a finite fiber consisting of k -points, and the images of three points determine its coefficients.

In the first case $u/v = \alpha/(\zeta\beta)$. Taking ℓ th powers and using $p = \ell n + \epsilon$ gives $(x/y)^\epsilon = \zeta^\ell$, so x/y is constant. In the second case $uv = \alpha\beta/\zeta$, and the same argument gives $(xy)^\epsilon = \zeta^\ell$, so xy is constant. Either case makes the chosen component a torus graph, contrary to its selection. This proves the claimed independence.

The height contradiction. Denote the four units by f_0, f_1, f_2, f_3 . Any three are independent over K^p , and their ordinary three-row Wronskians are nonzero and agree up to sign. For an explicit check, $u' \neq 0$ and

$$W(1, u, w) = u'w'' - u''w' = (u')^2(w'/u')'.$$

If this were zero, $w'/u' \in K^p$ and then $w - (w'/u')u \in K^p$, contradicting independence. Here K^p is the constant field of a nonzero derivation on K . Multiplying a Wronskian with derivative orders 0, 1, 2 by $(dt)^3$ gives a well-defined nonzero rational cubic differential Ω on C .

At a point $P \in S$, choose the three functions that omit a function of smallest valuation. Differentiation in a local parameter lowers valuation by at most one, so

$$\mathrm{ord}_P \Omega \geq \sum_{i=0}^3 \mathrm{ord}_P f_i - \min_i \mathrm{ord}_P f_i - 3.$$

Outside S , all local derivatives are regular, giving $\mathrm{ord}_P \Omega \geq 0$. Since each f_i has a degree-zero divisor supported on S , summing yields

$$H := - \sum_P \min_i \mathrm{ord}_P f_i \leq 3(2g_C - 2 + |S|) = 3\chi.$$

The tuple contains a nonzero constant and u , while v is a ratio of two coordinates. The degree of any such ratio is at most the projective height, so $H \geq n \max(a, b)$. By (51),

$$n \max(a, b) \leq 6ab + 3c(a + b) \leq 6(\min(a, b) + c) \max(a, b).$$

This contradicts $n > 6(B - 1 + c)$. Every component is therefore a torus graph. Each has bidegree $(1, 1)$, so the degree- B map has B geometric deck transformations and is Galois. Each deck map preserves T , which is nonempty. Evaluating it at a point of T shows that its scaling constant belongs to μ_n . Thus the entire deck group preserves μ_n .

Small defect forces free action. The rotation subgroup has order $d \mid n$. If there are inversions, the whole group has order $B = 2d$; otherwise $B = d$. Thus $B \mid 2n$. Since T is nonempty, $B \leq n$, so dyadic n forces $B \mid n$ even without a small-defect assumption. A nonzero finite point is fixed by no nonidentity rotation and by at most one inversion, so every nonfree orbit has $B/2$ points. Fibers of a Galois quotient are orbits; thus any nonfree orbit in μ_n lies entirely outside T . The inequality $c < B/2$ rules this out. If the group has only rotations, its action is already free. In either case all domain orbits have B points, proving full balance and $B \mid n$. Corollary 2.11 now applies. $\square$

The cutoff $c < B/2$ is strict. On $\mu_{15} \subset \mathbb{F}_{181}$, $R = X + X^{-1}$ has seven two-point fibers and one fixed point. Here $B = 2$, $c = 1 = B/2$, and $15 > 6(B - 1 + c)$, but $B \nmid 15$.

The exact checks in `research/structured_domains/frobenius_index/` verify the pinned arithmetic, nonzero Wronskians on sample non-diagonal collision conics in both congruence classes, and small finite-field deck groups. A separate exhaustive enumeration checks 27,273,130 disjoint pairs of possible complete fibers on five small domains, including non-Galois examples outside the theorem's hypotheses. These computations supplement the proof; they do not establish the general classification. The additional near-balance checker tests 26,880 exact arithmetic inequalities and 1,823,695 rational pencils. Its qualifying positive-defect cases all have torus deck groups; planted ramified examples also check the strict free-action cutoff.

F A limit on binary-affine support transfer

The large equal-prefix classes in this paper need not have a simple binary-linear description. The following bound gives one precise obstruction to such a description. It concerns affine families of *indicator vectors*, with addition over $\mathbb{F}_2$, even though the evaluation points lie in a field of different characteristic.

Theorem F.1 (Binary-affine locator classes). *Let $D = \{\alpha_1, \dots, \alpha_n\}$ be distinct elements of a field of characteristic different from two. Let $\mathcal{A} \subseteq \{0, 1\}^n$ be an affine subspace over $\mathbb{F}_2$ of dimension r . Suppose every word has integer weight t , and the corresponding monic degree- t locator polynomials have the same coefficients in degrees $t - 1, \dots, t - s$, where $0 \leq s < t$. Then*

$$r \leq \left\lfloor \frac{\min\{t, n - t\}}{s + 1} \right\rfloor, \quad |\mathcal{A}| \leq 2^{\lfloor \min\{t, n - t\} / (s + 1) \rfloor}.$$

The dimension bound is attained over suitable odd prime fields.

Proof. The case $r = 0$ is immediate. Parameterize the family injectively by $u \in \mathbb{F}_2^r$, writing the j th indicator as

$$x_j(u) = b_j + \langle a_j, u \rangle \pmod{2}, \quad a_j \in \mathbb{F}_2^r, \quad \epsilon_j = (-1)^{b_j}.$$

The nonzero column labels a_j span $\mathbb{F}_2^r$. For a nonzero label a , let C_a^+ and C_a^- be its coordinate classes with $b_j = 0$ and $b_j = 1$, respectively. The constant integer weight condition, expanded in the real-valued characters $\chi_a(u) = (-1)^{\langle a, u \rangle}$, gives $|C_a^+| = |C_a^-| =: h_a$. In every word exactly h_a coordinates of their union are selected. If c nonzero labels occur, then

$$r \leq c, \quad \sum_{a \neq 0} h_a \leq \min\{t, n - t\}.$$

We show that each occurring label has $h_a \geq s + 1$.

Let $L_u(X)$ denote the locator at u . Equality of the specified coefficients means, in the Laurent series field at infinity, that

$$L_u/L_0 = 1 + O(X^{-s-1}).$$

For a fixed nonzero a , take the product $\prod_u L_u(X)^{\chi_a(u)}$. Since $\sum_u \chi_a(u) = 0$, this product is also $1 + O(X^{-s-1})$. Character orthogonality over the integers gives

$$\sum_u \chi_a(u) x_j(u) = \begin{cases} -\epsilon_j 2^{r-1}, & a_j = a, \\ 0, & a_j \neq a. \end{cases}$$

Writing $F_a^\pm$ for the locators of $C_a^\pm$, we obtain

$$\left(\frac{F_a^-}{F_a^+} \right)^{2^{r-1}} = 1 + O(X^{-s-1}).$$

The ratio has leading term one, since both polynomials have degree h_a . Raising such a Laurent series to an integer power invertible in the field preserves its first nonzero nonconstant term. Thus $F_a^-/F_a^+ = 1 + O(X^{-s-1})$. The two disjoint nonempty root sets give different monic polynomials, so their nonzero difference has degree at most $h_a - s - 1$. Hence $h_a \geq s + 1$. Summing and using $r \leq c$ proves the bound.

For sharpness put $B = s + 1$ and use r disjoint pairs of full B -element fibers of X^B . Select one fiber in each pair independently, then add fixed selected and unselected coordinates as needed. Each fiber locator is $X^B - a$, so the varying locator product has the same first $B - 1$ coefficients. The common fixed factor preserves that prefix. The fiber-lifting argument of Lemma 4.2 supplies arbitrarily large odd primes for the required finite collection of fibers. With $r = \lfloor \min\{t, n - t\}/B \rfloor$, the remaining fixed-coordinate counts $t - rB$ and $n - t - rB$ are nonnegative, giving equality. $\square$

Finite consequences. For the $(n, t, s) = (64, 34, 2)$ integer class, every affine binary subfamily has at most $2^{10} = 1024$ members. For $(157, 68, 5)$ the bound is $2^{11} = 2048$. The much larger certified classes therefore require nonlinear families of support indicators. Likewise, choosing 272 of 511 packet labels with fourteen fixed leading coefficients permits at most $2^{15} = 32768$ members in an affine binary subfamily. This does not exclude nonlinear support families or other transfers from binary-field constructions.

The characteristic-two exception. On $\mathbb{F}_{16}$, let a range over an affine binary hyperplane not containing zero and take the supports

$$\{x \in \mathbb{F}_{16} : \text{Tr}_{\mathbb{F}_{16}/\mathbb{F}_2}(ax) = 1\}.$$

Their indicators form an affine binary family of dimension three, and every support has size eight. The monic locator is a normalized trace polynomial plus a constant, with possible degrees only 8, 4, 2, 1, 0. Its first three coefficients below the leading term are zero. Thus $(n, t, s, r) = (16, 8, 3, 3)$, violating the bound $r \leq 2$ that would follow in odd characteristic. Here taking a power of two can raise a Laurent vanishing order, exactly the step excluded above.

The checker in `research/binary_affine_locator/` enumerates 103,416 small affine binary spaces and verifies all 8,407 constant-weight families, including an odd-characteristic example with characteristic no greater than the prefix length. It also checks twenty sharpness examples and the $\mathbb{F}_{16}$ exception. These finite checks supplement the proof.

Corollary F.2 (Subquadratic transfer of punctured binary subspaces). *Let $V = \mathbb{F}_2^m$, $N = 2^m$, and $m \geq 4$. Injectively label $V \setminus \{0\}$ by elements of a field of characteristic different from two. Let $\mathcal{C}$ be a collection of codimension-two linear subspaces $W \leq V$ whose labeled punctured supports $W \setminus \{0\}$ have monic locators with the same first $N/8 - 1$ coefficients below the leading term. Put*

$$M_N = \left\lfloor \frac{1 + \sqrt{4N - 7}}{2} \right\rfloor.$$

Then

$$|\mathcal{C}| \leq \left\lfloor \frac{(N-1)M_N}{3} \right\rfloor = O(N^{3/2}).$$

Proof. Fix a nonzero functional $a \in V^*$. The codimension-two subspaces with $a \in W^\perp$ are $W_b = \ker a \cap \ker b$, indexed by nonzero $b \in Q := V^*/\langle a \rangle$, a binary group of size $N/2$. Let $S_a \subseteq Q \setminus \{0\}$ index the members of $\mathcal{C}$ of this form. On $\ker a$, the binary indicator of W_b is $1 + b(x)$; outside $\ker a$ it is zero. This dependence on b is affine and injective, and puncturing the origin preserves affine dimensions.

If S_a contained an affine binary plane, its four punctured supports would form an affine family of dimension two and constant weight $t = N/4 - 1$. Theorem F.1, with $s = N/8 - 1$, would give

$$2 \leq \left\lfloor \frac{t}{s+1} \right\rfloor = 1,$$

a contradiction. Thus S_a contains no affine plane. In a binary group, equality of the sums of two different unordered pairs of distinct elements gives four distinct elements forming an affine plane. Hence all $\binom{|S_a|}{2}$ pair sums are distinct and nonzero, so

$$\binom{|S_a|}{2} \leq N/2 - 1, \quad |S_a| \leq M_N.$$

Every $W \in \mathcal{C}$ has three nonzero elements in $W^\perp$. Double-counting gives $3|\mathcal{C}| = \sum_{a \neq 0} |S_a| \leq (N-1)M_N$. $\square$

In characteristic two, every binary additive domain of size N has all $(N-1)(N-2)/6$ codimension-two subspaces in this prefix class. Indeed, their degree- $N/4$ subspace locators are linearized: the next possible degree is $N/8$. Dividing by X gives exactly the punctured prefix above. The linearized form follows inductively from $L_{U \oplus \langle b \rangle} = L_U^2 + L_U(b)L_U$. The corollary therefore rules out preserving a positive asymptotic fraction of this quadratic family under any odd-characteristic relabeling. It does not exclude different locator identities or other transfers from binary-field constructions.

The additional checker tests 24,825 prefix classes across 390 injective relabelings, including characteristic three with prefix length three. It verifies the dual-incidence counts and absence of repeated pair sums, and checks 14,364 affine-plane parameterizations. The full binary families over $\mathbb{F}_{16}$ and $\mathbb{F}_{32}$ give negative controls of sizes 35 and 155, exceeding the respective odd-characteristic bounds 20 and 62.

A sublinear bound within every fixed multiplicative orbit

The preceding bound still permits a linear-sized shared-prefix family, which could be enough for the padding reduction. A stronger obstruction applies within a fixed multiplicative orbit of binary subspaces.

Theorem F.3 (Transfer within a multiplicative orbit). *Let $K = \mathbb{F}_{2^m}$, $N = 2^m$, and let $U \subset K$ be an r -dimensional binary subspace, with $1 \leq r < m$. For $a \in K^*$ put*

$$H_a = \{v \in K : \text{Tr}_{K/\mathbb{F}_2}(auv) = 0 \text{ for all } u \in U\}.$$

Injectively label $K \setminus \{0\}$ by elements of a field of characteristic different from two. Let $\mathcal{C}$ be a collection of distinct H_a whose punctured locators share their first s coefficients below the leading term, where $0 \leq s \leq N/2^r - 2$. Define

$$b = \left\lceil \log_2 \frac{(2^r - 1)N}{s + 1} \right\rceil, \quad h = |\{c \in K^* : cU = U\}|,$$

and assume $b \leq m$. Then

$$|\mathcal{C}| \leq \frac{4}{h} N^{1-2^{1-b}}.$$

For fixed r and a fixed positive prefix fraction $(s + 1)/N$, this is $o(N)$. In particular, for $r = 2$, $s = N/8 - 1$, and $m \geq 6$,

$$|\mathcal{C}| \leq \frac{4}{h} N^{15/16}, \quad h \in \{1, 3\}.$$

The case $h = 3$ is the family of $\mathbb{F}_4$ -linear hyperplanes.

Proof. Let $A \subset K^*$ consist of all normals for members of $\mathcal{C}$. Nondegeneracy of the trace pairing shows that $H_a = H_{a'}$ exactly when $aU = a'U$, so $|A| = h|\mathcal{C}|$. Newton identities express each power sum in the corresponding elementary symmetric functions without division. Thus the moments

$$M_j(a) = \sum_{v \in H_a \setminus \{0\}} x_v^j$$

are independent of $a \in A$ for $1 \leq j \leq s$; the zeroth moment is the common size $N/2^r - 1$. This step does not require the label-field characteristic to exceed s .

Suppose A contains an affine binary b -flat $a_\epsilon = a_0 + \sum_{i=1}^b \epsilon_i w_i$ with independent w_i . For $\beta \in \mathbb{F}_2^b \setminus \{0\}$ define

$$c_\beta(v) = 2^{-b} \sum_{\epsilon \in \mathbb{F}_2^b} (-1)^{\beta \cdot \epsilon} \mathbf{1}_{v \in H_{a_\epsilon}}.$$

All its moments vanish: $\sum_{v \neq 0} c_\beta(v) x_v^j = 0$ for $0 \leq j \leq s$. The character expansion of the annihilator of U gives, with $\text{Tr} = \text{Tr}_{K/\mathbb{F}_2}$,

$$c_\beta(v) = 2^{-r} \sum_{\substack{u \in U \setminus \{0\} \\ \text{Tr}(uw_i v) = \beta_i \ (1 \leq i \leq b)}} (-1)^{\text{Tr}(ua_0 v)}.$$

For fixed $u \neq 0$, multiplication by u preserves independence of the w_i , so the displayed fiber has $N/2^b$ points. Therefore c_β is supported on at most $(2^r - 1)N/2^b \leq s + 1$ points.

At least one such vector is nonzero over the label field. Otherwise invertibility of the binary Walsh transform would make all 2^b incidence rows identical. A fixed H_a can have at most $2^r - 1$ normals: for any fixed nonzero $u_0 \in U$, their distinct products au_0 all belong to its same annihilator. Our choice of b has $2^b > 2^r - 1$, a contradiction. On the support of a nonzero c_β , the square Vandermonde matrix of moments from zero to one less than the support size is invertible, because the labels are distinct. Its vanishing moments give the final contradiction. Thus A has no affine binary b -flat.

An elementary induction bounds an affine- j -flat-free set B in a binary group of size q by $4q^{1-2^{1-j}}$. For $j = 1$ it has at most one point. For $j \geq 2$, pair its elements along each nonzero direction v . The occupied pairs form an affine- $(j-1)$ -flat-free set in the quotient of size $q/2$. With $\alpha = 2^{2-j}$,

$$|B|(|B| - 1) = \sum_{v \neq 0} |B \cap (B + v)| \leq 8(q-1)(q/2)^{1-\alpha}.$$

Consequently $|B| \leq 1 + 2 \cdot 2^{\alpha/2} q^{1-\alpha/2} \leq 4q^{1-\alpha/2}$. Apply this with $j = b$ and $q = N$ to A .

Finally, $\{c \in K : cU \subseteq U\}$ is a finite subfield: it is closed under addition and multiplication, and every nonzero element acts bijectively on U . Its binary degree divides r . For $r = 2$ it has size two or four, so $h = 1$ or 3 ; the latter case makes U a scalar copy of $\mathbb{F}_4$. $\square$

This rules out a linear-sized transfer from any bounded number of fixed-codimension multiplicative orbits. At codimension two, a cN -sized shared-prefix family would have to occupy at least $(c/4)N^{1/16}$ orbits. Arbitrary codimension-two families can span linearly many orbits, so the general linear-sized-transfer question remains open.

The exact checker verifies the character support formula on all 63 affine binary five-flats avoiding zero at $N = 64$, and 20 such flats at $N = 256$. It separately checks multiplicative orbits with stabilizers one and three, and a codimension-three orbit with seven-dimensional affine sections. Odd-field prefix tests, a characteristic-two negative control, and a small complete odd-field transfer complement these checks. They supplement the universal proof.

Short prefixes determine subspaces under multiplicative relabeling

For the natural multiplicative relabeling, a stronger statement applies to all binary subspaces simultaneously, without restricting their orbits.

Theorem F.4 (Multiplicative relabeling rigidity). *Let $K = \mathbb{F}_{2^m}$, $N = 2^m$, and let $\chi : K^* \rightarrow E^*$ be an injective multiplicative character into a field of characteristic different from two. Label $v \neq 0$ by $\chi(v)$, or by $a\chi(v) + b$ with $a \neq 0$. For each r -dimensional binary subspace $U \subset K$, with $1 \leq r < m$, let H_U be its annihilator under the absolute trace pairing, and put $t = 2^r - 1$. The first $2t - 1$ coefficients of the punctured locator of H_U determine U uniquely whenever that prefix fits in the locator. In characteristic zero or characteristic greater than t , the first t coefficients suffice.*

The latter cutoff is sharp: if $r \mid m$, the $(N-1)/t$ subspaces $U = c\mathbb{F}_{2^r}$ have punctured locators with their first $t-1$ coefficients all zero under the unshifted labeling.

Proof. Write $\psi(v) = (-1)^{\text{Tr}_{K/\mathbb{F}_2}(v)}$ and

$$G_j = \sum_{v \in K^*} \psi(v) \chi(v)^j, \quad 1 \leq j < N-1.$$

Every G_j is invertible. Indeed, substituting $x = uy$ in $G_j G_{-j}$ gives

$$G_j G_{-j} = \sum_{u \neq 0} \chi(u)^j \sum_{y \neq 0} \psi((u+1)y) = N.$$

The inner sum is $N-1$ at $u = 1$ and -1 otherwise, and the nontrivial multiplicative character χ^j has total sum zero. The integer N is nonzero in E .

Expand the indicator of H_U with additive characters and substitute $y = uv$ for each nonzero $u \in U$. The zero- u term vanishes, giving

$$\sum_{v \in H_U \setminus \{0\}} \chi(v)^j = \frac{G_j}{2^r} \sum_{u \in U \setminus \{0\}} \chi(u)^{-j}. \quad (53)$$

Thus equal locator prefixes, and hence equal power sums, give equal moments of the two t -element dual label sets $\{\chi(u)^{-1} : u \in U \setminus \{0\}\}$. When the characteristic is zero or exceeds t , reverse Newton identities recover the dual set from its first t moments.

In every other allowed characteristic, subtract the indicator measures of the two dual sets. Their zeroth moments agree, and the difference is supported on at most $2t$ distinct points. Moments zero through $2t - 1$ and an invertible square Vandermonde minor force the difference to vanish. Injectivity of χ then recovers U . All moments used have index less than $N - 1$, since $r < m$. For affine relabeling, the binomial expansion recovers the original χ -moments inductively; its diagonal coefficient is $a^j \neq 0$.

For sharpness, if $U = c\mathbb{F}_{2^r}$, its annihilator is $\mathbb{F}_{2^r}$ -linear. Its nonzero elements are a union of multiplicative cosets of $\mathbb{F}_{2^r}^*$. Their labels are full cosets of the t th roots of unity, so the locator is a polynomial in X^t . Its first $t - 1$ coefficients vanish. The scalar stabilizer has size t , giving $(N - 1)/t$ distinct subspaces. $\square$

At codimension two, three coefficients suffice in characteristic greater than three, and five suffice in every odd characteristic. Thus a fixed positive shared-prefix fraction cannot survive this multiplicative relabeling, even for two supports, once N is large enough. The sharpness example has only a constant prefix length; its agreement gap shrinks with N and does not provide the desired fixed-gap list.

The exact checker exhausts 26,114 subspaces in six prime-field fixtures, checking both multiplicative and affine-multiplicative prefixes. A separate $\mathbb{F}_{330}$ fixture checks all 155 codimension-two subspaces of $\mathbb{F}_{32}$, without the large-characteristic Newton assumption. It checks 126,361 primal-dual moment identities and 27 Gauss products. A codimension-three example at $N = 512$ has 73 locators with six common zero coefficients and pairwise distinct seventh coefficients. Arbitrary nonmultiplicative relabelings across many subspace orbits remain outside these results.

G Boundedly many roots imply linear full-support MCA

A polynomial solution variety of quadratic degree need not produce quadratically many exceptional labels on a received line. We prove an $O(n)$ bound for every candidate family with a bounded number of distinct roots, allowing both the roots and their multiplicities to vary.

Let $\mathcal{D}$ be a set of n distinct field elements and let f, g be arbitrary words on it. For a candidate P of degree at most D , set

$$S(P, z) = \{x \in \mathcal{D} : f(x) + zg(x) = P(x)\}.$$

Call (P, z) bad at threshold $A > D$ if $|S(P, z)| \geq A$ and no polynomials F_0, G_0 of degree at most D agree with f, g , respectively, on all of $S(P, z)$. Equivalently, g has no degree- D interpolant on this full support, since otherwise $F_0 = P - zG_0$ works. A label is bad if it has at least one bad candidate in the family.

Theorem G.1 (Bounded root count). *Fix $r \geq 1$ and $0 < \rho_- \leq \rho_+ < a \leq 1$. Suppose $\rho_- \leq D/n \leq \rho_+$ and $A \geq an$. There are constants $C_{r,a}$, $c_{r,a}$, and $n_0(r, a, \rho_-)$ such that, for $n \geq n_0$, every received line has at most $C_{r,a}n$ bad labels witnessed by polynomials with at most r distinct roots over an algebraic closure, if the characteristic is zero or $p > c_{r,a}n$. Nonzero scalar coefficients, root positions, and root multiplicities may all vary. Including the zero polynomial costs at most n additional labels.*

The constants below are explicit but very large. No general first-order MCA theorem or improved better.codes certificate follows. We may work over an algebraic closure when discussing roots: a degree- D interpolant on more than D base-field points descends to the base field.

G.1 Pencils and clusters with small residual degree

An affine polynomial pencil $P_z = U + zV$, with $\deg U, \deg V \leq D$, contains at most n bad labels. Indeed a bad full support contains a coordinate where $(f, g) \neq (U, V)$; agreement there forces $g \neq V$ and determines the unique label $z = (U - f)/(g - V)$.

Suppose a selected family of distinct bad labels has candidates HQ , where H is fixed with at most s distinct roots, and $\deg Q \leq K$. We can take $K \leq D - \deg H$. If three graph points (z, Q) are not collinear, their line-incidence determinant is a nonzero polynomial of degree at most K . Their common agreement support has size at most $K + s$, after dividing by H away from its roots. Assume $K + s \leq a^3 n/16$. For $L \geq 4/a$, counting ordered triples through coordinates gives a lower bound $a^3 n L^3/8$: if b_x counts agreeing candidates, use $(b_x)_3 \geq \max(b_x - 2, 0)^3$ and convexity. Noncollinear triples contribute at most $(a^3 n/16)L^3$. Each ordered pair determines a unique graph line, which contains at most n selected bad labels by the pencil observation. Collinear triples therefore contribute at most $n^2 L^2$. Consequently

$$L \leq 16n/a^3. \quad (54)$$

The same bound covers $L < 4/a$. This argument needs no characteristic hypothesis.

Set $\delta = a^3/(128r)$. Fix any set S of at most $3r$ algebraic roots and consider candidates whose total root multiplicity outside S is at most $r\delta D$. Partition their exponents on S into boxes of side $w = \lfloor \delta D \rfloor + 1$. Within each box factor the product H given by its lower-corner exponents. The residual degree is at most

$$|S|(w - 1) + r\delta D \leq 4r\delta D \leq a^3 n/32.$$

If $3r \leq a^3 n/32$, equation (54) applies to each box. There are at most

$$B_0 = (1 + \lceil 1/\delta \rceil)^{3r}$$

boxes, so this cluster contains at most Cn selected bad labels, where

$$C = 16B_0/a^3. \quad (55)$$

The bound is uniform in S . Factors may be considered over the algebraic closure; a pencil through two original base-field candidates has base-field coefficients.

G.2 Four candidates with private high-multiplicity roots

Lemma G.2. *Let $P_1, \dots, P_4$ be nonzero polynomials of degree at most D , each with at most r distinct roots. Suppose each P_i has a root of multiplicity at least δD which is not a root of any other P_j . For an integer $m \geq 1$, put*

$$N = \binom{m+3}{3}, \quad J = m+1, \quad M = N - J.$$

Assume

$$\delta D > \binom{N-1}{2}(4r-1), \quad \text{char } \mathbb{F} = 0 \quad \text{or} \quad p > mD.$$

For two independent homogeneous linear forms F, G in four variables, the gcd of $F(P_1, \dots, P_4)$ and $G(P_1, \dots, P_4)$, after removing factors supported on the union of the roots of the P_i , has degree at most

$$\frac{24D}{m+5} + \frac{(4r-1)(M-1)}{2}. \quad (56)$$

Proof. The N homogeneous degree- m monomials in the P_i are linearly independent. For otherwise take a minimal relation with q terms $f_1, \dots, f_q$. A ratio of two distinct monomials has an order of absolute value at least δD at some private root, so $q = 2$ is impossible. Any $q - 1$ terms have a nonzero ordinary Wronskian: their degrees are less than p , and row reduction to distinct leading degrees gives a nonzero Vandermonde leading coefficient. The same argument holds in characteristic zero.

Let S be the union of the roots of the P_i and A_0 the common gcd of the f_i . All Wronskians obtained by omitting one term are constant multiples. Omit a term of largest degree for the global degree bound, and a term of smallest order separately at each root for the local bounds. Comparing them yields

$$\max_i \deg f_i - \deg A_0 \leq \binom{q-1}{2} (|S| - 1).$$

The left side is at least the rational degree of any ratio f_i/f_j , and hence at least δD . Since $q \leq N$ and $|S| \leq 4r$, this contradicts the assumed threshold.

The degree- m part of the homogeneous ideal (F, G) has dimension M : its quotient is a polynomial ring in two variables, of degree- m dimension J . Evaluation in the P_i is injective by the preceding paragraph. Write d_α for the degrees of the evaluated monomials. A basis of the evaluated ideal has sum of degrees at most the sum of the largest M of the d_α . Indeed, order the formal monomials by degree and row-reduce from highest degree; each basis element has degree at most the degree of its distinct pivot. Thus its nonzero Wronskian W satisfies

$$\deg W \leq \sum_\alpha d_\alpha - \binom{M}{2}.$$

At a root $s \in S$, let $a_\alpha(s)$ be the monomial orders. Intersecting a codimension- J subspace with the span of monomials of order at least each integer k gives an adapted basis whose sum of orders is at least $\sum_\alpha a_\alpha(s) - J \max_\alpha a_\alpha(s)$. The dimension of each such intersection loses at most J ; summing these filtration dimensions proves the bound. Cancellations can only increase orders. A constant basis change preserves the order of W , so, using

$$\sum_{s \in S} \sum_\alpha a_\alpha(s) = \sum_\alpha d_\alpha, \quad \sum_{s \in S} \max_\alpha a_\alpha(s) \leq m \sum_{i=1}^4 \deg P_i,$$

we obtain

$$\sum_{s \in S} \text{ord}_s W \geq \sum_\alpha d_\alpha - Jm \sum_i \deg P_i - |S| \binom{M}{2}.$$

Let T be the gcd in the statement. It divides every evaluated ideal element. The identity $\text{Wr}(Th_1, \dots, Th_M) = T^M \text{Wr}(h_1, \dots, h_M)$ contributes at least $M \deg T$ additional zeros outside S . Combining the degree and order bounds gives

$$\deg T \leq \frac{4JmD}{M} + \frac{(4r-1)(M-1)}{2}.$$

Since $M = m(m+1)(m+5)/6$, this is (56). □

This is an elementary specialization of the high-power gcd phenomenon. A broader characteristic-zero result is Guo–Sun–Wang [38, Theorem 11]; the finite-characteristic statement above follows directly from the Wronskian proof.

G.3 Counting the remaining quadruples

Select one nonzero bad candidate at each bad label. Call a root heavy if its multiplicity is at least δD . A quadruple without a private heavy root for every member has some member P_i whose heavy roots all lie in the union of the other three root sets. Fix the other three candidates. Their root union has size at most $3r$, and every root of P_i outside it has multiplicity less than δD . The cluster bound (55) gives at most Cn choices for this selected bad label. Summing over the four possible members, there are at most $4CnL^3$ such ordered quadruples. This includes candidates with no heavy roots.

For every other quadruple at distinct labels, common agreement coordinates satisfy the two independent linear forms

$$\begin{aligned}(z_2 - z_3)P_1 + (z_3 - z_1)P_2 + (z_1 - z_2)P_3 &= 0, \\ (z_2 - z_4)P_1 + (z_4 - z_1)P_2 + (z_1 - z_2)P_4 &= 0.\end{aligned}$$

Lemma G.2, together with the at most $4r$ roots in the union, bounds their common support. Choose

$$m \geq 1536/a^4 - 5, \quad \frac{(4r-1)(M-1)}{2} + 4r \leq a^4 n/64,$$

as well as the lemma's hypotheses and $3r \leq a^3 n/32$. Their common support then has size at most $a^4 n/32$. For $L \geq 6/a$, fourth incidence gives

$$\frac{a^4 n L^4}{16} \leq \frac{a^4 n L^4}{32} + 4Cn^2 L^3.$$

Therefore

$$L \leq \frac{128C}{a^4} n = \frac{2048}{a^7} (1 + \lceil 128r/a^3 \rceil)^{3r} n. \quad (57)$$

The small- L case also satisfies this bound. At fixed positive rate, all degree and length thresholds eventually hold, and a sufficiently large fixed lower bound on p/n ensures $p > mD$. The zero polynomial contributes at most n more labels by the same single-coordinate argument as for a fixed pencil. This proves Theorem G.1.

G.4 A quadratic-degree solution surface

Fix $e \geq 1$ in characteristic zero or $p > 2e$. The nonzero polynomial solutions of degree at most $2e$ to

$$e^2 P^2 P''' + 3e(1-e)PP'P'' + (1-e)(1-2e)(P')^3 = 0 \quad (58)$$

are exactly $P = cH^e$ with $\deg H \leq 2$. At a root of multiplicity m , the coefficient of the lowest possible term, after removing a nonzero local scalar, is $m(m-e)(m-2e)$. Thus every root has multiplicity e or $2e$. Conversely substitution verifies every cH^e ; equivalently this is the denominator-cleared equation $(P^{1/e})''' = 0$. The monic base descends by uniqueness.

The projective map $[H] \mapsto [H^e]$ from the plane of quadratics is injective and base-point free, with injective tangent maps since $e \neq 0$ in the field. Its image has degree e^2 . Explicitly, cut by $P(u) = P(v)$ and $P(w) = P(v)$ for three distinct points u, v, w . There is no nonzero solution with $H(v) = 0$, so normalize $H(v) = 1$. The values $H(u), H(w)$ independently range over the e -th roots of unity, giving e^2 distinct reduced points by quadratic interpolation. This fixed third-order, jet-degree-three equation has an actual solution variety of degree $\Theta(n^2)$ at fixed positive rate, but Theorem G.1 gives only $O(n)$ bad labels in its large-characteristic regime.

The checkers in `research/bounded_root_mca/` verify unequal-degree Wronskian bookkeeping and cluster counting, including every full support of 1,025 bad labels in a length-4,096 example with at most two roots per candidate. The earlier `research/power_family_mca/` package verifies the quadratic-power classification and degree sections. These checks supplement the proofs. Families with a growing number of distinct roots and the fixed-gap quadratic lower-bound target remain open here.

H Actual first-order components have linear positive-dimensional degree

The degree of a Taylor-reconstruction cover can exceed the degree of the actual polynomial-solution locus. The following bound separates its curve and surface components. We then exhibit quadratically many actual isolated regular solutions, even at derivative degree one, and prove that this example still has only constantly many nearby labels at a fixed positive agreement gap. The full first-order MCA conjecture remains unresolved.

Theorem H.1 (Positive-dimensional actual components). *Let E be algebraically closed of characteristic zero or $p > D$, with $D \geq 1$. Let $Q(X, z, u, v) \neq 0$ have total degree at most $B \geq 1$ in (u, v) and degree at most H in z , with no restriction on its X -degree. In affine (z, P) coefficient space, let $\mathcal{V}$ be the Zariski closure of the actual regular solutions*

$$\deg P \leq D, \quad Q(X, z, P, P') = 0, \quad Q_v(X, z, P, P') \neq 0.$$

Put

$$\tau = \max\{0, 2D - 3\}, \quad b = 1 + \tau(B - 1), \quad T = b + \tau H, \quad K = H + BT, \quad q = B + H.$$

Every component of $\mathcal{V}$ has dimension at most two. If J_i is the sum of the projective degrees of its i -dimensional irreducible components, then

$$J_1 \leq qK, \quad J_2 \leq qT.$$

In particular $J_1 + J_2 = O_{B,H}(D)$. The reduced generic challenge fiber of the surface components has total degree at most $\deg_v Q$. The bounds concern actual components defined over E before choosing a generic Taylor center; they do not apply to arbitrary Taylor images.

Proof. If the regular locus is empty, the claims are immediate. Otherwise adjoin a Taylor center t transcendental over E . Put $S = Q_v$, $A_0 = Q_X + vQ_u$, and

$$\mathcal{D} = \partial_X + v\partial_u - (A_0/S)\partial_v,$$

holding z constant. On $S \neq 0$, implicit differentiation reconstructs every derivative through order D . For $j \geq 2$, write this derivative as N_j/S^{2j-3} , with $N_2 = -A_0$. If $e = 2j - 3$, the recurrence is

$$\begin{aligned} N_{j+1} &= S^2(\partial_X N_j + v\partial_u N_j) - A_0 S \partial_v N_j \\ &\quad - e N_j (S(\partial_X S + v\partial_u S) - A_0 \partial_v S). \end{aligned}$$

Induction gives jet degree at most $1 + (2j - 3)(B - 1)$ and challenge degree at most $(2j - 3)H$. The common-denominator Taylor polynomial through degree D is $F(X)/S(t, z, u, v)^\tau$, where every coefficient of F has jet degree at most b and challenge degree at most τH . Division by the Taylor factorials is valid since $p > D$.

The actual locus in initial data. Substitute this polynomial and its derivative into Q and clear $S^{\tau B}$. In a term of jet degree $e \leq B$, the numerator has jet degree at most

$$eb + \tau(B - e)(B - 1) \leq Bb$$

and challenge degree at most $H + \tau BH$. Thus every X -coefficient of the residual has total (z, u, v) degree at most K . Include the initial equation $Q(t, z, u, v) = 0$, of total degree at most q . On $S(t, z, u, v) \neq 0$, these equations describe exactly the initial data of actual degree $\leq D$ solutions. Reconstruction and the initial-jet map are inverse regular maps on these open loci. Consequently all actual regular components have dimension at most two. Their closures acquire no irreducible component supported entirely on the boundary of this open locus.

Curve components. Choose λ, c independently generic after adjoining t , and cut the actual locus by

$$P(t) + \lambda z = c.$$

Each projective curve component has finitely many points at infinity, defined over E . At any such point the coefficient direction and the z direction are not all zero, so $\sum_i a_i t^i + \lambda z \neq 0$. The chosen hyperplane avoids these points. Its generic intersection also avoids the finite separant boundary and the finite intersections of each curve with other actual components, including surfaces.

The function $P(t) + \lambda z$ on each curve is separable. The coefficient functions and z generate the curve's function field, so their differentials cannot all vanish; independence of t, λ prevents cancellation. Its generic fiber is therefore reduced. The section selects exactly J_1 distinct affine points in the regular reconstruction locus, away from the surface components.

In initial data the section substitutes $u = c - \lambda z$ into $Q(t, z, u, v)$, a nonzero plane equation of degree at most q . Each residual restricts to degree at most K . A plane component on which all residuals vanish, unless contained in $S = 0$, maps to a positive-dimensional part of this generic hyperplane section of the actual locus, and hence to a section of a surface component. None of the selected J_1 points lies on it. Nor do they lie on components contained in $S = 0$. On every remaining plane component at least one residual is nonzero. A generic linear combination of the residuals vanishes identically on none of them. Bézout's bound for its intersections with their union gives $J_1 \leq qK$.

Surface components. Use the same generic section. It contains no curve at infinity of an actual surface: a generic form $P(t) + \lambda z$ is nonzero on each such curve. Thus the affine section closures retain the full surface degrees. On the regular initial-data locus the projection to (z, u) is generically separable, because $Q_v \neq 0$. A generic line $u + \lambda z = c$ therefore gives generically reduced sections.

Their initial-data curves are distinct components of the plane equation $Q(t, z, c - \lambda z, v) = 0$ and have summed degree at most q . The projective reconstruction coordinates can be written as

$$S^\tau, \quad zS^\tau, \quad \text{the coefficient numerators of } F.$$

After the substitution, all have degree at most T : the first two have degrees at most $T - 1, T$, and the remaining numerators have degree at most $b + \tau H = T$. Homogenizing to degree T , the degree of the rational image of each source curve is at most T times its degree, as seen by pulling back a generic hyperplane. Summing gives $J_2 \leq qT$.

Finally fix the generic challenge and work over its algebraic closure. At an independent generic Taylor center, regular reconstruction identifies every positive-dimensional actual solution component with a component of the plane equation in (u, v) . Projection to u has total degree at most $\deg_v Q$.

The hyperplane $P(t) = u_0$ avoids all curve points at infinity by the same transcendence argument, and its generic fiber is reduced. It therefore counts the full coefficient-space degree of these curves, proving the generic-fiber assertion. $\square$

Linear degree cannot be replaced by a constant. Let $D \geq 2$, in characteristic zero or $p > D + 1$, and put

$$R(X, z) = X^{D+1} + zX - 1, \quad Q = RY_1 - R_X Y_0 + Y_0^2.$$

Here $B = 2$, $H = 1$. Every nonzero polynomial solution satisfies $(R/P)' = 1$. The rational function $R/P - X$ has degree at most $D + 1$ and derivative zero, so it is constant. Indeed, in positive characteristic its degree is less than p , whereas a nonconstant rational p th power has degree at least p . Hence all nonzero solutions are

$$z = a^{-1} - a^D, \quad P_a(X) = \sum_{i=0}^{D-1} a^i X^{D-i} + a^{-1}, \quad a \neq 0.$$

Conversely these are solutions, since $R(a, z) = 0$ and $P_a = R/(X - a)$. The coefficient of X^{D-1} recovers a . After multiplying the projective coordinates by a , they span all monomials $1, a, \dots, a^{D+1}$, so the curve has degree $D + 1$. The additional zero-solution line gives $J_1 = D + 2$.

Linear surface growth also occurs: $(X - z)P' - DP = 0$ has the actual family $P = c(X - z)^D$. Two generic affine hyperplanes restrict to $A(z) + cB(z) = 0$ and $C(z) + cE(z) = 0$, with A, C of degree at most one and B, E of degree at most D . Their elimination polynomial $AE - CB$ generically has degree $D + 1$ and no extraneous roots. Thus this surface has degree $D + 1$, although its generic challenge fiber is a line.

Corollary H.2 (Nonaffine actual curves contribute linearly). *Fix a received line $f + zg$ on n distinct evaluation points and an agreement threshold $D < A \leq n$. Among the actual curve components in Theorem H.1, omit those that are affine codeword graphs $P = F + zG$, with $\deg F, \deg G \leq D$. The remaining components contain at most*

$$\frac{n - D}{A - D} qK$$

polynomial-challenge pairs with at least A agreements. In particular, their exceptional-label contribution is $O_{B,H,\eta}(n)$ when $A - D \geq \eta n$ and $D \leq n$.

Proof. On an irreducible curve C , call a coordinate persistent if its agreement hyperplane contains C . If more than D coordinates were persistent, interpolation at $D + 1$ of them would force $P = F + zG$ identically on C . A vertical curve would then be a point; otherwise C would be that affine codeword graph. Thus every retained curve has at most $h \leq D$ persistent coordinates.

Each of its other $n - h$ agreement hyperplanes meets C in at most $\deg C$ distinct affine points. Every selected nearby pair lies on at least $A - h$ of these hyperplanes. Its number is therefore at most

$$(\deg C) \frac{n - h}{A - h} \leq (\deg C) \frac{n - D}{A - D}.$$

Sum over the retained curves and use $J_1 \leq qK$. $\square$

Scope and verification of the component bounds. The theorem gives no linear bound on the number of isolated joint solutions; the next construction shows that such a bound is false. Nor does it bound by a constant the number of persistent affine codeword graphs that can each contribute accidental agreements. These obstacles prevent an inference of general $O(n/|F|)$ first-order MCA from the component-degree bounds alone. The checks in `research/first_order_actual_components/` verify the reconstruction recurrence on five equations, 576 finite Taylor fixtures, and 40 sharp-order linear-equation curves. A second checker exhausts 33,849 polynomial–challenge pairs for the nonlinear curve family, verifies 39 split reduced hyperplane sections, and detects extra solutions outside that example’s stricter characteristic hypothesis $p > D + 1$. These checks supplement the geometric proof; they do not certify the generic-section argument by themselves.

H.1 Quadratically many isolated regular first-order solutions

Theorem H.3 (Quadratic isolated count at derivative degree one). *Let $D \geq 1$, in characteristic zero or $p > D + 2$, and let R be a monic squarefree polynomial of degree $N = D + 1$ with nonzero roots. The first-order equation*

$$(z - X^2)(RP' - R'P + P^2) + 2XRP - 2R^2 = 0 \quad (59)$$

has exactly $\binom{D+2}{2}$ polynomial–challenge solutions of degree at most D , all regular. They are indexed by unordered multisets $\{a, b\}$ of roots of R , including $a = b$, and are

$$z = ab, \quad P_{a,b} = \frac{R}{X - a} + \frac{R}{X - b}.$$

Thus its reduced actual joint solution locus consists of exactly that many isolated points. Its jet degree is two, derivative degree one, and challenge degree one. For every D , these points can all have distinct labels over a prime field of size $O(D^3)$.

Proof. Put $U = (X^2 + z)P - 2XR$ and $V = R - XP$. Direct differentiation gives

$$U'V - UV' = (z - X^2)(RP' - R'P + P^2) + 2XRP - 2R^2.$$

Since $R(0) \neq 0$, V is not identically zero. The rational function U/V has degree at most $D + 2$ and derivative zero, so it is constant. In positive characteristic this uses the strict bound $D + 2 < p$: a nonconstant rational p th power has degree at least p . Writing the constant as $-u$ gives

$$(X^2 - uX + z)P = (2X - u)R.$$

For $H = X^2 - uX + z$, this says $P = RH'/H$. Every root of H must belong to R : a root of multiplicity one or two gives a nonzero simple-pole residue in H'/H , since the characteristic is not two. Conversely every such multiset gives the displayed solution. Its partial-fraction residues distinguish the multisets. The separant is $(z - X^2)R$, nonzero as a polynomial in X at every solution, so all solutions are regular.

For distinct prime-field labels, choose the roots to form a multiplicative Sidon set including diagonal products. If k roots have already been chosen and S is their $k(k + 1)/2$ products, a new root x is excluded by at most

$$(k + 2)|S| + k + 1$$

field elements: those with $xa \in S$ for an old root a , those with $x^2 \in S$, the old roots, and zero. Products among the new pairs cannot collide otherwise. For $k \leq N - 1$ this bound is at most $(N^3 + N)/2$. A prime larger than this bound and $D + 2$ therefore suffices, and a prime of size $O(N^3)$ exists. This constructs all the solutions and their pairwise distinct labels in the prime field. $\square$

The solution count by itself therefore cannot prove linear first-order MCA. The following agreement bound also shows why the example is not a quadratic proximity-gap lower bound.

Proposition H.4 (Constant fixed-gap count for the isolated family). *Over a finite field satisfying Theorem H.3, fix a received line $f + zg$ on n distinct coordinates and an agreement threshold $D < A \leq n$. Let L be the number of distinct labels with an A -agreeing candidate from the displayed family. Then*

$$L \leq \frac{(D+2)n}{A-D}.$$

More strongly, let r_0 be the number of domain points that are roots of R , put $m = n - r_0$ and $t = A - r_0$, and suppose $t^3 > 4m^2$. Then

$$L \leq \frac{2m}{t - (4m^2)^{1/3}}.$$

In particular $L = O_\eta(1)$ for $A - D \geq \eta n$ and sufficiently large n . Neither bound requires that all family labels be distinct.

Proof. Every candidate has degree D , hence at least $A - D$ of its agreements have nonzero value. At a root c of R , only the N multisets containing c have nonzero value. At a point x outside the roots, write $\alpha = f(x)/R(x)$ and $\beta = g(x)/R(x)$. For fixed a , agreement is

$$\frac{1}{x-a} + \frac{1}{x-b} = \alpha + \beta ab.$$

Clearing $x - b$ gives a nonzero polynomial of degree at most two in b . If it vanished identically, its quadratic coefficient would give $\beta = 0$ since $a \neq 0$, its linear coefficient would give $\alpha = 1/(x - a)$, and its constant coefficient would still be nonzero. Thus there are at most $2N$ ordered matches. Diagonal matches satisfy $2 = (\alpha + \beta a^2)(x - a)$, a nonzero cubic equation, so there are at most three of them. The number of unordered matches is consequently at most $\lfloor (2N + 3)/2 \rfloor = N + 1$. Counting nonzero agreement incidences proves the first bound, even for the number of nearby pairs.

For the stronger bound, choose one nearby candidate at each of the L distinct labels and put

$$W_{a,b} = P_{a,b}/R = \frac{1}{X-a} + \frac{1}{X-b}.$$

No three of the labeled functions $(ab, W_{a,b})$ at distinct labels are affine-collinear. To prove this, suppose their affine relation has coefficients c_1, c_2, c_3 , all nonzero, with $\sum c_i = \sum c_i z_i = 0$. A root belonging to only one support cannot have its nonzero residue canceled. Each root in the union therefore belongs to at least two supports, so the union has at most three roots. If it has three, the supports must be the three edges of a triangle. Their residue matrix has determinant 2, so they are linearly independent. If it has two roots a, b , the three distinct multisets must be $\{a, a\}, \{a, b\}, \{b, b\}$. Their only affine relation is $W_{a,a} - 2W_{a,b} + W_{b,b} = 0$, but the same coefficients give $a^2 - 2ab + b^2 = (a - b)^2 \neq 0$ on the labels. One root cannot support three distinct candidates. This proves the claim.

For any three distinct labels, take $(c_1, c_2, c_3) = (z_2 - z_3, z_3 - z_1, z_1 - z_2)$. The nonzero rational function $\sum c_i W_i$ has a denominator of degree at most six. Using the product of the three monic quadratic denominators, its numerator has degree at most four: each summand has leading term $2X^5$, which cancels because $\sum c_i = 0$. Thus three candidates can agree with the received line at at most four common coordinates outside the roots of R .

Each selected candidate has at least t agreements on these m coordinates. If h_x counts the agreeing candidates at x , then $\sum h_x \geq tL$ and

$$\sum_x \binom{h_x}{3} \leq 4 \binom{L}{3} \leq \frac{4L^3}{6}.$$

For every nonnegative integer h , $6\binom{h}{3} \geq (h-2)_+^3$. Convexity therefore gives

$$m \left(\frac{tL}{m} - 2 \right)_+^3 \leq 4L^3.$$

Taking cube roots and rearranging proves the bound (the case $tL \leq 2m$ also satisfies it). Finally $r_0 \leq D+1$, so $t \geq \eta n - 1$ and $m \leq n$ at a fixed gap. The denominator is positive and of order n for sufficiently large n . $\square$

Exact checks. The isolated-family checker exhausts 17,410 polynomials, covering 180,384 polynomial–challenge pairs, and checks every pair of local received values at each coordinate of four small-field fixtures. It verifies the rational derivative identity and constructs prime-field Sidon examples through $D = 64$ with 2,145 distinct regular isolated labels. It also checks the nonzero degree-four triple numerators and the integer convexity inequality. A characteristic three negative control has an extra solution outside the theorem’s strict characteristic hypothesis. These finite checks supplement the classification and incidence proofs above.

H.2 A positive density of nonsingular agreements gives linear MCA

The preceding component counts do not by themselves bound all bad labels linearly. Conditioning Taylor reconstruction at an agreeing coordinate gives a different sufficient condition that also covers surface components. It isolates singular agreement coordinates as the remaining obstruction to this argument.

Theorem H.5 (Nonsingular-agreement bound). *Let F have characteristic zero or $p > D$, with $D \geq 1$, and let $Q(X, z, u, v) \neq 0$ have jet degree at most $B \geq 1$ and challenge degree at most H . Its X -degree is unrestricted. Put*

$$\tau = \max\{0, 2D - 3\}, \quad b = 1 + \tau(B - 1), \quad T_0 = b + \tau H, \quad K = H + BT_0, \quad q = B + H.$$

Fix a received line $f + zg$ on n distinct coordinates, an agreement threshold $D < A \leq n$, and an integer $L \geq 1$. Consider labels admitting a degree- $\leq D$ polynomial solution $Q(X, z, P, P') = 0$ whose full agreement support S satisfies $|S| \geq A$, has no degree- $\leq D$ polynomial agreeing with g on all of S , and contains at least L coordinates x with

$$Q_v(x, z, P(x), P'(x)) \neq 0.$$

The number of these labels is at most

$$\frac{n}{L} \left(qK + qT_0 \frac{n-D}{A-D} + qn \right). \tag{60}$$

In particular it is $O_{B,H,\eta,\lambda}(n)$ when $A - D \geq \eta n$ and $L \geq \lambda n$, with $\eta, \lambda > 0$ fixed.

Proof. Fix a domain point x and substitute $u = f_x + zg_x$ into the initial equation. The resulting polynomial in (z, v) has degree at most q . If it vanishes identically, so does its v derivative, and there are no nonsingular initial data to count at this coordinate. Otherwise work on the open locus $S_0 = Q_v(x, z, f_x + zg_x, v) \neq 0$ of this plane curve.

The Taylor reconstruction in Theorem H.1 applies at the fixed center x on this open locus. Its common denominator is S_0^τ , of degree at most $T_0 - 1$; its coefficient numerators and zS_0^τ have degree at most T_0 . Thus its projective coordinate degree is at most T_0 . Exact solutions are selected by residual equations of degree at most K . Reconstruction and the initial-jet map are inverse regular maps on this open exact-solution locus.

The isolated points of the residual locus number at most qK . Indeed, on every component of the initial plane curve not wholly satisfying the residuals, some residual is nonzero. A generic linear combination over an algebraic closure is nonzero on each such component; Bézout gives the bound. Components contained in $S_0 = 0$ are omitted, and points on a retained curve are accounted for with that curve.

The retained source curves are components of the initial plane equation, so there are at most q and their degrees sum to at most q . Their reconstructed image curves have total degree at most qT_0 , by pulling back a generic hyperplane under the rational map of degree at most T_0 . Initial-jet recovery ensures none of these images collapses to a point.

For every nonaffine image curve, the incidence argument of Corollary H.2 gives at most its degree times $(n - D)/(A - D)$ nearby pairs. That argument applies to these curves whether or not they are components of the full joint locus. The total contribution is at most $qT_0(n - D)/(A - D)$.

Each remaining image curve is a graph $P = F_0 + zG_0$ and has at most n bad labels. A support consisting only of persistent agreements has the common witnesses F_0, G_0 ; a bad support therefore contains an accidental agreement at a nonpersistent coordinate, which forces one label. If the graph does not descend to F , it has at most one F -rational point at distinct challenge labels: two such points would recover F_0, G_0 over F . This also satisfies the bound n . There are at most q graphs.

Hence this fixed coordinate occurs nonsingularly in a bad nearby witness for at most $qK + qT_0(n - D)/(A - D) + qn$ labels. Select one qualifying witness per label. Each supplies at least L such coordinate-label incidences. Summing over all n coordinates and dividing by L proves (60). $\square$

Corollary H.6 (Value-independent separant). *Suppose*

$$Q(X, z, u, v) = R(X, z)v - \mathcal{A}(X, z, u), \quad R \neq 0,$$

under the same degree and characteristic hypotheses, and put $r = \deg_X R$. If $A > \max\{D, r\}$, all its full-support bad labels number at most

$$H + \frac{n}{A - r} \left(qK + qT_0 \frac{n - D}{A - D} + qn \right). \quad (61)$$

This is $O(n)$ for fixed B, H and fixed positive lower bounds on $(A - D)/n$ and $(A - r)/n$. The value degree of $\mathcal{A}$ is arbitrary. At each fixed label with $R(X, z) \not\equiv 0$, its ordinary list size at threshold A is at most $\lfloor (n - r)/(A - r) \rfloor$.

Proof. There are at most H labels with $R(X, z) \equiv 0$, using any one nonzero X -coefficient of R . At all others, at most r domain points are zeros of R , so every nearby candidate has at least $A - r$ nonsingular agreements. Apply the theorem with $L = A - r$.

For the ordinary list bound, two distinct solutions at the same label have a nonzero difference Δ satisfying $R\Delta' = C(X)\Delta$. At a zero of Δ where R is nonzero, the multiplicity $1 \leq m \leq D < p$ makes the left side have order $m - 1$ and the right side order at least m , a contradiction. Thus a

value at a nonsingular coordinate determines at most one solution. If $r_z \leq r$ domain points are singular, nearby candidates have disjoint regular agreement supports of size at least $A - r_z$. Their number is at most $(n - r_z)/(A - r_z) \leq (n - r)/(A - r)$. $\square$

Sharp linear order in a nonlinear subclass. For a nonzero polynomial W , the equation

$$WP' - W'P - P(P - W) = 0$$

has exactly the polynomial solutions $0, W$. To see this, write $Z = P/W$, giving $Z' = Z(Z - 1)$. If $Z \neq 0$, then $U = 1 - 1/Z$ satisfies $U' = U$. In characteristic zero, a nonzero rational logarithmic derivative U'/U vanishes at infinity and cannot equal one. In characteristic p , the p -fold ordinary derivative is zero on $F(X)$, whereas $U' = U$ would make it equal U . Hence $U = 0$.

Let $n = 12s$, $D = 3s$, $A = 6s$, and choose a prime $p > 12s$. On domain $\{0, \dots, n - 1\}$ let W be the locator of the first $3s$ points. Put $f = g = 0$ on these roots and on a second block of $3s$ points; put $f = W, g = 0$ on a third block of $3s$. At the final $3s$ coordinates, indexed by $j = 1, \dots, 3s$, put $f = 3jW, g = W$. Each candidate cW , $c \in \{0, 1\}$, has $6s$ persistent agreements and acquires one more at each label $z = c - 3j$. These $6s = n/2$ labels are distinct modulo p . On each resulting support, g has $6s > D$ zeros and one nonzero value, so it has no degree- $\leq D$ explanation. Every other label has only persistent supports and is good. Thus exactly $n/2$ labels are bad, with jet degree two, challenge degree zero, code rate tending to $1/4$, and agreement gap $1/4$.

Scope and checks. The theorem does not assert that interpolation equations have many nonsingular agreement coordinates. At fixed gap its bound is $O(n^2/L)$; a superlinear example must concentrate on witnesses with a vanishing fraction of such coordinates. The general first-order conjecture remains open. The checks in `research/quasilinear_first_order/` enumerate small-field polynomial solutions, verify regular-coordinate uniqueness and list incidences, include implicit equations with actual solution surfaces, and check the exact $n/2$ lower construction at five lengths. These finite checks supplement the reconstruction and incidence proof; no independent review or literature-priority claim is made.

I Sharp isolated-solution growth at fixed derivative order

The quadratic isolated family of Appendix H extends to every fixed derivative order. Here we give a matching general upper bound and an explicit family attaining its exponent. These are bounds on the actual polynomial-solution locus, before any agreement filter. The attaining family still has a constant nearby-label count at a fixed positive gap, so this sharpness does not establish sharp MCA error.

Theorem I.1 (Isolated regular solution upper bound). *Let $d \geq 1$, $D \geq d$, and let E be algebraically closed of characteristic zero or $p > D$. Let $Q(X, z, u_0, \dots, u_d)$ have total jet degree at most $B \geq 1$ and challenge degree at most H , with unrestricted X -degree. Let J_0 count the isolated points of the reduced closure of actual regular pairs*

$$Q(X, z, P, P', \dots, P^{(d)}) = 0, \quad \deg P \leq D, \quad Q_{u_d}(X, z, P, \dots, P^{(d)}) \neq 0.$$

Put

$$\tau = \max\{0, 2(D - d) - 1\}, \quad b = 1 + \tau(B - 1), \quad K = H + B(b + \tau H), \quad q = B + H.$$

Then $J_0 \leq qK^{d+1} = O_{d,B,H}(D^{d+1})$.

Proof. If $J_0 = 0$ there is nothing to prove. Every isolated point of this closure belongs to the actual regular locus, since that locus is dense. Choose a Taylor center t at which the separant is nonzero at every isolated regular solution; there are finitely many such points and only finitely many excluded centers for each. Write

$$S = Q_{u_d}, \quad A_0 = Q_X + \sum_{i < d} u_{i+1} Q_{u_i}, \quad \mathcal{D} = \partial_X + \sum_{i < d} u_{i+1} \partial_{u_i} - (A_0/S) \partial_{u_d}.$$

The quotient-rule recurrence in the proof of Theorem H.1, with this derivation, gives for the j th derivative, $j > d$, denominator $S^{2(j-d)-1}$, jet numerator degree at most $1 + (2(j-d) - 1)(B-1)$, and challenge degree at most $(2(j-d) - 1)H$. The initial derivatives $u_0, \dots, u_d$ need no denominator. Thus the degree- D Taylor reconstruction, with common denominator S^τ , has coefficient numerators of jet degree at most b and challenge degree at most τH . Taylor factorials are invertible.

Substitute this polynomial and its first d derivatives into Q , clear $S^{\tau B}$, and take coefficients in X . Every residual has total initial-data degree at most K : a term of jet degree $e \leq B$ has numerator jet degree at most $eb + \tau(B-e)(B-1) \leq Bb$ and challenge degree at most $H + \tau BH$. Include the initial equation $Q(t, z, u_0, \dots, u_d) = 0$, of degree at most q . On $S(t, z, u_0, \dots, u_d) \neq 0$, these equations describe precisely the initial jets of actual degree- $\leq D$ solutions. The initial-jet and reconstruction maps are inverse regular maps there, so they preserve the isolated points being counted.

The initial hypersurface has dimension at most $d+1$ in $d+2$ variables. Cut successively by $d+1$ generic linear combinations of the residuals, each of degree at most K . At each stage, discard positive-dimensional components on which every residual vanishes; none can contain a target isolated point in the regular open locus. On each remaining component through a target point, some residual is nonzero, and a generic combination cuts its dimension. Choose the combinations simultaneously for the finitely many components at each step. This process retains every target isolated point. Bézout bounds the sum of the degrees of the retained intersections by qK^{d+1} . This argument does not require that the whole unpruned intersection be zero-dimensional: other positive-dimensional components may remain away from the target points. $\square$

Theorem I.2 (Attaining every isolated-count exponent). *Fix $r = d+1 \geq 2$ and $D \geq 1$, in characteristic zero or $p > D+r$. Let R be monic and squarefree of degree $N = D+1$, with nonzero roots. Put*

$$\phi_j = X^j \ (1 \leq j < r), \quad \phi_r = X^r + z, \quad F_j = R\phi_j' - P\phi_j,$$

and let Q_r be the Wronskian of $F_1, \dots, F_r$, with derivative rows $0, \dots, r-1$ and z held constant. Then Q_r is a nonzero differential polynomial of order d , jet degree r , challenge degree one, and degree one in its highest derivative. Its degree- $\leq D$ polynomial solutions are exactly

$$P = R \frac{H'}{H}, \quad z = H(0),$$

where H ranges over monic degree- r polynomials whose root multisets lie among the roots of R . All are regular, and there are exactly

$$\binom{D+r}{r} = \Theta_d(D^{d+1})$$

isolated polynomial-challenge pairs, with no positive-dimensional solution component. For fixed d , all labels can be distinct over prime fields of size $O_d(D^{2d+1})$.

Proof. The separant has the explicit value

$$\frac{\partial Q_r}{\partial P^{(r-1)}} = (-1)^r C_r R^{r-1} (X^r + (-1)^{r-1} z), \quad C_r = \prod_{j=1}^{r-1} j!. \quad (62)$$

Indeed, the highest jet occurs only in the last derivative row, with coefficient row $-(\phi_1, \dots, \phi_r)$. Move that row to the top. In the other rows the successive highest derivatives of the ϕ_j have coefficient R , giving $(-1)^r R^{r-1} \text{Wr}(\phi_1, \dots, \phi_r)$. The monomial Vandermonde formula gives

$$\text{Wr}(\phi_1, \dots, \phi_r) = C_r (X^r + (-1)^{r-1} z).$$

This is nonzero for every constant z . The pure P^r coefficient is $(-1)^r \text{Wr}(\phi_1, \dots, \phi_r)$, so the jet degree is exactly r . All claimed degree bounds follow; the X -degree is also $O_r(D)$, although the upper bound does not need it.

Each specialized F_j has degree at most $D+r < p$. In this range a polynomial Wronskian vanishes exactly when its columns are linearly dependent over constants. To see the nontrivial direction, echelon any independent polynomial basis to distinct leading degrees. The leading Wronskian coefficient is their nonzero leading-coefficient product times the Vandermonde of these distinct degrees, all less than p . The same argument works in characteristic zero.

The first $r-1$ columns are independent for every polynomial P . Otherwise $RH' - PH = 0$ for a nonzero $H \in \text{span}(X, \dots, X^{r-1})$. Its root at zero has multiplicity between one and $r-1$, so H'/H has a nonzero simple pole there, contradicting $R(0) \neq 0$ and polynomial P . Consequently $Q_r = 0$ if and only if a unique monic polynomial

$$H = X^r + \sum_{j=1}^{r-1} h_j X^j + z$$

satisfies $RH' - PH = 0$. Its root multiplicities are at most $r < p$; polynomiality forces every root to belong to R , and conversely every such root multiset works. Partial-fraction residues distinguish all the multisets. The leading coefficient of P is r , so its degree is exactly D . Equation (62) proves regularity, and the classification proves finiteness and the count.

To separate the labels, greedily choose N nonzero roots whose products of r factors, with repetitions allowed, are all different. Suppose $k \geq 1$ roots have been chosen. Products of any fixed smaller number of factors are also distinct, by padding with one old root. A new product has form $x^j c$, with $1 \leq j \leq r$ and c an old product of $r-j$ factors. Its collision with an old r -fold product excludes at most j choices of x . Collisions between $x^i c$ and $x^j c'$ for $i < j$ exclude at most $j-i$ nonzero choices. Equal-power collisions are already excluded. Writing $M_s(k) = \binom{k+s-1}{s}$, including $M_0(k) = 1$, a bound on all forbidden choices, including old roots and zero, is

$$k+1 + \sum_{j=1}^r j M_{r-j}(k) M_r(k) + \sum_{1 \leq i < j \leq r} (j-i) M_{r-i}(k) M_{r-j}(k) = O_r(k^{2r-1}).$$

Start from one nonzero root and choose a prime larger than this bound for every $k < N$ and larger than $D+r$. A prime of size $O_r(N^{2r-1})$ suffices. All resulting H , P , and distinct labels $H(0) = (-1)^r \prod a$ lie in that prime field. $\square$

Proposition I.3 (The attaining family still has constant nearby count). *For the family in Theorem I.2, let a received line have L distinct labels with candidates agreeing on at least A of n*

coordinates, where $1 \leq A \leq n$. Let r_0 count domain roots of R , and put $m = n - r_0$, $t = A - r_0$, and $C = 3r - 2$. If $t^3 > 8Cm^2$, then

$$L \leq \max \left\{ \frac{4m}{t}, \frac{8m^3(r-1)}{t^3 - 8Cm^2} \right\}.$$

In particular $L = O_{r,\eta}(1)$ at fixed $A - D \geq \eta n$ and sufficiently large n . This conclusion remains valid for any assignment of one label to each candidate, counting distinct labels.

Proof. Choose one candidate at each selected label, and normalize it to $W = P/R = H'/H$. Its residues at the roots of R belong to $\{0, \dots, r\}$, and their sum is r . Distinct candidates have different residue vectors. On an affine line of labeled functions (z, W) through two of them at distinct labels, choose a residue coordinate where their vectors differ. That residue is an injective affine function of z along the line and has only $r + 1$ allowed values. Thus every such line contains at most $r + 1$ candidates. There are at most $(r - 1)\binom{L}{2}/3$ collinear triples.

A noncollinear triple gives a nonzero rational relation $\sum_i c_i W_i$, with $\sum_i c_i = \sum_i c_i z_i = 0$. On the product of its three degree- r denominators, the numerator has degree at most $3r - 2$: each W_i starts with r/X at infinity, so the highest coefficient cancels. Such a triple has at most $C = 3r - 2$ common agreeing coordinates outside the roots of R . A collinear triple has at most m .

Let h_x count agreements at each outside coordinate. Each selected candidate contributes at least t agreements. The same convexity estimate as in Proposition H.4 gives

$$m \left(\frac{tL}{m} - 2 \right)_+^3 \leq 6 \sum_x \binom{h_x}{3} \leq CL^3 + m(r - 1)L^2.$$

If $L \geq 4m/t$, the left side is at least $t^3 L^3 / (8m^2)$; rearranging proves the second displayed alternative. Otherwise the first alternative holds. Finally $r_0 \leq D + 1$ gives $t \geq \eta n - 1$ and $m \leq n$, proving the constant fixed-gap bound. $\square$

Collinear triples cannot simply be omitted. For $r = 3$, take two roots $a = 1$, $b = -2$ and $H_j = (X - a)^{3-j}(X - b)^j$ for $j = 0, 1, 3$. Their logarithmic derivatives are affine in j , and their labels $-1, 2, 8$ are also affine at these three indices. Over sufficiently large characteristic this is a distinct-label collinear triple. The extra L^2 term in the proof handles such configurations.

Verification and limits. The package `research/isolated_solution_sharpness/` checks the separant through order five, exhausts 760 polynomials covering 5,270 polynomial–challenge pairs at orders one through three, and includes characteristic-three and characteristic-five negative controls. Prime-field examples have 969 distinct isolated labels at order two, 495 at order three, and 462 at order four. A separate checker verifies the reconstruction recurrence on four equations of orders two through five and 360 finite Taylor fixtures. The checks support the written proofs; they do not replace the generic intersection argument. No fixed-gap superlinear MCA lower bound follows: the polynomially many isolated solutions in the sharpness family fail the required nearby-label condition.

I.1 A rational envelope has only linearly many bad labels

The agreement obstruction above extends beyond a fixed residue alphabet. The following elementary incidence bound allows arbitrarily many candidates and arbitrary characteristics. It gives another reason why a large algebraic solution count need not make the proximity bound sharp.

Proposition I.4 (Rational-envelope bound). *Let $\mathcal{D}$ consist of n distinct points of a field F , and fix polynomials $S, R \in F[X]$ with $R \neq 0$. Consider any family of polynomials P of degree at most D satisfying*

$$\frac{P - S}{R} = \frac{a_P}{b_P}, \quad \deg a_P, \deg b_P \leq r, \quad b_P \neq 0.$$

Let h be the number of zeros of R on $\mathcal{D}$. For an integer agreement threshold $1 \leq A \leq n$, suppose

$$A - h \geq \beta n > 0, \quad n \geq 48r/\beta^3.$$

On every received line $f + zg$, the number of full-support MCA-bad labels witnessed by this family is at most

$$\frac{16(n - A + 1)}{\beta^3}.$$

Here badness means that g has no degree-at-most- D interpolant on the candidate's entire agreement set. There is no characteristic hypothesis.

Proof. First, any polynomial pencil $U + zV$, with $\deg U, \deg V \leq D$, has at most $M = n - A + 1$ bad labels. Let c count its persistent agreements, where $(f, g) = (U, V)$. Each other coordinate agrees at at most one label. If $c < A$, the number of nearby labels is at most $\lfloor (n - c)/(A - c) \rfloor \leq n - A + 1$. If $c \geq A$, a bad label must have a nonpersistent agreement, giving at most $n - c \leq n - A$ labels.

Choose one bad candidate P_z at each of L distinct labels and put $W_z = (P_z - S)/R$. Delete the h zeros of R and normalize the received line by subtracting S and dividing by R . Each candidate retains at least βn agreements. Its reduced denominator is nonzero at every retained coordinate, since $W_z = (P_z - S)/R$ is regular there.

A noncollinear triple of graph points (z, W_z) over $F(X)$ has a nonzero affine-incidence determinant. Clearing its three denominators gives a polynomial of degree at most $3r$, so the triple has at most $3r$ common retained agreements. A graph line through two selected points corresponds, on multiplying by R and adding S , to a polynomial pencil of degree at most D : interpolate the two actual polynomials at their distinct scalar labels. It therefore contains at most M selected bad labels.

Let b_x count retained agreements at coordinate x , and set $b_x = 0$ at deleted coordinates. If $L \geq 4/\beta$, convexity gives

$$\sum_x b_x(b_x - 1)(b_x - 2) \geq n(\beta L - 2)^3 \geq \beta^3 n L^3 / 8.$$

Noncollinear ordered triples contribute at most $3rL^3$. Each ordered pair determines one graph line with at most M choices of a third selected label; these triples contribute at most nML^2 . Consequently

$$\beta^3 n L^3 / 8 \leq 3rL^3 + nML^2.$$

Since $3r \leq \beta^3 n / 16$, this implies $L \leq 16M/\beta^3$. The case $L < 4/\beta$ satisfies the same bound. Over an infinite field apply the argument to any finite selection of bad labels. $\square$

For example, if $\deg R \leq D + c$ and $A - D \geq \eta n$, one can take $\beta = \eta - c/n$ when positive. For fixed $c, r, \eta > 0$ the bound is therefore $O_\eta(n)$. It applies to $P = RH'/H$ with $\deg H \leq r$ without a restriction on the residues, and even permits $r \leq \beta^3 n / 48$. A bounded union of such families still has a linear bound. The hypothesis does not hold for arbitrary positive-rate candidates merely by setting $R = 1$, since then the necessary residual degree r is generally too large.

The checker `research/prime_field_tightness/check_rational_envelope.py` verifies 220 rational triple determinants over $\mathbb{F}_{101}$ and all full agreement supports in a pencil attaining $n - A + 1 = 13$ bad labels. These checks supplement the incidence proof.

I.2 Many nearby labels force a rational path to be affine

There is a stronger restriction when the candidate itself is a rational function of the challenge of bounded degree. No bound on the coefficient heights in X is needed.

Theorem I.5 (Rational-path rigidity). *Let $D < A \leq n$ be integers, and let $f + Zg$ be a received line on n distinct points of a field F . Suppose $T(X, Z) \in F(X)(Z)$ has numerator and denominator degrees in Z at most $h \geq 1$. Either*

$$T(X, Z) = U(X) + ZV(X), \quad \deg U, \deg V \leq D,$$

or the number of scalar labels z at which $T(X, z)$ is defined, is a polynomial of degree at most D , and has at least A agreements with $f + zg$ is at most

$$\frac{(8h + 10)n}{A - D}.$$

In the affine case there are at most $n - A + 1$ full-support MCA-bad labels. These statements hold in every characteristic.

Proof. Write $T = a/b$ with coprime $a, b \in F[X, Z]$, $b \neq 0$, and both Z -degrees at most h . For every polynomial specialization P_z under consideration, $a(X, z) = b(X, z)P_z(X)$ as polynomials. Put

$$F_x(Z) = a(x, Z) - b(x, Z)(f(x) + Zg(x)).$$

Call a coordinate persistent if $F_x \equiv 0$, and let H count them. At a persistent coordinate $b(x, Z) \not\equiv 0$, since otherwise $X - x$ would divide both a and b . Thus all but at most h selected labels agree there. At any other coordinate agreement requires a root of the nonzero polynomial F_x , of degree at most $h + 1$. This argument includes points where specialization makes both $a(x, z)$ and $b(x, z)$ zero.

Put $\delta = A - D$. If $H \leq D + \delta/2$, each nearby label needs at least $\delta/2$ nonpersistent agreements. Their total incidence is at most $(h + 1)n$, giving at most $2(h + 1)n/\delta$ nearby labels.

Now suppose $H > D + \delta/2$, and choose L nearby labels. If T is affine in Z over $F(X)$, two polynomial values give the stated polynomial affine form; the case $L \leq 1$ already satisfies the numerical bound. Otherwise any graph line over $F(X)$ meets the graph of T in at most $h + 1$ labels: the numerator of $T - U - ZV$ is nonzero and has Z -degree at most $h + 1$.

The selected graph points (z, P_z) therefore have at most $(h - 1)L(L - 1)$ ordered collinear triples. A noncollinear triple has a nonzero affine-incidence determinant, a polynomial in X of degree at most D , so it shares at most D persistent agreeing coordinates. Every persistent coordinate has at least $L - h$ agreements. For $L \geq h + 2$, counting ordered triples gives

$$H(L - h - 2)^3 \leq DL^3 + H(h - 1)L^2.$$

Using $(L - h - 2)^3 \geq L^3 - 3(h + 2)L^2$ yields

$$(H - D)L \leq (4h + 5)H, \quad L \leq \frac{(4h + 5)H}{H - D} < \frac{(8h + 10)n}{\delta}.$$

The smaller- L case satisfies the same bound. In the affine alternative, the sharp polynomial-pencil bound from the proof of Proposition I.4 gives $n - A + 1$ bad labels. $\square$

Corollary I.6 (Degree required by a rational witness formula). *Suppose a rational function $T(X, Z)$ selects a degree-at-most- D nearby witness at each of $L > n - A + 1$ distinct full-support MCA-bad labels, with agreement threshold $A > D$. If its numerator and denominator degrees in Z are at most h , then*

$$h \geq \frac{(A - D)L}{8n} - \frac{5}{4}.$$

Proof. The affine alternative can cover at most $n - A + 1$ of the selected bad labels. Rearrange the other alternative of Theorem I.5. A degree-zero formula is a constant polynomial pencil whenever it has a polynomial value, so it cannot cover the hypothesized set either. $\square$

For the complete-coverage examples, $L = p - 1$ and $A - D = \eta n + 1$. Any rational formula selecting witnesses throughout the nearby set consequently has challenge degree $\Omega(\eta p)$; when $\eta = \Theta(1/\log p)$ this is $\Omega(p/\log p)$. This quantifies a failure of low-degree witness coherence. It is a degree bound, not a computational lower bound: efficient witness recovery may use branching rather than one rational formula.

The reciprocal-gap dependence in the rigidity theorem is necessary up to constants for fixed h . The construction and proof in `research/prime_field_tightness/RATIONAL_PATH_RIGIDITY.md` give a non-affine projective path with order $1/\eta$ nearby bad labels at fixed rate. The exact checker considers lengths 160, 320, 640, all at rate $1/4$ and capacity gap $1/8$: the path has respectively 41, 81, 161 polynomial members, but exactly ten nearby labels, all full-support bad. This count is for the specified path, not for all codewords on the line.

J Positive-dimensional actual components save one degree power

Appendix I proves that isolated actual solutions can require a D^{d+1} count. Positive-dimensional actual components satisfy a smaller bound, extending Appendix H to every fixed derivative order. We retain the projective degree in coefficient space throughout.

Theorem J.1 (Degrees by actual component dimension). *Use the hypotheses and notation of Theorem I.1, and additionally put $T = b + \tau H$. Let J_j be the sum of the projective degrees of the j -dimensional irreducible components of the reduced closure of the actual regular joint solution locus. All components have dimension at most $d + 1$, and*

$$J_0 \leq qK^{d+1}, \quad J_j \leq qK^{d+1-j}T^{j-1} \quad (1 \leq j \leq d+1).$$

Thus every positive-dimensional stratum has total degree $O_{d,B,H}(D^d)$. For an ordinary equation with the challenge fixed, let I_j count the degrees of its actual regular components of dimension j . Then

$$I_0 \leq B(Bb)^d, \quad I_j \leq B(Bb)^{d-j}b^{j-1} \quad (1 \leq j \leq d).$$

In particular the positive-dimensional part of a reduced generic challenge fiber has degree $O_{d,B}(D^{d-1})$. This last assertion excludes the isolated points of that fiber.

Proof. The initial-jet description and Taylor reconstruction are as in Theorem I.1. They give dimension at most $d + 1$ and the stated J_0 bound. We explain why a special family of linear sections counts the full actual coefficient-space degree.

Evaluation projections count actual degrees. Let C be an actual j -dimensional component, defined before any Taylor center is adjoined. Choose independent generic parameters t_i, λ_i for $0 \leq i < j$, and use the affine linear forms

$$\ell_i(P, z) = P(t_i) + \lambda_i z.$$

At a projective point at infinity, the coefficient direction of P and the z direction cannot all be zero. Consequently $P(t) + \lambda z$ is not the zero polynomial in (t, λ) at that point. The incidence of a point of C at infinity and j vanishing forms has dimension at most $(j - 1) + j = 2j - 1$, less than the

2j-dimensional parameter space. Thus the projective projection center misses the closure of C . The resulting map to $\mathbb{P}^j$ is finite: its hyperplane pullback is the ample hyperplane bundle on C . Its degree is $\deg C$.

This map is generically separable. The coefficient functions of P and z generate the function field of C , so their differentials span its j -dimensional cotangent space. The evaluation forms span these linear functions; hence the differentials of j independent generic evaluation forms are independent. Their generic affine fiber is therefore reduced and consists of exactly $\deg C$ points. The same choice works simultaneously for the finitely many actual j -components. Generic target values $\ell_i = c_i$ avoid their singular loci, intersections with other actual components, and the separant boundary. The latter is proper after a generic t_0 is chosen, since each actual component has a dense regular locus. These sections therefore select exactly J_j distinct regular points.

Count the section in initial data. Use t_0 as Taylor center. The first section is $u_0 = c_0 - \lambda_0 z$, leaving $d + 1$ initial-data variables. The initial equation has degree at most q . The coefficient residuals have degree at most K . Each of the other $j - 1$ sections, pulled back through reconstruction and cleared by S^τ , has degree at most T : its coordinates are coefficient numerators of degree at most T , zS^τ of degree at most T , and S^τ of degree at most $T - 1$.

These further sections cut the initial hypersurface properly on $S \neq 0$, even before the residuals are imposed. Indeed the Taylor map there has a linear inverse given by its derivatives through order d and z , so it is an immersion. Evaluation hyperplanes span all its coefficient coordinates and cut each positive-dimensional image properly for generic parameters and target values. After the $j - 1$ further sections, the local source dimension at each selected point is at most $d + 1 - j$.

Now impose $d + 1 - j$ generic linear combinations of the residuals. Their full common zero set already isolates every selected point: those points lie on no higher-dimensional actual component and the actual j -components have just been cut to points. The pruning argument of Theorem I.1 therefore retains and isolates all selected points. Bézout bounds their number by $qT^{j-1}K^{d+1-j}$, as claimed. Components on $S = 0$ contribute no selected point and need not be interpreted as actual solutions.

Fixed challenge. Work over an algebraic closure of the coefficient field with that challenge held constant. There are $d + 1$ initial variables, initial equation degree at most B , coefficient-numerator degree at most b , and residual degree at most Bb . Use the evaluation forms $P(t_i)$ without a z coordinate. The same projection argument applies: a nonzero coefficient direction cannot vanish at every t . The first section is $u_0 = c_0$. The identical count gives $B(Bb)^{d-j}b^{j-1}$ for $j \geq 1$ and $B(Bb)^d$ for isolated points. Apply this in particular over the algebraic closure of the generic challenge field. $\square$

Corollary J.2 (Nonaffine actual curves at every order). *For a received line on n distinct coordinates and $D < A \leq n$, the actual curve components in Theorem J.1, other than affine codeword graphs $P = F + zG$, contain at most*

$$\frac{n - D}{A - D} qK^d$$

pairs with at least A agreements. This is $O(n^d)$ at fixed jet and challenge degrees, derivative order, and positive gap.

Proof. The agreement-hyperplane count in Corollary H.2 uses only actual curve degrees. Each retained curve has at most D persistent agreement coordinates, and its nearby-pair count is at most its degree times $(n - D)/(A - D)$. Sum and apply $J_1 \leq qK^d$. $\square$

Theorem J.3 (Sharp positive-dimensional and fiber exponents). *For every $r \geq 1$ and $e \geq 1$, put $D = re$ and assume characteristic zero or $p > D + r - 1$. There is an equation of derivative order r , jet degree r , challenge degree one, and X -degree bounded in terms of r alone, whose reduced actual regular solution closure is irreducible of dimension $r + 1$ and degree at least e^r . Its reduced challenge fibers have degree e^{r-1} . Thus the exponents D^d and D^{d-1} in Theorem J.1 cannot be lowered uniformly, even with a single actual component.*

Proof. Set $\phi_j = X^j$ for $1 \leq j < r$ and $\phi_r = X^r + z$, and define

$$Q = \text{Wr}(eP\phi'_1 - P'\phi_1, \dots, eP\phi'_r - P'\phi_r).$$

This is homogeneous of jet degree r , and its separant is

$$Q_{P(r)} = (-1)^r C_r (eP)^{r-1} (X^r + (-1)^{r-1} z), \quad C_r = \prod_{j=1}^{r-1} j!.$$

The proof is the same row calculation as Equation (62). Its coefficient of P^r is $e^r \text{Wr}(\phi'_1, \dots, \phi'_r)$, which is nonzero; hence the jet degree is exactly r . The separant proves exact order r , linear highest-derivative degree, and challenge degree one. Every nonzero polynomial solution is regular.

Each specialized Wronskian column has degree at most $D + r - 1 < p$. The polynomial Wronskian criterion from Theorem I.2 therefore says that $Q = 0$ precisely when some nonzero

$$H \in \text{span}(X, \dots, X^{r-1}, X^r + z)$$

satisfies $ePH' - P'H = 0$. For nonzero P , this makes $(P/H^e)' = 0$. The rational degree of this quotient is at most $D < p$, so $P = cH^e$ for a nonzero constant c . Conversely every such power satisfies the equation. The zero polynomial also satisfies it; for $r > 1$ it is singular and belongs to the closure of the nonzero regular locus.

Write $H = \sum_{i=0}^r h_i X^i$ projectively. The condition on H is exactly $h_0 = zh_r$. The power map

$$\mathbb{P}^r \longrightarrow \mathbb{P}^D, \quad [H] \longmapsto [H^e]$$

is base-point-free and generically one-to-one, with hyperplane pullback $\mathcal{O}(e)$. It is birational onto its image: on the monic degree- r chart its coefficients are recovered successively from the coefficients of H^e , dividing by e , which is invertible. Its image consequently has degree e^r .

The projective subfamily $h_0 = zh_r$ is closed over the affine z -line. Its power-map image is closed by properness, and its relative affine cone is exactly the actual solution closure just classified. The family of hyperplanes $h_0 = zh_r$ is a projective bundle, so that cone is irreducible of dimension $r + 1$. On the dense monic chart, z is recovered from the constant coefficient of H . Projection forgetting z is thus birational onto the affine cone of the whole power-map image, of degree e^r . Linear projection cannot increase degree, giving the claimed joint-degree lower bound. For each fixed z , the hyperplane of H is a $\mathbb{P}^{r-1}$; its power-map image, and hence its affine cone, has degree e^{r-1} . $\square$

Proposition J.4 (Sharp isolated fibers and affine-graph counts). *Fix $d \geq 1$ and $D \geq d$, in characteristic zero or $p > D + d$. There is a challenge-independent equation of order d and jet degree $d + 1$ with exactly*

$$\binom{D + d + 1}{d} = \Theta_d(D^d)$$

polynomial solutions of degree at most D , all regular and isolated. Adjoining a free challenge turns these into that many actual affine codeword-graph curve components of degree one. Thus both the I_0 and J_1 exponents of Theorem J.1 are sharp as algebraic counts.

Proof. Let R be monic and squarefree of degree $N = D + 1$, and take the Wronskian of

$$R(X^j)' - PX^j, \quad j = 0, \dots, d.$$

Its separant in $P^{(d)}$ is $(-1)^{d+1} R^d \prod_{j=1}^d j!$, a nonzero polynomial independent of P . Each column has degree at most $D + d < p$, so the Wronskian vanishes exactly when $RH' - PH = 0$ for some nonzero polynomial H of degree at most d . Hence $P = RH'/H$. For nonconstant H , its roots must belong to R ; conversely every multiset of at most d roots gives a solution. Constant H gives $P = 0$. The residues are integers between zero and $d < p$ and distinguish all the multisets. The number is

$$\sum_{s=0}^d \binom{N+s-1}{s} = \binom{N+d}{d}.$$

Every solution is regular by the separant formula. The finite classification proves the isolated count; a free challenge gives one constant-polynomial graph per solution. $\square$

Proposition J.5 (Linear MCA for the logarithmic-derivative family). *For fixed d , let R be monic and squarefree of degree $D + 1$, in characteristic zero or $p > d$. Consider the family*

$$\mathcal{P} = \{RH'/H : H \text{ monic, } \deg H \leq d, \text{ all roots of } H \text{ belong to } R\}.$$

This includes the zero polynomial. For $D < A \leq n$ with $A - D \geq \eta n$, every received line has $O_{d,\eta}(n)$ full-support bad labels witnessed by $\mathcal{P}$, for sufficiently large n . A candidate can be reused at many challenges; no prior assignment of labels is assumed.

Proof. Select one bad candidate at each of L distinct bad labels, and write $W = P/R$. At each root of R its residue belongs to $\{0, \dots, d\}$. On any affine line of graph points (z, W) where W varies, one residue is an injective affine function of z , so the line has at most $d + 1$ selected points. If W is constant, the candidate P is fixed. It has at most n bad labels: a bad full support must contain an agreeing coordinate with $g(x) \neq 0$, which determines z ; otherwise $(P, 0)$ explains that support. Thus every graph line has at most $M = \max\{n, d + 1\}$ selected points.

A noncollinear triple gives a nonzero rational relation of numerator degree at most $C = 3d - 1$, after clearing its three denominators of degree at most d . It has at most C common agreeing coordinates outside roots of R . There are at most $(M - 2)\binom{L}{2}/3$ collinear triples, by counting pairs. Let r_0 count domain roots of R , and put $m = n - r_0$ and $t = A - r_0$. Triple incidence and convexity give

$$m \left(\frac{tL}{m} - 2 \right)_+^3 \leq CL^3 + m(M - 2)L^2.$$

For sufficiently large n , $t \geq \eta n - 1$ and $m \leq n$ imply $t^3 > 8Cm^2$. If $L < 4m/t$ the claim is immediate. Otherwise the left side is at least $t^3 L^3 / (8m^2)$, giving

$$L \leq \frac{8m^3(M - 2)}{t^3 - 8Cm^2} = O_{d,\eta}(n).$$

This counts genuinely bad labels rather than all nearby challenges, which may fill the field on a constant candidate graph. $\square$

Verification and MCA scope. The exact checker in `research/actual_higher_order_components/` checks the complete solution classification on four small-field fixtures, the separant through order five, and 40 reduced sections of the projected power images and fixed-challenge fibers. These sections have exactly e^r and e^{r-1} points. It includes a small-characteristic failure and explicitly checks lower-degree boundary powers, which must not be lost by using only monic degree- r charts. The geometric upper bound is a separate written argument. A second checker exhausts 7,203 polynomials for the fixed-equation family and verifies its separant through order four.

The positive-dimensional degree saving alone does not prove an $O(n^d)$ full-MCA bound. Isolated actual joint points can number $\Theta(D^{d+1})$, and actual affine codeword-graph curve components can each contribute accidental labels. Their generic challenge fibers are isolated points, outside the positive-dimensional fiber bound. The sharpness family here has at most r distinct roots per candidate and is already covered by the linear-MCA result of Appendix G in its sufficiently large characteristic regime. The fixed-gap quadratic MCA lower bound remains open.

K A nonlinear first-order family near the characteristic cutoff

The bounded-root result of Appendix G requires characteristic larger than a suitable constant times the length. For the inverse Bernoulli family below, a direct classification also handles characteristic just above the message degree. The extra Frobenius constants available in that range do not produce large lists.

Theorem K.1 (Inverse Bernoulli solution count). *Let $s \geq 2$, $D \geq 1$, and let E be algebraically closed of characteristic zero or $p > \max\{D, s+1\}$. For every nonzero polynomial $A(X)$, the equation*

$$P^{s-1}P' = A, \quad \deg P \leq D,$$

has at most $2s$ polynomial solutions, all regular. In characteristic zero, or when $p > sD$, it has at most s solutions. The bound $2s$ is attained for suitable degrees and primes with $p \equiv 1 \pmod{s}$.

Proof. Every solution has a positive degree $d < p$ (with the obvious characteristic-zero interpretation), and $\deg A = sd - 1$. If its leading coefficient is a , the leading coefficient of A is da^s . Thus all solutions have the same degree and their leading coefficients differ by s th roots of unity. Normalize one leading coefficient to one. It suffices to count monic solutions of the resulting equation and multiply by s . For two such solutions P, Q ,

$$(P^s - Q^s)' = 0.$$

In characteristic zero their difference of powers is constant. Factoring it as $\prod_{\zeta^s=1} (P - \zeta Q)$ shows that it cannot be a nonzero constant: a factor with $\zeta \neq 1$ has degree d . If it is zero, monicity gives $P = Q$. The same argument applies when $p > sD$. Otherwise, in positive characteristic, a distinct pair has

$$P^s - Q^s = C(X)^p, \quad 1 \leq \deg C \leq s-1.$$

At least three factors. Suppose $s \geq 3$. Cancel the monic gcd G of P, Q and put $U = P/G$, $V = Q/G$, both monic of the same degree $h \geq 1$. The rational map U/V has degree $h < p$ and is separable. The preimages of the s distinct s th roots of unity all lie among the roots of C and infinity. These are at most s points; the s fibers are disjoint and nonempty. Consequently each fiber has exactly one point, of ramification index h .

For completeness, an elementary degree count forces $h = 1$. Make a generic fractional linear change of the X coordinate so that all those fiber points are finite and the new infinity maps to

none of the s values, nor to zero or infinity. The new numerator and denominator have degree h . At each of the s points their derivative numerator $U'V - UV'$ vanishes to order $h - 1$, since $h < p$. Its degree is at most $2h - 2$. Hence $s(h - 1) \leq 2h - 2$, which for $s \geq 3$ gives $h = 1$.

Return to the original coordinate, and write $U = X - u$, $V = X - v$, with $u \neq v$. The polynomial $S = U^s - V^s$ has $s - 1$ distinct roots, none equal to u or v . These exhaust the roots of C , which therefore has degree $s - 1$ and is squarefree. Comparing valuations in $G^s S = C^p$ proves

$$p = sg + 1, \quad G = S_0^g, \quad d = (s - 1)g + 1,$$

where S_0 is the monic normalization of S and $g = (p - 1)/s$. There are no other roots of G . The strict assumption $p > s + 1$ ensures $g \geq 2$.

Direct differentiation, using $sg = -1$ in the field, gives

$$A = \text{nonzero constant} \cdot (UV)^{s-1} S_0^{p-2}.$$

At a root of multiplicity m in any other solution T , the multiplicity in $T^{s-1}T'$ is $sm - 1$, since $m < p$. Thus T can choose roots only from S_0 , with multiplicity g , and from u, v , with multiplicity one. If it chooses t of the former and k of the latter, its degree requires

$$tg + k = (s - 1)g + 1, \quad 0 \leq t \leq s - 1, \quad 0 \leq k \leq 2.$$

For $g \geq 2$ this forces $t = s - 1$ and $k = 1$, so $T = P$ or Q . There are at most two monic solutions. Conversely the displayed U, V, S_0 construction supplies two, and their s scalar multiples attain the bound whenever $p = sg + 1$ with $g \geq 2$.

Two factors. For $s = 2$, a distinct monic pair satisfies $(P - Q)(P + Q) = c(X - a)^p$. Since $P + Q$ has degree d and leading coefficient two, its two factors give

$$P = (X - a)^d + v(X - a)^{p-d}, \quad Q = (X - a)^d - v(X - a)^{p-d}, \quad v \neq 0.$$

In particular $d > p/2$. If $p - d \leq d - 2$, the coefficient of X^{d-1} in P fixes a , and then v is fixed. If $p - d = d - 1$, then $p = 2d - 1 \geq 5$ and a is a root of P of multiplicity $d - 1 > d/2$, so it is again unique. Every monic solution has at most one distinct partner. The displayed pair and its negatives attain four solutions. Finally, the separant P^{s-1} is nonzero for every solution of a nonzero- A equation, proving regularity. $\square$

The strict characteristic condition has a real exception. For $s = p - 1$ and $D = p - 1$, the polynomials

$$P_a = \frac{X^p - X}{X - a}, \quad a \in \mathbb{F}_p, \quad P_a^{p-2} P'_a = -(X^p - X)^{p-2},$$

and all their nonzero $\mathbb{F}_p$ scalar multiples give $p(p - 1)$ solutions. Here the jet degree s grows with p . This does not contradict Theorem K.1, which assumes $p > s + 1$.

Corollary K.2 (Linear full MCA for inverse Bernoulli families). *Under the theorem's hypotheses, for nonzero A the solution family witnesses at most $2sn$ full-support bad labels on any received line. For $A = 0$, the corresponding count at agreement threshold $a \geq 2$ is at most $\binom{n}{2}/(a - 1)$. In particular both counts are $O(n)$ at fixed s and fixed positive agreement fraction.*

Proof. A fixed candidate has at most n bad labels: each bad full support contains an agreeing coordinate with nonzero received direction, which determines the label; otherwise $(P, 0)$ explains that support. Apply Theorem K.1 when $A \neq 0$.

When $A = 0$, the degree guard makes every solution constant. Each coordinate gives an affine line $c = f(x) + zg(x)$ in the (z, c) plane. A bad support of size at least a contains at least two distinct line equations, since otherwise constant polynomials F, G explain the whole support. Among its coordinates there are at least $a - 1$ pairs with distinct line equations. Each such pair intersects at at most one point. Select one candidate per bad label and count pairs to obtain the stated bound. $\square$

Checks and scope. The exact standard-library checker in `research/inverse_bernoulli/` exhausts 164,803 monic polynomials across thirteen fixtures, checks six sharpness examples, and verifies the exceptional family at $p = 3, 5, 7, 11$. The argument is specific to inverse Bernoulli equations. It does not settle the first-order conjecture or give a quadratic MCA construction at any higher order.

L An exact cubic complete-fiber bound on the pinned subgroup

The near-complete-fiber classification leaves maps with many uncovered points outside its range. The following finite, computer-assisted result closes one such proposed packet construction.

Theorem L.1 (Exact cubic maximum). *Let $p = 2130706433$ and $S = \mu_{256} \subset \mathbb{F}_p^*$. Every rational map R of degree three, with coefficients in any extension of $\mathbb{F}_p$, has at most three fibers consisting of three distinct points of S . This bound is attained. Consequently a fixed map $R(X^{1024})$ has at most three complete degree-3072 fibers on μ_{262144} and supplies at most eight unions of such fibers.*

Proof. Write the monic locator of a complete fiber as $X^3 - e_1X^2 + e_2X - e_3$. Its coefficients belong to $\mathbb{F}_p$, with $e_3 \in \mu_{256}$. All these coefficient triples lie on one affine line: they are the monic members of the numerator–denominator pencil of R . Two full fibers determine that pencil over $\mathbb{F}_p$, regardless of the original coefficient field. Two locators on a degree-three pencil have no common root, since a common factor would reduce the degree of the map.

First suppose e_3 is constant. Scaling all roots in μ_{256} normalizes $e_3 = 1$, because cubing is bijective on this group. There are exactly

$$M = \frac{(256 - 1)(256 - 2)}{6} = 10795$$

three-element subsets with product one. Enumerating all $\binom{M}{2} = 58,260,615$ pairs of their (e_1, e_2) points and counting collinear points, with common-factor pencils excluded, gives maximum three and an explicit attaining line.

Otherwise e_3 parametrizes the line injectively. Normalize one fiber’s product to one and write

$$e_1 = a_1 + s_1(e_3 - 1), \quad e_2 = a_2 + s_2(e_3 - 1).$$

The associated map is

$$\frac{X^3 + (s_1 - a_1)X^2 + (a_2 - s_2)X}{s_1X^2 - s_2X + 1}.$$

It is degenerate exactly when the denominator vanishes at one of the three roots of the anchor locator $X^3 - a_1X^2 + a_2X - 1$. Every line with at least three complete fibers is found by taking one of the M product-one anchors, enumerating all $255M$ triples at the remaining product levels,

and identifying repeated pairs of slopes (s_1, s_2) . For every nondegenerate repeated pair, lookup among the normalized triples counts its complete fibers at all 256 product levels. The exhaustive enumeration of

$$255M^2 = 29,715,666,375$$

anchor/triple pairs gives maximum three, also attained. Together the two cases prove the first assertion. The map $X \mapsto X^{1024}$ is 1024-to-one from μ_{262144} onto S , so a complete degree-3072 fiber of the composition corresponds exactly to a complete three-point fiber of R . Three disjoint packets have at most 2^3 unions. $\square$

Reproducibility and scope. The source, mathematical reduction, all 85 disjoint anchor-range checkpoints, witnesses, and resource reports are in `research/structured_domains/general_cubic_packets/`. The general enumeration found 22,692 nondegenerate candidate occurrences and took 524.46 seconds, with sampled peak process-group RSS 176,192 KiB under a 384 MiB limit. A separate Python checker verified every checkpoint witness on the entire domain, validated the source hash and complete disjoint coverage, independently exhausted all pencils on $\mu_8/\mathbb{F}_{97}$ and $\mu_{16}/\mathbb{F}_{193}$, and reran the full pinned constant-product enumeration. The optional `-full` flag repeats the general enumeration too.

Eight unions, even with one fixed core appended, are far short of the required bank size 274980728111395088. This excludes the fixed composed-cubic packet route; maps varying with the candidate and constructions using partial fibers remain outside the conclusion. No better.codes score improvement follows from this exclusion.

M Sharp list bounds for polynomial Riccati solutions

The cross-ratio identities of Riccati equations, combined with intersection multiplicities in positive characteristic, give an ordinary-list bound even when all agreement coordinates are singular. For a challenge-independent nonlinear equation it also gives a linear full-support exceptional-label bound. The parametrization and its gcd form are classical; see Gasull, Torregrosa, and Zhang [4, Lemmas 6–7]. We prove the finite-characteristic statements and the agreement consequences used here.

Theorem M.1 (Riccati lists and solution counts). *Let $D \geq 1$, and let E have characteristic zero or $p > D$. Consider one fixed equation*

$$a(X)P' = b_0(X) + b_1(X)P + b_2(X)P^2, \quad a \neq 0, \quad (63)$$

with coefficients in $E[X]$, with no restriction on their degrees. On any n distinct evaluation points, the number of distinct solutions of degree at most D agreeing with an arbitrary received word in at least $A > D$ coordinates is at most

$$\left\lfloor \frac{n}{A - D} \right\rfloor. \quad (64)$$

This bound is attained over prime fields at unbounded lengths and positive limiting rate and agreement gap. If $b_2 \neq 0$, the total number of degree-at-most- D polynomial solutions is at most

$$C_D = \begin{cases} D + 2, & \text{in characteristic zero or } p > D + 1, \\ 2D + 2, & \text{if } p = D + 1. \end{cases} \quad (65)$$

The bound $D + 2$ is sharp. The boundary bound $2D + 2$ is attained at $p = 2, D = 1$.

Proof. Total candidate count. Assume $b_2 \neq 0$. Work at a transcendental Taylor center t , where $a(t)b_2(t) \neq 0$, and write the equation as $P' = F(X, P)$, with $F = \alpha + \beta u + \gamma u^2$ and $\gamma \neq 0$. Put $\Delta = \partial_X + F\partial_u$ and $F_j = \Delta^j u$. For $j < p$, or in characteristic zero, F_j has u -degree $j+1$ and leading coefficient $j!\gamma^j$. This follows inductively: $F\partial_u$ gives the highest degree and ∂_X gives a lower degree. In characteristic zero or $p > D+1$, every solution satisfies

$$F_{D+1}(t, P(t)) = P^{(D+1)}(t) = 0.$$

This is a nonzero polynomial of degree $D+2$ in $P(t)$. Distinct polynomial solutions have distinct values at the transcendental t , proving the first count.

If $p = D+1$, use instead the truncated reconstruction

$$T(U, X) = \sum_{j=0}^D F_j(t, U) \frac{(X-t)^j}{j!}.$$

Every candidate is $T(P(t), X)$. Its U -degree is $D+1$, with leading coefficient $\gamma(t)^D(X-t)^D$. The residual $a\partial_X T - b_0 - b_1 T - b_2 T^2$ therefore has U -degree exactly $2D+2$, with nonzero leading coefficient $-b_2(X)\gamma(t)^{2D}(X-t)^{2D}$. Any nonzero X -coefficient of this residual is a nonzero polynomial in U of degree at most $2D+2$ vanishing at every candidate value $P(t)$. This proves the boundary count. In particular, in positive characteristic the total count is at most $2D+2 \leq 2p$.

Lists by intersection multiplicities. For $M \geq 2$ distinct candidates and a domain coordinate x , put

$$C_x = \sum_{i < j} \text{ord}_x(P_i - P_j).$$

The root bound gives $\sum_x C_x \leq D \binom{M}{2}$. If h candidates match a given received value at x , we claim

$$(h-1)(M-1) \leq 2C_x. \quad (66)$$

If $b_2 = 0$, differences solve one homogeneous linear equation. The ratio of two nonzero differences has derivative zero and rational degree at most $D < p$, and hence is constant. The same is true in characteristic zero. Thus the candidates have form $P_0 + cW$, and their values at x are either all distinct or all equal; (66) follows immediately.

Now suppose $b_2 \neq 0$. Differences satisfy

$$a(P_i - P_j)' = (b_1 + b_2(P_i + P_j))(P_i - P_j),$$

so every cross ratio has derivative zero. In characteristic zero it is constant and different from $0, 1, \infty$. A colliding pair and two candidates outside its value class would force that cross ratio to be one. Therefore either all values are distinct or one value occurs at least $M-1$ times. For $h \in \{0, 1, M-1, M\}$, the inequality follows from $C_x \geq \binom{h}{2}$.

In positive characteristic a cross ratio need not be constant. However, any zero of a nonzero rational function with derivative zero has order divisible by p . Choose a colliding pair and two candidates outside its value class. The denominator of their cross-ratio identity is nonzero at x , and the numerator of the cross ratio minus one, up to sign, is the product of the two within-pair differences. If the two outside values were distinct, this product would have positive order at most $D < p$, impossible. Consequently every collision pattern other than those already handled has exactly two value classes, both of size at least two.

Let the matching class have size h and the other class have size $b = M - h$. Let μ, ν be the minimum positive vanishing orders of differences within these classes. The same cross-ratio

argument for every choice of two within-class pairs gives $\mu + \nu \geq p$. Write $A_0 = \binom{h}{2}$, $B_0 = \binom{b}{2}$, so $C_x \geq A_0\mu + B_0\nu$. If $h \leq b$, already

$$2(A_0 + B_0) - (M - 1)(h - 1) = (b - 1)(b - h + 1) \geq 0.$$

If $h \geq b$, minimizing with $\mu, \nu \geq 1$ and $\mu + \nu \geq p$ gives $C_x \geq A_0 + (p - 1)B_0$, and hence

$$2C_x - (M - 1)(h - 1) \geq (b - 1)(pb - M + 1) \geq 0,$$

because $b \geq 2$ and the total-count bound gives $M \leq 2p$. An unmatched received value has $h = 0$. This proves (66) in all cases. Summing it gives $AM \leq n + DM$, proving (64). The case $M = 1$ is immediate, and the same proof allows real agreement thresholds.

For sharpness of $D + 2$, choose a split and squarefree of degree $D + 1$. The equation $aP' = a'P - P^2$ has the $D + 2$ solutions 0 and $a/(X - r)$, one for each root r of a . At the boundary $p = 2, D = 1$, $(X^2 + X)P' = P^2 + P$ has exactly 0, 1, $X, X + 1$ as its four solutions. $\square$

Proposition M.2 (Sharp prime-field lists). *Fix $M \geq 3$ and a prime $p \equiv 1 \pmod{2M}$. Put*

$$s = \frac{p - 1}{2M}, \quad n = p - 1, \quad D = (M - 2)s, \quad A = Ms.$$

There is a nonlinear equation (63) over $\mathbb{F}_p$ and a word on $\mathbb{F}_p^$ whose list of degree-at-most- D solutions at agreement threshold A has exactly $M = n/(A - D)$ elements.*

Proof. Partition the image of X^s on $\mathbb{F}_p^*$ into $a_1, \dots, a_M, b_1, \dots, b_M$ and set

$$R(Y) = \prod_{i=1}^M (Y - a_i), \quad T(Y) = Y^{M-1}, \quad C_i(Y) = \frac{R(Y)}{Y - a_i} - T(Y), \quad P_i(X) = C_i(X^s).$$

Leading terms cancel, so $\deg P_i \leq D$. On the fiber over a_j choose received value $-T(a_j)$, and on the fiber over b_j choose $C_j(b_j)$. Exactly $M - 1$ candidates match a root fiber, since the remaining difference is $R'(a_j) \neq 0$. The values $C_i(b_j)$ are all distinct, since $R(b_j) \neq 0$ and the denominators $b_j - a_i$ are distinct. Each P_i therefore agrees in exactly $(M - 1)s + s = A$ coordinates, and the polynomials are distinct.

The polynomials $U_i = R/(Y - a_i)$ satisfy $RU'_i = R'U_i - U_i^2$. Translating $U = C + T$ gives

$$RC' = (R'T - T^2 - RT') + (R' - 2T)C - C^2.$$

Substitute $Y = X^s$ and multiply the right side by sX^{s-1} . This is one polynomial Riccati equation for all the P_i , with nonzero quadratic coefficient. Since $p > 2D$, the preceding theorem shows that no additional solution can enter this list. Primes in the progression $1 \pmod{2M}$ give unbounded lengths for fixed M . The code rate $(D + 1)/n$ tends to $(M - 2)/(2M)$ and the capacity gap $(A - D - 1)/n$ tends to $1/M$. $\square$

Corollary M.3 (Fixed-equation full-support exceptional labels). *Under Theorem M.1, assume $b_2 \neq 0$. For a received line $f + zg$ on n distinct points and $A > D$, the number of labels possessing a solution of the fixed equation with at least A agreements and bad full agreement support is at most*

$$\frac{2n(n + C_D)}{A - D}. \tag{67}$$

Here a support is bad if no polynomial of degree at most D interpolates g on the entire support. Thus the bound is $O(n)$ when $A - D$ is a fixed positive fraction of n .

Proof. Put $\delta = A - D$. A candidate's persistent coordinates have $g_i = 0$ and $P(x_i) = f_i$. Call a candidate heavy if it has at least $D + \delta/2$ persistent coordinates. These candidates form a list for f at that threshold, so there are at most $2n/\delta$ by (64). Each has at most n bad labels: a bad support must contain an agreement at a coordinate with $g_i \neq 0$, and this coordinate determines a unique label. Otherwise $g = 0$ on the support, contradicting badness.

A light candidate needs more than $\delta/2$ nonpersistent agreements at every nearby label. It has at most n such coordinate-label incidences altogether, so it has at most $2n/\delta$ nearby labels. There are at most C_D candidates in total, with C_D as in (65). Adding the heavy and light bounds proves the claim. $\square$

Characteristic guard and scope. The constant-cross-ratio argument alone cannot justify the weaker guard $p > D$; intersection multiplicities are essential. In characteristic three, put $D = 2$, $U = 1 + X + X^2$, $V = X + X^2$, $W = 1 + X$, and $N = UV$. The equation with $a = N$, $b_0 = 0$, $b_1 = N'$, and $b_2 = U'V - UV'$ has the solutions $0, U, V, W$. Their values at zero are $0, 1, 0, 1$, contradicting the collision property. Their cross ratio is a nonconstant rational function with zero derivative. Exact enumeration gives five degree-at-most-two solutions, also exceeding $D + 2$. This is consistent with the boundary value $C_D = 2D + 2$ in (65).

The failure of the $D + 2$ count at the boundary is not confined to small primes. For every prime p , the equation

$$(X^p - X)P' = -P - P^2, \quad \deg P \leq p - 1,$$

has exactly $p + 2 = D + 3$ solutions over every extension of $\mathbb{F}_p$: 0 , -1 , and $(X^p - X)/(X - r)$ for $r \in \mathbb{F}_p$. Indeed, evaluating at $\mathbb{F}_p$ forces each value of P to be 0 or -1 . Let S be the set of coordinates with value -1 and $f = \prod_{r \in S} (X - r)$. Lagrange interpolation gives $P = (X^p - X)f'/f$. Substitution reduces the equation to $f'' = 0$. Since $\deg f \leq p$, a nonconstant proper divisor f of $X^p - X$ with $f'' = 0$ must be linear: its only possible nonzero coefficients have degrees $0, 1, p$. Thus S is empty, a singleton, or all of $\mathbb{F}_p$, proving the assertion. This family does not prove that $D + 3$ is a universal upper bound at the boundary.

The ordinary-list bound applies separately at each label of a challenge-dependent Riccati equation. The full-support bound (67) requires a single fixed equation and does not extend by this argument to challenge-dependent families. The checks in `research/riccati_cross_ratio/` enumerate 183,581 candidate polynomials, test 616 received lines, and verify nine prime-field equality fixtures for Proposition M.2, together with the characteristic-three negative control. A further boundary-characteristic checker enumerates 1,084,807 polynomials for 549 equations, checks 2,790 coordinate-family instances of the multiplicity inequality, and includes nonconstant cross ratios and the sharp characteristic-two boundary count. An independent extension-field implementation exhausts all 729 degree-at-most-two polynomials over $\mathbb{F}_9$ and checks 85 received lines, including bad full-support labels at $p = D + 1$. These restricted results do not settle the general first-order conjecture.

N Quadratic ordinary correlated-agreement failures in large characteristic

Corollary 2.7 concerns full agreement-set MCA at a fixed exact gap. A separate construction excludes even ordinary correlated agreement, still with quadratically many nearby labels, but with a gap that is bounded below rather than fixed exactly. The key is to descend a nearest-polynomial orbit to its own evaluation domain.

Theorem N.1. *For unbounded lengths n there are primes p and Reed–Solomon codes over $\mathbb{F}_{p^2}$, of exact rate $3/13$, with a received affine line having at least $\lceil n^2/4096 \rceil$ nearby labels at an agreement threshold T , but no ordinary correlated agreement of size T . The capacity gap satisfies $T/n - 3/13 > 3/26$, the characteristic exceeds the message degree, and the radius is strictly below the characteristic-based Elias radius for all sufficiently large members.*

This rules out a linear ordinary-CA bound uniform over gaps bounded below in the large-characteristic field class. The actual ratio T/n is not asserted to be constant, nor is p/n bounded above. This is neither a prime-ambient-field result nor a first-order-regime tightness statement.

An eight-point certificate. Let ζ be a primitive eighth root of unity. No affine polynomial agrees at three distinct points of μ_8 with $W_2(Y) = (1 + Y^4)/2 - Y^2$ in characteristic $p > 17$, $p \equiv 1 \pmod{8}$. For completeness, write $v_j = 1 + \zeta^{4j} - 2\zeta^{2j}$. For $0 \leq i < j < h < 8$, the norm of

$$(\zeta^j - \zeta^i)(v_h - v_i) - (\zeta^h - \zeta^i)(v_j - v_i)$$

from $\mathbb{Q}(\zeta)$ to $\mathbb{Q}$ has the following complete distribution, obtained by reduction modulo $Z^4 + 1$ and resultants:

Norm	64	256	576	1088	3136	4096	6400
number of triples	12	8	8	8	4	8	8.

All 56 determinants are nonzero and their norms have prime factors only 2, 3, 5, 7, 17. Reduction at $p > 17$ preserves nonvanishing. Exact verification and individual determinants are in `research/dickson_fixed_gap/verify_two_orbit.py` and its JSON output.

Proof of Theorem N.1. Fix an integer $R \geq 40$, and let B_R be the product of odd primes at most R . By the Chinese remainder theorem and Dirichlet’s theorem, there are arbitrarily large primes with

$$p \equiv 9 \pmod{16}, \quad p \equiv -1 \pmod{B_R}.$$

Put $k = (p - 1)/4$. Then $v_2(k) = 1$, $k \equiv 1 \pmod{3}$, and no odd prime at most R divides k . On $\mathbb{F}_p^\times$ use $W_k(X) = (1 + X^{2k})/2 - X^k$. Its true maximum agreement M with degree- $< k$ polynomials is at least $3k/2$ by Proposition 2.5, and at most $2k$ by degree. The group $H = \mu_k$ preserves the word and acts on its nearest list by $P(X) \mapsto P(hX)$.

Take one nearest polynomial and let r be its orbit size. This divides k . If $r \leq R$, then $r \in \{1, 2\}$, forcing $P(X) = a + bX^{k/2}$. Passing to $Y = X^{k/2}$, whose eight fibers each have size $k/2$, the preceding certificate bounds its agreement by k , a contradiction. Thus $r > R$, and $3 \nmid r$.

Descent. Put $d = k/r$. The stabilizer is μ_d , so $P(X) = V(X^d)$ with $\deg V < r$. The map $X \mapsto X^d$ sends $\mathbb{F}_p^\times$ onto μ_{4r} with fibers of size d , and $W_k(X) = W_r(X^d)$. Consequently $M = dm$, where $m \geq 3r/2$ is the agreement of V with W_r on μ_{4r} . Any closer degree- $< r$ polynomial would lift to a closer degree- $< k$ polynomial, impossible. The μ_r -orbit of V has exactly r members, all nearest at m . We now have a true nearest list of size r on $4r$ coordinates, in the same prime field. Its maximum stays m over $\mathbb{F}_{p^2}$: matching r base-field coordinates and values forces base-field coefficients by interpolation.

Anchors and rate. Choose $t \in \{1, 2\}$ with $t \equiv r \pmod{3}$. At each of t steps, choose an agreement anchor incident to the largest number of remaining candidates, then divide candidate and word minus the anchor value by the anchor’s linear factor. At step $j = 0$ or 1 , the incidence fraction is $(m - j)/(4r - j) \geq 1/3$. At least $\ell \geq r/9$ quotients remain, with degrees $< K = r - t$ and exactly $m - t$ old agreements. No polynomial of degree $< K$ has more: restoring the anchors would contradict the old maximum. Set

$$q = (r - 10t)/3, \quad n = 4r - t + q = 13(r - t)/3, \quad T = m - t + 1.$$

Here $q \geq r/6 > 0$, $q \leq m - r + 1$, $K/n = 3/13$, and

$$\frac{T - K}{n} = \frac{m - r + 1}{n} > \frac{3}{26}.$$

Labels and absence of common agreement. Put $Q = p^2$. On the $Q - 4r$ points outside the original domain, pairwise root counting and Cauchy–Schwarz give average distinct-value count at least

$$\mu = \frac{\ell(Q - 4r)}{Q - 4r + (\ell - 1)(K - 1)} \geq \ell/2,$$

since $p \geq 4r + 1$ and $\ell \leq r$. Choose the q points with largest value counts s_x , give them direction one, and retain direction zero on the old domain. Independent uniform translations of the new value sets have expected union size

$$Q \left[1 - \prod_x (1 - s_x/Q) \right] \geq Q[1 - (1 - \mu/Q)^q] \geq q\ell/4.$$

The last inequality uses $1 - u \leq e^{-u}$, $\mu \geq \ell/2$, and $q\ell/(2Q) \leq 1/4$. Each counted label adds a new agreement to a selected quotient, giving at least T agreements. The count is at least

$$q\ell/4 \geq r^2/216 \geq n^2/4056 > n^2/4096.$$

A correlated witness pair F, G of degree $< K$ with $G = 0$ has joint agreements only on the old core, with maximum $m - t = T - 1$. If $G \neq 0$, it has at most $K - 1$ old zeros and at most q new joint agreements, again at most $r - t - 1 + (m - r + 1) = m - t = T - 1$. Thus ordinary correlated agreement of size T is absent.

Finally $p \geq 4r + 1 > K - 1$ and $H_p(1 - T/n) \leq 1 - T/n + 1/\log_2 p < 1 - K/n$ for large p . Taking R unbounded gives unbounded r and n . $\square$

The proof does not show that the descended domain is short relative to p . If its orbit sizes could be chosen with $r/p \rightarrow 0$, the same descent and the prime-field exact-halving compiler would give superlinear full-support MCA failures over prime fields at exact rate $1/8$ and gap $1/16$. This ratio condition remains unproved. Details and independent small-field checks are in [research/ordinary_ca_superlinear/](#).

N.1 Fixing both the rate and the gap

Allowing larger finite extensions permits an exact positive gap as well. The following operations preserve the true nearest-codeword boundary; their field extensions, rather than just fresh coordinates, are essential.

Lemma N.2 (Boundary-preserving padding). *Let a word w on N distinct points of a finite field F have maximum agreement $M \geq K$ with degree- $< K$ polynomials. Keep any selected nearest list. The following operations preserve that list as nearest witnesses.*

1. If $a \in F$ is outside the domain and $b \notin F$ lies in a finite extension, appending (a, b) preserves dimension K and maximum M .
2. If $a \notin F$, replace every old value by $(x - a)w(x)$, append $(a, 0)$, and increase dimension to $K + 1$. The maximum becomes $M + 1$; replace each selected polynomial P by $(X - a)P$.

Proof. Maximum agreement cannot increase upon extending the coefficient field, since K matching base-field coordinates and values force coefficients into F . For the first operation, a polynomial with $M + 1$ agreements would match the new point and at least $M \geq K$ old ones. Its coefficients would belong to F , contradicting its new value b at $a \in F$.

For the second operation, the displayed polynomials attain $M + 1$. Suppose a degree- $\leq K$ polynomial attained $M + 2$. If it matches at a , dividing by $X - a$ contradicts the old maximum. Otherwise take a set S of $M + 2$ old matches. Modulo the degree- $\leq K$ evaluation space on S , the interpolation condition reads $[xw] - a[w] = 0$. This is a linear system over F , so $a \notin F$ forces both classes to vanish. There are polynomials A, B of degree at most K interpolating xw, w on S . But $A - XB$ has degree at most $K + 1$ and at least $M + 2 \geq K + 2$ roots. Thus $A = XB$, so $\deg B < K$, contradicting the old maximum. $\square$

Corollary N.3 (Quadratic ordinary-CA failures at a fixed exact gap). *For unbounded lengths n there are finite fields of characteristic $p > n/8 - 1$ and Reed–Solomon codes of exact rate $1/8$, with a received line having at least $\lceil n^2/8192 \rceil$ nearby labels at agreement $3n/16$, but no ordinary correlated agreement of that size. The capacity gap is exactly $1/16$. The radius is strictly below the characteristic-based Elias radius for all sufficiently large members.*

Proof. Use the descended source from the proof of Theorem N.1: length $4r$, dimension r , true maximum $m \in [3r/2, 2r]$, and r selected nearest polynomials over $\mathbb{F}_p$, with $p \geq 4r + 1$ and $r > 40$. Put

$$\Delta = m - r + 1, \quad u = 13\Delta - 3r, \quad s = 2\Delta - r + 1.$$

These are nonnegative integers. First apply the first operation of Lemma N.2 u times, and then the second operation s times. One may initially pass to $\mathbb{F}_{p^2}$ to provide all u fresh first-stage coordinates; each subsequent new value or common-zero coordinate is chosen outside the entire current field. A tower of quadratic extensions suffices. The resulting source has

$$N' = 15\Delta + 1, \quad K' = 2\Delta + 1, \quad M' = 3\Delta,$$

and retains r selected nearest polynomials.

An agreement anchor is incident to at least $\ell \geq M'r/N' \geq r/6$ candidates. Dividing by its linear factor gives an old core of length 15Δ , dimension 2Δ , and true maximum $3\Delta - 1$. Append $q = \Delta$ challenge coordinates with direction one, retaining direction zero on the core. Enlarge the field if necessary to cardinality $Q \geq 2N'r$. The same root-count and random-translation argument gives at least $q\ell/4 \geq \Delta r/24$ nearby labels. The final parameters are

$$n = 16\Delta, \quad K = 2\Delta, \quad T = 3\Delta.$$

A zero explaining direction has common agreement at most $3\Delta - 1$ on the core. A nonzero direction has at most $2\Delta - 1$ core zeros and Δ padding agreements, again at most $3\Delta - 1$. Thus ordinary correlated agreement at threshold T is absent.

Since $\Delta \leq r + 1$ and $\Delta \geq 4$, we have $r \geq 3\Delta/4$, giving

$$J \geq \Delta r/24 \geq \Delta^2/32 = n^2/8192.$$

Also $K - 1 = 2\Delta - 1 \leq 2r + 1 < p$. The exact gap $1/16$ yields the Elias inequality for large p , as before. $\square$

This proves that quadratic dependence on length is sometimes unavoidable for ordinary correlated-agreement exceptions even at a fixed positive capacity gap, in the full large-characteristic field

class. The extension degree is reduced to two below; Corollary N.3 is still not asserted over prime fields. Theorem N.1 gives a different exact rate with a gap only bounded below. Neither assertion establishes first-order-regime tightness.

Independent exhaustive checks of both preservation operations, including a negative control showing failure when the common-zero point is chosen inside the current field, are in `research/ordinary_ca_superlinear/verify_normalization.py`.

A complete finite census also identifies the nearest list for the full Dickson word over $\mathbb{F}_{41}$: its maximum agreement is 15 on 40 coordinates, and exactly 210 degree- < 10 polynomials attain it, in 21 multiplicative orbits of size 10. The four-coset anchor cover reduces completeness to 52,252,200 determining supports; all witnesses are independently replayed in `research/dickson_nearest_threshold/`. This finite result does not establish maximum agreement $3/8$ for an unbounded family.

N.2 Quadratic extensions via random directions

Proposition N.4. *The field in Corollary N.3 can be taken to be $\mathbb{F}_{p^2}$, and the exceptional count can be strengthened to $\lceil n^2/192 \rceil$.*

Proof. We combine a common-zero block with random padding directions. Throughout, a word on N distinct points of $F = \mathbb{F}_q$ has true maximum agreement $M \geq K$ with degree- $< K$ polynomials, and $\Delta = M - K + 1$.

A common-zero block. For $1 \leq s \leq \Delta$, choose s distinct elements A of $E = \mathbb{F}_{q^2} \setminus F$ with $A \cap A^q = \emptyset$; there are $(q^2 - q)/2$ conjugate pairs available. Let $Z = \prod_{a \in A} (X - a)$. Multiply the old word by Z , append A with received value zero, and increase dimension to $K + s$. The maximum becomes exactly $M + s$. Indeed, suppose Q of degree $< K + s$ has more agreements. Divide out its j matching new factors, leaving R of degree $< K + h$ and a remaining locator Z_0 of degree $h = s - j$. On at least

$$\ell = M + h + 1 = K + \Delta + h \geq K + 2h$$

old coordinates, $R(x) = Z_0(x)w(x)$. Coefficientwise q -Frobenius preserves these equations, so $RZ_0^{(q)} - R^{(q)}Z_0$ has ℓ roots and degree $< K + 2h$; it is zero. Thus the reduced rational function R/Z_0 lies in $F(X)$, by uniqueness of its expression with monic denominator. That denominator divides both Z_0 and $Z_0^{(q)}$, which are coprime by the choice of A . Therefore R/Z_0 is a polynomial of degree $< K$ with $\ell > M$ agreements, a contradiction. Multiplying every selected old nearest polynomial by Z attains $M + s$.

Random padding directions. Suppose a core of length N_c and dimension K has true maximum $T - 1$, and put $\Delta = T - K > 0$. On any t fresh coordinates over a field of size H , choose the direction independently and uniformly nonzero, retaining direction zero on the core. A nonzero degree- $< K$ polynomial with T direction agreements has $z < K$ old zeros and $T - z$ new matches. The probability that any such polynomial exists is at most

$$\sum_{z=0}^{K-1} \binom{N_c}{z} \binom{t}{T-z} \frac{H^{K-z}}{(H-1)^{T-z}} \leq 2^{N_c+t} H^{-\Delta} (1 - 1/H)^{-t}.$$

Indeed, the old zeros leave at most H^{K-z} polynomials, and each prescribed new value is matched with probability at most $1/(H-1)$. If this bound is less than one, a good direction exists. No nonzero explaining direction can then give ordinary CA at threshold T , regardless of the intercept. The zero explaining direction has joint agreements only on the core, where its intercept has at most $T - 1$.

Apply the zero block to the descended source. By Corollary O.2, $r/2+1 \leq \Delta = m-r+1 \leq 2r/3+1$. For large r , the $s = 2\Delta - r + 1$ common zeros satisfy $1 \leq s \leq \Delta$, so one quadratic extension suffices. The resulting source has length $3r + 2\Delta + 1$, dimension $2\Delta + 1$, and maximum 3Δ . One anchor retains

$$\ell \geq \frac{3\Delta r}{3r + 2\Delta + 1}$$

nearest quotients on a core of length $N_c = 3r + 2\Delta$, dimension $K = 2\Delta$, and maximum $T - 1 = 3\Delta - 1$. Append $t = 14\Delta - 3r$ coordinates of $\mathbb{F}_{p^2}$, giving $n = 16\Delta$. With $H = p^2 \geq 16r^2$, the random-direction failure bound tends to zero: its logarithm is $-\Omega(r \log r) + O(r)$.

Choose the t fresh coordinates with largest evaluation diversity among the ℓ quotients. Pairwise root counting gives average diversity at least $\ell/2$ for large r . Fix a good direction there. At a new coordinate with direction $d_x \neq 0$, the counted labels are $(P_i(x) - f_0(x))/d_x$. Nonzero scaling preserves diversity, and random intercepts give independent uniform translations. Since $t \leq 7r$ and $t\ell/(2H) \leq 7/32 < 1$, their expected union is at least $t\ell/4$. Every counted label has $T - 1$ old matches and a new match. For large r , $r \leq 2\Delta - 2$ gives $t \geq 8\Delta$, while $r \geq 3(\Delta - 1)/2$ gives $\ell \geq 2\Delta/3$ once $\Delta \geq 13$. Thus there are at least

$$t\ell/4 \geq 4\Delta^2/3 = n^2/192$$

nearby labels and no ordinary CA. The rate and gap are exactly $1/8$ and $1/16$, and $p > K - 1$. $\square$

The zero-block checks include a conjugate-pair negative control. Independently, a complete $\mathbb{F}_{11}$ test checks all 1,000 nonzero padding directions, all line labels, and all 121^2 ordinary-CA witness pairs for a selected line. Its 11 labels are all nearby at threshold three, but its maximum ordinary CA is two. The finite checks and parameter inequalities are in `research/ordinary_ca_superlinear/verify_random_direction.py`; the asymptotic conclusion follows from the proof above.

N.3 A family of exact rates and gaps

Corollary N.5. *Fix $b \in \{2, 3, 4, 5\}$ and an integer $d \geq b + 7$. There are unbounded-length families over $\mathbb{F}_{p^e}$ of exact rate b/d and capacity gap $1/d$ with at least*

$$\left\lceil \frac{(b+1)n^2}{4(b+6)d^2} \right\rceil$$

nearby challenges and no ordinary correlated agreement at that threshold. The characteristic exceeds the message degree, $e = 2^{b-1}$ is fixed, and the radius is strictly below characteristic-based Elias for sufficiently large p . In particular, $b = 2, d = 9$ gives rate $2/9$, gap $1/9$, at least $\lceil n^2/864 \rceil$ nearby challenges, over $\mathbb{F}_{p^2}$. The case $b = 5, d = 12$ gives rate $5/12$, gap $1/12$, at least $\lceil n^2/1056 \rceil$ nearby challenges, over $\mathbb{F}_{p^{16}}$.

Proof. Corollary O.2 bounds the original full-field source maximum by $5k/3$. Descent by the stabilizer size k/r therefore gives $m \leq 5r/3$. Consequently

$$r/2 + 1 \leq \Delta = m - r + 1 \leq 2r/3 + 1 \leq r$$

for sufficiently large r . This also preserves the characteristic condition for the additional case $b = 5$.

Add $s = b\Delta - r + 1$ common zeros, in at most $b - 1$ blocks of size at most Δ , using the quadratic block operation above. This is possible since $\Delta \leq r - 1$ for large r , so $s \leq (b - 1)\Delta$. Enlarge to

$\mathbb{F}_{p^{2b-1}}$ if necessary; interpolation preserves the nearest maximum under enlargement. The new source has

$$N' = 3r + b\Delta + 1, \quad K' = b\Delta + 1, \quad M' = (b+1)\Delta.$$

An agreement anchor retains

$$\ell \geq \frac{(b+1)\Delta r}{3r + b\Delta + 1} \geq \frac{(b+1)r}{b+6},$$

where $r \leq 2\Delta - 2$. After division, the core has length $3r + b\Delta$, dimension $K = b\Delta$, and maximum $T - 1 = (b+1)\Delta - 1$. Append $t = (d-b)\Delta - 3r \geq \Delta$ fresh coordinates, giving $n = d\Delta$. With field size $H \geq p^2 \geq 16r^2$, the random-direction failure bound from Proposition N.4 has logarithm $-\Omega(r \log r) + O_{b,d}(r)$, so a good direction exists. It excludes ordinary CA for every choice of new intercepts.

For a uniform count, choose Δ padding coordinates of largest diversity and fill the other positions arbitrarily. Their average diversity is at least $\ell/2$ for large r , and $\Delta\ell/(2H) < 1$. Independent intercept translations on these counted coordinates yield

$$J \geq \Delta\ell/4 \geq \frac{(b+1)\Delta r}{4(b+6)} \geq \frac{(b+1)n^2}{4(b+6)d^2},$$

using $r \geq \Delta$. The other padding coordinates cannot invalidate any counted near label or the direction certificate. Finally, $b\Delta - 1 \leq 5(2r/3 + 1) - 1 < 4r + 1 \leq p$ for large r , and the exact positive gap $1/d$ gives strict Elias for sufficiently large p . $\square$

The exact integer identities and extension-degree accounting are recorded in `research/ordinary_ca_superlinear/RANDOM_DIRECTION_QUADRATIC.md`, with independent proof audit and finite checks alongside it. The extension degree is now fixed independently of length. The examples still do not lie above the first-order agreement curve or in prime ambient fields.

N.4 Characteristic-zero obstructions at prime and prime-power cyclic orders

The following restriction concerns a characteristic-zero cyclic source; it does not restrict exceptional split characteristics in which such a source may exist only after reduction. The prime-order statement also follows from the more general coset-word argument recorded in `research/cyclotomic_dickson_word/PRIME_ORDER_COSET_WORD_OBSTRUCTION.md`. We give a direct proof that extends to the prime-power bound below.

Proposition N.6. *Let $r > 5$ be prime, let μ_{4r} be the $4r$ th roots of unity in an algebraically closed field of characteristic zero, and put*

$$W_r(X) = \frac{1 + X^{2r}}{2} - X^r.$$

Every polynomial of degree below r agrees with W_r on at most $r+2$ points of μ_{4r} . Thus this length- $4r$, dimension- r source has maximum agreement fraction at most $1/4 + 1/(2r)$, and capacity gap at most $1/(2r)$.

Proof. Suppose P agrees on a subset S of size $M \geq r+3$. Write

$$F = X^{2r} - 2X^r + 1 - 2P = A_S Q, \quad A_S = \prod_{z \in S} (X - z).$$

The monic quotient Q has degree $k = 2r - M \leq r - 3$. Its coefficients are determined successively by the coefficients of F in degrees $2r, \dots, M$, all above r . Those coefficients are $1, 0, \dots, 0$, so monic

division gives $Q \in \mathbb{Z}[\zeta_{4r}][X]$ and hence $P \in \mathbb{Z}[\zeta_{4r}, 1/2][X]$. In particular, this conclusion requires no Vandermonde inverse at a ramified prime.

Reduce at a prime above r . The roots in μ_{4r} collapse to the four distinct fourth roots of unity, with exactly r original roots in each reduction bucket. If the selected multiplicities are n_c , then $\bar{F}$ is divisible by $\prod_c (X - c)^{n_c}$. Consequently

$$\bar{F}' = -2\bar{P}' \quad \text{is divisible by} \quad \prod_{n_c > 0} (X - c)^{n_c - 1}.$$

If nonzero, this derivative would have degree at least $M - 4 \geq r - 1$, whereas $\deg \bar{P}' \leq r - 2$. Thus $\bar{P}' = 0$, and $\deg \bar{P} < r$ makes $\bar{P}$ constant. At the four fourth roots, $\bar{W}_r$ takes the values $0, 2, i, -i$, in some order. These are distinct in characteristic $r > 5$: a collision of 2 with $\pm i$ would force characteristic 5. Only one bucket can therefore be occupied, contradicting $M > r$. The agreement-fraction statement follows by division by $4r$. $\square$

Corollary N.7. *If $r = \ell^a$ with $a \geq 1$ and prime $\ell > 5$, the same characteristic-zero source has maximum agreement at most $r + 2r/\ell$, and capacity gap at most $1/(2\ell)$.*

Proof. Suppose instead that $M > r + 2r/\ell$. The monic-division argument above again makes P integral at ℓ . Reduction collapses the selected roots into four buckets of multiplicities $n_c \leq r$ with $\sum_c n_c = M$. Inductively, after j Frobenius roots, write

$$\bar{F} = F_j^{\ell^j}, \quad F_j = X^{2r/\ell^j} - 2X^{r/\ell^j} + 1 - 2P_j, \quad \deg P_j < r/\ell^j.$$

For $0 \leq j < a$, the derivative is $-2P_j'$ and has root multiplicity at least $\lceil n_c/\ell^j \rceil - 1$ in each occupied bucket. If it were nonzero, its degree would therefore be at least

$$\sum_c \lceil n_c/\ell^j \rceil - 4 \geq M/\ell^j - 4 > r/\ell^j - 2,$$

where the last inequality uses $M - r > 2\ell^{a-1} \geq 2\ell^j$. This contradicts $\deg P_j' \leq r/\ell^j - 2$. The residue field is perfect, so taking an ℓ th root gives the next stage. At $j = a$, P_a is constant and agrees with W_1 at every occupied fourth root. The four values $0, 2, i, -i$ are distinct, so only one bucket is occupied, contradicting $M > r$. $\square$

Thus growing prime orders, and prime-power orders whose prime bases also grow, cannot supply a fixed positive capacity margin in this characteristic-zero source, and hence cannot supply a fixed positive margin above the first-order agreement threshold. This does not exclude powers of a fixed prime or arbitrary composite orders, nor a separate positive-characteristic source at another split prime.

More generally, the same induction with $4s$ buckets shows the following reduction, rather than an exclusion: if $r = \ell^a s$ with ℓ an odd prime, $a \geq 1$, $\gcd(\ell, s) = 1$, and $M > r + (4s - 2)\ell^{a-1}$, then a polynomial of degree below s agrees with W_s on at least $\lceil M/\ell^a \rceil$ points of μ_{4s} in the residue field of characteristic ℓ . This field is a prime field when $\ell \equiv 1 \pmod{4s}$; in general it is an extension field. The reduction preserves or increases the normalized capacity gap, but does not rule out the smaller source. Independent proof audits and the detailed reduction are recorded in `research/ordinary_ca_superlinear/PRIME_ORDER_INDEPENDENT_AUDIT.md` and `research/ordinary_ca_superlinear/PRIME_POWER_FROBENIUS_SOURCE_REDUCTION.md`.

N.5 Prime-field singleton exceptions above first order

Theorem N.8 (A prime-field quadratic line with $\Omega(n^{3/2})$ singleton exceptions). *For every integer $L \geq 25$, put*

$$n_0 = L(L-1), \quad t = n_0 + 1, \quad n = 2n_0 + 1, \quad K = 3, \quad A = 2L - 2, \quad T = 2L - 1.$$

For every sufficiently large prime p , there is an n -point domain $\mathcal{D} \subset \mathbb{F}_p$ and a received line $F + zG$, entirely over $\mathbb{F}_p$, such that

$$\text{agr}_3(F) = \text{agr}_3(G) = \text{CA}_3(F, G) = A < T.$$

Exactly

$$M = L(L(L-1) + 1) + 1 > \frac{1}{4}n^{3/2}$$

nonzero labels have agreement at least T , and each qualifying list consists of exactly one quadratic. Moreover,

$$na_1(3/n) < T < \sqrt{2n}, \quad T^2 = 2n - 1,$$

so the threshold is above first order and strictly below Johnson, with characteristic exceeding the degree cap.

An explicit sufficient field bound is as follows. Let $a_1, \dots, a_L$ be the first L primes, set $Q = a_L^2$, and take any prime

$$p > \max\{2Q, 5L^4\}.$$

Proof. Work over $\mathbb{F}_p$, using the indicated integer representatives. Define

$$P_i(U) = \frac{U^2}{a_i^2} + a_i^2 \quad (1 \leq i \leq L).$$

Their pairwise differences factor as

$$P_i - P_j = \frac{a_j^2 - a_i^2}{a_i^2 a_j^2} (U^2 - a_i^2 a_j^2).$$

The n_0 numbers $\pm a_i a_j$, for $i < j$, are distinct over the integers by unique factorization, and remain distinct modulo $p > 2Q$. At these core coordinates set

$$f(\pm a_i a_j) = a_i^2 + a_j^2, \quad g(\pm a_i a_j) = 0.$$

Each P_i has exactly A core matches. Any other quadratic has at most L core matches: each match is shared by two of the P_i , whereas its difference with each P_i has at most two roots.

Choose t fresh nonzero coordinates x_j greedily, setting

$$f(x_j) = x_j^4, \quad g(x_j) = x_j^3.$$

Require all labels

$$\lambda_{ij} = \frac{P_i(x_j) - x_j^4}{x_j^3}$$

to be distinct and outside $\{0, 1\}$. This is possible under the stated field bound. After j fresh coordinates have been chosen, avoid the n_0 core coordinates, zero, and the j previous coordinates. Avoiding labels zero and one excludes at most $8L$ additional points. Avoiding the jL previous labels

excludes at most $4jL^2$ points, since each condition $P_i(U) - U^4 - \lambda U^3 = 0$ has degree four. At one fresh coordinate, two labels could coincide only if $P_i = P_k$ there, already excluded by core avoidance. For $j \leq t - 1 = n_0$, the total number of forbidden points is at most

$$n_0 + 1 + j + 8L + 4jL^2 \leq 2n_0 + 1 + 8L + 4n_0L^2 < 5L^4 < p.$$

Thus the greedy construction supplies all t coordinates and Lt distinct labels.

For a fixed label λ , a quadratic outside the bank has at most L core matches and four fresh matches, since $H(U) - U^4 - \lambda U^3$ has degree four. Thus it has at most $L + 4 \leq A$ matches. Each bank member has exactly A core matches and can gain one fresh match only at its designated label. Consequently $f + \lambda g$ has agreement T , with a singleton list at threshold T , exactly at the Lt labels λ_{ij} . Every other label has nearest agreement A .

The direction g has the zero polynomial as its unique nearest witness, with n_0 matches. Every nonzero quadratic has at most two core matches and three fresh matches with g , hence at most five. Its list at threshold T is therefore also a singleton. Common agreement of f, g is exactly A : the pair $(P_i, 0)$ attains it on the core, while any nonzero quadratic explanation of g has at most five matches in total. For the zero explanation of g , common matches are confined to the core and number at most A .

Finally set

$$F = f, \quad G = f + g.$$

Both individual agreements equal A , because zero and one are not among the λ_{ij} , and this invertible source change preserves common agreement. For $z \neq -1$,

$$F + zG = (1 + z) \left(f + \frac{z}{1 + z} g \right).$$

The Lt original exceptional labels become the distinct nonzero labels $z = \lambda_{ij}/(1 - \lambda_{ij})$. The additional label $z = -1$ gives $-g$, whose only threshold witness is zero. These are precisely the $M = Lt + 1$ exceptions, all with singleton lists at threshold T .

The population bound follows from $n < 2L^2$ and $M > L^3 - L^2 > L^3/\sqrt{2}$ for $L \geq 25$. Also $T^2 = 2n - 1$. For the first-order comparison, the low-rate bound gives

$$na_1(3/n) \leq \sqrt{3n/2} + (3n/8)^{1/4} < \frac{7}{4}L + \sqrt{L} < 2L - 1 = T$$

when $L \geq 25$. This proves all claims. $\square$

Here the rate $3/n$ tends to zero, and the common-agreement threshold is one coordinate below T . Thus the theorem gives prime-field $\Omega(n^{3/2})$ exceptional labels above first order, with singleton qualifying lists, but does not give a fixed-rate or fixed-gap family.

The core attains the classical finite Johnson list bound. On the core domain of size $N_0 = L(L - 1)$, the maximum list size at agreement $A = 2(L - 1)$ for degree-at-most-two polynomials is exactly L , even when the received word is allowed to vary. Indeed, for any list of size ℓ , choose exactly A matching coordinates per candidate and let r_x count their selected incidences at coordinate x . Distinct quadratics coincide at at most two coordinates, so

$$\frac{\ell^2 A^2}{N_0} - \ell A \leq \sum_x r_x(r_x - 1) \leq 2\ell(\ell - 1).$$

For $L > 2$, this is the usual finite Johnson bound

$$\ell \leq \frac{N_0(A-2)}{A^2 - 2N_0} = L.$$

The displayed quadratic bank attains equality: every core coordinate has exactly two matches and every pair of bank polynomials coincides at exactly two core coordinates. This is an equality example for the classical counting bound, not a stronger decoding upper bound; the received-line construction uses additional coordinates to turn these core candidates into distinct singleton exceptional challenges.

An explicit finite first-order certificate. For Theorem N.8, write $n = 2L(L-1) + 1$ and $T = 2L - 1$, with $L \geq 25$. The derivative-weighted support of [43, Eq. (60)] with

$$m = 4, \quad B_\partial = 2, \quad B = B_{\text{jet}} = 2T, \quad H = B_Z = 48B$$

has exact coefficient count $G = 3B^2 = 24n - 12$ and local-rank bound $R = 3 + 6 + 8 + 6 = 23$. Since $n > 24$, one has $G > (47/2)n$ and hence

$$(H - B + 1)G > (47B + 1)\frac{47}{2}n > (H + 1)23n.$$

This certifies the graded challenge test of [43, Proposition 5.10]. The reconstruction guard is only $p > \max\{K - 1, B_\partial\} = 2$.

For a uniform quantitative comparison, set both incidence thresholds to $\lfloor T/2 \rfloor = (B - 2)/4$. The three ratios in [43, Lemma 5.6] are at most $3n/B$, $4n/B$, and $5n/B$. Since the message degree is two, the remaining parameters satisfy

$$F_{\text{reg}} \leq 5B, \quad S \leq 3B, \quad H_s = 144B, \quad J_1 \leq 778B^2.$$

Using $B < n$, $B^2 = 8n - 4 < 8n$, and $\Psi_2(S) \leq 24B$, the fixed-word list budget is at most $18n$, and the full-agreement-set MCA budget is at most

$$9336n^2 + 25n^2 + 6924n + 13827nB < 31000n^2.$$

Thus the prime-field lower bound of order $n^{3/2}$ belongs to an explicit finite first-order regime; no rate-dependent hidden constant or extra characteristic assumption is needed.

Corollary N.9 (Positive-density singleton bad challenges). *For arbitrarily large primes p , there are prime-field lines of length $n = \Theta(p^{2/3})$ with message dimension 3, both endpoints and common agreement equal to A , and a threshold $T = A + 1$ above the first-order curve and below the Johnson threshold, for which more than $p/36$ challenges have singleton lists at threshold T .*

Proof. Use the quadratic bank and core of Theorem N.8, with L bank polynomials, $N_0 = L(L-1)$ core coordinates, and $t = N_0 + 1$ fresh coordinates. Thus $n = 2N_0 + 1$, $A = 2L - 2$, and $T = 2L - 1$. Let a_L be the largest bank parameter and put

$$R = \max\{2a_L^2, N_0 + 8L + 2 + 8Lt\}.$$

Choose a prime $R < p < 2R$. On fresh coordinates the canonical labels are $\phi_i(x) = (P_i(x) - x^4)/x^3$. Excluding zero, the core, and all roots $\phi_i(x) \in \{0, 1\}$ leaves $m \geq p - N_0 - 1 - 8L$ allowed coordinates. Each allowed coordinate has L distinct labels. For fixed i, j, x , the equation $\phi_j(y) = \phi_i(x)$ has at

most four solutions. Consequently the allowed domain contains at most $2mL^2$ unordered colliding pairs of incidences at different coordinates.

For a uniformly random t -subset, the expected number C of such pairs is at most $2L^2t(t-1)/(m-1) < Lt/4$. Since a label with multiplicity $r \geq 2$ consumes at most $2\binom{r}{2}$ incidences, some subset has more than $Lt/2$ labels occurring exactly once. Each gives a singleton threshold list: its bank polynomial has $A+1$ matches, whereas any nonbank quadratic has at most $L+4 < T$ matches for $L \geq 6$. The exclusion of labels 0, 1 permits the same two-far-endpoint reparametrization as in the theorem, preserving these singleton labels and common agreement A .

The bank parameters may be the first L primes, so $a_L = O(L \log L)$ and eventually $R = N_0 + 8L + 2 + 8Lt \leq 9Lt$. Thus $p < 18Lt$, and $Lt/2 > p/36$. Also $p = \Theta(L^3)$ and $n = \Theta(L^2)$. The threshold comparison follows from the theorem; in particular $T^2 = 2n - 1$. The rate vanishes and the source gap is one coordinate; no fixed-rate or fixed-gap assertion is made. $\square$

Reproducible finite examples. For $L = 25$, the parameters are $n = 1201$, $K = 3$, $A = 48$, and $T = 49$. Over $\mathbb{F}_{2000003}$, an exact certificate realizes all 15026 singleton exceptions in Theorem N.8. A separate checker verifies the domain, labels, witnesses, and structural exclusion of every other quadratic; the certificate and replay are in `research/binary_characteristic_transfer/prime_L25/`. A collision-permitting domain over $\mathbb{F}_{20011}$ has 12630 finite bad labels. Including the additional projective label, there are exactly 12631 bad challenges and exactly 10798 singleton threshold lists (53.96% of the entire field). These counts use exact modular arithmetic and the same degree-based completeness argument, rather than an exhaustive scan of all p^3 quadratics. The domain and independent replay are in `research/binary_characteristic_transfer/pooled_prime_labels/`.

A larger verified domain has $L = 363$, $n = 262813$, $p = 47700761$, and threshold $T = 725$. It has exactly 31301643 bad challenges, including exactly 19622801 singleton threshold lists (41.14% of the field). The winning domain, direct polynomial replay, and bounded CPU-search receipt are archived in the same directory.

Proposition N.10 (An elliptic triple core). *Let $\ell \geq 53$ be prime and put*

$$B = \ell - 1, \quad N_0 = \frac{B(B-4)}{3}, \quad A = B - 4, \quad T = B - 3,$$

$$n = \frac{T^2 + 1}{2}, \quad t = n - N_0 = \frac{B^2 - 10B + 30}{6}.$$

For infinitely many primes p , there is an n -point domain in $\mathbb{F}_p$ and an affine received line $F + zG$, over $\mathbb{F}_p$, with

$$\text{agr}_3(F) = \text{agr}_3(G) = \text{CA}_3(F, G) = A.$$

Exactly $M = Bt + 1$ nonzero labels have agreement at least T , and each qualifying list consists of one quadratic. The threshold satisfies

$$na_1(3/n) < T < \sqrt{2n}, \quad T^2 = 2n - 1.$$

As ℓ tends to infinity,

$$\frac{M}{n^{3/2}} \longrightarrow \frac{\sqrt{2}}{3},$$

a factor $4/3$ above the leading constant of Theorem N.8.

Proof. Fix a nonsingular elliptic curve over $\mathbb{Q}$, embedded as a plane cubic with identity O , and choose a point of order ℓ over the algebraic closure. Let H be its cyclic subgroup. Use the B points $H \setminus \{O\}$ as parameter points, retaining only lines through three distinct such points whose sum is O .

For a fixed nonzero $R \in H$, a second point S ranges over $H \setminus \{O, R\}$. Exclude $S = -R$, $S = -2R$, and $S = -R/2$, which would make the third point zero or repeat one of the first two. These three exclusions are distinct because $\ell > 3$. The remaining $\ell - 5 = B - 4$ choices are paired by $S \leftrightarrow -R - S$. Thus each parameter point lies on $(B - 4)/2$ retained lines, and there are $B(B - 4)/6$ such lines. Each contains exactly three parameter points, by intersection with the nonsingular cubic.

Choose a projective line at infinity avoiding all parameter points and all intersections of two retained lines. Choose horizontal and vertical directions avoiding the retained lines as well. A rational projective coordinate change satisfying these finitely many nonvanishing conditions exists by Zariski density. In the resulting affine chart, write the parameter points as (a_i, b_i) and set

$$P_i(U) = a_i U^2 + b_i \quad (1 \leq i \leq B).$$

The retained lines now have distinct finite nonzero slopes. A retained line $b = ma + c$ supplies two evaluation coordinates $U = \pm\sqrt{-m}$ with received value c . Adjoin all these square roots to the number field of the construction. The coordinates are distinct, and each has exactly three matching bank polynomials.

This core therefore has N_0 coordinates and exactly A matches for each P_i . Any other quadratic has at most $\lfloor 2B/3 \rfloor$ core matches: each match contributes three roots to its differences with the bank, while these B nonzero differences contribute at most $2B$ roots in total.

Pass to a finite Galois number field containing the construction. Discard the finitely many primes at which a needed denominator, distinctness condition, or nonzero coefficient degenerates, including characteristic two. At every remaining completely split prime all coordinates and incidences reduce to the prime field. There are infinitely many such primes; choose one also exceeding $5B^4$.

Set $g = 0$ on the core, and denote its received word by f . Append t fresh nonzero coordinates x , setting $f(x) = x^4$ and $g(x) = x^3$. Choose them greedily so all labels

$$\lambda_{i,x} = \frac{P_i(x) - x^4}{x^3}$$

are distinct and outside $\{0, 1\}$. In addition to zero and previous coordinates, avoid every pairwise bank intersection, including those not retained in the core; there are at most $B(B - 1)$ such coordinates. After $j < t < B^2$ choices, the number of forbidden points is at most

$$B(B - 1) + 1 + j + 8B + 4jB^2 < 5B^4 < p.$$

The last two terms exclude labels zero and one and all earlier labels, using degree-four root bounds. Within a fresh coordinate, label collisions are already excluded by avoiding bank intersections.

For every λ , a nonbank quadratic matches $f + \lambda g$ at most

$$\lfloor 2B/3 \rfloor + 4 \leq B - 4 = A$$

times. Each bank member gains exactly one fresh match at each of its t designated labels and none at any other label. Thus precisely Bt labels have agreement T , with singleton threshold lists, and every other label has nearest agreement A . The zero polynomial has N_0 matches with g ; any nonzero quadratic has at most two core and three fresh matches with g . Since $N_0 > T > 5$, the list for g at threshold T consists only of zero. Also $CA_3(f, g) = A$: the pair $(P_i, 0)$ attains it, and a nonzero explanation of g has at most five matches.

As in Theorem N.8, set $F = f$ and $G = f + g$. Their individual agreements are A , and common agreement is preserved. The Bt finite exceptions transform by $z = \lambda/(1 - \lambda)$; the additional exception $z = -1$ gives $-g$, whose only threshold witness is zero. This proves the exact count and all singleton assertions.

Finally, $T^2 = 2n - 1$. Since $T \geq 49$, the low-rate estimate gives

$$na_1(3/n) \leq \sqrt{3n/2} + (3n/8)^{1/4} < \frac{7}{8}T + \frac{3}{4}\sqrt{T} < T.$$

The asymptotic constant follows from $n \sim B^2/2$ and $M \sim B^3/6$. $\square$

The elliptic core improves a leading constant, not the exponent or the source separation: the rate still tends to zero and $T - A = 1$. Its number-field specialization guarantees infinitely many suitable primes for each ℓ , rather than every sufficiently large prime, and gives no comparable effective field-size bound. The use of algebraic torsion points avoids any claim of unbounded rational torsion.

N.6 Growing source separation with quadratic witnesses

Theorem N.11. *For every sufficiently large prime p there are a domain of size n in $\mathbb{F}_p$, two received words F, G , and integers A, T such that both words and their common agreement with degree-at-most-two polynomials are exactly A , while at least $(1/e - o(1))p$ nonzero parameters z have a singleton list of polynomials agreeing with $F + zG$ on at least T coordinates. Moreover,*

$$T^2 < 2n, \quad T > na_1(3/n), \quad T - A \sim \frac{\log n}{4 \log \log n}, \quad \frac{p}{n} \sim 2e \log n \log \log n.$$

Thus the number of these parameters is superlinear in n and the absolute source separation grows, with the polynomial degree fixed at two. The rate and the relative source separation tend to zero.

Proof. We first construct the quadratic intersection bank over a field of size $O(L^2)$ or larger. Choose an odd prime $q \in [L, 2L]$ and put $b_i = 2qi + (i^2 \bmod q)$ for $0 \leq i < L$. Equality of two pair sums of the b_i forces equality of the sums of the indices, since the difference of the two remainder sums has absolute value less than $2q$. Reduction modulo q then forces equality of the unordered index pairs. Also $0 \leq b_i < 8L^2$. If $p - 1 > 32L^2$ and γ generates $\mathbb{F}_p^*$, set $a_i = \gamma^{b_i}$ and $P_i(X) = X^2/a_i^2 + a_i^2$. The $L(L - 1)$ points $\pm a_i a_j$ for $i < j$ are distinct, and the a_i^2 are distinct. Assign the value $a_i^2 + a_j^2$ at these points. Every P_i agrees on exactly $A = 2L - 2$ core points; a nonbank quadratic agrees on at most L , by counting its roots against all L bank polynomials.

Put $d = \lfloor \log p / (4 \log \log p) \rfloor$ and $L = \lfloor (p^d d!)^{1/(2d+1)} \rfloor$. Define $T = 2L - 2 + d$, $n = \lceil (T^2 + 1)/2 \rceil$, and $t = n - L(L - 1)$. Stirling's formula gives

$$t = L^2 + O(Ld + d^2), \quad \frac{L(t/p)^d}{d!} \longrightarrow 1, \quad \frac{p}{n} \sim 2e \log p \log \log p.$$

In particular the core exists and we can choose t additional distinct nonzero coordinates outside it. Set $g = 0$ on the core and $g(x) = x^3$ on the additional coordinates. Independently at each additional coordinate, choose $f(x)$ uniformly outside $\{P_i(x), P_i(x) - g(x) : 0 \leq i < L\}$; on the core use the assigned word.

Fix $\lambda \notin \{0, 1\}$. Let $X_{i,\lambda}$ count additional agreements of P_i with $f + \lambda g$, and let $Y_\lambda = \#\{i : X_{i,\lambda} \geq d\}$. For each i , at most $6(L - 1)$ coordinates forbid its desired value: the equations causing a prohibition are $P_i - P_j = \lambda X^3$ or $P_i - P_j = (\lambda - 1)X^3$. At all remaining coordinates its probability

is between $1/p$ and $1/(p - 2L)$. Different incumbents have distinct desired values. For every fixed positive integer k , selecting disjoint d -sets of agreements for k specified incumbents gives the bounds

$$\frac{(t - 6L - kd)_{kd}}{(d!)^k p^{kd}} \left(1 - \frac{k}{p - 2L}\right)^{t - kd} \leq \Pr(X_{i_1, \lambda}, \dots, X_{i_k, \lambda} \geq d) \leq \frac{(t)_{kd}}{(d!)^k (p - 2L)^{kd}}.$$

Here $(u)_v$ denotes a falling factorial. For the lower bound require exactly the selected matches and no others for these incumbents. Multiplication by $(L)_k$ shows $\mathbb{E}(Y_\lambda)_k = 1 + o(1)$ uniformly in λ . For every fixed positive integer m , the pointwise Bonferroni inequality

$$\mathbf{1}_{Y=1} \geq \sum_{j=1}^{2m} \frac{(-1)^{j-1} (Y)_j}{(j-1)!}$$

therefore gives a uniform lower limit of $\sum_{r=0}^{2m-1} (-1)^r / r!$ for $\Pr(Y_\lambda = 1)$. Letting m tend to infinity after p proves $\Pr(Y_\lambda = 1) \geq 1/e - o(1)$.

A union bound over all at most p^4 quadratic/parameter pairs shows that the probability any nonbank quadratic has $L - 1$ additional matches is at most

$$p^4 \binom{t}{L-1} (p - 2L)^{-(L-1)} = o(1).$$

Consequently some realization has no such pair and has at least $(1/e - o(1))p$ singleton threshold lists. Every nonbank polynomial then has total agreement at most A , uniformly in the parameter. The blacklist ensures that f and $f + g$ each have agreement exactly A . Any nonzero quadratic agrees with g at at most two core points and three additional points. The zero polynomial confines common agreements to the core, whose maximum is A . Thus the common agreement is exactly A .

Finally set $F = f$ and $G = f + g$. For $z \neq -1$, the substitution $\lambda = z/(1 + z)$ and multiplication by $1 + z$ preserve the lists, and the counted parameters become nonzero. At $z = -1$, the word $-g$ additionally has singleton threshold list $\{0\}$. The asserted asymptotics follow from the parameter choice. Finally $T/\sqrt{n} \rightarrow \sqrt{2}$, whereas $na_1(3/n)/\sqrt{n} \rightarrow \sqrt{3/2}$, proving the first-order inequality. $\square$

Remark N.12. The same argument permits any fixed $0 < c \leq 1/2$ with $d = \lfloor c \log p / \log \log p \rfloor$. In this case

$$T - A \sim \frac{c \log n}{\log \log n}, \quad \frac{p}{n} \sim \frac{e}{2c} (\log n)^{1/(2c)-1} \log \log n,$$

and the singleton fraction remains at least $1/e - o(1)$. In particular $c = 1/2$ gives separation asymptotic to $\log n / (2 \log \log n)$ and $\Omega(n \log \log n)$ singleton parameters. The factorial-moment bounds above require $t/p \rightarrow 0$, not $dt/p \rightarrow 0$, so this endpoint is included.

An exact finite existence certificate. For $p = 18120497$, the parameters $n = 4503001$, $K = 3$, $A = 2998$, and $T = 3001$ admit at least 4680342 singleton exceptional challenges, more than n . Both source agreements and common agreement are exactly A . Exact rational factorial-moment bounds and a nonbank union bound certify existence; this certificate does not specify a sampled received word. The independent replay is in `research/prime_quadratic_line/finite_gap3/`.

An explicit example with independent computational replay. At the same parameters, a concrete monic polynomial h of degree 13 for the fresh received values gives exactly 5127684 singleton exceptional challenges after the endpoint change, and 6437753 nonempty threshold lists in total. Both source agreements and common agreement remain 2998. Here every nonbank

quadratic has at most $1500 + 13 = 1513$ agreements, uniformly on the line, by its core incidence bound and the degree of $h + \lambda X^3 - Q$. Thus the list count reduces to the explicit bank. The saved coefficients, deterministic domain rule, and separately implemented exact replay are in [research/prime-quadratic-line/algebraic_fresh_certificate/](https://research.prime-quadratic-line/algebraic_fresh_certificate/).

N.7 Polynomial source separation by correlated covers

Theorem N.13. *Fix $0 < \alpha \leq 1/3$. For infinitely many lengths n , there are a prime $p = \Theta(n^{3(1-\alpha)/2})$, an evaluation domain in $\mathbb{F}_p$, an RS code of dimension $k = \Theta(n^\alpha)$, and words F, G whose individual and common agreements are exactly A , such that at least $(1/e - o(1))p$ nonzero parameters z have a singleton list at agreement threshold T . Here*

$$T - A = \Theta(n^\alpha), \quad T^2 < (k - 1)n, \quad T > na_1(k/n).$$

For $0 < \alpha < 1/3$ the exceptional count is superlinear in n . The code rate and the relative source separation both tend to zero.

Proof. Put $b = 2\alpha/(1 - \alpha) \in (0, 1]$. We first choose infinitely many primes p and integers s such that

$$2s \mid p - 1, \quad L = \lfloor p^{1/3} \rfloor, \quad s = \Theta(L^b), \quad s \leq L/100.$$

For completeness, the required consequence of the Bombieri–Vinogradov theorem follows by summing prime counts in the progressions $1 \bmod 2s$ for $s \in [\eta X^{b/3}, 2\eta X^{b/3}]$, where $\eta > 0$ is a sufficiently small fixed constant, and restricting primes to $[X, 2X]$. The sum of the main terms is $\gg X/\log X$, since $\sum_s 1/\varphi(2s) \geq \sum_s 1/(2s) \gg 1$. The sum of the absolute errors is $O(X/\log^2 X)$: all these moduli are below $X^{1/2}/\log^B X$ for every fixed B and sufficiently large X . Thus some such prime/divisor pair exists for every sufficiently large X . We use the prime-counting formulation in [44, Theorem 7.1].

Let $H = (\mathbb{F}_p^*)^s$. It has even order $(p - 1)/s > 32L^2$, so $-1 \in H$, and each $y \in H$ has exactly s preimages under $x \mapsto x^s$. The Sidon-exponent construction in the proof of Theorem N.11, now inside the cyclic group H , gives $a_1, \dots, a_L$ with distinct squares and distinct core nodes $\pm a_i a_j$. Put

$$P_i(Y) = Y^2/a_i^2 + a_i^2, \quad Q_i(X) = P_i(X^s), \quad k = 2s + 1.$$

Include the full fibers of the $L(L - 1)$ core nodes, setting $g = 0$ and $f = a_i^2 + a_j^2$ on the fiber of $\pm a_i a_j$. Each bank polynomial Q_i has exactly $A = 2s(L - 1)$ core matches. Every nonbank polynomial of degree at most $2s$ has at most sL core matches: each match is counted twice among its root sets against the L bank polynomials, each of which has size at most $2s$.

Choose $t = L^2 - L + 1$ other nodes $y \in H$, and include their full fibers. Independently for each fresh node, choose z_y uniformly outside $\{P_i(y), P_i(y) - y^3 : 1 \leq i \leq L\}$, and set $f(x) = z_y$, $g(x) = y^3 = x^{3s}$ on its fiber. Then

$$n = s(2L^2 - 2L + 1), \quad T = s(2L - 1) = A + s, \quad (k - 1)n - T^2 = s^2.$$

Fix a nonbank polynomial Q and a parameter λ . Agreement of Q with $f + \lambda g$ exceeding A requires at least $s(L - 2) + 1$ fresh coordinate matches, hence matches on at least $L - 1$ fresh fibers. On each fiber the probability of any match is at most $s/(p - 2L)$, since there are at most s possible values of $Q(x) - \lambda y^3$. Independence across fibers and a union bound give

$$\Pr(\exists Q \text{ nonbank}, \lambda : \text{agr}(f + \lambda g, Q) > A) \leq p^{2s+2} \binom{t}{L-1} \left(\frac{s}{p-2L} \right)^{L-1} = o(1).$$

Indeed, the logarithm of the displayed bound is at most $[6s + (b - 2 + o(1))L] \log L$; this tends to minus infinity both for $b < 1$ and for $b = 1$ with $s \leq L/100$.

Fix $\lambda \notin \{0, 1\}$. The desired bank values $P_i(y) - \lambda y^3$ are distinct at each fresh node. For each bank at most $6L$ fresh nodes forbid the desired value, by the root bound for the nonzero cubics $P_i - P_j - \lambda Y^3$ and $P_i - P_j - (\lambda - 1)Y^3$. The total number S_λ of bank-fiber hits is therefore a sum of independent Bernoulli variables, with mean

$$\mathbb{E}S_\lambda = \frac{Lt - O(L^2)}{p}(1 + o(1)) = 1 + o(1)$$

and maximum success probability at most $L/(p - 2L) = o(1)$, uniformly in λ . Consequently $\Pr(S_\lambda = 1) = 1/e + o(1)$. On the nonbank-good event, $S_\lambda = 1$ means exactly one codeword has agreement at least T , and its agreement is exactly T . Taking expectations over labels and subtracting the global failure probability shows that one realization has at least $(1/e - o(1))p$ such labels and no nonbank improvement anywhere on the line. The excluded fresh values give both sources $f, f + g$ agreement exactly A .

A nonzero polynomial V of degree at most $2s$ agrees with g at at most $2s$ core coordinates and $3s$ fresh coordinates. For large L this is less than A . The choice $V = 0$ restricts common agreements to the core, where a bank polynomial attains exactly A . Thus common agreement is exactly A as well. Set $F = f, G = f + g$ and use $\lambda = z/(1 + z)$ for $z \neq -1$; scaling by $1 + z$ preserves lists. The counted labels yield nonzero z . The additional parameter $z = -1$ gives $-g$, with unique threshold witness zero, although it is not needed for the asserted count.

Finally $n = \Theta(L^{b+2})$ and $p \sim L^3$, giving the exponents in the statement. Since $s \rightarrow \infty$, $T/\sqrt{kn/2} \rightarrow \sqrt{2}$, whereas $a_1(k/n) = \sqrt{k/(2n)}(1 + o(1))$. This proves the threshold comparison. $\square$

Remark N.14. For example, $\alpha = 1/5$ gives a source separation $\Theta(n^{1/5})$ and $\Theta(n^{6/5})$ singleton exceptional challenges. The cover multiplies lengths, source separation, and polynomial degrees together: its relative separation is exactly $1/(2L^2 - 2L + 1)$, and its rate is asymptotic to $1/L^2$. This is a polynomial absolute-gap extension of the construction, not fixed-rate or fixed-relative-gap sharpness of the general upper bounds.

N.8 An upper bound for witnesses on a conic

Theorem N.15. *Let $\mathcal{D}$ be N distinct points of a field $\mathbb{F}$, and let $1 \leq k \leq N$. For words $F, G \in \mathbb{F}^{\mathcal{D}}$, define their common agreement*

$$A = \max_{\deg P, \deg Q < k} |\{x \in \mathcal{D} : F(x) = P(x), G(x) = Q(x)\}|.$$

Fix polynomials H, U, V of degree less than k , with U, V not both zero, and a finite set $S \subseteq \mathbb{F}^$. Put*

$$Q_\theta = H + \theta U + \theta^{-1}V, \quad b = |\{x \in \mathcal{D} : U(x) = V(x) = 0\}| \leq k - 1.$$

For every integer $A < T \leq N$, the number of labels $\lambda \in \mathbb{F}$ for which some Q_θ , $\theta \in S$, agrees with $F + \lambda G$ on at least T coordinates is at most

$$\frac{2N(N-1)}{(T-b)(T-A)} \leq \frac{2N(N-1)}{(T-k+1)(T-A)}. \quad (68)$$

No singleton-list hypothesis is required.

Proof. Interpolation on any k coordinates gives $A \geq k$, so the denominators are positive. We bound qualifying pairs (θ, λ) , which also bounds labels even when bank polynomials repeat. At a coordinate with $G(x) \neq 0$, the agreement equation is the Laurent graph

$$\lambda = a_x \theta + b_x \theta^{-1} + c_x, \quad (a_x, b_x, c_x) = \frac{(U(x), V(x), H(x) - F(x))}{G(x)}.$$

Two distinct graphs meet at most twice: multiplying their difference by θ gives a nonzero polynomial of degree at most two.

Remove the set $Z = \{U = V = G = 0, F = H\}$, of size z ; these coordinates match every pair. Fix a qualifying pair. Group its remaining matches into one class with $G = 0$, classes of identical nonconstant graphs, and one class of constant graphs of size h . Matching constant graphs all have value λ , and $z + h \leq b$. Write m_i for the sizes of the first two kinds of classes and $w = \sum_i m_i$, so the total agreement is $s = z + h + w \geq T$.

Every $m_i + z \leq A$. For the zero-direction class this follows from the simultaneous witnesses $(Q_\theta, 0)$. For a nonconstant graph with $a_x = a \neq 0$, use $(H - c_x U/a, U/a)$; if $a = 0$, use $(H - c_x V/b_x, V/b_x)$. These polynomial pairs also explain Z . Hence $\sum_i m_i^2 \leq (A - z)w$. Coordinate pairs from different classes, excluding Z and pairs within the constant class, number at least

$$hw + \frac{w^2 - \sum_i m_i^2}{2} \geq \frac{w(s + h - A)}{2} \geq \frac{(T - b)(T - A)}{2}.$$

A fixed unordered coordinate pair is counted at most twice over all qualifying parameter pairs. For two nonzero-direction coordinates this is the distinct-graph intersection bound. For a zero-direction coordinate paired with a nonzero-direction coordinate, the former requires

$$U(x)\theta^2 + (H(x) - F(x))\theta + V(x) = 0.$$

This is nonzero outside Z , has at most two roots, and the other coordinate then determines λ . Pairs of two zero-direction coordinates are never counted, nor are pairs of two constant graphs. Thus the total budget is at most $2\binom{N}{2} = N(N - 1)$, proving the first inequality. The second follows from $b \leq k - 1$. $\square$

Corollary N.16. *If $T - k \geq \varepsilon N$ for $\varepsilon > 0$, an inherited conic bank supplies at most $2N/[\varepsilon(T - A)]$ qualifying labels. The bound applies after any common affine-linear transformation of the witness space that leaves the transformed witnesses in the stated RS code. In particular it covers common polynomial pullbacks and multipliers, linear re-encoding, and puncturing. It uses the actual common agreement after the transformation.*

If the transformation collapses the bank to one polynomial, the stronger bound $N/(T - A)$ follows by assigning each nonzero-direction coordinate to its unique matching label.

For ordinary shortening, either at most two bank parameters survive, or the shortening constraints hold identically in θ : multiplication by θ turns each constraint into a quadratic with at least three roots. Subtracting the prescribed interpolation polynomial and dividing by the shortening factor then preserves the conic form. Consequently these uniform transformations cannot yield superlinear inherited-bank populations at a fixed positive capacity margin. This does not constrain new witnesses outside the transformed bank.

Proposition N.17. *Fix integers $L \geq 25$ and $s \geq 1$. For every prime p satisfying*

$$2s \mid p - 1, \quad (p - 1)/s > 5L^4 + 32L^2,$$

there is an RS code and a line with

$$k = 2s + 1, \quad N = s(2L^2 - 2L + 1), \quad A = 2s(L - 1), \quad T = s(2L - 1),$$

whose two sources and common agreement are exactly A , and which has exactly $L(L^2 - L + 1) + 1$ nonzero exceptional labels at threshold T . Every exceptional threshold list is a singleton; all other threshold lists are empty. More generally, at every point on this line and every agreement threshold strictly greater than A , the list has size at most one. In the original normalization all finite exceptions are witnessed by the fixed conic bank $Q_\theta(X) = \theta + \theta^{-1}X^{2s}$. Its exact count $L(L^2 - L + 1) \sim L^3$ compares with the upper bound $(4 + o(1))L^3$ in (68), while $p/L^3 \rightarrow \infty$. Also $(k - 1)N - T^2 = s^2 > 0$, and along sequences with $L, s \rightarrow \infty$, $T > Na_1(k/N)$ eventually.

Proof. Let $J = (\mathbb{F}_p^*)^s$. This cyclic group has even order greater than $32L^2$. The Sidon-exponent construction from Theorem N.11, applied inside J , gives $a_1, \dots, a_L \in J$ with distinct squares and distinct core nodes $\pm a_i a_j$. Put $P_i(Y) = Y^2/a_i^2 + a_i^2$ and $Q_i(X) = P_i(X^s)$. Include all s preimages of every core node, assigning $f = a_i^2 + a_j^2$ and $g = 0$ on its fiber. Each bank polynomial has exactly A core matches. Every nonbank polynomial of degree at most $2s$ has at most sL core matches: each match counts twice against the bank, and each of its L differences has at most $2s$ roots.

Choose $t = L^2 - L + 1$ further nodes $y \in J$ greedily and include their full fibers. On these fibers set $f(X) = X^{4s} = y^4$ and $g(X) = X^{3s} = y^3$. Require all Lt labels

$$\lambda_{i,y} = \frac{P_i(y) - y^4}{y^3}$$

to be distinct and outside $\{0, 1\}$. At stage $j \leq L(L - 1)$, exclude core nodes, previously selected nodes, and zeros of the degree-four polynomials $P_i(Y) - Y^4 - \lambda Y^3$, for forbidden labels 0, 1 and all jL previous labels. The number excluded is at most

$$L(L - 1) + j + 8L + 4jL^2 < 5L^4.$$

Within one new node, two bank labels coincide only at a core node, already excluded. Thus the choice exists in J . Each chosen base node has exactly s distinct preimages because $s \mid p - 1$.

Every nonbank polynomial has at most $4s$ fresh matches at any label, since $X^{4s} + \lambda X^{3s} - Q(X)$ has degree exactly $4s$. Its total agreement is therefore at most $s(L + 4) \leq A$. Each of the Lt labels has exactly one bank witness, with agreement $A + s = T$; every other original label has agreement exactly A . A nonzero degree-at-most- $2s$ polynomial agrees with g on at most $2s$ core and $3s$ fresh coordinates, fewer than A . The zero polynomial confines common agreement to the core, whose maximum is A . Thus the actual common agreement is A .

Change sources to $F = f$ and $G = f + g$. Both have agreement exactly A . For $z \neq -1$, use $\lambda = z/(1 + z)$ and multiply witnesses by $1 + z$. All Lt counted labels become nonzero finite labels; the additional parameter $z = -1$ gives $-g$, with singleton threshold witness zero. This proves the exact count and the line-wide list statement. Substitution into (68) gives $(4 + o(1))L^3$. Finally $k/N \sim L^{-2}$ and, when $s \rightarrow \infty$, $T/\sqrt{kN/2} \rightarrow \sqrt{2}$, while $a_1(k/N) \sim \sqrt{k/(2N)}$. $\square$

Dirichlet's theorem supplies arbitrarily large eligible primes for every prescribed L, s . No upper bound on their size, nor a constant exceptional fraction, is claimed. For example, taking $s = \Theta(L^b)$ gives gap $\Theta(N^{b/(b+2)})$ and count $\Theta(N^{3/(b+2)})$. The upper-bound comparison is in the *original* normalization: after the endpoint change the witnesses are parameter-dependent scalar multiples of the bank, and the count is transported by the parameter bijection, with one additional label. This is restricted conic-bank order sharpness, not general first-order tightness or a fixed-rate lower bound.

Nor do singleton lists on this line imply small lists for the whole code. On the same domain, keep the core word and assign a different fresh fiber to each of the L bank polynomials, giving that fiber the polynomial's value. Since $t \geq L$, all bank polynomials then reach $A + s = T$. The code therefore has a threshold list of size at least L at another center. The line-wide singleton assertion does not replace a code-wide list-size hypothesis in a general proximity theorem.

N.9 A low-rate bridge to the first-order regime

Proposition N.18 (A low-rate large-characteristic line with singleton nearest lists). *Let $p \geq 5$ be prime, let $\mathbb{B} = \mathbb{F}_{p^5}$, and let $\mathbb{F} = \mathbb{F}_{p^{15}}$. Choose $\theta \in \mathbb{F} \setminus \mathbb{B}$; then $1, \theta, \theta^2$ are linearly independent over $\mathbb{B}$. On the evaluation domain $\mathbb{B}$, put*

$$N = p^5, \quad K = p, \quad f = X^{p^4-1} + \theta X^{p^3-1} + \theta^2 X^{p^2-1}, \quad g = X^{p^2-1}.$$

For degree- $< K$ witnesses, both individual source agreements and their common agreement equal $C = p^2 - 1$. Exactly

$$M = \begin{bmatrix} 5 \\ 2 \end{bmatrix}_p = p^6 + p^5 + 2p^4 + 2p^3 + 2p^2 + p + 1$$

nonzero challenges $z \in \mathbb{F}$ have nearest agreement $A = p^3 - 1$. Exactly M further nonzero challenges have nearest agreement p^2 ; all other challenges have nearest agreement $C = p^2 - 1$. At every challenge in either of the first two classes, the list at threshold p^2 consists of exactly one polynomial. Consequently, at every integer threshold $p^2 < T \leq A$, the qualifying labels are exactly the first M labels and each qualifying list is a singleton.

For $p \geq 53$, exactly these $M > N^{6/5}$ challenges are bad at the lower advertised threshold $A_0 = p^3 - p^2$. This threshold is below $\sqrt{N(K-1)}$ and above the first-order curve $Na_1(K/N)$. An explicit first-order certificate uses

$$m = 4, \quad B_\partial = 2, \quad B_{\text{jet}} = 4p^2, \quad B_Z = 192p^2.$$

In particular its counting characteristic guard $p > \max\{K-1, B_\partial\}$ holds. The finite certificate and squarefree counting formulas of [43, Proposition 5.10 and Lemma 5.6] give the uniform upper bound $68N$ for the list at threshold A_0 of every received word on this evaluation set. For every affine received line, they give at most $300000N^2$ exceptional challenges for full-agreement-set MCA recovery at A_0 .

Proof. For every three-dimensional $\mathbb{F}_p$ -subspace $W \subseteq \mathbb{B}$, write its monic locator as

$$L_W = X^{p^3} + aX^{p^2} + cX^p + vX.$$

Its linear coefficient v is nonzero. Define

$$z_W = c^p - a^{p+1} + \theta a - \theta^2, \quad h_W = -((\theta - a^p)c + v^p)X^{p-1} - (\theta - a^p)v.$$

The locator compiler identity is

$$X(f + z_W g - h_W) = L_W^p + (\theta - a^p)L_W.$$

On $\mathbb{B}$ the second factor in $L_W(L_W^{p-1} + \theta - a^p)$ never vanishes. Since $h_W(0) \neq 0$, the agreement set is exactly $W \setminus \{0\}$. The coordinates of z_W in $1, \theta, \theta^2$ determine a, c . Two such locators with the same

a, c differ only by a multiple of X ; their subspaces intersect in a nonzero vector, so they coincide. Thus the displayed M labels are distinct and nonzero.

We prove exhaustiveness already at threshold p^2 , including arbitrary witness coefficients in $\mathbb{F}$. Projecting an agreement equation onto its θ^2 coordinate gives a polynomial of degree at most $p^2 - 1$ with at least p^2 roots. Hence the θ^2 coordinates of z, h are respectively $-1, 0$. Write $z = b + \theta a - \theta^2$ and $h = h_0 + \theta h_1$, with coefficients in $\mathbb{B}$, and set

$$L = X^{p^3} + aX^{p^2} - Xh_1, \quad T = X^{p^4} + bX^{p^2} - Xh_0.$$

The polynomial

$$\frac{T - L^p + a^p L}{X} = (b + a^{p+1})X^{p^2-1} + X^{p-1}h_1^p - h_0 - a^p h_1$$

has degree at most $p^2 - 1$ and vanishes at every agreement coordinate. At a matched zero this follows from $h_0(0) = h_1(0) = 0$, so division by X loses no required root. It is identically zero. The distinct exponents $p - 1 + pj$ force h_1 to have only degrees $0, p - 1$. Coefficient comparison gives precisely

$$L = X^{p^3} + aX^{p^2} + cX^p + vX, \quad z = c^p - a^{p+1} + \theta a - \theta^2,$$

and the displayed formula for h_W , whether or not L is a three-space locator.

If $v \neq 0$, zero is not matched and the agreement count is $|\ker_{\mathbb{B}} L| - 1$. At least p^2 agreements therefore forces L to be a three-space locator, giving the first class. If $v = 0$, then $L = J^p$, where

$$J = X^{p^2} + a^{1/p}X^p + c^{1/p}X.$$

Zero is now matched, and at least p^2 agreements forces J to be the locator of a two-dimensional subspace U . Each such U supplies a distinct secondary label, with exactly p^2 agreements; there are $\binom{5}{2}_p = M$ of them.

The two label classes are disjoint. Otherwise their identical a, c would give $L_W = J_U^p + vX$ with $v \neq 0$. Their kernels intersect trivially, so $\mathbb{B} = U \oplus W$. The $\mathbb{F}_p$ -linear map $x \mapsto -J_U(x)^p/v$ is the projection onto W along U , of operator trace 3. But its linearized polynomial has no X term. For a map $\sum_{i=0}^4 a_i x^{p^i}$ on $\mathbb{B}$, the operator trace is $\text{Tr}_{\mathbb{B}/\mathbb{F}_p}(a_0)$: after extending scalars, multiplication is diagonal in the five embedding coordinates, while every nonidentity Frobenius power permutes these coordinates without a fixed point. Thus this projection has trace zero, a contradiction since $p \neq 3$. The coefficient comparison and injectivity show that each label in either class has a singleton list at threshold p^2 .

The source upper bounds follow by projecting f onto its θ^2 coordinate and by using the degree of g : either source has at most $p^2 - 1$ agreements with a strict witness. For the matching common lower bound, choose any two-dimensional subspace U . Remainders of the additive heads modulo L_U have degree at most p ; division by X gives simultaneous degree- $< p$ witnesses on $U \setminus \{0\}$. Hence both individual agreements and common agreement equal C . The same simultaneous witnesses give every word on the line at least C agreements. Exhaustiveness above therefore gives nearest agreement exactly C at all remaining challenges, completing the three-level classification.

Assume now $p \geq 53$. For completeness, the stated finite first-order comparison involves no characteristic-dependent hidden constant. Write $d = p - 1$ and $B = 4p^2$. The derivative-weighted support

$$X^x Y_0^u Y_1^v, \quad v \leq 2, \quad u + v \leq B, \quad x + du + (d - 1)v < 4A_0$$

has

$$G = 24p^5 - 24p^4 - 6p^3 + 18p^2 + p - 3 > \frac{47}{2}N$$

coefficients. Its local-rank bound is 23, from Taylor layers 3, 6, 8, 6. With $H = 48B$, the graded challenge test follows from

$$(H - B + 1)G > (47B + 1)\frac{47}{2}N > (H + 1)23N.$$

The finite squarefree counting formulas of [43, Eqs. (38), (54), and (55)] apply with derivative cap two, even though $B > p$. To record a reproducible coarse budget, their parameters satisfy

$$\lambda < 2p^2, \quad F_{\text{reg}} \leq 28p^3, \quad S < 12p^2, \quad H_s = 576p^2, \quad J_1 \leq 30776p^6.$$

Choose both incidence thresholds $\lfloor A_0/2 \rfloor$. The two regular terms are at most $184656p^{10}$ and $112p^{10}$; the ordinary tail is at most $82956p^7 + 15588p^4$. These give the asserted bounds. Finally $A_0^2 < N(K - 1)$, while the low-rate curve estimate

$$a_1(p^{-4}) \leq \frac{p^{-2}}{\sqrt{2}} + \frac{p^{-3}}{2^{3/4}\sqrt{3}} < p^{-2}(1 - 1/p)$$

holds for $p \geq 53$. □

Scope. The baseline population is an explicit specialization of the existing fixed-characteristic locator compiler of [42]; the refinement is its exhaustive singleton-list classification. The ambient field is an extension field, the rate tends to zero, and the absolute first-order margin tends to zero. Singleton classification holds throughout $p^2 < T \leq A$, including the below-Johnson advertised threshold A_0 . This is a statement about the displayed received line, not a bound on lists of arbitrary received words. The ratio $A_0/(Na_1(K/N))$ tends to $\sqrt{2}$, but this does not establish fixed-rate or fixed-margin tightness.

N.10 Projectivization gives a quadratic code and exponent three halves

Theorem N.19 (A large-characteristic quadratic line with $\Omega(N^{3/2})$ exceptional labels). *Let $p \geq 5$ be prime, put $\mathbb{B} = \mathbb{F}_{p^5}$ and $\mathbb{F} = \mathbb{F}_{p^{15}}$, and choose $\theta \in \mathbb{F} \setminus \mathbb{B}$. Set*

$$n = p^4 + p^3 + p^2 + p + 1, \quad N = 2n, \quad \mathcal{D} = \mu_N \subset \mathbb{B}^*, \quad K = 3,$$

and, on $\mathcal{D}$, define

$$E_j(T) = T^{2(p^j - 1)/(p - 1)}, \quad f = E_4 + \theta E_3 + \theta^2 E_2, \quad g = E_2.$$

For degree- < 3 witnesses, both individual source agreements and common agreement are exactly

$$C = 2(p + 1).$$

Exactly

$$M = \begin{bmatrix} 5 \\ 2 \end{bmatrix}_p = p^6 + p^5 + 2p^4 + 2p^3 + 2p^2 + p + 1 > \frac{1}{4}N^{3/2}$$

nonzero labels $z \in \mathbb{F}$ have nearest agreement $A = 2(p^2 + p + 1)$. Each has a unique nearest quadratic; every other label has nearest agreement exactly C . In particular, at every integer threshold $C < T_0 \leq A$, precisely these M labels qualify and their lists are singletons.

The advertised threshold $A_0 = 2p^2 + p$ satisfies

$$Na_1(K/N) < A_0 < \sqrt{N(K - 1)},$$

where a_1 is the first-order curve. The characteristic guard $p > K - 1$ holds.

Proof. Since $N \mid p^5 - 1$, the stated evaluation domain exists. Put $\ell = (p - 1)/2$. The map

$$X \mapsto T = X^\ell$$

is an ℓ -to-one surjection $\mathbb{B}^* \rightarrow \mathcal{D}$, and $E_j(X^\ell) = X^{p^j-1}$. A quadratic in T pulls back to a polynomial of degree at most $p - 1$.

For a three-dimensional $\mathbb{F}_p$ -subspace $W \leq \mathbb{B}$, write its monic locator as

$$L_W = X^{p^3} + aX^{p^2} + bX^p + cX, \quad c \neq 0.$$

The additive-locator identity gives the label and witness

$$z_W = b^p - a^{p+1} + \theta a - \theta^2, \quad h_W = (a^p b - c^p - \theta b)T^2 + a^p c - \theta c.$$

Before the map $X \mapsto X^\ell$, the exact support is $W \setminus \{0\}$. Therefore its size after the map is $(p^3 - 1)/\ell = A$. The labels are nonzero by independence of $1, \theta, \theta^2$. Equality of two labels first fixes a, b ; the corresponding locators then differ by a multiple of X and share a nonzero root, so coincide. This gives exactly the stated M distinct constructed labels.

For completeness, use the all-witness locator classification in the proof of Proposition N.18: for every odd p , any degree-at-most- $p - 1$ witness with more than p^2 agreements for the original sources is the unique canonical witness of a three-space W . If an arbitrary quadratic $h(T)$ has more than C matches here, its pullback has at least

$$\ell(C + 1) = \frac{p - 1}{2}(2p + 3) > p^2$$

matches in $\mathbb{B}^*$. Thus the classification applies, forcing $h(X^\ell) = uX^{p-1} + v$ and hence $h(T) = uT^2 + v$. It recovers precisely z_W, h_W . This proves both exhaustiveness and uniqueness above C , including witnesses with a priori nonzero odd term.

The original simultaneous explanations on every two-dimensional subspace are of the form $uX^{p-1} + v$: reduce the additive heads modulo its locator and divide by X . They therefore descend to quadratics and give simultaneous agreement $(p^2 - 1)/\ell = C$ here. The degree of g is C , so its agreement cannot exceed C . For f , project any proposed quadratic witness onto the θ^2 component over $\mathbb{B}$. More than C matches would make g equal to a quadratic, impossible. Consequently both individual agreements and common agreement equal C , and every nonexceptional label has agreement exactly C .

For the quantitative claims, $n < p^4 p / (p - 1) \leq (5/4)p^4$, so $N < (5/2)p^4$ and $N^{3/2} < 4p^6 < 4M$. The exact Johnson deficit is

$$2N - A_0^2 = 3p^2 + 4p + 4 > 0.$$

The rate $\rho = 3/(2n)$ lies in the low-rate branch. Its standard bound $a_1(\rho) \leq \sqrt{\rho/2} + (\rho/2)^{3/4}/\sqrt{3}$ gives

$$Na_1(\rho) \leq \sqrt{3n} + \frac{(3n)^{1/4}}{\sqrt{2}} < 2p^2 + p = A_0,$$

since $3n < (15/4)p^4 < 4p^4$. Finally $C < A_0 < A$, so the exact classification applies at this below-Johnson threshold. $\square$

This family improves the population exponent of the unprojectivized large-characteristic example to $3/2$, while retaining singleton nearest lists. Its rate tends to zero and its absolute proximity gap vanishes. The evaluation and challenge fields are extension fields; this is not a fixed-rate or prime-field construction.

Uniform finite first-order comparison. The characteristic hypothesis can be checked without rate-dependent hidden constants. For the parameters of Theorem N.19, take the derivative-weighted support of [43, Eq. (60)] with

$$m = 4, \quad B_\partial = 2, \quad B = B_{\text{jet}} = 2A_0, \quad H = B_Z = 116B.$$

Its coefficient count is $G = 3B^2$ and its local-rank bound is $R = 3 + 6 + 8 + 6 = 23$. Direct expansion gives

$$5G - 116N = 8(p - 5)^4 + 168(p - 5)^3 + 1148(p - 5)^2 + 2648(p - 5) + 308 > 0.$$

Thus $(H - B + 1)G > (H + 1)23N$, which suffices for the graded challenge test of [43, Proposition 5.10]. Reconstruction requires only $p > \max\{2, B_\partial\}$.

For clarity, the resulting bounds are uniform over every received word and every affine received line on this evaluation set. In the notation of [43, Lemma 5.6], one has

$$\lambda \leq 3N/B, \quad F_{\text{reg}} \leq 5B, \quad S \leq 3B, \quad H_s = 348B, \quad J_1 \leq 1866B^2.$$

Choose both incidence thresholds $L = L_0 = \lfloor A_0/2 \rfloor$. The remaining incidence ratios are at most $4N/B$ and $5N/B$. Using $B < N$, $B^2 < 8N$, and $\Psi_2(S) \leq 24B$, the list budget is at most $18N$, and the full-agreement-set MCA exceptional budget is at most

$$22392N^2 + 25N^2 + 16716N + 33411NB < 75000N^2.$$

The lower bound $M > \frac{1}{4}N^{3/2}$ therefore lies in the same explicit large-characteristic first-order setting, though it does not match this quadratic upper bound.

N.11 A Frobenius quadratic bank in characteristic comparable to the square root of length

Proposition N.20. *For every prime $p \geq 41$, there is a length- $N = 9p^2$, dimension-three Reed–Solomon code over $E = \mathbb{F}_{p^4}$ and two received words whose individual and common agreements are all exactly $2p$, with the following property. At threshold $T = 4p$, their affine linear combination has exactly*

$$M = (p + 1)(p - 1)^2$$

exceptional parameters with singleton lists, and one further exceptional parameter with a list of size two. All other parameters have empty lists. The threshold is above the DKT first-order curve and below the exact Johnson threshold. In particular $M = \Theta(N^{3/2})$, the source and common agreement gaps are exactly $2p = \Theta(\sqrt{N})$, and $p = \sqrt{N}/3$ exceeds the message degree.

Proof. Put $B = \mathbb{F}_{p^2}$ and choose $s \in E^*$ such that $\eta = s^2 \notin B$. Let

$$D_0 = \{x \neq 0 : x^2 \in B^*\}, \quad D_1 = sD_0.$$

These are disjoint sets of size $2(p^2 - 1)$. For each $a \in B$ with $a^{p+1} = 1$, let

$$I_a = \{y^p - ay : y \in B\}.$$

The I_a are precisely the $p + 1$ one-dimensional $\mathbb{F}_p$ subspaces of B . On D_0 set $f(x) = x^{2p}$ and $g(x) = 0$; on D_1 set $f(x) = \eta(x^2/\eta)^p$ and $g(x) = 1$.

We use the elementary quadratic classification of the Frobenius word. On D_0 , its canonical quadratics are $Q_{a,b} = aX^2 + b$, with $a^{p+1} = 1$ and $b \in I_a$. They have $2p$ matches if $b \neq 0$, and $2p - 2$

if $b = 0$. Every noncanonical quadratic, including quadratics with coefficients in E , has at most p matches for $p \geq 5$. For completeness, write an arbitrary quadratic as $aX^2 + cX + b$. The core is the union of two B -lines uB^* , where $u = 1$ or $u = v$ with v^2 a nonsquare in B . On either line, put $x = ut$ and

$$(A, C, B_0) = (au^{2-2p}, cu^{1-2p}, bu^{-2p}).$$

The matching equation is $t^{2p} = At^2 + Ct + B_0$. If these three coefficients are not all in B , a B -linear functional on E vanishing on B gives a nonzero polynomial of degree at most two vanishing at every match on this component. If neither component has all three coefficients in B , there are therefore at most four matches. If one component does, then $a, b \in B$ and $cu \in B$. When $c \neq 0$, the other component has no matches, because its nonzero linear term lies outside B , whereas all its other terms lie in B . On the coefficientwise B -valued component, Frobenius and substitution give

$$C^p t^p = H(t), \quad H(t) = (1 - A^{p+1})t^2 - A^p Ct - A^p B_0 - B_0^p.$$

Every match is consequently a zero of

$$H(t)^2 - C^{2p}(At^2 + Ct + B_0).$$

If this polynomial is nonzero, it has at most four roots. If it is zero, H has degree at most one, and the nonzero polynomial $C^p t^p - H(t)$ has at most p roots. Finally, if $c = 0$ and a component is coefficientwise B -valued, then $a, b \in B$ on both components. The substitution $y = x^2$ reduces the entire core to $y^p - ay = b$ on B^* , with two preimages per nonzero y . A norm-one a has kernel size p and image I_a ; a non-norm-one a gives an invertible map. This proves both the canonical counts and the asserted bound $\max(p, 4) = p$ for all other quadratics.

For the word $f + \lambda g$, the same classification on D_1 , after $x = sz$ and division by η , makes $Q_{a,b}$ canonical precisely when

$$b - \lambda = \eta v, \quad v \in I_a.$$

Thus a quadratic canonical on both blocks has parameter $\lambda = b - \eta v$ with $b, v \in I_a$. Since $1, \eta$ are linearly independent over B , the planes $I_a + \eta I_a$ intersect pairwise only at zero. Every nonzero parameter in their union determines a unique a, b, v . If both b and v are nonzero, the corresponding quadratic has exactly $4p$ matches; there are M such distinct parameters. If exactly one is zero it has $4p - 2$ matches, and at parameter zero all doubly canonical quadratics have $4p - 4$ matches. A quadratic canonical on just one block has zero matches on the other: its norm-one coefficient is already fixed, and the other constant term fails the relevant image condition. A quadratic canonical on neither block has at most $2p$ matches in total.

We also need a sharper bound for a quadratic canonical on neither block. The four physical B -lines forming $D_0 \cup D_1$ are distinct, since $\eta \notin B$. Write the quadratic as $aX^2 + cX + b$. On a component uB^* of block j , with $\eta_0 = 1$ and $\eta_1 = \eta$, normalizing the received leading coefficient shows that coefficientwise B -valuedness requires $cu/\eta_j \in B$, or equivalently $c \in uB$. If $c \neq 0$, this can hold on at most one of the four lines. The block containing that line has at most p matches by the noncanonical bound; the other block has at most four by projection. If $c = 0$ and $a \in B$, the linearized equations give at most two matches per block unless it is canonical there. If $c = 0$ and $a \notin B$, projection gives at most four per block. Thus every quadratic canonical on neither block has at most $\max(p + 4, 8) = p + 4$ active matches for $p \geq 5$.

Choose distinct α, β outside the union of the parameter planes, whose size is $1 + (p+1)(p^2 - 1) < p^4 - 1$. Append $5p^2 + 4$ distinct points outside $D_0 \cup D_1$, avoiding all roots of $X^3 - Q$ for the following three families: the $p(p+1)$ quadratics canonical on D_0 , and the $p(p+1)$ quadratics canonical on

D_1 at each of the two parameters α, β . Put $f(x) = x^3$, $g(x) = 0$ at these new points. There are at most $9p(p+1)$ forbidden roots, and the required padding exists since

$$p^4 \geq 18p^2 + 9p$$

for $p \geq 5$. Every D_0 -canonical quadratic gains zero matches, and any other quadratic gains at most three. Consequently the only parameters with agreement at least $4p$ are precisely the M parameters above, each with its unique canonical witness: a quadratic canonical only on D_1 has at most $2p+3 < 4p$ matches, and one canonical on neither block has at most $p+7 < 4p$.

At each of α, β , a quadratic canonical on either block gains zero padding matches by construction and has at most $2p$ active matches. Every other quadratic has at most $p+7 \leq 2p$ matches for $p \geq 7$. A D_0 -canonical quadratic $Q_{a,b}$ with $b \neq 0$ attains $2p$ matches. Thus the two sources $r = f + \alpha g$ and $s_0 = f + \beta g$ each have agreement exactly $2p$. Their common agreement is at most either individual agreement, and the same $Q_{a,b}$ explains both words on its $2p$ matching core coordinates, so common agreement is also exactly $2p$. For $z \neq -1$,

$$r + zs_0 = (1+z) \left(f + \frac{\alpha + \beta z}{1+z} g \right).$$

This transports all M singleton exceptions bijectively. The omitted parameter β is not exceptional. At $z = -1$ the word is $(\alpha - \beta)g$: its list at threshold $4p$ consists exactly of the two constant polynomials 0 and $\alpha - \beta$, since every other quadratic agrees on at most four points.

Finally the exact Johnson threshold is $\sqrt{2N} = \sqrt{18}p > 4p$. The finite first-order estimate gives

$$Na_1(3/N) \leq \sqrt{3N/2} + (3N/8)^{1/4} < \frac{15}{4}p + \frac{3}{2}\sqrt{p} < 4p$$

for $p \geq 41$. All stated conclusions follow. $\square$

Compared with Theorem N.19, this retains an exceptional population of order $N^{3/2}$ and an agreement gap of order $\sqrt{N}$, while increasing the characteristic from order $N^{1/4}$ to order $\sqrt{N}$. The ambient field here has size $p^4 = N^2/81$, and the exceptional probability is of order $N^{-1/2}$. The rate and fractional agreement gap still tend to zero. This is an extension-field result, not a prime-alphabet construction or a tightness result for the general DKT bound. The threshold comparison above does not reuse the preceding construction's finite interpolation certificate.

Uniform finite first-order certificate. For the scaled-fiber parameters $N = 9p^2$, $K = 3$, and $A = 4p$, there is also an explicit finite certificate, independent of the limiting rate curve. In the derivative-weighted support of [43, Eq. (60)], take

$$m = 6, \quad B_\partial = 3, \quad B = B_{\text{jet}} = 12p, \quad H = B_Z = 40B.$$

Thus the support consists of $X^x Y_0^a Y_1^b$ with $b \leq 3$, $a + b \leq B$, and $x + 2a + b < 24p$. The exact coefficient count in [43, Eq. (61)] is

$$G = \sum_{b=0}^3 \sum_{t=b}^B (2B - 2t + b) = \sum_{b=0}^3 B(B - b + 1) = 4B^2 - 2B = 576p^2 - 24p.$$

The six local-rank contributions in [43, Eq. (62)] are 4, 8, 12, 15, 14, 9, so $R = 62$. Every column has remaining challenge budget at least $H - B$, and every local row has budget at most H . Consequently the graded challenge test of [43, Proposition 5.10] follows from

$$\begin{aligned} (H - B + 1)G - (H + 1)NR &= B(144p^2 - 936p) + 18p^2 - 24p \\ &> 0 \quad (p \geq 7). \end{aligned}$$

In particular $G > NR$. This constructs a sound first-order interpolant, nonzero at every challenge specialization, uniformly over all evaluation sets of this size and all affine received lines. Its actual derivative variable degree is at most three, so the reconstruction and counting hypothesis is only

$$\text{char}(\mathbb{F}) = p > \max\{K - 1, B_\partial\} = 3.$$

Thus the stated family satisfies an explicit finite first-order test; no characteristic bound depending on the growing total jet cap is needed. The corresponding uniform list and full-agreement-set MCA bounds follow from the counting transfers of [43, Lemma 5.6]; we do not optimize their numerical constants here.

N.12 Native cubic-extension coverage with small common agreement

Theorem N.21. *For every prime $p \geq 4099$, there is a dimension-three Reed–Solomon code on $N = 2p^2 - p - 1$ distinct points of $E = \mathbb{F}_{p^3}$ and a received line $f + \lambda g$ with ordinary common agreement at most $p + 2\sqrt{p}$, such that every $\lambda \in E$ has a nearby codeword at agreement threshold*

$$A = 2p - \lceil 4\sqrt{p} \rceil - 2.$$

Exactly $p^3 - p$ parameters have singleton lists, and the remaining p parameters each have a list of exactly p codewords. The threshold satisfies

$$Na_1(3/N) < A < \sqrt{2N}.$$

Thus every parameter is exceptional for ordinary common agreement, and the singleton fraction is $1 - p^{-2}$. The common-agreement gap is $\Theta(p) = \Theta(\sqrt{N})$ and the characteristic exceeds the message degree. This is an extension-field result at vanishing rate; the two affine endpoints are themselves nearby. Equivalently, the endpoints are f and $f + g$, with affine mixtures $(1 - \lambda)f + \lambda(f + g)$; this invertible source change preserves ordinary common agreement.

Proof. Choose a two-dimensional $\mathbb{F}_p$ -subspace $W \subset E$ and a square $\eta = s^2 \in E \setminus \mathbb{F}_p$. Put

$$D_0 = \{x \neq 0 : x^2 \in W\}, \quad D_1 = \{x \neq 0 : x^2 \in \eta W\} \setminus D_0.$$

The multiplicative stabilizer of W is $\mathbb{F}_p^*$; together with zero it is a subfield of E , and it cannot be E . Consequently $W \cap \eta^{-1}W$ is a line; write it as $\mathbb{F}_p u_0$. Since $\eta \notin \mathbb{F}_p$,

$$W = u_0 \text{span}_{\mathbb{F}_p}(1, \eta).$$

The quadratic character of E restricts to the quadratic character of $\mathbb{F}_p$. Each nonzero $\mathbb{F}_p$ -line therefore has exactly half its elements square in E . It follows that $|D_0| = p^2 - 1$, the overlap of the two untrimmed blocks has size $p - 1$, and $|D_0 \cup D_1| = N$. Define

$$(f, g) = (x^{2p}, 0) \quad \text{on } D_0, \quad (f, g) = (\eta(x^2/\eta)^p, 1) \quad \text{on } D_1.$$

Canonical fibers and their sizes. For a projective direction $[u]$ in W , let

$$a_u = u^{p-1}, \quad L_u(y) = y^p - a_u y, \quad I_u = L_u(W).$$

The kernel is $\mathbb{F}_p u$, so I_u is a line. For a basis w_1, w_2 of W , put $\delta = w_1 w_2^p - w_1^p w_2 \neq 0$. The alternating identity $u L_u(y) = u y^p - u^p y$ gives $I_u = (\delta/u)\mathbb{F}_p$. The even quadratic $Q = a_u X^2 + b$ has canonical fibers on both blocks precisely when

$$b, v \in I_u, \quad \lambda = b - \eta v.$$

A zero fiber has exactly $p - 1$ physical square-root points. A nonzero affine fiber $y_0 + u\mathbb{F}_p$ has between $p - 2\sqrt{p}$ and $p + 2\sqrt{p}$ such points. Indeed, $y_0/u \notin \mathbb{F}_p$ has degree three, so $\text{Norm}_{E/\mathbb{F}_p}(t + y_0/u)$ is a separable cubic. The identity between quadratic characters and the norm reduces the count to p plus a signed quadratic character sum of this cubic. Hasse's bound for its smooth genus-one double cover bounds that sum by $2\sqrt{p}$. Multiplication by the square η preserves the counts.

In the normalized fresh variable $y = x^2/\eta$, the deleted overlap is $\mathbb{F}_p u_0 \setminus \{0\}$. A canonical fresh fiber is entirely deleted only for $[u] = [u_0]$ and $v = 0$; every other fiber loses at most two physical points. Hence every doubly canonical candidate except this deleted-fiber case has at least A matches. The deleted-fiber case has at most $p + 2\sqrt{p} < A$.

Parameter geometry. For nonzero λ , its representing directions are the projective points of

$$W \cap (\lambda/\delta)^{-1} \text{span}_{\mathbb{F}_p}(1, \eta).$$

This intersection has dimension one or two. It has dimension two exactly when $\lambda \in (\delta/u_0)\mathbb{F}_p^*$: the stabilizer calculation above shows there are precisely these $p - 1$ scalars. Thus outside the line $\Lambda_0 = (\delta/u_0)\mathbb{F}_p$, every parameter has exactly one doubly canonical quadratic. Its fresh fiber cannot be the deleted one. Each nonzero parameter in Λ_0 has all $p + 1$ directions; exactly its u_0 direction has $v = 0$ and is excluded, leaving p qualifying quadratics. At zero, all directions have $b = v = 0$; again exactly u_0 is excluded. The other p directions give $2p - 2 \geq A$ matches.

Excluding all other quadratics. First consider one full block $\{x \neq 0 : x^2 \in W\}$ with word x^{2p} . An even quadratic $aX^2 + c$ gives $y^p - ay = c$ on W . Unless its kernel direction lies in W , this has at most one solution in W , hence at most two physical matches. If its kernel direction lies in W , its nonempty fibers are exactly the canonical ones above. In particular every even quadratic has at most $p + 2\sqrt{p}$ matches.

We claim every non-even quadratic has at most eight matches on this block. Write $Q = aX^2 + bX + c$, $b \neq 0$. If $Q = (uX + w)^2$ over E , then $u, w \neq 0$ and the two Frobenius equations are $x^p = \pm(ux + w)$. Since $\text{Norm}(-u) = -\text{Norm}(u)$, at most one has an affine $\mathbb{F}_p$ -line of solutions; the other has at most one. The affine line cannot contain zero. Write it as $z + tv$ with $z/v \notin \mathbb{F}_p$. If more than two of its points had squares in W , projecting $(z + tv)^2$ to the one-dimensional quotient E/W would give an identically zero quadratic in t . This would put z^2, zv, v^2 in W , impossible because $1, z/v, (z/v)^2$ are independent over $\mathbb{F}_p$. Thus the square case has at most three matches.

If Q is not a square in $E[X]$, put $a_i = a^{p^i}, b_i = b^{p^i}, c_i = c^{p^i}$ and

$$R = X^2 - a_2 a_1 Q - a_2 c_1 - c_2, \quad d = a_2 b_1,$$

$$F = R^2 + (d^2 - b_2^2 a_1)Q - b_2^2 c_1, \quad G = -2dR - b_2^2 b_1.$$

At a match, writing $y = x^p, z = x^{p^2}$ and using $y^2 = Q(x), z^2 = Q^{(p)}(y), x^2 = Q^{(p^2)}(z)$ gives $b_2 z = R(x) - dy$ and then $F(x) + G(x)y = 0$. Therefore the match is a root of $F^2 - G^2 Q$, of degree at most eight. This polynomial is nonzero. Otherwise nonsquareness of Q forces $F = G = 0$. If $a = 0$, however, $G = -b_2^2 b_1 \neq 0$; and if $a \neq 0$, $G = 0$ forces R constant, contrary to its nonzero linear coefficient $-a_2 a_1 b$. This proves the claim.

The same classification applies on the second block after $x = st$ and division by η , with the constant shifted by λ . Non-even quadratics therefore have at most sixteen active matches. An even quadratic canonical on only one block has zero matches on the other and at most $p + 2\sqrt{p}$ in all; an even quadratic canonical on neither has at most four. Since $A > \max(16, p + 2\sqrt{p})$, the parameter geometry lists every qualifying codeword, proving the exact list counts.

Ordinary common agreement. A direction quadratic other than the constants zero and one matches g on at most four points. Direction zero restricts simultaneous agreements to D_0 , and

direction one to D_1 . The one-block classification bounds either by $\max(p + 2\sqrt{p}, 8) = p + 2\sqrt{p}$. This proves the common-agreement bound. The endpoints f and $f + g$ are themselves nearby, as follows already from the all-parameter claim.

Thresholds and finite applicability. The stated onset implies $A \geq 2p - 5\sqrt{p}$ and

$$Na_1(3/N) \leq \sqrt{3N/2} + (3N/8)^{1/4} < \frac{7}{4}p + \sqrt{p} < A.$$

Also $A < 2p - 1 < \sqrt{2N}$ and $A > \max(16, p + 2\sqrt{p})$. One can certify first-order interpolation directly using the derivative-weighted support of [43, Proposition 5.10]: choose multiplicity six, derivative cap three, total jet cap $B = 3A$, and challenge height $H = 40B$. Its coefficient count is $G = 4B^2 - 2B$ and its local-rank bound is 62. For $p \geq 4099$, $29p/5 \leq B \leq 6p$, and

$$39G - 40 \cdot 62N \geq \frac{7196}{25}p^2 - 468p > 0.$$

Since $G > 62N$, this makes $(H - B + 1)G - (H + 1)62N > 0$, a sufficient strict graded row test. The reconstruction guard is only $p > 3$. This certifies applicability, not a matching numerical exceptional-count upper bound. The rate and normalized common-agreement gap vanish. $\square$

An exact physical-fiber census independently checks the list spectrum at $p = 47$ and $p = 53$: respectively, there are 103776 singleton parameters and 47 lists of size 47, and 148824 singleton parameters and 53 lists of size 53. These checks validate the construction and counts at the stated thresholds; the uniform finite-certificate onset above remains $p \geq 4099$. The replay scripts, separate quadratic-root proof, and independent certificate audit accompany the source.

N.13 One-fiber puncturing: an exact spectrum with far inputs

Theorem N.22. *For every prime $p \geq 4099$, there is a dimension-three Reed–Solomon code over $E = \mathbb{F}_{p^3}$, of length*

$$2p^2 - 2p - 1 \leq n \leq 2p^2 - p - 1,$$

and a received affine line with the following exact list-size spectrum at agreement threshold

$$T = 2p - \lceil 4\sqrt{p} \rceil - 4 :$$

exactly $p^3 - 2p$ parameters have singleton lists, p parameters have lists of size p , and p parameters have empty lists. The empty-list parameters form an affine $\mathbb{F}_p$ -line, and their received words all have individual agreement at most $p + 2\sqrt{p}$. Choose any two distinct such words as endpoints. Their ordinary common agreement is at most $p + 2\sqrt{p}$, and a uniform E -valued affine mixture is nearby with probability $1 - p^{-2}$ and has a singleton list with probability $1 - 2p^{-2}$. Moreover

$$na_1(3/n) < T < \sqrt{2n}, \quad T - (p + 2\sqrt{p}) \geq p - 6\sqrt{p} - 5.$$

Thus the source separation is $\Theta(\sqrt{n})$ and the characteristic is $\Theta(\sqrt{n})$. There are p distinct collinear far inputs available; these are not affinely independent. The rate and normalized source separation tend to zero.

Proof. Start with Theorem N.21, at length $N = 2p^2 - p - 1$ and threshold $A = 2p - \lceil 4\sqrt{p} \rceil - 2$. Retain its notation $\Lambda_0 = (\delta/u_0)\mathbb{F}_p$. Its proof classifies all competitors: outside Λ_0 there is exactly one doubly canonical quadratic, and every other quadratic has at most $\max(p + 2\sqrt{p}, 16) = p + 2\sqrt{p}$

matches. At each parameter in Λ_0 , exactly p canonical quadratics meet threshold A , and all other quadratics have at most the same smaller bound.

Put $u_1 = \eta u_0 \in W$. Then

$$I_{u_1} = \eta^{-1}\Lambda_0, \quad \eta I_{u_1} = \Lambda_0, \quad I_{u_1} \neq \Lambda_0.$$

The nonzero core fibers of L_{u_1} partition D_0 minus its zero fiber. Their total physical size is $p^2 - p$, over $p - 1$ possible nonzero values. Choose $b_* \in I_{u_1}^*$ whose core fiber C_* has size at most p , and delete C_* from the domain. This gives the asserted length bounds. Put

$$Q_* = a_{u_1}X^2 + b_*, \quad \Lambda_* = b_* + \Lambda_0.$$

The two parameter lines Λ_*, Λ_0 are disjoint. For every $\lambda \in \Lambda_*$, the unique doubly canonical witness is Q_* , because $\lambda = b_* - \eta v$ for some $v \in I_{u_1}$. Deleting C_* removes its entire core agreement; its remaining fresh fiber has at most $p + 2\sqrt{p}$ matches. Every other quadratic was already below that bound, so these p parameters are far at threshold $T = A - 2 > p + 2\sqrt{p}$.

For any generic parameter outside Λ_* , let Q_λ be its unique high witness. If its leading coefficient differs from a_{u_1} , a common core match with Q_* is a root of the nonzero quadratic $Q_\lambda - Q_*$, so at most two matches are removed. If its leading coefficient is a_{u_1} , its constant term differs from b_* ; the corresponding core fibers are disjoint, and no matches are removed. Hence Q_λ retains at least $A - 2 = T$ matches. All its competitors remain below T . There are exactly $p^3 - 2p$ such generic parameters, and their lists are singleton.

At $\lambda \in \Lambda_0$, consider the p qualifying quadratics from the unpunctured classification. The one of direction u_1 has constant term zero, because $I_{u_1} \cap \Lambda_0 = \{0\}$; its core fiber is disjoint from C_* . Every other one has a different leading coefficient and loses at most two matches. All p therefore remain above T , and no other quadratic can enter. This proves the exact three-part list-size spectrum.

The restricted direction g is still nonzero, since no fresh points were deleted. Thus the p parameters in Λ_* give distinct far received words. Their pairwise ordinary common agreement is bounded by either individual agreement. For any two distinct parameters $\alpha, \beta \in \Lambda_*$, their affine mixtures correspond to $\lambda = \alpha + t(\beta - \alpha)$, an affine bijection of E . The exact spectrum therefore gives nearby probability $1 - p^{-2}$ and singleton probability $1 - 2p^{-2}$. More generally, uniform E -valued weights summing to one on any $t \geq 2$ distinct selected inputs induce a uniform scalar parameter: the corresponding affine functional has fibers of size $|E|^{t-2}$. In contrast, $\mathbb{F}_p$ -valued affine weights keep the parameter inside Λ_* , where every mixture is far. The challenge field in the asserted probabilities is essential.

For the threshold comparison, $p \geq 4099$ gives $T \geq 2p - 5\sqrt{p}$, while

$$na_1(3/n) \leq \sqrt{3n/2} + (3n/8)^{1/4} < \frac{7}{4}p + \sqrt{p} < T.$$

Using $n \geq 2p^2 - 2p - 1$ and $T \leq 2p - 4\sqrt{p}$ gives

$$2n - T^2 \geq 16p\sqrt{p} - 20p - 2 > 0.$$

The ceiling in T gives the stated source-separation bound. A uniform finite first-order certificate is also immediate: $m = 6$, derivative cap three, $B = 3T$, $H = 40B$ give $G = 4B^2 - 2B$ and local-rank bound 62. Since $29p/5 \leq B \leq 6p$ and $n \leq 2p^2$,

$$39G - 40 \cdot 62n \geq \frac{7196}{25}p^2 - 468p > 0.$$

As $G > 62n$, this implies the strict graded row test $(H - B + 1)G > (H + 1)62n$, with reconstruction guard $p > 3$. This is an extension-field result in the first-order regime, not fixed-rate or prime-alphabet tightness of a general upper bound. $\square$

The same exact list profile also holds at the interior threshold $T_0 = \lceil 19p/10 \rceil$ for every $p \geq 4099$. Indeed, $p + 2\sqrt{p} < T_0 \leq T$, so lowering the threshold admits no additional witness. The nearby probability remains $1 - p^{-2}$, and the source separation is at least $9p/10 - 2\sqrt{p}$. Since $n = 2p^2 - O(p)$, both the distance above the first-order curve and the distance below the Johnson agreement threshold are $\Theta(\sqrt{n})$. The finite certificate at T_0 follows from the larger-length calculation in Theorem N.25.

N.14 An asymptotically sharp companion for affine witnesses

Proposition N.23 (Sharp triple packing for affine witnesses). *Let $x_1, \dots, x_N$ be distinct field elements and let f, g be received words. Suppose their common agreement with degree- < 2 witnesses is at most C , where $2 \leq C < A \leq N$. Then the number of pairs (z, h) with $\deg h < 2$ and at least A agreements with $f + zg$ is at most*

$$\frac{N(N-1)(N-2)}{A(A-1)(A-C)}.$$

In particular this bounds the number of qualifying challenges.

Proof. Write $h(X) = \alpha X + \beta$. Coordinate i gives the affine plane

$$\alpha x_i + \beta - zg_i = f_i$$

in parameter space (α, β, z) . Any two distinct coordinates give independent constraints and determine a line on which α, β are affine functions of z . The coordinate planes containing this entire line are exactly the common agreements of the two corresponding affine witnesses for f, g , so there are at most C of them.

For any qualifying (z, h) , choose an ordered pair among A of its agreement coordinates. At least $A - C$ further agreement coordinates have planes not containing the resulting line. These give at least $A(A-1)(A-C)$ ordered triples of independent constraints. An independent triple determines at most one (z, h) , and there are at most $N(N-1)(N-2)$ ordered coordinate triples. Double counting proves the claim. $\square$

Corollary N.24 (Asymptotic equality in the triple bound). *Let $p \geq 5$ be prime, put*

$$n = 1 + p + p^2 + p^3 + p^4, \quad A = 1 + p + p^2, \quad C = p + 1,$$

and use the domain $\mu_n \subseteq \mathbb{F}_{p^5}$. Over $\mathbb{F}_{p^{15}}$, choose θ of degree three over $\mathbb{F}_{p^5}$ and take

$$f(Y) = Y^{1+p+p^2+p^3} + \theta Y^{1+p+p^2} + \theta^2 Y^{1+p}, \quad g(Y) = Y^{p+1}.$$

Their common and individual agreements for affine witnesses equal C . Exactly

$$M = \begin{bmatrix} 5 \\ 2 \end{bmatrix}_p = n(p^2 + 1)$$

nonzero challenges have agreement greater than C ; at every such challenge the qualifying list is a singleton, with exact agreement A . Moreover

$$M = \frac{n(n-1)(n-C)}{A(A-1)(A-C)}.$$

Thus the universal upper bound of Proposition N.23 is asymptotically attained with constant one, and $M \sim n^{3/2}$.

Proof. Projectivize Proposition N.18 by $Y = X^{p-1}$. Nonzero vectors of $\mathbb{F}_{p^5}$ have fibers of size $p-1$, and its three-space agreement sets become projective planes of size A . The witnesses become affine polynomials in Y . Conversely an affine witness pulls back to a degree- $< p$ witness, so the exact classification applies. The secondary agreements in that proposition include zero; removing it and projectivizing gives exactly C , not an additional high-agreement class. The second source has degree C , and its projection in the first source bounds both individual agreements by C . Projectivizing the simultaneous witnesses on any two-space minus zero attains this bound.

There is also exact equality in the line-specific triple count. Any two projective domain points span a projective line with C points. Their common interpolation pencil therefore lies in exactly C coordinate planes. Hence there are exactly $n(n-1)(n-C)$ independent ordered triples. Each such triple spans one projective plane, associated with its unique three-space locator and challenge. Equivalently, the displayed identity follows directly from

$$n-1 = p(p+1)(p^2+1), \quad n-C = p^2A, \quad A-1 = p(p+1), \quad A-C = p^2.$$

The ratio to the universal upper bound is $(n-C)/(n-2)$, which tends to one. $\square$

The affine example is a sharp finite packing statement, rather than an above-first-order asymptotic example: its agreement lies below the rate-only curve $na_1(2/n)$ for all sufficiently large p . The quadratic pullback in Theorem N.19 retains the exponent $3/2$ while moving above that curve and below exact Johnson at its advertised threshold.

N.15 A norm-one list with a matched finite upper bound

The next example concerns the size of one list, rather than the number of exceptional parameters on a received line. Its incidence parameters are those of the classical projective-plane design, with each coordinate doubled. We make no novelty claim for that geometry.

Theorem N.25. *Let $p \geq 53$ be prime, let $E = \mathbb{F}_{p^3}$, and put*

$$L = p^2 + p + 1, \quad N = 2L, \quad G = \ker \text{Norm}_{E/\mathbb{F}_p}, \quad D = \{x \in E^* : x^2 \in G\}.$$

The dimension-three Reed–Solomon code on D has a received word $f(x) = x^{2p+2}$ whose maximum agreement is $A = 2p+2$ and whose complete list at that threshold consists of the $L = N/2$ polynomials

$$Q_a(X) = aX^2 - a^{p^2+1}, \quad \text{Norm}_{E/\mathbb{F}_p}(a) = -1.$$

There are exactly $2L$ further quadratics

$$(uX - u^{p^2+1})^2, \quad \text{Norm}(u) \in \{1, -1\},$$

each with exactly $p+1$ matches. Every quadratic outside these two banks has at most four matches. Thus the complete list has size $3L$ for $5 \leq T' \leq p+1$, size L for $p+2 \leq T' \leq 2p+2$, and is empty above $2p+2$. The threshold- A list attains the Johnson counting bound exactly. At the lower threshold

$$T = \lceil 19p/10 \rceil$$

the complete list of f is still exactly the displayed $N/2$ witnesses, and the same code satisfies

$$Na_1(3/N) < T < \sqrt{2N},$$

$$\frac{N}{2} \leq \max_{w \in E^D} \#\{Q \in E[X] : \deg Q \leq 2, |\{x \in D : w(x) = Q(x)\}| \geq T\} \leq 22N.$$

Thus the lower bound is within a factor 44 of an explicit uniform finite first-order list bound at the same parameters. Moreover $T/\sqrt{N} \rightarrow 19/(10\sqrt{2})$, strictly between the limiting first-order value $\sqrt{3/2}$ and the Johnson value $\sqrt{2}$. Both margins therefore have order $\sqrt{N}$. Here the rate tends to zero, $p = \Theta(\sqrt{N})$, and $|E| = \Theta(N^{3/2})$; this is not a fixed-rate or prime-alphabet assertion.

Proof. The group G has odd order L and lies in the squares of E^* , so $|D| = 2L$. For $\text{Norm}(a) = -1$, substitute $y = a^{p^2}z$ in $y^{p+1} = ay - a^{p^2+1}$. The result is

$$z^{p+1} - z + 1 = 0, \quad z^p = 1 - 1/z.$$

Neither 0 nor 1 is a root. The fractional linear map $1 - 1/z$ has third iterate the identity, so all roots lie in E . The polynomial is squarefree: its derivative is $z^p - 1$, and a common zero would leave residual 1. Its $p+1$ roots satisfy $z^{p^2} = 1/(1-z)$ and hence $\text{Norm}(z) = -1$. Consequently all $p+1$ values of y have norm one. Each has two square roots in D , proving exactly A matches for each of the L distinct quadratics displayed above.

For any quadratic Q , the polynomial $f - Q$ has degree A , so no quadratic has more than A matches. If the complete threshold- A list has size m , let r_x count its matching witnesses at x . Distinct quadratics share at most two matching coordinates. Therefore

$$\sum_x r_x = mA, \quad \sum_x r_x^2 \leq mA + 2m(m-1).$$

Cauchy's inequality gives

$$m \leq \frac{N(A-2)}{A^2 - 2N} = L, \quad A^2 - 2N = 4p.$$

This proves completeness and exact Johnson saturation. Equality also shows that every coordinate belongs to $p+1$ supports and every pair of supports shares two coordinates.

We next show that every quadratic outside this bank has at most $p+1$ matches. We use the following elementary semilinear observation. For $v \neq 0$, a nonzero root of $z^{p+1} = uz + v$ satisfies $z^p = u + v/z$. The threefold Frobenius-twisted composition of this fractional linear map is the identity exactly when $u \neq 0$ and $v = -u^{1+p^2}$: this follows by multiplying the matrices $\begin{pmatrix} u^{p^i} & v^{p^i} \\ 1 & 0 \end{pmatrix}$ for $i = 2, 1, 0$. Otherwise it has at most two fixed points. The identity case has exactly $p+1$ roots by the same squarefreeness and Frobenius argument as above. When $v = 0$, there is at most one nonzero root.

For an even quadratic $Q = aX^2 + c$, apply this observation to $y = x^2$. More than four matches requires $a \neq 0$ and $c = -a^{1+p^2}$. All its y -roots have norm $-\text{Norm}(a)$, so any match on D then requires $\text{Norm}(a) = -1$. For $Q = aX^2 + bX + c$ with $b \neq 0$, put

$$r(X) = 1 - a^p Q(X) - c^p X^2.$$

At a matching $x \in D$, the norm-one condition gives $x^2 Q(x)^p = 1$, hence $r(x) = b^p x^{p+2}$. Thus

$$r(X)^2 - b^{2p} X^2 Q(X)$$

vanishes at every match. If this polynomial is nonzero, there are at most four matches. If it is zero, X divides r and

$$H(X) = \frac{r(X)}{b^p X}$$

is a polynomial of degree at most one, with $H^2 = Q$. Write $H = uX + v$. Since $b \neq 0$, both u and v are nonzero. Comparing the constant and linear coefficients in $r = b^p XH$ gives

$$(u^p v)^2 = 1, \quad v^p = -u^{p^2+1}.$$

Hence $v = -u^{p^2+1}$ and $\text{Norm}(u)^2 = 1$. Conversely these conditions imply the full quartic identity, with $r/(b^p X) = uX - u^{p^2+1}$. The original unsquared identity fixes the sign: at every match,

$$x^{p+1} = ux - u^{p^2+1}.$$

This equation has exactly $p + 1$ roots in E , all with norm $-\text{Norm}(u) \in \{1, -1\}$, by the semilinear observation. All lie in D and are matches, so the agreement is exactly $p + 1$. Writing $H_u(X) = uX - u^{p^2+1}$, the $2L$ resulting squares are distinct: equal squares would give $H_{u'} = \pm H_u$; the minus case forces $u' = -u$ but has incompatible constant coefficients because $p^2 + 1$ is even and p is odd. Their nonzero linear coefficients separate them from the original even bank. Every remaining quadratic has at most four matches. This proves the stated complete profile, and in particular the original bank remains the complete list at T .

For the uniform upper bound at T , use the derivative-weighted support of [43, Proposition 5.10] with

$$m = 6, \quad B_\partial = 3, \quad B = B_{\text{jet}} = 3T, \quad H = B_Z = 40B.$$

Its coefficient count is $G_{\text{coef}} = 4B^2 - 2B$ and its six local-rank contributions are 4, 8, 12, 15, 14, 9, totaling $R = 62$. For $p \geq 53$, $57p/10 \leq B \leq 6p$, and

$$39G_{\text{coef}} - 40NR \geq \frac{2711}{25}p^2 - 5428p - 4960 > 0.$$

Thus $G_{\text{coef}} > NR$ and the graded support surplus is at least $B(39G_{\text{coef}} - 40NR) + (G_{\text{coef}} - NR) > 0$. This certificate applies uniformly to evaluation domains and received words. Reconstruction requires only $p > \max(2, 3)$.

In [43, Eq. (54), Lemma 5.11], message degree 2 and $(B, M) = (3T, 3)$ give $F_{\text{reg}} = 7B - 12$ and $S = 5B - 9$. Writing $\lambda = (N - 2)/(T - 2)$, the fixed-word bound is therefore

$$\lambda F_{\text{reg}} + S = 21N + 30\lambda + 15T - 51 \leq 21N + 90p - 51 \leq 22N.$$

Here $T - 2 \geq p + 1$ gives $\lambda \leq 2p$, and $N - 90p + 51 = 2p^2 - 88p + 53 > 0$. The exhibited complete list supplies the lower bound at T . Finally, $T < 2p < \sqrt{2N}$, while the first-order curve estimate gives

$$Na_1(3/N) \leq \sqrt{3N/2} + (3N/8)^{1/4} < \frac{7}{4}p + \sqrt{p} < T$$

using $T \geq 19p/10$, $\sqrt{p} > 20/3$, and $p \geq 53$. □

The matched comparison is codewide, not just a bound along one received line. Its normalized threshold and normalized separation from the first-order curve both vanish with the rate; it makes no claim of constant-factor tightness for exceptional-label or common-agreement bounds.

Corollary N.26. *With the code and word of Theorem N.25, let $p \geq 401$. At every integer threshold*

$$\lceil \sqrt{3}p \rceil \leq T \leq 2p + 2,$$

the displayed word has exactly $N/2$ nearby quadratics, and the maximum list size over all received words is at most $137N$. Consequently, for every fixed $c \in (\sqrt{3}, 2)$, at threshold $T = \lceil cp \rceil$ the maximum list size is $\Theta(N)$. For sufficiently large p these thresholds lie strictly between the first-order curve and the Johnson threshold, with both margins of order $\sqrt{N}$.

Proof. The classification in the theorem already gives the exact $N/2$ list throughout this band. For the uniform upper bound, take

$$m = 16, \quad B_\partial = 8, \quad B = B_{\text{jet}} = 8T, \quad H = B_Z = 100B.$$

The derivative-weighted coefficient count and local rank are

$$G = 9B^2 - 27B, \quad R = 852.$$

Since $8\sqrt{3}p \leq B \leq 16p + 16$, we have

$$99G - 100NR \geq 672p^2 - 213168p - 213168 > 0 \quad (p \geq 401).$$

The graded surplus $(H - B + 1)G - (H + 1)NR$ is therefore positive. The characteristic guard is $p > 8$. Equation (54) and Lemma 5.11 of [43] give $F_{\text{reg}} = 17B - 72$ and $S = 15B - 64$. With $\lambda = (N - 2)/(T - 2) \leq 2p$, the list bound is

$$\lambda F_{\text{reg}} + S = 136N + 200\lambda + 120T - 336 \leq 136N + 640p - 96 \leq 137N.$$

Finally, $Na_1(3/N) = \sqrt{3}p + O(\sqrt{p})$ and $\sqrt{2N} = 2p + O(1)$. Thus every fixed $c \in (\sqrt{3}, 2)$ has the asserted two strict margins for all sufficiently large p . The finite band in the first statement need not itself lie wholly between those two boundaries. $\square$

Corollary N.27. *Fix $\sqrt{2} < c < 2$. For all sufficiently large primes p depending on c , the maximum list size for the norm-one code at $T = \lceil cp \rceil$ satisfies*

$$\frac{N}{2} \leq L_{\max}(T) \leq C_c N, \quad C_c = O((c^2 - 2)^{-3}).$$

The upper bound holds on every length- N , dimension-three Reed–Solomon domain in characteristic p . The lower bound uses the same word and exactly the same $N/2$ witnesses as Theorem N.25.

Proof. Put $\delta = c^2 - 2$ and choose the fixed integers

$$h = \lceil 8/\delta \rceil, \quad s = \lceil 8h/\delta \rceil, \quad \kappa = \lceil 4c^2/\delta \rceil.$$

Use multiplicity $m = 2h$, derivative cap s , joint cap $B = hT$, and challenge cap $H = \kappa B$. For sufficiently large p , $B \geq 2h + s$ and $p > s$. The exact derivative-weighted counts are

$$G = (s + 1)B(B + 1 - s/2), \quad R = \frac{h(h + 1)(6s + 2h + 7)}{6}.$$

For the first identity, the coefficient allowance at a fixed derivative exponent b sums to $B(B - b + 1)$. For the second, the local block at index $a \in \{0, \dots, 2h - 1\}$ has rank $q(a + s + 2 - q)$, where $q = \min(a + 1, 2h - a)$; summing gives R . The assumed joint cap makes all nonzero blocks active.

The leading coefficient condition is controlled by

$$\gamma := \frac{2R}{(s + 1)h^2} = 2 + \frac{2}{h} + \frac{(h + 1)(2h + 1)}{3h(s + 1)} \leq 2 + \frac{\delta}{2} < c^2.$$

Indeed, $2/h \leq \delta/4$ and the final term is at most $2h/(s + 1) \leq \delta/4$. The exact expression $(\kappa - 1)G - \kappa NR$ has leading coefficient in p

$$(s + 1)h^2(\kappa(c^2 - \gamma) - c^2) > 0;$$

the rounding of T contributes only $O_c(p)$. Hence it is positive for all sufficiently large p . It follows that $G > NR$ and the full graded surplus

$$(H - B + 1)G - (H + 1)NR = B((\kappa - 1)G - \kappa NR) + (G - NR)$$

is strictly positive. This checks the finite support certificate, with reconstruction guard $p > \max(2, s)$.

The fixed-word transfer gives

$$F_{\text{reg}} = (2s + 1)B - s(s + 1), \quad S = (2s - 1)B - s^2.$$

Since $\lambda = (N - 2)/(T - 2) = O_c(p)$, its list bound is

$$\lambda F_{\text{reg}} + S = h(2s + 1)N + O_c(p) \leq (h(2s + 1) + 1)N$$

for sufficiently large p . Thus one can take $C_c = h(2s + 1) + 1 = O(\delta^{-3})$. Finally $p + 1 < T \leq 2p + 2$ eventually, so the complete list of the exhibited word is still exactly $N/2$. $\square$

This wider range is a finite-degree consequence of the same DKT support machinery, using message degree two. For $\sqrt{2} < c \leq \sqrt{3}$ it is not the strict interior above the rate-only curve $Na_1(3/N)$. The result concerns growing lists at vanishing rate over extension fields; it does not change the stated first-order curve or supply prime-alphabet tightness.

N.16 Fixed rate with growing characteristic and almost every native challenge

The locator compiler of [42] also admits a growing-characteristic specialization. The improvement below uses its full seed domain, padding inside that domain, and an exact five-dimensional label calculation. Write agr_J for agreement with polynomials of strict degree less than J , and CA_J for common agreement with two such polynomials.

Theorem N.28. *Fix $0 < \rho < 1$. For every sufficiently large prime p , put*

$$E = \mathbb{F}_{p^5}, \quad N = p^5, \quad J = \lfloor \rho N \rfloor, \quad \Delta = \left\lfloor \frac{1 - \rho}{2} p^3 \right\rfloor.$$

There are words $f, g : E \rightarrow E$ with

$$\text{agr}_J(g) = \text{CA}_J(f, g) = J$$

such that at least $N - 3$ values $\lambda \in E$ satisfy $\text{agr}_J(f + \lambda g) \geq J + \Delta$. For an infinite set of primes, of relative density $1/3$ among the primes, the same construction has this property for every $\lambda \in E$. A sufficient onset is

$$p \geq 37, \quad p \geq \frac{4}{1 - \rho}, \quad p^3 \geq \frac{2}{\rho}, \quad \frac{(1 - \rho)^2 p^3}{64} > \log 4 + 6 \log p.$$

The characteristic is $N^{1/5}$ and the additive gap is $\Theta_\rho(N^{3/5})$. The alphabet remains an extension field, the normalized gap tends to zero, and $p \ll J$. Thus this is not a prime-alphabet or large-characteristic first-order counterexample. The first source f is itself close for every admissible prime, since the canonical label zero always occurs. Thus the native pair has only its direction individually far.

More explicitly, at actual rate J/N the radius certified for these challenges is $1 - J/N - \Delta/N$, while the common-agreement loss is $\Delta/N = \Theta_\rho(N^{-2/5})$. This is a smaller distance below capacity

than the inverse-logarithmic margin of the fixed-rate counterexamples in [21, Theorem 1]. The theorem therefore does not improve that capacity-margin benchmark. Moreover, Crites–Stewart [23, Corollary 1] already imply an asymptotically stronger existence statement on this same full domain: choose the direction $g(X) = X^J$, whose agreement is exactly J . Their result gives a center with every challenge nearby at agreement $J + \Theta_\rho(N/\log N)$. Common agreement with this direction is exactly J , by interpolation on J coordinates and the direction’s root-count upper bound. Thus the theorem above supplies a structured algebraic realization and a specific finite onset, not a new existential coverage or capacity-margin regime. The two-far-source norm-quotient specialization in Theorem N.31 gives a larger margin and two far sources, but has common agreement above J rather than exactly J . The full domain has size p^5 ; the theorem supplies no corresponding construction on short multiplicative evaluation domains.

Lemma N.29. *For each three-dimensional $\mathbb{F}_p$ -subspace $W \subset E$, write*

$$L_W(X) = X^{p^3} + a_1 X^{p^2} + a_2 X^p + a_3 X, \quad z_W = a_1 + a_2^p - a_1^{p+1}.$$

If $p \geq 37$, then

$$\{z_W\} = E \setminus \{z \in \mathbb{F}_p : z^3 + 2z^2 + 3z + 1 = 0\}.$$

Proof. Let $U = W^\perp$ for the trace pairing, and for a basis u, v of U write $P_{ij} = u^{p^i} v^{p^j} - u^{p^j} v^{p^i}$, with indices modulo five. Complementary Moore minors give $a_1 = P_{24}/P_{34}$ and $a_2 = P_{14}/P_{34}$. Consecutive Moore minors are nonzero at every rational two-space. Frobenius and the Plücker relation $P_{02}P_{34} - P_{03}P_{24} + P_{04}P_{23} = 0$ give

$$z_W = (P_{24} - P_{23})/P_{34}.$$

Evaluate the determinant of

$$u \mapsto (u - zu^{p^2}, u^p - u^{p^2})$$

on the two-space U^{p^2} . It equals $P_{23} - P_{24} + zP_{34}$ on the original basis of U . The map is injective for $z \neq 1$, so absence of that label is exactly maximum scatteredness of its five-dimensional image. Montanucci–Zanella’s Proposition 4.3 [46] excludes maximum scatteredness when $z \notin \mathbb{F}_p$; the theoretical bound in that proof applies for $p \geq 37$. Thus it remains to analyze scalar z . The value 1 occurs by taking any three-space in the kernel of $\text{Tr}_{E/\mathbb{F}_p}$.

Here is a direct scalar calculation, avoiding the full classification of scattered linear sets. Over an algebraic closure, let indices be cyclic modulo five and put

$$\ell_j = e_j \wedge e_{j+2} - e_j \wedge e_{j+1} - ze_{j+1} \wedge e_{j+2}.$$

These five forms are independent. The primal Grassmannian section encoding the label and the dual section $D = \text{Gr}(2, 5) \cap \mathbb{P}\langle \ell_0, \dots, \ell_4 \rangle$ have equal numbers of $\mathbb{F}_p$ -points. This elementary incidence identity, valid also for improper sections, is [47, Lemma 2.3]: count pairs consisting of a two-space and an alternating form vanishing on it. Rank-two and rank-four forms have incidence counts differing by p^4 , while a two-space in the primal section contributes p^4 extra forms.

For completeness the relevant geometry of D can be checked explicitly. Its equations are the five principal Pfaffians of $\sum t_j \ell_j$. The coordinate vertices belong to D and Frobenius cycles them. At e_0 , the Jacobian in $t_1, \dots, t_4$ is

$$\begin{pmatrix} 0 & z^2 & z & 0 \\ 0 & -z & -1 & 0 \\ 0 & z & 1 & 0 \\ 0 & -1 & -z^2 & 1-z \\ z-1 & 1 & z & 0 \end{pmatrix}.$$

For $z \neq 1$ it has rank three. On $t_0 = t_1 = 0$ the Pfaffians force $t_2 t_4 = (1 - z)t_2 t_3 = (z - 1)t_3 t_4 = 0$, leaving only three vertices. A component of dimension at least two would meet this plane at a vertex, contradicting the local dimension one there. Hence the section is proper. The height-three Pfaffian resolution has terms $R(-5), R(-3)^5, R(-2)^5, R$; consequently D is a pure, connected degree-five curve of arithmetic genus one.

Its tangent direction at e_0 is

$$v = (0, z + 1, 1, -z, z^2 + z + 1).$$

Substitution in the five Pfaffians gives $(zH, -H, H, -H, z^2 H)$, where $H = z^3 + 2z^2 + 3z + 1$. If D has no rational point, the component through this smooth vertex has Frobenius orbit of length five: an invariant component would contain all five spanning vertices and have normalization of genus at most one, hence a rational point by Hasse–Weil. This is the smooth-orbit argument of [47, Lemma 3.1]. Degree five then forces five lines, so the tangent line is contained in D and $H = 0$. Conversely, if $H = 0$, all five conjugate tangent lines lie in D . The minor testing the first line against its Frobenius image is $-(z^2 + z + 1)$, which is nonzero since H modulo $z^2 + z + 1$ is z . These two lines are disjoint. Their orbit has length five and exhausts D ; a rational point would belong to all conjugates, which is impossible. Thus precisely the roots of H are absent. Since $H(1) = 7$, the previously separated value causes no exception for $p \geq 37$. $\square$

Proof of Theorem N.28. For each W , coefficient elimination gives

$$L_W^p + (1 - a_1^p)L_W = X^{p^4} + X^{p^3} + z_W X^{p^2} + R_W, \quad \deg R_W \leq p, \quad R_W(0) = 0.$$

Thus $f_0 = X^{p^4-1} + X^{p^3-1}$ and $g_0 = X^{p^2-1}$ have a witness $h_W = -R_W/X$, of degree at most $p - 1$, on $W \setminus \{0\}$ at challenge z_W . This is the locator compiler of [42], with its degrees displayed explicitly.

Set $w = J - p^2 + 1$ and choose a uniformly random w -subset $B \subset E$. For each W , the hypergeometric variable $X_W = |B \cap W|$ has mean at most ρp^3 and satisfies

$$\Pr\{X_W > (\rho + \varepsilon)p^3\} \leq \exp(-\varepsilon^2 p^3/4), \quad \varepsilon = (1 - \rho)/4.$$

For example this follows from $\mathbb{E}(\binom{X_W}{j}) \leq \binom{p^3}{j}(w/N)^j$ and the exponential moment bound. There are $\binom{5}{2}_p < 4p^6$ subspaces. The stated onset therefore gives a single B with the required intersection bound for every W .

Let $L_B = \prod_{b \in B} (X - b)$ and set $f = L_B f_0$, $g = L_B g_0$ on E . The padded witness $L_B h_W$ has degree at most $w + p - 1 = J - p^2 + p < J$. Its matches include $B \cup (W \setminus \{0\})$, so their number is at least

$$w + p^3 - 1 - |B \cap W| \geq J + (1 - \rho - \varepsilon - 1/p)p^3 \geq J + (1 - \rho)p^3/2.$$

The polynomial g is monic of degree exactly J . Root counting bounds its agreement with every strict-degree witness by J , and interpolation attains J . Interpolating explanations for both words on any J coordinates also proves $\text{CA}_J(f, g) = J$. Lemma N.29 now supplies at least $N - 3$ challenges. Finally H is irreducible over $\mathbb{Q}$ by the rational-root test and has discriminant -23 . Its Galois group is S_3 , so Chebotarev gives density $1/3$ for primes where it remains irreducible. On those primes it has no scalar root, giving every native challenge. Padding can create additional good challenges; no exact post-padding exclusion is asserted. $\square$

N.17 Full-field norm quotients with almost every challenge

The quotient construction of Krachun, Kazanin, and Haböck [21, Appendix A] can be applied directly in the quotient variable. A pole outside its image suffices; it need not be the image of an evaluation-field point. Combining this observation with a character-sum estimate gives the following native-field specialization.

Lemma N.30. *Let $\mathcal{A}$ be an s -element subset of a finite abelian group H of order M . Suppose every nontrivial character χ of H satisfies $|\sum_{a \in \mathcal{A}} \chi(a)| \leq B < s$. For $0 < r < s$, put $\theta = r/s$ and*

$$\varepsilon = (s + 1) \exp(-\theta(1 - \theta)(s - B)).$$

If $(M - 1)\varepsilon < 1$, every element of H is the product of an r -element subset of $\mathcal{A}$.

Proof. Fix a nontrivial character and write $v_a = \chi(a)$. For $t > 0$ and real ϕ , the identity

$$|1 + te^{i\phi}v_a|^2 = (1 + t)^2 - 2t(1 - \operatorname{Re}(e^{i\phi}v_a))$$

and $1 - u \leq e^{-u}$ give

$$\left| \prod_{a \in \mathcal{A}} (1 + te^{i\phi}v_a) \right| \leq (1 + t)^s \exp\left(-\frac{t(s - B)}{(1 + t)^2}\right).$$

Cauchy's coefficient bound, with $t = r/(s - r)$, bounds the absolute value of the degree- r coefficient by $t^{-r}(1 + t)^s \exp(-\theta(1 - \theta)(s - B))$. The value r is a mode of the binomial distribution with parameters $s, r/s$, so $t^{-r}(1 + t)^s \leq (s + 1)\binom{s}{r}$. Consequently every nontrivial character contributes at most $\varepsilon\binom{s}{r}$ in absolute value. Character orthogonality therefore bounds the number of r -subsets with any prescribed product from below by

$$\frac{1}{M} \binom{s}{r} (1 - (M - 1)\varepsilon) > 0.$$

□

Theorem N.31. *Fix $0 < \rho < 1$ and an integer $d \geq 2$. For every sufficiently large prime p , put*

$$q = p^d, \quad E = \mathbb{F}_q, \quad m = \frac{q - 1}{p - 1}, \quad J = \lfloor \rho q \rfloor.$$

There are words $f, g : E \rightarrow E$ such that

$$\operatorname{agr}_J(f) = \operatorname{agr}_J(g) = \operatorname{CA}_J(f, g) = J + m - 1,$$

and every $\lambda \in E \setminus \{0\}$ satisfies

$$\operatorname{agr}_J(f + \lambda g) \geq J + 2m - 1.$$

Here the code has strict degree bound J , the capacity margin is $(2m - 1)/q = \Theta(p^{-1})$, and the source-agreement gap is $m/q = \Theta(p^{-1})$.

For an explicit sufficient condition, write $J - 1 = (r - 2)m + w$ with $0 \leq w < m$, and set $\theta = r/(p - 2)$. It is enough that $2 \leq r < p - 2$ and

$$\log((p^d - 2)(p - 1)) < \theta(1 - \theta)(p - 4 - (d - 1)\sqrt{p}).$$

Proof. Choose $b \in E$ of degree exactly d over $\mathbb{F}_p$. Write $N(X) = X^m$, so that N is the norm on E , with image $\mathbb{F}_p$. Fix $a_0 \in \mathbb{F}_p^*$, choose a w -element subset $B \subset N^{-1}(a_0)$, and put $R(X) = \prod_{x \in B} (X - x)$. In particular R is monic of degree $w < m$. Define

$$f(X) = R(X) \frac{N(X)^r - b^r}{N(X) - b}, \quad g(X) = -\frac{R(X)}{N(X) - b}.$$

The first expression is a monic polynomial of degree $(r-1)m + w = J + m - 1$; the denominator in the second expression never vanishes on E .

For every polynomial h of degree less than J , root counting bounds the agreement of f with h by $J + m - 1$. The same bound applies to g : its agreement equation is $R + (N - b)h = 0$, a nonzero polynomial of degree at most $J + m - 1$. It is nonzero because $\deg R < m$. These two bounds are attained simultaneously. Choose any $r-1$ distinct tags in $\mathbb{F}_p^* \setminus \{a_0\}$, interpolate both functions $(Y^r - b^r)/(Y - b)$ and $-1/(Y - b)$ on those tags by polynomials of degree at most $r-2$, substitute $Y = N(X)$, and multiply by $R(X)$. Both resulting witnesses have degree at most $(r-2)m + w = J - 1$. They match on the selected $r-1$ norm fibers and on B , a disjoint set of size $(r-1)m + w = J + m - 1$. This proves all three stated equalities.

It remains to represent every element of E^* as a product

$$\prod_{a \in S} (b - a), \quad S \subset \mathbb{F}_p^* \setminus \{a_0\}, \quad |S| = r.$$

For every nontrivial multiplicative character χ of E^* , Katz's affine-line estimate [22, Theorem 1] bounds $|\sum_{a \in \mathbb{F}_p} \chi(b - a)|$ by $(d-1)\sqrt{p}$. Deleting 0, a_0 gives a bound $(d-1)\sqrt{p} + 2$ on the remaining $p-2$ terms. Lemma N.30 consequently proves the required surjectivity under the displayed sufficient condition. For fixed $\rho, d, \theta \rightarrow \rho$, so this condition holds for all sufficiently large p .

For $\lambda \neq 0$, choose such an S with product $-\lambda$, and set

$$V_S(Y) = \prod_{a \in S} (Y - a), \quad P_S(Y) = Y^r - V_S(Y), \quad h_S(X) = R(X) \frac{P_S(N(X)) - P_S(b)}{N(X) - b}.$$

The quotient is polynomial and $\deg h_S \leq w + (r-2)m = J - 1$. Since $P_S(b) = b^r - V_S(b) = b^r + \lambda$, direct subtraction gives

$$f(X) + \lambda g(X) - h_S(X) = \frac{R(X)V_S(N(X))}{N(X) - b}.$$

Its zeros are exactly B and the r norm fibers indexed by S . These sets are disjoint and have total size $w + rm = J + 2m - 1$. In particular $X = 0$ supplies no hidden extra agreement. $\square$

In particular, put $T = J + 2m - 1$ and $\delta = 1 - T/q$. Both sources have distance exactly $\delta + m/q$ from the code, but every affine combination $(1-t)f + tg$ with $t \notin \{0, 1\}$ has distance at most δ . Indeed, divide by $1-t$ and apply the theorem with $\lambda = t/(1-t) \neq 0$. Thus the nearby fraction on this affine line is exactly $1 - 2/q$. The normalized common-agreement loss m/q is asymptotically half the capacity margin $(2m-1)/q$.

For $d = 5$, this gives a larger margin, of order $q^{-1/5}$, than the order- $q^{-2/5}$ margin of Theorem N.28, and both sources are individually far. It does not preserve that theorem's exact common agreement J : here common agreement is $J + m - 1$. Both margins are smaller than the inverse-logarithmic margin in [21]. Crites-Stewart [23, Corollary 1] already give every affine challenge nearby, with a far direction, at that larger margin over arbitrary fields and domains. Thus near-complete coverage is not itself a new regime. The additional property here is the structured construction

with two individually far sources and exact source and common agreement; no priority claim for that refinement is made. The alphabet is an extension field, $p \ll J$, and the evaluation domain has all p^d points. No construction on the short domains used by the practical benchmark follows from this theorem.

N.18 Two far endpoints at inverse-logarithmic margin

A logarithmic-size selection of norm fibers gives the inverse-logarithmic capacity margin of the quotient construction while retaining almost every native challenge. The evaluation domain is chosen for the construction. Crites-Stewart [23, Corollary 1] already give every affine challenge nearby at this order of margin, over arbitrary fields and domains, with a far direction. The point of the following specialization is its two far endpoints and exact agreement profile; no novelty claim for that refinement is made.

Lemma N.32. *Let $E = \mathbb{F}_q$, where $q = p^d$ and $d \geq 2$, and let $b \in E$ have degree d over $\mathbb{F}_p$. Fix a set $T \subset \mathbb{F}_p$ of e excluded tags, put $U = \mathbb{F}_p \setminus T$, and define*

$$N = p - e, \quad \mu = \frac{(d-1)\sqrt{p} + e}{p - e}.$$

For $1 \leq s \leq N$, choose an s -element subset G of U uniformly without replacement. For any real $\beta > \mu$,

$$\Pr \left[\exists \chi \neq 1 : \left| \sum_{a \in G} \chi(b - a) \right| > \beta s \right] \leq 4(q - 2) \exp \left(-\frac{s(\beta - \mu)^2}{4} \right),$$

where the quantifier ranges over all nontrivial multiplicative characters of E^ .*

Proof. Katz's affine-line estimate [22, Theorem 1] and deletion of the e tags give $|N^{-1} \sum_{a \in U} \chi(b - a)| \leq \mu$ for every nontrivial χ . We prove the needed concentration estimate without assuming independence of the sampled tags.

For a real population $x_1, \dots, x_N \in [-1, 1]$ with mean $\bar{x}$, let S be the sum over a uniform s -element subset. For real t ,

$$\mathbb{E} e^{tS} = \frac{e_s(e^{tx_1}, \dots, e^{tx_N})}{\binom{N}{s}} \leq \left(\frac{1}{N} \sum_{i=1}^N e^{tx_i} \right)^s.$$

Here e_s denotes the elementary symmetric polynomial. Its displayed bound follows by averaging pairs of positive arguments with fixed sum: with all other arguments fixed, it has the form $Ay_i y_j + B(y_i + y_j) + D$ with $A \geq 0$. Equivalently, its maximum on the nonnegative simplex occurs when all arguments are equal. Indeed, for $s \geq 2$ a maximizing tuple has at least s positive arguments, so averaging any unequal pair strictly increases the value; the case $s = 1$ is immediate.

If X is one uniform population entry, the second derivative of $\log \mathbb{E} e^{t(X - \bar{x})}$ is the variance under an exponential tilt, and is at most 1 because $X \in [-1, 1]$. Its value and first derivative vanish at $t = 0$. Thus $\mathbb{E} e^{t(S - s\bar{x})} \leq e^{st^2/2}$. Chernoff's bound, applied with both signs of t , yields

$$\Pr[|S - s\bar{x}| > su] \leq 2e^{-su^2/2}.$$

Apply this separately to the real and imaginary parts of $\chi(b - a)$. A complex deviation of modulus greater than sv has a coordinate of absolute value greater than $sv/\sqrt{2}$, so its probability is at most $4e^{-sv^2/4}$. Set $v = \beta - \mu$ and take the union bound over the $q - 2$ nontrivial characters. $\square$

Corollary N.33. Fix $0 < \eta \leq 1/2$, set $\kappa = \eta(1 - \eta)$ and

$$C = \max\{128, 8/\kappa\}, \quad s = \lceil C \log q \rceil.$$

Fix $a_0 \in \mathbb{F}_p^*$, and suppose

$$q \geq C + 2, \quad s \leq p - 2, \quad \frac{(d-1)\sqrt{p} + 2}{p-2} \leq \frac{1}{4}.$$

With probability at least $1 - 4/q$, a uniform s -element subset $G \subset \mathbb{F}_p^* \setminus \{a_0\}$ has the following property simultaneously for every integer r with $\eta \leq r/s \leq 1 - \eta$: every element of E^* is a product of exactly r distinct elements of $\{b - a : a \in G\}$.

Proof. Use Lemma N.32 with $e = 2$ and $\beta = 1/2$. The failure probability is at most

$$4(q-2)e^{-s/64} \leq 4q^{1-C/64} \leq 4/q.$$

For any successful G , Lemma N.30 applies with $B = s/2$. Its total nontrivial-character error is at most

$$(q-2)(s+1)e^{-\kappa s/2} \leq \frac{s+1}{q^3} \leq \frac{1}{q} < 1.$$

For the last inequality, use $s+1 \leq C \log q + 2 \leq Cq + 2 \leq q^2$, which follows from $q \geq C + 2$. One simultaneous character-bias event suffices for every indicated r ; no further probability union bound is needed. $\square$

Theorem N.34. Fix $0 < \rho < 1$, an integer $d \geq 2$, and

$$C > \max\left\{64, \frac{4}{\rho(1-\rho)}\right\}.$$

For every sufficiently large prime p , put

$$q = p^d, \quad E = \mathbb{F}_q, \quad m = \frac{q-1}{p-1}, \quad s = \lceil C \log q \rceil, \quad n = (s+1)m, \quad J = \lfloor \rho n \rfloor.$$

There is an n -point domain $D \subset E$ and words $f, g : D \rightarrow E$ with

$$\text{agr}_J(f) = \text{agr}_J(g) = \text{CA}_J(f, g) = J + m - 1,$$

such that every $\lambda \in E \setminus \{0\}$ satisfies

$$\text{agr}_J(f + \lambda g) \geq J + 2m - 1.$$

The capacity margin and source-agreement gap are respectively

$$\frac{2m-1}{n} = \frac{2}{s+1} - \frac{1}{n}, \quad \frac{m}{n} = \frac{1}{s+1},$$

both of order $1/\log n$. The alphabet and characteristic obey

$$q = \Theta\left((n/\log n)^{d/(d-1)}\right), \quad p = \Theta\left((n/\log n)^{1/(d-1)}\right).$$

Proof. Choose $b \in E$ of degree d over $\mathbb{F}_p$ and fix $a_0 \in \mathbb{F}_p^*$. For every nontrivial multiplicative character χ of E^* , Katz's estimate [22, Theorem 1] gives

$$\left| \frac{1}{p-2} \sum_{a \in \mathbb{F}_p^* \setminus \{a_0\}} \chi(b-a) \right| \leq \frac{(d-1)\sqrt{p}+2}{p-2} \leq \frac{1}{4}$$

for sufficiently large p . Choose an s -element subset G of this population uniformly without replacement. Lemma N.32 and a union bound show that, with probability at least $1 - 4(q-2)\exp(-s/64) > 0$, it satisfies

$$\left| \sum_{a \in G} \chi(b-a) \right| \leq s/2 \quad \text{for every nontrivial } \chi.$$

Fix such a G , write $N(X) = X^m$, and set $D = N^{-1}(G \cup \{a_0\})$. Each nonzero norm fiber has size m , so $|D| = (s+1)m = n$.

Write $J-1 = (r-2)m + w$, where $0 \leq w < m$, and put $\theta = r/s$. Then $\theta \rightarrow \rho$ and $2 \leq r < s$ for sufficiently large p . Lemma N.30, with $\mathcal{A} = \{b-a : a \in G\}$ and $B = s/2$, shows that all elements of E^* are products of exactly r distinct elements of $\mathcal{A}$, provided

$$(q-2)(s+1)\exp(-\theta(1-\theta)s/2) < 1.$$

Our choice of C makes both displayed probability conditions hold eventually. Also $s < p-2$ eventually, so the sampling is defined.

Choose a w -point set $B_0 \subset N^{-1}(a_0)$, put $R = L_{B_0}$, and use the norm quotient words

$$f = R \frac{N^r - b^r}{N - b}, \quad g = -\frac{R}{N - b} \quad \text{on } D.$$

The denominator is nonzero on E . The polynomial f is monic of degree $J+m-1$, and the agreement equation for g is the nonzero polynomial $R + (N-b)h = 0$, of degree at most $J+m-1$ for every $\deg h < J$. These source bounds are attained simultaneously by interpolating the two quotient-variable functions on any $r-1$ tags of G , substituting N , and multiplying by R . The resulting degree-at-most- $J-1$ witnesses match on those fibers and B_0 , giving exactly $J+m-1$ common agreements.

For $\lambda \neq 0$, choose $S \subset G$, $|S| = r$, with $V_S(b) = -\lambda$, where $V_S(Y) = \prod_{a \in S} (Y-a)$. Writing $P_S(Y) = Y^r - V_S(Y)$, the polynomial

$$h_S = R \frac{P_S(N) - P_S(b)}{N - b}$$

has degree at most $w + (r-2)m = J-1$ and satisfies

$$f + \lambda g - h_S = \frac{RV_S(N)}{N - b}.$$

Its zeros on D are exactly B_0 and the r norm fibers indexed by S , totaling $w + rm = J + 2m - 1$. Finally, $n = \Theta(p^{d-1} \log p)$ proves the parameter assertions. $\square$

For finite parameters the proof requires only $s < p-2$, $2 \leq r < s$, and the three explicit inequalities

$$\frac{(d-1)\sqrt{p}+2}{p-2} \leq \frac{1}{4}, \quad 4(q-2)e^{-s/64} < 1, \quad (q-2)(s+1)e^{-\theta(1-\theta)s/2} < 1.$$

At threshold $T = J + 2m - 1$, exactly $q - 2$ of the affine mixtures $(1 - t)f + tg$, $t \in E$, are nearby; the two endpoints are farther by m/n . This gives almost-complete coverage at an inverse-logarithmic margin by standard quotient and character-sum ingredients. The domain is a chosen union of norm fibers, without a subgroup requirement. The alphabet remains an extension field, and the domain length is $\Theta(p^{d-1} \log p)$, rather than p^d . No prescribed short evaluation domain is obtained.

N.19 Prime alphabets on selected multiplicative fibers

The same quotient construction applies over a prime alphabet when the image of the power map is a sufficiently large subgroup. This gives the following exact two-source profile on a chosen union of cosets.

Lemma N.35 (Fixed-cardinality products and completion). *Let $\mathcal{A}$ be a P -element subset of a finite abelian group Γ of order M , and suppose*

$$\left| \frac{1}{P} \sum_{a \in \mathcal{A}} \chi(a) \right| \leq \varepsilon \quad (\chi \neq 1).$$

Choose integers $0 < r_0 < s_0$ and $t \geq 0$ with $s_0 + 2t \leq P$ and $a = 1 - \binom{s_0}{2}/P > 0$. Set

$$L_0 = \binom{s_0}{r_0}, \quad h_0 = \frac{M-1}{aL_0} \sum_{u \geq 0} \binom{r_0}{u} \binom{s_0 - r_0}{u} \varepsilon^{2u} \leq \frac{(M-1)(1+\varepsilon)^{s_0}}{aL_0}.$$

For $0 \leq j < t$, define

$$P_j = P - s_0 - 2j, \quad \gamma_j = \frac{P\varepsilon + s_0 + 2j}{P_j}, \quad h_{j+1} = \frac{h_j^2 + \gamma_j^2 h_j}{1 - 1/P_j}.$$

If $Mh_t < 1$, there is an $(s_0 + 2t)$ -element subset of $\mathcal{A}$ whose $(r_0 + t)$ -element subset products cover Γ .

Proof. Draw s_0 independent uniform elements of $\mathcal{A}$, initially allowing repetitions. Consider the L_0 products indexed by the r_0 -subsets of the positions, and let E_0 be their ordered collision energy. For two index sets with u positions on each side of their symmetric difference, character orthogonality gives collision probability

$$\frac{1}{M} \sum_{\chi} |\hat{\mu}(\chi)|^{2u} \leq \frac{1}{M} + \frac{M-1}{M} \varepsilon^{2u},$$

where μ is uniform on $\mathcal{A}$; for $u = 0$ the right side is exactly one. For each index set there are $\binom{r_0}{u} \binom{s_0 - r_0}{u}$ partners of this type. Hence the nonnegative quantity $\Delta = ME_0/L_0^2 - 1$ satisfies

$$\mathbb{E}\Delta \leq \frac{M-1}{L_0} \sum_{u \geq 0} \binom{r_0}{u} \binom{s_0 - r_0}{u} \varepsilon^{2u}.$$

The sum is at most $(1+\varepsilon)^{s_0}$, by retaining just the terms with equal exponents in $(1+\varepsilon)^{r_0}(1+\varepsilon)^{s_0-r_0}$.

The draws are distinct with probability at least a . Since Δ is nonnegative pointwise, its conditional expectation given distinctness is at most $(\mathbb{E}\Delta)/a$. Some distinct seed therefore has $\Delta \leq h_0$. Cauchy–Schwarz bounds its product-image size below by $M/(1+\Delta)$, so its missing fraction is at most $\Delta/(1+\Delta) \leq h_0$.

Now use the classical random-translate completion strategy [1, Theorem 2], keeping cardinality fixed by adding pairs of alternatives. After j steps, $s_0 + 2j$ population elements have been used.

Every nontrivial character mean of the remaining population is at most γ_j by the triangle inequality. Let B be the complement of the current product image, of density h . For two independent uniform fresh elements u, v , normalized Fourier expansion and Parseval give

$$\mathbb{E} \frac{|Bu \cap Bv|}{M} = \sum_{\chi} |\widehat{1_B}(\chi)|^2 |\widehat{\mu_j}(\chi)|^2 \leq h^2 + \gamma_j^2(h - h^2) \leq h^2 + \gamma_j^2 h.$$

Conditioning $u \neq v$ costs at most $1/(1 - 1/P_j)$. Choose a distinct pair attaining no more than that conditional mean. Replace the product family by its translates by u and by v . Every represented product now selects exactly one of these two fresh elements, so its cardinality increases by one. The missing fraction obeys the stated recurrence, uniformly over all previous choices. After t steps it is less than $1/M$, hence zero by integrality. All final products use $r_0 + t$ distinct elements among $s_0 + 2t$ tags. $\square$

Corollary N.36 (The leading entropy threshold). *Fix $0 < \rho < 1$, $\alpha, \kappa > 0$, and $C > 1/h(\rho)$, where $h(\rho) = -\rho \log \rho - (1 - \rho) \log(1 - \rho)$. Suppose $M \rightarrow \infty$, $P \geq M^\alpha$, and $\varepsilon \leq M^{-\kappa}$ in the preceding lemma. Let $s = C \log M + O(1)$ be an integer, and let $r/s \rightarrow \rho$. For all sufficiently large M , some s -element subset of $\mathcal{A}$ has every element of Γ as an exact r -element subset product.*

Proof. Put $a_* = Ch(\rho) - 1 > 0$, $\nu = \min\{\alpha, \kappa\}$, and choose

$$\lambda = \min\{a_*/4, \nu/2\} > 0, \quad t \geq \lceil 1/\lambda \rceil$$

as fixed constants. Apply the lemma with $s_0 = s - 2t$ and $r_0 = r - t$. Then $0 < r_0 < s_0$ eventually and

$$L_0 = M^{Ch(\rho)+o(1)}, \quad a = 1 - o(1), \quad h_0 \leq M^{-a_*+o(1)} \leq M^{-2\lambda}.$$

For the fixed number of steps, deleting $O(\log M)$ elements gives $\gamma_j \leq M^{-\nu+o(1)}$, in particular $\gamma_j^2 \leq M^{-2\lambda}$ eventually, and $1 - 1/P_j \geq 1/2$. If $h_j \leq M^{-2\lambda}$, the recurrence implies

$$h_{j+1} \leq 4M^{-2\lambda} h_j \leq M^{-\lambda} h_j$$

for sufficiently large M . Thus $Mh_t \leq M^{1-(t+2)\lambda} < 1$, as required. $\square$

The counting requirement $\binom{s}{r} \geq M$ makes the leading constant $1/h(\rho)$ necessary, so this threshold is asymptotically optimal at leading order for the stated populations.

Theorem N.37. *Fix $0 < \rho < 1$, $\beta > 12/5$, and $C > 1/h(\rho)$, where $h(\rho) = -\rho \log \rho - (1 - \rho) \log(1 - \rho)$. There are infinitely many powers of two m and primes $p \equiv 1 \pmod{m}$, with $p = \Theta(m^\beta)$, for which the following holds. Put*

$$s = \lceil C \log p \rceil, \quad n = (s + 1)m, \quad J = \lfloor \rho n \rfloor.$$

There is an n -point domain $D \subset \mathbb{F}_p^$ and words $f, g : D \rightarrow \mathbb{F}_p$ such that*

$$\text{agr}_J(f) = \text{agr}_J(g) = \text{CA}_J(f, g) = J + m - 1,$$

and every $\lambda \in \mathbb{F}_p^$ satisfies*

$$\text{agr}_J(f + \lambda g) = J + 2m - 1.$$

Thus the capacity margin is $2/(s + 1) - 1/n$, the source-agreement gap is $1/(s + 1)$, and both are $\Theta(1/\log n)$. The alphabet is prime, with

$$p = \Theta\left((n/\log n)^\beta\right), \quad p/J \longrightarrow \infty.$$

The domain is a union of $s + 1$ cosets of the subgroup of order m . If $C < 2/h(\rho)$, the nearby radius is strictly below the Elias radius.

Proof. Use the prime-selection step of Krachun, Kazanin, and Haböck [21, Theorem 1 and Appendix A], following Thorner–Zaman [24, Corollary 3.1 and (3.2)], to choose $m \rightarrow \infty$ and $p \equiv 1 \pmod{m}$ with $p = \Theta(m^\beta)$. Let $H = (\mathbb{F}_p^*)^m$, of order $L = (p-1)/m$, and choose $b \in \mathbb{F}_p^* \setminus H$ and $a_0 \in H$.

For every nontrivial multiplicative character χ of $\mathbb{F}_p^*$, orthogonality over the m characters trivial on H gives

$$\sum_{a \in H} \chi(b-a) = \frac{1}{m} \sum_{\psi|_H=1} \chi(b)\psi(b)J(\psi, \chi), \quad J(\psi, \chi) = \sum_{t \in \mathbb{F}_p} \psi(t)\chi(1-t).$$

All characters, including the trivial one, are extended by zero at zero. Each Jacobi sum has absolute value at most $\sqrt{p}$: the exceptional cases $\psi = 1$ or $\psi\chi = 1$ have absolute value 1, and otherwise the Gauss-sum identity gives $\sqrt{p}$. Consequently

$$\left| \frac{1}{L-1} \sum_{a \in H \setminus \{a_0\}} \chi(b-a) \right| \leq \frac{\sqrt{p}+1}{L-1} \rightarrow 0,$$

because $L/\sqrt{p} = \Theta(m^{\beta/2-1}) \rightarrow \infty$.

Apply Corollary N.36 to $\Gamma = \mathbb{F}_p^*$ and the population $\mathcal{A} = \{b-a : a \in H \setminus \{a_0\}\}$. Here $M = p-1$, while

$$P = L-1 = p^{1-1/\beta+o(1)}, \quad \varepsilon \leq p^{-(1/2-1/\beta)+o(1)}.$$

Both exponents are positive. Write $J-1 = (r-2)m + w$, $0 \leq w < m$. Then $2 \leq r < s$ eventually, $r/s \rightarrow \rho$, and $s = C \log(p-1) + O(1)$. The corollary supplies s distinct tags $G \subset H \setminus \{a_0\}$ for which every element of $\mathbb{F}_p^*$ is a product of exactly r distinct factors $b-a$. Its constant number of completion pairs lies within these same s tags. Set $Y = X^m$ and $D = Y^{-1}(G \cup \{a_0\})$, so that $|D| = (s+1)m$.

Choose w points B_0 in the reserved fiber $Y^{-1}(a_0)$ and put $R = L_{B_0}$. Define

$$f = R \frac{Y^r - b^r}{Y - b}, \quad g = -\frac{R}{Y - b}.$$

The denominator has no root in $\mathbb{F}_p$, since $b \notin H$ and $b \neq 0$. Root counting bounds both source agreements by $J+m-1$; interpolation on $r-1$ tags of G , followed by substitution of Y and multiplication by R , attains this bound simultaneously. Thus the two individual agreements and their common agreement are exactly $J+m-1$.

For $\lambda \neq 0$, choose $S \subset G$ of size r with $V_S(b) = -\lambda$, where $V_S(Z) = \prod_{a \in S} (Z-a)$. Put $P_S(Z) = Z^r - V_S(Z)$. The witness

$$h_S = R \frac{P_S(Y) - P_S(b)}{Y - b}$$

has degree at most $J-1$ and residual $f + \lambda g - h_S = RV_S(Y)/(Y-b)$. Its zeros are the core and the r selected fibers, exactly $w + rm = J + 2m - 1$ points of D . Conversely, for every polynomial h of degree less than J ,

$$(Y-b)(f + \lambda g - h) = R(Y^r - b^r - \lambda) - (Y-b)h$$

is monic of degree $J+2m-1$, since the second term has degree at most $J+m-1$. Root counting therefore gives the matching upper bound on agreement for every λ , and equality when $\lambda \neq 0$.

Finally, $n = \Theta(m \log m)$ and $p = \Theta(m^\beta)$ give the stated field-size and characteristic asymptotics. Writing $\eta = (2m - 1)/n$, we have $\eta \log p \rightarrow 2/C$ and

$$H_p(1 - J/n - \eta) = 1 - J/n + \frac{h(\rho) - 2/C + o(1)}{\log p}.$$

This is less than $1 - J/n$ when $C < 2/h(\rho)$, proving the strict comparison with Elias. $\square$

Lemma N.35 gives an explicit finite energy-and-completion criterion. A separate sufficient finite route is the character-bias and Cauchy argument of Lemmas N.32 and N.30, applying the former lemma's concentration proof to this population. For any prime p and divisor $m \geq 2$ of $p - 1$, put $L = (p - 1)/m$ and $\theta = r/s$. This route succeeds when $s < L - 1$, $2 \leq r < s$, and

$$\frac{\sqrt{p} + 1}{L - 1} \leq \frac{1}{4}, \quad 4(p - 2)e^{-s/64} < 1, \quad (p - 2)(s + 1)e^{-\theta(1-\theta)s/2} < 1.$$

At threshold $J + 2m - 1$, exactly $p - 2$ of the mixtures $(1 - t)f + tg$, $t \in \mathbb{F}_p$, are nearby, and both endpoints are farther by m/n . The two-far-input pattern already appears in Corollary 4.16; no priority claim is made for that pattern. The parameter combination here is a polynomial prime alphabet, inverse-logarithmic capacity and source gaps, and almost-complete coverage strictly below Elias when $1/h(\rho) < C < 2/h(\rho)$. The domain is chosen as a union of multiplicative cosets and need not be a subgroup. The theorem makes no claim for a prescribed evaluation domain or for a fixed positive capacity margin.

Corollary N.38 (Two-point fibers over every large prime). *Fix $0 < \rho < 1$ and $C > 1/h(\rho)$, where $h(\rho) = -\rho \log \rho - (1 - \rho) \log(1 - \rho)$. For every sufficiently large odd prime p , put*

$$s = \lceil C \log p \rceil, \quad n = 2(s + 1), \quad J = \lfloor \rho n \rfloor.$$

There are an n -point domain $D \subset \mathbb{F}_p^$, which is a union of pairs $\{x, -x\}$, and words $f, g : D \rightarrow \mathbb{F}_p$ with*

$$\text{agr}_J(f) = \text{agr}_J(g) = \text{CA}_J(f, g) = J + 1, \quad \text{agr}_J(f + \lambda g) = J + 3 \quad (\lambda \in \mathbb{F}_p^*).$$

At radius $\theta = 1 - J/n - \eta$, where $\eta = 3/n$, both endpoints are at distance $\theta + 2/n = \theta + (2/3)\eta$, and all $p - 2$ non-endpoint affine mixtures $(1 - z)f + zg$ are at distance exactly θ . If $C < 3/[2h(\rho)]$, this radius is strictly below Elias. Here

$$n = 2C \log p + O(1), \quad p = \exp(n/(2C) + O(1)), \quad p/J \rightarrow \infty.$$

Proof. Let $H = (\mathbb{F}_p^*)^2$, choose a nonsquare b , and reserve $a_0 \in H$. The Jacobi-sum argument in the proof of Theorem N.37 gives every nontrivial character mean of $\mathcal{A} = \{b - a : a \in H \setminus \{a_0\}\}$ absolute value at most

$$\varepsilon = \frac{\sqrt{p} + 1}{P}, \quad P = \frac{p - 3}{2}.$$

Thus $P = p^{1+o(1)}$ and $\varepsilon \leq p^{-1/2+o(1)}$ over all odd primes; no prime-progression selection is needed. Write $J - 1 = 2(r - 2) + w$, with $w \in \{0, 1\}$. Then $r/s \rightarrow \rho$. Corollary N.36, with $M = p - 1$, supplies an s -set $G \subset H \setminus \{a_0\}$ such that the products of exactly r factors $b - a$ cover $\mathbb{F}_p^*$.

Apply the quotient construction with $Y = X^2$ and $D = Y^{-1}(G \cup \{a_0\})$. Let R be the monic locator of w points of $Y^{-1}(a_0)$, and set

$$f = R \frac{Y^r - b^r}{Y - b}, \quad g = -\frac{R}{Y - b}.$$

The denominator is root-free. Root counting gives each source agreement at most $J + 1$, while interpolation on any $r - 1$ tags, together with the core roots, attains $J + 1$ simultaneously. For $\lambda \neq 0$, choose $S \subset G$, $|S| = r$, with $V_S(b) = -\lambda$. The same quotient witness as in that theorem has residual $RV_S(Y)/(Y - b)$ and hence $J + 3$ agreements. For every polynomial h of degree less than J ,

$$(Y - b)(f + \lambda g - h) = R(Y^r - b^r - \lambda) - (Y - b)h$$

is monic of degree $J + 3$, because the second term has degree at most $J + 1$. This proves the matching upper bound. Scaling by $1 - z$ transfers the nonzero-pencil statement to every $z \notin \{0, 1\}$.

Finally $\eta \log p \rightarrow 3/(2C)$, so

$$H_p(1 - J/n - \eta) = 1 - J/n + \frac{h(\rho) - 3/(2C) + o(1)}{\log p} < 1 - J/n$$

when $C < 3/[2h(\rho)]$. This is the claimed strict Elias comparison. $\square$

A 48-coordinate certificate strictly below Elias. There is also an explicit instance over $\mathbb{F}_{65537}$ with $m = 2$, $s = 23$, $n = 48$, and $J = 24$. Take $b = 3$, reserve $a_0 = 1$, and choose the following square tags:

$$G = \{25703, 16562, 45729, 40011, 19330, 14018, 39598, 39860, \\ 36881, 40752, 56378, 11224, 11940, 40844, 22047, 23180, \\ 33705, 65375, 48969, 314, 27487, 55681, 29802\}.$$

Every element of $\mathbb{F}_{65537}^*$ is a product of exactly 13 distinct factors from $\{3 - a : a \in G\}$, as certified by exact finite enumeration. On $D = \{x : x^2 \in G \cup \{1\}\}$, put

$$f = (X - 1) \frac{X^{26} - 3^{13}}{X^2 - 3}, \quad g = -\frac{X - 1}{X^2 - 3}.$$

The preceding algebra gives exact source and common agreement 25, and exact agreement 27 for every nonzero pencil label. Thus both endpoints have distance $23/48$, while all 65535 interior affine mixtures have distance $7/16$. The distance below capacity is exactly $1/16$, and the proximity-gap loss is exactly $1/24$. This radius is strictly below Elias: the entropy inequality $H_{65537}(7/16) < 1/2$ is equivalent to the integer inequality

$$65536^7 16^{16} < 65537^8 7^7 9^9,$$

which holds exactly. The certificate supplies the domain and a 13-subset witness for each of the 65536 nonzero products. Independent dynamic programming, meet-in-the-middle enumeration, and direct field multiplication agree on coverage; each product occurs between 3 and 36 times among the $\binom{23}{13} = 1144066$ subsets. The replay files are in `research/norm_quotient_coverage/below_elias_prime_fixture`.

The challenge field matters. Scalar extension of this particular certificate to $E = \mathbb{F}_{65537^d}$ gives exact agreement 25 at every $\lambda \in E \setminus \mathbb{F}_{65537}$, not 27. More generally, for any F -linear code C , its scalar extension C_E , and $f, g \in F^D$, one has

$$\text{agr}_{C_E}(f + \lambda g) = \text{CA}_C(f, g) \quad (\lambda \in E \setminus F).$$

Indeed, extend $1, \lambda$ to an F -basis of E and project an extension-code witness onto these two coordinates. The projections belong to C and agree simultaneously with f, g wherever that witness agrees with $f + \lambda g$. Conversely, combining maximizing common witnesses attains equality. Thus the finite example has exactly 65536 nearby pencil labels even over E , a fraction $65536/65537^d$. This conclusion accounts for all extension-field witnesses.

A separate character-sum certificate. The accompanying certificate gives $p = 65537$, $m = 2$, $b = 3$, $a_0 = 1$, and $s = 1420$ explicitly listed square tags. It has $n = 2842$ and $J = 1421$, with

$$\text{agr}_J(f) = \text{agr}_J(g) = \text{CA}_J(f, g) = 1422, \quad \text{agr}_J(f + \lambda g) = 1424 \quad (\lambda \neq 0).$$

Here $r = 712$, $w = 0$, and the words are $f = (X^{1424} - 3^{712})/(X^2 - 3)$ and $g = -1/(X^2 - 3)$ on the listed domain. The denominator is root-free because 3 is a nonsquare.

Coverage is certified with integer arithmetic. Take discrete logarithms base 3 of the 1420 factors $3 - a$ in the cyclic group of order 65536, and let $c_3(t)$ count ordered triples of these exponents with sum t . The supplied convolution has

$$\sum_t c_3(t)^2 = 125113995212860.$$

Character orthogonality then gives

$$\sum_{\chi \neq 1} \left| \sum_{a \in G} \chi(3 - a) \right|^6 = 1052619325992960 < 710^6.$$

Thus every nontrivial character sum is below $s/2 = 710$. For $\theta = 712/1420$, the Cauchy exponent is $u = 63012/355$; the rational inequality $u^5/120 > 93125235 = (p - 2)(s + 1)$ proves the required positivity since $e^u > u^5/120$. Two independent integer convolution implementations reproduce every coefficient. In particular, this verifies all 65536 nonzero pencil labels without enumerating their witnesses. The normalized capacity margin is $3/2842$, and the common-agreement loss is $2/2842$. The domain, coefficient data, and replay programs are supplied in `research/norm_quotient_coverage/prime_finite_fixture`.

N.20 Finite extension alphabets on selected base-field domains

The same compiler has finite instances at length 2^{18} and dimension 2^{17} over cubic, quartic, quintic, and sextic alphabets. Their evaluation points remain in the prime subfield. The assertion is the existence of a selected union of fibers; it does not specify a list of tags or establish the profile on a prescribed NTT subgroup.

Proposition N.39. *Let p be prime, $d \geq 2$, $E = \mathbb{F}_{p^d}$, and $m \geq 2$ divide $p - 1$. Choose $b \in E$ of degree d over $\mathbb{F}_p$, and put*

$$P = (p - 1)/m - 1, \quad M = p^d - 1, \quad \varepsilon = \frac{d\sqrt{p} + 1}{P}.$$

Let $2 \leq r \leq s \leq P$ and $0 \leq w < m$, and set

$$n = (s + 1)m, \quad J = (r - 2)m + w + 1.$$

Suppose an integer $t \geq 0$, with $s_0 = s - 2t$ and $r_0 = r - t$, satisfies the finite criterion of Lemma N.35 for the displayed P, M, ε , in particular $Mh_t < 1$. Then there are an n -point set $D \subset \mathbb{F}_p^$ and words $f, g : D \rightarrow E$ such that*

$$\text{agr}_J(f) = \text{agr}_J(g) = \text{CA}_J(f, g) = J + m - 1,$$

and

$$\text{agr}_J(f + \lambda g) = J + 2m - 1 \quad (\lambda \in E^*).$$

The domain is a union of $s + 1$ fibers of $X \mapsto X^m$. Thus all $|E| - 2$ non-endpoint affine mixtures of the two far sources are at exactly the tested distance. The capacity margin and source separation are, respectively, $(2m - 1)/n$ and m/n .

Proof. Let $H = (\mathbb{F}_p^*)^m$ and reserve $a_0 \in H$. For every nontrivial character χ of E^* and every character ψ of $\mathbb{F}_p^*$, Katz's character-sum estimate on the finite etale algebra $E \times \mathbb{F}_p$, at its regular element $(b, 0)$, gives [22, Theorem 2 and its proof]

$$\left| \sum_{a \in \mathbb{F}_p^*} \chi(b-a)\psi(a) \right| \leq d\sqrt{p}.$$

The algebra has dimension $d+1$; the full-degree hypothesis makes the d conjugates of b distinct and nonzero, and the nontriviality of χ ensures geometric nontriviality. Averaging over the m characters trivial on H , and then removing a_0 , bounds every nontrivial character mean of

$$\mathcal{A} = \{b - a : a \in H \setminus \{a_0\}\} \subset E^*$$

by ε . This population has exactly P distinct elements. The finite completion lemma supplies s distinct tags $G \subset H \setminus \{a_0\}$ whose exact r -element products of factors $b - a$ cover E^* .

Set $Y = X^m$, $D = Y^{-1}(G \cup \{a_0\})$, and let R be the monic locator of w points in the reserved fiber $Y^{-1}(a_0)$. Define

$$f = R \frac{Y^r - b^r}{Y - b}, \quad g = -\frac{R}{Y - b}.$$

The denominator has no base-field root. The monic polynomial f has degree $J + m - 1$. For $\deg h < J$, the nonzero cleared numerator $-R - (Y - b)h$ has degree at most $J + m - 1$ as well. Interpolating the two quotient values on any $r - 1$ tags, and multiplying by R , gives simultaneous degree-at-most- $(J - 1)$ witnesses with exactly $w + (r - 1)m = J + m - 1$ common agreements. This proves the source and common-agreement equalities.

For $\lambda \neq 0$, choose $S \subset G$, $|S| = r$, with $V_S(b) = -\lambda$, and put $P_S(Z) = Z^r - V_S(Z)$. The witness

$$h_S = R \frac{P_S(Y) - P_S(b)}{Y - b}$$

has degree at most $w + (r - 2)m = J - 1$ and residual $RV_S(Y)/(Y - b)$, which vanishes exactly on the core and the r selected fibers. This gives $w + rm = J + 2m - 1$ agreements. For any other degree-below- J polynomial h ,

$$(Y - b)(f + \lambda g - h) = R(Y^r - b^r - \lambda) - (Y - b)h$$

is monic of degree $J + 2m - 1$, since the second term has degree at most $J + m - 1$. Root counting proves the matching upper bound. $\square$

Four exact finite instances. Put

$$p_G = 2^{64} - 2^{32} + 1, \quad p_B = 2013265921, \quad p_K = 2130706433.$$

All four cases below have $n = 262144$, $J = 131072$, and $w = m - 1$. Write $A = J + m - 1$ and $T = J + 2m - 1$ for the exact source/common and nonzero-pencil agreements. The parameter t counts completion pairs.

E	m	t	A	T	capacity margin	source separation
$\mathbb{F}_{p_G^3}$	1024	4	132095	133119	2047/262144	1/256
$\mathbb{F}_{p_B^4}$	1024	6	132095	133119	2047/262144	1/256
$\mathbb{F}_{p_B^5}$	1024	22	132095	133119	2047/262144	1/256
$\mathbb{F}_{p_K^6}$	512	0	131583	132095	1023/262144	1/512

For the first three cases, $(s, r) = (255, 129)$; for the last, $(s, r) = (511, 257)$. The proof uses the upper bounds $\sqrt{p_G} < 2^{32}$, $\sqrt{p_B} < 44870$, and $\sqrt{p_K} < 46160$. Exact rational evaluation of the seed majorant $(1 + \varepsilon)^{s_0}$ and the completion recurrence gives, in table order,

$$Mh_t < 2^{-21}, \quad Mh_t < 2^{-2}, \quad Mh_t < 2^{-2}, \quad Mh_t < 2^{-86}.$$

Every recurrence step is rounded upward to denominator 2^{512} . These inequalities are integer comparisons, not floating-point tests. Primality is certified by Proth witnesses 7, 11, 3, respectively: $p_G - 1 = (2^{32} - 1)2^{32}$, $p_B - 1 = 15 \cdot 2^{27}$, and $p_K - 1 = 127 \cdot 2^{24}$.

The script and exact receipts are supplied in `research/norm_quotient_coverage/verify_extension_native_short_domains.py` and `extension_native_short_domains_verified.json`. They certify existence of the selected tag sets through the finite lemma, without enumerating a domain or an extension field. The characteristic exceeds J in every case. These statements do not provide the prescribed-domain agreement threshold 139782. Their contribution is a finite existence guarantee at the displayed fields and length, not a larger asymptotic distance from capacity than prior inverse-logarithmic counterexamples. They establish neither a better codes improvement nor tightness of the prime-field upper bounds.

Counting ceiling for this construction. This compiler assigns one challenge label to each r -element subset of its s tags. Covering E^* therefore requires

$$|E| - 1 \leq \binom{s}{r} \leq 2^s.$$

Consequently its capacity margin satisfies

$$\frac{2m - 1}{n} < \frac{2}{s + 1} \leq \frac{2}{1 + \log_2(|E| - 1)}.$$

At the fixed length $n = 2^{18}$, the condition $n = (s + 1)m$ makes m a power of two. If $m \geq 2048$, then $s \leq 127$, whereas the exact integer comparison

$$\binom{127}{63} < 2013265921^4 - 1$$

holds. The right side is the smallest native multiplicative-group size in the table. Thus none of the four fields can have full native coverage with $m \geq 2048$ through these subset witnesses, regardless of how the tags are chosen or coverage is proved. The first three rows already attain the largest permitted $m = 1024$ and capacity margin 2047/262144 in this model. The last row has counting room for at most that same margin. This is a limitation of the displayed witness family, not an upper bound on all nearby codewords or on other counterexample constructions.

N.21 Targeted products and norm denominators

The product argument can target a specified subfield rather than merely a large, unspecified part of the auxiliary field. This permits a higher-degree denominator while keeping a single scalar challenge.

Lemma N.40 (Products in a specified target set). *Let $\mathcal{A}$ be a P -element subset of a finite abelian group Γ of order M , and suppose every nontrivial character mean has absolute value at most $\varepsilon \leq 1$. Let $1 \leq r \leq s \leq P$, $L = \binom{s}{r}$, and $a = 1 - \binom{s}{r}/P > 0$. Set*

$$b_u = \binom{r}{u} \binom{s-r}{u}, \quad V = \frac{M-1}{L} \sum_u b_u \varepsilon^{2u} + \frac{(M-1)(M-2)}{L} \sum_u b_u \varepsilon^{r+u}.$$

Here $0 \leq u \leq \min\{r, s - r\}$ and $\varepsilon^0 = 1$. For every specified $B \subset \Gamma$, some s distinct elements of $\mathcal{A}$ have at least

$$|B| - \lfloor |B|V/a \rfloor$$

distinct elements of B among their exact r -element subset products. In particular they cover B if $|B|V/a < 1$.

Proof. Draw s independent uniform elements X_i of $\mathcal{A}$. For $z \in \Gamma$, let Z_z count the r -subsets of positions whose product is z , and put $W_z = MZ_z/L$ and $a_\chi = \mathbb{E}\chi(X_i)$. Orthogonality gives

$$W_z - 1 = \frac{1}{L} \sum_{|S|=r} \sum_{\chi \neq 1} \chi(z)^{-1} \prod_{i \in S} \chi(X_i).$$

For two supports with $|S \setminus T| = |T \setminus S| = u$, a term in the ordinary square of this real quantity has expectation

$$(\chi\psi)(z)^{-1} a_{\chi\psi}^{r-u} a_\chi^u a_\psi^u.$$

There are Lb_u ordered support pairs. Among ordered pairs $\chi, \psi \neq 1$, the $M - 1$ with $\chi\psi = 1$ contribute absolute value at most ε^{2u} each; the remaining $(M - 1)(M - 2)$ contribute at most ε^{r+u} each. Consequently $\mathbb{E}(W_z - 1)^2 \leq V$, uniformly in z . On the event $Z_z = 0$ the square equals one. The expected number of missing elements of B is therefore at most $|B|V$. The probability that all draws are distinct is at least a ; conditioning this nonnegative missing count on distinctness costs at most $1/a$. Some distinct sample has missing count at most $|B|V/a$, and rounding that integer proves the claim. No independence between subset products is assumed. $\square$

Proposition N.41 (Norm denominators over the native subfield). *Let p be prime, $d \geq 2$, $e \mid d$, and*

$$E = \mathbb{F}_{p^d}, \quad K = \mathbb{F}_{p^{d/e}}.$$

Let $m \geq 2$ divide $p - 1$, choose $b \in E$ of degree d over $\mathbb{F}_p$, and set

$$P = (p - 1)/m - 1, \quad M = p^d - 1, \quad \varepsilon = (d\sqrt{p} + 1)/P.$$

Choose $e + 1 \leq r \leq s \leq P$, $0 \leq w < m$, and define

$$n = (s + 1)m, \quad J = w + (r - e - 1)m + 1.$$

Suppose $\varepsilon \leq 1$, $a > 0$, and $V/a < 1$ in Lemma N.40. There are an n -point domain $D \subset \mathbb{F}_p^$ and K -valued words f, g on D such that*

$$\text{agr}_J(f) = \text{CA}_J(f, g) = J + m - 1, \quad J + m - 1 \leq \text{agr}_J(g) \leq J + em - 1,$$

while at least

$$|K| - 1 - \lfloor (|K| - 1)V/a \rfloor$$

distinct $\lambda \in K^$ satisfy*

$$\text{agr}_J(f + \lambda g) = J + (e + 1)m - 1.$$

All agreements are with the dimension- J Reed–Solomon code over the native alphabet K . The capacity margin is $((e + 1)m - 1)/n$; the common-agreement separation is em/n , and the guaranteed separation of both sources is m/n .

Proof. Let $H = (\mathbb{F}_p^*)^m$ and reserve $a_0 \in H$. The mixed character-sum bound used in Proposition N.39 gives nontrivial character bias at most ε for

$$\mathcal{A} = \{b - a : a \in H \setminus \{a_0\}\} \subset E^*.$$

Apply Lemma N.40 with the specified target set $B = K^*$ to obtain s distinct tags G . Put $Y = X^m$ and $D = Y^{-1}(G \cup \{a_0\})$, and let R be the monic locator of w points in the reserved fiber. The polynomial

$$T(Z) = \text{Norm}_{E/K}(Z - b) = \prod_{j=0}^{e-1} (Z - b^{p^{(d/e)j}}) \in K[Z]$$

is monic of degree e and has no root in $\mathbb{F}_p$. Define the native- K words

$$f = RY^{r-e}, \quad g = \frac{R}{T(Y)}.$$

The polynomial f is monic of degree $J + m - 1$. For any degree-below- J polynomial h , the nonzero polynomial $R - T(Y)h$ has degree at most $J + em - 1$. This bounds the agreement of g . Interpolate each of Z^{r-e} and $1/T(Z)$ on any $r - e$ tags of G , then substitute Y and multiply by R . The two resulting polynomials have degree at most $J - 1$ and agree with their respective sources on the same $w + (r - e)m = J + m - 1$ points. Root counting for $f - h$ proves both the exact source and common-agreement claims.

For each represented $\lambda \in K^*$, take $S \subset G$ with $|S| = r$ and $V_S(b) = \lambda$, where $V_S(Z) = \prod_{a \in S} (Z - a) \in \mathbb{F}_p[Z]$. Since λ is fixed by $\text{Gal}(E/K)$, T divides $V_S - \lambda$ in $K[Z]$. Thus

$$h_S = R \frac{T(Y)Y^{r-e} + \lambda - V_S(Y)}{T(Y)} \in K[X]$$

has degree at most $w + (r - e - 1)m = J - 1$: its two monic degree- r terms cancel before division by the degree- e denominator. Its residual is $RV_S(Y)/T(Y)$, which vanishes on exactly the core and the r selected fibers. This supplies $w + rm = J + (e + 1)m - 1$ agreements. For every other degree-below- J polynomial h , the cleared numerator

$$R(T(Y)Y^{r-e} + \lambda) - T(Y)h$$

is monic of this degree, whereas the second term has degree at most $J + em - 1$. Root counting proves the matching upper bound. $\square$

A finite tradeoff at KoalaBear. Take $p = 2130706433$, auxiliary degree $d = 6$, $n = 262144$, $J = 131072$, $m = 1024$, $s = 255$, $w = 1023$, and $r = 128 + e$. Using $\sqrt{p} < 46160$, exact rational evaluation gives $V/a < 2^{-22}$ for each $e \in \{1, 2, 3, 6\}$. Each row therefore has at least $(|K| - 1) - \lfloor (|K| - 1)/2^{22} \rfloor$ good native labels.

K	e	$\text{CA}_J(f, g)$	$\text{agr}_J(g)$ upper bound	exact near agreement	capacity numerator
$\mathbb{F}_{p^6}$	1	132095	132095	133119	2047
$\mathbb{F}_{p^3}$	2	132095	133119	134143	3071
$\mathbb{F}_{p^2}$	3	132095	134143	135167	4095
$\mathbb{F}_p$	6	132095	137215	138239	7167

The capacity numerators are divided by n . These are selected-domain existence results, not explicit tag lists or assertions about a prescribed NTT subgroup. The auxiliary extension is not the code

alphabet. For a matched comparison, the one-pole construction over $\mathbb{F}_p, \mathbb{F}_{p^2}, \mathbb{F}_{p^3}$ admits respectively $m = 4096, 2048, 2048$ in the same pointwise-moment test, with capacity numerators 8191, 4095, 4095. The norm construction changes the tradeoff between label coverage, common agreement, and the weaker of the two source separations; it does not uniformly improve these rows. The exact arithmetic and comparison are recorded in `research/norm_quotient_coverage/verify_target_subset_norm_compiler.py`.

The loss ratios alone are also obtainable from the one-pole construction by increasing its decoding dimension. If its fiber size is $m' = em$, increasing the dimension by $t = (e - 1)m$ gives the same bounds $J + m - 1$, $J + em - 1$, and $J + (e + 1)m - 1$ for common agreement, rational-source agreement, and nearby agreement, respectively. For fixed e and growing fibers this recovers the norm construction's asymptotic loss tradeoff, including its below-Elias parameter window. The norm construction is therefore not asserted to improve that asymptotic tradeoff; its remaining distinctions here are the targeted product estimate and the finite divisibility and coverage certificates.

O A constant first-order list for the quadratic indicator

The quadratic-residue indicator gives an elementary obstruction for the entire nearest-list problem of the full-field Dickson word, not only for the explicit candidate family used in the lower bound. It imposes no such restriction on general received words or domains.

Proposition O.1. *Let $k \geq 1$ and let F contain all $4k$ th roots of unity. Suppose F has characteristic zero or characteristic $p > 8k/3$. Put $w(x) = (1 + x^{2k})/2$ on $\mu_{4k} \subset F$. The polynomials of degree at most k agreeing with w at more than $5k/3$ points are exactly $0, 1$ and*

$$P_{u,v}(X) = \frac{X^k - u}{v - u}, \quad u^2 = -1, \quad v^2 = 1.$$

Each of these six polynomials has exactly $2k$ agreements. The same classification holds over every extension coefficient field. With the strict degree bound $\deg P < k$, only $0, 1$ remain.

Proof. For a nonconstant candidate P of degree $\ell \leq k$, let z, o count its correct zeros on $x^{2k} = -1$ and correct ones on $x^{2k} = 1$. Put $A = z + o$ and $h = 2k - A$. Root counting gives $A \leq 2\ell \leq 2k$, and the hypothesis gives $0 \leq h < k/3$. Let U, V be the monic locators of the two agreement sets, and write

$$P = UR, \quad P - 1 = VS, \quad X^{2k} + 1 = UC, \quad X^{2k} - 1 = VD.$$

Set $a = RS$ and $b = CS - DR$. Subtraction gives $Pb + DR = 2a$. Since $\deg a = 2\ell - A \leq h$ and $\deg(DR) = \ell + h$, we have $\deg b \leq h$. Also $b \neq 0$: otherwise $D = 2S$, contradicting $\deg D = 2k - o > \ell - o = \deg S$. Substitution yields

$$bP^2 - (b + 2a)P + a(X^{2k} + 1) = 0.$$

Divide a, b by their gcd, retaining these names, and set $\alpha = \deg a$, $\beta = \deg b$, $Y = 2bP - b - 2a$, and $E = b^2 + 4a^2$. Then $\alpha, \beta \leq h$, $\gcd(a, b) = 1$, and

$$Y^2 + 4abX^{2k} = E. \tag{69}$$

Suppose $E \neq 0$. Since $\deg E \leq 2h < 2k$, we have $\deg Y = k + (\alpha + \beta)/2$. Also $\gcd(ab, E) = 1$. The common gcd of the terms in (69) is therefore X^{2t} for some $t \geq 0$: if E vanishes at zero, its order is less than $2k$ and equals the even order of Y^2 . After division the identity is pairwise coprime. In

positive characteristic all degrees are less than p , since $2k + 2h < 8k/3 < p$. In either characteristic polynomial abc gives

$$2k + \alpha + \beta - 2t \leq (\deg Y - t) + \alpha + \beta + (\deg E - 2t),$$

where the possible root zero contributes one to the radical degree and cancels the minus one in abc. Hence $k \leq (\alpha + \beta)/2 + \deg E - t \leq 3h$, a contradiction. Here the required form of abc follows directly from the nonzero Wronskian $F'G - FG'$ for a coprime identity $F + G = H$: the repeated-root parts of all three polynomials divide this Wronskian. A nonconstant term excludes a zero Wronskian in characteristic zero; in positive characteristic the degree bound less than p does so.

Thus $E = 0$. As F contains a square root of -1 , we have $b = \lambda a$ with $\lambda = \pm 2\sqrt{-1}$. Cancellation gives

$$(2\lambda P - \lambda - 2)^2 = -4\lambda X^{2k},$$

so unique factorization implies $P = cX^k + d$, with $c \neq 0$. Such a polynomial agrees with the zero value on at most one quartic coset, and with the one value on at most one quartic coset. Each coset has size k , so $A > 5k/3 > k$ forces both matches and $A = 2k$. The zero coset has $X^k = u$ with $u^2 = -1$, and the one coset has $X^k = v$ with $v^2 = 1$, giving precisely the four displayed candidates. The constant cases are immediate. Finally $A \geq k + 1$ base-field agreements force base-field coefficients by interpolation, proving the extension assertion. $\square$

The cutoff fraction $5/12$ is strictly below the quarter-rate first-order threshold $(3 + \sqrt{133})/31$. Thus this received word cannot witness growing first-order lists, even though its nearest agreement fraction is $1/2$.

Corollary O.2. *Under the hypotheses of Proposition O.1, and with characteristic different from five, the Dickson received word $W(x) = (x^k - 1)^2/2$ on μ_{4k} has agreement at most $\lfloor 5k/3 \rfloor$ with every degree- $< k$ polynomial, even with extension coefficients. In particular its maximum agreement fraction is at most $5/12$, below the quarter-rate first-order threshold.*

Proof. Since $W = w - X^k$, a candidate exceeding the cutoff would give a monic degree- k polynomial $P + X^k$ in Proposition O.1. The leading coefficient $c = 1/(v - u)$ of each nonconstant listed polynomial satisfies $4c^4 + 1 = 0$, since $(v - u)^4 = -4$. It cannot equal one unless the characteristic is five. $\square$

In particular this applies to every prime-field cyclic domain $\mu_{4k} \subseteq \mathbb{F}_p^*$ with $p > 5$, including the full domain when $p = 4k + 1$. Enlarging the candidate bank for this cyclic word cannot reach the first-order regime. This does not rule out arbitrary puncturing or changing the word or evaluation domain. As a finite check, exhaustive enumeration of all 17^5 polynomials of degree at most four over $\mathbb{F}_{17}$ finds no polynomial with seven agreements and exactly the stated six with eight. The proof above, rather than this census, establishes the uniform bound.

P Linear exception counts from leading-degree separation

A first-order identity can force a linear exceptional-label bound even when every agreement coordinate is singular. The criterion below separates the nonlinear terms from the highest coefficient equations. Those equations determine the nonconstant message coefficients with one denominator of degree $O(D)$. The residual solution locus then has bounded degree in one remaining message parameter and degree $O(D)$ in the challenge. The conclusion concerns actual solutions of the identity: an application to all nearby Reed–Solomon candidates must first establish that they all satisfy such an equation.

Theorem P.1 (Leading-degree separation). *Let $\mathbb{F}$ be any field, $D \geq 1$, and $\ell \geq 2$. Consider*

$$T(X, z)P' + \sum_{j=0}^{\ell} A_j(X, z)P^j = 0, \quad \deg P \leq D. \quad (70)$$

Suppose $A_\ell \neq 0$, $\deg_X T = t$ with leading coefficient $\tau(z) \neq 0$, and

$$\deg_X A_1 \leq t - 1, \quad \deg_X A_0 \leq t + D - 1, \quad \deg_X A_j + jD < t \quad (2 \leq j \leq \ell, A_j \neq 0). \quad (71)$$

Assume all coefficient polynomials have z -degree at most an integer $a \geq 0$, and that every pivot

$$d_i(z) = i\tau(z) + (A_1)_{t-1}(z), \quad 1 \leq i \leq D,$$

is a nonzero polynomial. Put $M = a(\ell D + 1)$ and $L = aD + 1$. On any n distinct evaluation points, for any received line $f + zg$ and integer $A > D$, outside a set of at most

$$a(D + 1) + 2\ell M + \frac{n(M + \ell L)}{A - D} + \ell n \quad (72)$$

labels, every solution of (70) with at least A agreements admits degree-at-most- D witnesses F_0, G_0 such that $P = F_0 + zG_0$ and its full agreement set is

$$\{x : f(x) = F_0(x), \quad g(x) = G_0(x)\}.$$

In particular the bound is $O_{a,\ell,\eta}(n)$ if $A - D \geq \eta n$ and $D \leq n$.

Proof. Write $P = c + \sum_{i=1}^D p_i X^i$. By (71), the nonlinear terms have degree less than t . Equating coefficients of X^{t+i-1} , in decreasing i , gives a triangular linear system with pivot d_i . It does not involve c . Every off-diagonal coefficient and inhomogeneous term has z -degree at most a .

Set $\Delta_i = \prod_{h=i}^D d_h$ and $\Delta = \Delta_1$. Descending induction gives

$$p_i = N_i / \Delta_i, \quad \deg_z N_i \leq a(D - i + 1).$$

Indeed, all higher denominators divide Δ_{i+1} , so multiplying the next equation by that product gives precisely the stated degree. Passing to the single denominator Δ yields

$$P = c + \frac{W(X, z)}{\Delta(z)}, \quad W(0, z) = 0, \quad \deg_z W, \deg_z \Delta \leq aD. \quad (73)$$

Discard the at most aD roots of Δ . On the complement this is an exact parametrization of the actual candidate locus by (z, c) . No division by the integer i is required: invertibility of the specified pivot is the only characteristic condition needed.

After substituting (73) and multiplying by Δ^ℓ , the remaining coefficient equations are the coefficients in X of

$$\Delta^{\ell-1} T W' + \sum_{j=0}^{\ell} A_j (c\Delta + W)^j \Delta^{\ell-j}.$$

Each has bidegree at most (M, ℓ) in (z, c) . Choose a nonzero coefficient $h(z)$ of A_ℓ in X and also discard its at most a roots. The corresponding residual has leading c -coefficient $h\Delta^\ell$, nonzero on this open set.

Work over an algebraic closure and take the gcd U of all nonzero residuals. Its factors give their common curve components. Since U divides the chosen residual, its summed bidegrees are at most

(M, ℓ) . Every vertical factor is supported on a root of $h\Delta$ and has already been excluded. There are therefore at most ℓ reduced horizontal components.

Divide the residuals by U . Their gcd is one. If one quotient is a nonzero constant, there are no common zeros off U . Otherwise choose a nonzero quotient and a generic constant linear combination of the other quotients sharing none of its irreducible factors. Such a combination exists because each factor is avoided by some quotient, and the algebraic closure is infinite. Both polynomials have bidegrees at most (M, ℓ) . Their proper intersection in $\mathbb{P}_z^1 \times \mathbb{P}_c^1$ has total multiplicity at most $2M\ell$. Every isolated point of the original locus lies off its curve locus and in this intersection. Thus at most $2M\ell$ additional labels come from isolated points. This argument removes common factors before using intersection bounds; linear independence alone would not suffice.

At an evaluation coordinate x , agreement is the equation

$$\Delta c + W(x, z) - \Delta(f(x) + zg(x)) = 0, \quad (74)$$

of bidegree at most $(L, 1)$. For a curve component of bidegrees (s, r) , unless agreement is identically true, its resultant with (74) has degree at most $s + rL$. Away from $\Delta = 0$, each root gives at most one agreeing point because the agreement equation is linear in c . Summing over components gives at most $M + \ell L$ nonpersistent agreeing points per coordinate.

If a component has at most D persistent agreement coordinates, each A -near point on it supplies at least $A - D$ nonpersistent incidences. Their labels number at most $n(M + \ell L)/(A - D)$. If it has more than D persistent coordinates, interpolation forces its entire candidate polynomial to be a base-field affine codeword graph $F_0 + zG_0$. Outside at most n labels having an extra agreement, its full agreement set is exactly the coefficientwise common agreement set of those witnesses. There are at most ℓ components. Adding the discarded labels, isolated points, and curve contributions proves (72). $\square$

Corollary P.2 (The unnormalized two-branch family). *Let T be the monic squarefree locator of the n evaluation points, assume $n > 4D$, and suppose the characteristic is zero or $p > D$. Let $U_z, V_z \in \mathbb{F}[X]$ be affine in z , each of X -degree at most $n - 1$. For solutions of*

$$TP' = (P - U_z)(P - V_z), \quad \deg P \leq D, \quad (75)$$

on any chosen received line and threshold $A > D$, the full common-witness property holds outside at most

$$18D + 10 + \frac{n(8D + 4)}{A - D} + 2n$$

labels. Thus the family has only $O_\eta(n)$ exceptions when $A - D \geq \eta n$.

Proof. Any nonzero coefficient of $U_z V_z$ above X -degree $n + D - 1$ is a polynomial in z of degree at most two, and every solution label must be its root. Otherwise the generic product degree is at most $n + D - 1$. If both generic branch degrees exceed D , each is at most $n - 2$. Apply Theorem P.1 with $A_2 = -1$, $A_1 = U_z + V_z$, $A_0 = -U_z V_z$, $\ell = 2$ and $a = 2$. The pivots are the nonzero constants $1, \dots, D$. Substitution in (72) gives the displayed bound.

If a branch, say U_z , has degree at most D , it is an affine codeword pencil. Any solution $P \neq U_z$ agrees with U_z at at most D domain points and therefore with V_z at at least $n - D$ points. Such a solution is unique at each label, since two would agree at $n - 2D > D$ points. If there are at least three such labels, their candidates satisfy the same affine relation as their labels on at least $n - 3D > D$ common agreement points with V_z . It is therefore a polynomial identity. Every such candidate lies on the affine codeword pencil determined by the first two; if there are at most two labels, discard them. All solutions lie on at most two affine codeword pencils outside at most

two labels. For any received line, each pencil has the full-witness property outside at most n extra-agreement labels. This gives $2n + 2$, which is smaller than the displayed bound. $\square$

Scope and the nonlinear boundary. The general theorem imposes no nonsingular-agreement requirement and no separate characteristic restriction beyond its explicit nonzero pivots. It does impose leading-degree separation, fixed challenge and value degrees, and a derivative coefficient independent of P . Under $p > D$ there can be at most one identically zero pivot, since $d_j - d_i = (j - i)\tau$; that resonant case is not included. The quadratic isolated family of Theorem H.3 has derivative coefficient $(z - X^2)R$ of degree $D + 3$ and a P^2 coefficient of X -degree two. Its nonlinear term can reach degree $2D + 2$, so it violates the strict separation condition. Its quadratic isolated count is not a contradiction. The conclusion here is a family-specific upper bound for actual nearby solutions, not a general linear first-order MCA theorem.

Checks. Independent symbolic elimination verifies the recurrence and residual bidegrees through $D = 4$, including nonmonic moving coefficients, value degrees two through four, and persistent nonlinear components. An independent finite-field check exhausts 29968 polynomial-label pairs and checks actual agreements, including persistent codeword pencils. The scripts and reports are in `research/two_branch_recurrence/`. The proofs above, rather than the finite fixtures, establish the quantified statements.

Q Singular covers and a linear bound for monic Riccati families

This section closes a broad first-order template: bounded challenge-degree identities with constant nonzero quadratic value coefficient have only linearly many bad nearby labels above agreement $5D/3$, even when the singular coordinates move with the challenge. The conclusion concerns *actual solutions of one identity*. Applying it to an entire received line requires proving that every nearby candidate under consideration satisfies that identity.

Throughout, $D \geq 1$, the characteristic is zero or greater than D , and the domain has n distinct coordinates. Fix a received line $f + zg$ and an integer $D < A \leq n$. A solution pair (z, P) , $\deg P \leq D$, is *bad* if it has at least A agreements and no degree-at-most- D polynomial agrees with g on its full agreement support. This is precisely the obstruction to full-support common witnesses.

Theorem Q.1 (A fixed cover for singular agreements). *Let $Q(X, z, u, v) \neq 0$ have total jet degree at most $B \geq 1$ and challenge degree at most $H \geq 0$. Suppose a fixed set S of s domain coordinates and a set of at most E exceptional labels have the following property: outside those labels, every bad solution of $Q(X, z, P, P') = 0$ has at most e singular agreements outside S , where singular means $Q_v(x, z, P(x), P'(x)) = 0$. Put*

$$\kappa = \min \left\{ s, \left\lfloor \frac{2s + D}{3} \right\rfloor \right\}, \quad L = A - e - \kappa,$$

and assume $L \geq 1$. Define

$$\tau = \max\{0, 2D - 3\}, \quad b_0 = 1 + \tau(B - 1), \quad T_0 = b_0 + \tau H, \quad K_0 = H + BT_0, \quad q_0 = B + H,$$

$$C = q_0 K_0 + q_0 T_0 \frac{n - D}{A - D} + q_0 n.$$

The number of bad labels is at most

$$E + \frac{n}{L} C + n. \tag{76}$$

The last n may be omitted if $\kappa = s$. In particular this is linear in n for fixed B, H , $A - D \geq \eta n$, $L \geq \lambda n$, and $E = O(n)$. The singular-cover assumption is uniform over the actual candidates; the separant need not be independent of P .

Proof. Theorem H.5 bounds labels admitting a bad witness with at least L nonsingular agreements by nC/L . Outside those and the exceptional labels, select one bad witness per remaining label. Each has at least $A - (L - 1) - e = \kappa + 1$ agreements within S . If $\kappa = s$, none exist. Otherwise, writing $b = \kappa + 1 \leq s$, we have $3b - 2s > D$. Two selected candidates at the same label would agree on at least $2b - s \geq 3b - 2s > D$ points and hence coincide. At three distinct labels, their S -agreement sets have an intersection of size at least $3b - 2s > D$. Their polynomials consequently satisfy the same affine relation as their labels identically. If there are at least three labels, all selected candidates lie on the affine codeword pencil $P_z = F_0 + zG_0$ determined by the first two; otherwise there are at most $2 \leq n$ labels.

For one such pencil, a coordinate at which $F_0 = f$ and $G_0 = g$ agrees persistently. Every other coordinate agrees at at most one label. A bad pair must have at least one such extra agreement, because otherwise F_0, G_0 are full-support common witnesses. There are at most n labels with an extra agreement. This proves the bound. $\square$

For example, a separant $R_0(X)R_1(X, z)$ with $\deg_X R_1 \leq e$ and $\deg_z R_1 \leq h$ satisfies the hypothesis with S the domain roots of R_0 and $E \leq h$. Use one nonzero coefficient of R_1 to exclude labels at which it vanishes identically. A different incidence argument below allows *all* roots to move.

Theorem Q.2 (Bounded-challenge monic Riccati families). *Let $T(X, z) \neq 0$, and let T, R, S have challenge degree at most a fixed integer $h \geq 0$, with unrestricted X -degrees. Consider*

$$T(X, z)P' - P^2 + R(X, z)P - S(X, z) = 0, \quad \deg P \leq D. \quad (77)$$

If $A - D \geq \eta n$ and $A - 5D/3 \geq \lambda n$ for fixed positive η, λ , then the number of labels possessing a bad A -near solution is $O_{h, \eta, \lambda}(n)$. The same conclusion holds with any constant nonzero coefficient of P^2 in place of -1 .

Proof. Write $t = \deg_X T$, with leading coefficient $\tau(z) \neq 0$, and $r = \deg_X R$ (allowing $r = -\infty$).

Low derivative degree: $t \leq 2D$. Let S_0 be the domain coordinates where $T(x, z)$ vanishes identically in z , and let $s = |S_0|$. A nonzero coefficient of T as a polynomial in z shows $s \leq t \leq 2D$. At each coordinate outside S_0 , the polynomial $T(x, z)$ has at most h roots. Thus the total number of singular coordinate-label incidences outside S_0 is at most hn . Set $e = \lfloor \lambda n/2 \rfloor$. At most $hn/(e + 1) \leq 2h/\lambda$ labels have more than e such incidences. Every other label satisfies the uniform cover hypothesis of Theorem Q.1, since the separant is T . Moreover

$$L = A - e - \min\{s, \lfloor (2s + D)/3 \rfloor\} \geq A - e - 5D/3 \geq \lambda n/2.$$

Use $B = 2, H = h$ in that theorem.

High linear degree: $r \geq t > 2D$. A nonzero coefficient of S above degree $r + D$ restricts solutions to at most h labels. Otherwise equate coefficients of X^{r+i} in order $i = D, \dots, 0$. This is a triangular system determining every coefficient of P , with pivot $R_r(z)$. The derivative term involves only higher message coefficients, and P^2 is absent since $r > 2D$. Outside at most h roots of the pivot, Cramer's rule gives one possible family $P = W/\Delta$, where

$$\deg_z W, \deg_z \Delta \leq K = h(D + 1).$$

An agreement numerator has degree at most $K + 1$ in z . If at most D coordinates agree persistently, every nearby label supplies at least $A - D$ nonpersistent incidences, giving at most $n(K + 1)/(A - D)$ labels. If more than D coordinates persist, interpolation forces an affine codeword pencil, with at most n bad labels. Further residual equations can only remove candidates.

High derivative degree without an identically zero pivot. Now $t > 2D$ and $r \leq t - 1$. A nonzero coefficient of S above $t + D - 1$ again restricts solutions to at most h labels. Otherwise the top coefficient pivots are

$$d_i(z) = i\tau(z) + R_{t-1}(z), \quad 1 \leq i \leq D.$$

If each is a nonzero polynomial, Theorem P.1 applies with nonlinear degree $2D < t$, giving the required linear bound. Its exceptional set already includes roots of the nonconstant pivots.

An identically resonant pivot. It remains to treat $R_{t-1} = -j\tau$ for some $1 \leq j \leq D$. The characteristic assumption makes this j unique, and every other pivot is $(i - j)\tau$. Exclude the roots of τ . Descending coefficient elimination leaves $a_j = d$ and $a_0 = c$ free. At the resonant equation, the compatibility condition is rational in z and independent of c, d . Its numerator has degree at most $K = h(D + 1)$; if nonzero, this already bounds the labels. Otherwise the remaining equations give

$$P = c + \frac{Q_z(X) + dV_z(X)}{\delta(z)}, \quad \delta = \tau^D, \quad \deg_z(\delta, Q, V) \leq K. \quad (78)$$

Here $Q_z(0) = V_z(0) = 0$, $\deg_X Q \leq D$, $\deg_X V = j$, and the leading coefficient of V is δ . These degree bounds follow by descending induction: each step multiplies by coefficient polynomials of degree at most h and introduces at most one factor of τ in the common denominator. At most D steps are used; the bound K includes the compatibility numerator as well.

Translate X by a transcendental element w , working over $\mathbb{F}(w)$, and absorb constant terms into c . This preserves all degree bounds and the original agreement incidences. The translated Q, V again have zero constant terms; their derivatives at zero are the pretranslation derivatives at w . Since $j \neq 0$ in the field,

$$v(z) = T(w, z)V'_z(0) \neq 0.$$

Put $M = 2K + h$. Multiplication of the residual identity by δ^2 gives coefficient equations of parameter degree at most two in (c, d) and z -degree at most M . Its constant coefficient, after division by δ , is

$$-\delta c^2 + R(w, z)\delta c + T(w, z)Q'_z(0) + v(z)d - S(w, z)\delta = 0.$$

Outside the roots of v , this solves $d = N(z, c)/v(z)$, where $\deg_c N = 2$, its leading coefficient is δ , and $\deg_z N \leq M$. The total number of roots excluded by δ and v is at most $K + M$.

Substitute $d = N/v$ in every cleared residual and multiply by v^2 . The resulting plane equations in (z, c) have bidegree at most $(3M, 4)$. They are not all zero: the X^{2j} residual before substitution is

$$-\delta^2 d^2 + \alpha(z)c + \beta(z)d + \gamma(z).$$

There is no cd term because $\deg_X V = j < 2j$, and c^2 occurs only in the constant coefficient. Its substituted c^4 coefficient is therefore $-\delta^4$. In particular a vertical common component can occur only at an already excluded root of δ .

For completeness, the plane incidence bound can be applied directly. Take the common gcd of all the plane residuals. Its curve components have total bidegree at most $(3M, 4)$. After removing that gcd, a generic linear combination of the residual quotients has no common component with one fixed nonzero quotient. Intersection in $\mathbb{P}^1 \times \mathbb{P}^1$ bounds the isolated points by $24M$. Relabel the

translated domain coordinates and their received values. An agreement equation, after substitution, is

$$v(\delta c + Q_z(x) - \delta(f(x) + zg(x))) + V_z(x)N(z, c) = 0.$$

It has bidegree at most $(J, 2)$, where $J = M + K + 1$. Thus the total number of nonpersistent incidences per coordinate on the residual curves is at most $6M + 4J$. On curves with at most D persistent coordinates, each nearby point supplies at least $A - D$ such incidences. More than D persistent coordinates force an affine codeword graph by interpolation. Reconstruction is injective: c is the translated constant coefficient and d is recovered from the leading coefficient of $P - c - Q/\delta$. Hence there are at most four such graph components. Each contributes at most n bad labels. Interpolation at the persistent original domain coordinates, before translation, gives this pencil over the original field. The resonant case is therefore bounded by

$$K + 25M + \frac{n(6M + 4J)}{A - D} + 4n = O_{h,\eta}(n). \quad (79)$$

This proves all cases. Replacing $-P^2$ by a constant nonzero multiple changes the displayed nonzero leading coefficients, but none of the degree or incidence estimates. $\square$

At quarter rate the first-order threshold has $A/D \rightarrow 4(3 + \sqrt{133})/31 > 5/3$. Thus bounded-challenge monic Riccati identities cannot supply superlinear bad-label examples there at fixed agreement slack. This does not cover arbitrary implicit first-order identities, or equations with a nonconstant quadratic value coefficient. It also does not prove that an interpolating identity of this form contains every nearby candidate.

Corollary Q.3 (Polynomial gauges of low degree). *Consider*

$$T(X, z)P' - B(X, z)P^2 + R(X, z)P - S(X, z) = 0, \quad \deg P \leq D,$$

where $T, B \neq 0$, all challenge degrees are bounded by a fixed h , and B divides $T\partial_X B$ in $\mathbb{F}[X, z]$. Put $a = \deg_X B$. Assume characteristic zero or greater than $D + a$. If

$$A - \frac{5D + 4a}{3} \geq \lambda n$$

for fixed $\lambda > 0$, there are only $O_{h,\lambda}(n)$ bad labels.

Proof. Write $E = D + a$. The substitution $Y = BP$ gives

$$TY' - Y^2 + \left(R - \frac{T\partial_X B}{B}\right)Y - BS = 0, \quad \deg_X Y \leq E.$$

The divisibility assumption makes all coefficients polynomial, of bounded challenge degree. If $\deg_X T \leq 2E$, apply the fixed-cover argument directly to the original degree- D candidates. The persistent singular set has size at most $2E$, and the total singular incidences outside it are at most hn . Its triple-intersection threshold is at most $(4E + D)/3 = (5D + 4a)/3$, so the stated slack gives a linear bound.

If $\deg_X T > 2E$, use the high-degree recurrence and plane-incidence arguments of Theorem Q.2 for Y . The characteristic assumption here is needed through degree E , not just through degree D . The received values become $B(x, z)(f(x) + zg(x))$, of bounded challenge degree. Thus the agreement incidence degrees remain bounded in the message parameter and $O(E)$ in z . Remove coordinates where $B(x, z)$ vanishes identically; there are at most a . On a source curve with at most D persistent

remaining coordinates, each actual nearby candidate contributes at least $A - a - D$ nonpersistent incidences. This has positive linear slack under the hypothesis, while $E = O(n)$.

If more than D remaining coordinates persist, fix $D + 1$ of them and exclude the roots of their nonzero $B(x, z)$ values, at most $h(D + 1)$ labels. At all other labels, division by these values shows that the original polynomial P agrees with the original affine received line on the fixed $D + 1$ coordinates. Interpolation forces all such actual candidates onto one affine codeword pencil, giving at most n bad labels. There are only a bounded number of source curves; isolated points and recurrence exceptions are already $O(E)$. Labels at which $B(X, z)$ vanishes identically contribute at most h more. This proves the claim without assuming that every transformed candidate is divisible by B . $\square$

Proposition Q.4 (A linear lower example with all agreements singular). *Let $\mathbb{F}_q$ have characteristic greater than D , with n distinct domain coordinates, and choose $D < s < n$. Put $A = s + 1$ and fix a subset S_0 of s coordinates. If*

$$2^n q^{-(s-D)} (1 - 1/q)^{-(n-s)} < 1, \quad (80)$$

there is a received line with no degree-at-most- D direction polynomial agreeing on A coordinates, but with at least $n - s$ distinct nearby labels whose selected witnesses have entirely singular agreement supports for one implicit identity of jet and challenge degree two.

Proof. Write $R_0(X) = \prod_{x \in S_0} (X - x)$. At coordinates in S_0 set $f(x) = 0, g(x) = x^D$. At each remaining coordinate x_i , choose $u_i \in \mathbb{F}_q^*$ and set

$$f(x_i) = -u_i x_i, \quad g(x_i) = x_i^D + u_i.$$

At the label $z = x_i$, the candidate $P_z = zX^D$ agrees exactly on $S_0 \cup \{x_i\}$. Its full support cannot have a common direction: more than D core coordinates force that direction to be X^D , which fails at x_i .

Let $U = v - zDX^{D-1}$ and

$$Q(X, z, u, v) = R_0(X)(X - z)U + U^2.$$

Every displayed candidate satisfies $Q = 0$. On it, the separant is $Q_v = R_0(X)(X - z)$, which vanishes at every agreement coordinate and has at most one zero outside the fixed set S_0 . The fixed-cover hypothesis therefore holds with $e = 1$ and no exceptional labels for these witnesses. Indeed every degree-at-most- D solution is $zX^D + c$: writing $a = P' - zDX^{D-1}$ gives $a(a + R_0(X)(X - z)) = 0$, and degree comparison excludes the second factor. We make no assertion that only $c = 0$ can be nearby for an arbitrary choice of the u_i .

To exclude ordinary common agreement on any A coordinates, choose the u_i independently and uniformly in $\mathbb{F}_q^*$. Subtract X^D from a prospective direction polynomial. The zero polynomial matches exactly $s < A$ core coordinates. For a nonzero polynomial H of degree at most D , fix an A -subset containing m core coordinates. If $m > D$, matching is impossible. Otherwise at most q^{D+1-m} polynomials vanish on the selected core, and each matches the $A - m$ independent outside values with probability at most $(q - 1)^{-(A-m)}$. The probability of any match on this subset is at most $q^{D-s}(1 - 1/q)^{-(n-s)}$. Union bounding over at most 2^n subsets and using (80) proves existence. Without an A -agreeing direction polynomial, ordinary common agreement is impossible regardless of f . $\square$

For example, $n = 100m$, $D = 25m - 1$, $s = 48m - 1$ gives quarter rate, $A/n = 0.48$ above its first-order curve, and $52m + 1$ displayed bad labels. Condition (80) holds for sufficiently large prime fields. Here the fixed-cover theorem has positive linear slack, so its linear order is sharp. This is a linear lower example, not a quadratic lower bound for general first-order proximity gaps.

R Hermite–Johnson control of quadratic and cubic derivative identities

A persistent singular agreement of a quadratic derivative equation often prescribes both the value and the first derivative. Using these two conditions in interpolation removes the small-singular-cover restriction: the relevant threshold is the order-two Hermite Johnson threshold. The conclusion below concerns actual solutions of one common identity, and controls their original full agreement supports. It does not assert that every first-order interpolant has the required quadratic form.

The Hermite Johnson threshold and interpolation with derivative data are classical; see Kopparty [18, Section 1.2] and the univariate results cited there. Algebraic solution covers and full-support incidence arguments also appear in Jeronimo [19]. We give the explicit specialization and degree bounds needed for the linear exceptional count below; we do not claim a new decoding radius.

Lemma R.1 (Polynomial roots under specialization). *Let $Q \in \mathbb{F}[X, z, Y]$ be nonzero, with $\deg_X Q \leq T$, $\deg_z Q \leq H$, and $1 \leq \deg_Y Q \leq B$. Assume characteristic zero or characteristic $p > B$, and fix $D \geq 0$. Put*

$$C = B + H, \quad N = T + BD, \quad I = \max\{0, 2N - 1\}C^2, \quad J = C(1 + \max\{0, 2D - 1\}C).$$

Over $\overline{\mathbb{F}}$ there are at most $(2B + 1)H$ exceptional challenge values such that, outside them, all pairs (z, P) satisfying

$$\deg P \leq D, \quad Q(X, z, P(X)) \equiv 0$$

are contained in at most I exceptional solution pairs and at most C irreducible curves of actual polynomial solutions. In the affine space of z and the coefficients of P , these curves have total degree at most J . Each retained curve is parametrized by an initial-value plane curve through a map regular outside the exceptional challenge values.

Proof. Work over the algebraically closed constant field $K = \overline{\mathbb{F}}$. Factor Q in $K[X, z, Y]$, remove its Y -independent content, and retain each positive- Y -degree irreducible factor once. Call the resulting product R . Its separate degrees do not increase. Since $p > B$ when the characteristic is positive, each retained factor is separable in Y , so R is squarefree over $K(X, z)$. The removed content can specialize to zero identically in X at at most H challenge values: exclude the roots of any one of its nonzero X -coefficients. Outside these values, substitution into Q vanishes if and only if substitution into R vanishes. If there are no positive- Y -degree factors, this already proves the claim.

Choose a constant $\alpha \in K$ for which the Y -leading coefficient of $R(\alpha, z, Y)$ and $\text{Res}_Y(R, R_Y)(\alpha, z)$ are nonzero polynomials in z . Such a constant exists because only finitely many constants annihilate all z -coefficients of either nonzero polynomial in X, z . Their z -degrees are at most H and $(2B - 1)H$, respectively. Exclude their roots as well. The combined exceptional set has size at most $(2B + 1)H$.

Write

$$F(z, u) = R(\alpha, z, u), \quad A(z, u) = R_Y(\alpha, z, u).$$

At every remaining point of $F = 0$, one has $A \neq 0$. The nonvertical irreducible components of this plane curve have total degree at most C , and there are at most C of them. Vertical components lie over the excluded leading-coefficient values. On each nonvertical component the formal implicit solution has the form

$$P(\alpha + t) = u + \sum_{j \geq 1} a_j t^j, \quad a_j = \frac{N_j(z, u)}{A(z, u)^{2j-1}}, \quad \deg N_j \leq (2j - 1)C. \quad (*)$$

Here and below degrees in (z, u) are total degrees. For completeness, expand using Hasse coefficients, $R(\alpha + t, z, u + v) = \sum c_{ab}(z, u)t^a v^b$. All c_{ab} have degree at most C ; on $F = 0$ the constant term vanishes and the linear term in v is Av . After solving the coefficient of t^j for a_j , any remaining term with t -degree a and v -degree b has denominator exponent $e = 2(j - a) - b + 1 \leq 2j - 1$ and numerator degree at most eC , by induction. Bringing it to denominator A^{2j-1} proves $(*)$. The term with $b = 0$ has $a = j$ and denominator exponent one. This recursion divides by A only; it requires no factorial inverses.

Retain a component if its generic formal solution is a polynomial of degree at most D . On every other component, some coefficient a_j with $D < j \leq N$ is nonzero. Indeed, otherwise the truncation P_D through degree D agrees with the formal solution through degree N . The polynomial $R(\alpha + t, z, P_D(\alpha + t))$ then vanishes through order N and has degree at most $T + BD = N$, so it vanishes identically. Uniqueness of the formal implicit root would make the original series this degree- D polynomial, a contradiction. Choose one such j on each discarded component. Every specialized polynomial root of degree at most D on it must satisfy $N_j = 0$. This numerator is not identically zero on that component. Bézout's theorem, summed over components, bounds these initial points by $\max\{0, 2N - 1\}C^2 = I$. Each regular initial point determines at most one polynomial solution, so this counts solution pairs, not merely auxiliary points.

For $D \geq 1$, put $\delta = A^{2D-1}$. On a retained component the map to $(z, u, a_1, \dots, a_D)$ has homogeneous coordinates represented by

$$\delta, \quad z\delta, \quad u\delta, \quad N_j A^{2D-2j} \quad (1 \leq j \leq D).$$

Their degrees are at most $1 + (2D - 1)C$. A rational map given by forms of degree at most L sends a plane curve of degree d to a curve of degree at most dL , by intersection with a generic hyperplane. Summing proves the bound J ; base points can only reduce it. For $D = 0$ the map is the identity on (z, u) and the bound is C . Changing from Taylor coefficients at α to coefficients in X is an invertible linear transformation. The maps are regular wherever $A \neq 0$, and every such point of a retained component is an actual polynomial solution: substitution of the truncated polynomial is an identity in its function field, hence at all points where its coefficients are regular. Their closures also consist of actual solutions, since the substitution equations are polynomial. Finally the maps remember both z and u , and therefore are generically injective. This proves all assertions. $\square$

Fix n distinct coordinates over F , a received line $f + zg$, and integers $1 \leq D < A \leq n$. A pair (z, P) , $\deg P \leq D$, is *bad* if its full value-agreement support

$$S(z, P) = \{x : P(x) = f_x + zg_x\}$$

has size at least A and no degree-at-most- D polynomial agrees with g on all of $S(z, P)$. Consider a nonzero identity

$$Q(X, z, u, v) = a_2(X, z, u)v^2 + a_1(X, z, u)v + a_0(X, z, u), \quad (81)$$

of total degree at most $B_0 \geq 1$ in (u, v) and degree at most H_0 in z . Its coefficient degrees in X are unrestricted. Write $d = B_0 + H_0$ and, at each domain coordinate, define

$$w_x(z) = f_x + zg_x, \quad \alpha_{i,x}(z) = a_i(x, z, w_x(z)), \quad \Delta_x = \alpha_{1,x}^2 - 4\alpha_{2,x}\alpha_{0,x}.$$

Each $\alpha_{i,x}$ has degree at most d . The following are fixed sets of coordinates, defined by polynomial identities in z :

$$S_{\text{ord}} = \{x : \alpha_{0,x} = \alpha_{1,x} = \alpha_{2,x} = 0\}, \\ S_{\text{jet}} = \{x : \alpha_{2,x} \neq 0, \Delta_x = 0\}.$$

Put $u = |S_{\text{ord}}|$ and $s = |S_{\text{jet}}|$. The ordinary core is an explicit limitation: on it the specialized equation vanishes identically in both the derivative variable and the challenge, and supplies no additional derivative information.

Theorem R.2 (Quadratic derivative identities above the Hermite Johnson threshold). *Fix positive constants $\rho_{\min}, \eta, \varepsilon$ and bounded B_0, H_0 . Suppose*

$$\frac{D}{n} \geq \rho_{\min}, \quad A - D \geq \eta n, \quad A - u \geq \sqrt{\frac{Ds}{2}} + \varepsilon n. \quad (82)$$

Then the number of bad labels admitting an actual solution $Q(X, z, P, P') = 0$ is

$$O_{\rho_{\min}, \eta, \varepsilon, B_0, H_0}(n).$$

The characteristic may be zero, or $p > \max\{D, 2, B\}$, where B is the bounded interpolation Y -degree constructed below. In particular, there is a choice of this characteristic threshold depending only on the displayed fixed parameters, apart from the requirement $p > D$.

If S_{ord} is empty, it suffices to have $A \geq \sqrt{Dn/2} + \varepsilon n$. More generally, $A - u \geq \sqrt{D(n - u)/2} + \varepsilon n$ is sufficient. No bound on the size of the persistent Hermite core is required.

We give finite inequalities underlying the asymptotic statement. For integers $L \geq 1$ and $e \geq 0$, set

$$b = A - u - (L - 1) - e.$$

If $b > s$, no singular-heavy witnesses remain after the two exceptional sets counted below. Otherwise assume $1 \leq b \leq s$. Choose an integer $M \geq 1$ and put

$$T = Mb - 1, \quad B = \left\lfloor \frac{T}{D} \right\rfloor, \quad V = \sum_{j=0}^B (T - Dj + 1), \quad c_M = \left\lfloor \frac{(M+1)^2}{4} \right\rfloor. \quad (83)$$

Whenever $V > sc_M$, the nonnegative integer

$$H = \left\lfloor \frac{sc_M B(d+1)}{V - sc_M} \right\rfloor \quad (84)$$

is sufficient. These parameters give an explicit finite bound as follows. In Theorem H.5, use B_0, H_0 and write

$$C_{\text{reg}} = qK + qT_0 \frac{n - D}{A - D} + qn$$

for its parenthesized bound, with q, K, T_0 as defined there. For the interpolation parameters set

$$\begin{aligned} C &= B + H, \\ E_{\text{alg}} &= (2B + 1)H + \max\{0, 2(T + BD) - 1\}C^2, \\ J_{\text{alg}} &= C(1 + \max\{0, 2D - 1\}C). \end{aligned}$$

Then the number of bad labels is at most

$$\frac{n}{L}C_{\text{reg}} + \frac{2dn}{e+1} + E_{\text{alg}} + J_{\text{alg}} \frac{n - D}{A - D} + Cn. \quad (85)$$

Here the algebraic terms use Lemma R.1; the last three terms may be omitted when $b > s$.

Proof. Separating regular and exceptional singular agreements. Theorem H.5 removes at most $(n/L)C_{\text{reg}}$ labels possessing a bad witness with at least L agreements where $Q_v \neq 0$. At all remaining labels every bad witness has at most $L - 1$ such agreements.

On S_{jet} , outside the roots of $\alpha_{2,x}$, a singular agreement forces

$$P'(x) = -\frac{\alpha_{1,x}(z)}{2\alpha_{2,x}(z)}. \quad (86)$$

There are at most d excluded labels per such coordinate. Outside $S_{\text{ord}} \cup S_{\text{jet}}$, if $\alpha_{2,x} \neq 0$, singular agreement forces the nonzero polynomial Δ_x to vanish, at at most $2d$ labels. If $\alpha_{2,x} = 0$ but $\alpha_{1,x} \neq 0$, it forces $\alpha_{1,x} = 0$, at at most d labels. If both vanish identically, the nonzero $\alpha_{0,x}$ must vanish at an actual agreement, again at at most d labels. Thus the total number of exceptional singular coordinate–label incidences, including the poles in (86), is at most $2dn$. At most $2dn/(e+1)$ labels have more than e such coordinates. Every remaining bad witness has at least b singular matches in S_{jet} with the derivative prescription (86) defined.

Hermite interpolation with bounded challenge degree. Seek $R(X, z, Y) \neq 0$ with $\deg_z R \leq H$, $\deg_Y R \leq B$, and weighted (X, Y) -degree at most T for weights $(1, D)$. Its coefficient space has dimension $(H+1)V$. At each $x \in S_{\text{jet}}$, expand

$$R\left(x+t, z, w_x(z) - \frac{\alpha_{1,x}(z)}{2\alpha_{2,x}(z)}t + U\right)$$

and impose vanishing of the coefficients of $t^a U^j$ for $a+2j < M$, after multiplying the expansion by $(2\alpha_{2,x}(z))^B$. There are

$$\sum_{j=0}^{\lfloor (M-1)/2 \rfloor} (M-2j) = c_M$$

such conditions per coordinate. Each cleared coefficient has challenge degree at most $H + B(d+1)$: a source Y -power j contributes a numerator of degree at most $j(d+1)$ and a clearing factor of degree at most $(B-j)d$. Consequently there are at most $sc_M(H + B(d+1) + 1)$ scalar linear equations. Equations (83)–(84) give

$$(H+1)V > sc_M(H + B(d+1) + 1),$$

so a nonzero interpolant exists over F .

At a good singular match, the difference between $P(x+t)$ and its prescribed affine first jet is divisible by t^2 . The local conditions therefore imply a zero of multiplicity at least M in $R(X, z, P(X))$. This argument is valid in every characteristic and uses no factorial denominators. Its X -degree is at most $T < Mb$, so b distinct such matches force

$$R(X, z, P(X)) = 0. \quad (87)$$

The interpolant has positive Y -degree: a nonzero Y -independent polynomial satisfying the local conditions would have at least M zeros counted with multiplicity over $F(z)$, contrary to $T < Mb \leq Ms$.

Specialization and original full supports. Lemma R.1 applies to (87). It accounts for at most E_{alg} exceptional labels or isolated solution pairs, including polynomial roots that exist only at special challenges. All other solutions lie on at most C coefficient curves, including the challenge coordinate, of total degree at most J_{alg} .

Each original coordinate agreement is the hyperplane $P(x) = f_x + zg_x$ in this coefficient space. If more than D distinct such hyperplanes contain a curve, interpolation forces that curve to lie on an affine codeword pencil $P = F_0 + zG_0$. A nonaffine curve therefore has at most D persistent agreement hyperplanes. Bézout and incidence counting bound its labels with at least A agreements by its degree times $(n - D)/(A - D)$.

On an affine pencil, let I be the persistent value-agreement set. A bad witness must agree at some coordinate outside I : otherwise G_0 agrees with g on its full support. Since this support has more than D distinct F -coordinates, interpolation then forces $G_0 \in F[X]$, even if the pencil was initially defined over $\overline{F}$. Each coordinate outside I contributes at most one label, so each pencil has at most n bad labels. There are at most C pencils. This proves (85). In particular this last argument uses *all original value agreements*, not merely the Hermite matches used to construct R .

Fixed-margin parameters. Under (82), take $L = \max\{1, \lfloor \varepsilon n/4 \rfloor\}$ and $e = \lfloor \varepsilon n/4 \rfloor$. Then $b \geq \sqrt{Ds/2} + \varepsilon n/2$, while both L and $e + 1$ are bounded below by positive multiples of n . If $b > s$ we are done. Otherwise, as M increases,

$$V = \frac{M^2 b^2}{2D} + O(Mn), \quad sc_M = \frac{sM^2}{4} + O(Ms).$$

The fixed margins and $D/n \geq \rho_{\min}$ permit constant M , and then constant B, H , with $V > sc_M$. These choices also cover the integer roundings for small n . Now $T = O(n)$, $E_{\text{alg}}, J_{\text{alg}} = O(n)$, and every term in (85) is $O(n)$ with the stated dependencies. $\square$

The ordinary-core qualification cannot be suppressed by squaring the margin: the signed quantity $A - u$ must be positive. When $u = 0$, the argument controls even a persistent Hermite core occupying the entire domain, and even rational challenge dependence in its derivative data. It identifies an explicit remaining mechanism: many coordinates where the whole derivative equation vanishes on the received line, rather than merely a large set of persistent quadratic singularities. In particular, an interpolation source constructed to vanish identically on every received fiber has $u = n$. The theorem gives no bound for such a raw source without an additional reduction to a different identity.

Sharpness of the linear order. Proposition Q.4 supplies matching $\Omega(n)$ examples within this class. Its equation has leading v^2 coefficient one, so $S_{\text{ord}} = \emptyset$, and bounded value and challenge degrees. With $n = 100m$, $D = 25m - 1$, and $A = 48m$, the margin $A > \sqrt{Dn/2}$ is fixed, while the construction has $52m + 1$ bad labels, all of whose displayed agreements are singular. Thus the order $O(n)$ cannot be improved to $o(n)$ uniformly even in the quadratic, empty-ordinary-core case.

R.1 Cubic derivative equations also prescribe a unique singular jet

A nonzero cubic has at most one distinct repeated root. When such a root persists with the challenge, it is rational in the coefficients. This extends the preceding argument without selecting among several possible derivative branches.

Corollary R.3 (Cubic derivative identities). *Retain the domain, received line, degree bound, and definition of bad pairs from Theorem R.2. Let $Q(X, z, u, v) \neq 0$ have derivative degree at most three, total (u, v) degree at most B_0 , and challenge degree at most H_0 , with unrestricted X -degree. Put $h = B_0 + H_0$ and*

$$F_x(z, v) = Q(x, z, f_x + zg_x, v) \in F[z, v].$$

Define S_{ord} by $F_x = 0$ identically. Define S_{jet} to consist of those other coordinates where the nonzero polynomial $F_x \in F(z)[v]$ has a repeated root over an algebraic closure of $F(z)$. Write $u = |S_{\text{ord}}|$ and $s = |S_{\text{jet}}|$.

Suppose $D/n \geq \rho_{\min} > 0$ and

$$A - D \geq \eta n, \quad A - u \geq \sqrt{Ds/2} + \varepsilon n$$

for fixed positive η, ε . In characteristic zero, or characteristic $p > \max\{D, 3, B\}$ for the bounded interpolation degree B below, there are $O_{\rho_{\min}, \eta, \varepsilon, B_0, H_0}(n)$ bad labels admitting an actual solution $Q(X, z, P, P') = 0$.

More precisely, retain L, e, b, M, T, B, V, c_M from (83), with $V > sc_M$, and replace (84) by

$$H = \left\lfloor \frac{sc_M B(2h+1)}{V - sc_M} \right\rfloor. \quad (88)$$

Then the finite bound (85) holds with its singular-incidence term $2dn/(e+1)$ replaced by $4hn/(e+1)$, and with its algebraic terms evaluated at these new B, H, T . The regular-incidence term still uses the original B_0, H_0 .

Proof. Every coefficient of F_x has challenge degree at most h . We show that outside at most $4h$ exceptional labels at each nonordinary coordinate, every singular match belongs to S_{jet} and has one prescribed rational derivative, of numerator and denominator degrees at most $2h$. All degree classifications below are generic over $F(z)$; the exceptional sets explicitly include degree drops.

First suppose

$$F_x = av^3 + bv^2 + cv + d, \quad a \neq 0.$$

Set

$$\begin{aligned} \Delta &= b^2c^2 - 4ac^3 - 4b^3d - 27a^2d^2 + 18abcd, \\ \Delta_0 &= b^2 - 3ac, \quad V_0 = 9ad - bc. \end{aligned}$$

Their challenge degrees are at most $4h, 2h, 2h$, respectively. If $\Delta \neq 0$ as a polynomial in z , every singular specialization lies among its at most $4h$ roots. This includes leading-coefficient zeros: at $a(z) = 0$ the expression becomes $b(z)^2(c(z)^2 - 4b(z)d(z))$, which vanishes at a repeated quadratic root and at any further degree drop admitting a singular root.

If $\Delta = 0$ identically and $\Delta_0 \neq 0$, exclude the roots of $a\Delta_0$, at most $3h$ labels. At every remaining label the unique repeated root is

$$r_x(z) = \frac{V_0(z)}{2\Delta_0(z)}. \quad (89)$$

Indeed, writing the cubic over an algebraic closure as $a(v-r)^2(v-t)$ with $r \neq t$ gives $\Delta_0 = a^2(r-t)^2$ and $V_0 = 2a^2r(r-t)^2$. If both Δ and Δ_0 vanish identically, exclude the roots of a . The remaining polynomial is $a(v + b/(3a))^3$, and its unique repeated root is $-b/(3a)$. This follows from $c = b^2/(3a)$ and $d = b^3/(27a^2)$, using characteristic different from two and three. These are exactly the persistent repeated-root cases of generic cubic degree; all their prescriptions have rational height at most $2h$.

If the generic degree is two, write $F_x = bv^2 + cv + d$, with $b \neq 0$. A nonzero discriminant $c^2 - 4bd$ confines every singular specialization, including degree drops, to at most $2h$ labels. An identically zero discriminant prescribes the unique repeated root $-c/(2b)$ outside the at most h roots of b . For generic degree one, singular agreement requires the nonzero linear coefficient to vanish. For a nonzero constant polynomial in v , an actual agreement requires that constant coefficient to vanish. Each case contributes at most h exceptional labels and no persistent jet core. The identically zero case is precisely S_{ord} .

Thus the exceptional coordinate-label incidence budget is $4hn$, uniformly over all candidates. After the regular-incidence deletion and deletion of at most $4hn/(e+1)$ further labels, every

remaining bad witness has at least $A - u - (L - 1) - e$ matches in S_{jet} with $P'(x) = r_x(z)$ defined. Write $r_x = N_x/E_x$ with both degrees at most $2h$. In the local Hermite expansion multiply by E_x^B . Each cleared constraint now has challenge degree at most $H + B(2h + 1)$, so (88) gives

$$(H + 1)V > sc_M(H + B(2h + 1) + 1).$$

The multiplicity argument, specialization-safe algebraic-root lemma, and original full-support incidence count are unchanged. The same fixed-margin choices of L, e, M make B, H bounded and give the claimed linear bound. $\square$

The assertion is not that every cubic has a repeated derivative root: generically squarefree specializations contribute only the bounded exceptional incidences. It is the *uniqueness and rationality* of a persistent repeated root that permits this extension. Degree four is the first degree allowing two distinct repeated derivative roots, as in $(v - r_1)^2(v - r_2)^2$. The corollary leaves the ordinary core, the agreement margin, and containment in one common bounded-degree identity as explicit hypotheses; it is not an unconditional obstruction to all proximity-gap counterexamples.

A primitive equation with a full ordinary core. The ordinary core can contain a genuine growing list of actual polynomial solutions, even when the differential equation is primitive and irreducible. The following identity places the existing full-length Dickson list inside a cubic first-order equation; it does not provide a new list lower bound.

Proposition R.4. *Let $p \equiv 1 \pmod{8}$ be prime, $k = (p - 1)/4$, and $D = k - 1$. There is a primitive irreducible polynomial $Q \in \mathbb{F}_p[X, u, v]$ of total (u, v) -degree three and v -degree one, and a word w on $\mathbb{F}_p^\times$, such that:*

1. $Q(X, P, P') = 0$ for $2k$ distinct degree- D polynomials P , each agreeing with w at exactly $3k/2$ coordinates; together with $P = 0$, these are all degree-at-most- D polynomial solutions, even over every extension of $\mathbb{F}_p$;
2. $Q(x, w(x), v)$ is the zero polynomial in v for every $x \in \mathbb{F}_p^\times$.

The equation has no challenge dependence. Thus its ordinary core is the entire domain for the constant received line $f = w$, $g = 0$.

Proof. Put $e = (p + 1)/2$, $q = X^{2k}$, and $S = 1 + q$. Use the polynomials $G_a, P_a = G_a - X^k$ from Proposition 2.5, with $a \in \mathbb{F}_p^\times$ modulo sign. That proposition gives their number, degrees, and agreement counts with $w(x) = (1 + x^{2k})/2 - x^k$. Define

$$\mathcal{Q}(X, G, V) = 4X(2G^2 - 1 - X^{2k})V + G(4G^2 - 1 - 3X^{2k}), \quad Q(X, u, v) = \mathcal{Q}(X, u + X^k, v + kX^{k-1}).$$

We prove $\mathcal{Q}(X, G_a, G'_a) = 0$. With $t^2 = X$, write

$$J = \frac{(a + t)^e + (a - t)^e}{2}, \quad G = \frac{(a + t)^e - (a - t)^e}{2t} = G_a(X).$$

Frobenius, using $a^p = a$ and $2e = p + 1$, gives

$$J^2 + XG^2 = a^2 + X^{2k+1}, \quad 2JG = aS.$$

Differentiation, with $e = 1/2$ in $\mathbb{F}_p$, gives $2(a^2 - X)(2XG' + G) = aJ - XG$. Consequently

$$a^2[4G(2XG' + G) - S] = 2XG(4XG' + G), \quad a^2(S^2 - 4G^2) = 4XG^2(q - G^2).$$

Eliminating a^2 and cancelling the nonzero polynomial $2XG$ yields

$$(4XG' + G)(S^2 - 4G^2) - 2G(q - G^2)[4G(2XG' + G) - S] = 0.$$

The left side factors as $(2G^2 - S)\mathcal{Q}(X, G, G')$. Since G is monic of degree k , the degree- $2k$ coefficient of $2G^2 - S$ is $2 - 1 = 1$. Cancellation proves the claimed identity.

The same calculation includes $a = 0$, which gives $P = 0$. To prove completeness over any extension field, write a prospective monic solution as $G = X^k + \sum_{r=1}^k c_r X^{k-r}$. For $1 \leq r \leq k$, the coefficient of X^{3k-r} in $\mathcal{Q}(X, G, G')$ is

$$4(1-r)c_r + F_r(c_1, \dots, c_{r-1}).$$

Indeed its linear multiplier is $20k - 4r + 9 = 4(1-r)$ in $\mathbb{F}_p$; any nonlinear occurrence of c_r has larger coefficient deficit. The terms of degree at most k do not contribute. Only $r = 1$ is resonant, with $F_1 = 0$, so c_1 uniquely determines all the other coefficients. Set $\tau = -8c_1$ and define

$$\tilde{G}_\tau(X) = \sum_{j=0}^k \binom{2k+1}{2j+1} \tau^{k-j} X^j.$$

Its next-to-leading coefficient is $-\tau/8$. For each $r \leq k$, its coefficient of X^{3k-r} in the residual is a polynomial of degree at most r in τ . It vanishes at the $2k+1$ values consisting of zero and all nonzero squares in $\mathbb{F}_p$, by the identities already proved for G_a . It therefore vanishes identically, and the recurrence forces $G = \tilde{G}_\tau$. The constant term of the differential equation now gives

$$0 = G(0)(4G(0)^2 - 1) = \frac{1}{2}\tau^k(\tau^{2k} - 1).$$

All roots of $\tau^{2k} - 1$ are already the nonzero squares in $\mathbb{F}_p$. Thus no extension contributes further solutions, proving the exact count $2k+1$ after shifting back to $P = G - X^k$.

For $W(X) = S/2 - X^k$, direct substitution gives the stronger fiber identity

$$\mathcal{Q}(X, W(X), v) = (X^{p-1} - 1)[2X(v + kX^{k-1}) + X^{2k}/2].$$

This proves the ordinary-core assertion. Finally, the coefficients of V and of 1 in $\mathcal{Q}$ are coprime in $\mathbb{F}_p[X, G]$. Indeed X does not divide the latter, G does not divide $2G^2 - S$, and

$$(4G^2 - 1 - 3q) - 2(2G^2 - S) = 1 - q$$

shares no factor with $2G^2 - S$, whose leading G -coefficient is the nonzero constant 2. Gauss's lemma and linearity in V give irreducibility. The displayed affine change of variables preserves it. $\square$

Here $n = p - 1$, the rate is $1/4$, and the agreement fraction is $3/8$. The construction depends on the Frobenius identity at this characteristic; it does not supply a short-domain family by enlarging p while keeping k fixed. Nor does it show that an ordinary-core hypothesis is necessary for a linear bad-label bound: the displayed received line has $g = 0$, so $(P, 0)$ explains the full support of every witness and there are no bad labels. Its agreement fraction is below the first-order threshold at rate $1/4$, so it is not a tightness example for that threshold.

The sharp agreement transition inside the solution bank. The complete solution family in Proposition R.4 has a precise list-size transition. This is a statement about that family, not about all Reed–Solomon codewords.

Theorem R.5. *Let $p \equiv 1 \pmod{8}$ be prime, $n = p - 1$, $k = n/4$, and let $\mathcal{B} = \{0\} \cup \{P_a : a \in \mathbb{F}_p^\times / \{\pm 1\}\}$ be the complete degree-at-most- $k - 1$ solution bank of Proposition R.4. For every received word on $\mathbb{F}_p^\times$, including words over an extension field, and every $0 < \varepsilon \leq 5/8$, at most $\lceil \varepsilon^{-2} \rceil$ members of $\mathcal{B}$ have agreement at least $(3/8 + \varepsilon)n$.*

Conversely, for every $L \geq 1$ and every prime $p \equiv 1 \pmod{8}$ with

$$p \geq 2^{36} L^2,$$

there is a received word with L distinct members of $\mathcal{B}$ each having agreement at least $(3/8 + 1/(32\sqrt{L}))n$. The entire domain remains ordinary core for the same differential equation. Thus, allowing sufficiently large primes depending on ε , the worst-case list size within this bank is $\Theta(\varepsilon^{-2})$ above agreement $3/8$. At agreement exactly $3/8$, the word in Proposition R.4 has $n/2$ bank members.

Proof. Write χ for the quadratic character with $\chi(0) = 0$. Fix L nonzero bank members with representatives $a_1, \dots, a_L$ of distinct squares. Shift each received symbol by x^k , so that agreements concern $G_a = P_a + X^k$.

At a nonsquare x , the identities in the preceding proposition give $J_a G_a = 0$ and $J_a^2 + x G_a^2 = a^2 - x$. Thus $G_a(x) = 0$ exactly when $\chi(a^2 - x) = 1$; otherwise $G_a(x)^2 = a^2/x - 1$. Every nonzero value therefore occurs for at most one selected parameter. If $N_0(x)$ is the zero-bucket size, then

$$\sum_{x \text{ nonsquare}} N_0(x) = Ln/4. \quad (90)$$

Indeed, for each $a \neq 0$,

$$\sum_{x \text{ nonsquare}} \chi(a^2 - x) = \frac{1}{2} \left(\sum_{x \neq 0} \chi(a^2 - x) - \sum_x \chi(x(a^2 - x)) \right) = 0,$$

because both sums in parentheses equal -1 .

At a square $x = s^2 \neq 0$, Euler's criterion gives

$$G_a(s^2) = \frac{(a + s)\chi(a + s) - (a - s)\chi(a - s)}{2s}.$$

Away from $s = \pm a$, equal character signs give the shared value $+1$ or -1 ; opposite signs give $\pm a/s$. Values outside $\{+1, -1\}$ again occur at most once. At $s = \pm a$ the value is $\chi(2a) = \chi(a)$. Let $N_+(x), N_-(x)$ count the actual shared buckets. Then

$$\sum_{x \text{ square}} N_+(x) = \sum_{x \text{ square}} N_-(x) = Ln/8. \quad (91)$$

To include endpoints correctly, consider

$$I_{a,\tau}(s) = \frac{(1 + \tau\chi(a + s))(1 + \tau\chi(a - s))}{4}, \quad \tau \in \{+1, -1\}.$$

Its sum over all s is $n/4$. Removing $s = 0$ subtracts one precisely when $\tau = \chi(a)$; at each of $s = \pm a$, replacing $I_{a,\tau}$ by the true bucket indicator adds $1/2$ in precisely that case. These corrections cancel. Dividing by two proves (91).

Put

$$S(s) = \sum_{i=1}^L (\chi(s + a_i) + \chi(s - a_i)).$$

This is even in s , and $N_+(s^2) - N_-(s^2) = S(s)/2 + E(s^2)/2$, where $E(a_i^2) = \chi(a_i)$ and E vanishes elsewhere. The $2L$ shifts are distinct. The exact correlations

$$\sum_s \chi(s + b)^2 = p - 1, \quad \sum_s \chi(s + b)\chi(s + c) = -1 \quad (b \neq c)$$

give

$$\sum_s S(s)^2 = 2Lp - 4L^2.$$

The second identity follows, for example, by counting $y^2 = (s + b)(s + c)$ through a change of variables to two factors with fixed nonzero product. Using $\sum_x |E(x)| = L$, Cauchy–Schwarz, and (91), we obtain

$$\sum_{x \text{ square}} \max\{N_+(x), N_-(x)\} \leq \frac{Ln}{8} + \frac{1}{8}\sqrt{n(2Lp - 4L^2)} + \frac{L}{4}.$$

Allowing singleton buckets costs at most one at each coordinate. Together with (90), this bounds the total agreements of the selected candidates with any word by

$$\frac{3Ln}{8} + n + \frac{L}{4} + \frac{1}{8}\sqrt{n(2Lp - 4L^2)}. \quad (92)$$

Symbols outside $\mathbb{F}_p$ contribute no agreements, so the bound also holds for extension-field received words.

If all L candidates have agreement at least $(3/8 + \varepsilon)n$, then, since $L \leq n/2$ and $p \leq 2n$,

$$\varepsilon \leq \frac{1}{L} + \frac{1}{4n} + \frac{1}{8}\sqrt{\frac{2Lp - 4L^2}{L^2n}} \leq \frac{9}{8L} + \frac{1}{4\sqrt{L}}.$$

If $L \geq \varepsilon^{-2}$, the right side is at most $\varepsilon(9\varepsilon/8 + 1/4) \leq 61\varepsilon/64$, a contradiction. Adding the zero solution proves the upper bound.

For the lower bound, put $M = 2L$ and choose $a_1, \dots, a_M$ independently and uniformly in $\mathbb{F}_p$. We shall retain L good parameters while keeping the majority word built from all M . For independent uniform signs ξ_i, η_i , let σ be the sign of their sum, resolving a tie as $+1$. The Boolean indicator

$$F_i = \frac{(1 + \sigma\xi_i)(1 + \sigma\eta_i)}{4}$$

has mean

$$\mu_M = \frac{1}{4} + \frac{\mathbb{E}|\sum_i(\xi_i + \eta_i)|}{4M} = \frac{1}{4} + \frac{1}{2} \frac{\binom{2M}{M}}{4^M}. \quad (93)$$

The first equality follows from permutation symmetry and $\sigma \sum_i(\xi_i + \eta_i) = |\sum_i(\xi_i + \eta_i)|$; the second follows by telescoping binomial probabilities. Writing $b_M = \binom{2M}{M}/4^M$, one has $b_M \geq 1/(2\sqrt{M})$: this holds at $M = 1$, and $(b_{M+1}\sqrt{M+1}/(b_M\sqrt{M}))^2 = (2M+1)^2/(4M(M+1)) > 1$.

For $s \neq 0$, evaluate the multilinear extension of F_i at the $2M$ signs $\chi(a_j + s), \chi(a_j - s)$ and call the result $Z_i(s)$. A zero sign is interpreted as an independent uniform sign followed by expectation,

so $0 \leq Z_i(s) \leq 1$. Put $Y_i = n^{-1} \sum_{s \neq 0} Z_i(s)$. For uniform a , the completed pair of signs at $s \neq 0$ has law

$$\pi(\epsilon_1, \epsilon_2) = \frac{1}{4}(1 - \epsilon_1 \epsilon_2 / p).$$

Indeed both first moments vanish and their product moment is $-1/p$. Its total variation distance from the uniform law is $1/(2p)$. Independence of the parameters therefore gives

$$|\mathbb{E}Y_i - \mu_M| \leq M/(2p). \quad (94)$$

We next bound the variance independently of M . Fix nonzero s, t with $t \neq \pm s$. For one uniform parameter a , let Ω be the joint law of the completed sign pairs at s and t , using independent completions of zero signs. Its two marginals are π . The four shifts are distinct. Exact pair correlations and the product-character bound [2, Lemma 17] give, entrywise,

$$|\Omega - \pi \otimes \pi| \leq \frac{4/p + 16/\sqrt{p} + 1/p^2}{16} \leq \frac{21}{16\sqrt{p}}.$$

Here the four cross-pair terms contribute $4/p$, the four triple terms and one quadruple term contribute $(4 \cdot 3 + 4)/\sqrt{p}$, and the product of the internal pair moments contributes $1/p^2$. The first moments vanish. Since $\min \pi \geq 1/6$, the 4-by-4 matrix obtained by dividing each centered entry by the square root of its two marginal probabilities has Frobenius norm at most $32/\sqrt{p}$.

This bounds the norm of the conditional-expectation operator on mean-zero functions by $\rho = 32/\sqrt{p} < 1$. For M independent parameter blocks the same bound holds: decompose each $L^2(\pi)$ into constants and mean-zero functions. The uncentered conditional operator preserves both spaces, acting as the identity on constants and with norm at most ρ on the other space. Its tensor product has norm at most ρ on every sector containing a mean-zero factor. Consequently the covariance of the two Boolean majority indicators is at most $\rho/4 \leq 8/\sqrt{p}$ in absolute value. Independent completions at s, t , conditional on the parameters, show that this covariance equals that of $Z_i(s), Z_i(t)$. The pairs $t = \pm s$ cost at most $2/n$ in the averaged variance. Thus

$$\text{Var}(Y_i) \leq 8/\sqrt{p} + 2/n \leq 10/\sqrt{p}. \quad (95)$$

Under the assumed bound $p \geq 2^{36}L^2 = 2^{34}M^2$, both $M/(2p)$ and M/n are at most $1/(32\sqrt{M})$. Chebyshev's inequality and (95) give

$$\Pr \left\{ |Y_i - \mathbb{E}Y_i| > \frac{1}{32\sqrt{M}} \right\} \leq \frac{10240M}{\sqrt{p}} \leq \frac{5}{64}.$$

The expected number of failing indices is at most $5M/64$; Markov's inequality bounds the probability of at least $M/2$ failures by $5/32$. The probability that some $a_i = 0$ or some $a_i = \pm a_j$ is at most $M^2/p \leq 2^{-34}$. There is therefore a choice of parameters with distinct nonzero squares and at least L good indices, each satisfying $Y_i \geq \mu_M - 1/(16\sqrt{M})$ by (94).

For this choice define the shifted word $W(x)$ to be zero at nonsquares and, for $x = s^2$, the sign of $\sum_j (\chi(s + a_j) + \chi(s - a_j))$, again with tie +1. The sum is even in s , so W is well defined. Put $w(x) = W(x) - x^k$ and do not change this word when retaining only the L good candidates. Each candidate has exactly $n/4$ nonsquare agreements. Away from the at most $2M$ endpoints $s = \pm a_j$, its square agreement indicator is exactly $Z_i(s)$. Removing these endpoints and dividing by two gives total agreement fraction at least

$$\frac{1}{4} + \frac{Y_i}{2} - \frac{M}{n} \geq \frac{3}{8} + \frac{b_M}{4} - \frac{1}{32\sqrt{M}} - \frac{1}{32\sqrt{M}} \geq \frac{3}{8} + \frac{1}{16\sqrt{M}} \geq \frac{3}{8} + \frac{1}{32\sqrt{L}}.$$

Both $W = 0$ at nonsquares and $W = \pm 1$ at squares make the original differential equation vanish identically in its derivative variable; thus the ordinary core remains the entire domain. Finally, for $0 < \varepsilon \leq 1/64$ take $L = \lfloor 1/(1024\varepsilon^2) \rfloor \geq 1/(2048\varepsilon^2)$ and choose a sufficiently large prime $p \equiv 1 \pmod{8}$. This proves the matching order. $\square$

Quadratic rational first-integral pencils. A bounded list theorem also applies when the denominator of a first integral depends on the candidate value. The square-value input below is the classical function-field theorem of Pasten and Wang [11, Theorem 3]; the argument keeps track of the height needed to exclude its positive-characteristic exceptions.

Theorem R.6. *Let $\mathbb{K}$ be a field of characteristic $p > \max\{2, 10D\}$, where $0 \leq D < n$, and let $N, T \in \mathbb{K}(X)[u]$ be coprime with $\max\{\deg_u N, \deg_u T\} = 2$. Consider the polynomial sections*

$$\mathcal{B} = \{P \in \mathbb{K}[X] : \deg P \leq D, N(X, P) = cT(X, P) \text{ for some } c \in \mathbb{K}\}.$$

On any n distinct evaluation coordinates in $\mathbb{K}$, every received word has at most

$$\max\{40, \lfloor 1/\eta \rfloor\}$$

members of $\mathcal{B}$ agreeing on at least $D + \eta n$ coordinates, for every $\eta > 0$. The received symbols may lie in an extension field. This is a bound for the indicated section family, not for all Reed–Solomon codewords.

Proof. Extend constants to an algebraic closure k and put $K = k(X)$. A section has a unique label: otherwise $N(P) = T(P) = 0$, contrary to coprimality. Each label has at most two sections. Thus a selected list with fewer than twenty-one labels has size at most forty. Suppose henceforth that twenty-one labels occur.

Reconstructing a low-height pencil. Choose five sections P_i with distinct labels c_i . They are distinct elements of K . The matrix with rows

$$(P_i^2, P_i, 1, -c_i P_i^2, -c_i P_i, -c_i)$$

has rank five over K . Indeed, another kernel vector defines $(\tilde{N}, \tilde{T})$ of degrees at most two whose cross-difference $\tilde{N}T - \tilde{T}N$ has degree at most four and vanishes at all five P_i . It is zero; coprimality and the exact degree-two maximum make the pairs proportional. The signed maximal minors therefore give a polynomial coefficient representative $(n_2, n_1, n_0, t_2, t_1, t_0)$ for (N, T) with degrees bounded respectively by

$$(4D, 5D, 6D, 4D, 5D, 6D).$$

Replace (N, T) by this common polynomial representative. Consequently the coefficients of the discriminant

$$F(C) = (n_1 - Ct_1)^2 - 4(n_2 - Ct_2)(n_0 - Ct_0)$$

have X -degree at most $10D$. The polynomial F has degree at most two in C and is not a square in $K(C)$: the coprime rational map N/T has degree two, so $N - CT$ is irreducible over $K(C)$ as a polynomial in u .

Every section gives

$$F(c_i) = (2(n_2 - c_i t_2)P_i + n_1 - c_i t_1)^2.$$

This remains valid for a linear specialized fiber. Choose c_0 with $F(c_0) = s^2 \neq 0$, possible since F has at most two zeros, and form

$$H(Z) = \frac{Z^2 F(c_0 + 1/Z)}{F(c_0)}.$$

It is monic quadratic. Its projective coefficient vector is $[F(c_0) : F'(c_0) : [C^2]F]$, so its function-field height is at most $10D < p$. At the twenty distinct constants $Z_i = (c_i - c_0)^{-1}$ it takes square values.

Applying the square-value theorem. For a monic quadratic over the genus-zero field K , the Pasten–Wang theorem has threshold $M > 11 \cdot 2 - 3 = 19$. After removing constant polynomial factors, it requires every remaining irreducible factor to be separable and outside $K^p[Z]$; with twenty square values it makes every such factor have multiplicity two. These hypotheses hold here. Degree at most two and $p > 2$ give separability. Monic factor heights add, by multiplicativity of the local Gauss valuation, so every factor has height below p . A monic factor in $K^p[Z]$ with nonconstant coefficients would have a nonconstant coefficient which is a p th power, with height at least p , a contradiction. If H has nonconstant coefficients, the multiplicity conclusion and its total degree two force it to be a square; a positive-degree constant factor cannot remain in that case. But the reciprocal substitution and its square multiplier would then make F a square in $K(C)$. Therefore $H \in k[Z]$, and undoing the substitution yields

$$F(C) = s^2 q(C), \quad q \in k[C], \quad \deg q \leq 2.$$

The polynomial q is nonsquare, hence squarefree of degree one or two over the algebraically closed field k .

A constant conic and its sections. Write $N - CT = a(C)u^2 + b(C)u + d(C)$. In $K(u)$ put $c = N(u)/T(u)$ and $v = (2a(c)u + b(c))/s$. Then

$$v^2 = q(c), \quad K(u) = K(c, v).$$

The smooth projective constant conic $v^2 = q(c)$ is rational: parametrizing by lines through a constant point gives its function field $k(t)$. Hence $K(u) = K(t)$, so u is a degree-one rational function of t ,

$$u = \frac{A(X)t + B(X)}{C(X)t + E(X)}, \quad AE - BC \neq 0.$$

Every polynomial section corresponds to a constant parameter in $\mathbb{P}^1(k)$. Indeed its c is constant and its v has square $q(c) \in k$, hence is constant too. The isomorphism of smooth projective curves handles degree-drop fibers and parameters at infinity. Distinct sections give distinct parameters. Polynomial members of this Möbius family need not form an affine line.

Local collisions and agreement. At each evaluation coordinate x , scale the 2×2 coefficient matrix over the discrete valuation ring at $X = x$ to make its entries integral with at least one unit. Its reduction has rank one or two. Rank two gives distinct values at all selected parameters. In rank one, all parameters except the unique kernel point have the same projective image. That image is finite since at least one nonkernel selected section is a polynomial. Thus all but at most one selected value coincide.

Let L be the selected list size, let m_x be the number of selected candidates matching the received symbol at x , and let b_x count all agreeing candidate pairs there. The local description gives $m_x \in \{0, 1, L - 1, L\}$, and in every case

$$(m_x - 1)(L - 1) \leq 2b_x.$$

For $m_x = L - 1$ this follows from $b_x \geq \binom{L-1}{2}$; for $m_x = L$ use $b_x = \binom{L}{2}$; the other cases are immediate. Every distinct polynomial pair agrees at at most D coordinates. Summing, and writing a for the minimum agreement count, gives

$$(La - n)(L - 1) \leq 2 \sum_x b_x \leq DL(L - 1), \quad L(a - D) \leq n.$$

Thus $a - D \geq \eta n$ implies $L \leq 1/\eta$. Together with the forty-section case, this proves the stated bound. Extending the constants did not change any agreement count or polynomial degree, so the conclusion holds over the original field. $\square$

First integrals with isolated critical points. The multiplicity argument extends to every fixed value degree. The hypothesis below concerns the original critical ideal, including fibers above repeated roots of the denominator.

Theorem R.7. *Let $b \geq 2$, $1 \leq D < n$, and let $\mathbb{K}$ have characteristic zero or characteristic $p > b$. Suppose $0 \neq H \in \mathbb{K}[X]$ and*

$$F(X, u) = u^b + \sum_{j=0}^{b-1} a_j(X)u^j, \quad \deg a_j \leq (b-j)D.$$

Assume that the original ideal

$$I = (F_u, HF_X - H'F)$$

is zero-dimensional after extending constants to an algebraic closure; the unit ideal is allowed. On any n distinct evaluation coordinates, every received word has at most

$$\left\lfloor b^2 + \frac{b + 3b^2(b-1)^2 D/n}{\eta} \right\rfloor$$

polynomials P of degree at most D satisfying $F(X, P) = cH$ for some constant c and agreeing with the word on at least $D + \eta n$ coordinates, for every $\eta > 0$. Coordinates and received symbols may lie in an extension field.

Proof. Extend constants to an algebraically closed field k containing the coordinates and symbols. Each section has a unique label c , and each label has at most b sections. If no nonzero label occurs the claim is immediate. Otherwise $N = \deg H \leq bD$ by the weighted coefficient caps. Put $B = HF_X - H'F$ and $\ell = \dim_k k[X, u]/I$. The bidegrees of F_u, B are at most $((b-1)D, b-1)$ and $(N + bD - 1, b)$. Homogenize each nonzero generator in $\mathbb{P}^1 \times \mathbb{P}^1$ to its own actual bidegree. Neither contains a boundary divisor; a common projective curve would therefore meet the affine chart and contradict zero-dimensionality. The proper intersection product gives

$$\ell \leq (b-1)(N + 2bD - 1). \quad (96)$$

A nonzero constant generator gives length zero directly.

For a section P write $A_P = F_u(X, P)$. Differentiation gives $B(X, P) = -HA_PP'$. Consequently $A_P \neq 0$, since otherwise the graph of P would be contained in the critical scheme, and

$$k[X, u]/(I, u - P) \simeq k[X]/(A_P).$$

After localization, $\text{ord}_x A_P$ is therefore at most the critical scheme's local length at $(x, P(x))$.

We use one valuation estimate at finite points and infinity. In a fixed valued algebraic closure factor

$$F_u(X, P + Z) = b \prod_{i=1}^{b-1} (Z - \delta_i), \quad A_P = b \prod_i (-\delta_i),$$

and choose δ of maximal valuation. All δ_i are nonzero. If $d = v(A_P)$ then $v(\delta) \geq d/(b-1)$. Writing the derivative as $\sum_j c_j Z^j$, each summand of $c_j \delta^j$ is, up to sign, A_P times a product of ratios δ/δ_i . Its valuation is at least d . Integrating this polynomial in Z thus gives

$$v(F(X, P + \delta) - F(X, P)) \geq d + v(\delta) \geq \frac{bd}{b-1}. \quad (97)$$

Only $1, \dots, b$ are inverted, so the characteristic assumption suffices. The argument permits repeated critical roots and ramified valuations.

At infinity this chooses a root of minimal degree. If $\deg A_P < (b-1)N/b$, the correction has degree less than N and the fixed critical value $F(X, P + \delta)/H$ has residue c at infinity. There are at most $b-1$ fixed critical roots of F_u , hence at most $b-1$ such exceptional labels. For every other label,

$$\deg A_P \geq \lceil (b-1)N/b \rceil.$$

At a root x of H of multiplicity m_x , the inequality $\text{ord}_x A_P > \lfloor (b-1)m_x/b \rfloor$ similarly forces one of those critical values to have residue c at x . For each fixed x , at most $b-1$ labels can have this excess. Fixing one valuation extension at each place makes this count independent of how the sections choose critical roots.

Let r_H be the number of distinct roots of H , and suppose $r_H > D$. Select one section for each nonzero label not exceptional at infinity, and charge it

$$C(P) = \sum_{H(x) \neq 0} \text{ord}_x A_P + \sum_{H(x)=0} \max \left\{ 0, \text{ord}_x A_P - \left\lfloor \frac{(b-1)m_x}{b} \right\rfloor \right\}.$$

Off $H = 0$ a critical point determines the label F/H , so different selected labels have disjoint charged support. At a root of H at most $b-1$ labels have positive excess. The local length bound gives $\sum_P C(P) \leq (b-1)\ell$. Also

$$\begin{aligned} C(P) &\geq \left\lceil \frac{(b-1)N}{b} \right\rceil - \sum_{H(x)=0} \left\lfloor \frac{(b-1)m_x}{b} \right\rfloor \\ &\geq r_H - \lfloor N/b \rfloor \geq r_H - D > 0, \end{aligned}$$

since each summand in the last sum is at most $m_x - 1$. Adding the exceptional labels and the zero label, and allowing at most b sections per label, bounds the entire section bank by

$$|\mathcal{B}| \leq b^2 + \frac{b(b-1)\ell}{r_H - \lfloor N/b \rfloor} \leq b^2 + \frac{3b^2(b-1)^2 D}{r_H - D}. \quad (98)$$

Here (96) and $N \leq bD$ give the last step.

Now select any finite list of L sections with at least $a \geq D + \eta n$ agreements each. At a coordinate outside $H = 0$, the received symbol fixes one label and matches at most b sections. Thus

$$La \leq Lr_H + bn, \quad r_H - D \geq \eta n - bn/L.$$

For $L \leq \max\{b^2, b/\eta\}$ the stated bound is immediate. Otherwise $r_H > D$, and (98) implies

$$(L - b^2)(\eta - b/L) \leq 3b^2(b-1)^2 D/n.$$

Expanding and dropping the positive term b^3/L gives the claimed bound after taking its floor. Applying this to every finite subset also excludes an infinite list. $\square$

Corollary R.8. *Keep the weighted coefficient and section hypotheses of Theorem R.7, but do not assume its critical ideal is zero-dimensional. Assume characteristic zero or $p > b(b-1)D$. If a list exceeds the bound in that theorem, then, after extending constants to an algebraic closure k , there is a $c \in k$ such that $F(X, u) - cH(X)$ has a repeated nonconstant factor in $k(X)[u]$.*

Proof. The theorem forces a positive-dimensional critical component. It cannot be vertical because F_u has nonzero constant leading coefficient b . Its generic root r generates an extension $L/k(X)$ of degree $d \leq b-1$, separable in positive characteristic since $d < p$. The derivation extends to L , and the critical relations give

$$\left(\frac{F(X, r)}{H} \right)' = 0.$$

The excessive list includes a nonzero-label section, so $N \leq bD$. The monic weighted equation $F_u/b = 0$ makes r integral at finite places and bounds its pole at each place above infinity by D times the ramification index. Therefore $v = F(X, r)/H$ has pole-divisor height

$$h_L(v) \leq dN + d \max\{bD - N, 0\} = dbD \leq b(b-1)D.$$

In characteristic zero a function with zero derivative is constant. In characteristic p , the kernel of this nonzero derivation is L^p ; a nonconstant p th power has height at least p , contradicting the strict bound. Hence $v = c \in k$. The minimal polynomial of r divides both F_u and $F - cH$, proving the repeated-factor assertion. $\square$

The corollary is a structural condition, not a classification of the exceptional pencils. It does not assert that the repeated factor is linear or that its label descends to the original constant field.

Weighted monic-cubic first integrals. The next result treats polynomial sections of a cubic first integral whose denominator depends only on the evaluation coordinate. It bounds lists for arbitrary received words, including words supported on its critical fibers.

Theorem R.9. *Let $\mathbb{K}$ have characteristic zero or characteristic $p > \max\{3, 2D\}$, where $0 \leq D < n$. Let $0 \neq H \in \mathbb{K}[X]$ and*

$$F(X, u) = u^3 + a_2(X)u^2 + a_1(X)u + a_0(X), \quad \deg a_i \leq (3-i)D \quad (0 \leq i \leq 2).$$

Consider the section family

$$\mathcal{B} = \{P \in \mathbb{K}[X] : \deg P \leq D, \quad F(X, P) = cH \text{ for some } c \in \mathbb{K}\}.$$

On any n distinct evaluation coordinates, every received word has at most

$$\left\lfloor 9 + \frac{3 + 108D/n}{\eta} \right\rfloor$$

members of $\mathcal{B}$ agreeing on at least $D + \eta n$ coordinates, for every $\eta > 0$. Coordinates and received symbols may lie in an extension field.

This concerns the displayed first-integral family, not arbitrary cubic differential equations or all Reed–Solomon codewords. In particular, the characteristic and weighted coefficient hypotheses must both be checked; this theorem does not cover the Dickson differential-equation bank merely because its equation has cubic jet degree.

Proof. Extend constants to an algebraically closed field k containing all coordinates and received symbols, and put $K = k(X)$. Each section has a unique label c , and each label has at most three sections. The case $D = 0$ follows because distinct constants have disjoint agreement sets. Assume $D \geq 1$. If no nonzero label occurs, there are at most three sections. Otherwise, writing $N = \deg H$, the weighted degree caps imply $N \leq 3D$.

The critical scheme. Set

$$B = HF_X - H'F, \quad I = (F_u, B) \subset k[X, u].$$

If I is zero-dimensional, including the unit-ideal case, Theorem R.7 with $b = 3$ gives the asserted bound directly. Its characteristic hypothesis $p > 3$ follows from the present one. We therefore assume that I has a positive-dimensional component.

Positive-dimensional critical loci. Since F_u has constant leading u -coefficient 3, a positive-dimensional common component cannot be vertical. A root r of such a component satisfies

$$F_u(X, r) = 0, \quad (F(X, r)/H)' = 0.$$

We use the elementary height fact that a rational function of height less than p with zero derivative is constant: it belongs to K^p , and a nonconstant p -th power has height at least p . We also need a distinct-root budget. If A/H is a nonconstant p -th power with $\deg A, \deg H \leq 3D$, write

$$A = Ga^p, \quad H = Gb^p, \quad \gcd(a, b) = 1, \quad M = \max\{\deg a, \deg b\} \geq 1.$$

Then $\deg G + pM \leq 3D$, and the number n_0 of distinct roots of H satisfies

$$n_0 \leq \deg G + \deg b \leq 3D - (p-1)M \leq 3D - p + 1 \leq D.$$

In this case a received symbol off $H = 0$ fixes a label and matches at most three candidates. A list of L candidates with at least $a \geq D + \eta n$ agreements thus satisfies $La \leq LD + 3n$, and hence $L \leq 3n/(a - D) \leq 3/\eta$. We may finish immediately whenever such a nonconstant Frobenius ratio occurs.

If $r \in K$, monic integrality makes r polynomial, and the weighted caps give $\deg r \leq D$, so $\deg F(X, r) \leq 3D$. In positive characteristic a nonconstant ratio $F(X, r)/H$ is covered by the distinct-root budget just proved. Otherwise the ratio is constant, say c_0 , as it always is in characteristic zero. The repeated root of the cubic $F - c_0H$ then gives

$$F - c_0H = (u - R)^2(u - S), \quad R, S \in k[X], \quad \deg R, \deg S \leq D. \quad (99)$$

For completeness, an irreducible quadratic critical component cannot produce an additional large family under our sharper characteristic bound. Depress the cubic, writing

$$F = z^3 - 3gz + b, \quad z = u + a_2/3, \quad \deg g \leq 2D, \quad \deg b \leq 3D.$$

Here g is a nonsquare in K . Both conjugate critical values have zero derivative. Their trace gives $(b/H)' = 0$. A nonconstant ratio in positive characteristic is again covered by the distinct-root budget, since $\deg b \leq 3D$. In all remaining cases $b = \beta H$ for a constant β . Their difference gives $3g'H - 2gH' = 0$. A section with $\lambda = c - \beta \neq 0$ would satisfy $z(z^2 - 3g) = \lambda H$. Differentiating and eliminating H' yields

$$(z^2 - g)(2gz' - g'z) = 0.$$

The first factor is nonzero because g is a nonsquare, and $z \neq 0$. Thus $(g/z^2)' = 0$. Its height is at most $2D < p$, so g/z^2 is constant, contradicting nonsquareness over the algebraically closed constant

field. All sections therefore have label β , giving at most three. The same argument in characteristic zero uses the usual constant field of K instead of a height bound.

It remains to analyze (99). If $R = S$, any section outside this fiber makes H a constant times a polynomial cube, and all sections lie in an affine family $R + th$ with $\deg h \leq D$. If $R \neq S$, put $h = R - S$, $Y = (P - R)/h$, and $q = H/h^3$. Then

$$Y^2(Y + 1) = (c - c_0)q. \quad (100)$$

If q is constant, Y is constant and again all sections lie in $R + th$. In an affine family of degree-at-most- D polynomials, different members can match the same received symbol only at zeros of h , of which there are at most D . A list of L members with at least $a > D$ agreements therefore satisfies $L(a - D) \leq n$.

If q is nonconstant, Lemma R.10, with $r = 2$ and $s = 1$, permits at most two distinct nonzero labels $\lambda = c - c_0$ in (100). Its three-cover fiber product has genus three, so the argument also excludes inseparable parametrizations. There are at most six sections outside the repeated fiber and at most two within it, for a total of eight. This completes the positive-dimensional case.

The bounds obtained in this case—at most $3/\eta$ in the Frobenius case, at most $1/\eta$ for an affine family, or a bank of at most eight sections—are all bounded by the stated expression. This completes the proof. $\square$

A distinction for the Dickson family. For the complete nonzero Dickson bank, $n = p - 1$, $D = n/4 - 1$, and $n/2$ candidates agree with the displayed received word on $3n/8$ coordinates. If this bank lay in a common first-integral family as above, taking $\eta = 1/8 + 1/n$ would give

$$\frac{n}{2} \leq \left\lfloor 249 - \frac{2784}{n + 8} \right\rfloor.$$

This fails for $n \geq 488$, hence for every admissible prime $p \equiv 1 \pmod{8}$ with $p \geq 521$. Thus its primitive cubic differential equation does not arise from a weighted monic-cubic first integral of this form. This does not exclude higher-degree first integrals or rational denominators depending on the candidate value.

Constant fibers supported on two roots. The following treats constant fibers supported on at most two roots in every fixed value degree. It does not classify all repeated fibers.

Lemma R.10 (Three scaled two-root covers). *Let k be algebraically closed of characteristic zero or characteristic $p > r + s$, where $r, s \geq 1$ are coprime. For nonconstant $q \in k(X)$, at most two distinct $\lambda \in k^\times$ admit a rational solution of*

$$Y^r(Y + 1)^s = \lambda q(X).$$

Proof. Put $b = r + s$, $h(Y) = Y^r(Y + 1)^s$, and $\kappa = (-1)^r r^r s^s / b^b$. The cover $t = h(Y)/\lambda$ has inertia cycles $(r)(s)$ over zero, a b -cycle over infinity, and a transposition over κ/λ . Indeed $h' = Y^{r-1}(Y + 1)^{s-1}(bY + r)$; these are all branch points, and ramification is tame. For $r = s = 1$, zero is unramified.

Its monodromy group is S_b . To see primitivity, first exclude a polynomial decomposition $h = A \circ B$ of degrees greater than one. If A has two distinct roots, their disjoint preimages under B are supported on $\{0, -1\}$. Thus $B - \alpha = cY^m$ and $B - \beta = c(Y + 1)^m$, up to interchange, whose difference cannot be constant for $m > 1$ in the stated characteristic. If A has one root, its degree divides both r and s , also impossible. Any intermediate field gives such a polynomial decomposition:

by Lüroth it is rational, and the unique point over infinity makes both maps polynomial after choosing that point as infinity. Hence the group is primitive. The graph whose edges are conjugates of its transposition has connected components forming a block system. It is therefore connected, and its edge transpositions generate S_b .

For three distinct nonzero λ_i , take the compositum of the three Galois closures. Its group is a subdirect subgroup of S_b^3 . Inertia at κ/λ_i supplies a transposition in the i th factor and identities in the others. Conjugation by the surjective i th projection generates that entire factor, proving the group is S_b^3 . The normalized fiber product is thus connected of degree $N = b^3$. At infinity it has b^2 inertia orbits; each of the three unique simple branch points has index b^2 . At zero the orbit count is

$$r^2 + s^2 + 3r + 3s.$$

Here the all- r and all- s triples contribute r^2, s^2 ; the mixed triples have orbit length rs , contributing $3r, 3s$. Riemann–Hurwitz gives

$$2g - 2 = 2b^2 - r^2 - s^2 - 3b = b^2 + 2rs - 3b \geq (b - 2)(b + 1) \geq 0.$$

Three rational solutions would embed this positive-genus function field into $k(X)$ via $t = q(X)$, contradicting Lüroth. This also excludes inseparable q ; no assumption $q' \neq 0$ is used. $\square$

Theorem R.11 (Two-root constant fibers). *Let $D \geq 1$ and let k be algebraically closed of characteristic zero or characteristic $p > b = r + s$, with $r, s \geq 1$. Suppose $H \in k[X] \setminus \{0\}$ and a polynomial $F(X, u)$ has a constant fiber*

$$F(X, u) - c_0 H(X) = (u - R(X))^r (u - S(X))^s, \quad \deg R, \deg S \leq D.$$

Put $d = \gcd(r, s)$. The degree-at-most- D polynomial sections $F(X, P) = cH$ either all lie in one affine polynomial pencil with coefficient degrees at most D , or number at most $2b/d + 2$. Consequently, on any n distinct evaluation points and for any received word, at most

$$\max \{2b/d + 2, \lfloor 1/\eta \rfloor\}$$

such sections agree in at least $D + \eta n$ coordinates, for $\eta > 0$.

Proof. If $R \neq S$, put $a = R - S$, $Y = (P - R)/a$, $q = H/a^b$, and $\lambda = c - c_0$. Then $Y^r(Y + 1)^s = \lambda q$. For nonconstant q and $d = 1$, Lemma R.10 allows at most two nonzero labels, each with at most b rational roots; the zero fiber adds only $Y = 0, -1$.

For general d , write $h_0(Y) = Y^{r/d}(Y + 1)^{s/d}$. If a nonzero-label solution Y_* exists, set $q_0 = h_0(Y_*)$, so $q = q_0^d/\lambda_*$. Every other nonzero-label solution has

$$\left(\frac{h_0(Y)}{q_0} \right)^d = \frac{\lambda}{\lambda_*},$$

so $h_0(Y) = \mu q_0$ for a constant $\mu \neq 0$. The function q_0 is nonconstant. Applying the lemma to the coprime exponents $r/d, s/d$ permits at most two μ , each with at most b/d roots. Including the zero fiber proves $2b/d + 2$.

If q is constant, every Y is constant, giving the pencil $P = R + ta$. If $R = S$, the identity is $(P - R)^b = (c - c_0)H$. One noncritical section P_* expresses H as a constant multiple of $(P_* - R)^b$, after which every section belongs to the pencil $R + t(P_* - R)$. If none exists, there is only $P = R$.

Finally, in a pencil $P_0 + tP_1$ with $P_1 \neq 0$ of degree at most D , at most D coordinates have $P_1 = 0$; at every other coordinate at most one member matches a fixed received symbol. If L members have at least $A > D$ agreements, then $LA \leq LD + n$, proving $L \leq n/(A - D)$. Extending constants proves the same upper bound for banks over smaller fields. $\square$

S Contact resources and their equality cases

Actual first-jet contacts give a nonnegative resource for each interpolant factor, in every characteristic. Its zero case is exactly a power of a received-word graph. These statements concern factors; relating their charges to normal-divisor costs remains a separate step.

Let K be any field, let $x_1, \dots, x_n \in K$ be distinct, and let $r_i \in K$ be received symbols. Fix $w \geq 2$, give (X, Y, R) weights $(1, w, w-1)$, and write wt for weighted degree. For $0 \neq F \in K[X, Y, R]$, define

$$a_i(F) = \text{ord}_t F(x_i + t, r_i + tR + t^2E, R), \quad \epsilon(F) = \text{wt}(F) - \frac{w}{n} \sum_i a_i(F).$$

The substitution is injective after inverting t , so contacts are finite and additive on products. Weighted degrees, and hence charges, are additive as well. The coefficient field may be $k(Z)$.

Lemma S.1 (Subset contact bound). *If $w \leq n$, $W = \text{wt}(F)$, and $0 \leq u_i \leq a_i(F)$ are integers, then*

$$\# \left\{ T \subseteq [n] : |T| = w, \sum_{i \in T} u_i > W \right\} \leq \left\lfloor \frac{W}{w-1} \right\rfloor. \quad (101)$$

Proof. For each qualifying T , interpolate the symbols by P_T of degree less than w , and put $H_T = \prod_{i \in T} (X - x_i)$. Contact makes $F(X, P_T + cH_T, P'_T + cH'_T)$ divisible by $\prod_{i \in T} (X - x_i)^{u_i}$. Its degree in X is at most W , so it vanishes identically, with c indeterminate. Consequently

$$R - P'_T - \frac{H'_T}{H_T} (Y - P_T)$$

is a factor over $K(X)$. Distinct T give distinct factors: their slopes have simple poles of residue one at precisely their respective nodes, in every characteristic. Their number is at most the total jet degree, which is at most $\lfloor W/(w-1) \rfloor$. $\square$

Theorem S.2 (Universal contact resource). *Suppose $n \geq 2w$. Write $r = \deg_R F$, $A_r = [R^r]F$, $q = \deg_Y A_r$, and $B = [Y^q]A_r \neq 0$, with $d = \deg_X B$. Then, in every characteristic,*

$$a_i(F) \leq q + 2 \text{ord}_{x_i} B, \quad \epsilon(F) \geq r(w-1) + \left(1 - \frac{2w}{n}\right) d \geq r(w-1) \geq 0. \quad (102)$$

Proof. Fix a node and translate Y by its received symbol. This preserves r, q, B . In the translated polynomial take the homogeneous (Y, R) -degree $q+r$ piece G , whose $Y^q R^r$ coefficient is $B(x_i + t)$. The R -degree bound forces $G = Y^q H$, where H is homogeneous of degree r and $H(t, 0, 1) = B(x_i + t)$. Distinct homogeneous pieces remain distinct homogeneous pieces in (R, E) after contact substitution, so they cannot cancel. Thus

$$a_i(F) \leq \text{contact}(G) = q + \text{contact}(H).$$

Expand $h(t, Y) = H(t, Y, 1) = \sum_j c_j(t)(Y - t)^j$. Homogeneity and the substitution $Y = tR + t^2E$ give

$$b := \text{contact}(H) = \min_j \{ \text{ord}_t c_j + 2j \}.$$

But $B(x_i + t) = h(t, 0) = \sum_j c_j(t)(-t)^j$. Each summand has order at least $\lceil b/2 \rceil$, since $\text{ord}_t c_j \geq 0$. Hence $b \leq 2 \text{ord}_{x_i} B$. Summing over the distinct nodes yields $\sum_i a_i(F) \leq nq + 2d$. Finally $\text{wt}(F) \geq r(w-1) + qw + d$, proving the claim. $\square$

Corollary S.3 (Additive primary and uniform bounds). *Assume $n \geq 2w$. If $Q_0 = c \prod_j F_j^{e_j} \neq 0$ has contact at least m at every node and weight $W_0 < mA$, then*

$$(w-1) \sum_j e_j \deg_R F_j \leq \sum_j e_j \epsilon(F_j) < m(A-w). \quad (103)$$

If instead $F \neq 0$ has weight at most $aw + e$, where $a, e \geq 0$ are integers, and contact at least a at every node, then

$$\deg_R F \leq \left\lfloor \frac{e}{w-1} \right\rfloor.$$

In particular $e < w-1$ forces complete independence of R , including in small characteristic.

Proof. The theorem and additivity give the first lower bound, while

$$\sum_j e_j \epsilon(F_j) = W_0 - \frac{w}{n} \sum_i a_i(Q_0) \leq W_0 - mw < m(A-w).$$

For the uniform statement, $(w-1) \deg_R F \leq \epsilon(F) \leq e$. □

Theorem S.4 (Zero-charge rigidity in every characteristic). *Suppose $n \geq 2w$. Then $\epsilon(F) = 0$ if and only if*

$$F = c(Y - P(X))^q, \quad c \in K^\times, \quad \deg P \leq w, \quad P(x_i) = r_i \text{ for all } i,$$

for an integer $q \geq 1$, or F is a nonzero scalar ($q = 0$). Consequently, weight at most aw and contact at least a at every node force this form with $q = a$.

Proof. The universal bound first gives $\deg_R F = 0$. Put $q = \deg_Y F$ and let $B(X)$ be its leading coefficient, of degree d . The coefficient of $R^q E^0$ in the contact substitution gives the stronger bound $a_i(F) \leq q + \text{ord}_{x_i} B$. Therefore

$$\epsilon(F) \geq \text{wt}(F) - qw - \frac{w}{n}d \geq \left(1 - \frac{w}{n}\right)d.$$

Equality forces $d = 0$, $\text{wt}(F) = qw$, and $a_i(F) = q$ at every node. The case $q = 0$ is scalar.

First suppose q is nonzero in K . Write $F = BY^q + C(X)Y^{q-1} + \dots$, with $B \in K^\times$ and $\deg C \leq w$. The coefficient of $t^{q-1}R^{q-1}$ gives $qBr_i + C(x_i) = 0$. Thus $P = -C/(qB)$ has degree at most w and matches all symbols. Translate Y by $P(X)$. Weight is preserved and the new polynomial has zero-word contact q : locally, the difference $P(x_i + t) - P(x_i)$ is absorbed by an invertible translation of R . If $C_j(X)$ is its Y^j coefficient for $j < q$, contact implies

$$\prod_i (X - x_i)^{q-j} \mid C_j, \quad \deg C_j \leq (q-j)w < (q-j)n.$$

Hence every such coefficient vanishes, proving the required form.

It remains to treat characteristic $p \mid q$. Differentiating the contact identity in R shows that F_Y has contact at least $q-1$, since the derivative is $t(F_Y)_{\text{sub}}$. If $F_Y \neq 0$, its weight is at most $(q-1)w$. Nonnegative charge forces equality and zero charge. The preceding characteristic-free argument for R -free zero-charge polynomials would then give $\deg_Y F_Y = q-1$, whereas the leading term differentiates to zero and $\deg_Y F_Y \leq q-2$. Thus $F_Y = 0$.

Differentiating in t now gives contact at least $q - 1$ for F_X . If $F_X \neq 0$, let j be its Y -degree and $B_j(X)$ its leading coefficient. All Y -exponents of F are divisible by p , and its leading coefficient is constant, so $j \leq q - p$. Put $d' = q - j \geq 2$. The R -free contact bound forces

$$\deg B_j \geq n(d' - 1), \quad \deg B_j \leq d'w - 1.$$

This is impossible, since $n \geq 2w$ implies $n(d' - 1) \geq d'w$. Therefore $F_X = 0$ too. Over the perfect closure of K , write $F = G^p$. Then G has weight $(q/p)w$ and contact q/p at the same nodes. Repeat until its Y -degree is nonzero in the field, apply the previous case, and raise back to the appropriate power. The resulting P descends to K : interpolation at any $w + 1$ original nodes uniquely recovers it from the original symbols.

Conversely, the displayed graph power has weight qw and exact contact q . For the final assertion, nonnegative charge and the two bounds force zero charge, weight aw , and contact a . $\square$

For nodes in k and affine symbols $r_i = f_i + Zg_i$, interpolation at $w + 1$ nodes gives $P = P_0 + ZP_1$, with $\deg P_0, \deg P_1 \leq w$. Thus over $k(Z)$ every positive-weight irreducible zero-charge factor is a received-line graph factor. If the original polynomial lies in $k[Z, X, Y, R]$, the scalar in its graph-power expression is its leading Y -coefficient and lies in $k[Z]$.

Finite scope and the remaining geometric step. At $n = 262144$, $w = 131071$, $m = 118$, and $A = 181275$, the total charge is strictly below $m(A - w) = 5924072$. The resulting total R -degree bound is 45, weaker than the existing source cap 36; it gives no benchmark improvement. Nor does the resource alone exclude the binding carrier flags $(\deg_R, \deg_{Y,R}, \deg_{Y,R,Z}) = (12, 55, 3261)$. Writing $T = Y - Z$ and $\Lambda = \prod_i (X - x_i)$, the polynomial

$$F = T^{55} + \Lambda^{10}(R^{12} + R + Z^{3261}) - \Lambda^9 \Lambda' T$$

has these flags, weight $55w$, and exact contact ten for the word Z . Its charge is $45w = 5898195$, within the budget, and $Q_0 = FT^{108}$ has contact 118 and weight $163w < 118A$. This is source compatibility, not a bad-label construction. A first-tail component is a codimension-two intersection component, not another primary factor: charging each such component separately requires a new inequality controlling their aggregate normal-divisor cost from the carrier's single charge.

T A second one-pole extension over prime fields

The following is a fixed-size positive construction. It does not assert an unbounded list, a superlinear line-exception count, or a benchmark improvement.

Theorem T.1. *There are a number field K , a set $\Omega \subset K$ of sixty distinct points, a word $w: \Omega \rightarrow K$, and six distinct polynomials of degree at most fourteen, each agreeing with w on at least thirty points. The same statement holds over $\mathbb{F}_p$ for arbitrarily large primes p . Thus the code dimension is fifteen, exactly one quarter of the length.*

Proof. We first construct a five-polynomial word of length twenty-eight together with a proper one-pole rational function having fourteen agreements.

Four quadratics and the first rational witness. Choose nonzero parameters $a_1, \dots, a_4$ with distinct squares and with all twelve signed pair products distinct. Put

$$H_i(X) = X^2/a_i^2 + a_i^2, \quad \Omega_0 = \{\pm a_i a_j : i < j\}.$$

At either signed node for the pair ij , give the received word the value $a_i^2 + a_j^2$. Exactly H_i, H_j agree there, and every H_i has six agreements. Write e_j for the elementary symmetric functions of the four parameters and set

$$\begin{aligned} D_1(X) &= e_1 e_4 - e_3 X, \\ N_1(X) &= e_1 X^3 + (e_3 - e_1 e_2) X^2 + (e_2 e_3 - e_1 e_4) X - e_3 e_4. \end{aligned}$$

The coefficient identity

$$N_1 - D_1 H_i = \frac{\prod_{j \neq i} (a_i + a_j)}{a_i^2} \prod_{j \neq i} (X - a_i a_j)$$

follows by using $a_i^4 - e_1 a_i^3 + e_2 a_i^2 - e_3 a_i + e_4 = 0$. Assume $e_3 \neq 0$ and that $\alpha = e_1 e_4 / e_3$ lies outside Ω_0 . Then N_1 / D_1 is proper and agrees at the six positive pair nodes. Indeed, a common zero of N_1, D_1 would, by the identity, lie among the positive incident edges for every vertex, which is impossible.

Choose c outside $\Omega_0 \cup \{\alpha\}$ and put $y = T^2 + c$. On the twenty-six-point core

$$\mathcal{C} = \{T : T^2 + c \in \Omega_0 \cup \{\alpha\}\}$$

use the word

$$W(T) = (y - \alpha)(a_i^2 + a_j^2) \quad (y = \pm a_i a_j), \quad W(T) = 0 \quad (y = \alpha).$$

The five polynomials

$$G_i(T) = (y - \alpha)H_i(y) \quad (1 \leq i \leq 4), \quad G_5(T) = -N_1(y)/e_3$$

have degree at most six. The first four have fourteen core agreements; G_5 has twelve, over the six positive pair nodes. They are distinct, since equality with G_5 would make the first rational witness polynomial. All fibers here have two distinct points.

An exact finite-field certificate for a second rational witness. The following data are in $\mathbb{F}_{97}$:

$$(a_1, a_2, a_3, a_4) = (48, 1, 31, 32), \quad \alpha = 53, \quad c = 0, \quad b = 81,$$

$$\begin{aligned} N(T) &= 34 + 54T + 93T^2 + 84T^4 + 3T^5 + 56T^6 + 70T^7, \\ G_5(T) &= 86 + 17T^4 + 93T^6. \end{aligned}$$

The proper function $N(T)/(T - b)$ agrees at the following twelve core points. The first row is T , the second is $y = T^2 + c$, and the third is $W(T)$:

T	76	57	90	79	8	88	93	39	56	68	33	47
y	53	48	49	33	64	81	16	66	32	65	22	75
W	0	18	92	78	25	64	54	90	9	78	60	20

These are one pole-fiber point and one point above every old base node except the positive edge $a_2a_3 = 31$. Direct substitution verifies the agreements. The pole is outside the entire core and $N(81) = 48 \neq 0$. Moreover, exactly five displayed points are agreements of G_5 , and

$$N(T) - (T - 81)G_5(T) = \prod_{t \in \{57, 79, 88, 56, 33\}} (T - t)(74T^2 + 2T + 8).$$

The remaining quadratic has discriminant $61 \neq 0$ and roots $24, 52$, both outside the core and the pole. Giving these two fresh coordinates the values of G_5 therefore produces the desired length-twenty-eight word over $\mathbb{F}_{97}$. To transfer this construction to characteristic zero, it remains essential to verify smooth lifting; the finite-field example alone would not suffice.

Polynomial equations and a nonsingular lifting certificate. Write

$$N(T) = E(y) + TO(y), \quad E(y) = \sum_{j=0}^3 u_j y^j, \quad O(y) = \sum_{j=0}^3 v_j y^j.$$

Let $\mathcal{Y}$ consist of α and all signed pair products except $+a_2a_3$. For $y = \alpha$ put $w_y = 0$; for $y = \pm a_i a_j$ put $w_y = (y - \alpha)(a_i^2 + a_j^2)$. Consider the following thirteen polynomial equations in fifteen variables, ordered as $(a_1, \dots, a_4, \alpha, c, b, u_0, \dots, u_3, v_0, \dots, v_3)$:

$$e_3\alpha - e_1e_4 = 0, \quad (E(y) + bw_y)^2 - (y - c)(O(y) - w_y)^2 = 0 \quad (y \in \mathcal{Y}). \quad (104)$$

Their solution at the displayed point is

$$(48, 1, 31, 32, 53, 0, 81, 34, 93, 84, 56, 54, 0, 3, 70).$$

The 13×13 Jacobian minor on the first thirteen columns is nonzero modulo 97. Exact modular differentiation and elimination certify rank thirteen; the reproducible certificate is `transversal_search/independent_lift.py` with its accompanying JSON. At all twelve nodes $O(y) - w_y \neq 0$, so the selected root is recovered by

$$T = -\frac{E(y) + bw_y}{O(y) - w_y}.$$

Consequently these norm equations do not introduce a spurious square-root choice near the certified point.

There is also an independently checked pointed certificate: use variables

$$(a_1, \dots, a_4, \alpha, c, b, N_0, \dots, N_7, T_0, \dots, T_{11})$$

and equations $e_3\alpha - e_1e_4 = 0$, followed, in the table order, by

$$T_j^2 + c - y_j = 0, \quad N(T_j) - (T_j - b)w_{y_j} = 0.$$

The minor on columns zero through twenty-four has determinant 21 modulo 97. Thus the lifting conclusion has a second certificate using ordinary polynomial evaluations rather than the norm equations.

Fix $v_2 = 3, v_3 = 70$ as integers in (104). Multivariate Hensel lifting gives a solution of the remaining thirteen equations in $\mathbb{Z}_{97}$, with invertible Jacobian. Its coordinates are algebraic over $\mathbb{Q}$: in the field they generate, differentiating the thirteen equations and using the invertible Jacobian

makes every coordinate differential zero; in characteristic zero this forces transcendence degree zero. They therefore lie in a number field embedded in $\mathbb{Q}_{97}$.

All required guards hold at the residue point and hence remain nonzero in this lift. They include distinct nonzero seed parameters and signed nodes, $e_3 \neq 0$, the first pole and critical-value exclusions, $O(y) - w_y \neq 0$, the new pole being off the whole core, and $N(b) \neq 0$. The five forced roots of $N - (T - b)G_5$ persist. Its remaining quadratic still has nonzero leading coefficient and discriminant and has no root in the core or at b : these are polynomial nonvanishing conditions, verified by the displayed finite-field factorization. Adjoin its two roots and the other core square roots to the number field. They supply two distinct fresh coordinates. Thus the five polynomials G_i each have at least fourteen agreements on a twenty-eight-point word, and the proper function $N/(T - b)$ also has at least fourteen.

The second pole-clearing extension. Choose a quadratic $\psi(U) = U^2 + d$ whose critical value avoids the length-twenty-eight domain and b , and split its finitely many required fibers in a number-field extension. Use the fifty-six preimages of that domain, the two preimages of b , and two further fresh points. The six polynomials are

$$(\psi - b)G_i(\psi) \quad (1 \leq i \leq 5), \quad N(\psi).$$

They have degree at most fourteen. On the first fifty-six points set the received word to $(\psi - b)$ times the preceding word, on the pole fiber set it to zero, and on the two fresh points give it the value $N(\psi)$. Every old candidate has at least $28 + 2$ agreements; the new one has at least $28 + 2$. The old candidates remain distinct, and equality of the new candidate with any old one would make $N/(T - b)$ polynomial, contrary to properness.

Finally, all data lie in one number field. Exclude the finitely many primes where denominators, distinctness, or any used nonzero guard fails. At every remaining completely splitting prime the data reduce into the prime field and preserve the sixty nodes, six distinct polynomials, degree bounds, and displayed agreements. There are arbitrarily large such primes, proving the theorem. $\square$

U An explicit seven-polynomial word and its pullbacks

Theorem U.1. *For every integer $m \geq 1$, there are a number field K_m , a set of $14m$ distinct evaluation points, and a word whose complete list of polynomials of degree at most $3m$ with at least $7m$ agreements consists of exactly seven polynomials, each with exactly $7m$ agreements. For each fixed m , the same holds over arbitrarily large prime fields. For $m = 1$ it holds over every prime field with $p \equiv 1 \pmod{7}$; for $m = 2$ it holds over every such field with $p > 10^{10}$. The latter case has length 28, dimension 7, and agreement 14, so has exactly quarter rate.*

Proof. We first give a construction over any field k of characteristic different from 2, 3, 7 containing a primitive seventh root ζ . Write

$$\eta = \zeta + \zeta^2 + \zeta^4, \quad \bar{\eta} = \zeta^3 + \zeta^5 + \zeta^6 = -1 - \eta.$$

Then $\eta^2 + \eta + 2 = 0$ and $\eta\bar{\eta} = 2$. Set

$$A(X) = X^3 - \eta X^2 + \bar{\eta}X - 1 = \prod_{j \in \{1,2,4\}} (X - \zeta^j), \quad P(X) = \eta X^3 + X^2 + X + \bar{\eta},$$

$$\alpha = \frac{\bar{\eta}}{\eta} = \frac{\eta - 1}{2}, \quad c = \frac{3(\eta + 3)}{4}.$$

The construction rests on the two polynomial identities

$$P(X) - X^5 = (\bar{\eta} - X)(X - 1)A(X), \quad (105)$$

$$P(\alpha X) - cX^5 = -A(X)(cX^2 + c\eta X + \bar{\eta}). \quad (106)$$

Both follow by multiplication and $\eta^2 = -\eta - 2$. All parameters needed below are nonzero, and $\alpha^7 \neq 1$. For completeness, under the involution $\eta \leftrightarrow \bar{\eta}$,

$$c\bar{c} = \frac{9}{2}, \quad (\alpha^7 - 1)(\bar{\alpha}^7 - 1) = \frac{7^3}{2^7}.$$

These identities prove the assertions in every allowed characteristic, including when the quadratic algebra splits.

Take the fourteen-point domain $\Omega = \mu_7 \sqcup \alpha\mu_7$ and define its received word by

$$w(\zeta^t) = \zeta^{5t}, \quad w(\alpha\zeta^t) = c\zeta^{5t} \quad (t \in \mathbb{Z}/7\mathbb{Z}).$$

The seven candidates are

$$P_i(X) = \zeta^{5i}P(\zeta^{-i}X) \quad (i \in \mathbb{Z}/7\mathbb{Z}).$$

Their constant coefficients are distinct. The identities above give four prescribed matches for each candidate on μ_7 and three on $\alpha\mu_7$. Each point of the first orbit has four prescribed matching candidates, and each point of the second has three. These incidences already use the entire pairwise root budget:

$$7\binom{4}{2} + 7\binom{3}{2} = 63 = 3\binom{7}{2}.$$

Indeed, two distinct cubics share at most three evaluation values. Consequently there are no further incidences: every candidate has exactly seven matches, and the two orbit multiplicities are exactly four and three.

The same argument handles pullbacks and proves completeness. Let $\psi(U) \in k[U]$ have degree m , and suppose each point of Ω has m distinct preimages in k . Pull back the domain, word, and seven candidates along ψ . Their degrees are at most $3m$, their agreement counts are $7m$, and the matching multiplicities remain four and three. If a different polynomial Q of degree at most $3m$ has at least $7m$ agreements, counting its intersections with the seven candidates gives

$$3\#\{\text{agreements of } Q\} \leq 7(3m).$$

Equality is necessary, so every agreement lies above $\alpha\mu_7$. Thus Q agrees at all $7m$ such points with $c\alpha^{-5}\psi(U)^5$. Their difference has degree $5m$, with nonzero leading coefficient, and cannot have $7m$ distinct roots. This proves completeness, even for candidates over an extension of k .

Over a number field containing ζ , choose any shift a outside Ω and take $\psi(U) = U^m + a$. Adjoining its finitely many fiber roots gives the required K_m and distinct fibers. Reduction at sufficiently large completely split primes of a normal closure preserves all these identities and distinctness; there are infinitely many such primes. For $m = 1$, the construction is already defined over every $\mathbb{F}_p$ with $p \equiv 1 \pmod{7}$.

For the asserted uniform prime bound when $m = 2$, it remains to find a shift $a \in \mathbb{F}_p$ such that all fourteen values $x - a$, $x \in \Omega$, are nonzero squares. Let χ be the quadratic character, with $\chi(0) = 0$. The distinct-root quadratic-character Weil bound [41, Lemma 2.1] gives

$$\#\{a : \chi(x - a) = 1 \text{ for every } x \in \Omega\} \geq \frac{p - 98305\sqrt{p}}{16384} - 14.$$

To see the constant directly, expand $2^{-14} \sum_a \prod_{x \in \Omega} (1 + \chi(x - a))$. A nonempty subset of size j contributes error at most $(j - 1)\sqrt{p}$, and $\sum_{j=1}^{14} \binom{14}{j} (j - 1) = 98305$. Removing the fourteen possible zero locations costs at most fourteen. The displayed lower bound is positive for $p > 10^{10}$. The map $\psi(U) = U^2 + a$ therefore supplies the required distinct split quadratic fibers. $\square$

This is a fixed list-size construction. It does not assert an unbounded list, a superlinear number of exceptional challenges, or a universal list bound of seven for arbitrary words at length twenty-eight.

V The sharp cubic list bound on fourteen coordinates

Theorem V.1. *There is an absolute constant p_0 with the following property. If k has characteristic zero or characteristic greater than p_0 , then for every fourteen-element set $\Omega \subset k$ and every word on Ω , at most seven distinct polynomials of degree at most three have at least seven agreements with the word. The bound is attained in characteristic zero and over arbitrarily large prime fields.*

The lower bound is Theorem U.1 with $m = 1$. The upper bound proved here is specific to cubics on fourteen coordinates; it does not assert a universal list bound of seven for the quarter-rate length-twenty-eight pullback. The proof below includes two finite exact certificates: an exhaustive incidence enumeration and a rational polynomial-identity certificate. Their algorithms and the mathematical conditions they certify are specified explicitly.

Proof. Saturation and incidence reduction. Suppose eight distinct cubics have at least seven agreements. For each coordinate let s_x be their matching multiplicity. Pairwise root counting and convexity give

$$\sum_x s_x \geq 56, \quad 84 \leq \sum_x \binom{s_x}{2} \leq 3 \binom{8}{2} = 84.$$

Thus every coordinate has multiplicity four, every candidate has exactly seven matches, and every pair has exactly three common matches. The fourteen matching sets form a 2-(8, 4, 3) design. Its blocks are distinct: fixing a block B , the intersections j with the other thirteen blocks satisfy $\sum j = 24$, $\sum \binom{j}{2} = 12$, and hence $\sum (j - 1)(j - 2)/2 = 1$. All summands are nonnegative, whereas a repeated block would contribute three.

Delete a candidate. The seven coordinates at which it matched become triple columns T ; the other seven remain quadruple columns. Write C for the complementary triples of these quadruples in the remaining seven candidates. Both T and C are three-regular collections of seven triples, and every pair has combined multiplicity two in T, C . The extension back to eight candidates is forced: append index eight to each T -triple and leave the complementary- C quadruples unchanged.

The finite incidence certificate. Let

$$\mathcal{F} = (123, 145, 167, 246, 257, 347, 356),$$

where strings denote subsets. Every 2-(8, 4, 3) design arising above has a deletion which, after relabeling, has $T = \mathcal{F}$ and one of the following four ordered choices for C :

$$\begin{array}{c|l} 4 & 123, 145, 167, 246, 257, 347, 356 \\ 5 & 123, 145, 167, 247, 256, 346, 357 \\ 6 & 123, 146, 157, 247, 256, 345, 367 \\ 7 & 124, 136, 157, 237, 256, 345, 467. \end{array} \tag{G1}$$

Here is the exhaustive algorithm verifying this finite assertion. Enumerate nondecreasing lists of seven of the 35 triples on seven vertices, retaining exactly those with vertex degrees three and pair degrees at most two. Repetitions are permitted. Encode the 21 pair degrees in base three. Two lists are compatible exactly when their digits sum to two, so a complementary-key lookup lists every ordered pair (T, C) . Quotient by simultaneous S_7 relabeling, without interchanging T, C . There are 12,780 labelled halves, 6,150 compatible ordered pairs, and eight ordered orbits. For each representative, form the forced eight-candidate extension and try all eight deletions. Each has a deletion of one of the four types (G1). The other four types have deletion sets $\{0, 3, 6\}$, $\{1, 5\}$, $\{2, 7\}$, and $\{0, 3, 6\}$; the four displayed types have deletion sets $\{4\}$, $\{1, 5\}$, $\{0, 3, 6\}$, and $\{2, 7\}$, respectively.

The implementations `designs.cpp` and `independent_design_set.py` independently verify the literal set of all 6,150 pairs, not only its cardinality. `eight_bank_removal_types.py` checks the deletion assertion. All these files and the explicit representatives are in the accompanying `fano_seven` certificate directory. The only enumeration pruning discards partial lists whose vertex or pair degrees already exceed the prescribed bounds, or whose vertices cannot reach degree three. Consequently no admissible incidence pattern is discarded.

We now exclude types 4–6 in characteristic different from two and classify type 7 in characteristic zero. We work over an algebraic closure and with binary cubics, so projective coordinate changes are permitted. Every pair already agrees at three prescribed distinct points. Therefore an additional pair equality is impossible, including an outsider sharing a quadruple's value.

Type 4. Normalize the seven quadruple coordinates, in the order of C , to $\infty, 0, 1, u, v, w, z$. All finite parameters are distinct and different from zero and one. On each T -triple, divide the three difference identities by the common triple-coordinate factor. The resulting three binary quadratics, whose roots are the prescribed quadruple coordinates, are linearly dependent. For triples 167, 246, 347 their determinants are

$$E_2 = -uz + vw, \quad E_3 = vw - vz - v + z, \quad E_5 = uz - vz + v - z.$$

All vanish, but $E_3 - E_5 - E_2 = 2(z - v)$ contradicts distinctness.

Type 5. Use the same normalization. The corresponding quadratic determinants for 123, 145, 167 give

$$-uvw - uvz + uv + uwz + vwz - wz = 0, \quad -uw + u + vz - v + w - z = 0, \quad uz - vw = 0.$$

Substitution of $u = vw/z$ into the first two forces $v(w + z - 1) = z$. With $s = w + z - 1$ we obtain $s \neq 0$, $u = w/s$, $v = z/s$. Subtract P_1 so that $P_1 = 0$. All quadruple-compatible families are

$$\begin{aligned} P_2 &= (X - w)(X - z)L_0, & P_3 &= s^2(X - u)(X - v)L_0, \\ P_4 &= (X - z)(X - v)L_1, & P_5 &= (X - w)(X - u)L_1, \\ P_6 &= (X - z)(X - u)L_2, & P_7 &= (X - w)(X - v)L_2, \end{aligned}$$

where $L_0 = pX + q$, $L_1 = \ell(X - 1) - s(p + q)$, and $L_2 = \ell X + sq$. Indeed the two forced quadruple roots determine each quadratic factor, and the equalities at $0, 1, \infty$ determine the remaining linear factors as displayed. In particular $\ell \neq 0$ and $p - \ell \neq 0$, by the outsider guards at infinity. Define

$$\begin{aligned} A &= \ell z + pzs + qs(s + 1), & B &= \ell w + pws + qs(s + 1), \\ C_z &= \ell(z - 1) + s(w - 1)p + s(s - 1)q, & C_w &= \ell(w - 1) + s(z - 1)p + s(s - 1)q, \\ F_{\pm} &= \pm(z - w)\ell + s(s + 1)p + 2s^2q. \end{aligned}$$

The factors A, B, C_z, C_w are nonzero: respectively they are, up to nonzero scalar factors, $(P_2 - P_4)(0)$, $(P_2 - P_5)(0)$, $(P_2 - P_6)(1)$, and $(P_2 - P_7)(1)$. For triple 246, divide the differences by their common factor $X - z$. The resultant of $s[(X - w)L_0 - (X - v)L_1]$ and $s[(X - w)L_0 - (X - u)L_2]$ equals $-(p - \ell)sAC_zF_+$. The analogous resultant for triple 257, after dividing by $X - w$, is $-(p - \ell)sBC_wF_-$. The required additional common roots force both resultants to vanish, so $F_+ = F_- = 0$. Their difference $2(z - w)\ell$ is nonzero, a contradiction.

Type 6. Relabel candidates by $(1, 2, 3, 4, 5, 6, 7) \mapsto (1, 2, 3, 4, 7, 5, 6)$. Then $C = \mathcal{F}$ and $T = (123, 147, 156, 245, 267, 346, 357)$. Normalize the quadruple coordinates to $\infty, 0, 1, u, v, w, z$ and subtract P_1 . Scale the nonzero quadruple value at zero to one, and write the values at one and infinity as K, L , both nonzero. Put

$$\begin{aligned} Q_2 &= (X - w)(X - z), & Q_3 &= (X - u)(X - v), & Q_4 &= (X - v)(X - z), \\ Q_5 &= (X - u)(X - w), & Q_6 &= (X - v)(X - w), & Q_7 &= (X - u)(X - z). \end{aligned}$$

The complete family is

$$\begin{aligned} P_i &= Q_i[(K/Q_i(1) - 1/Q_i(0))X + 1/Q_i(0)] & (i = 2, 3), \\ P_i &= Q_i[L(X - 1) + K/Q_i(1)] & (i = 4, 5), \\ P_i &= Q_i[LX + 1/Q_i(0)] & (i = 6, 7). \end{aligned}$$

The triple 123 and the two other triples containing index one force, with $s = w + z - 1$,

$$v = \frac{wz(1 - u)}{wz - us}, \quad L = \frac{K}{(1 - v)(1 - z)} - \frac{1}{uz} = \frac{K}{(1 - u)(1 - w)} - \frac{1}{vw}. \quad (\text{G2})$$

The first denominator cannot vanish, since then its defining equation would give $u = 1$. Substitution factors the difference of the two expressions for L into nonzero guarded factors times

$$(us - w)G, \quad G = Kw^2z(u - z) + (u - w)(w - 1)(z - 1)^2.$$

On $G = 0$, one obtains

$$K = \frac{(u - w)(w - 1)(z - 1)^2}{w^2z(z - u)}, \quad L = \frac{us - 2wz + w}{w^2z(z - u)}.$$

The leading coefficient of P_2 then equals L , an impermissible extra agreement at infinity. Thus $u = w/s$, $v = z/s$, where $s \neq 0, 1$.

For a pair, divide its difference by its two prescribed quadruple-root factors and take the determinant of the remaining linear forms on each triple. Set

$$\begin{aligned} A &= Kwz(w + z), & H &= (w - 1)(z - 1)(s - 1), \\ F_3 &= -Kw^2z + (w - 1)(z - 1)^2, & F_4 &= Kwz^2 - (w - 1)^2(z - 1). \end{aligned}$$

After (G2), the determinants for 245, 267, 346, 357 are respectively

$$\frac{\frac{2(z - w)(A + H)}{w^2z^2(w - 1)(z - 1)}, \quad \frac{2K(z - w)(A + H)}{wz(w - 1)^2(z - 1)^2}, \quad \Delta = w^2z^2(w - 1)^2(z - 1)^2,}{\frac{2s^2F_3(A - H)}{\Delta}, \quad \frac{-2s^2F_4(A - H)}{\Delta}},$$

All denominators are guarded and $H \neq 0$. Hence $A = -H$ and $F_3 = F_4 = 0$. The latter imply $(w - z)(w + z - 1) = (w - z)s = 0$, again impossible. The displayed determinant and resultant

identities in these three cases are ordinary polynomial identities, independently checked in the corresponding `ORBIT5` and `ORBIT6` certificate files.

Type 7: a complete normalized classification. Normalize its quadruple coordinates to $\infty, 0, 1, u, v, w, z$, subtract P_1 , and use amplitudes j, k, ℓ for the nonzero values at $0, 1, \infty$. Put

$$\begin{aligned} Q_2 &= (X - w)(X - z), & Q_3 &= (X - v)(X - z), & Q_4 &= (X - u)(X - v), \\ Q_5 &= (X - u)(X - z), & Q_6 &= (X - u)(X - w), & Q_7 &= (X - v)(X - w). \end{aligned}$$

Its complete quadruple-compatible family is

$$\begin{aligned} P_i &= Q_i[(k/Q_i(1) - j/Q_i(0))X + j/Q_i(0)] & (i = 2, 4), \\ P_i &= Q_i[\ell(X - 1) + k/Q_i(1)] & (i = 3, 6), \\ P_i &= Q_i[\ell X + j/Q_i(0)] & (i = 5, 7). \end{aligned}$$

Every denominator is a node guard. Dividing pair differences by their two quadruple-root factors leaves linear forms. On each of the seven triples the determinant of two such forms vanishes. In the order 123, 145, 167, 246, 257, 347, 356, divide these determinants by

$$-k, -j, -\ell, -P_2(u), [X^3]P_2 - \ell, j - P_3(0), -P_3(u), \quad (\text{G3})$$

respectively. Each divisor is a nonzero outsider–quadruple difference. The resulting seven expressions are linear in j, k, ℓ , giving a 7-by-3 matrix M whose entries are explicit rational functions of u, v, w, z . Clear only node-guard denominators.

Number rows from zero. Minor $(0, 1, 3)$ factors, up to nonzero guards, as

$$[(w - 1)(z - v)u + vw(w - v)][(v(w + z - 1) - wz)u + wz(1 - v)]. \quad (\text{G4})$$

On the second branch its denominator is nonzero, since otherwise $wz(1 - v) = 0$. The substituted minors $(0, 1, 2)$ and $(0, 3, 4)$ are, up to nonzero guards, the squares of

$$v(w + z - 1) - 2wz + z, \quad v(w + z - 1) - z.$$

Their simultaneous vanishing contradicts $2z(1 - w) \neq 0$. Every admissible realization therefore satisfies

$$u = \frac{vw(v - w)}{(w - 1)(z - v)}. \quad (\text{G5})$$

This assertion also holds after each relabeling preserving the ordered pair (T, C) and renormalization of its first three quadruple coordinates. There are twenty-one such relabelings; only validity of the listed relabelings, not completeness of that group list, is needed below.

Here is the exact identity certificate for the remaining elimination. Substitute (G5) in all 35 minors of M and in the relabeled versions of (G5). Cancel only nonzero node guards. Adjoin a variable h and the equation $h(v - z)(w - 1) - 1 = 0$. Rational polynomial linear combinations of these necessary equations give

$$(z - 1)(w^2 + w - z) = 0, \quad (z - 1)(wz - 2w + z - 1) = 0, \quad (z - 1)(v - z + 1) = 0. \quad (\text{G6})$$

For reproducibility, `orbit7_generic.py` constructs M by the operations just specified, `orbit7_branches.py` performs (G4)–(G5), and `orbit7_symmetry.py` checks each relabeling and projective renormalization. The identity file `orbit7_identity_certificate.json.gz` contains 380 nodes. An input node is one of these necessary equations. Every subsequent node is an explicit sum of polynomial

multiples of earlier nodes. An independent verifier, `orbit7_certificate_independent.py`, checks the inputs against the necessary-equation files and checks all 359 subsequent identities by sparse rational arithmetic, then checks the three output identities (G6). It uses no symbolic-algebra library. Thus correctness or completeness of a Groebner-basis implementation is not a trusted step of this certificate.

Since $z \neq 1$, equations (G5)–(G6) imply

$$w^3 + 2w^2 - w - 1 = 0, \quad u = w^2, \quad v = w^2 + w - 1, \quad z = w^2 + w.$$

On this cubic locus the minor of rows 0, 1 and columns j, k of M is $6w^2 - 2w - 2 \neq 0$. The matrix has rank two, and after scaling $j = 1$ its unique kernel vector is

$$(j, k, \ell) = (1, 3w^2 + 4w - 5, 3w^2 + 4w - 6).$$

Nonvanishing follows from irreducibility of the cubic; substitution checks the kernel vector in all seven rows. Thus every admissible type-7 realization has the displayed coefficients, up to the invertible transformations already made. Its first five triple coordinates are

$$a_1 = 1 - w^2, \quad a_2 = w^2 + 2w - 1, \quad a_3 = w^2 + 2w + 1, \quad a_4 = w^2 - 1, \quad a_5 = -w^2 - 2w + 3.$$

An eighth cubic must match every triple coordinate. The first three received values are zero, so it must be $\lambda \prod_{i=1}^3 (X - a_i)$. The fourth value determines $\lambda = (1 - 4w - 2w^2)/4$. At the fifth its difference from the received value is $486 - 390w - 270w^2$, which is nonzero modulo $w^3 + 2w^2 - w - 1$. These are direct substitutions in the displayed polynomials. This excludes the last possibility in (G1) and proves the characteristic-zero upper bound.

Finally the argument uses finitely many rational polynomial identities. Exclude primes dividing their coefficient denominators and the nonzero integer resultants separating $f(T) = T^3 + 2T^2 - T - 1$ from the rank minor and the interpolation obstruction. Outside this finite set, and characteristic two, every step remains valid over an algebraic closure. Taking p_0 larger than those finitely many primes proves the asserted characteristic range; no optimal value of p_0 is claimed. $\square$

W Rigidity of the seven-cubic banks and a one-pole obstruction

We work in characteristic zero. A seven-cubic bank means seven distinct sections of $\mathcal{O}_{\mathbf{P}^1}(3)$ agreeing with a word on fourteen distinct points in at least seven positions each. Equivalence permits projective changes of the evaluation coordinate, a common nonzero scaling, and addition of a common cubic section.

Proposition W.1. *Every seven-cubic bank is equivalent to a Galois conjugate of one of the two banks of incidence types 2 and 7 in the accompanying incidence certificate. In particular, type 2 has no additional realization component or nontrivial normalized infinitesimal deformation.*

Proof. Pair counting gives at least 63 common pair agreements and at most $3\binom{7}{2} = 63$. Equality forces seven triple columns, seven quadruple columns, seven agreements per candidate, and three per pair. Writing T for the triple columns and C for the complements of the quadruple columns gives the eight ordered incidence types enumerated in Appendix V. The guard-preserving identities in `NONFANO_ORBITS_01_OBSTRUCTION.md`, `ORBIT3_OBSTRUCTION.md`, and the type-4–6 proof of that appendix exclude all but types 2 and 7. The normalized type-7 classification is proved in that appendix. Here we complete the type-2 classification.

Its ordered blocks are

$$C = (137, 145, 167, 236, 247, 256, 345), \quad T = (123, 124, 156, 257, 346, 357, 467).$$

Normalize the quadruple coordinates to $\infty, 0, 1, u, v, z, d$. After subtracting a common cubic, all quadruple-compatible families are obtained by reducing $W = aX^3 + bX^4 + cX^5$ modulo each candidate's finite quadruple locator. Indeed, set the infinity word to zero, interpolate the other six values by a quintic, and subtract its quadratic part. A candidate matching infinity has a cubic locator and a quadratic remainder; the others have quartic locators and cubic remainders. The amplitude vector is nonzero, since the candidates are distinct.

The seven triple conditions give seven linear equations in (a, b, c) . The 156 equation has a -coefficient -1 , so eliminate a without a denominator. A two-by-two minor of the resulting six-by-two matrix is

$$-(d-1)(d-uv+uz-z)(-dvz+dv+uvz-uv+uz-vz).$$

Thus there are two branches. On $d = uv - uz + z$, another equation gives $u = vz/D$, where $D = v + z^2 - vz \neq 0$. A remaining minor is

$$\frac{vz(v-z)^2(v+z^2-2vz)}{D}.$$

It is nonzero: its last factor divided by D is $1 - u$. This branch is impossible.

On the other branch,

$$d = \frac{uv(z-1) + z(u-v)}{v(z-1)}, \quad u = \frac{vz^2(v-z)}{D_1}, \quad D_1 = v^2z - v^2 - vz^3 + 2vz^2 - vz + v - z^2.$$

The first denominator is a node guard; the second cannot vanish because its defining equation would give $vz^2(v-z) = 0$. Two necessary remaining equations are

$$\begin{aligned} E &= v^2z - 2v^2 + vz + v - z^2, \\ F &= v^3z^3 - v^3z^2 + v^3z - v^3 - 2v^2z^3 + 2v^2z^2 + v^2z - 2vz^3 + z^5. \end{aligned}$$

Direct polynomial elimination gives

$$\text{Res}_v(E, F) = z^6(z-1)^4 f(z), \quad f(z) = z^3 - 10z^2 + 3z + 1.$$

Since $z \neq 0, 1$, we have $f(z) = 0$. Reduction of E, F on this cubic gives $25v = -4z^2 + 27z + 7$, and hence

$$u = \frac{-z^2 + 13z + 3}{20}, \quad v = \frac{-4z^2 + 27z + 7}{25}, \quad d = \frac{-3z^2 + 29z + 4}{20}.$$

The amplitude matrix has rank two: its rows 0, 1, columns 0, 1 minor is 1 modulo the good prime $z \mapsto 4$ over 83, and the archived exact bank supplies a nonzero kernel vector over $\mathbf{Q}[z]/(f)$. That vector has $a \neq 0$, so it is unique after $a = 1$. Each triple contains a pair belonging to no other triple; its two quadruple roots leave a unique possible third root. Thus the triple coordinates are also fixed. The archived exact bank verifies existence and all guards at the three conjugates.

For completeness, the seven normalized equations in (u, v, z, d, b, c) have a six-by-six Jacobian minor equal to 14 at $(31, 12, 4, 70, 69, 25)$ modulo 83. The normalized tangent space is therefore zero in characteristic zero. The exact resultants, reductions, rank minor, and Jacobian are independently replayed by `orbit2_rigidity_audit.py`; no modular nonlinear inconsistency inference is used. $\square$

Theorem W.2. *Let a seven-cubic bank be given over a characteristic-zero field. Let $\psi = [R : S] : \mathbf{P}^1 \rightarrow \mathbf{P}^1$ be a rational map of degree $e \geq 2$, with coprime binary forms R, S of degree e . Assume the fourteen selected fibers are separable and disjoint, and pull back the word by the section multiplier S^3 . There is no binary form N of degree $3e + 1$ and linear form ℓ whose zero b is outside the selected fibers, with $N(b) \neq 0$, such that N/ℓ has at least $7e$ agreements with this pulled-back word. The pole at infinity is included.*

This excludes the stated proper one-pole augmentation of a seven-cubic bank, not arbitrary eight-word constructions or arbitrary covers of larger-degree banks.

Proof. We may extend constants algebraically and use Proposition W.1. All transformations in that proposition preserve the asserted construction and its properness. We first give the cover-independent norm bound underlying the finite linear certificates.

Pole and norm bound. Suppose f selected fibers are fully matched. Let L interpolate their base word values, with $\deg L \leq f - 1$, and let Q be their monic locator; for $f = 0$ use $L = 0, Q = 1$. Put

$$v = \frac{N/(\ell S^3) - L(\psi)}{Q(\psi)}, \quad m = 3 - f.$$

Full-fiber cancellation gives

$$(v)_\infty \leq m\psi^*(\infty) + [b]. \quad (\text{W1})$$

The pair (ψ, v) is birational onto its image, even when e is composite. Otherwise both functions descend through a finite map ϕ of degree $t > 1$ to the normalization of the image. The pole excess of v over m times the pole order of ψ is exactly one at b , and is nonpositive everywhere else by (W1). Pole orders pull back by ramification indices. Thus every point of the fiber $\phi^{-1}(\phi(b))$ would have positive excess. This fiber must consist only of b , whose ramification index is then t , contradicting that its excess is one. The exact excess at b survives subtraction and division by the base interpolant and locator, because b is outside the selected fibers.

The monic minimal equation of v has degree e . At finite places its symmetric coefficients have no poles except possibly a simple pole at $\psi(b)$: the sum of the additional negative conjugate valuations is at most one. At infinity the j th symmetric coefficient has pole order at most jm , with an additional at most one if $\psi(b) = \infty$. Clearing the possible finite pole therefore gives a nonzero primitive relation

$$H(X, V) = \sum_{j=0}^e A_j(X) V^{e-j}, \quad \deg A_j \leq j(3 - f) + 1, \quad A_0 \neq 0. \quad (\text{W2})$$

This argument includes ramification and cancellation at poles. At a remaining matched fiber the value of v is the fixed scalar $y_x = (w_x - L(x))/Q(x)$. Two distinct matched preimages imply

$$H(x, y_x) = H_X(x, y_x) = H_V(x, y_x) = 0, \quad (\text{W3})$$

since the birational image then has two normalization branches there. Only the value equation is imposed at a single match.

There is also a uniform singularity budget. Compactify the total space of $\mathcal{O}_{\mathbf{P}^1}(m)$ to the Hirzebruch surface $\mathbb{F}_m$. Write C_0 for its negative section and F for a fiber, so $C_0^2 = -m$, $C_0 F = 1$, and $K = -2C_0 - (m + 2)F$. The irreducible closure of the primitive equation (W2) has class

$$D = eC_0 + (me + c)F, \quad c = \max_j (\deg A_j - mj) \in \{0, 1\}.$$

Indeed these coefficient degrees are precisely the transition rule for the fiber coordinate of $\mathcal{O}(m)$; choosing the actual maximum removes any artificially added fiber at infinity. Primitivity removes vertical components, and the nonzero leading coefficient makes $c \geq 0$. Adjunction gives

$$p_a(D) = 1 + \frac{D(D+K)}{2} = (e-1) \left(\frac{me}{2} + c - 1 \right) \leq \frac{me(e-1)}{2}. \quad (\text{W4})$$

Its normalization is $\mathbf{P}^1$. Hence the sum of finite δ -invariants is at most (W4), with no nonsingularity assumption at the boundary. If k_x distinct selected preimages match at x , they give k_x distinct branches at (x, y_x) , so that point contributes at least $\binom{k_x}{2}$.

All degrees at least four. Select exactly $7e$ matches, and let T, Q count those above the triple and quadruple base nodes. Each old cleared section differs from N by a nonzero degree- $(3e+1)$ binary form, since the witness is proper. Summing its root bound over the seven old candidates gives

$$T + Q = 7e, \quad 3T + 4Q \leq 7(3e+1), \quad T \geq 7(e-1).$$

Convexity on the seven triple fibers and (W4), with $m = 3$, imply

$$\frac{7(e-1)(e-2)}{2} \leq \sum_{x \text{ triple}} \binom{k_x}{2} \leq \frac{3e(e-1)}{2}.$$

This is impossible for every integer $e \geq 4$.

Quadratic covers. Move b to source infinity. The proper witness becomes a polynomial of exact degree seven, while the old sections have degree at most six. Four full fibers are impossible: subtract the cleared cubic interpolant through their values to obtain a degree-at-most-seven polynomial with eight roots. Select exactly fourteen matches. Full and empty fibers form disjoint sets F, E of the same size $f \leq 3$. Writing $I_i(x)$ for the base agreement indicator of candidate i , pairwise degree-seven root counting implies

$$\sum_{x \in F} I_i(x) \leq \sum_{x \in E} I_i(x) \quad (1 \leq i \leq 7). \quad (\text{W5})$$

Enumerating these finite subsets gives the counts in the table below. The relation (W2) has coefficient degree caps $(1, 4-f, 7-2f)$ and is imposed at $14-2f$ single fibers.

f	number of patterns				matrix columns			
	0	1	2	3	0	1	2	3
type 7	1	7	42	252	15	12	9	6
type 2	1	1	54	276	15	12	9	6

(W6)

All $f \geq 1$ matrices have full column rank. For $f = 0$ both fourteen-by-fifteen matrices have rank fourteen. For type 7 the received word has the form $(a + cX^7)/X^2$ in its Paley chart; hence the kernel is spanned by $X^2V - a - cX^7$, contradicting $A_0 \neq 0$. For type 2 the unique kernel relation is quadratic, but its discriminant has degree eight and is squarefree. It defines a function field of genus three, contradicting $k(X, v) = k(U)$. These exact rank and squarefreeness certifications are detailed below.

Cubic covers: the fiber budget. Let d and f denote the numbers of double and full matched fibers. Formula (W4), with $e = 3, m = 3$, gives $d + 3f \leq 9$. It forces $f \leq 2$ for a twenty-one-match witness: $f = 3$ would force $d = 0$ and hence at most twenty matches, and $f \geq 4$ is impossible. If $f = 2$, subtract the linear interpolant and divide by the two full fiber locators. Applying (W4) with

$m = 1$ gives $d \leq 3$. This argument includes every ramification partition above infinity, cancellation at the poles, and the extra pole at infinity itself.

Combining these inequalities with the seven pair intersection bounds leaves just two forms of support: every triple fiber double and every quadruple fiber single; or one triple fiber full, the other six double, one quadruple fiber empty, and the other six single. There are seven allowed full/empty pairs for type 7 and one for type 2. This finite assertion can be checked directly from the displayed incidence blocks: their triple and quadruple incidence matrices are invertible; two full fibers would force their nonconstant excess vector to be constant. The remaining one-full alternatives reduce to column differences of the two incidence matrices and fail unless the full triple lies in the empty quadruple. The complete integer census is also supplied with the certificates.

For the no-full pattern, (W2)–(W3) give 28 rows and 26 columns, with coefficient caps $(1, 4, 7, 10)$. For a one-full pattern they give 24 rows and 20 columns, with caps $(1, 3, 5, 7)$. Every one of these matrices has full column rank, contradicting (W2). This completes the proof. $\square$

Exact certificate provenance. The finite linear assertions above use reductions of the specified number-field banks, not a search over covers. For type 7 use the Paley chart with $\zeta_7 \mapsto 16$ in $\mathbf{F}_{29}$; for type 2 use $z \mapsto 4$ in $\mathbf{F}_{83}$ and the independently verified affine bank. Every node difference and interpolation denominator is a unit. The quadratic certificates contain all 302 and 332 matrices and selected minors; independent verifiers rebuild the incidence masks, enumerate (W5), and evaluate these minors by integer Bareiss elimination. The cubic certificates contain eight type-7 matrices and two type-2 matrices; the latter maximal minors are 38 and 1 modulo 83. Separate verifiers reconstruct (W2)–(W3).

In the type-2 quadratic no-full case, normalize the free kernel coordinate to one. Its discriminant reduces to

$$-8X^8 + 35X^7 + 36X^6 + 28X^5 - 13X^4 - 21X^3 - 3X^2 + 24X + 36 \quad \text{in } \mathbf{F}_{83}[X].$$

The supplied Bezout identity with its derivative proves squarefreeness. The unit rank minor makes the normalized characteristic-zero kernel integral at this prime by Cramer’s rule. Thus degree eight and squarefreeness lift, as do all nonzero minors. Scalar extension and Galois conjugation preserve the conclusions.

The files are in `quadratic_one_pole_route` and `cubic_cover_norm_gate`, with independent verification receipts and the source bank arrays. The orbit-2 elimination and Jacobian are in `fano_seven/orbit2_rigidity_audit.json`; its branch formulas are replayed from `nonfano2_projective.json` and the two branch files. All lifting arguments here are nonvanishing-minor or polynomial-identity arguments; none infers characteristic-zero inconsistency from a modular Groebner computation.

W.1 Proper denominators through cover degree nine

The same norm argument gives a finite extension of Theorem W.2. Here the denominator degree is limited by the intended quarter-rate construction; no assertion for arbitrary denominator degree is made.

Corollary W.3. *Let a characteristic-zero seven-cubic bank and a degree- e rational cover $\psi = [R : S]$ satisfy the hypotheses of Theorem W.2. Suppose*

$$2 \leq e \leq 9, \quad 0 \leq h \leq \lfloor e/2 - 1 \rfloor.$$

There is no proper witness N/D , with binary forms of degrees $3e + h$ and h , respectively, having at least $7e$ agreements with the pulled-back word S^3w , where D has no zero on the selected core and $\gcd(N, D) = 1$. For $h = 0$, exclude the seven original pulled-back polynomials from what is called a witness. Zeros of D may coincide with zeros of S and may have multiplicity.

Proof. Extend constants algebraically. By Proposition W.1, it suffices to treat the two archived banks; its transformations preserve these assertions. Put $v = N/(DS^3)$ and let t be the degree of the map to the normalization of the image of (ψ, v) . The positive part of the signed function

$$P \mapsto \text{pole}_P(v) - 3 \text{pole}_P(\psi)$$

is exactly the degree- h zero divisor of D . Pole orders, and hence this positive part, pull back through the intermediate map. Thus $t \mid e$ and, for $h > 0$, $t \mid h$. By Lüroth's theorem the intermediate curve is rational. Separability of the selected fibers means a descended match accounts for exactly t original matches. The descended problem therefore has parameters $(e/t, h/t)$ and at least $7e/t$ matches.

First take $h = 0$. Pairwise root counting against the seven old polynomials shows that all $7e$ matched points must be the full preimages of the seven triple buckets: if T, Q count matched points in triple and quadruple buckets, then $T + Q = 7e$ and $3T + 4Q \leq 21e$. Let W interpolate the seven triple-bucket word values. Each old cubic matches only three of those seven triple buckets. Thus a cubic through all seven values would be an eighth distinct cubic, contradicting the universal maximum-seven theorem. Consequently $4 \leq d = \deg W \leq 6$. The binary form

$$NS^{d-3} - W^{\text{hom}}(R, S)$$

has degree $de < 7e$ and vanishes at those $7e$ distinct points, so it is zero. This would give $v = W(\psi)$, whose pole divisor has degree $de > 3e$, contrary to $(v)_\infty \leq 3\psi^*(\infty)$.

For $h = 1$ use Theorem W.2. The only nontrivial intermediate-field cases still needed below are $(e, h) = (6, 2), (8, 2), (9, 3)$, descending respectively to $(3, 1), (4, 1)$, and $(3, 1)$, all excluded by that theorem. We may consequently assume the remaining maps (ψ, v) are birational.

Their primitive norm equation has the form

$$F(X, V) = \sum_{j=0}^e A_j(X) V^{e-j}, \quad \deg A_j \leq 3j + h.$$

The pole argument above proves these bounds even with repeated poles or poles over infinity. Its irreducible closure on the third Hirzebruch surface has arithmetic genus at most

$$g_{e,h} = (e-1)(3e/2 + h - 1).$$

A fiber containing m distinct matching points contributes at least $\binom{m}{2}$ to its singularity delta. Select $7e$ matches. Pairwise root bounds against the old degree- $(3e + h)$ polynomials give

$$T + Q = 7e, \quad 3T + 4Q \leq 7(3e + h), \quad T \geq 7(e - h).$$

Convexity, using $e \geq 2h + 2$, therefore gives

$$\delta \geq 7 \left\{ \binom{e-h}{2} + \binom{h}{2} \right\}.$$

For $h = 2$, this exceeds $g_{e,2}$ for every $e \geq 7$; at $e = 7$ the comparison is $77 > 69$. For $(e, h) = (9, 3)$ it is $126 > 124$. Only $(6, 2)$ and $(8, 3)$ remain.

For $(6, 2)$, the delta budget 50 forces the unique pattern with four matches in each triple bucket and two in each quadruple bucket. The norm flags give 84 coefficients. Vanishing of all Hasse partials of total order below the fiber multiplicity gives a 91×84 homogeneous matrix. For both banks it has full column rank: independently replayed maximal minors are nonzero modulo 29 and 83, respectively.

For $(8, 3)$ write triple multiplicities as $5 + u_j$, quadruple multiplicities as $3 + v_j$, and $s = T - 35$. The genus budget 98 gives

$$s \in \{0, 1, 2\}, \quad \sum u_j = s, \quad \sum v_j = -s, \quad \sum (u_j^2 + v_j^2) \leq 14 - 4s.$$

The seven old-candidate deficits are nonnegative integers summing to s . For each such deficit vector and each bounded integer vector u , invert the fixed seven-by-seven quadruple-incidence matrix to recover v ; retain precisely the integral vectors satisfying the bounds. This exhaustive rational-arithmetic enumeration gives 386 Paley patterns and 410 orbit-2 patterns. The norm flags have 144 coefficients. Every corresponding Hasse-jet matrix has rank 144 over the same good residue field. All 796 selected maximal minors were independently reconstructed and replayed.

For clarity, these are linear specialization certificates: the base coordinates and word values are integral at the chosen good prime, and each saved nonzero minor is the reduction of the same characteristic-zero minor. No modular nonlinear inconsistency inference is involved. The reproducible census, matrix construction, and independent minor checks are recorded in the companion `sextic_two_pole_norm_gate` and `octic_three_pole_route` directories. The resulting trivial kernels exclude a nonzero primitive norm relation in both remaining cases. $\square$

Clearing D gives degree at most $3e + h$ on $14e$ coordinates, so $h \leq \lfloor e/2 - 1 \rfloor$ is exactly the condition $3e + h + 1 \leq 14e/4$. The corollary excludes this specific augmentation mechanism through degree nine, not arbitrary eight-word constructions at quarter rate.

W.2 The uniform degree-ten Paley net has no rational integral member

We work over an algebraically closed field of characteristic zero. Put $\eta^2 + \eta + 2 = 0$, $\alpha = (\eta - 1)/2$, and $\gamma = 3(\eta + 3)/4$. On $\mathbb{F}_3$, with negative section C_0 and fiber f , consider the net $V \subset H^0(10C_0 + 34f)$ whose members have multiplicity at least four at $Q_i = (\zeta^i, \zeta^{5i})$ and at least six at $T_i = (\alpha\zeta^i, \gamma\zeta^{5i})$, $0 \leq i < 7$. Here ζ is a primitive seventh root. Its dimension is three. An exact basis F_0, F_1, F_2 has characters 1, 3, 5 under $(X, Y) \mapsto (\zeta X, \zeta^5 Y)$, normalized by the coefficients 25, 27, 20 of $Y^{10}, X^2Y^{10}, X^4Y^{10}$, respectively. These specifications and the prescribed jets determine the basis over $\mathbb{Q}(\eta)$.

Theorem W.4. *No member of this net is an integral curve with rational normalization. In particular, the Paley uniform degree-ten, four-pole norm net cannot supply an augmentation through an integral norm equation. This statement concerns this net, not all possible eight-word constructions or all degree-ten incidence patterns.*

Proof. We give the geometric argument and specify the finite certificates on which it depends. Reduction is at the prime above 29 with $\eta = 7$; thus $\alpha = 3$ and $\gamma = 22$. The exact basis and its reduction are checked coefficientwise, using two orbit representatives for the prescribed jets. The 217 prescribed jet conditions on 220 coefficients have rank 217 modulo 29. For each of the three indicated characters, the 31 representative jet equations in its 32 allowed monomials give one exact eigenvector over $\mathbb{Q}(\eta)$. The independently verified vectors have the stated normalizations and reduce to the modular kernel. The unit rank minor therefore proves that they span the entire net and

its integral model. Every nonzero member has actual class $10C_0 + 34f$: the leading coefficient is $25a + 27bX^2 + 20cX^4$, and the three weighted-infinity restrictions below have disjoint supports.

Besides the fourteen prescribed basepoints, the net has exactly two simple basepoints: $(X, Y) = (0, 0)$ and $(t, Z) = (0, 0)$ in the chart $t = 1/X$, $Z = Y/X^3$. One blowup at each of these sixteen points resolves the net. On the resulting smooth surface S , its moving line bundle L satisfies

$$L^2 = 14, \quad K_S L = 14, \quad c_2(S) = 20, \quad p_a(L) = 15.$$

These assertions hold over the DVR at the chosen prime: the closed relative base locus has empty special fiber, and properness then excludes basepoints on the geometric generic fiber.

Here are the boundary checks needed for the discriminant computation. On the affine chart, the nine two-by-two minors of the jet matrix have no common divisorial factor. Along the fourteen exceptional curves, the net restrictions are basepoint-free tangent-cone maps of degree at most six. Their coefficient matrices have rank three; the displayed three-column minors at the two orbit representatives are 5 and 20 modulo 29. The two remaining exceptional curves have degree-one maps. On C_0 the restriction is $[25 : 27X^2 : 20X^4]$, and on the original infinity fiber it is represented by

$$[23Z^8 + 10Z : 15Z^9 + 7Z^2 : 20Z^{10} + 2Z^3].$$

After removal of the common basepoint factor this is nonconstant of degree at most ten. All these degrees are below 29, so the maps are generically separable. Thus the jet matrix has rank at least two at the generic point of every boundary divisor. Its rank-one locus is finite on all of S .

The universal first-jet incidence on $S \times \mathbb{P}^2$ consequently has the expected dimension one: over jet-rank-two points its fiber is a point, and over the finitely many rank-one points its fiber is a line. Jet rank zero is excluded by basepoint-freeness. Its cycle is therefore the top Chern class, and its pushforward is an effective discriminant divisor of degree

$$c_2(J^1 L) = c_2(S) + 2K_S L + 3L^2 = 90.$$

This is a statement about a divisor with multiplicities; it assumes neither generic nodality nor that every component has Gauss degree one.

The seven old cubic graphs yield seven parameter lines. Each contributes multiplicity four. Indeed, the ramification determinant is simple along each old graph; the two normal sections on the first graph are the coprime quartics

$$24 + 14X + 2X^2 + 27X^3 + 12X^4, \quad 25 + 27X + 28X^2 + 9X^3 + 6X^4.$$

They define a separable degree-four map to its parameter line. Equivariance supplies the other six cases. Subtracting these contributions leaves an effective discriminant divisor of degree 62.

Modulo 29, the remaining ramification factor Γ has 135 terms and bidegree $(73, 21)$. It is irreducible, and its smooth rational point $(1, 10)$, with gradient $(27, 14)$, proves geometric irreducibility. There are 4,018 distinct certified Gauss-image points over $\mathbb{F}_{293}$. Each receipt includes a point of Γ , a rank-two jet witness, and its normalized parameter image. Interpolation in all seven character spaces of homogeneous degree 62 has total kernel dimension one: the character-five block has rank 287 of 288, and the other six blocks have full rank. Its equation $H(a, b, c)$ has 271 terms. Since $4,018 > 62^2$, Bézout's theorem implies that every interpolated relation contains the Gauss image. If that image had degree less than 62, all its degree-62 multiples would give a kernel of dimension larger than one. Thus the image has degree 62 and exhausts the residual discriminant cycle, including possible contributions from isolated rank-one points.

For completeness, this identifies the reduction of a characteristic-zero divisor, not just a modular curve. On the smooth relative surface times the parameter plane, the three jet equations and the uniformizer form a regular sequence, since the special incidence has expected codimension. The incidence is therefore flat over the DVR, and proper pushforward of its fundamental cycle commutes with specialization. Subtract the seven horizontal old-line contributions. The primitive equation of the remaining degree-62 divisor reduces to H up to a nonzero scalar.

An integral rational original member must have multiplicity at least 15 in this residual divisor. One minor boundary issue is relevant here. The rank-three tangent matrices prohibit multiplicity jumps at the fourteen prescribed points. At the two additional simple points the linear-part kernel is exactly $[0 : 1 : 0]$. The corresponding member F_1 has order exactly two at each, with quadratic terms $24Y^2$ and $7Z^2$ modulo 29. Thus its resolved divisor is reduced as well: if the original member were integral rational, its components would be its rational strict transform and the two exceptional curves, each once. For any of these reduced connected divisors with k rational components,

$$\sum \delta = p_a(L) + k - 1 \geq 15.$$

A generic pencil through it avoids its singular points with its other section, so its local total spaces are smooth. The discriminant intersection length is the sum of the local Milnor numbers, at least $\sum \delta$. Hence its parameter has discriminant multiplicity at least 15. An integral original member lies on none of the seven old-factor lines, so the same assertion holds for the residual divisor.

Finally explicit polynomial identities exclude every multiplicity-15 point of H over $\mathbb{F}_{29}$. On $a = 1$, a recorded linear combination of monomial multiples of Hasse derivatives of orders below 15 is 1. On $a = 0, c = 1$, a univariate Bézout identity in b gives 1 using restrictions of those derivatives. The remaining point $[0 : 1 : 0]$ has multiplicity six. These are coefficient identities, not pointwise finite-field tests, and have been independently replayed. The multiplicity-15 locus is closed in the proper relative parameter plane. Its empty geometric special fiber therefore implies an empty geometric generic fiber, even when a putative characteristic-zero parameter specializes onto an old-factor line. This contradicts the necessary multiplicity established above. $\square$

Certificate provenance. The reproducibility directory is `research/prime_line_strengthening/degree_ten_four_pole_route/`. The exact kernel and independent coefficient check are `exact_kernel.json` and `verify_exact_kernel.py`. In its `surface/` subdirectory, the base-locus, boundary, and Gauss receipts are `base_locus.json`, `base_locus_boundary.verified.json`, and `gauss_determinant.json`; the boundary supplement is regenerated by `verify_base_boundary.py`. The point and interpolation receipts are `gauss_samples.json` and `gauss_image_interpolation.json`, with independent source replay `verify_gauss_samples.py`. The final explicit identities are `hasse_multiplicity_gate.json`, replayed by `verify_hasse_gate.py` without elimination. The relative-cycle and exceptional-component audit is recorded in `RELATIVE_DISCRIMINANT_AUDIT.md`.

Nonprimitive norm powers. The theorem alone does not exclude a norm equation that is a power of an irreducible rational curve. This possibility is excluded by separate small linear certificates. If a nonzero polynomial G^t belongs to the uniform norm net, where $t \mid 10$ and $t > 1$, then

$$\deg_Y G \leq 10/t, \quad \text{wdeg}_{(1,3)} G \leq \lfloor 34/t \rfloor, \quad \text{mult}_{Q_i} G \geq \lceil 4/t \rceil, \quad \text{mult}_{T_i} G \geq \lceil 6/t \rceil.$$

The corresponding jet matrices modulo 29 have full column rank:

t	matrix size	rank	certified maximal minor
2	63×63	63	27
5	28×12	12	27
10	14×5	5	8

All entries are Hasse evaluations at the same fourteen points. These nonzero minors exclude the corresponding characteristic-zero kernels. Thus the net contains no such power, without assuming that a declared four-pole allowance is attained. For an augmentation whose eliminated norm equation is a scalar multiple of a pure power of its irreducible image equation, these three cases account for every nonbirational intermediate degree. The receipts, including selected original rows, are `surface/nonprimitive_norm_gate.py` and `surface/nonprimitive_norm_gate.json`. For a proper augmentation, separability of the selected fibers and regularity of the added function there ensure that taking the primitive part preserves these contacts. Norm transitivity gives a power of the minimal monic equation; Gauss's lemma gives the same power after primitive polynomial clearing. The weighted norm bound must also hold, as in the stated uniform four-pole construction.

The orbit-2 net. The same method closes the other rigid seven-cubic bank, but its base geometry and interpolation have to be checked separately. Let V_2 be the space of weighted-degree-at-most-34, Y -degree-at-most-10 polynomials having multiplicity at least four at the seven four-agreement points and at least six at the seven three-agreement points of the orbit-2 bank. Use its archived affine model over $\mathbb{Q}(q)$, where $q^3 - 10q^2 + 3q + 1 = 0$.

Theorem W.5. *No member of V_2 is an integral curve with rational normalization. Moreover, V_2 contains no nonzero polynomial G^t with $t \in \{2, 5, 10\}$. Consequently this bank also admits no proper uniform degree-ten four-pole augmentation whose primitive norm equation has the stated caps and contacts.*

Proof. Reduce at $q = 4$ above 83. The fourteen points remain distinct. The 217 prescribed Hasse equations have a 217-column minor equal to 3 modulo 83 among the 220 columns, independently checked by an exact integer determinant. Since there are exactly 217 equations, this proves rank 217 over the exact cubic field as well. Over the chosen DVR, solving those pivot variables with each of the three free variables set to one gives an integral kernel basis reducing to the saved modular basis. Thus no unverified lifting of a modular kernel is needed.

In this case there are exactly fourteen basepoints, not sixteen. After dividing two pairwise resultants by their forced degree-364 factor, their residual gcd is one. Specialized ordinate gcds at every known base coordinate contain only the prescribed point and power. The leading and infinity restrictions have gcd one, including the projective corner. All fourteen tangent-cone triples have rank three and no common projective zero. Hence one blowup at each point resolves the net, giving

$$L_2^2 = 16, \quad K_S L_2 = 12, \quad c_2(S) = 18, \quad p_a(L_2) = 15.$$

The rank-three tangent maps also prohibit every base-multiplicity jump: an integral original member remains integral on the resolved surface. The affine jet-minor gcd is one. The boundary restrictions are nonconstant of degrees at most 4 and 10, and the exceptional restrictions have degrees 4 or 6. All are below 83, so they are generically separable. The rank-one jet locus is finite on the whole surface.

The expected-dimensional jet incidence therefore pushes to a degree-90 discriminant cycle. Every old graph again contributes four: the ramification determinant is simple along it, and its two normal

quartics are coprime of maximum degree four. The residual effective cycle has degree 62. The remaining ramification factor Γ_2 is irreducible modulo 83, of bidegree (76, 21) with 986 terms. Its smooth rational point (1, 75), with gradient (64, 11), proves geometric irreducibility.

No cyclic symmetry is assumed in the interpolation. There are 4,102 distinct certified images over $\mathbb{F}_{83}[r]/(r^3 - r - 3)$. Full homogeneous-degree-62 interpolation has rank 2,015 on 2,016 columns; its unique relation H_2 has 1,990 terms and vanishes at all the certified points. Since $4,102 > 62^2$, the same Bézout and kernel-dimension argument shows that its image has degree 62 and exhausts the residual cycle with multiplicity one. The relative complete-intersection and proper cycle-specialization argument in the preceding proof identifies H_2 as the reduction of the characteristic-zero residual discriminant.

An integral rational member would have total delta 15, hence residual discriminant multiplicity at least 15, by the same generic-pencil Milnor-number argument. It lies off the old-factor lines. Explicit Hasse identities exclude that multiplicity everywhere in the parameter plane: on $a = 1$ a linear combination of degree-bounded derivative multiples is 1; on $a = 0, c = 1$ a univariate Bézout identity is 1; and $[0 : 1 : 0]$ is not on H_2 . The affine certificate comes from a full-rank 2,560-by-2,016 row span and all identities are independently replayed by coefficient summation. Properness of the closed relative multiplicity locus then excludes the characteristic-zero member.

For the nonprimitive cases, the same smaller jet matrices as above have full column rank modulo 83:

t	matrix size	rank	certified maximal minor
2	63×63	63	67
5	28×12	12	49
10	14×5	5	31

The selected original-row determinants were independently recomputed by integer Bareiss elimination. These unit minors exclude all three exact kernels. The primitive norm-power argument and properness hypotheses from the preceding paragraph now give the augmentation conclusion, without assuming that its pole allowance is attained. $\square$

The separate receipts are in `research/prime_line_strengthening/degree_ten_four_pole_route/orbit2_surface/`. The complete proof and scope are in `EXCLUSION.md`; `profile.json` and `old_graph_degrees.json` certify the base geometry and the seven degree-four contributions. `gauss_samples.verified.json` replays every source and jet-kernel witness; `gauss_image_interpolation.json` records the full interpolation. The explicit final identities and their independent replay are `hasse_multiplicity_gate.json` and `hasse_multiplicity_gate.verified.json`. The norm-power determinants are in `nonprimitive_norm_gate.json` and `nonprimitive_norm_gate.verified.json`.

W.3 An explicit rational eight-cubic seed

Theorem W.6. *Over $\mathbb{Q}$, and over every sufficiently large prime field, there is a word on 16 distinct coordinates whose maximum agreement with cubics is 7 and whose complete nearest list consists of exactly eight cubics. Over arbitrarily large prime fields there is also a length-32, dimension-7 Reed–Solomon word whose maximum agreement is 14 and whose complete nearest list consists of exactly eight codewords. This agreement exceeds the audited first-order curve at rate 7/32. The construction has fixed length and list size; no growing-list or asymptotic-tightness assertion follows.*

Proof. We first give the rational source as sections of $\mathcal{O}_{\mathbb{P}^1}(3)$, then specify an ordinary affine

Reed–Solomon chart. Number the coordinates from 0 to 15 and put

$$(z_j)_{j=0}^{15} = \left(\frac{11}{9}, \frac{11}{5}, \frac{77}{65}, 1, \frac{66}{65}, \frac{22}{25}, \frac{6}{5}, \frac{55}{52}, \frac{11}{10}, \frac{77}{95}, \infty, 0, \frac{33}{5}, \frac{33}{25}, \frac{99}{95}, \frac{33}{35} \right).$$

The eight sections $f_i(X) = \sum_{r=0}^3 c_{ir} X^r$ and their exact agreement sets are

i	c_{i0}	c_{i1}	c_{i2}	c_{i3}	S_i
0	693/25	-165/2	883/11	-6175/242	2, 4, 5, 9, 13, 14, 15
1	-99/14	90/7	-81/14	0	0, 3, 7, 10, 12, 14, 15
2	-198/25	69/5	-65/11	0	4, 6, 8, 10, 12, 13, 14
3	0	0	0	0	0, 3, 4, 9, 10, 11, 13
4	0	-1	1	0	2, 3, 5, 8, 10, 11, 14
5	0	-12/5	38/11	-130/121	2, 4, 7, 8, 11, 12, 15
6	0	6	-54/11	0	0, 1, 2, 6, 10, 11, 12
7	0	-6	116/11	-50/11	1, 3, 6, 8, 11, 13, 15

At infinity, a section is evaluated by its X^3 coefficient. Define $v_j = f_i(z_j)$ for any i with $j \in S_i$. Direct rational arithmetic verifies that this definition is independent of the choice of i , and that f_i agrees precisely on S_i .

To obtain sixteen finite nodes, use

$$u_j = \frac{1}{z_j - 3}, \quad q_i(U) = U^3 f_i(3 + U^{-1}).$$

The convention is $u_{10} = 0$. Set $w_j = u_j^3 v_j$ for finite z_j , and $w_{10} = v_{10} = 0$. These expressions define rational polynomials and values, and preserve all incidences. The nodes are distinct and remain distinct modulo 17.

Here is a finite certificate of completeness. In the residue field let

$$R(Y) = 3 + Y + 15Y^2 + 7Y^3, \quad P_i(Y) = 10^{-1} q_i((11 - Y)/3) + R(Y).$$

The nodes $11 - 3u_j$ reduce to $j + 1$, and the corresponding received values reduce, after this same gauge change, to

$$(9, 3, 12, 15, 0, 14, 8, 13, 11, 13, 11, 8, 10, 14, 10, 13).$$

Enumeration of all $\binom{16}{4} = 1820$ four-point cubic interpolants gives 1022 distinct polynomials. The maximum agreement is seven, and exactly the eight P_i attain it. This also proves completeness over the algebraic closure of the residue field: four matching points determine a cubic over $\mathbb{F}_{17}$.

Any characteristic-zero cubic with at least seven agreements on the rational source is determined by four of these rational points and is integral at 17, since their Vandermonde determinant is a unit. Its reduction is a known seven-match cubic. Its matching support must therefore be exactly that known support, and four-point uniqueness identifies it with the corresponding q_i . This proves the complete rational source claim. Since there are only finitely many four-point interpolants and nonzero evaluation differences, the same complete list persists modulo every sufficiently large prime.

For the square pullback it is essential to use the affine cover

$$U = \frac{11 - T^2}{3}, \quad T = \pm \sqrt{11 - 3u_j},$$

rather than $U = T^2$. The sixteen radicands reduce to $1, \dots, 16$ modulo 17. Thus all thirty-two nodes have distinct reductions in $\mathbb{F}_{289}$, and the eight rational sextics $q_i((11 - T^2)/3)$ reduce to

$$10(P_i(T^2) - R(T^2)).$$

After the common polynomial subtraction and nonzero scalar change, this is exactly the certified square pullback of the residue seed.

For clarity, its completeness certificate is as follows. A non-even degree-six polynomial has the form $E(Y) + TO(Y)$ with $Y = T^2$, $\deg E \leq 3$, $\deg O \leq 2$, and $O \neq 0$. At most two fibers can contribute both signs, so fourteen matches require at least twelve distinct hit fibers. At a hit fiber,

$$w^2 + B(Y)w + C(Y) = 0, \quad B = -2E, \quad C = E^2 - YO^2.$$

The eleven coefficients of B, C satisfy twelve linear equations. Among all $\binom{16}{12} = 1820$ choices, 1686 systems have rank eleven and are inconsistent, 132 have rank eleven and are consistent, and two have rank ten and are inconsistent. There are 84 distinct consistent solutions. Every one fails at least one necessary identity $C_0 = E_0^2$ or $C_6 = E_3^2$. Hence no non-even competitor exists, even over the algebraic closure. Even competitors are covered by the cubic census. The complete residue nearest list therefore consists of eight sextics, each with fourteen matches.

For a characteristic-zero competitor, interpolation on seven matching nodes makes its coefficients integral at the chosen unramified quadratic place above 17. Its reduction is one of those eight sextics with exactly fourteen matches. Seven-point uniqueness then identifies the competitor with the corresponding rational sextic. This argument also handles a nonzero odd part whose reduction vanishes. Thus the square pullback has a complete eight-element nearest list and maximum agreement fourteen.

Finally adjoin the finitely many square roots to a number field and exclude the finitely many primes where a coefficient, node difference, or nonzero evaluation difference of a seven-point interpolant degenerates. Arbitrarily large rational primes split completely in its Galois closure, yielding the claimed ordinary prime-field codes. $\square$

The first-order expression at $\rho = 7/32$, $a = 14/32$ is

$$(8 - \rho)a^2 - 6\rho a + \rho(4\rho - 5) = \frac{105}{8192} > 0.$$

The exact rational table and all 56 incidence identities are checked by `research/relaxed_eight_cubic_route/rational_seed.py`. The coefficientwise residue gauge and exact cover are independently recorded by `verify_rational_cover.py`. The complete source and pullback enumerations have independent interpolation and elimination replays in `independent_verify.py` and `verify_quadratic_pullback.py`.

W.4 Simultaneous padding of the rational cubic seed

Theorem W.7. *Over $K = \mathbb{Q}(\sqrt{17})$ there is a word on eighteen distinct nodes whose maximum agreement with a polynomial of degree at most three is seven, and whose complete nearest list has ten elements. The same parameters occur over infinitely many prime fields. A degree-five pullback gives at least ten degree-at-most-fifteen polynomials on ninety nodes, each agreeing in thirty-five positions, over a number field and infinitely many prime fields. Its rate is $16/90 = 8/45$ and its agreement $7/18$ lies strictly above the first-order boundary.*

Here “complete” in the first assertion also excludes extra polynomials over the algebraic closure. No completeness assertion is made for the pullback. This is a finite-list construction, not an unbounded-list family.

i	q_{i0}	q_{i1}	q_{i2}	q_{i3}
0	0	-9	-34	-32
1	0	$-\frac{702}{77}$	$-\frac{13206}{385}$	$-\frac{12384}{385}$
2	0	$\frac{27}{77}$	$\frac{129}{77}$	$\frac{144}{77}$
3	$-\frac{225}{242}$	$-\frac{1453}{242}$	$-\frac{1503}{121}$	$-\frac{1008}{121}$
4	$-\frac{1235}{132}$	$-\frac{3499}{66}$	$-\frac{16312}{165}$	$-\frac{3328}{55}$
5	0	$-\frac{225}{22}$	$-\frac{411}{11}$	$-\frac{9324}{275}$
6	$-\frac{2470}{11}$	$-\frac{14249}{11}$	$-\frac{136821}{55}$	$-\frac{87462}{55}$
7	$-\frac{475}{44}$	$-\frac{2843}{44}$	$-\frac{7032}{55}$	$-\frac{4599}{55}$

Table 2: The eight additional cubics $Q_i(U) = \sum_{j=0}^3 q_{ij}U^j$ with five matches on the rational source.

The rational source and its five-match cubics. Use the sixteen-node rational source and eight cubics $P_0, \dots, P_7$ of Theorem W.6, in its affine coordinate $U = 1/(X - 3)$. The complete set of additional cubics with at least five source matches consists of the following eight Q_i : All eight Q_i have exactly five matches. These finite completeness statements can be checked without solving any nonlinear equations: for each of the $\binom{16}{4} = 1820$ four-node subsets interpolate its unique cubic, deduplicate the coefficient vectors, and count its matches with the rational source table. Exact rational arithmetic gives

agreement	4	5	7
distinct interpolants	1500	8	8

This also certifies that no six-match or higher-than-seven-match cubic exists. Any cubic with at least four matches necessarily occurs in this interpolation census.

Two simultaneous promotions. Direct subtraction gives

$$Q_0 - Q_3 = -\frac{(16U + 9)(358U^2 + 125U - 25)}{242}.$$

Set $g(U) = 358U^2 + 125U - 25$. Its discriminant is $55^2 \cdot 17$, so its two roots

$$u_{\pm} = \frac{-125 \pm 55\sqrt{17}}{716}$$

are distinct, irrational, and outside the rational source domain. At either root, append the common received value

$$Q_0(u_{\pm}) = Q_3(u_{\pm}) = \frac{-104594u_{\pm} - 51075}{32041}.$$

Reduction modulo g verifies that none of $P_0, \dots, P_7$ equals this value at either root. Among $Q_0, \dots, Q_7$, precisely Q_0, Q_3 do. These are exact quadratic-remainder checks with Table 2 and the source coefficients. Since g is irreducible over $\mathbb{Q}$, a rational cubic has the prescribed value at one root if and only if it does at both.

Thus the eight P_i retain seven matches, Q_0, Q_3 increase from five to seven, and the other six Q_i retain five. Any polynomial over $\overline{\mathbb{Q}}$ with at least seven matches on the enlarged domain has at least five source matches. Four of these rational interpolation conditions force its coefficients

to be rational, so the source census applies. This proves both maximum agreement seven and completeness of the ten-element list.

All data and all nonzero remainder and interpolation determinants are algebraic. Outside finitely many primes their reductions remain nonzero; at completely split primes the nodes and values lie in the prime field. The same completeness argument applies. As an explicit check, reduction modulo 103 and enumeration of all $\binom{18}{4} = 3060$ interpolation subsets gives agreement histogram

agreement	4	5	6	7
distinct interpolants	2480	43	1	10

First-order pullback. None of the eighteen nodes equals 3. Pull back by $U = T^5 + 3$, taking all five roots above each node in a finite splitting field. The ninety nodes are distinct, and composition carries each of the ten cubics to degree at most fifteen with exactly thirty-five inherited matches. Good completely split prime specializations preserve these assertions.

For precision about the first-order comparison, in the low-rate branch $\rho < 11 - 3\sqrt{13}$ write $t = \sqrt{\rho/2}$ and let $v > 0$ satisfy $v^2(v + 3) = t$; the boundary is $a_1(\rho) = t(1 + v)$. Put $a = 7/18$ and $K = t^4 + 3at^2 - a^3$. Here $a > t$ and $K > 0$, so

$$a > a_1(\rho) \iff 4t^6 - K^2 > 0.$$

At $\rho = 8/45$ the right side equals

$$\frac{475871}{21257640000} > 0,$$

and the agreement margin is approximately 0.00113940394. Degree five is the smallest pullback degree giving this comparison: degrees one through three fail the high-rate test, while degree four is on the low-rate branch and gives $4t^6 - K^2 = -1557169/34828517376 < 0$.

W.4.1 Three simultaneous additions

The same fresh-five family admits a triangle of fresh intersections. We use Q_0, Q_2, Q_5 from Table 2, in the same affine coordinate as that table.

Theorem W.8. *Over $\mathbb{Q}(\sqrt{39})$ there is a received word on nineteen distinct coordinates whose complete nearest list among polynomials of degree at most three consists of eleven polynomials, each agreeing at exactly seven coordinates. Such examples also exist over arbitrarily large prime fields. A degree-thirteen pullback gives length 247, dimension 40, and agreement 91, with at least eleven polynomials in the threshold list; these parameters lie strictly above the first-order proximity threshold.*

Proof. Retain the sixteen coordinates and received values of the rational eight-word seed. Add the two rational coordinates

$$r = -\frac{80}{163}, \quad s = -\frac{135}{262},$$

with values $Q_0(r) = Q_2(r)$ and $Q_0(s) = Q_5(s)$, respectively. Finally add either root θ of

$$g(U) = U^2 + \frac{4175}{3826}U + \frac{4525}{15304},$$

with value $Q_2(\theta) = Q_5(\theta)$. Its discriminant is $117975/14638276 = 55^2 \cdot 39/3826^2$. The three added coordinates are pairwise distinct and avoid all sixteen old coordinates. The relevant polynomial

differences vanish at the stated coordinates, so each of Q_0, Q_2, Q_5 gains two matches and all eight original candidates retain their seven matches.

Completeness is checked by a finite exact interpolation certificate, not by counting only these eleven candidates. Any cubic agreeing at seven of the nineteen coordinates agrees at at least four old coordinates. Those coordinates and their received values are rational, so the cubic is uniquely determined over $\mathbb{Q}$ by one of the $\binom{16}{4} = 1820$ old four-subsets. The resulting 1516 distinct interpolants are evaluated at all old coordinates, at r, s , and modulo the irreducible quadratic g for the last coordinate. Exactly the eleven displayed candidates have seven matches; none has eight. The exact coefficients, supports, and exhaustive replay are recorded in `research/relaxed_eight_cubic_route/eleven_cubic_certificate.py` and its accompanying JSON certificate.

All distinctness conditions and failed interpolation equalities exclude only finitely many primes. Passing to primes that split completely in a normal closure of the coefficient field therefore gives the claimed arbitrarily large prime-field examples, including completeness of the source list.

For the pullback choose c away from the nineteen coordinates and substitute $U = T^{13} + c$. All nineteen fibers then have thirteen distinct roots in a number-field extension. Pulling back the eleven cubics gives eleven polynomials of degree at most 39, with 91 matches on 247 coordinates. After adjoining the fiber roots and again taking completely split good primes, these are also prime-field examples. Completeness is not asserted for the pulled-back list.

For the threshold comparison put $\rho = 40/247$, $a = 91/247 = 7/19$, $t^2 = \rho/2$, and $K = 3at^2 + t^4 - a^3$. These parameters are in the low-rate branch, and $a > t$, $K > 0$, and

$$4t^6 - K^2 = \frac{4666871}{1343677407241} > 0.$$

Thus a strictly exceeds the first-order threshold. Exact tests of the smaller positive pullback degrees show that thirteen is the first degree with this property for this source. $\square$

Within the eight fresh-five cubics, this padding is optimal for three added coordinates: the exact factor census of all pair differences shows that no fresh coordinate, including infinity, has three of these cubics taking the same section value. Each added coordinate can therefore supply at most two matches to this family, and each promoted cubic requires two. The triangle attains the resulting upper bound of three promotions. This is a statement about this finite family, not an asymptotic list-size lower bound.

W.4.2 Completeness under arbitrary quadratic pullback

Proposition W.9. *Let w be either the eighteen-coordinate ten-cubic word of Theorem W.7 or the nineteen-coordinate eleven-cubic word of Theorem W.8. Pull w back along any connected degree-two map $\mathbb{P}^1 \rightarrow \mathbb{P}^1$ in characteristic zero, assuming that every selected fiber is separable. Among degree-six sections, the complete nearest list is precisely the pullback of the original list: its members have fourteen agreements, and every other section has at most twelve. Thus the lengths are respectively 36 and 38, with list sizes 10 and 11.*

Proof. Work over an algebraic closure. A quadratic cover can be written in binary form as $Z^2 = B_2(X, V)$, with B_2 squarefree of degree two. Every section of the pullback of $\mathcal{O}(3)$ uniquely has the form $E_3 + ZO_2$, where the subscripts denote homogeneous degree. If $O_2 = 0$, source completeness proves the claim. Otherwise at most two selected fibers are full agreement fibers. The monic norm

$$H(X, w) = w^2 + B_3(X)w + C_6(X) = (w - E_3(X))^2 - B_2(X)O_2(X)^2$$

must consequently vanish at either twelve selected base points, or eleven selected base points including two full fibers. At a full fiber it also satisfies $H_w = H_X = 0$. These are linear conditions on the eleven coefficients of B_3, C_6 .

An exact exhaustive census solves every twelve-subset system and every eleven-subset system with the possible full pairs. After the closed necessary-condition screen described below, the remaining exact norms number 48 for the ten-word source and 61 for the eleven-word source. Every one has $E_3^2 - C_6$ equal to a nonzero scalar binary square; no consistent positive-dimensional space remains. This contradicts the required binary squarefree part of degree two for $B_2O_2^2$.

For reproducibility, the census first screens at primes 47 and 61. A unit coefficient minor fixes an integral norm; its necessary condition $E_3^2 - C_6 = B_2O_2^2$ survives specialization because the image of $\mathbb{P}^2 \times \mathbb{P}^2 \rightarrow \mathbb{P}^6$, $(B, O) \mapsto BO^2$, is closed. The screen retains square-branch degenerations and the zero sextic. Crucially, *every* coefficient-rank-deficient base system, even a modularly inconsistent one, is replayed in characteristic zero: there are 588 and 1154 such systems respectively. This covers possible nonintegral solutions and rank increases. Exact replay uses $\mathbb{Q}(\sqrt{17})$ and $\mathbb{Q}(\sqrt{39})$, respectively, retaining all systems until after solving and checking the full-fiber conditions. The scripts and full receipts are the `ten_quadratic_norm_screen/exact` and `eleven_quadratic_norm_screen/exact` files in `research/relaxed_eight_cubic_route/`. $\square$

This proposition gives robustness of these finite lists; it makes no first-order proximity claim for the immediate quadratic-pullback parameters.

W.5 Completeness of the cubic pullback

Proposition W.10. *The explicit rational eight-cubic source admits a degree-three pullback with parameters*

$$n = 48, \quad k = 10, \quad A = 21,$$

whose complete nearest list has exactly eight elements, each with twenty-one agreements. Exactly eight further degree-at-most-nine polynomials have fifteen agreements, and every remaining polynomial has at most fourteen. These assertions hold over a number field and its algebraic closure, and over arbitrarily large prime fields.

Proof. Use the rational affine source (u_j, w_j) from Theorem W.6 and pull back by

$$U = \frac{11 - T^3}{3}.$$

The domain comprises all three roots of $T^3 = 11 - 3u_j$ for every j . Its forty-eight points are distinct. The eight inherited polynomials have degree at most nine and exactly twenty-one agreements.

We first certify a residue-field statement. Modulo 17, the sixteen radicands reduce to $1, \dots, 16$. All forty-eight cube roots lie in $\mathbb{F}_{289}$ and are distinct. After a common polynomial subtraction and a nonzero scalar change, the received word is the full cubic pullback of the original sixteen-node $\mathbb{F}_{17}$ seed.

An exhaustive finite certificate shows that no degree-at-most-nine polynomial not descending through T^3 has sixteen matches on this residue word. Its coverage is short: partition the forty-eight points into two halves of twenty-four. Sixteen matches imply eight anchors in one half. For each eight-anchor subset, every interpolant is

$$q(T) + L(T)(a + bT),$$

where $\deg q \leq 7$ and L is the monic degree-eight anchor locator. Each of the forty remaining nodes gives a line $a + bT = (w - q(T))/L(T)$ in the parameter plane. Eight concurrent lines are equivalent to sixteen total matches. Every such concurrence yields a descended polynomial. The original census uses 248980 symmetry-reduced planes. An independent replay changes to alternating base-fiber halves, uses Newton interpolation, and quotients only by deck rotations, not Frobenius; it checks 490314 planes from 1470942 anchor subsets. It detects concurrence independently by taking each of the first 33 residual lines as reference and counting seven equal intersection slopes to the others: any set of eight among forty contains such a reference. Again there is no non-descended hit. Any algebraic-closure candidate with ten matches already has $\mathbb{F}_{289}$ coefficients by interpolation, so the residue conclusion holds over the algebraic closure.

The same parameter-plane enumeration at threshold fifteen has exactly nine non-descended candidates after closing the three reported representatives under deck transformations and Frobenius. An independent enumeration with the alternate halves and deck-only quotient has the same orbit closure. All nine have exactly fifteen residue matches. None lifts on the fixed rational cover: lift each node uniquely modulo 17^2 by $T^3 = 11 - 3u_j$, apply the residue gauge $C(T) \mapsto 10(C(T) - R(T^3))$ with $R(Y) = 3 + Y + 15Y^2 + 7Y^3$, and interpolate ten of the fifteen matching values. Every candidate has a nonzero residual at a remaining support point modulo 17^2 . Independent replay uses the last ten support nodes and Newton interpolation, instead of the first ten and Gaussian elimination. The residue nodes are unit-separated, so interpolation determines the same unique polynomial in every extension of the local field. Thus this fixed-cover obstruction also excludes ramified lifts.

Now let Q be a characteristic-zero degree-nine polynomial with at least fifteen matches. Ten matching points make its coefficients integral by a unit Vandermonde determinant. The preceding obstruction rules out a non-descended reduction. Its descended residue cubic has at least five and at most seven source matches. If it has seven, it is one of the known eight, and all fifteen actual matching indices of Q belong to that candidate's twenty-one-point residue support. The known lift agrees at all twenty-one corresponding points, hence agrees with Q at fifteen distinct points. The two degree-nine polynomials are equal.

Otherwise all actual matches lie over at most six base fibers. At least three are fully matched, since non-full fibers contribute at most two points. Write

$$Q(T) = E(T^3) + TO(T^3) + T^2V(T^3), \quad \deg E \leq 3, \quad \deg O, \deg V \leq 2.$$

Full agreement at all three distinct roots of a fiber forces O and V to vanish at its base value. Three full fibers imply $O = V = 0$. Therefore Q descends to a cubic with at least five agreements on the exact rational source. The exhaustive rational four-point census gives exactly sixteen such cubics: the eight original seven-match cubics and eight fresh five-match cubics. Their pullbacks have respectively twenty-one and fifteen matches. This proves both complete tiers and excludes every other polynomial at threshold fifteen.

For prime-field descent, take a finite Galois field containing the forty-eight coordinates. Exclude the finitely many primes at which node differences, coefficients, or nonzero evaluation differences of any ten-point interpolant degenerate. At completely split remaining primes every candidate with fifteen matches is still one of the finitely many ten-point interpolants, with its characteristic-zero agreement count unchanged. Arbitrarily large such primes give the asserted complete agreement profiles. $\square$

The independent finite census and its coverage proof are stored in `research/relaxed_eight_cubic_route/cubic_extension_search/`, including `independent16.cpp`, `independent15.cpp`, `lift15.independent.json`, and `INDEPENDENT_AUDIT.md`. This is a fixed-cover, fixed-list result, not a general cubic-cover exclusion or an unbounded amplification theorem.

At $\rho = 10/48 = 5/24$ and $a = 21/48 = 7/16$, the audited first-order expression is strictly positive:

$$(8 - \rho)a^2 - 6\rho a + \rho(4\rho - 5) = \frac{1409}{18432} > 0.$$

X A complete nine-word quartic list by incidence exchange

Theorem X.1. *There are arbitrarily large prime fields admitting a length-18 Reed–Solomon code of dimension 5 and a received word whose maximum agreement is 8 and whose complete nearest list has exactly nine elements. There are also arbitrarily large prime fields admitting length 216, dimension 49, and a received word with at least nine distinct codewords having agreement 96. Completeness is asserted only for the length-18 construction.*

Proof. Start with the eight-cubic seed over $\mathbb{F}_{17}$ from the preceding construction. Its nodes are $1, \dots, 16$, its received word is

$$(9, 3, 12, 15, 0, 14, 8, 13, 11, 13, 11, 8, 10, 14, 10, 13),$$

and the ascending cubic coefficient vectors are

$$\begin{array}{cccc} 1 & 10 & 7 & 2 \\ 16 & 11 & 12 & 4 \\ 13 & 4 & 7 & 5 \\ 3 & 1 & 15 & 7 \\ 0 & 10 & 2 & 10 \\ 15 & 12 & 5 & 12 \\ 14 & 15 & 0 & 14 \\ 16 & 1 & 13 & 15. \end{array}$$

Write these cubics as $P_0, \dots, P_7$. Each has exactly seven agreements with the displayed word. Set

$$N(X) = 2 + 15X + 8X^2 + 5X^3 + 13X^4, \quad b = 0.$$

The numerator equations $N(x) = (x - b)w(x)$ hold at $x \in \{1, 2, 5, 6, 7, 8, 9, 12\}$. Consider the 64 equations consisting of these eight equations and the 56 selected cubic agreements. Their variables are the 32 cubic coefficients, 16 nodes, 16 received values, b , and five coefficients of N .

Delete the equation $P_0(3) = w(3)$ and the equation $N(1) = (1 - b)w(1)$. The remaining 62 equations have Jacobian rank 62. An independently reconstructed original 62×62 minor has determinant 8 modulo 17. Fix the eight complementary variables to their integral representatives. Multivariate Hensel lifting produces a solution over $\mathbb{Z}_{17}$; the invertible square Jacobian shows that its coordinates are algebraic over $\mathbb{Q}$, hence belong to a number field.

The identity

$$N(X) - XP_0(X) = 11(X - 5)(X - 6)(X^2 + 10)$$

holds in the residue field. Let $\theta^2 = 7$ in $\mathbb{F}_{289}$. This is a fresh root, and the derivative of the displayed difference at θ is $6 + 15\theta \neq 0$. Thus the equation $N(\theta) = (\theta - b)P_0(\theta)$ lifts uniquely near this root over the unramified quadratic extension. Append the node b , with received value zero, and this lifted node θ , with received value $N(\theta)$. On the old nodes replace $w(x)$ by $(x - b)w(x)$. The nine quartics are

$$(X - b)P_0(X), \dots, (X - b)P_7(X), N(X).$$

All have eight selected agreements: the common node b supplies one to each old cubic, while θ replaces each of the two deleted incidences. The pole is proper since $N(b)$ reduces to 2, and all domain separations and polynomial distinctions are units.

We verify that the lifted list is complete and has maximum agreement exactly eight. Over $\mathbb{F}_{289}$, enumeration of all $\binom{18}{5} = 8568$ five-point interpolants produces 5628 distinct quartics. Exactly the nine displayed quartics have at least eight agreements; their actual agreement counts are

$$(9, 8, 8, 8, 8, 8, 8, 8, 9).$$

This census is valid over the algebraic closure as well, since any quartic with five matches is already determined by five residue-field points. It has been independently replayed using Newton divided differences in place of the original Lagrange interpolation.

The two extra agreements do not survive lifting. Evaluate the 64 original equations on their integer representatives, divide by 17, and reduce modulo 17. Two independently checked left-null vectors of the full Jacobian give an invertible system on the two deleted coordinates. Since all retained equations vanish, that system forces the deleted equation residuals divided by 17 to be respectively 13 and 2 modulo 17. Both are nonzero. Therefore each lifted listed quartic has exactly eight agreements.

Any characteristic-zero quartic with at least eight agreements is integral at this place, by interpolation at five unit-separated matching nodes. Its reduction is one of the nine listed residue quartics, with support of size at most nine. Its support and the corresponding listed lift's eight selected matches therefore overlap in at least $8 + 8 - 9 = 7$ nodes. Two quartics agreeing at seven distinct nodes are equal. This proves completeness and excludes agreement greater than eight.

Finally, pass to a finite Galois number field containing all coordinates. Outside finitely many primes, all domain and coefficient guards remain nonzero; completeness also persists, by retaining the finitely many nonzero evaluation differences of all five-point interpolants. Arbitrarily large rational primes split completely in this field, proving the prime-field assertion.

For the second assertion choose a constant c different from all eighteen source nodes and pull back by $X = U^{12} + c$. After adjoining the roots, the domain consists of 216 distinct points. The nine polynomials have degree at most 48 and each agrees on $12 \cdot 8 = 96$ points. Split-prime specialization again gives arbitrarily large prime fields. This step asserts at least nine codewords, not completeness of the pulled-back list. $\square$

At $\rho = 49/216$ and $a = 96/216 = 4/9$, the first-order expression used above is strictly positive:

$$(8 - \rho)a^2 - 6\rho a + \rho(4\rho - 5) = \frac{73}{34992} > 0.$$

These are fixed finite list constructions, not a growing-list amplification theorem. Reproducible independent certificates are stored in `research/relaxed_eight_cubic_route/verify_nine_exchange.py`, `verify_nine_complete.py`, and their JSON receipts; the latter verifies both the exhaustive decoder and the two left-null defect identities.

References

- [1] P. Erdős and A. Rényi. Probabilistic methods in group theory. *Journal d'Analyse Mathématique* 14 (1965), 127–138. https://renyi.hu/~p_erdos/1965-15.pdf.
- [2] J. Bourgain, M. Z. Garaev, S. V. Konyagin, and I. E. Shparlinski. *On the hidden shifted power problem*. arXiv:1110.0812v2, 2011. <https://arxiv.org/abs/1110.0812>.

- [3] T. Xylouris. *On Linnik's constant*. Acta Arithmetica 150 (2011), 65–91. <https://arxiv.org/abs/0906.2749>.
- [4] A. Gasull, J. Torregrosa, and X. Zhang. *The number of polynomial solutions of polynomial Riccati equations*. arXiv:1602.03503, 2016. <https://arxiv.org/abs/1602.03503>.
- [5] P. Corvaja and U. Zannier. *Greatest common divisors of $u - 1, v - 1$ in positive characteristic and rational points on curves over finite fields*. J. Eur. Math. Soc. 15 (2013), 1927–1942. <https://doi.org/10.4171/JEMS/409>.
- [6] A. Beauville. *Finite Subgroups of $\mathrm{PGL}_2(K)$* . Contemporary Mathematics 522 (2010), 23–29. [https://math.univ-cotedazur.fr/~beauvill/pubs/PGL\(2\).pdf](https://math.univ-cotedazur.fr/~beauvill/pubs/PGL(2).pdf).
- [7] J. Thaler. *Proofs, Arguments, and Zero-Knowledge*. Foundations and Trends in Privacy and Security, 4(2–4), 2022. <https://people.cs.georgetown.edu/jthaler/ProofsArgsAndZK.html>.
- [8] E. Ben-Sasson, D. Carmon, Y. Ishai, S. Kopparty, and S. Saraf. *Proximity Gaps for Reed–Solomon Codes*. ECCC TR20-083, revision 3, 2021. <https://eccc.weizmann.ac.il/report/2020/083/>.
- [9] E. Ben-Sasson, L. Goldberg, S. Kopparty, and S. Saraf. *DEEP-FRI: Sampling Outside the Box Improves Soundness*. ITCS 2020, LIPIcs 151, 5:1–5:32. <https://doi.org/10.4230/LIPIcs.ITCS.2020.5>.
- [10] O. Alrabiah, Z. Guo, V. Guruswami, R. Li, and Z. Zhang. *Random Reed–Solomon Codes Achieve List-Decoding Capacity With Linear-Sized Alphabets*. Advances in Combinatorics, 2025:8. <https://doi.org/10.19086/aic.2025.8>.
- [11] H. Pasten and J. T.-Y. Wang. *Extensions of Büchi's Higher Powers Problem to Positive Characteristic*. International Mathematics Research Notices 2015(11) (2015), 3263–3297. <https://doi.org/10.1093/imrn/rnu033>.
- [12] J. Brakensiek, Y. Chen, A. Putterman, Z. Zhang, and K. Z. Zheng. *Algorithmic List Decoding of Reed–Solomon Codes up to Capacity*. arXiv:2609.08005, 2026. <https://arxiv.org/abs/2609.08005>.
- [13] J. Brakensiek, S. Gopi, and V. Makam. *Generic Reed–Solomon Codes Achieve List-decoding Capacity*. arXiv:2206.05256. <https://arxiv.org/abs/2206.05256>.
- [14] R. M. Roth. *Higher-Order MDS Codes*. arXiv:2111.03210v2, 2021. <https://arxiv.org/abs/2111.03210v2>.
- [15] P. Borwein, T. Erdélyi, and G. Kós. *Littlewood-type problems on $[0, 1]$* . Theorem 2.7 and its proof. <https://people.tamu.edu/~terdelyi/papers-online/PLMS.pdf>.
- [16] V. Gandikota, B. Ghazi, and E. Grigorescu. *NP-Hardness of Reed-Solomon Decoding, and the Prouhet–Tarry–Escott Problem*. 2016. <https://arxiv.org/abs/1611.03069>.
- [17] A. Rudra. *List Decoding and Property Testing of Error Correcting Codes*. Ph.D. thesis, University of Washington, 2007, Section 6.4.3. <https://cse.buffalo.edu/faculty/atri/papers/coding/thesis-chaps/chap6.pdf>.

- [18] S. Kopparty. *List-Decoding Multiplicity Codes*. Theory of Computing 11(5):149–182, 2015. <https://theoryofcomputing.org/articles/v011a005/v011a005.pdf>.
- [19] F. G. Jeronimo. *Algorithmic List Decoding at Capacity and Optimal Proximity Gaps for Reed–Solomon Codes*. ECCC TR26-169, September 2026. <https://eccc.weizmann.ac.il/report/2026/169/>.
- [20] J. S. Milne. *Algebraic Number Theory*. Theorem 8.31 and Corollary 8.32. <https://www.jmilne.org/math/CourseNotes/ANTc.pdf>.
- [21] D. Krachun, S. Kazanin, and U. Haböck. *Failure of proximity gaps close to capacity*. IACR ePrint 2026/782. <https://eprint.iacr.org/2026/782>.
- [22] N. M. Katz. *An estimate for character sums*. Journal of the American Mathematical Society 2(2), 197–200, 1989. <https://web.math.princeton.edu/~nmk/old/estcharsums.pdf>.
- [23] E. Crites and A. Stewart. *On Reed–Solomon Proximity Gaps Conjectures*. IACR ePrint 2025/2046. <https://eprint.iacr.org/2025/2046>.
- [24] J. Thorner and A. Zaman. *Refinements to the prime number theorem for arithmetic progressions*. arXiv:2108.10878, 2021. <https://arxiv.org/abs/2108.10878>.
- [25] E. Ben-Sasson, D. Carmon, U. Haböck, S. Kopparty, and S. Saraf. *On Proximity Gaps for Reed–Solomon Codes*. IACR ePrint 2025/2055. <https://eprint.iacr.org/2025/2055>.
- [26] R. Goyal, V. Guruswami, Y. Sun, and M. Wootters. *Locality of Curve-Decoding and Improved Proximity Gaps*. arXiv:2607.08516, 2026. <https://arxiv.org/abs/2607.08516>.
- [27] S. Jo. *Reed–Solomon Mutual Correlated Agreement Beyond the Johnson Radius*. IACR ePrint 2026/1432, revised August 19, 2026. <https://eprint.iacr.org/2026/1432>.
- [28] U. Haböck, D. Levit, and S. Papini. *Circle STARKs*. IACR ePrint 2024/278. <https://eprint.iacr.org/2024/278>.
- [29] Proximity Prize. *OrbitPencil: validated upper-certificate construction*, repository revision cdb451f13fdc6c84f5fe363e77ee13a89bd30974. [Formal source](#).
- [30] Carmon et al. *S-two Whitepaper*. IACR ePrint 2026/532, March 2026 revision. <https://eprint.iacr.org/2026/532>.
- [31] StarkWare. *proving*, source snapshot cd7bc5f4697fb188a27e09f9242f1dd76df8afdc. <https://github.com/starkware-libs/proving/tree/cd7bc5f4697fb188a27e09f9242f1dd76df8afdc>.
- [32] StarkWare. Public production parameters and default conjectured soundness. [Production Cairo configuration](#); [nominal security counter](#); [documented 96-bit conjectured default](#).
- [33] StarkWare. Implemented challenge and grinding order, source snapshot cd7bc5f4697fb188a27e09f9242f1dd76df8afdc. [Polynomial-commitment verifier](#); [FRI verifier](#).
- [34] Y. Lin and R. Wong. *Global Asymptotics of the Discrete Chebyshev Polynomials*. arXiv:1207.2536, 2012. Equations (1.1)–(1.2). <https://arxiv.org/abs/1207.2536>.

- [35] X. Li and P. Ding. *General Forms of Finite Population Central Limit Theorems with Applications to Causal Inference*. arXiv:1610.04821, 2016. Theorem 1. <https://arxiv.org/abs/1610.04821>.
- [36] J. Hájek. *Limiting Distributions in Simple Random Sampling from a Finite Population*. Publications of the Mathematical Institute of the Hungarian Academy of Sciences, 5:361–374, 1960. <https://real.mtak.hu/200979/>.
- [37] S. Bobkov and M. Madiman. *The Entropy per Coordinate of a Random Vector Is Highly Constrained under Convexity Conditions*. IEEE Transactions on Information Theory, 57(8):4940–4954, 2011. Proposition III.1 and proof of Corollary III.2. <https://arxiv.org/abs/1006.2883>.
- [38] J. Guo, C.-L. Sun, and J. T.-Y. Wang. *On Pisot’s d -th root conjecture for function fields and related GCD estimates*. arXiv:2010.01896. <https://arxiv.org/abs/2010.01896>.
- [39] A. Cafure and G. Matera. *Improved explicit estimates on the number of solutions of equations over a finite field*. Finite Fields and Their Applications, 12(2):155–185, 2006. <https://doi.org/10.1016/j.ffa.2005.03.003>.
- [40] András Sárközy and Gábor N. Sárközy. On the size of partial block designs with large blocks. *Discrete Mathematics*, 305:264–275, 2005. [doi:10.1016/j.disc.2005.09.014](https://doi.org/10.1016/j.disc.2005.09.014).
- [41] Seoyoung Kim, Chi Hoi Yip, and Semin Yoo. Explicit constructions of Diophantine tuples over finite fields. *The Ramanujan Journal*, 65:163–172, 2024. [doi:10.1007/s11139-024-00888-5](https://doi.org/10.1007/s11139-024-00888-5).
- [42] Q. Dao, S. D. Kominers, J. Thaler, and K. Z. Zheng. *Counterexamples to Beyond-Johnson Proximity Gaps over Binary Fields*. Working manuscript, September 2026. https://github.com/GUJustin/binary_field_counterexamples.
- [43] Q. Dao, S. D. Kominers, and J. Thaler. *Reed–Solomon Codes Beyond Johnson: Efficient Decoding and Smaller Cryptographic Proofs*. Cryptology ePrint Archive, Report 2026/2056, September 2026. The 145-page version retrieved September 17, 2026. <https://eprint.iacr.org/2026/2056>.
- [44] A. J. Harper. Lecture notes 2 for Cambridge Part III course on Elementary Methods in Analytic Number Theory, 2015. Theorem 7.1. <https://warwick.ac.uk/fac/sci/math/people/staff/harper/elementarymethodsnotes2.pdf>.
- [45] B. E. Diamond and A. Gruen. *On the Distribution of the Distances of Random Words*. Cryptology ePrint Archive, Report 2025/2010; full version of CRYPTO 2026. <https://eprint.iacr.org/2025/2010>.
- [46] M. Montanucci and C. Zanella. *A class of linear sets in $\text{PG}(1, q^5)$* . Finite Fields and Their Applications 78 (2022), 101983. <https://doi.org/10.1016/j.ffa.2021.101983>.
- [47] S. Lu. *Completing the classification of maximum scattered linear sets in $\text{PG}(1, q^5)$* . Preprint, September 1, 2026. Only the explicitly checked incidence and tangent arguments are used here. https://shunqilu.github.io/publications/classification_maximum_scattered_linear_sets_PG1q5.pdf.